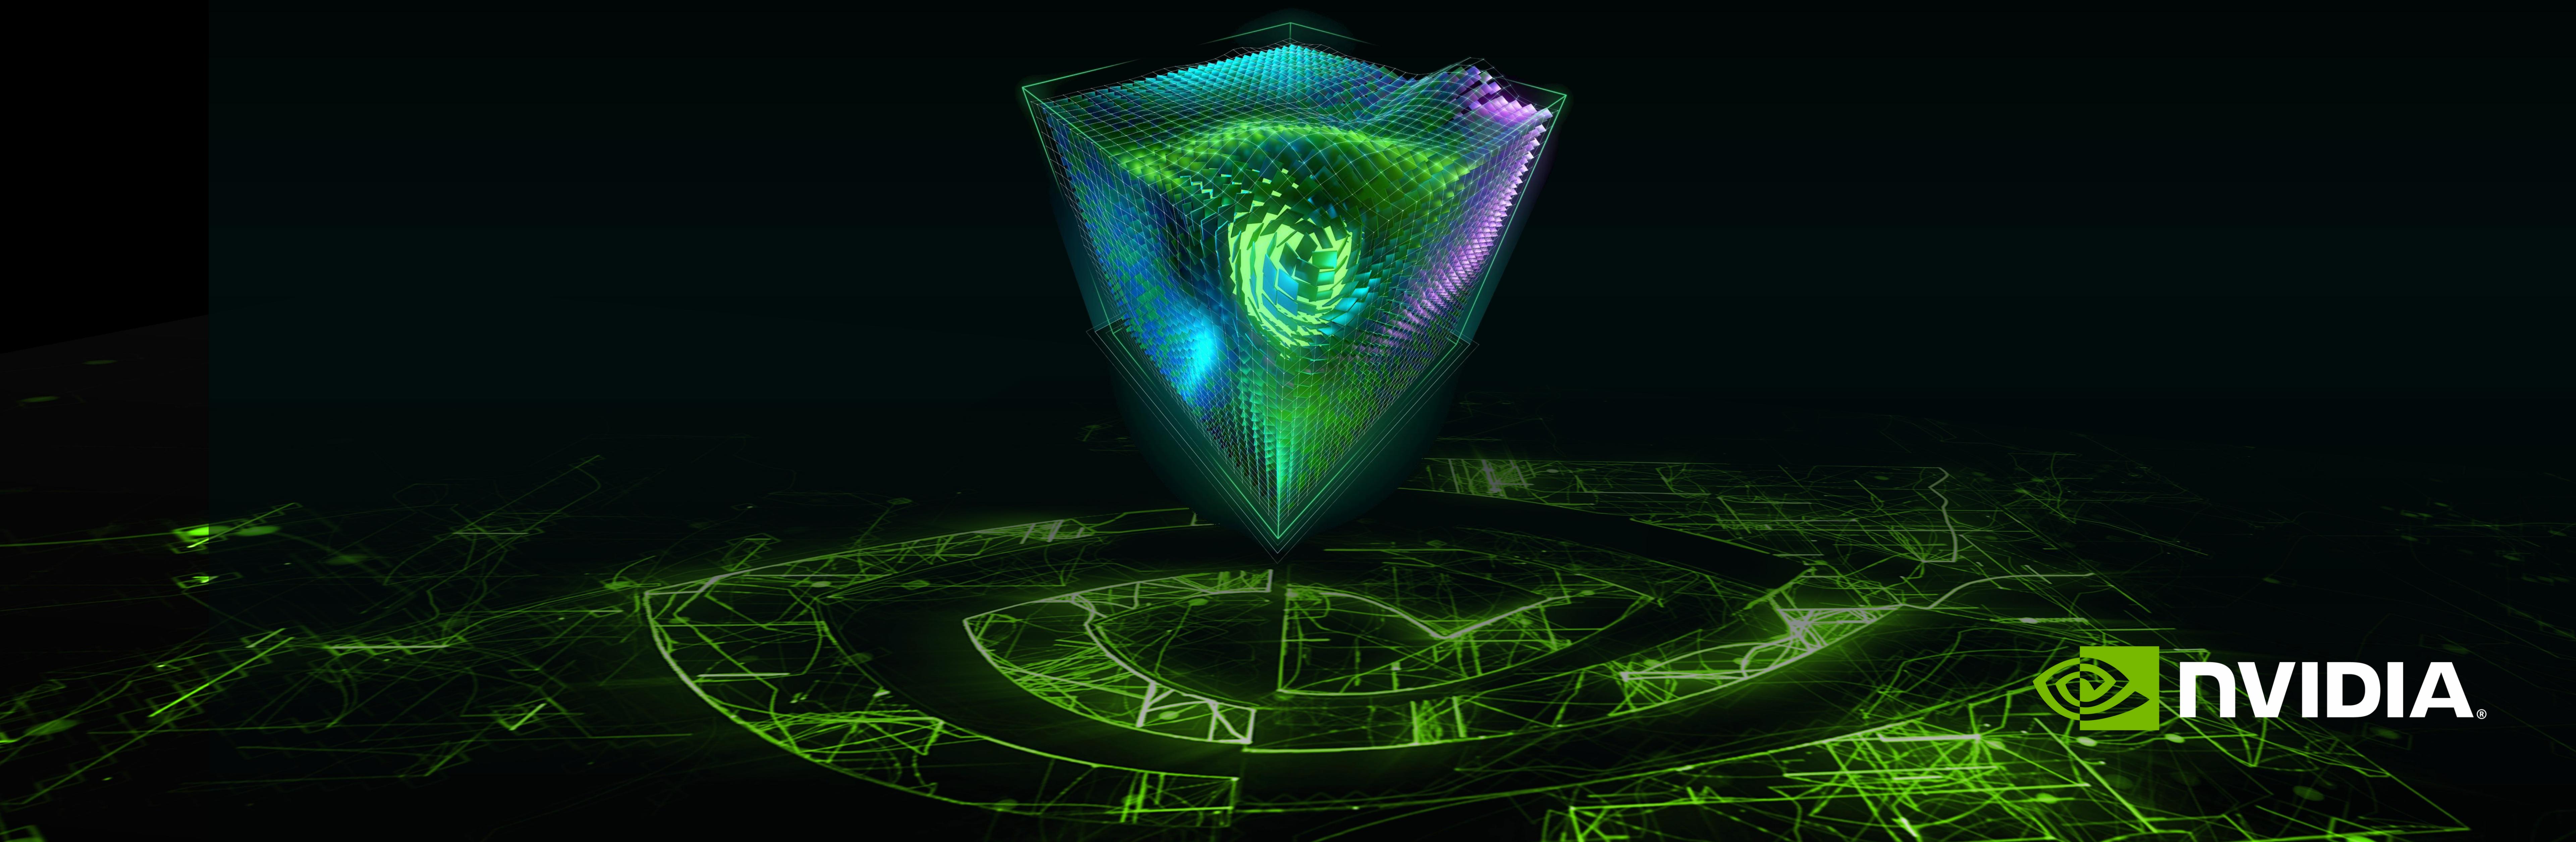


CUDA: NEW FEATURES AND BEYOND

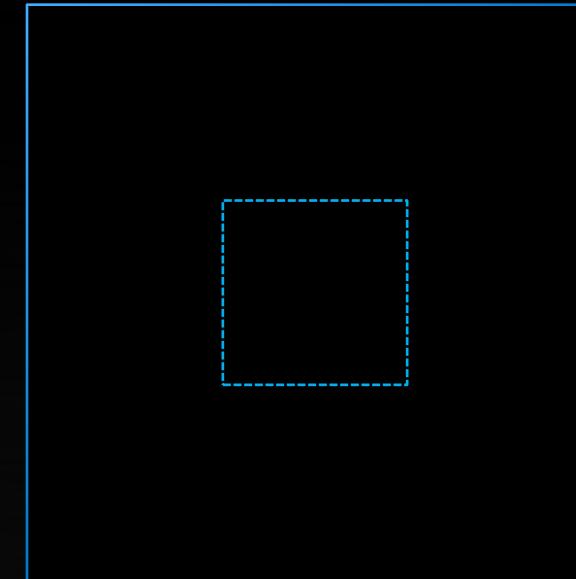
Stephen Jones, GTC 2022



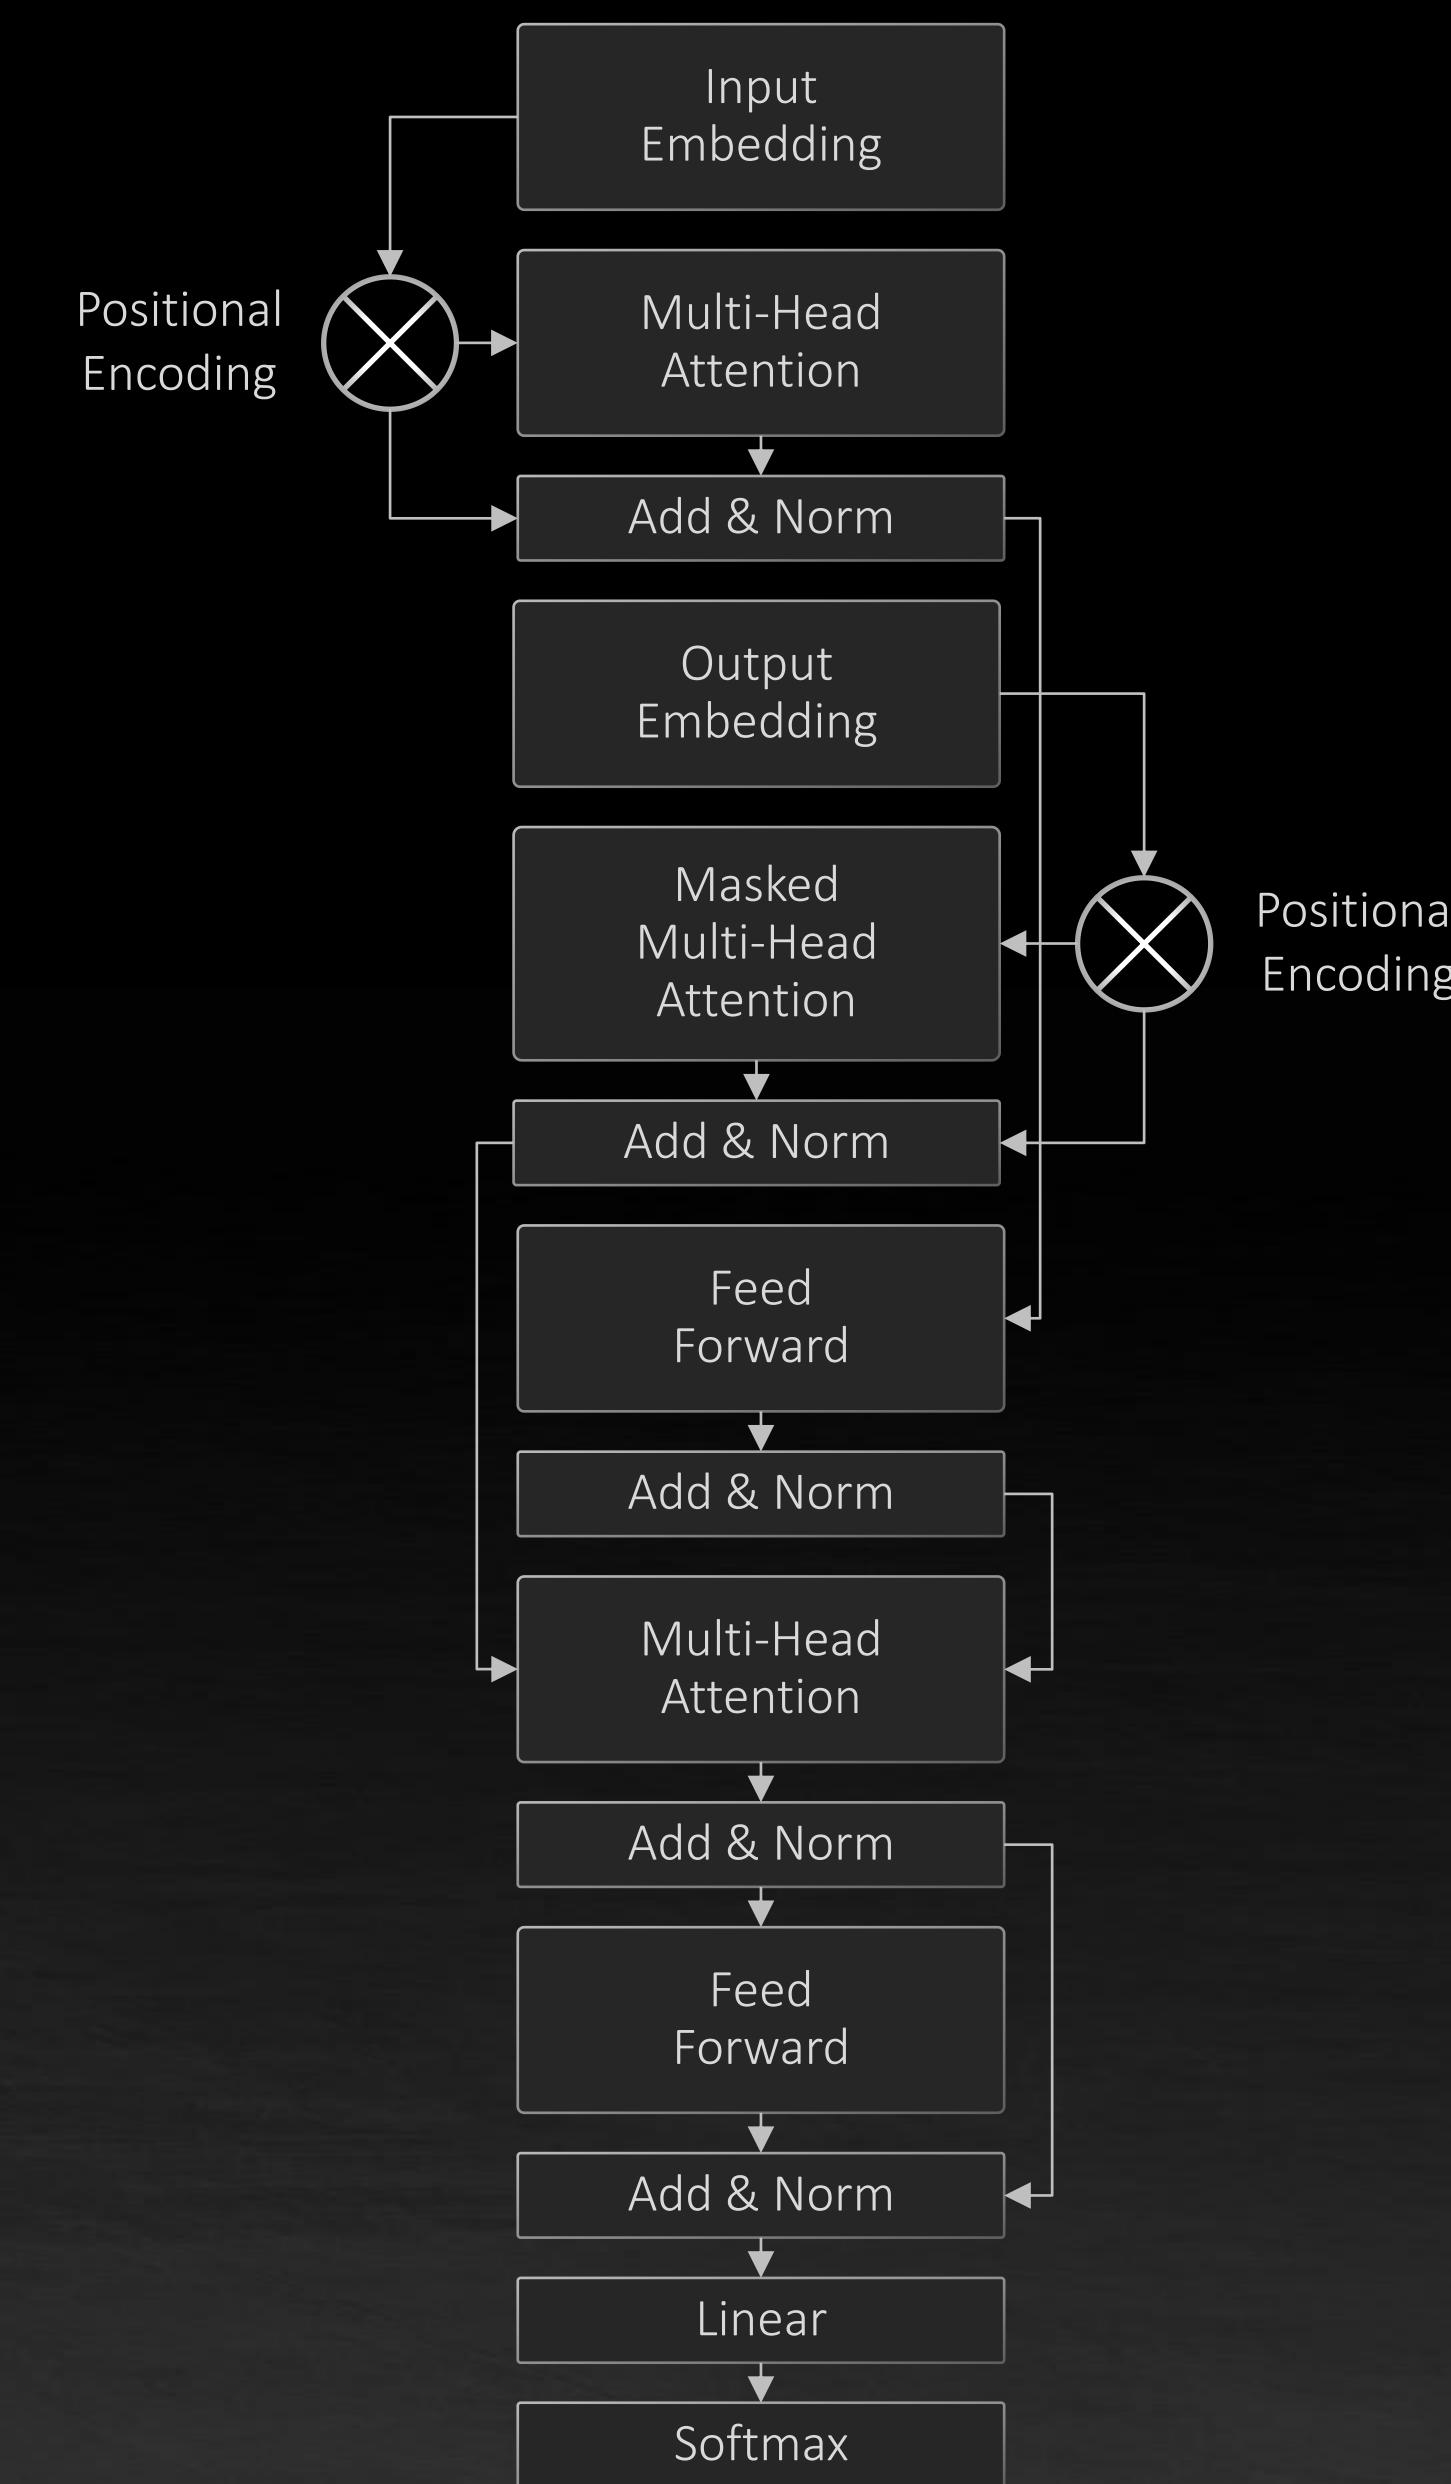
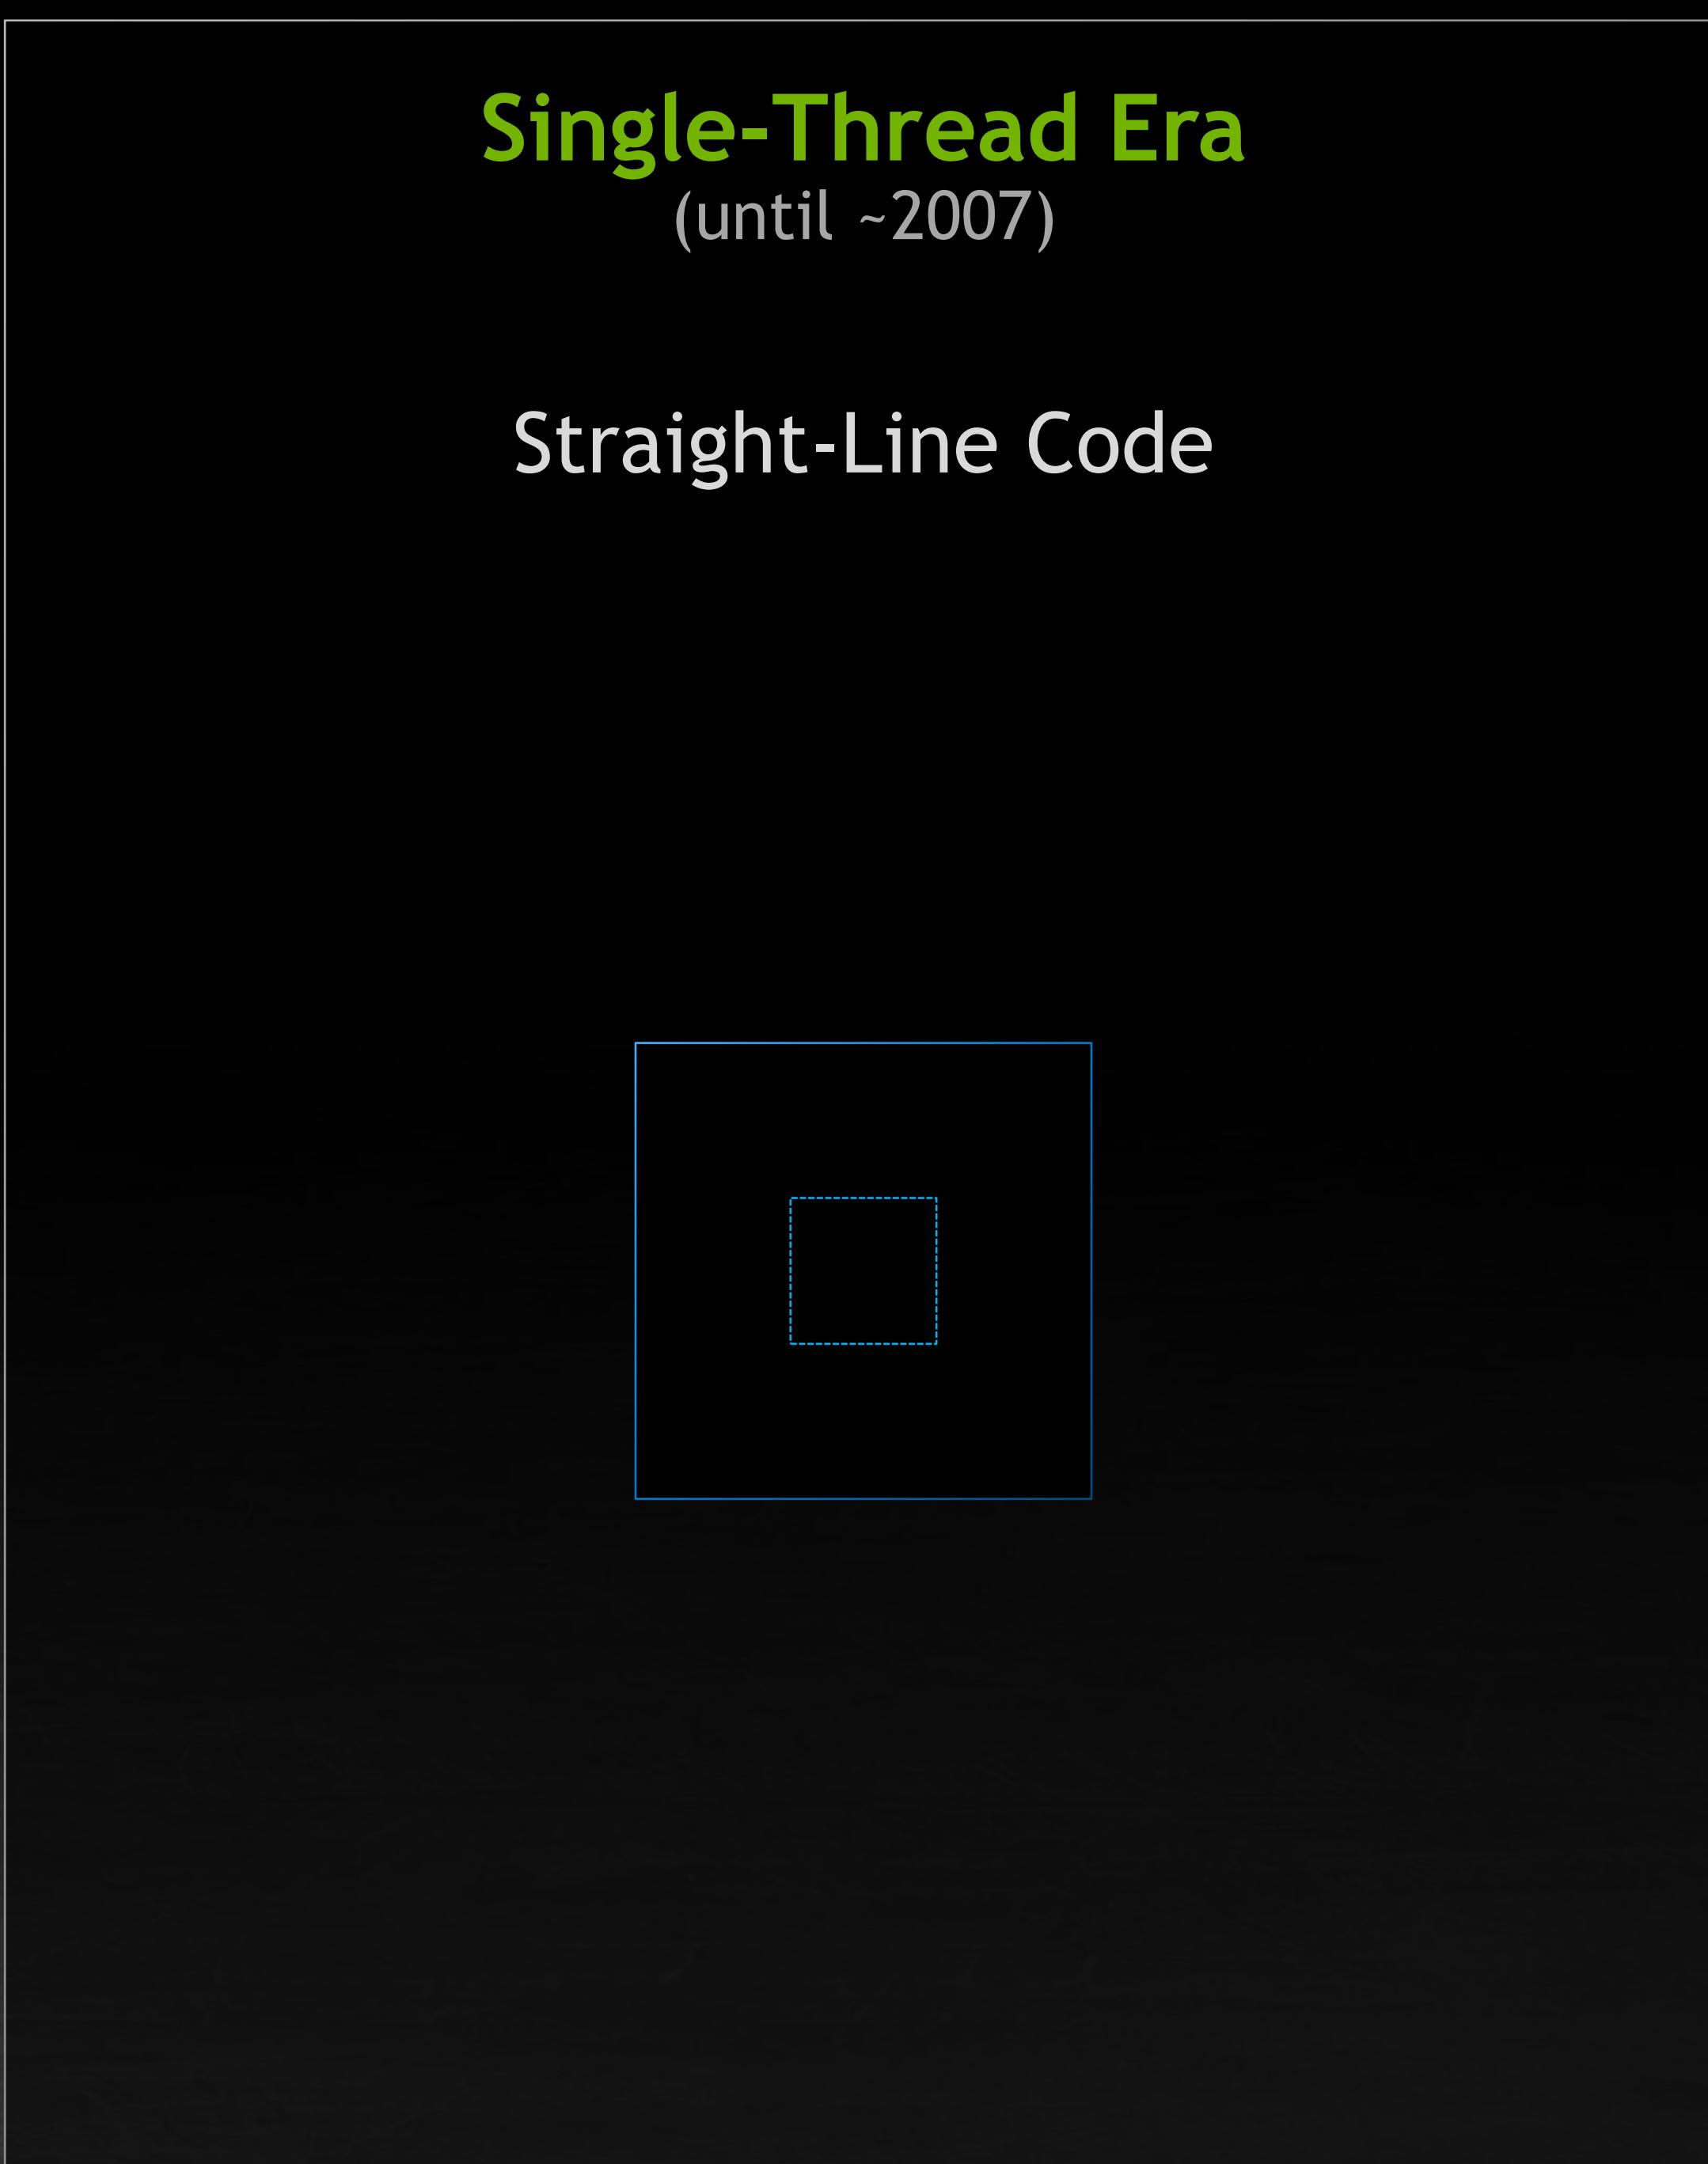
THE FIRST ERA OF SOFTWARE DEVELOPMENT

Single-Thread Era
(until ~2007)

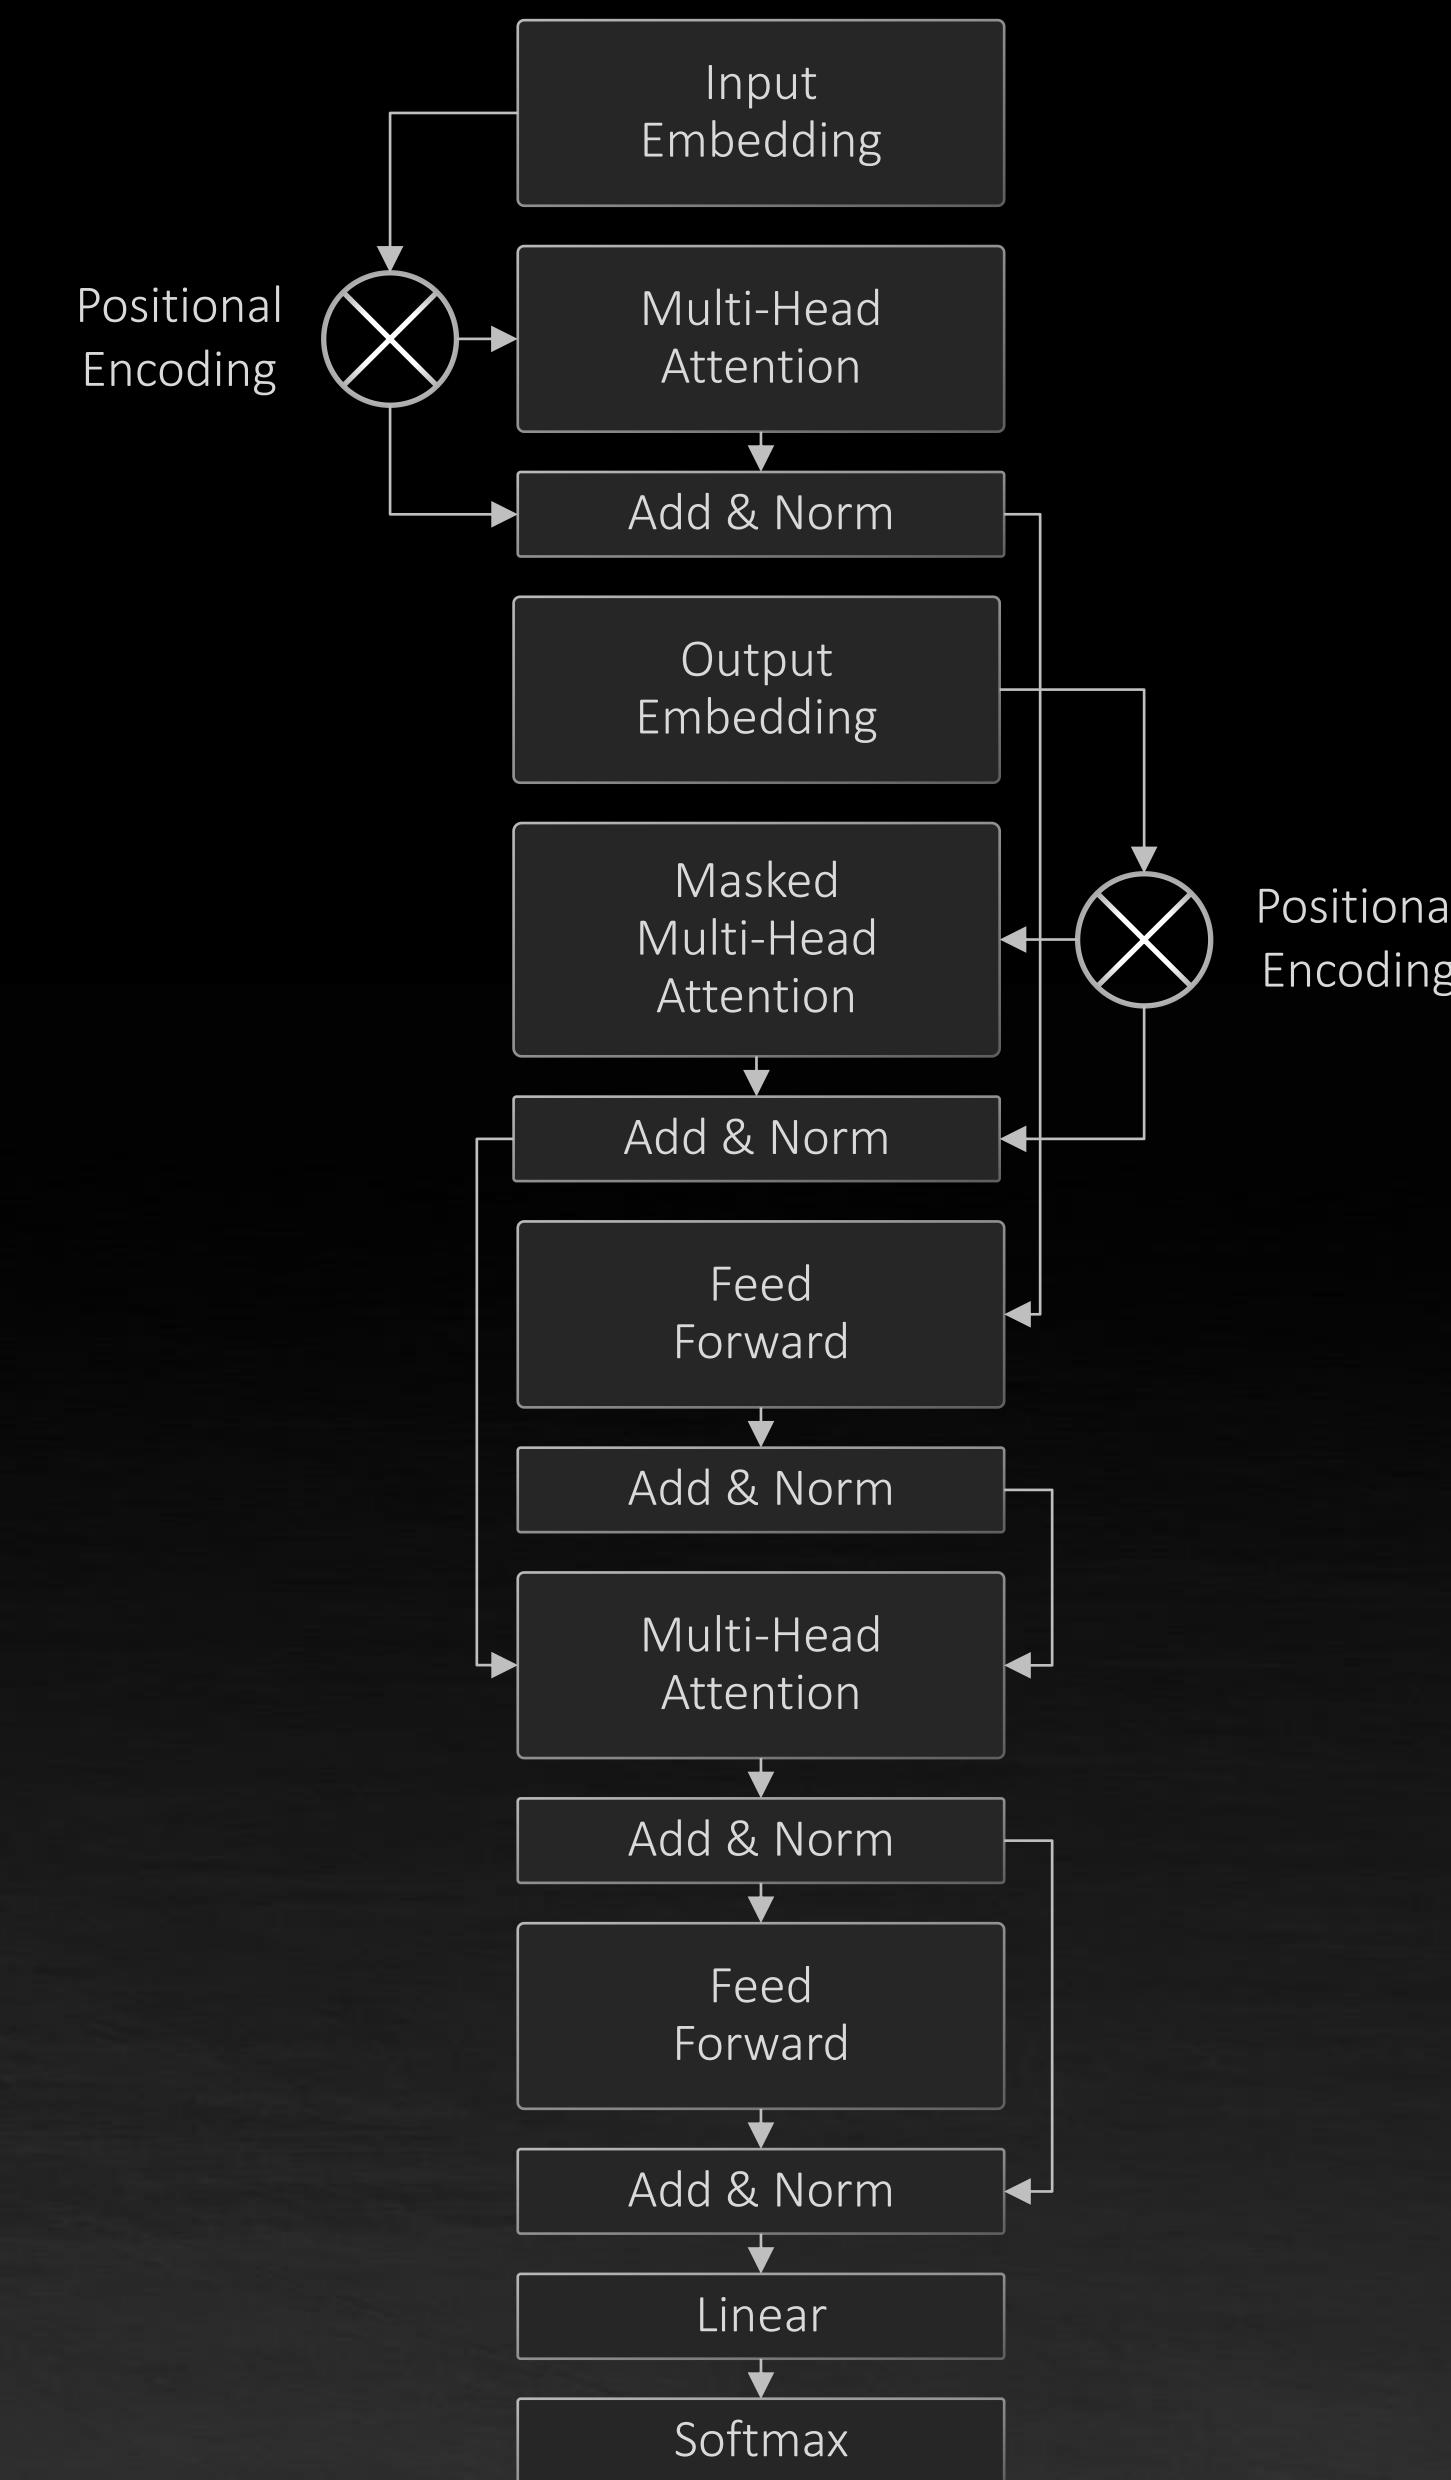
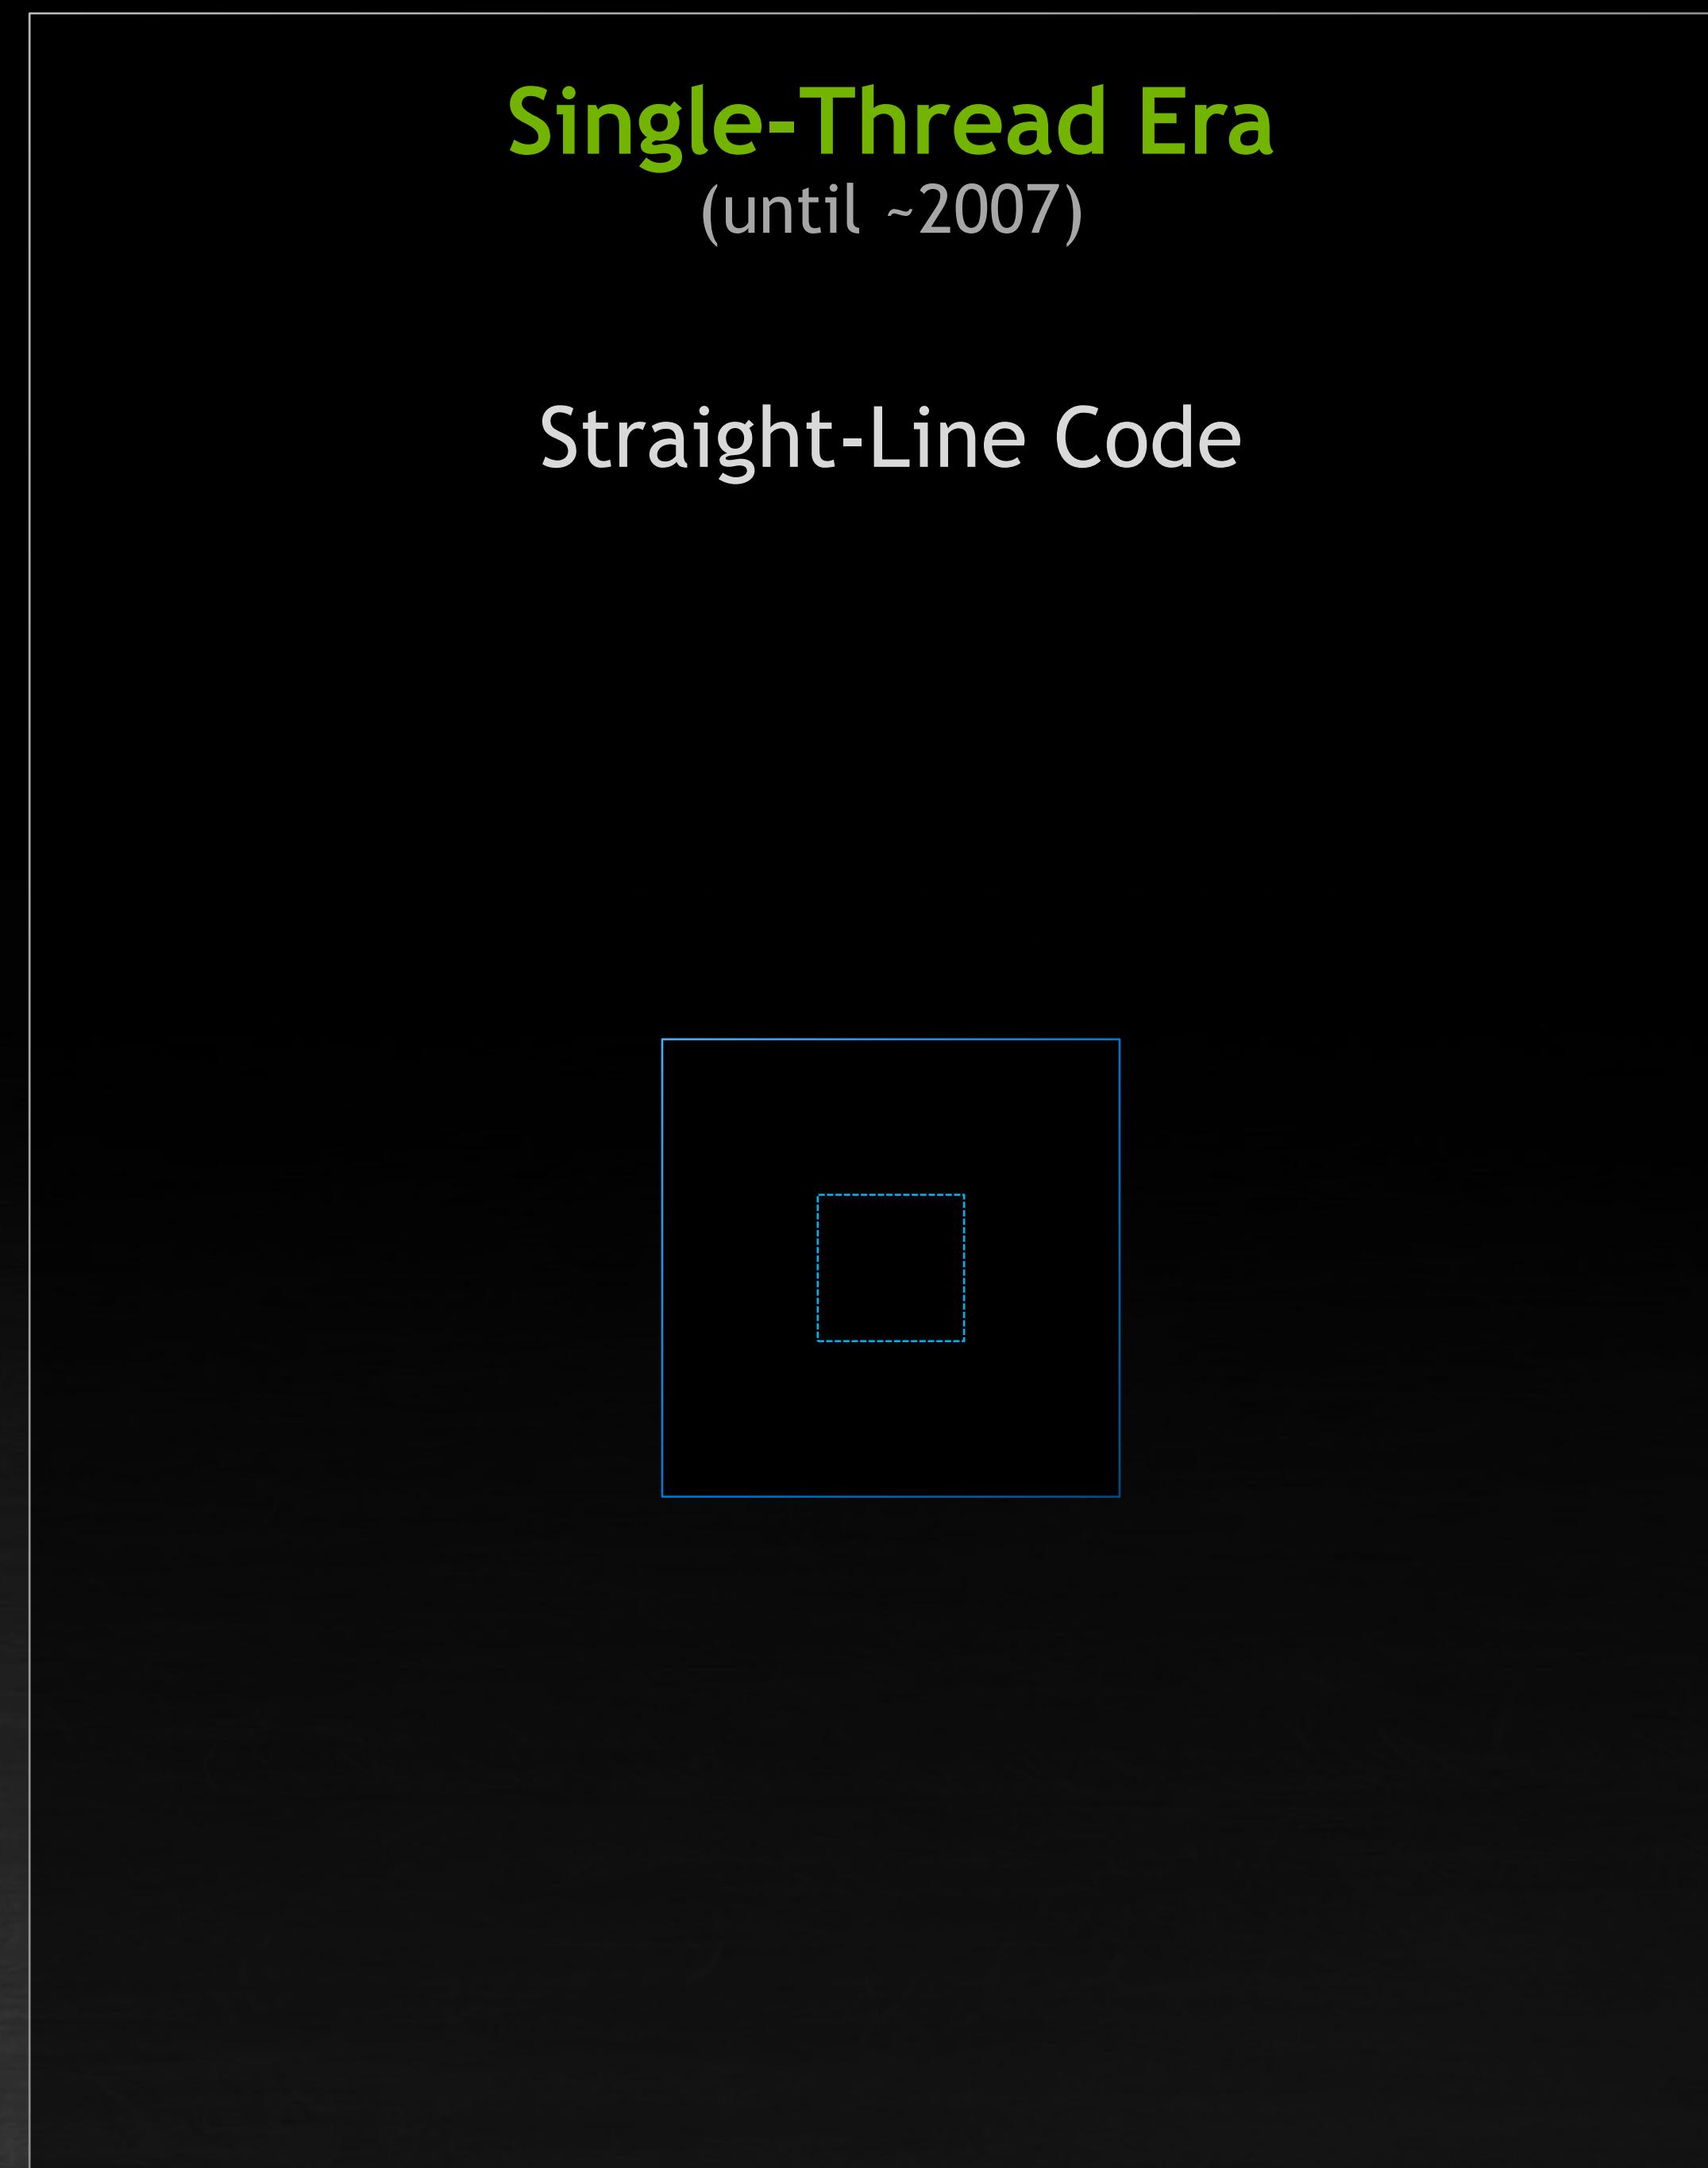
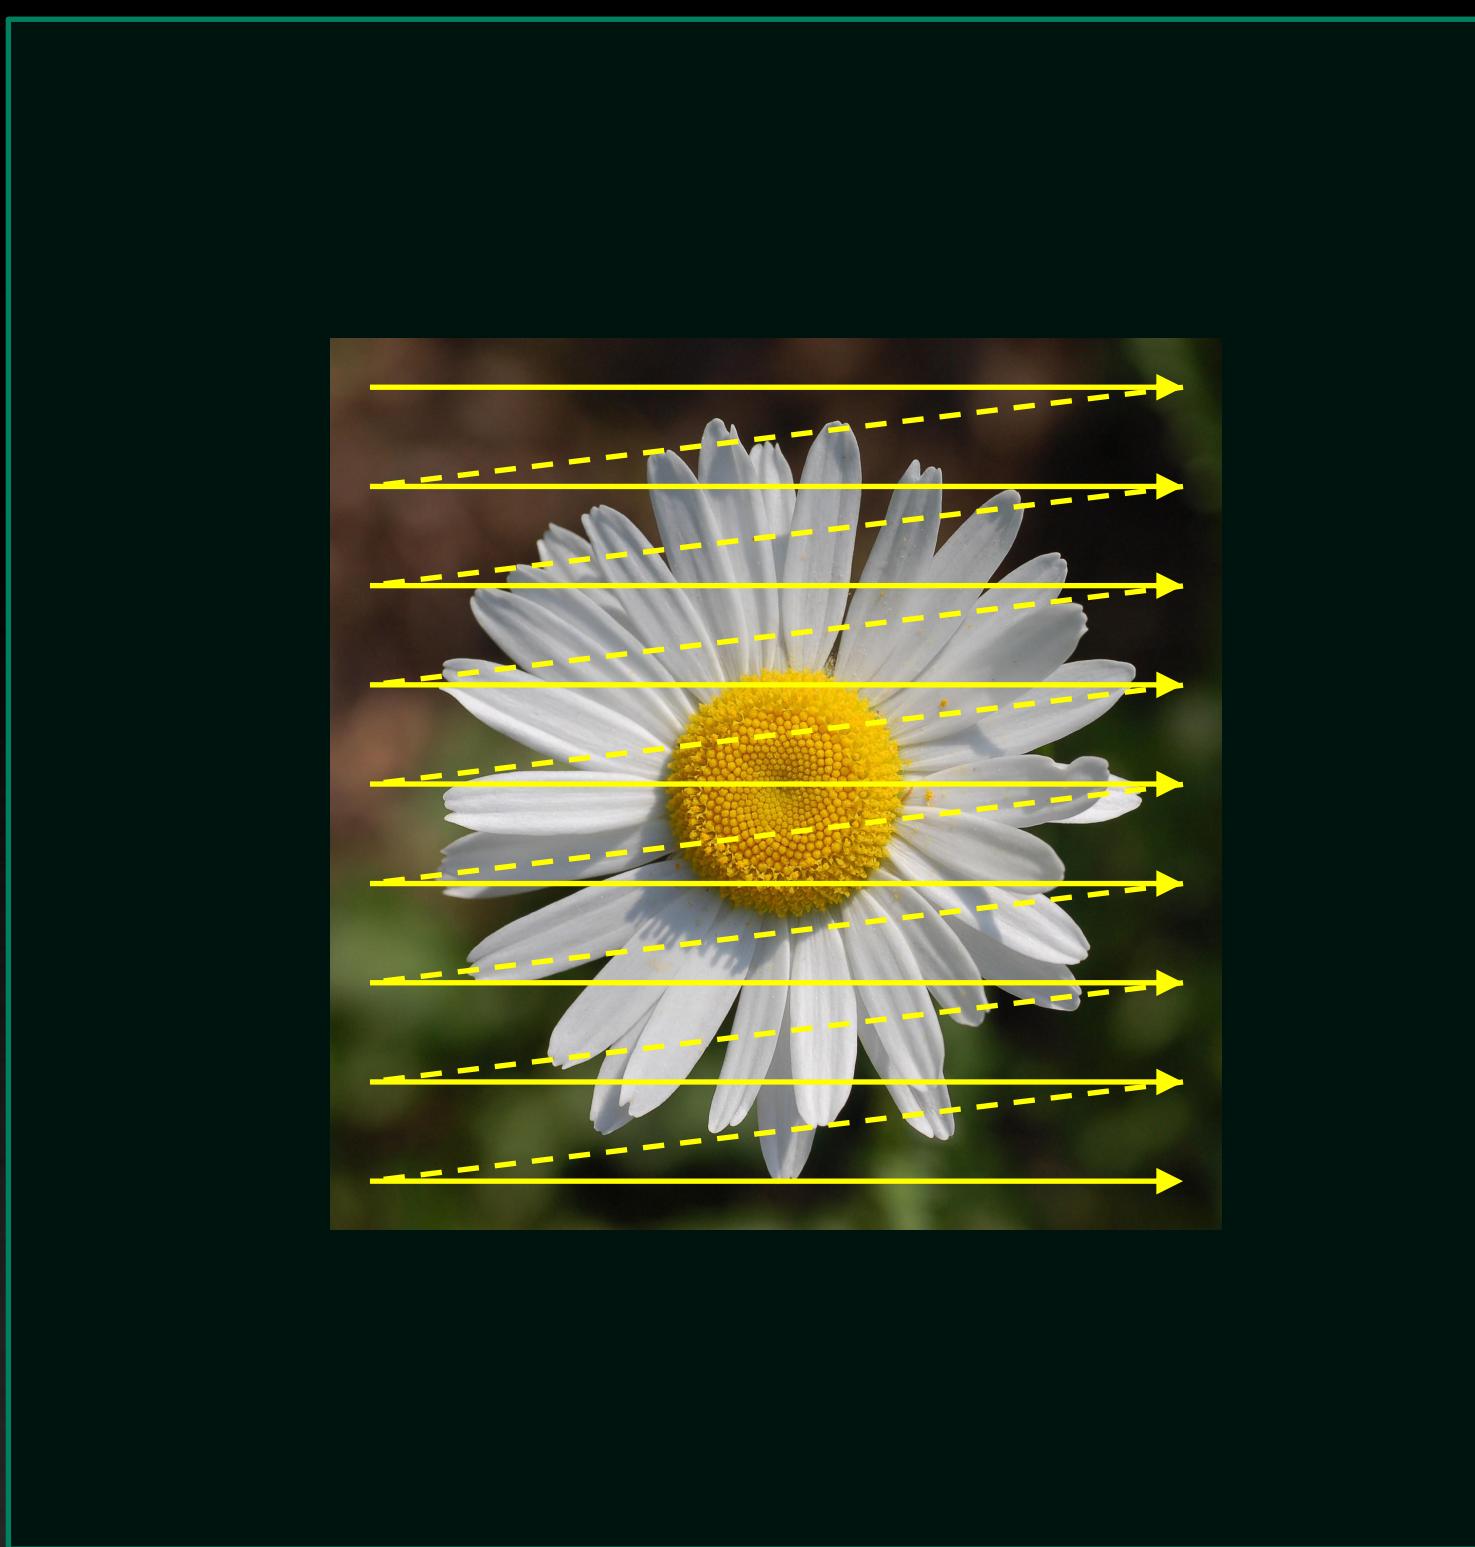
Straight-Line Code



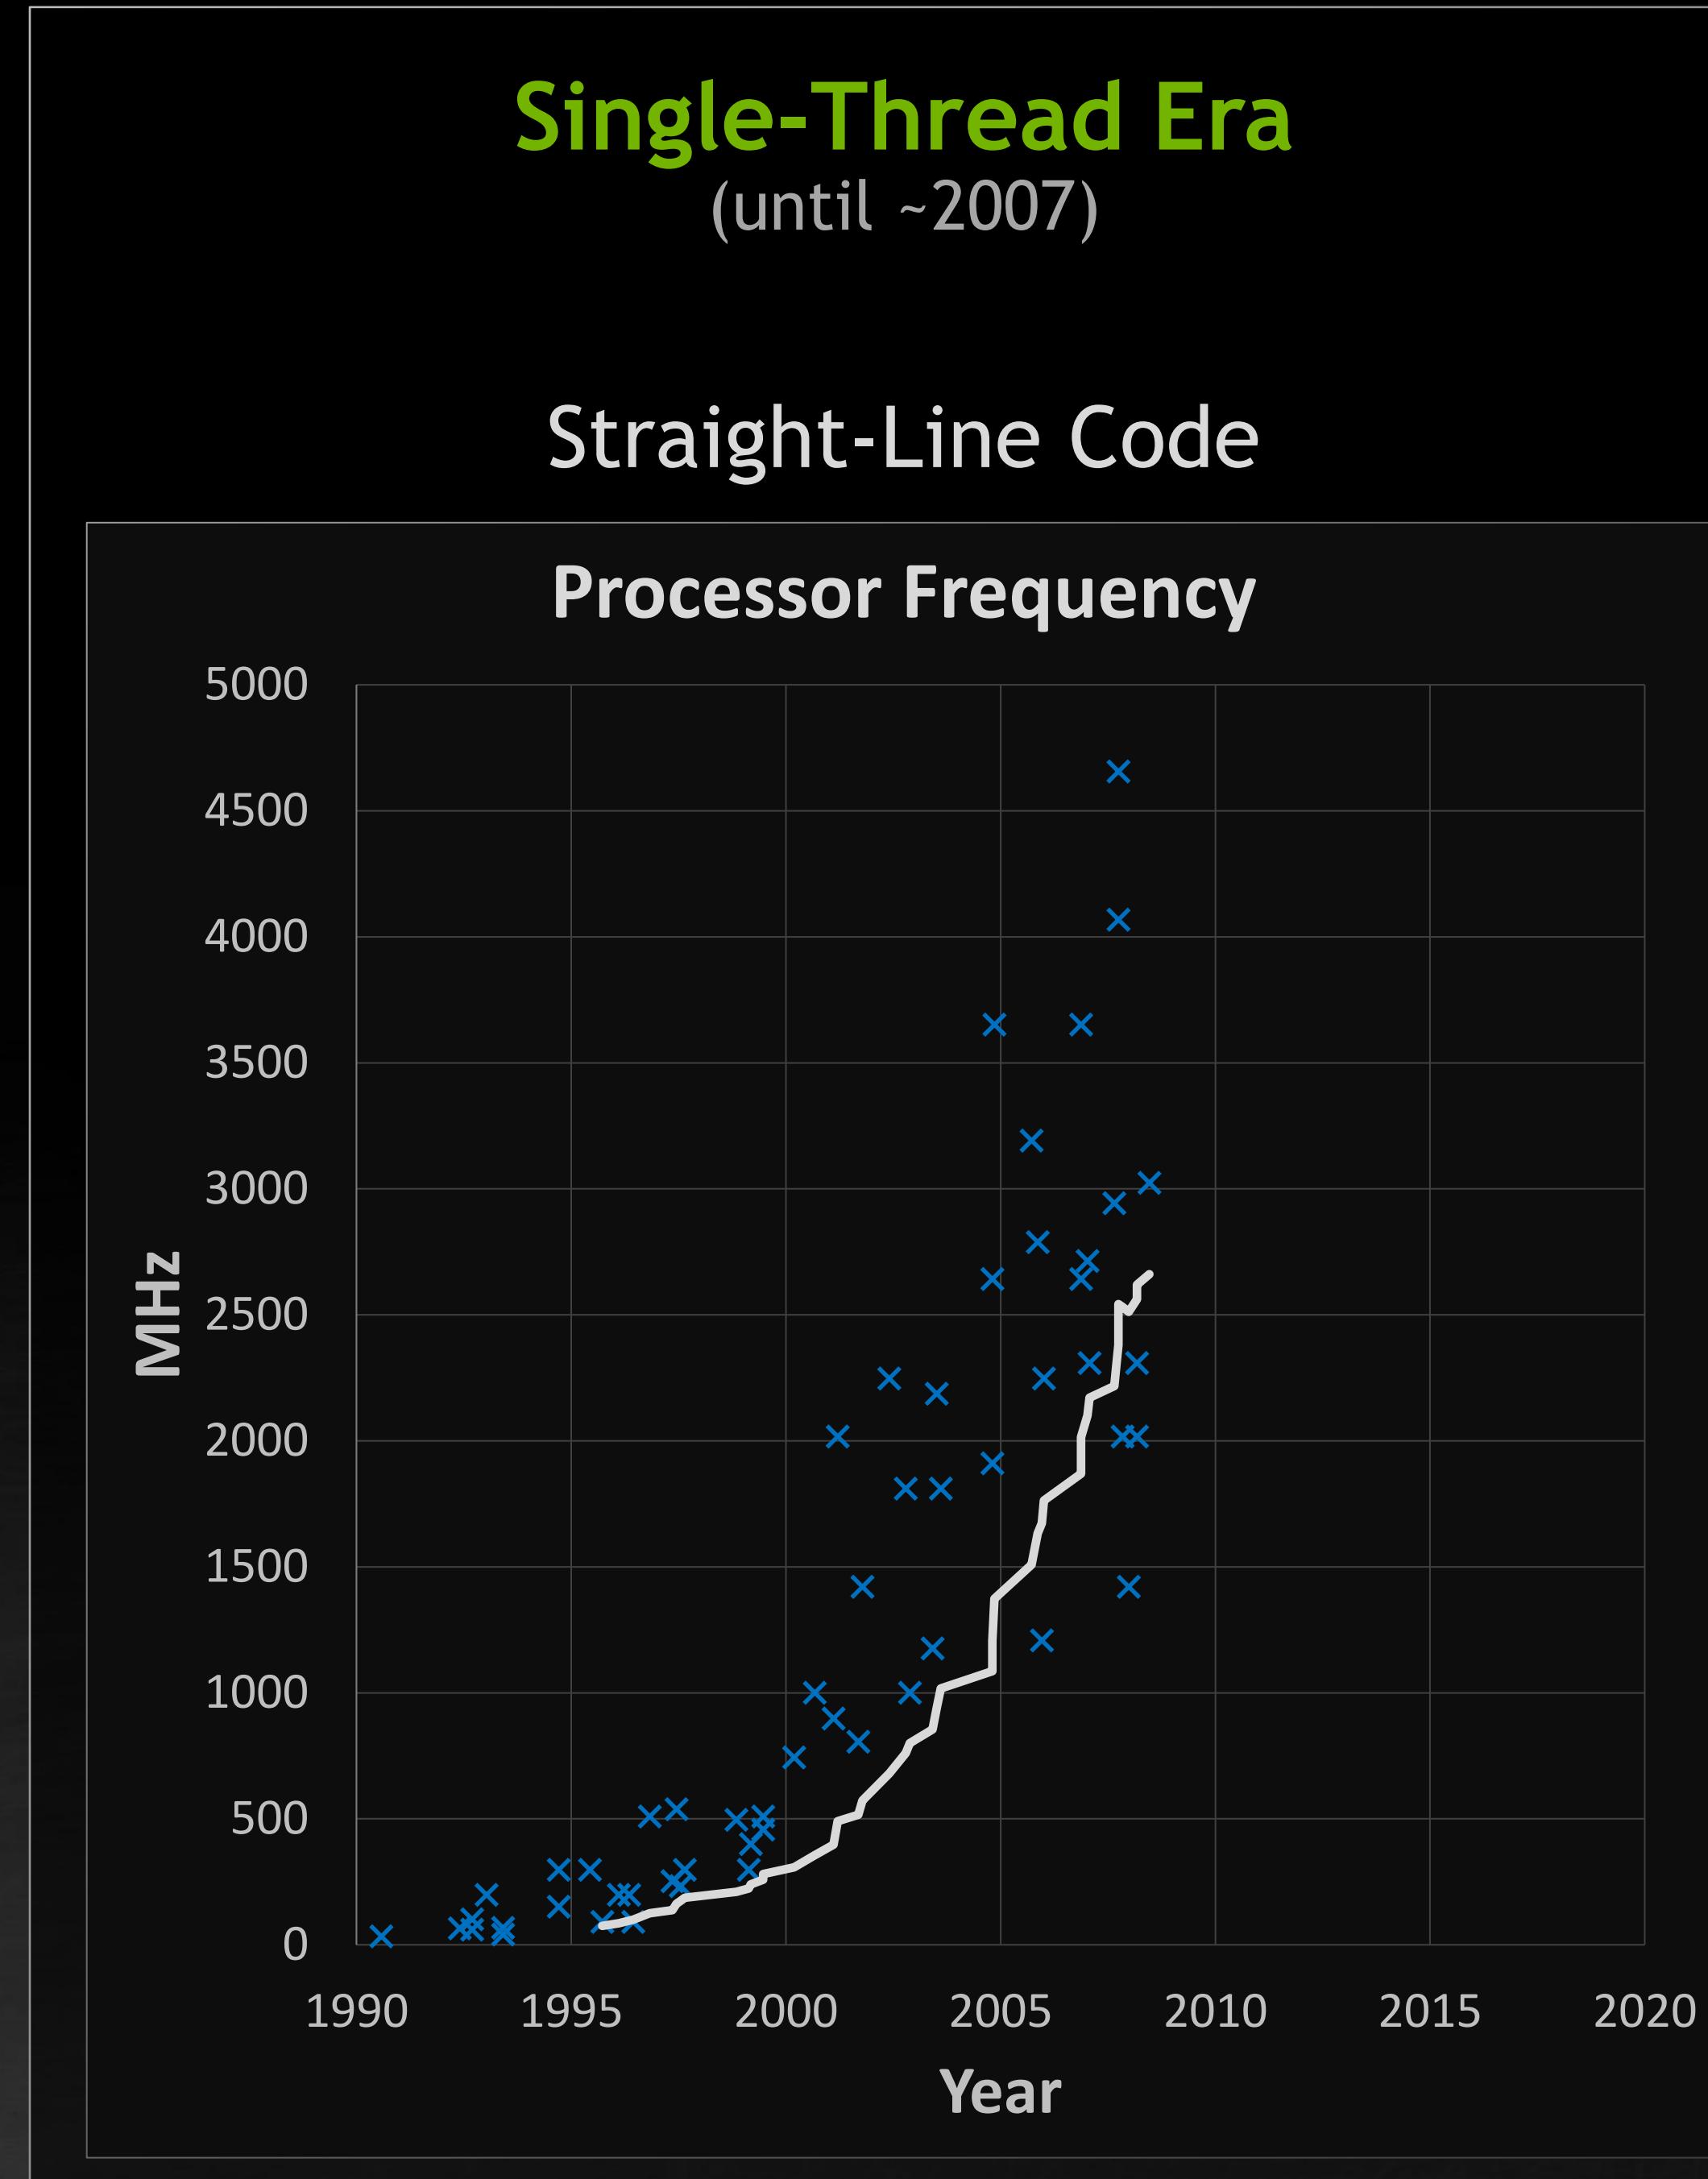
THE FIRST ERA OF SOFTWARE DEVELOPMENT



THE FIRST ERA OF SOFTWARE DEVELOPMENT

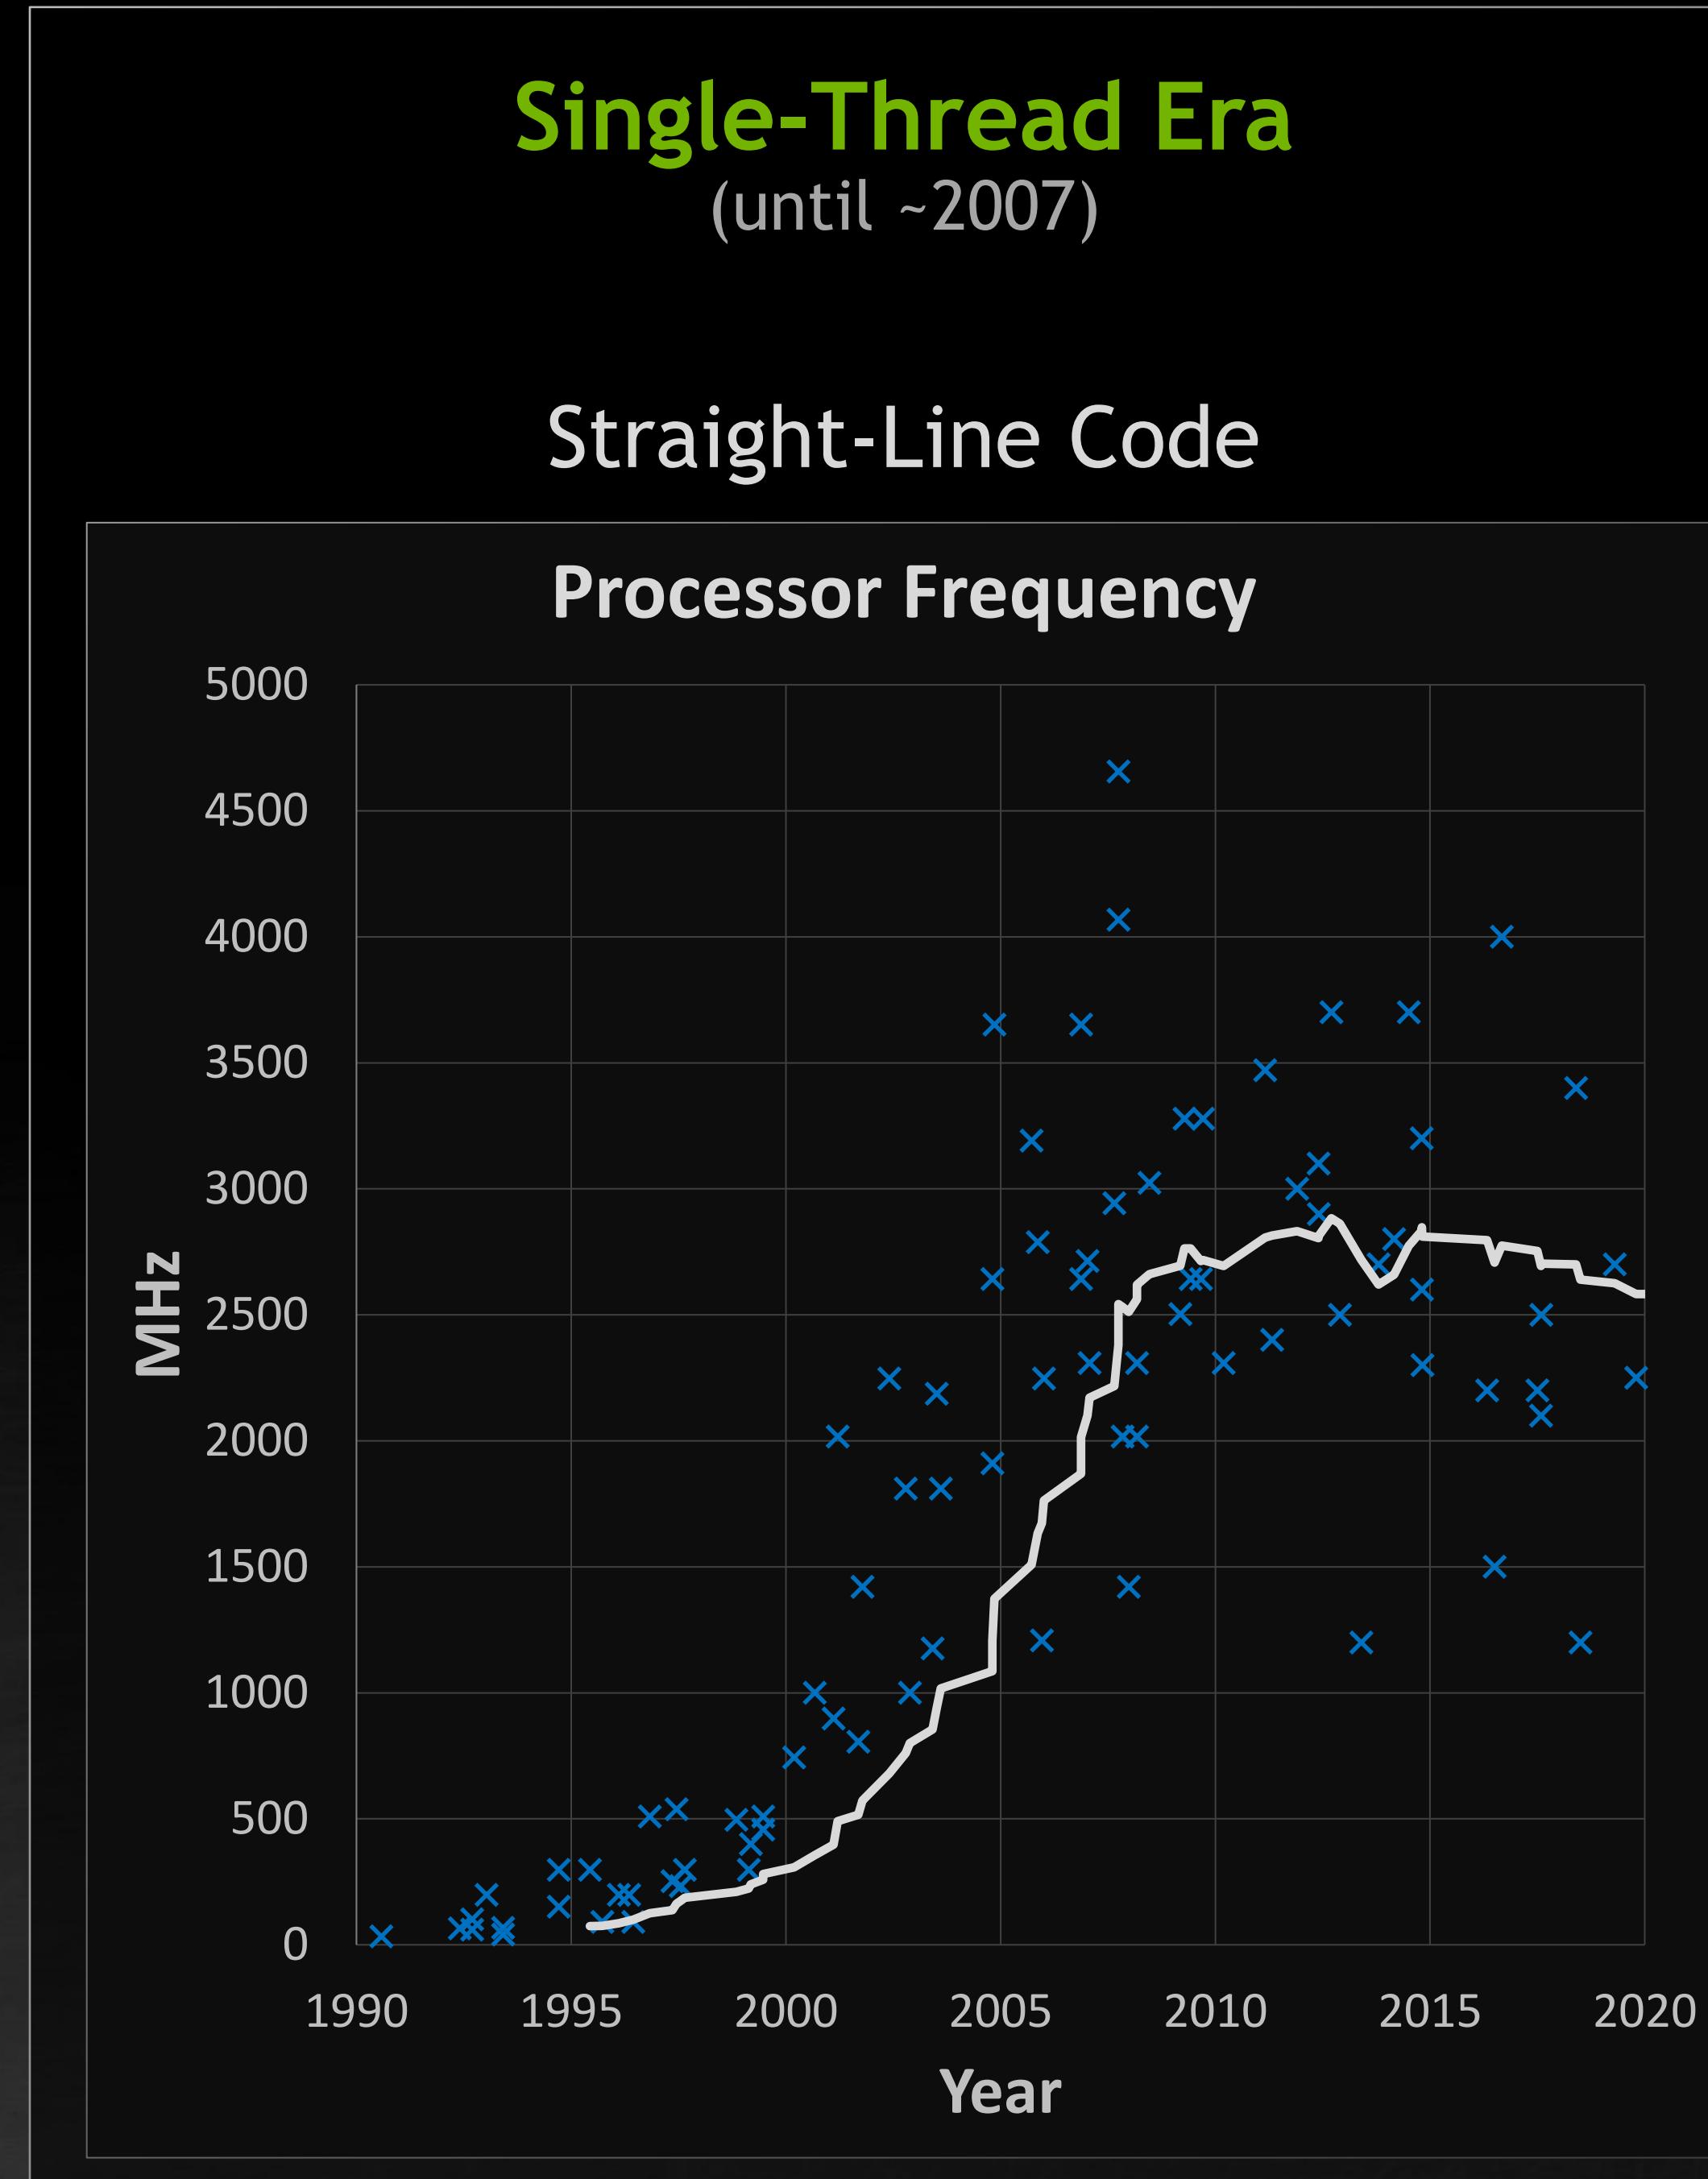


THE FIRST ERA OF SOFTWARE DEVELOPMENT

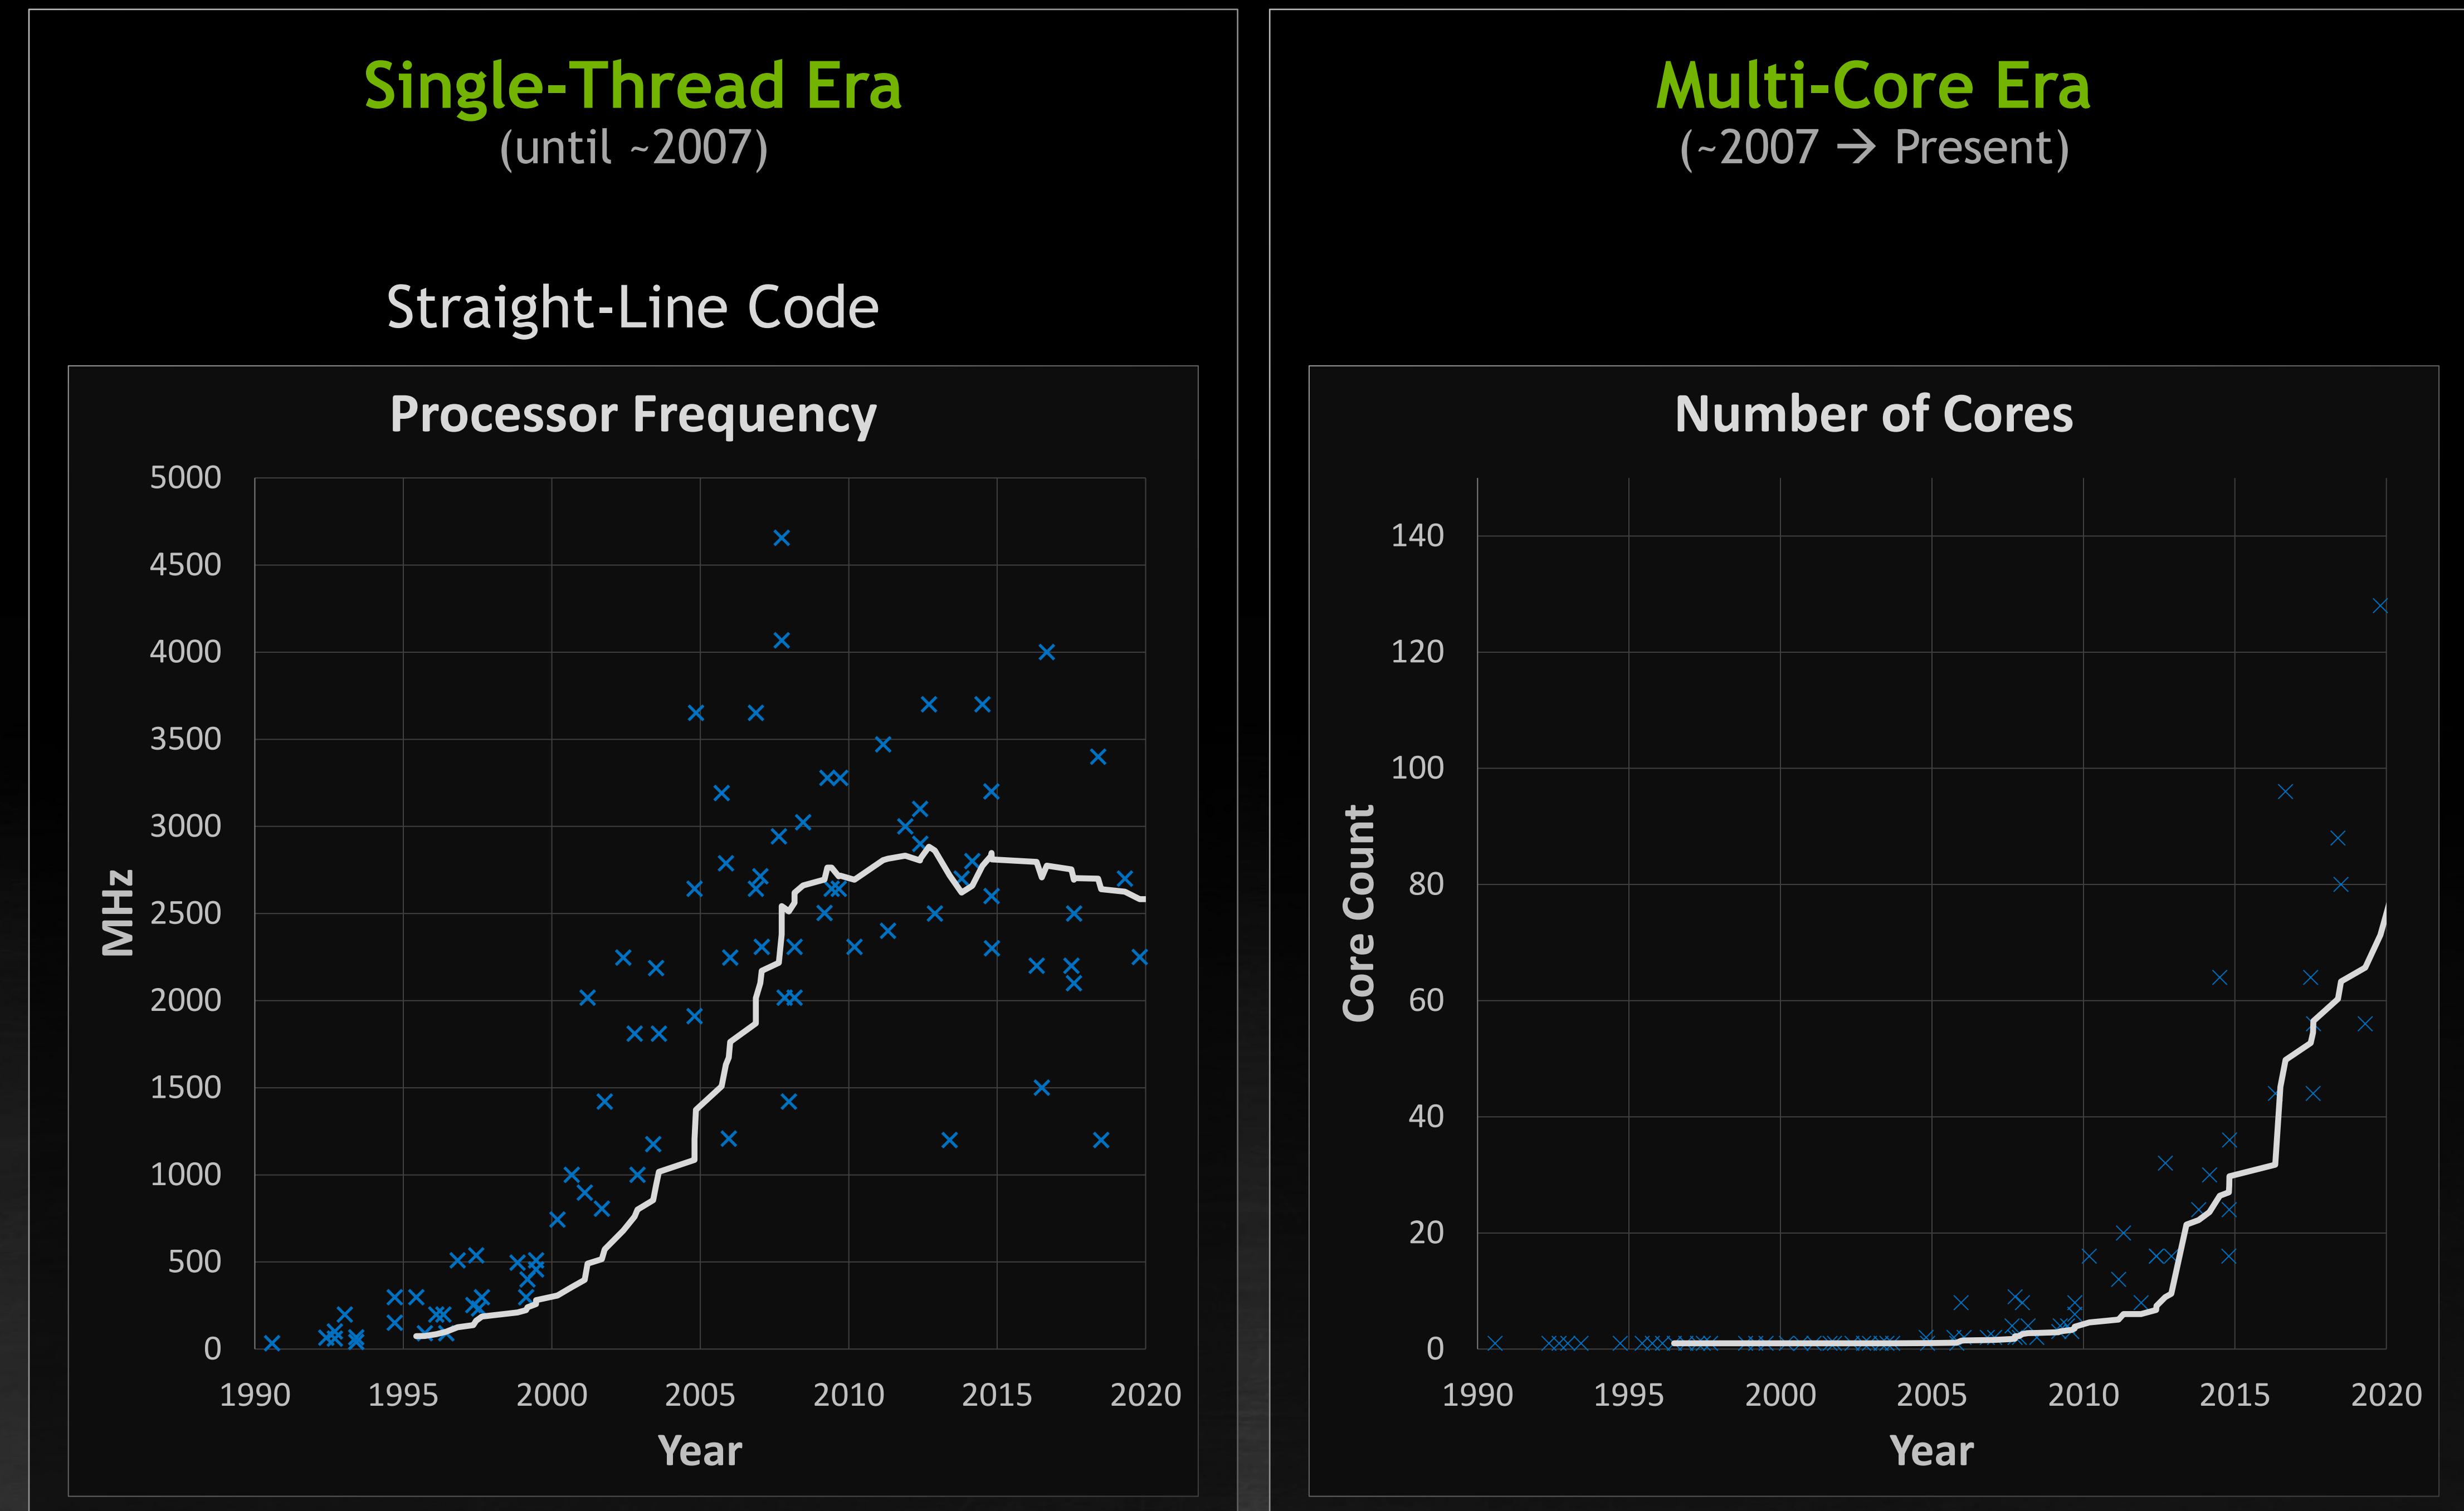


Source: Karl Rupp “40 years of microprocessor trend data”
<https://github.com/karlrupp/microprocessor-trend-data>

THE FIRST ERA OF SOFTWARE DEVELOPMENT



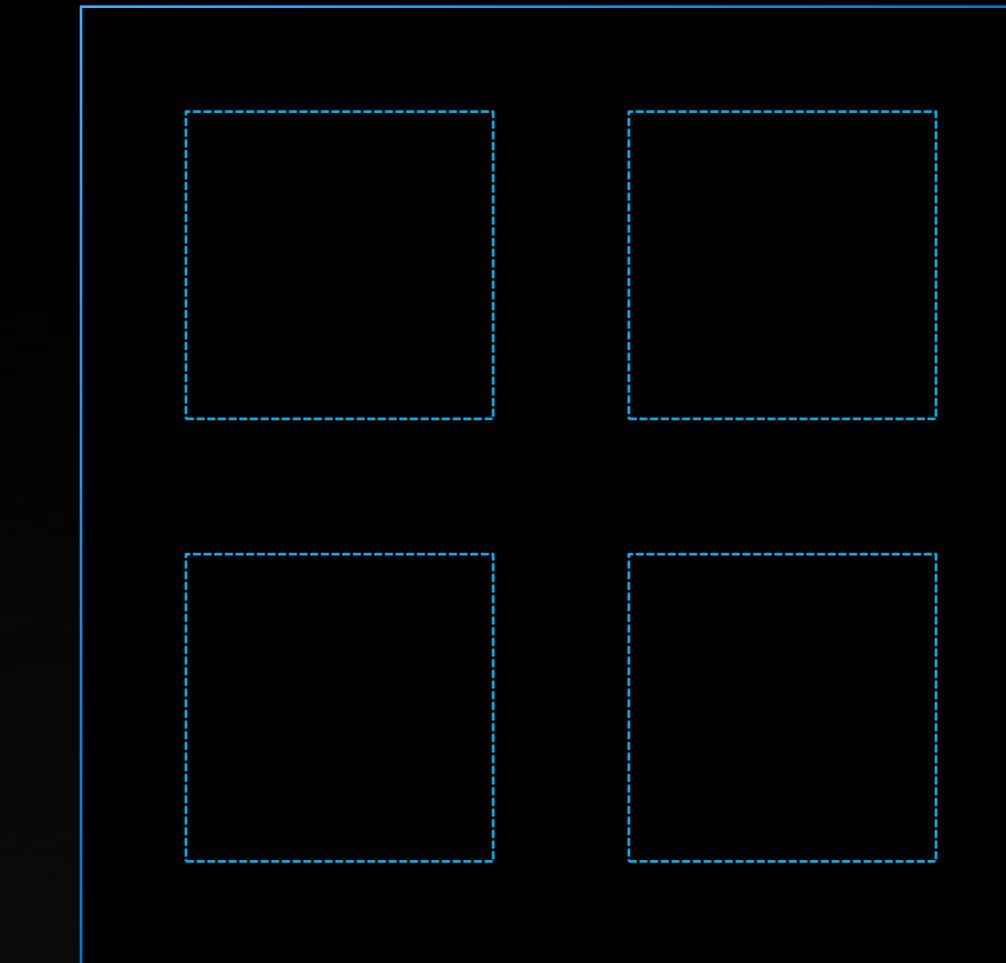
Source: Karl Rupp “40 years of microprocessor trend data”
<https://github.com/karlrupp/microprocessor-trend-data>



Source: Karl Rupp “40 years of microprocessor trend data”
<https://github.com/karlrupp/microprocessor-trend-data>

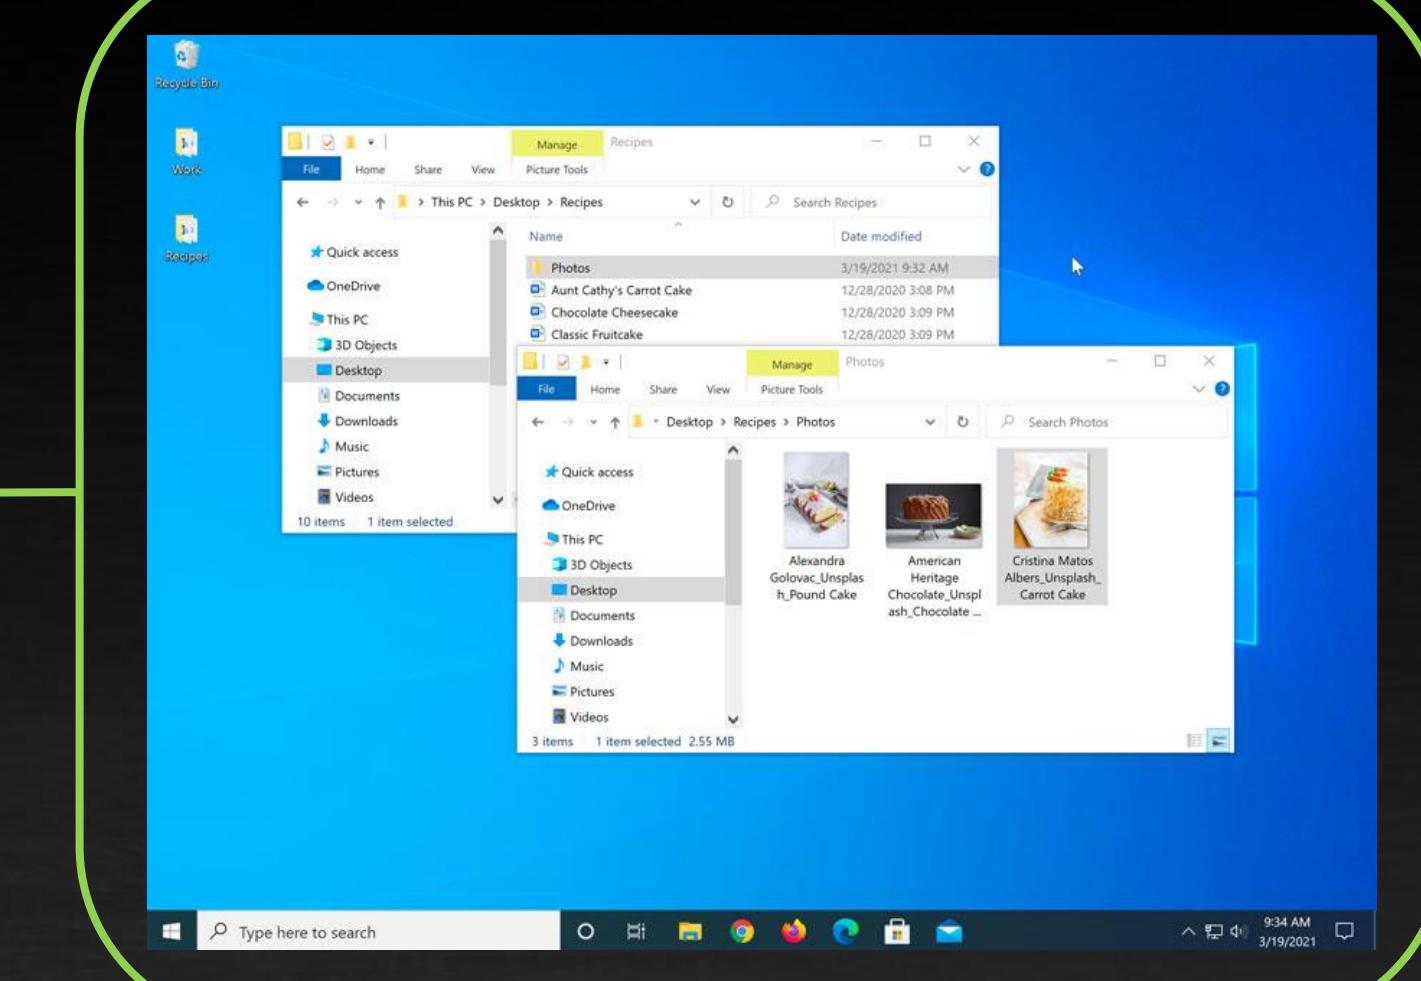
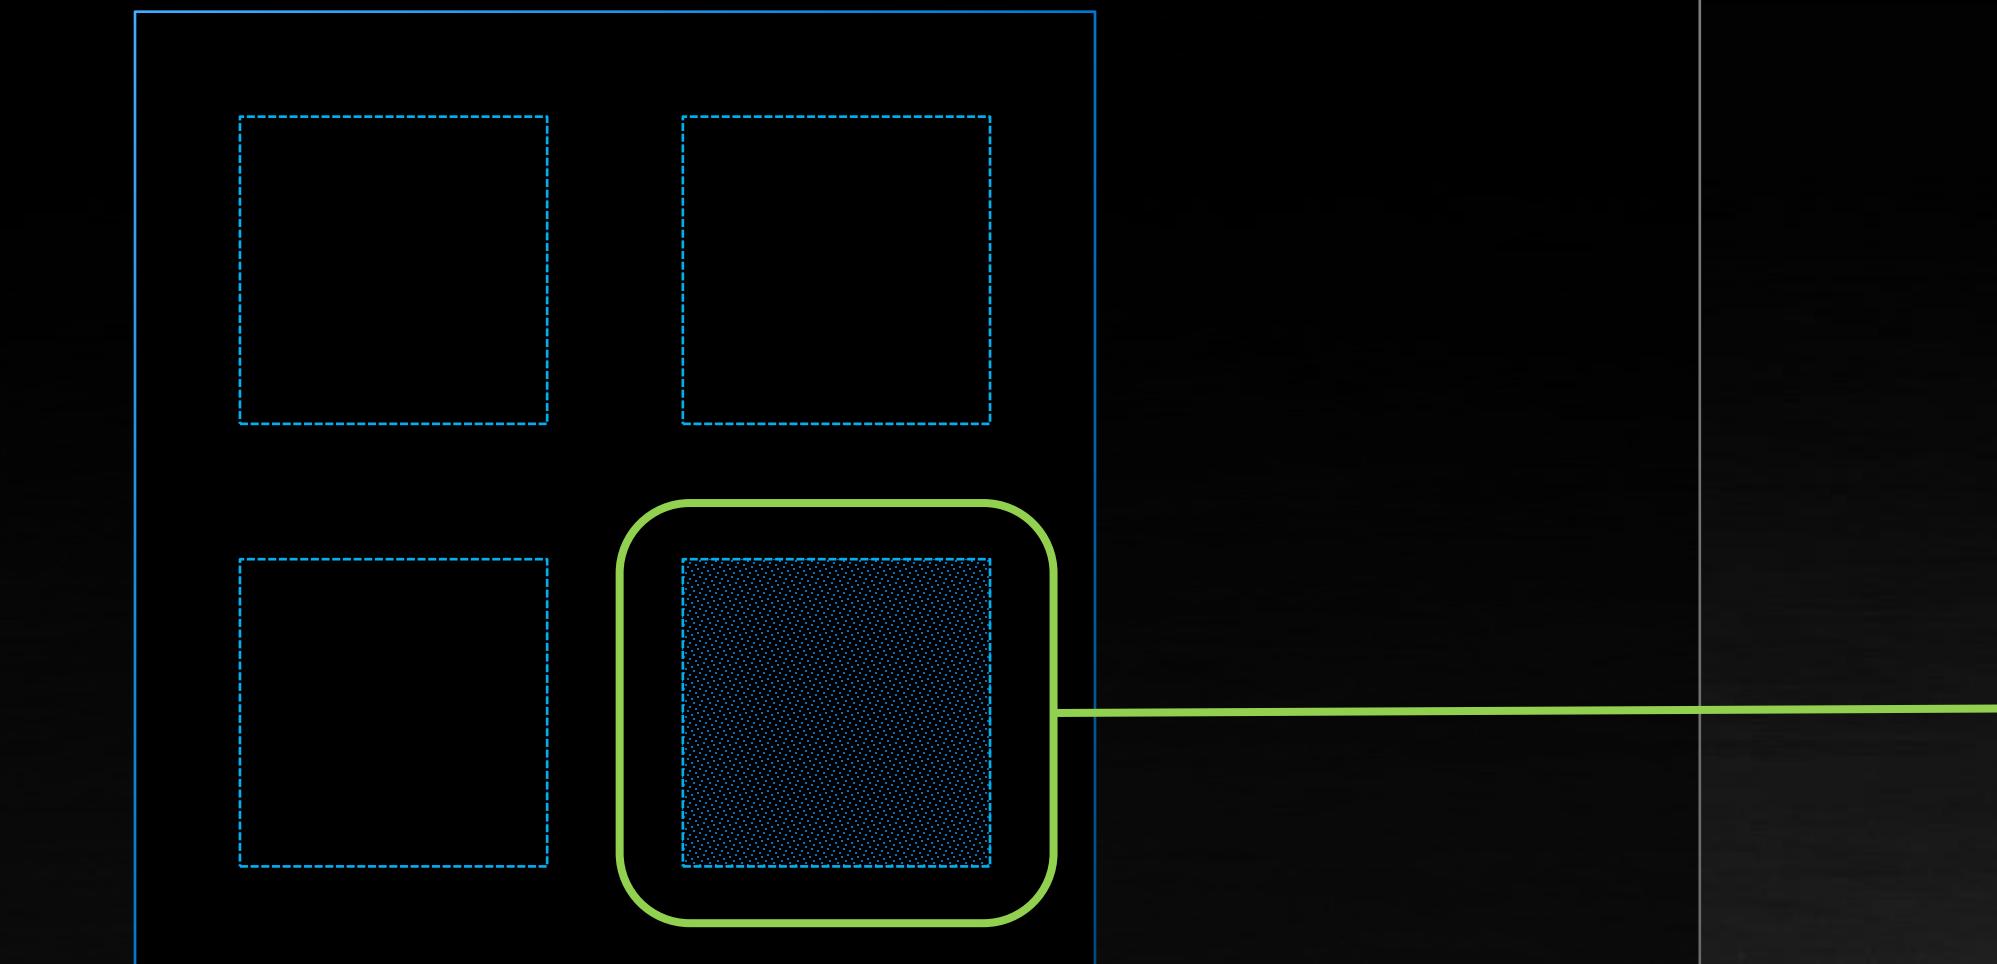
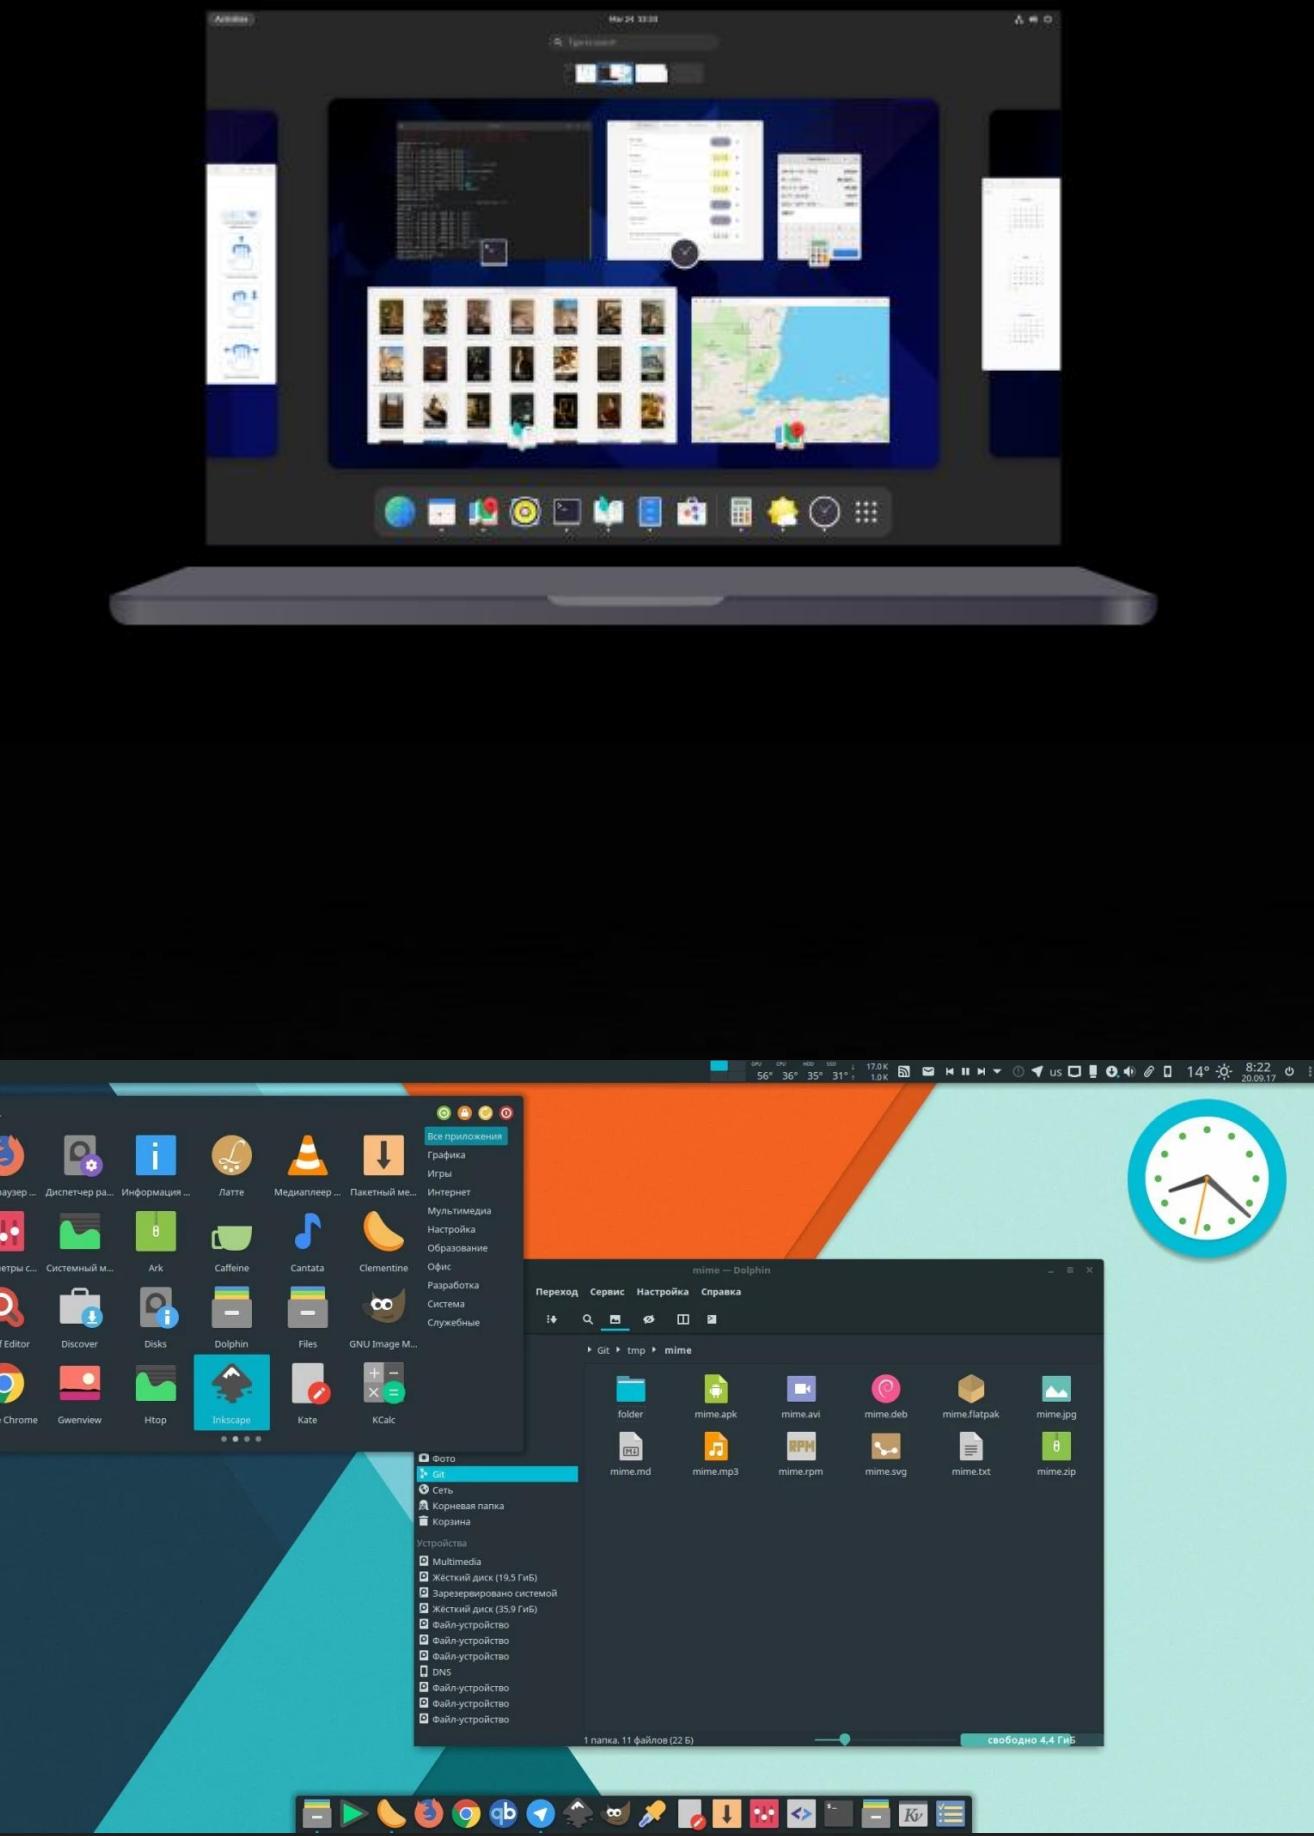
THE SECOND ERA OF SOFTWARE DEVELOPMENT

Multi-Core Era
(~2007 → Present)



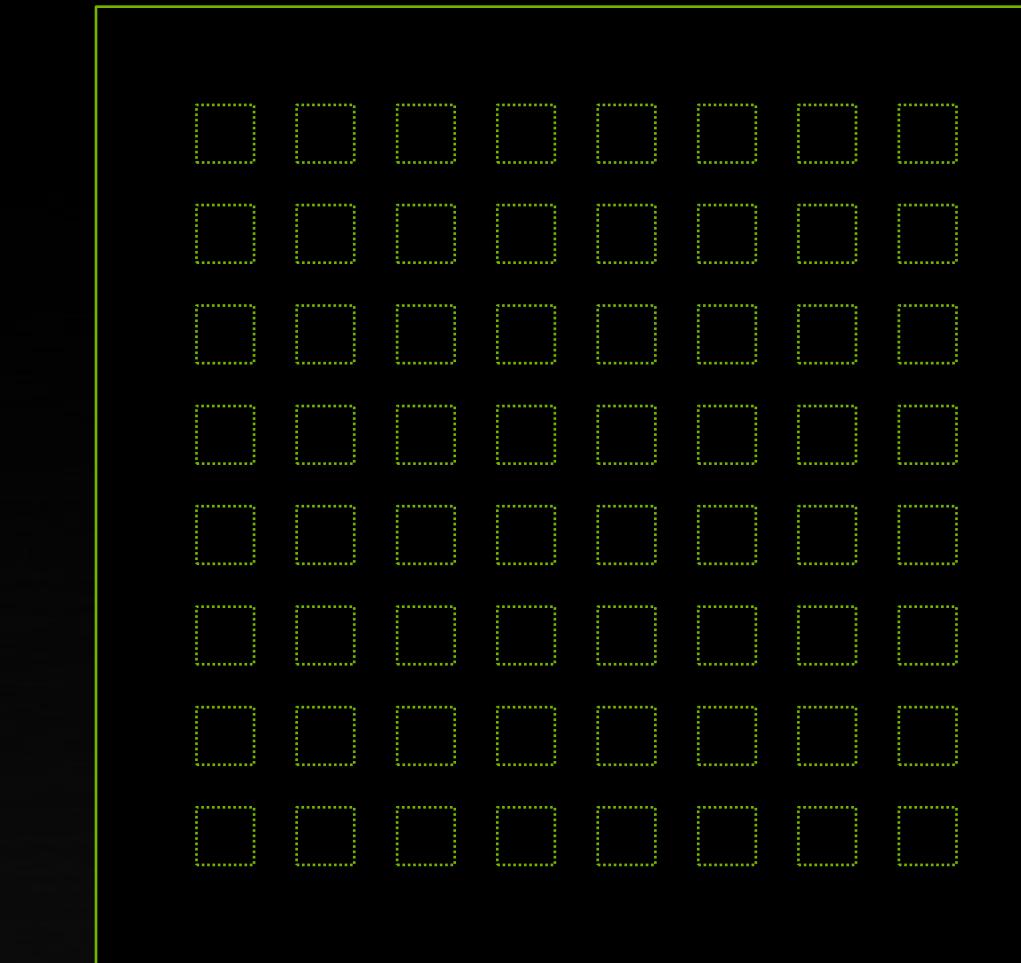
THE SECOND ERA OF SOFTWARE DEVELOPMENT

Multi-Core Era (~2007 → Present)



THE SECOND ERA OF SOFTWARE DEVELOPMENT

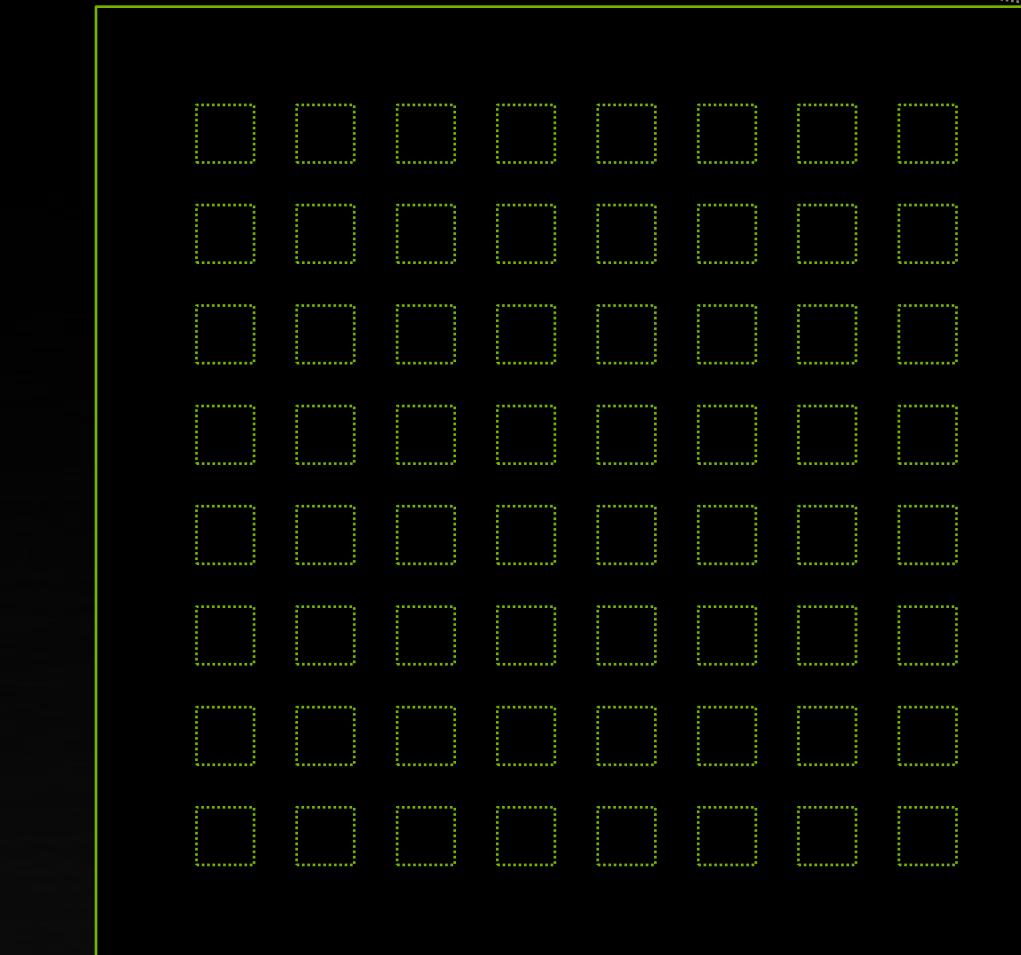
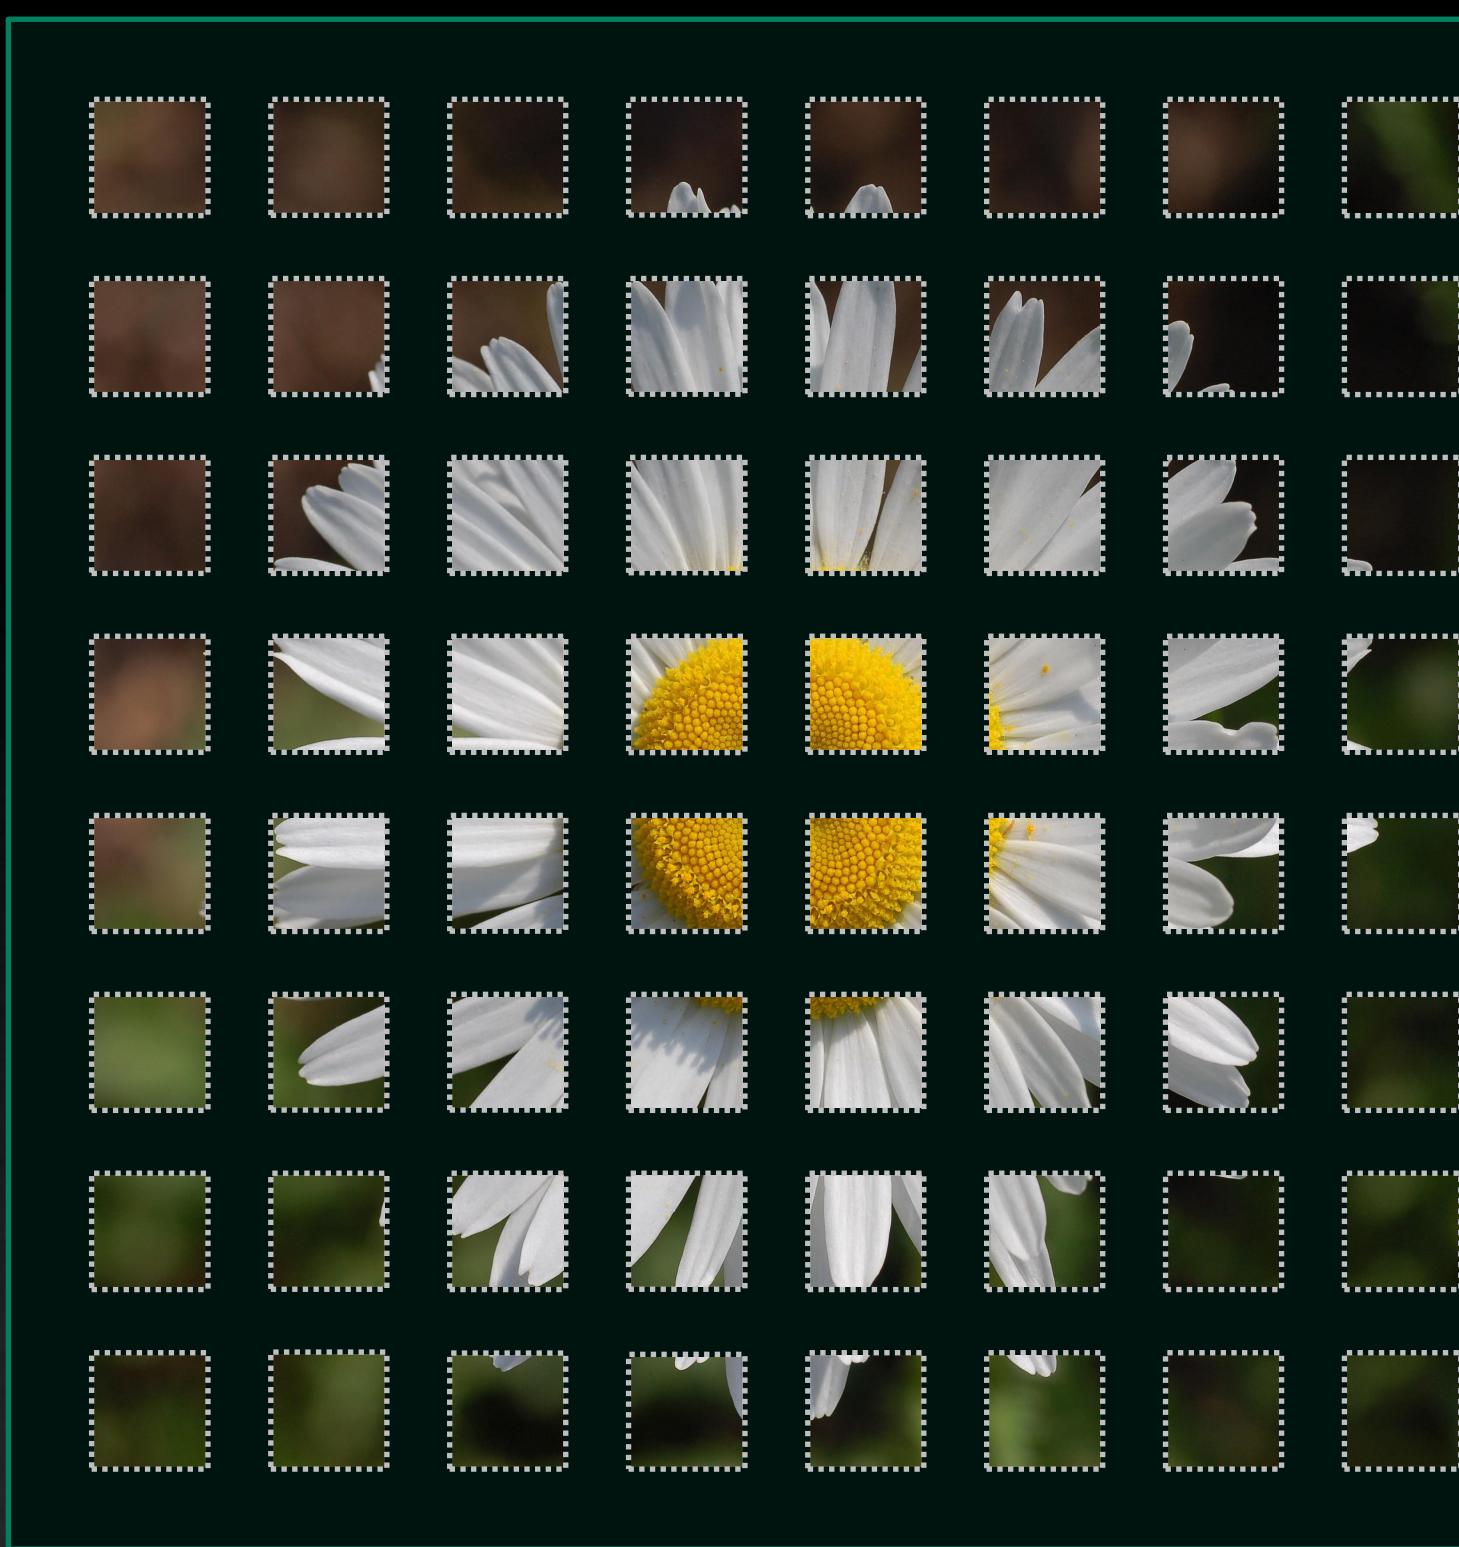
Multi-Core Era
(~2007 → Present)



THE SECOND ERA OF SOFTWARE DEVELOPMENT

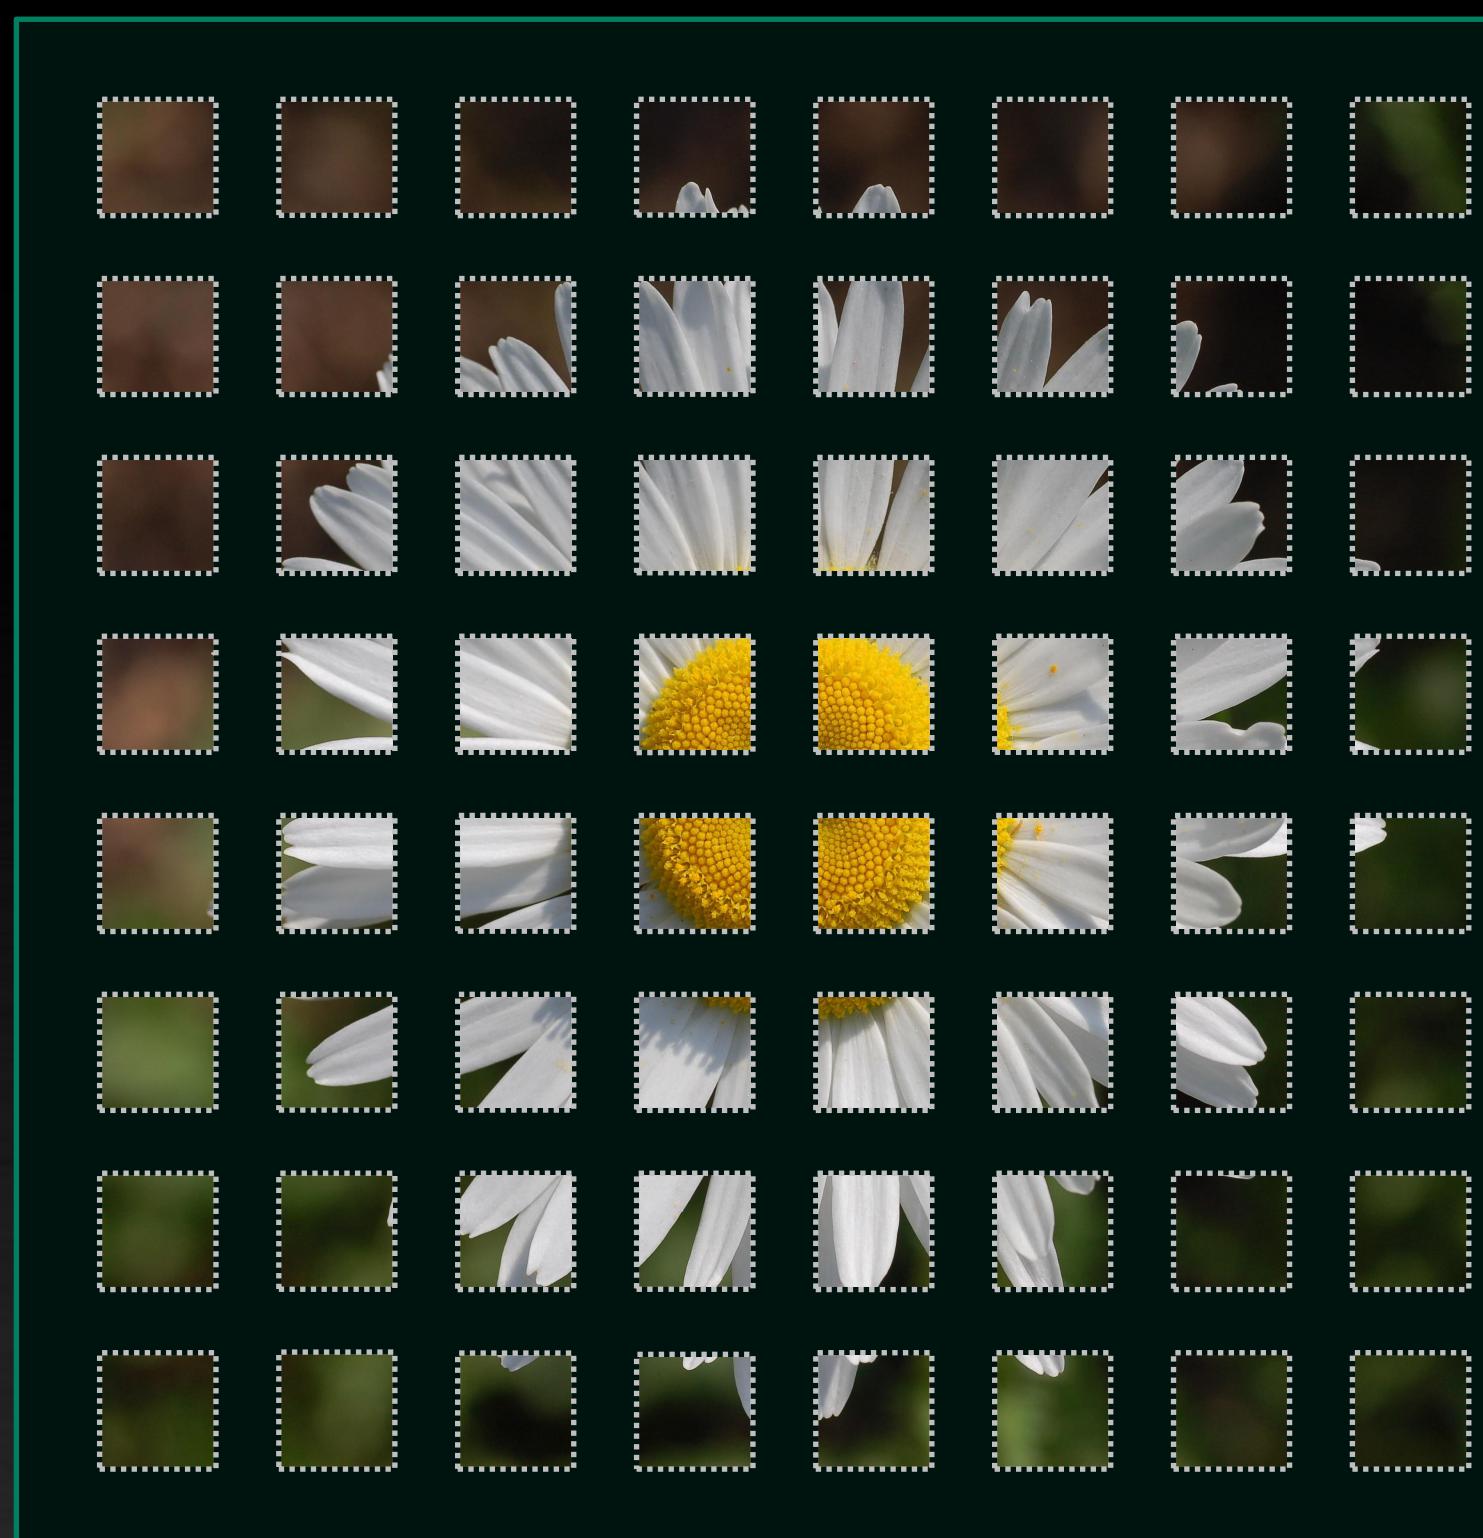
Multi-Core Era
(~2007 → Present)

Data Parallelism

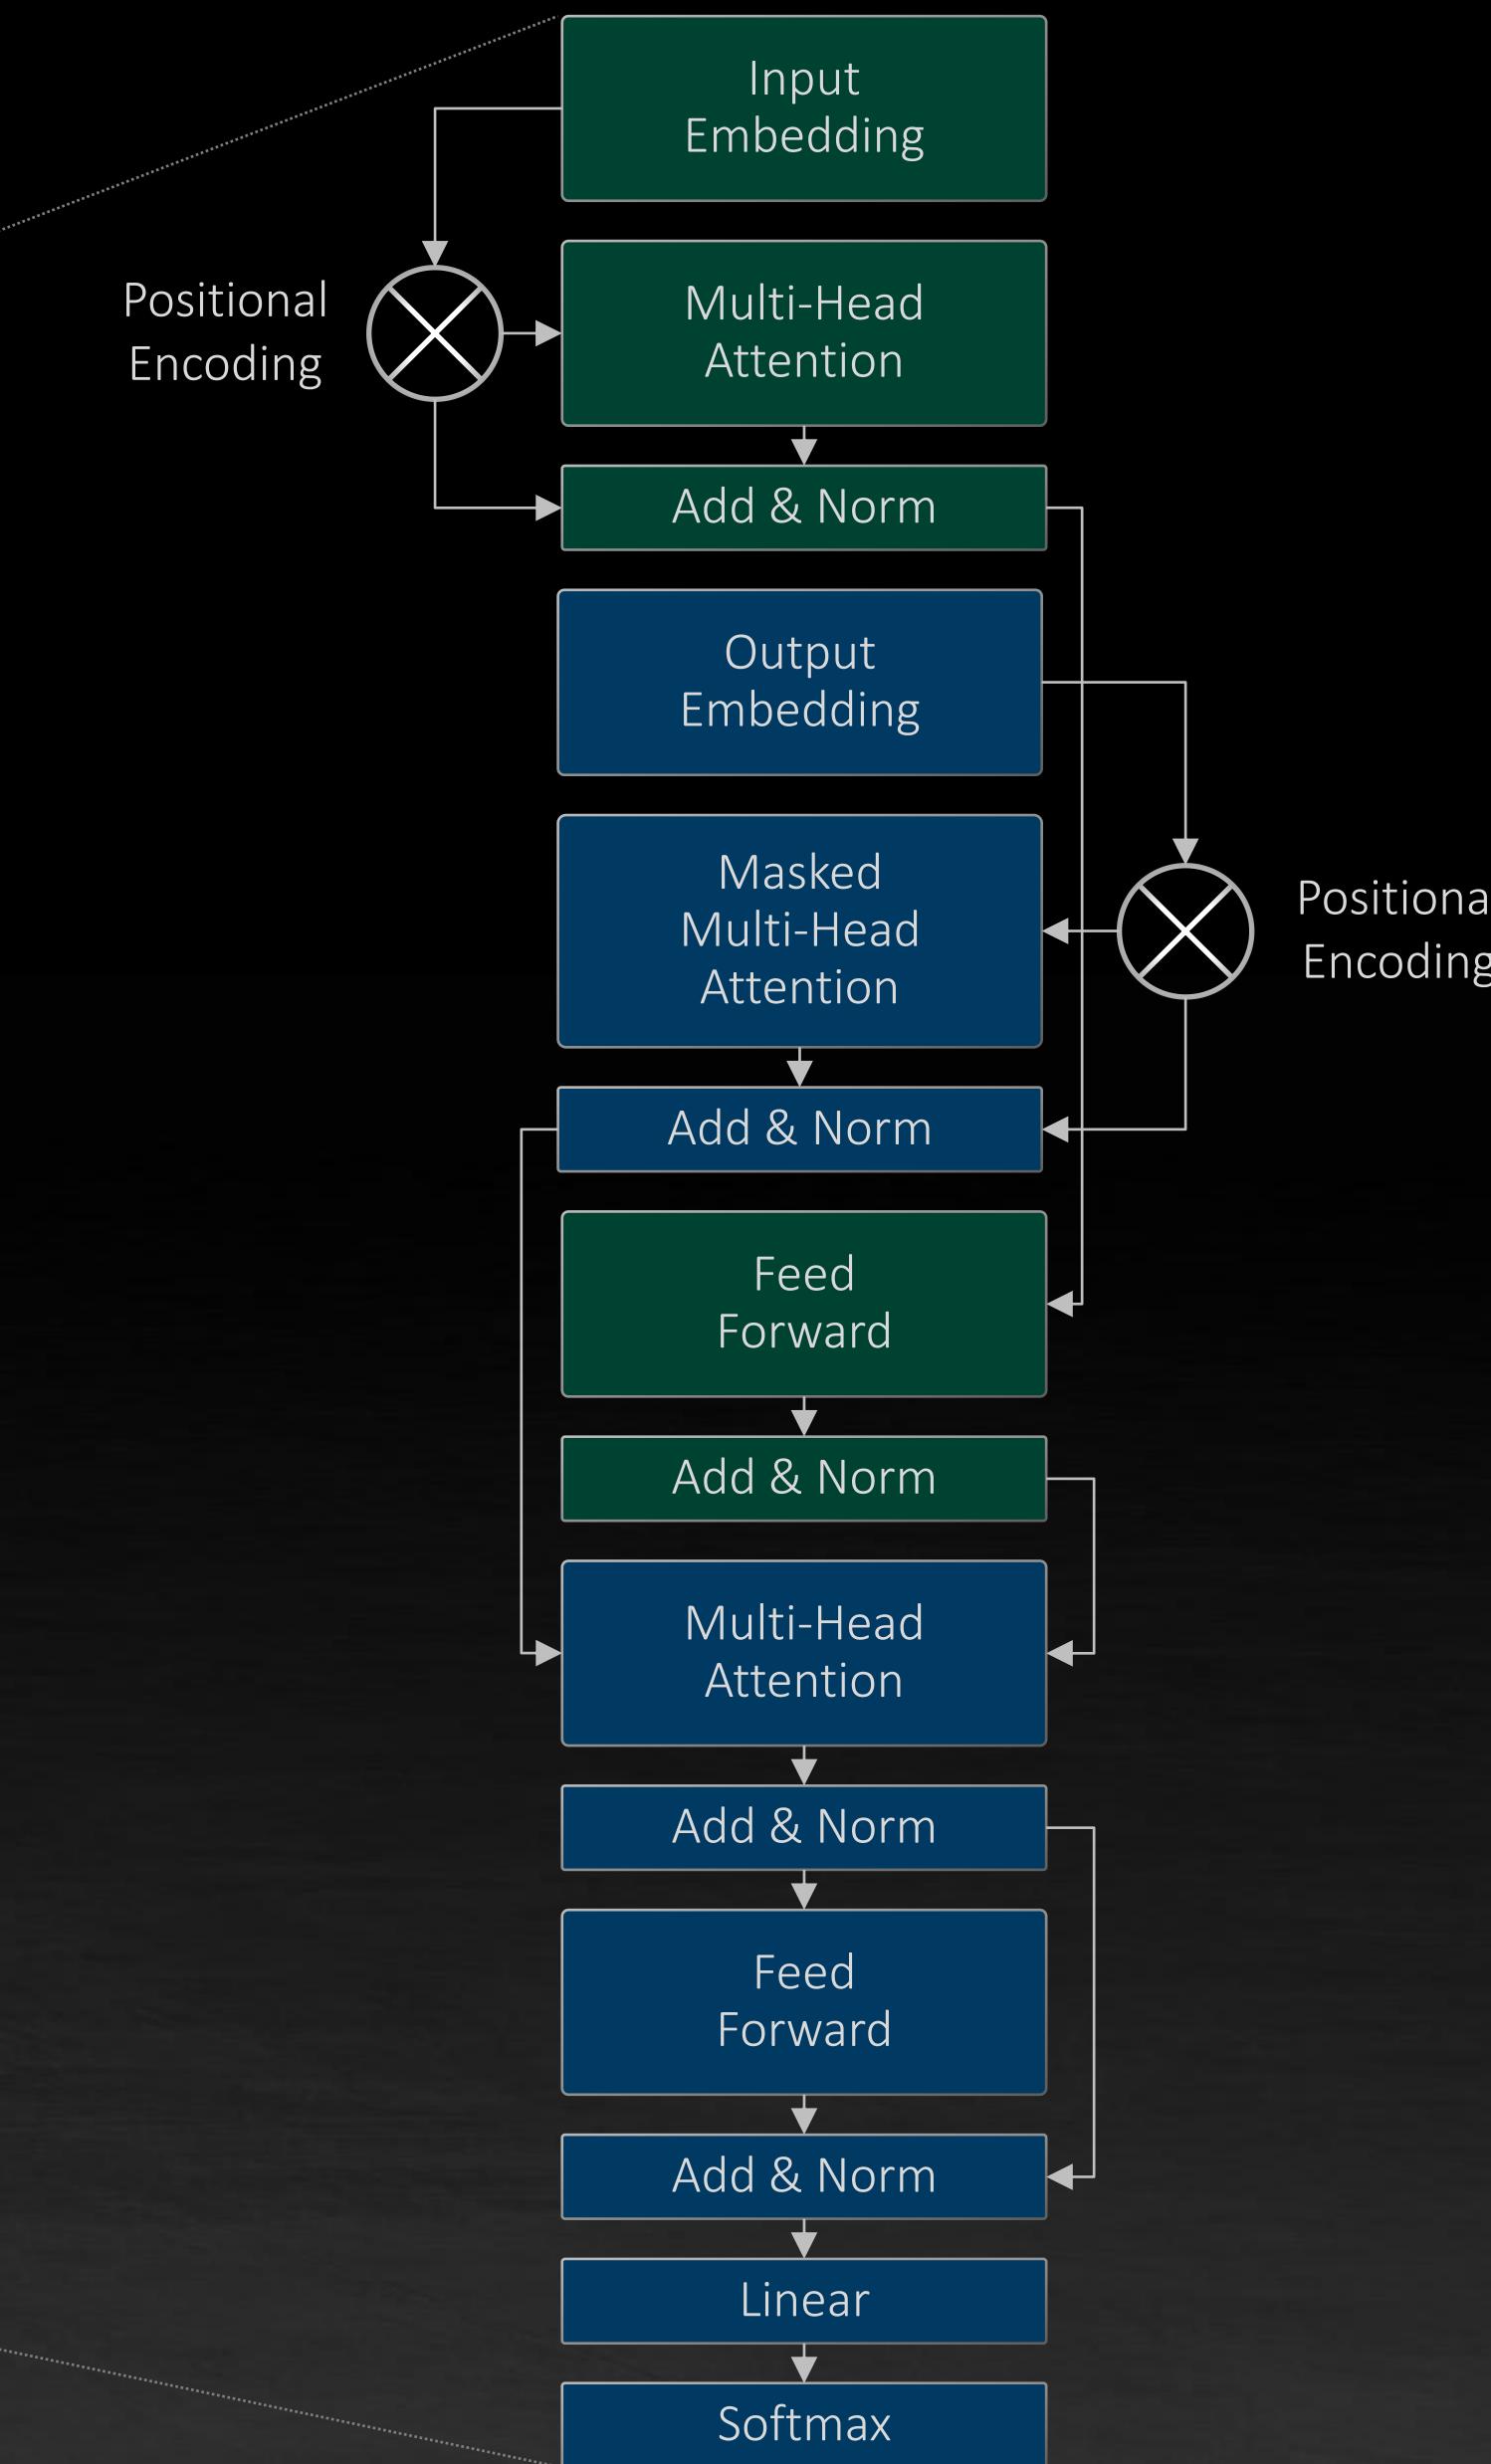
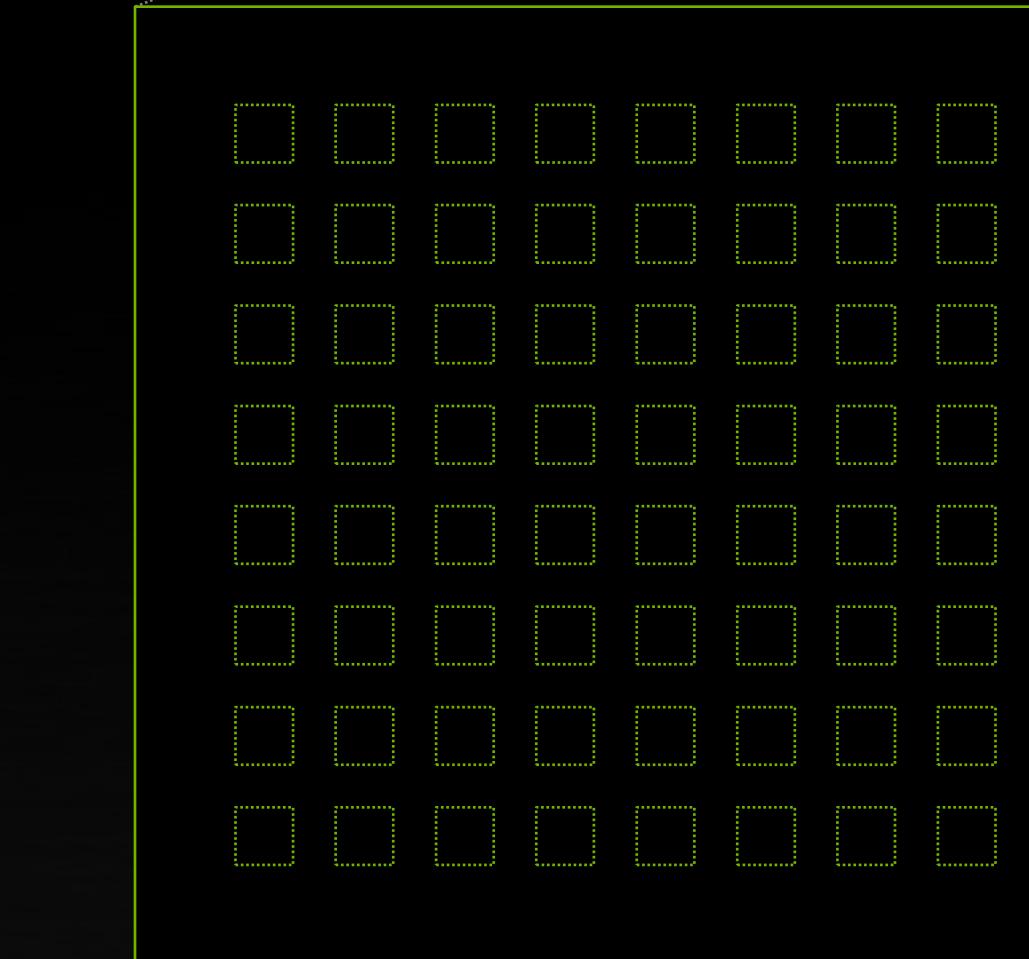


THE SECOND ERA OF SOFTWARE DEVELOPMENT

Data Parallelism

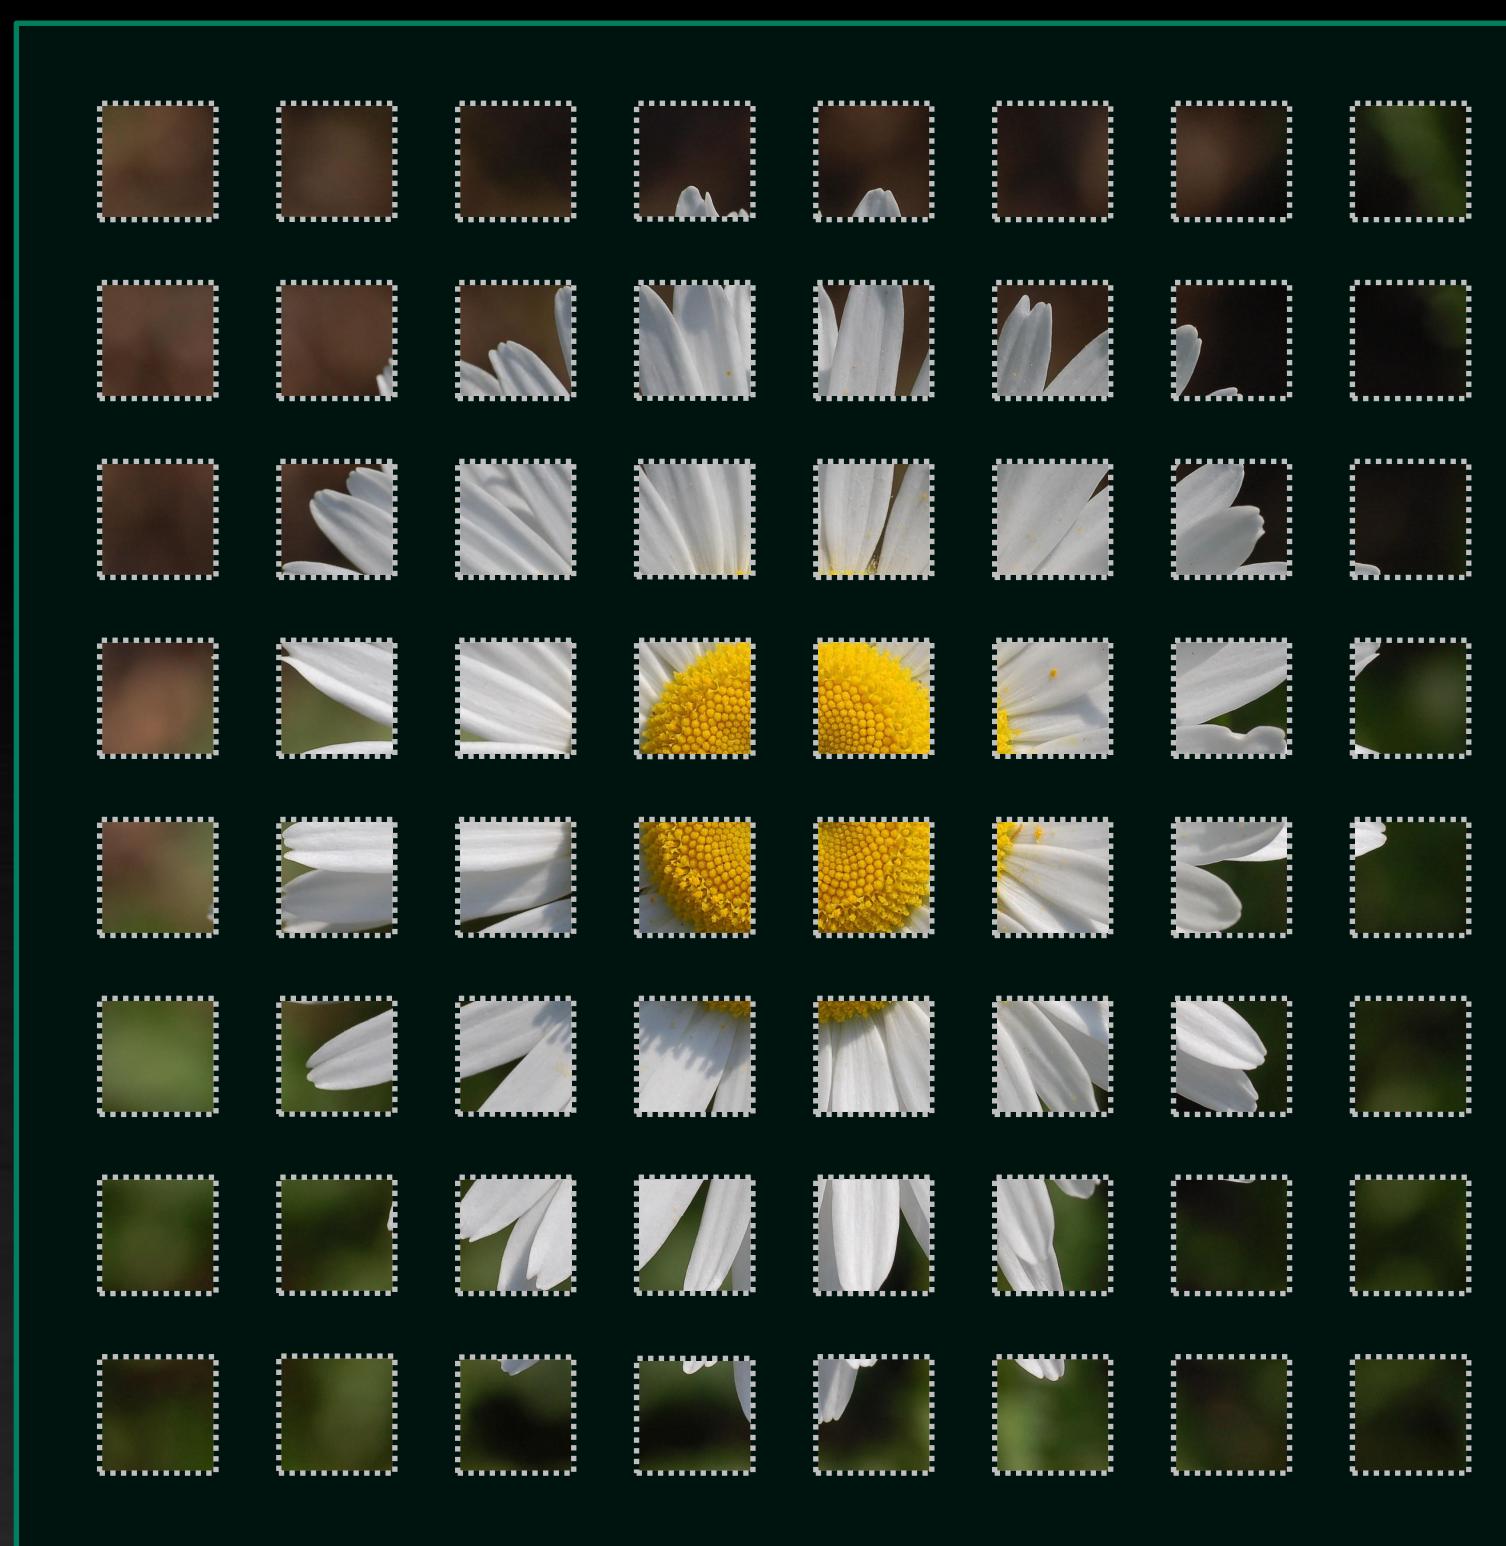


Multi-Core Era
(~2007 → Present)

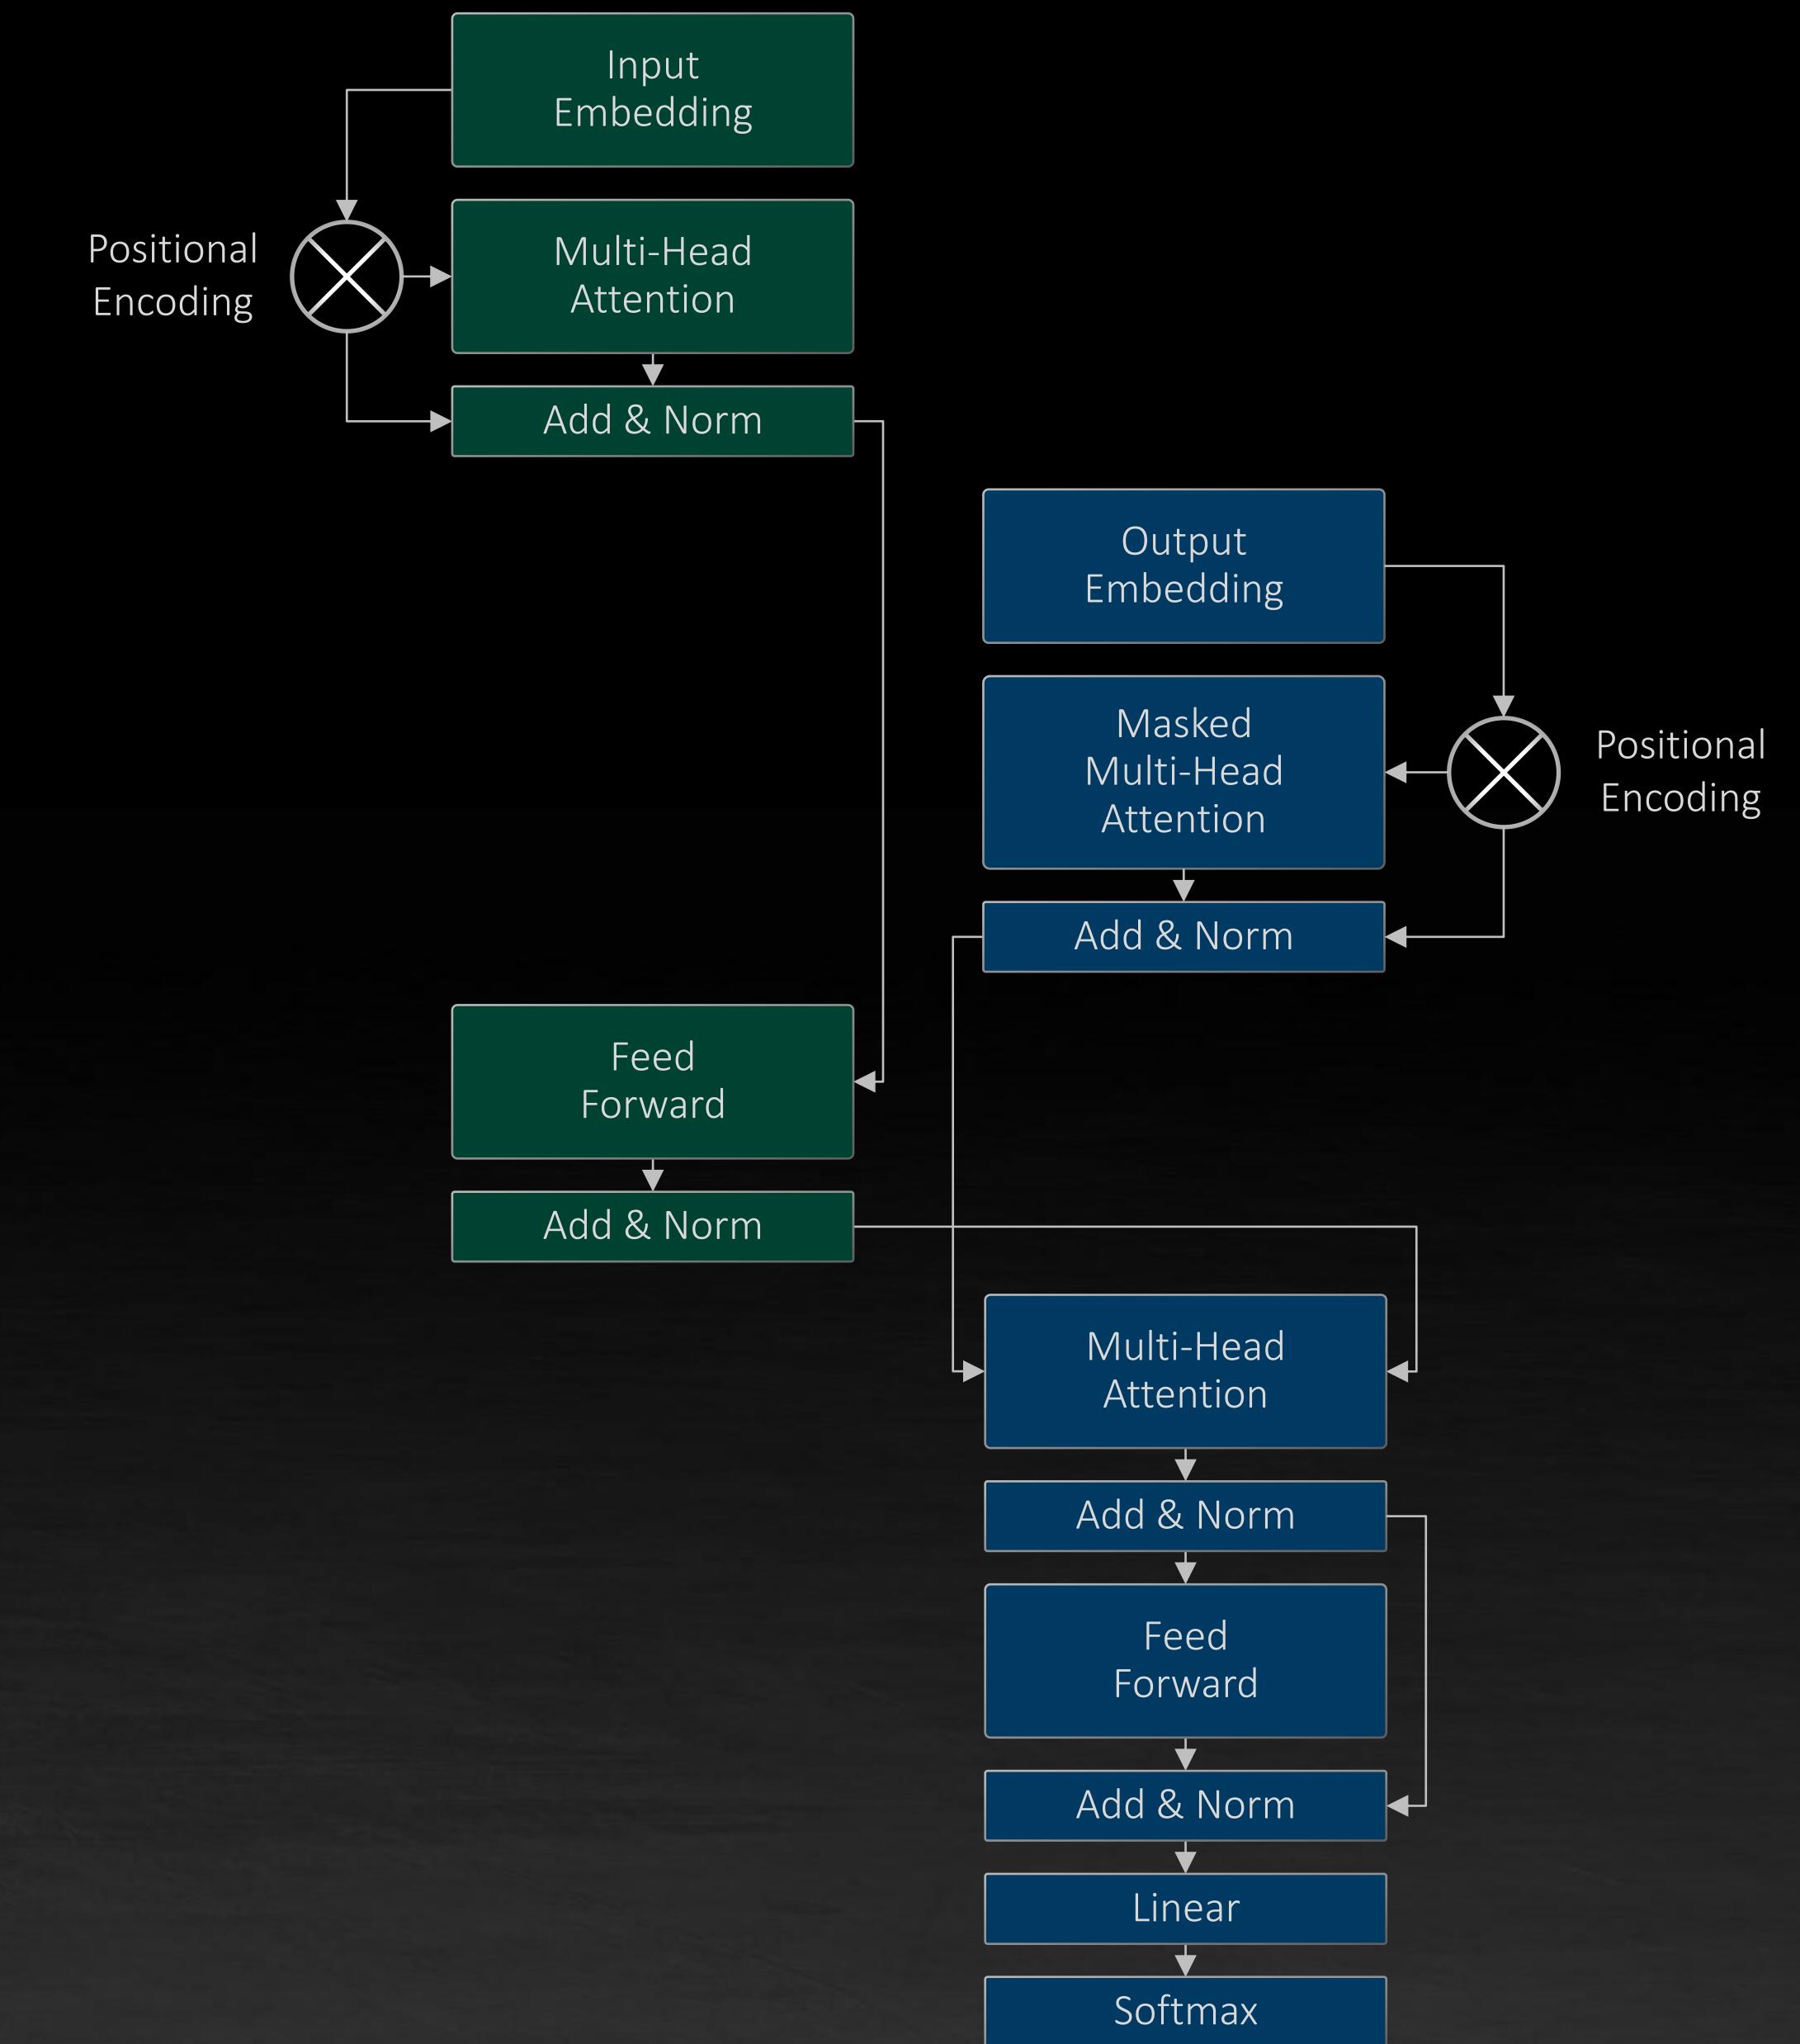
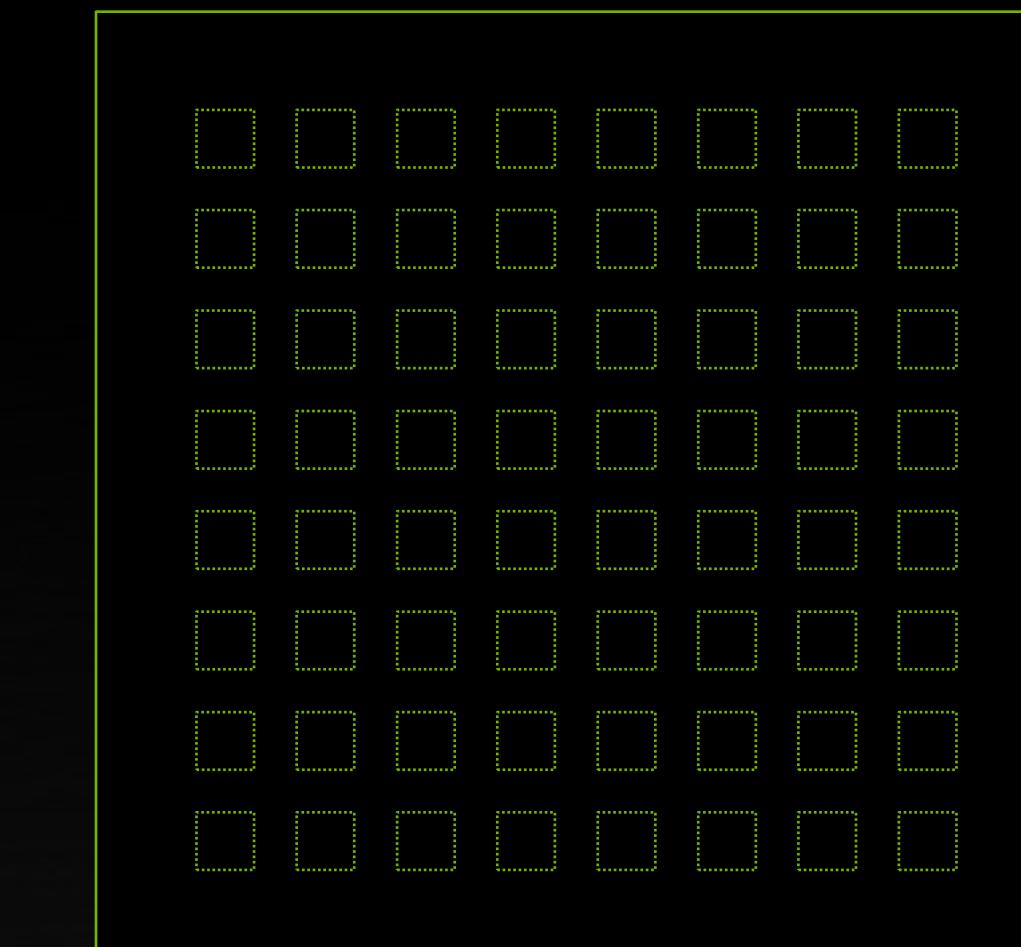


THE SECOND ERA OF SOFTWARE DEVELOPMENT

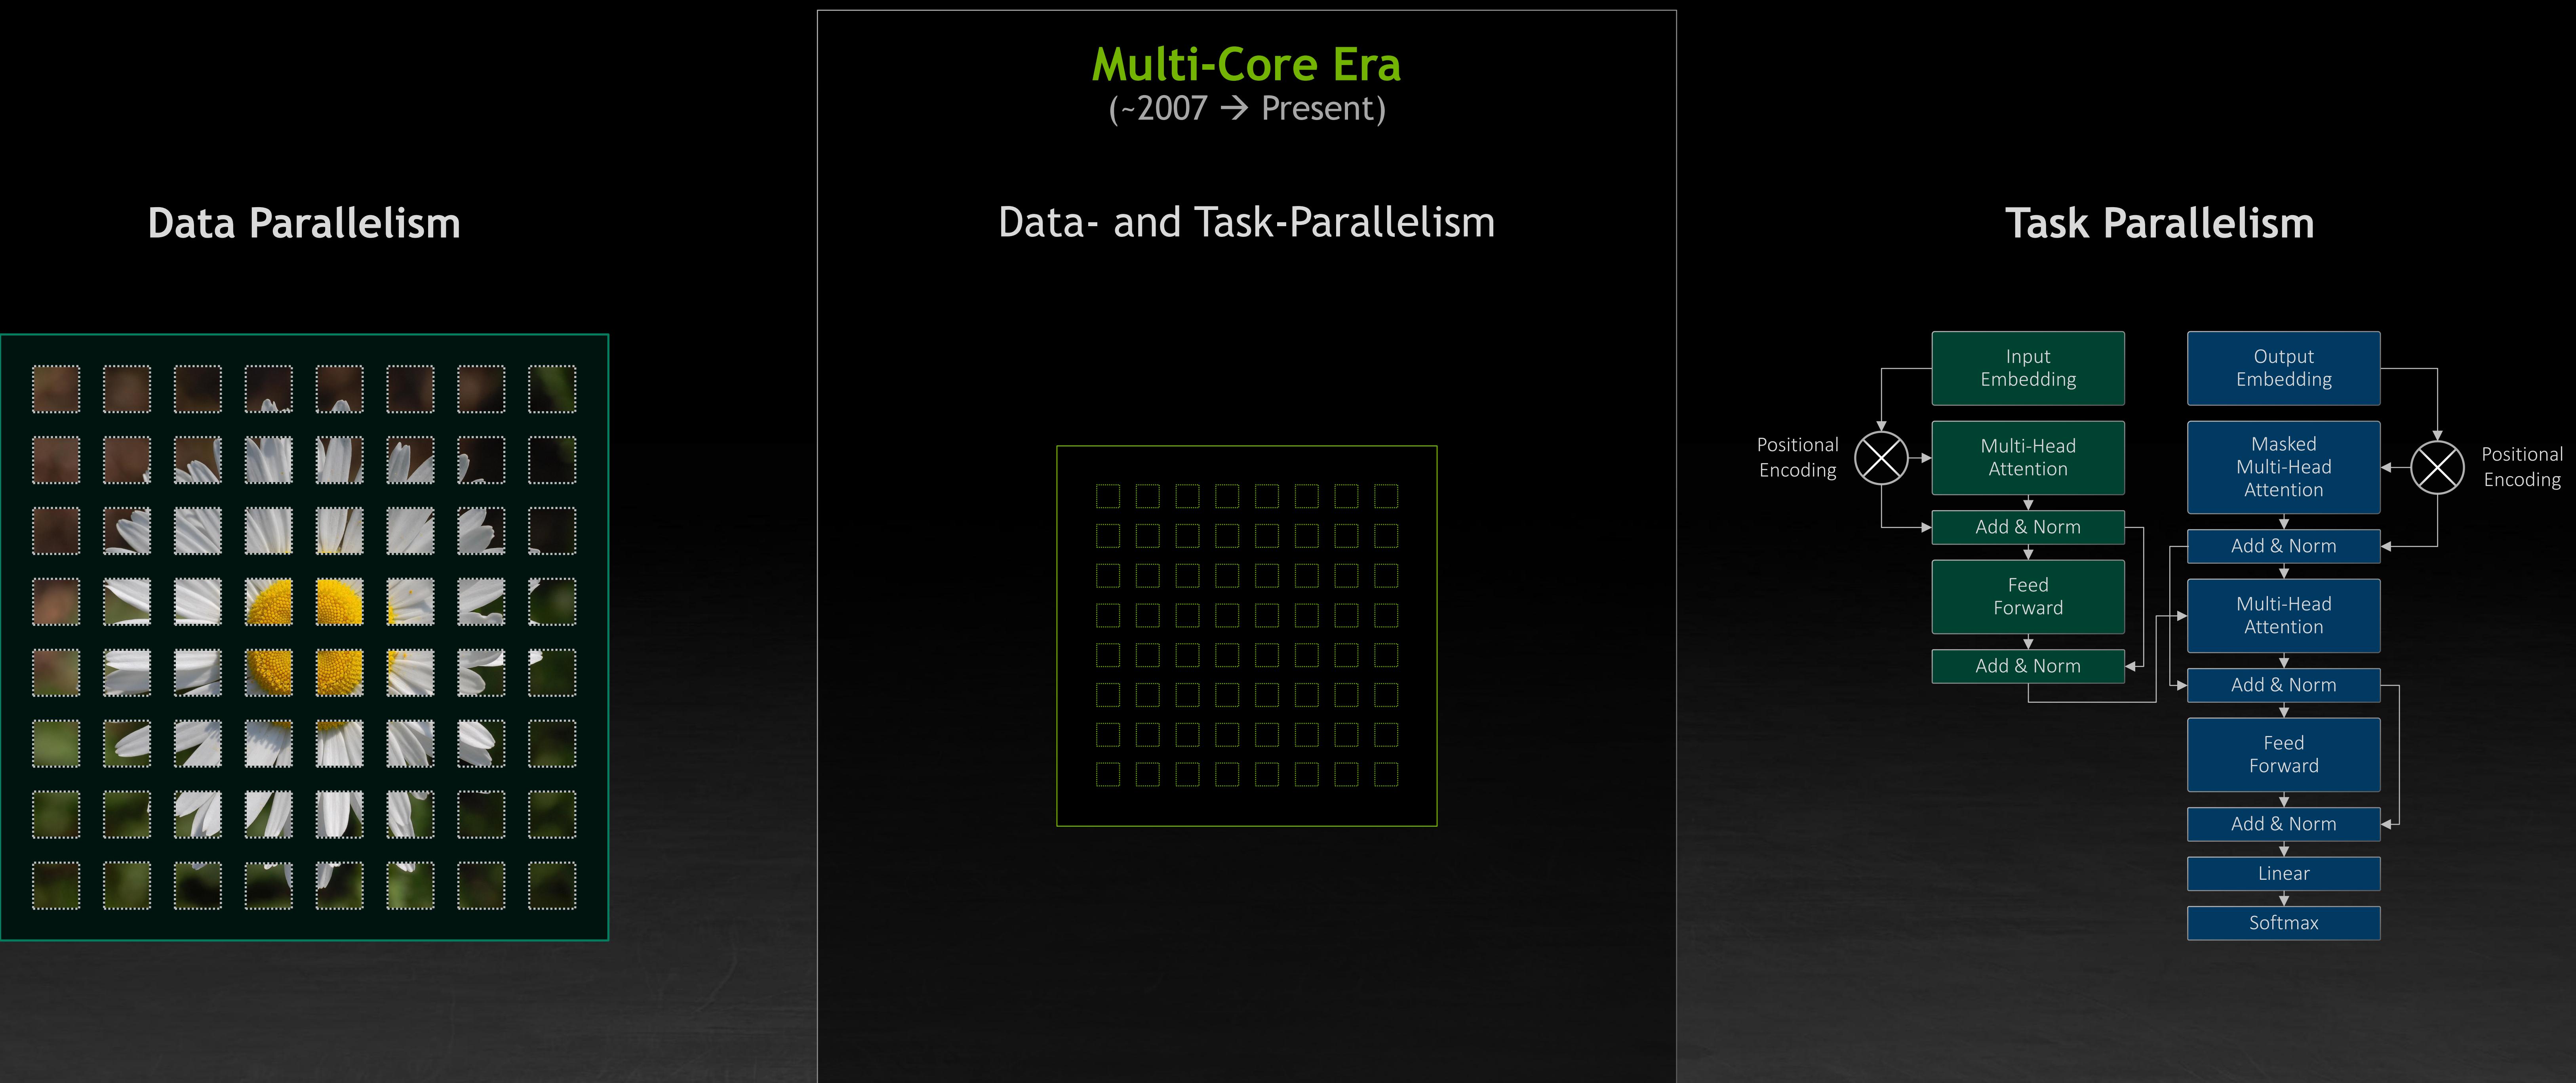
Data Parallelism



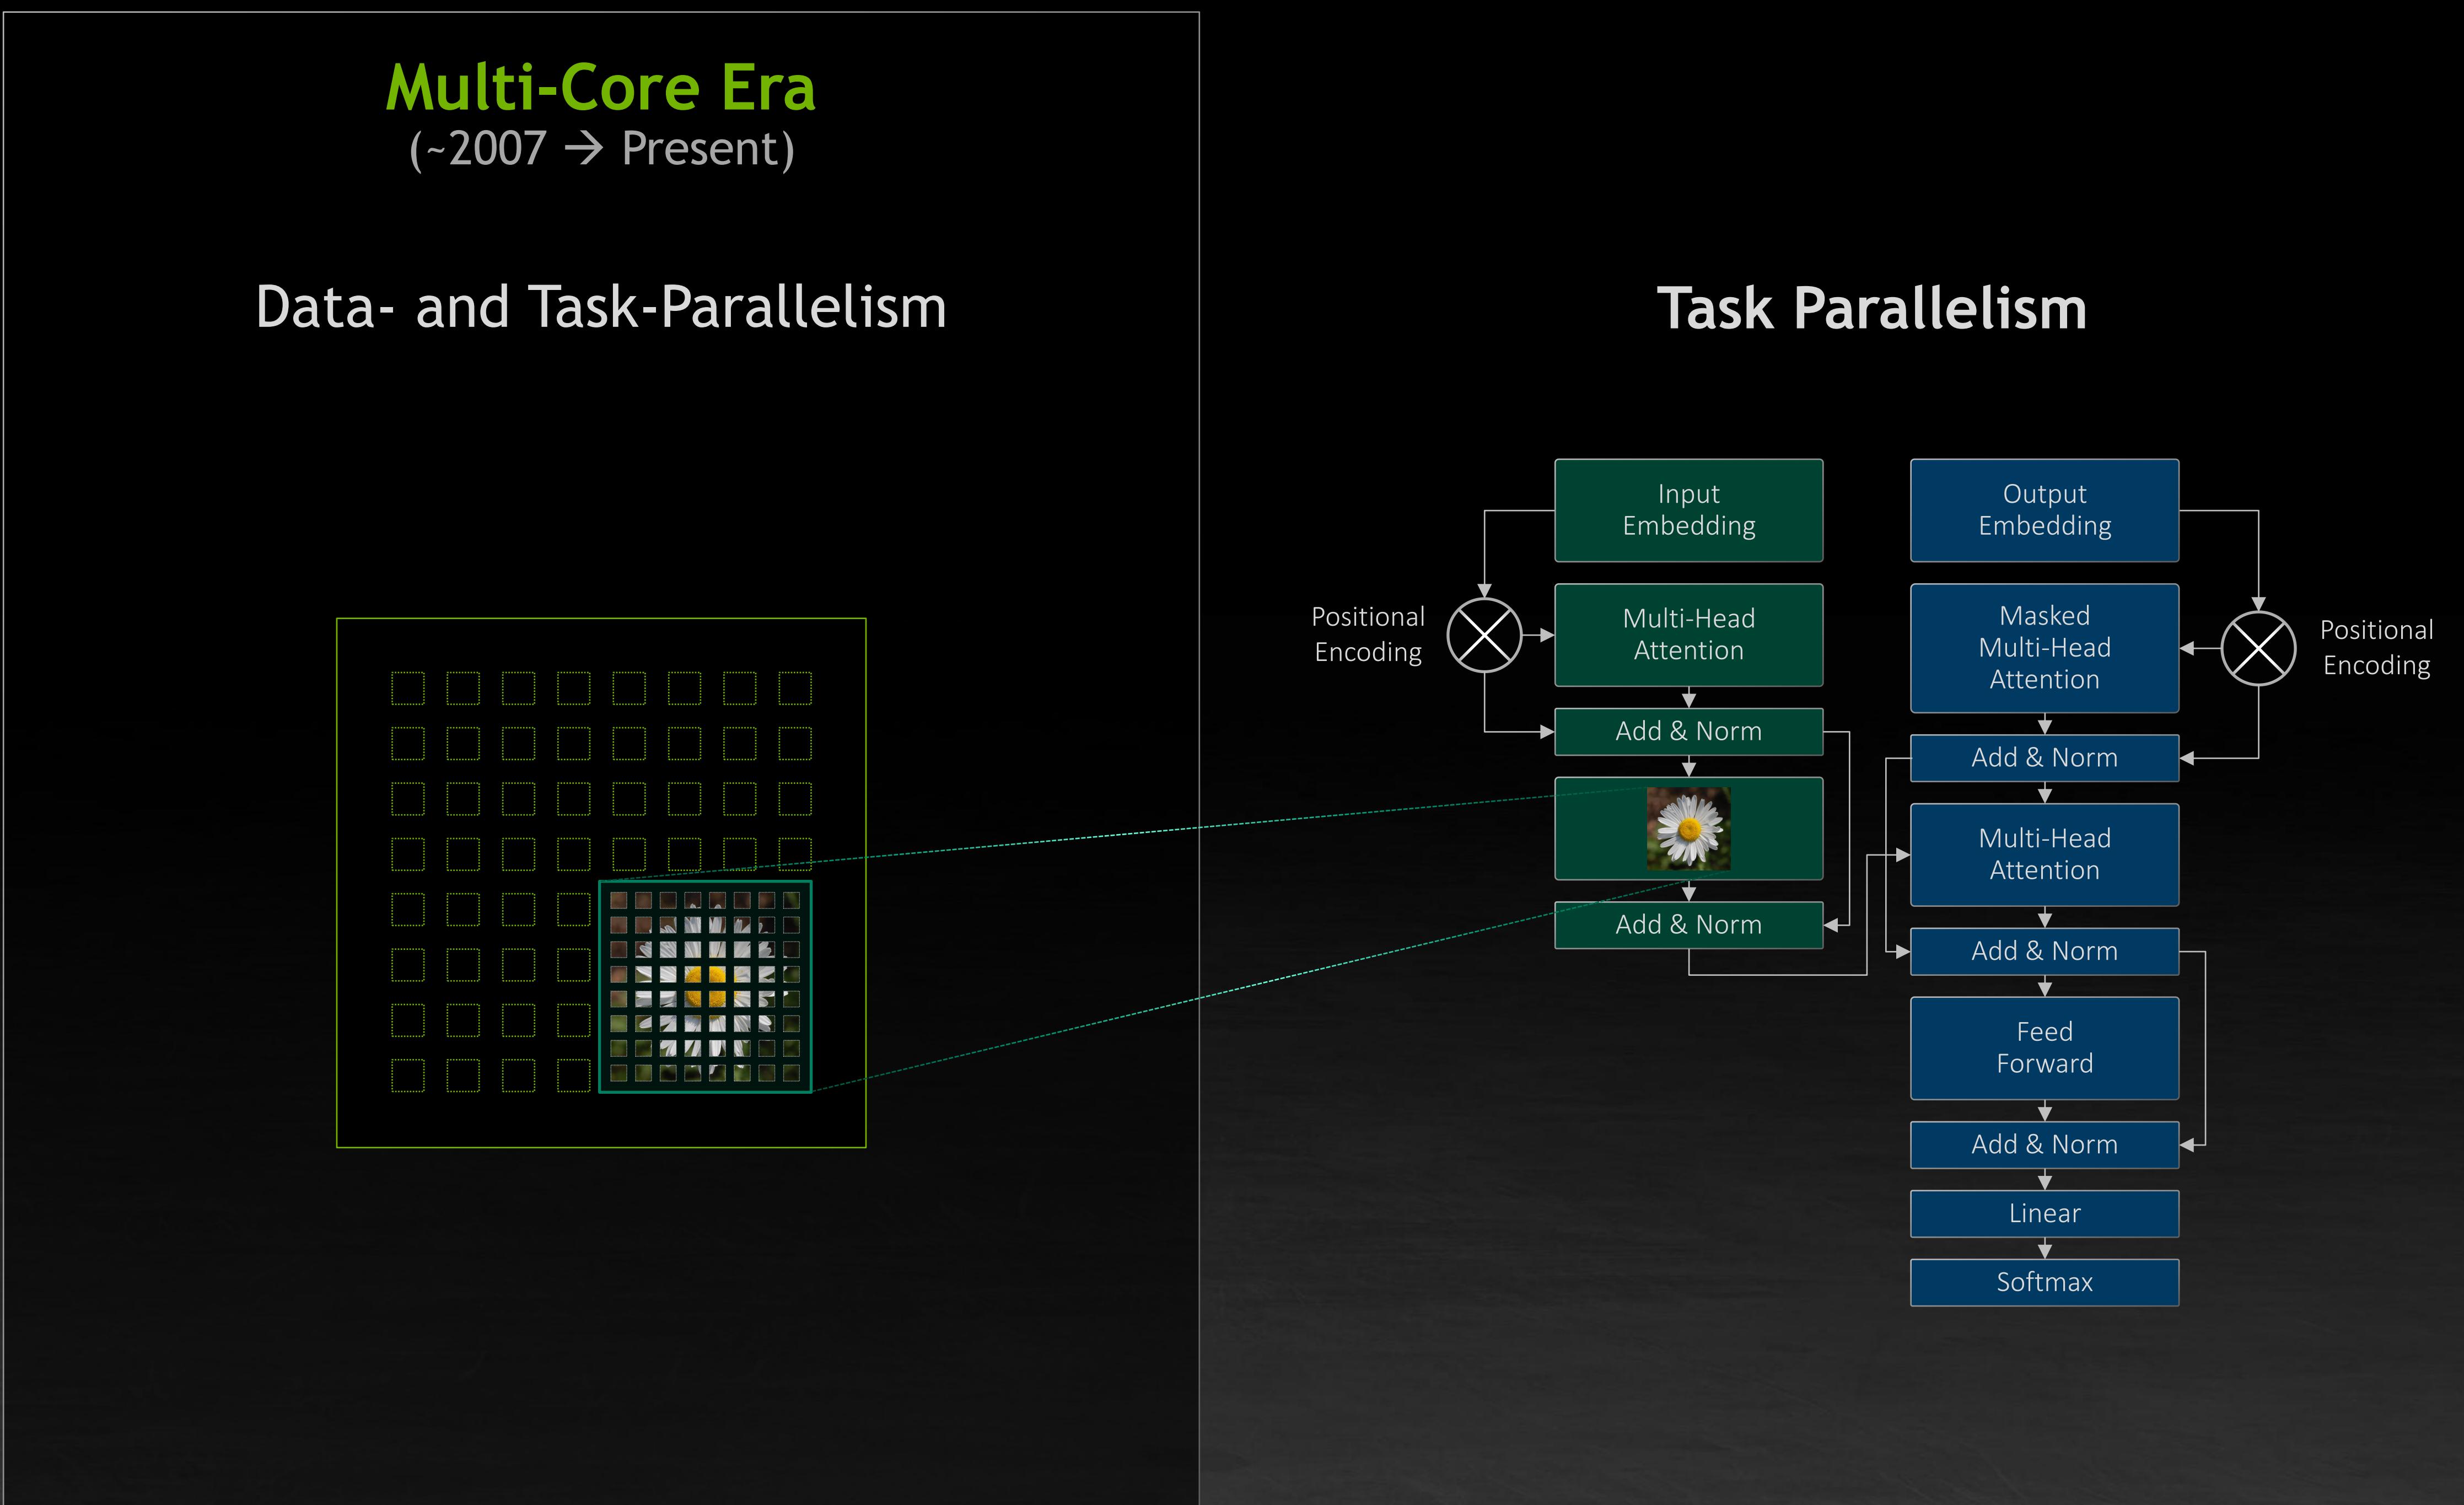
Multi-Core Era (~2007 → Present)

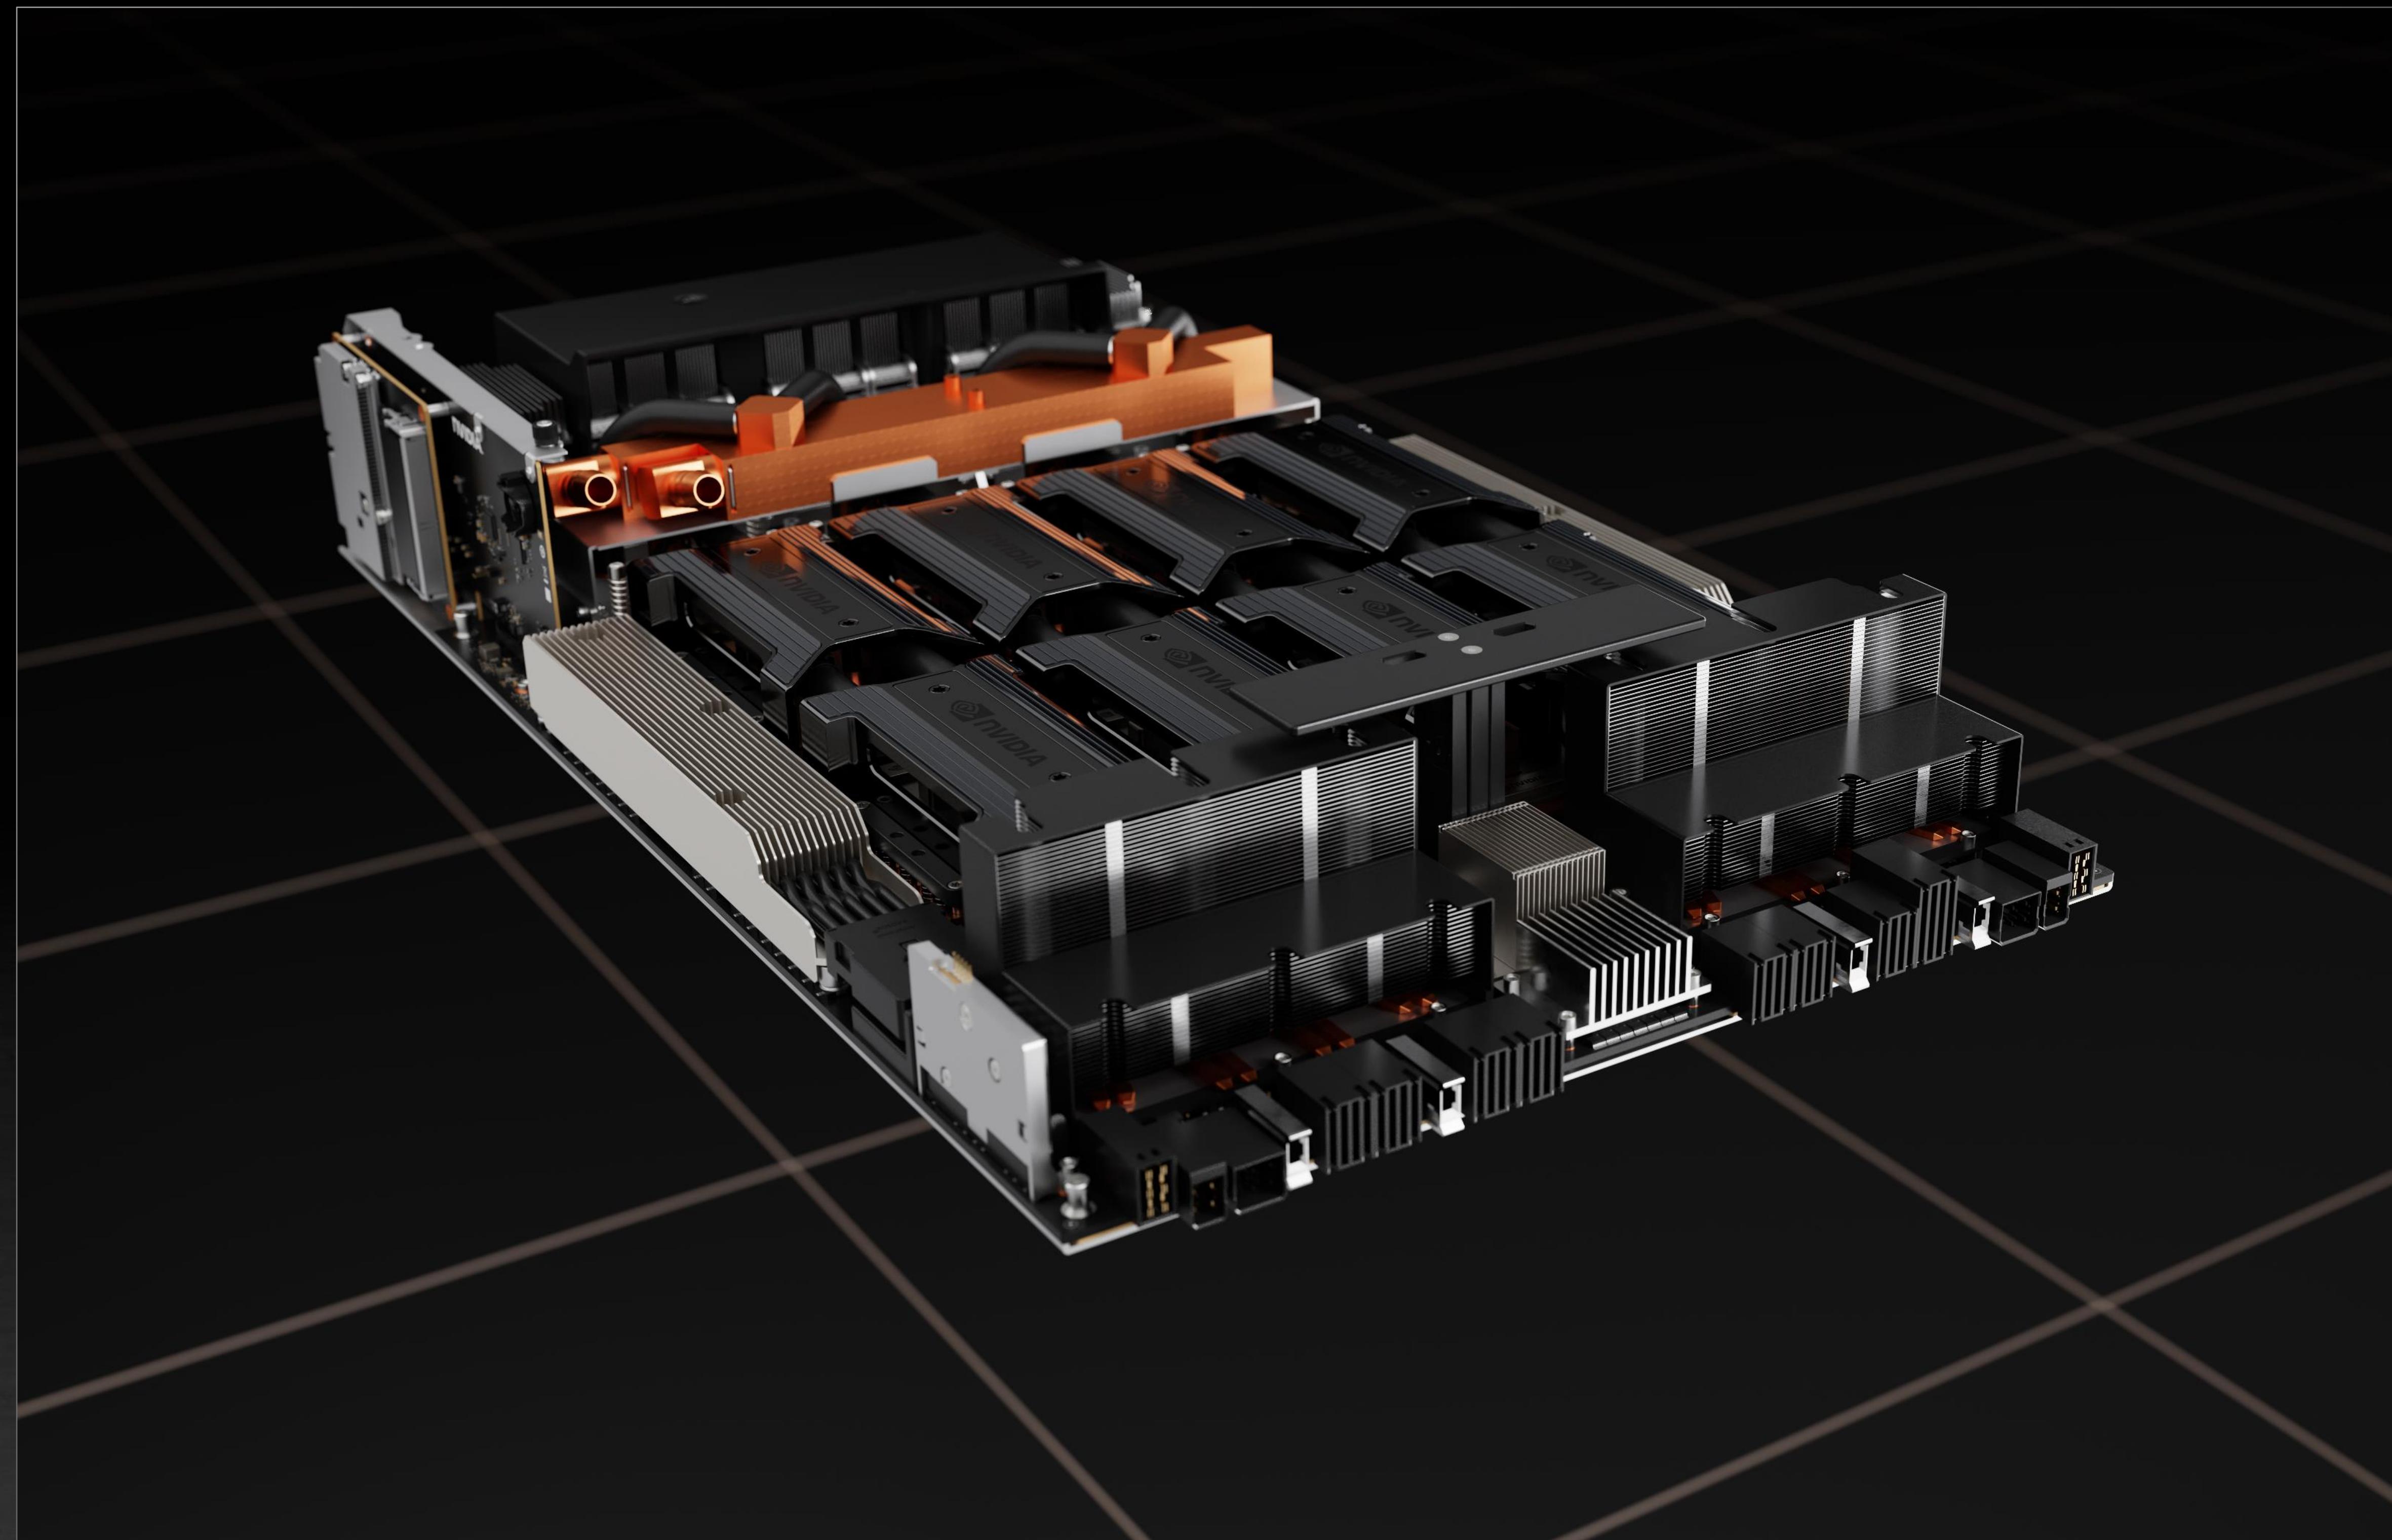


THE SECOND ERA OF SOFTWARE DEVELOPMENT

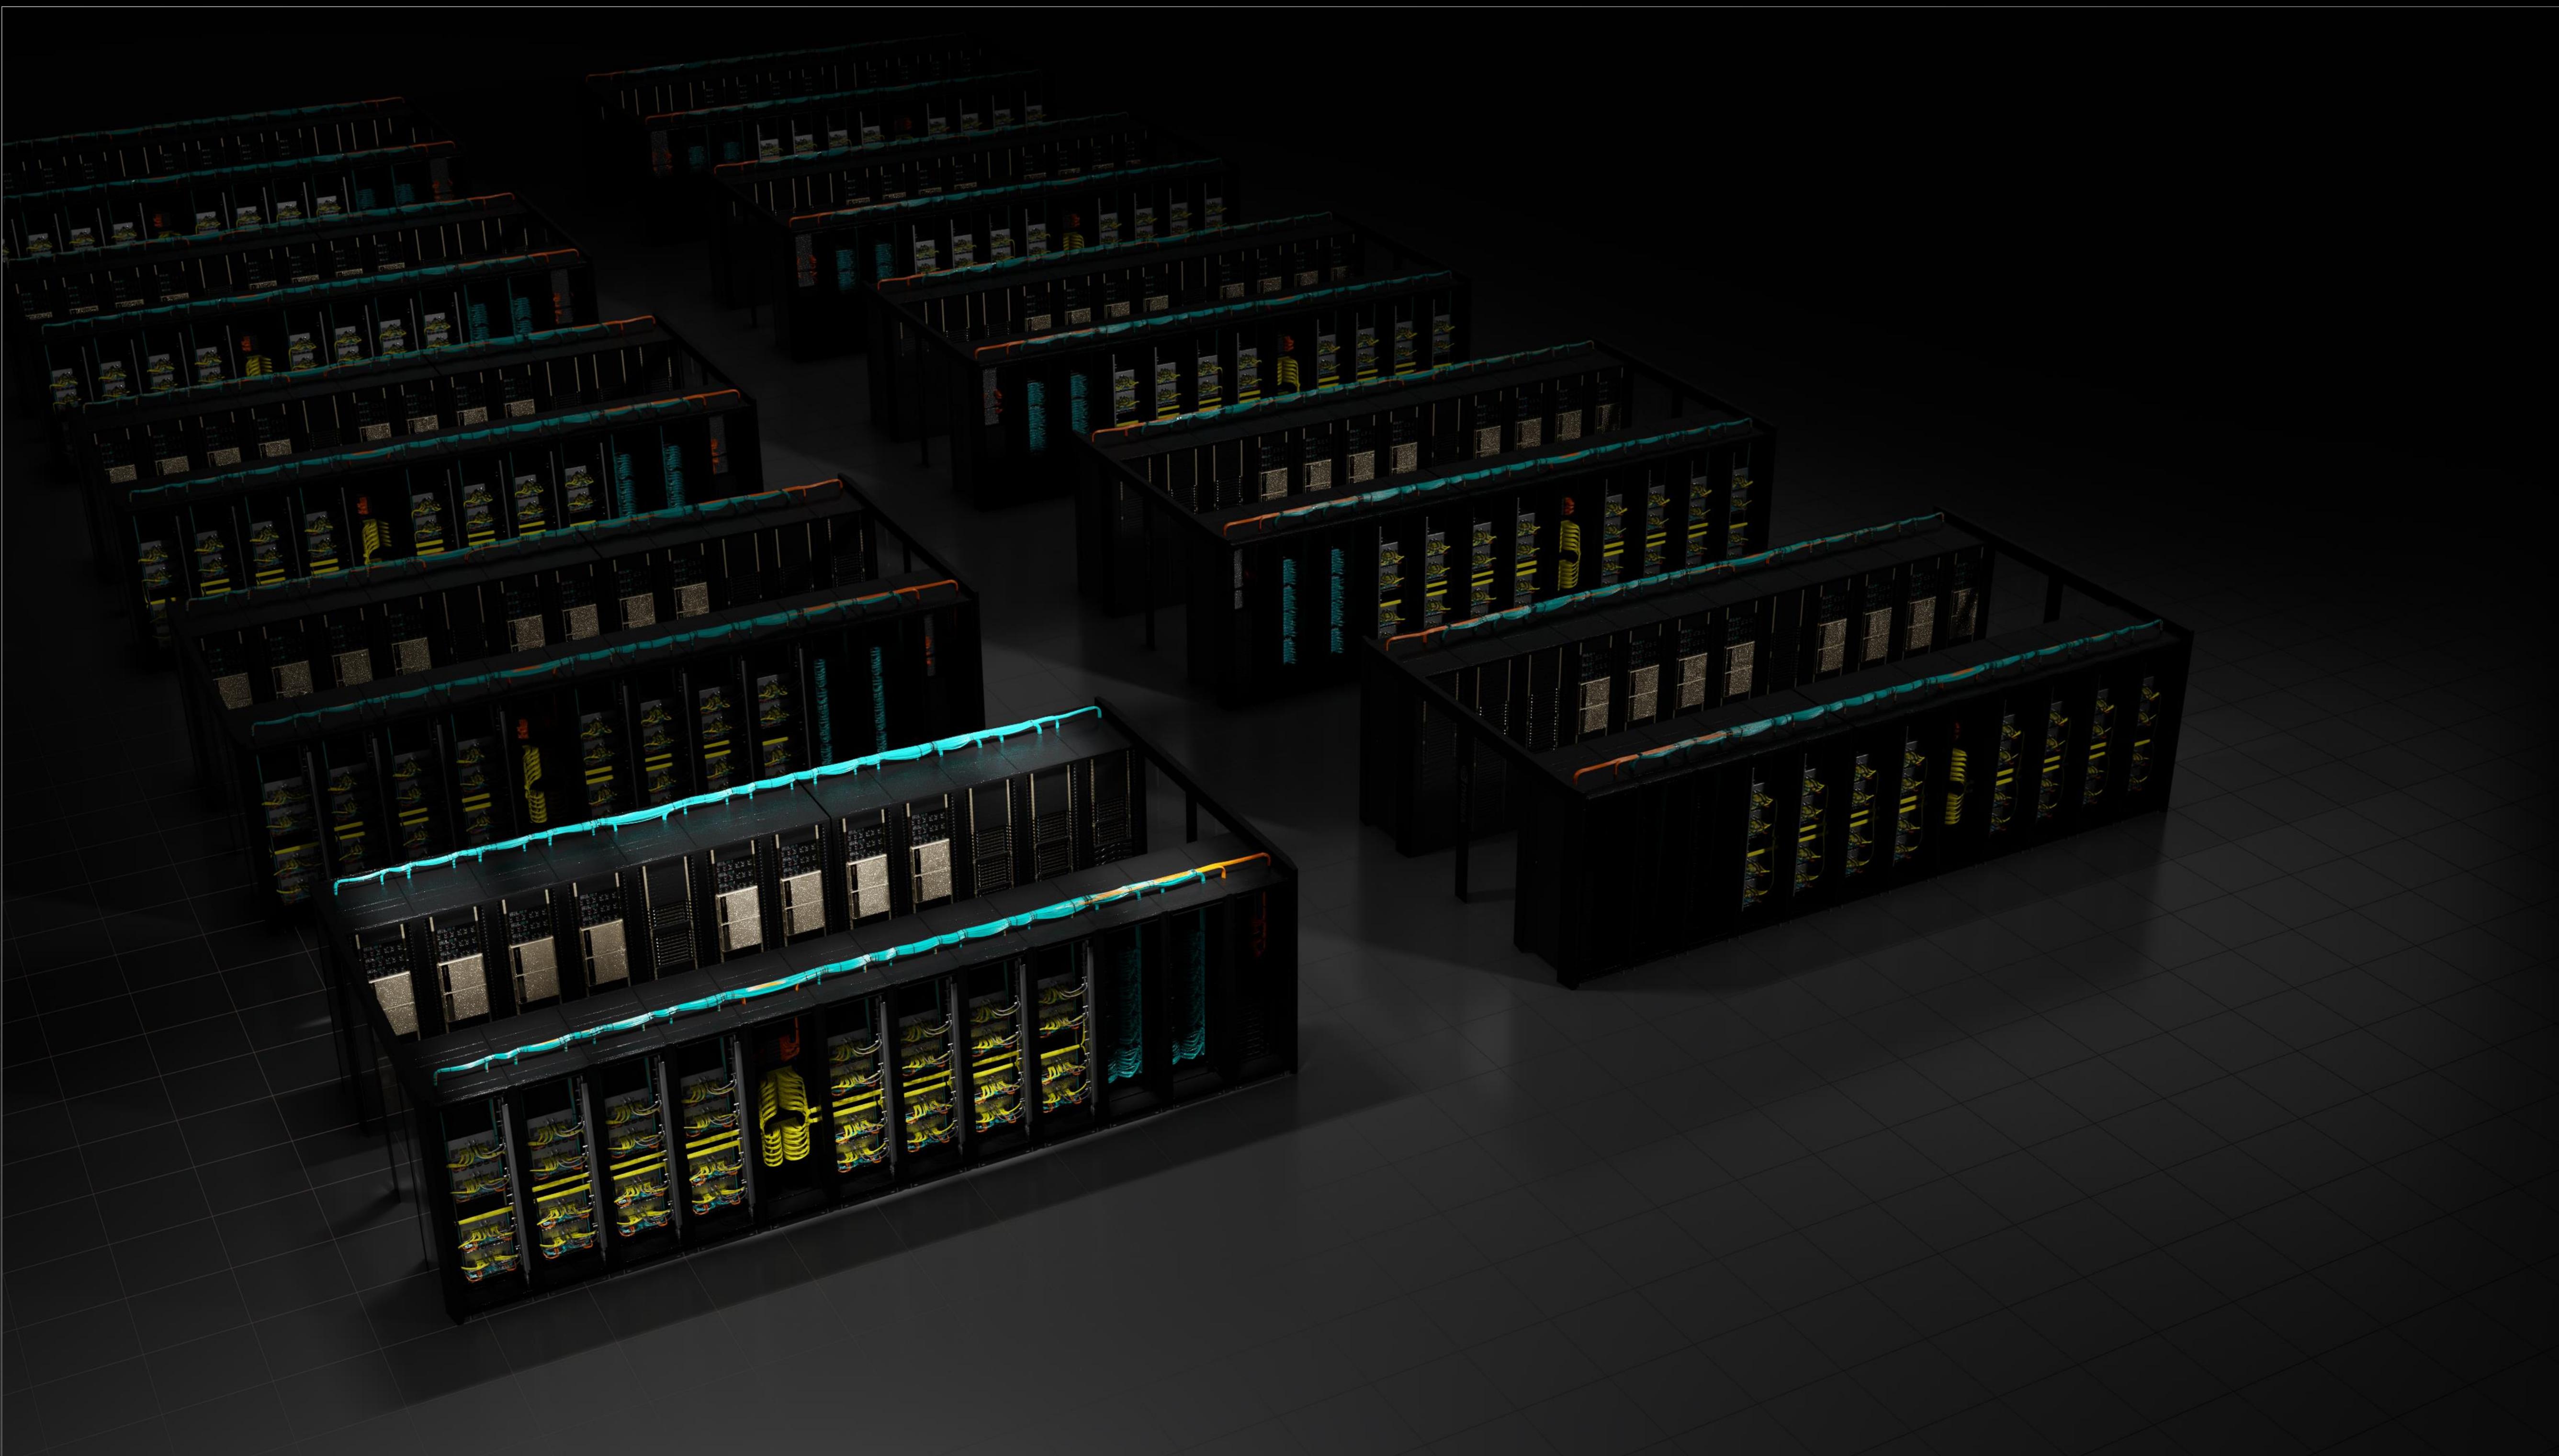


HIERARCHICAL PARALLELISM

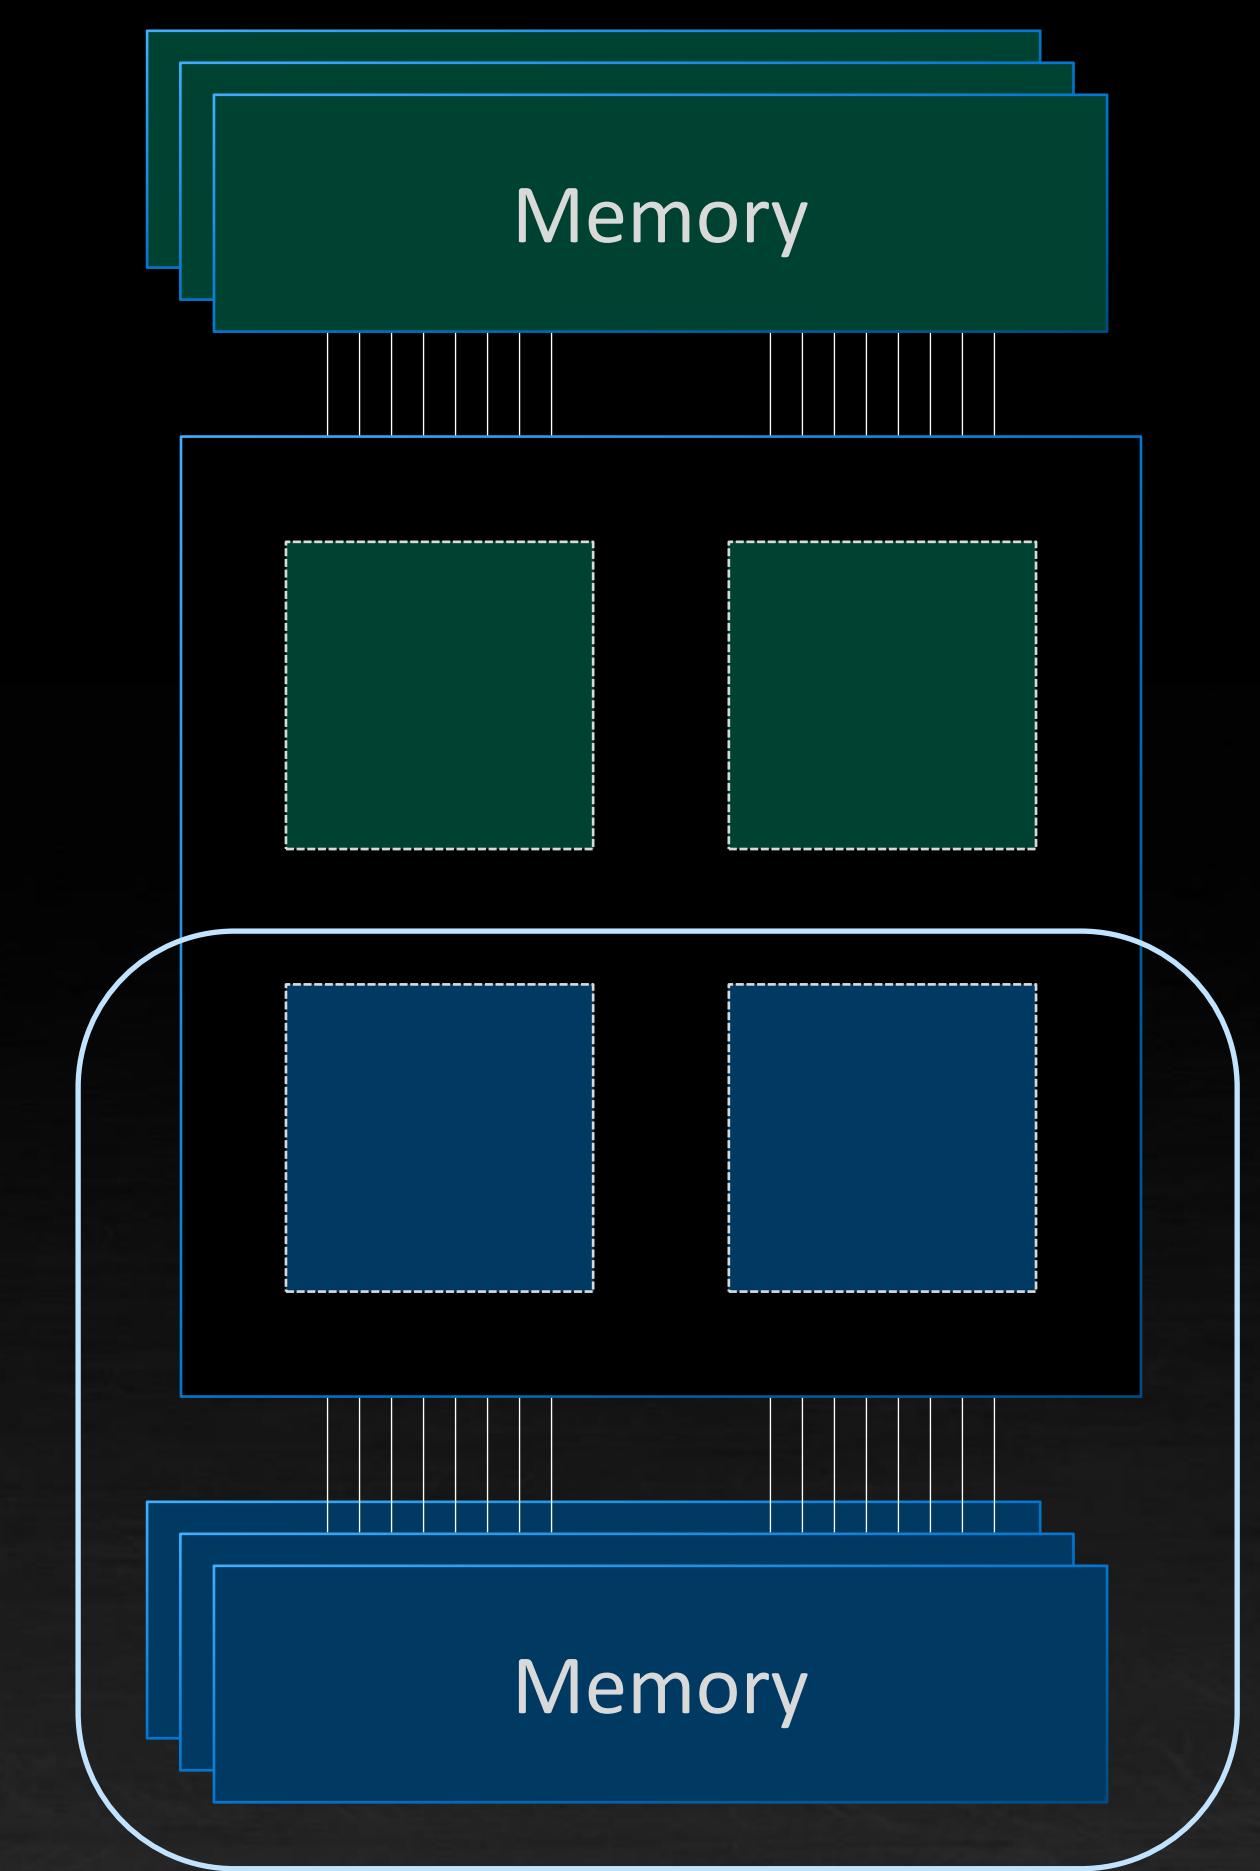




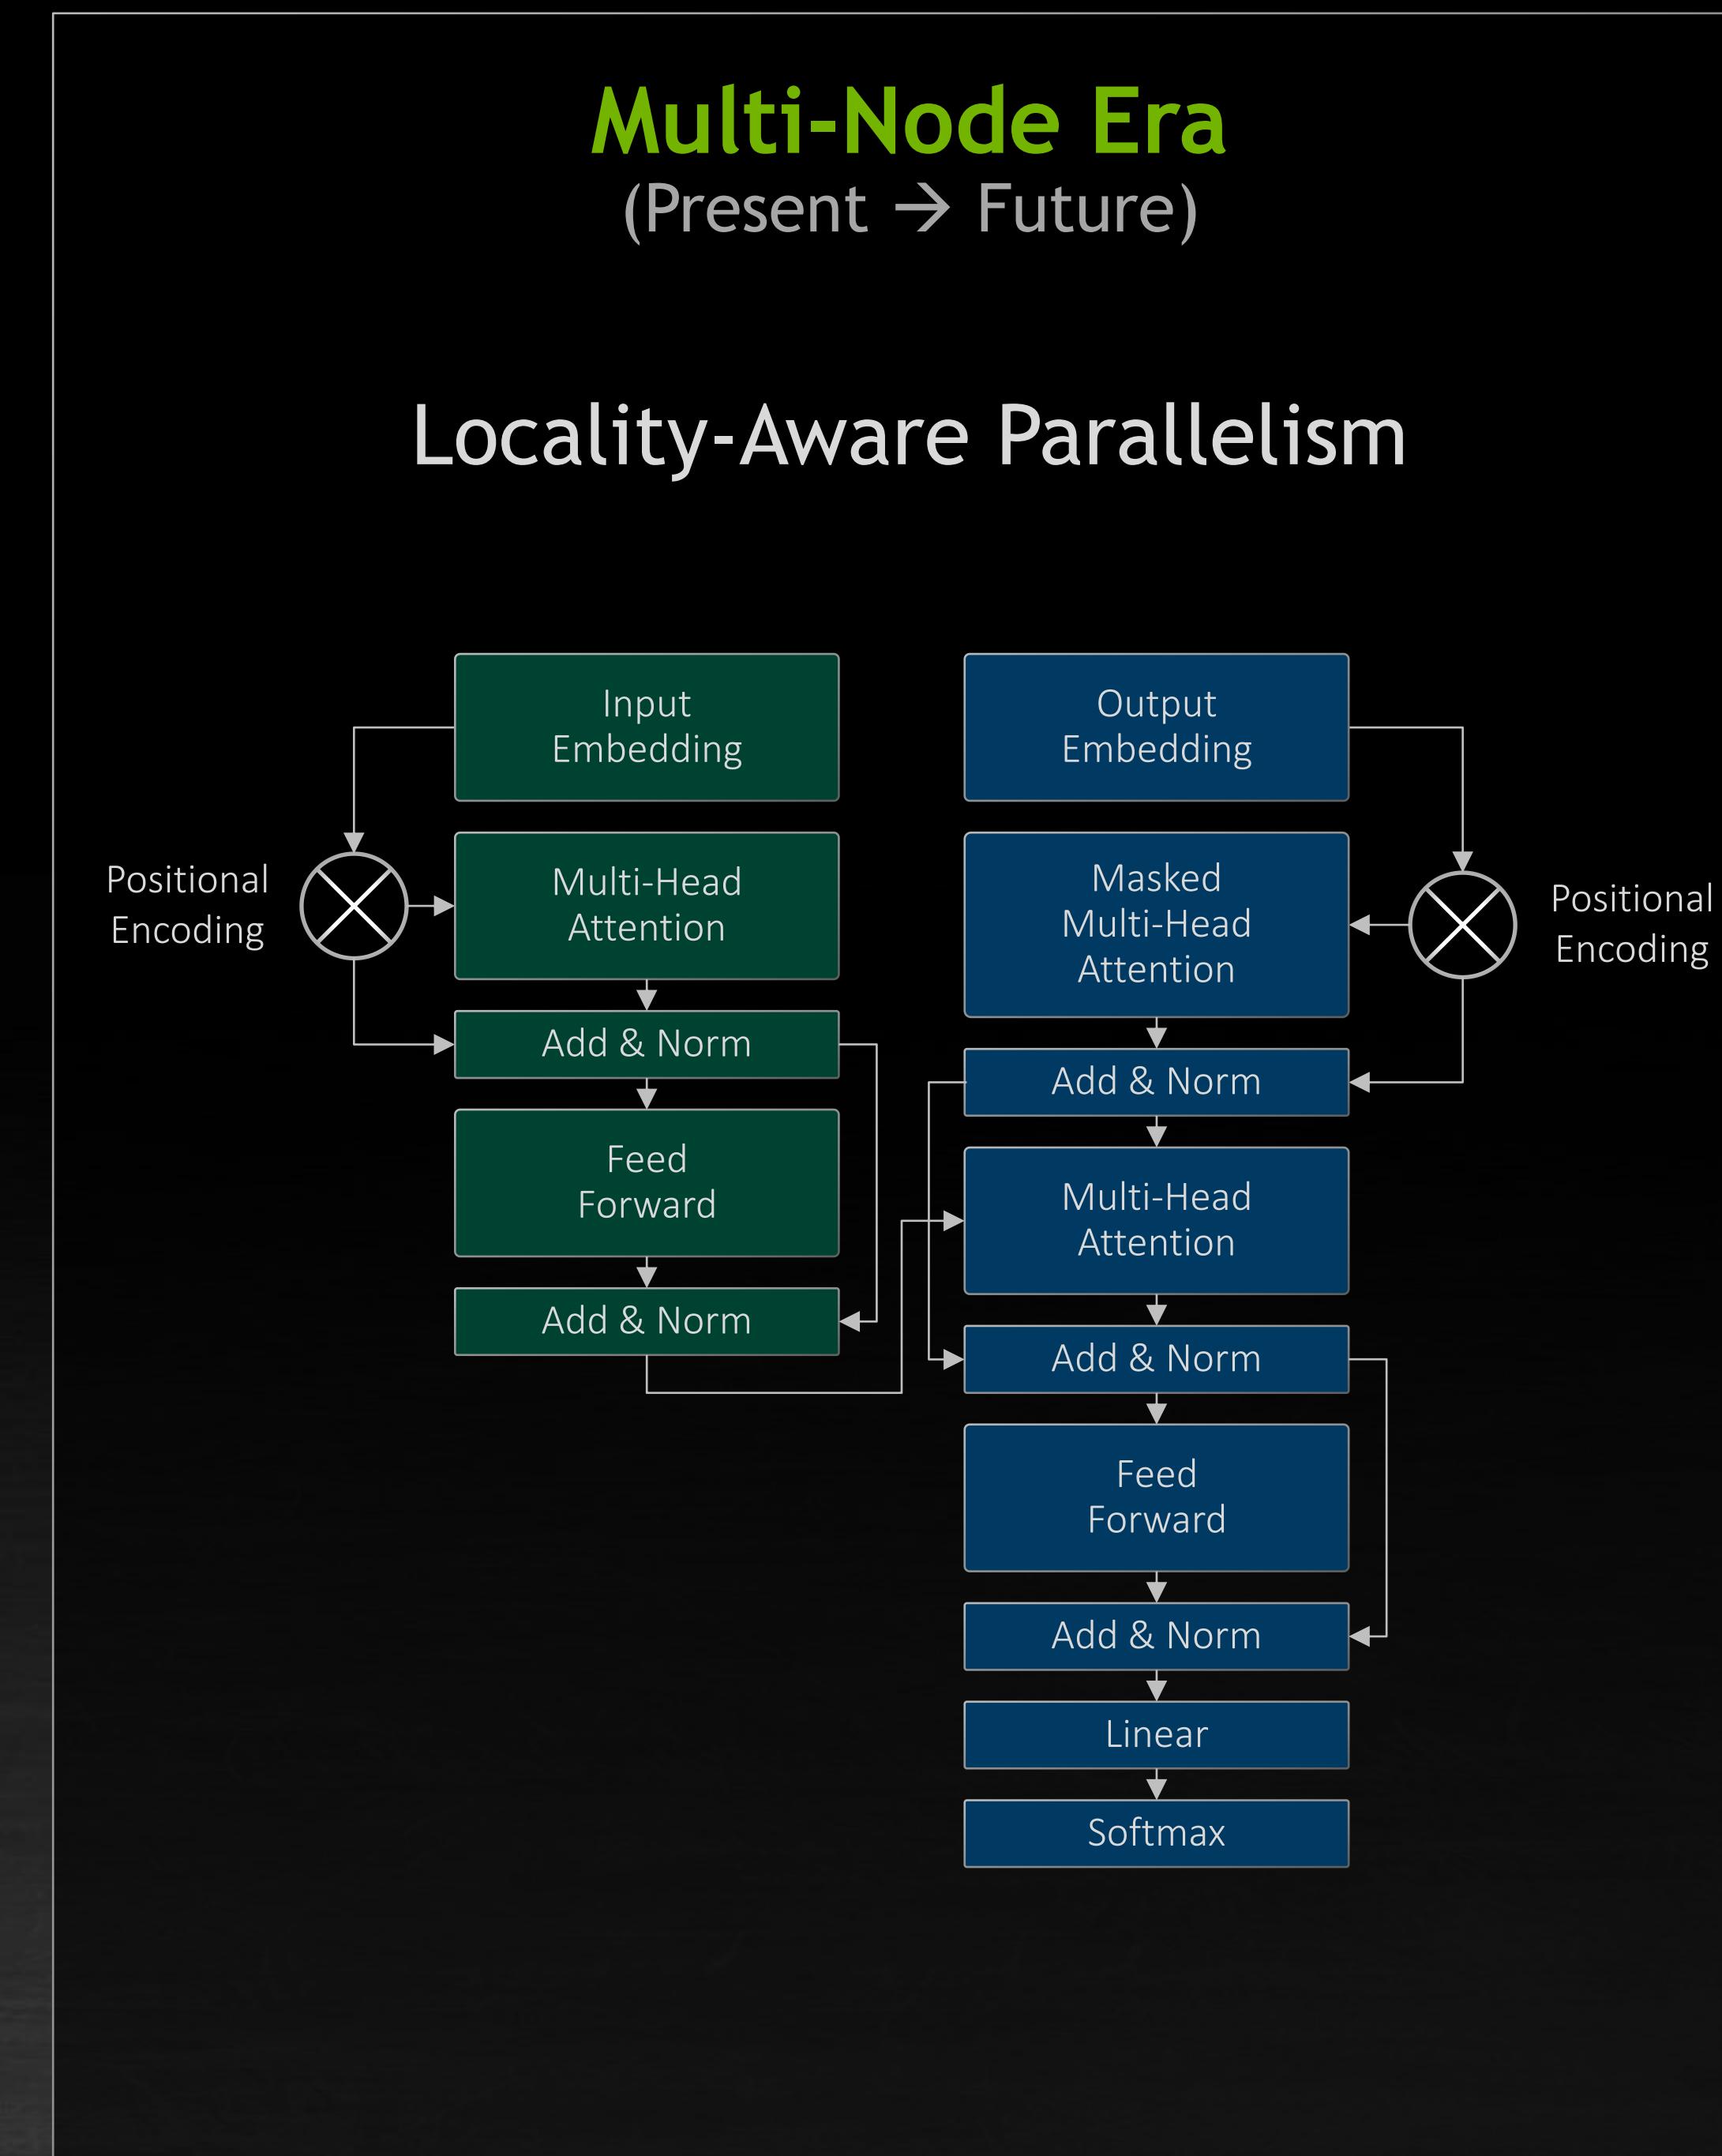
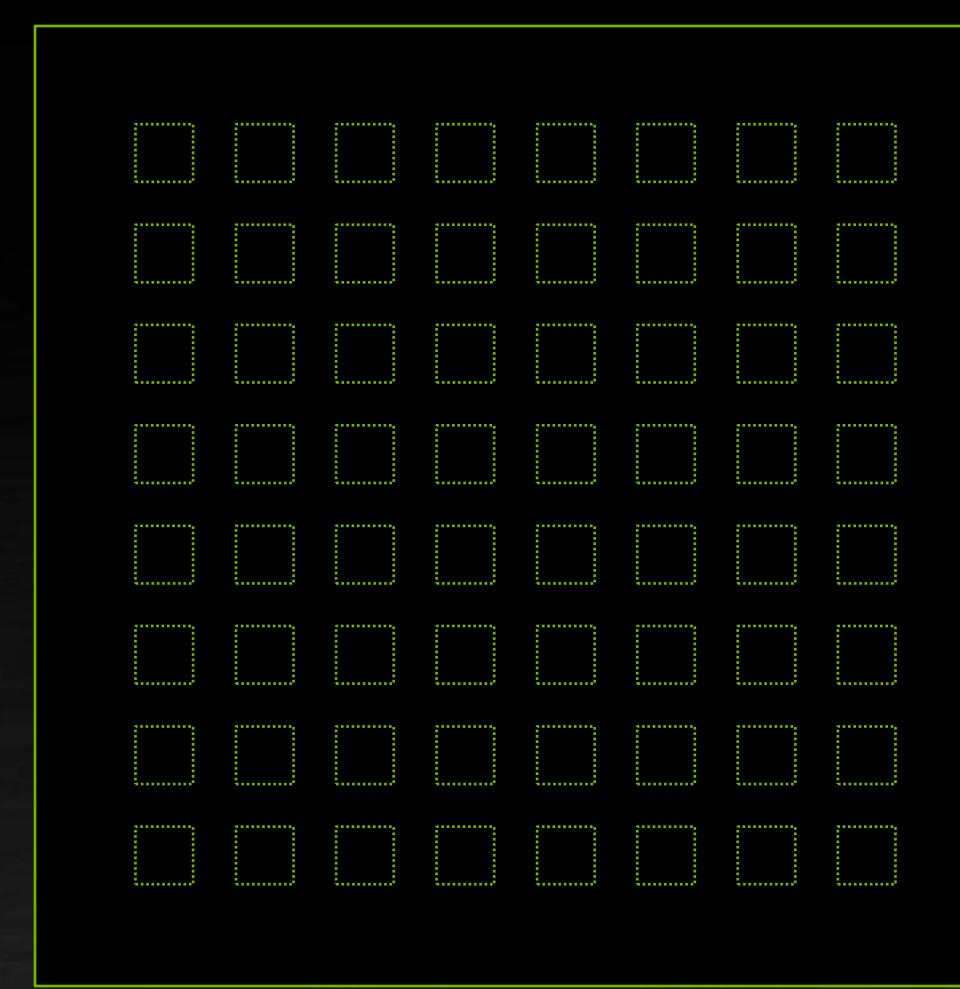
DATACENTER-SCALE COMPUTING



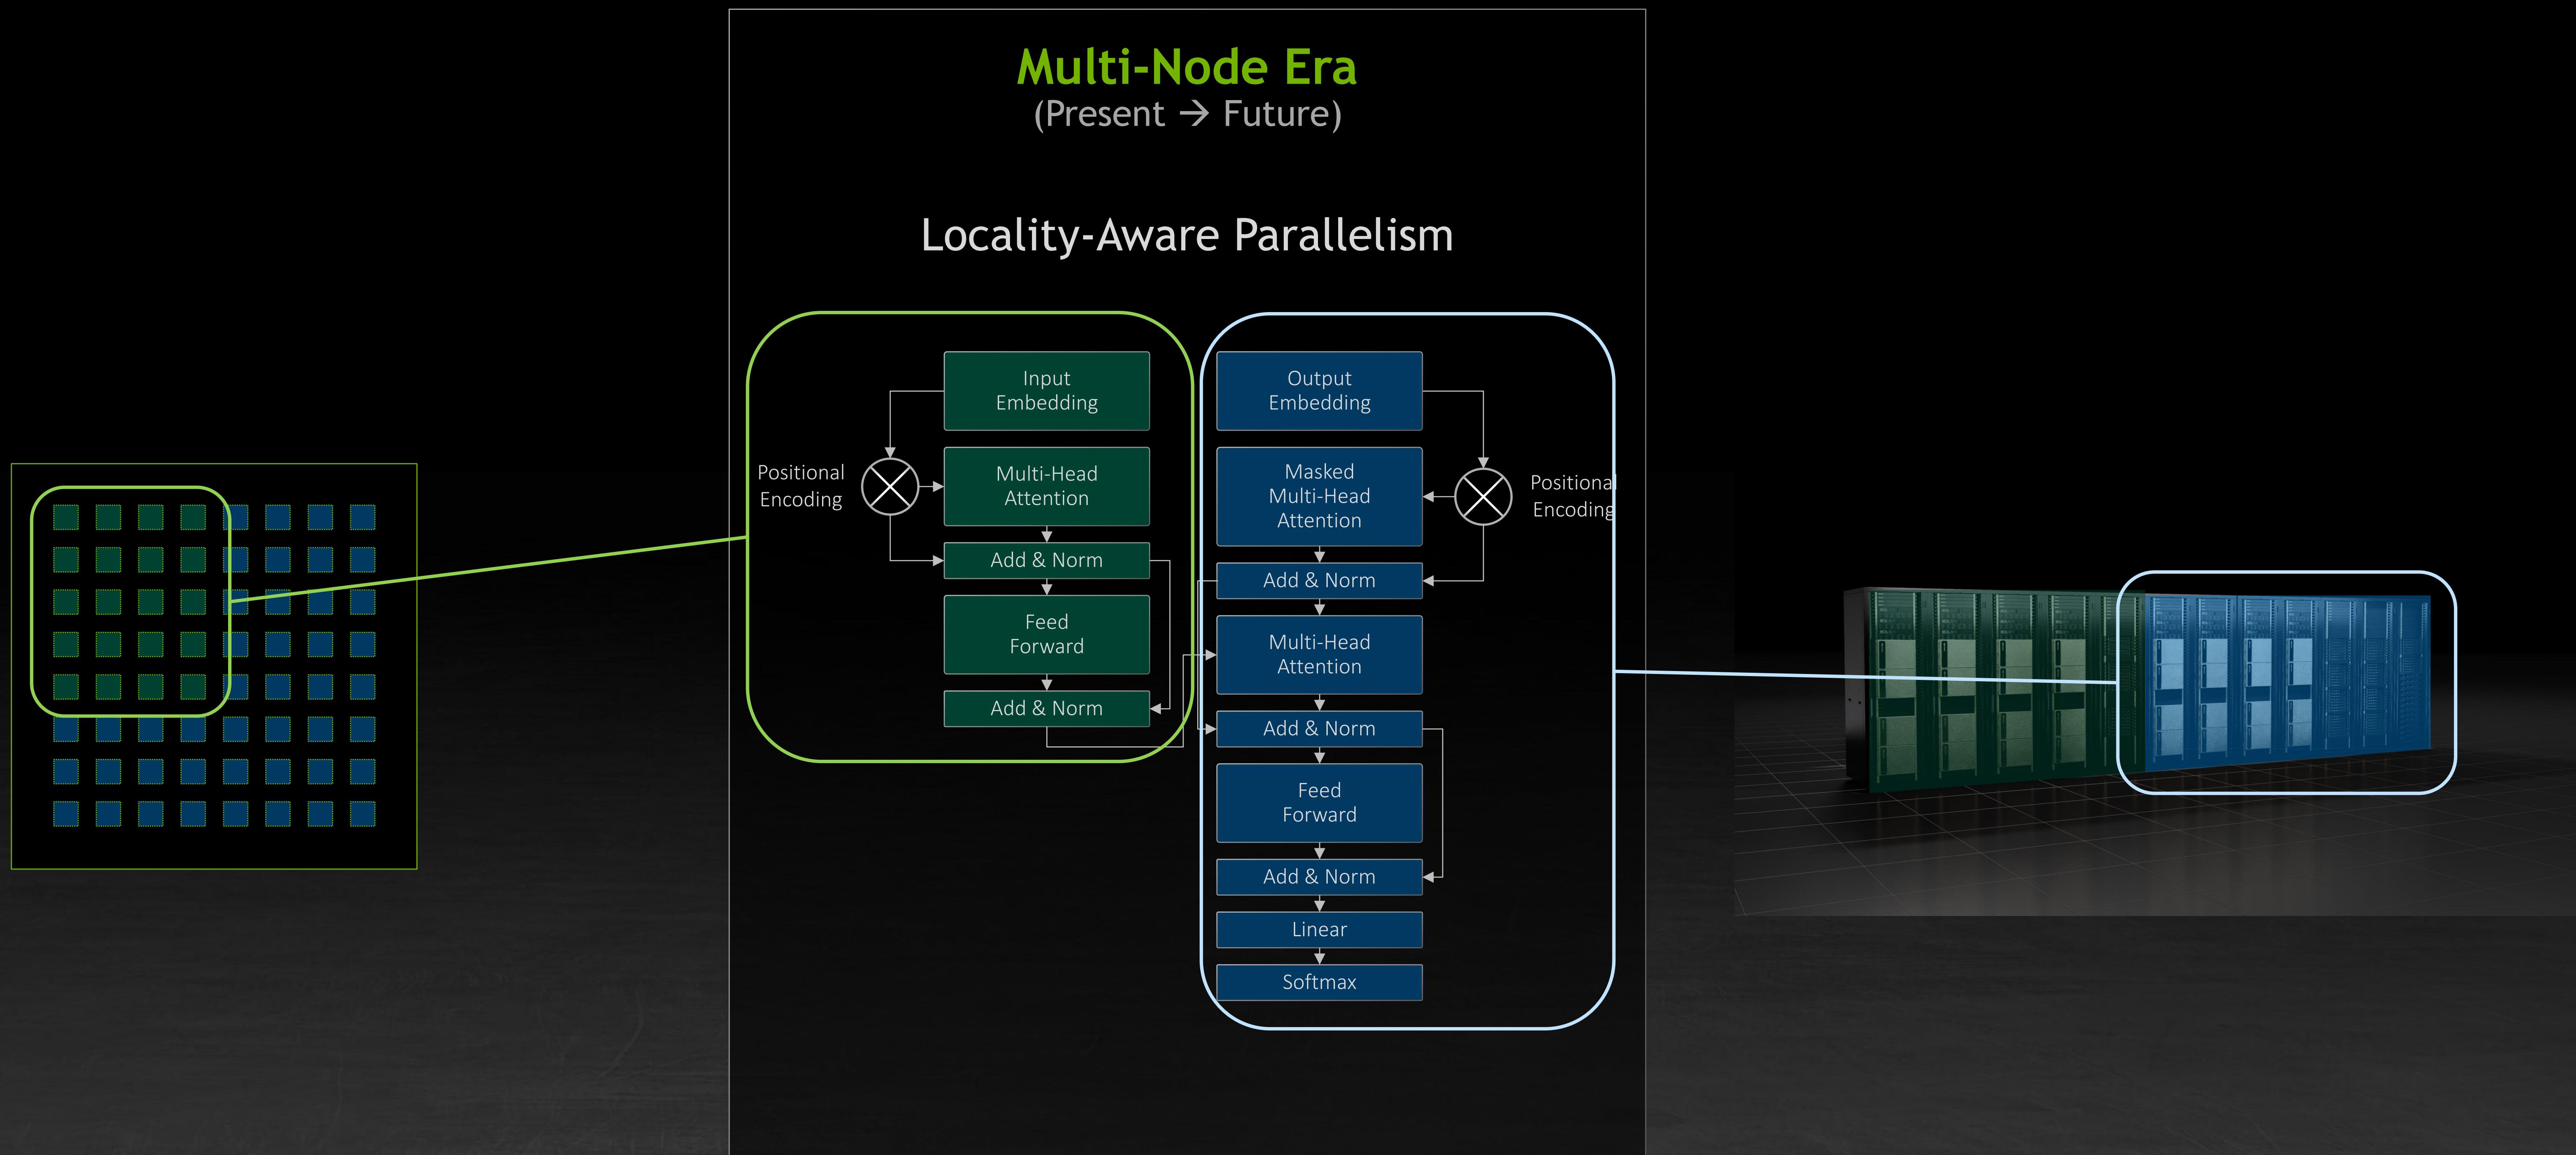
MANAGING LOCALITY IS NOT NEW



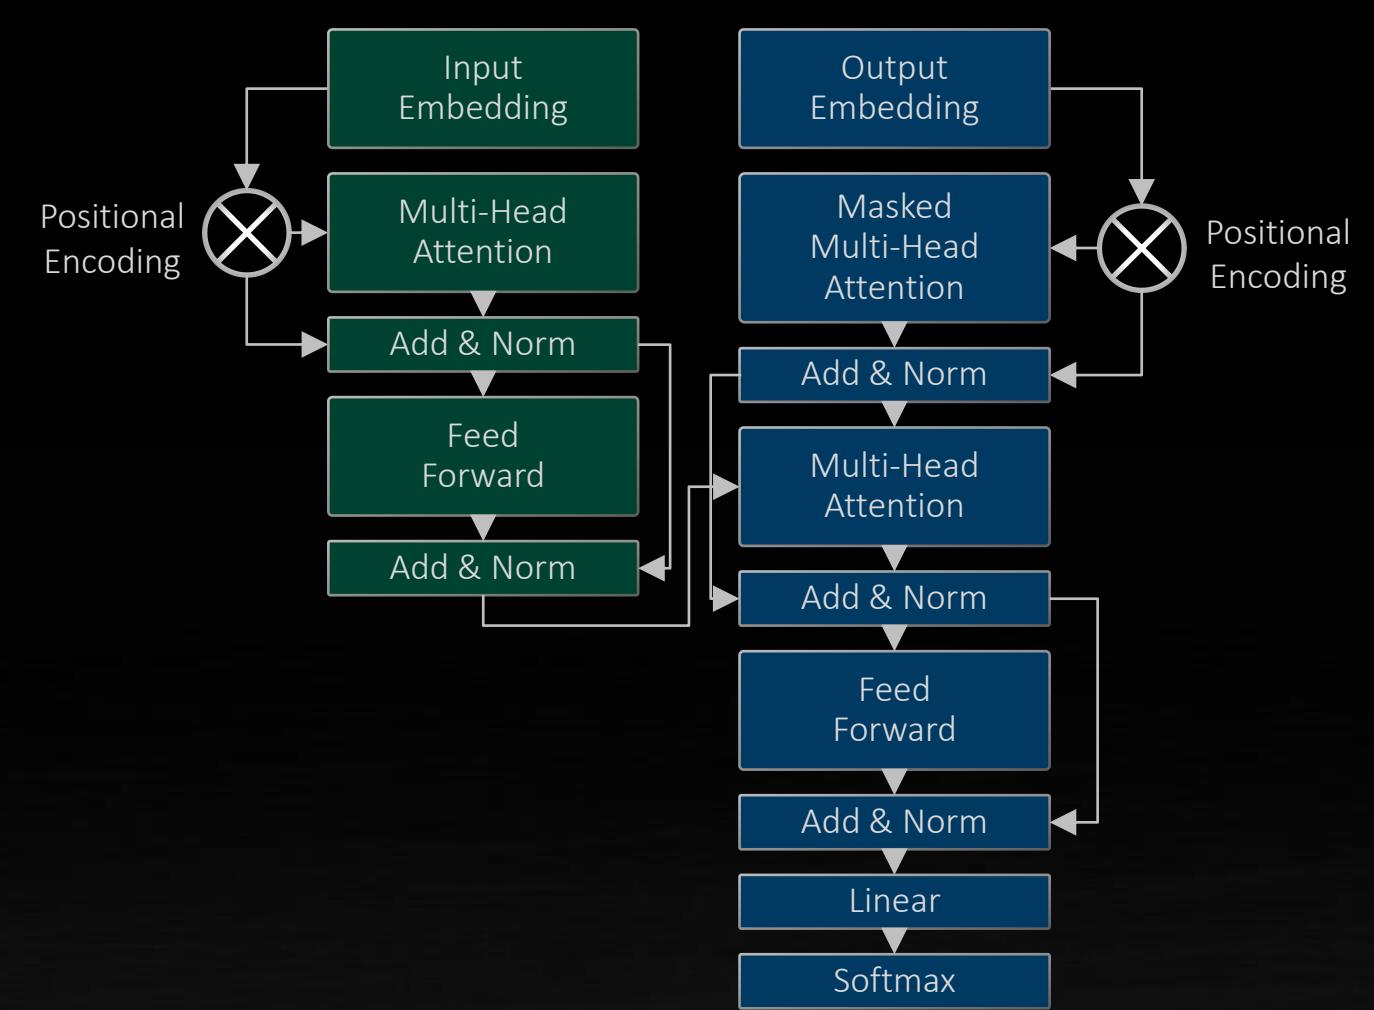
THE THIRD ERA OF SOFTWARE DEVELOPMENT



LOCALITY: THE THIRD ERA OF SOFTWARE DEVELOPMENT

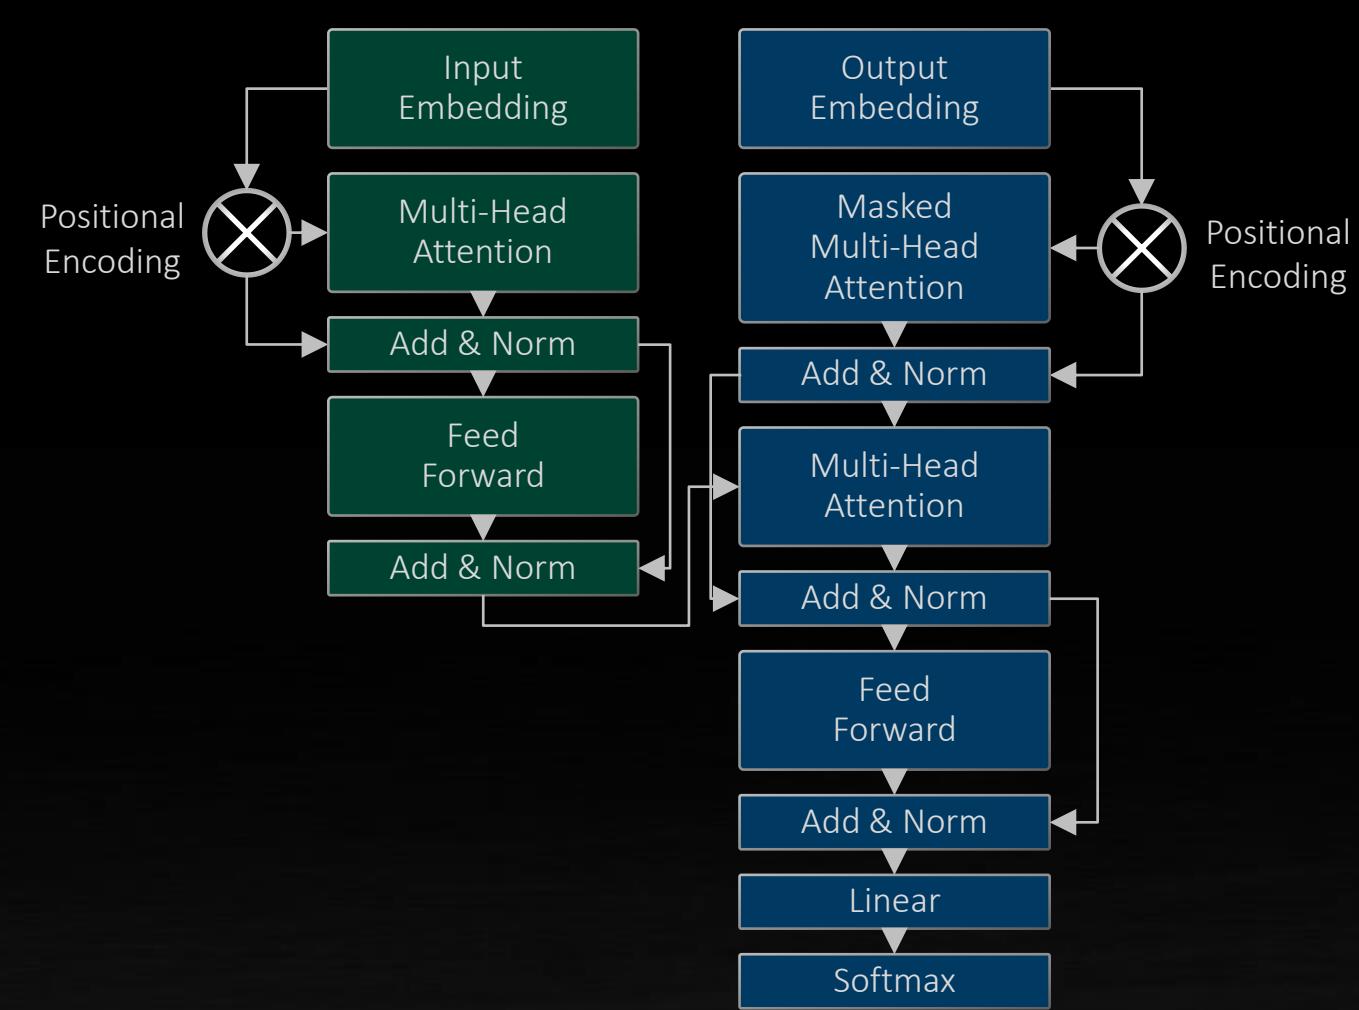


PROGRAMMING TO THE HIERARCHY



Application level

PROGRAMMING TO THE HIERARCHY

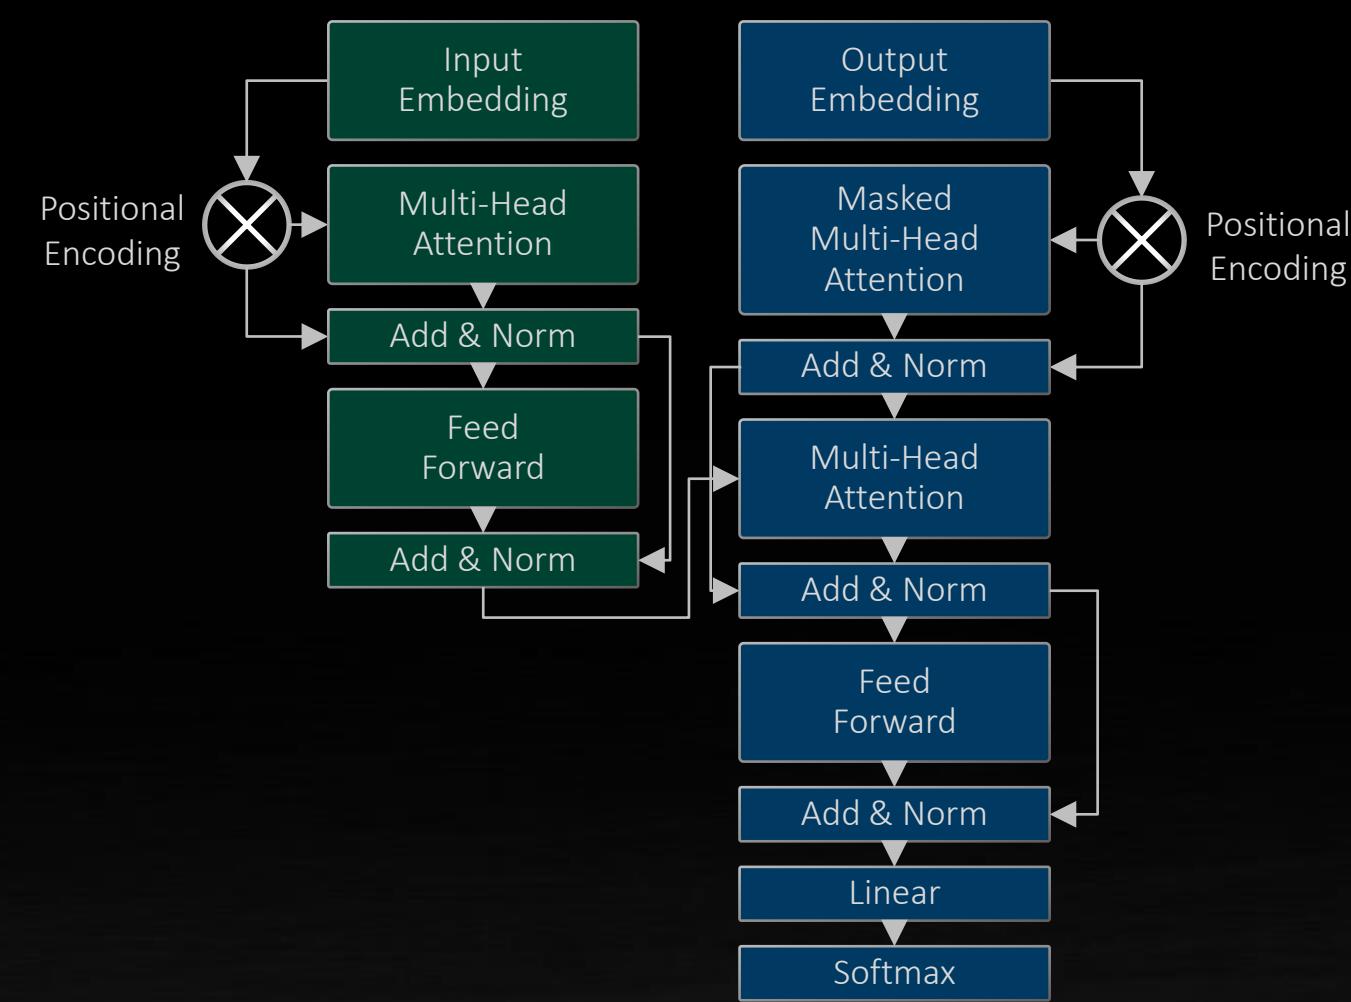


Application level



Framework level

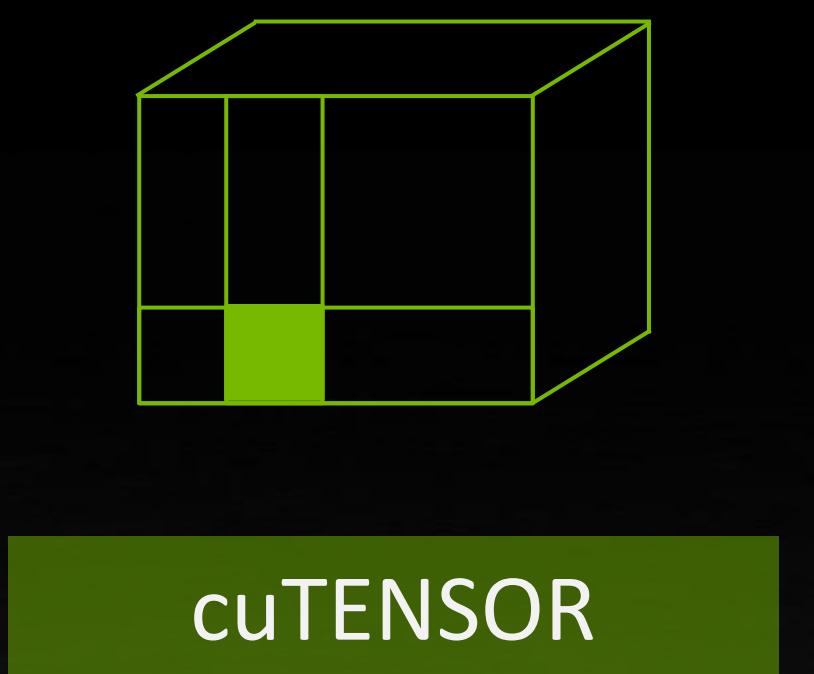
PROGRAMMING TO THE HIERARCHY



Application level

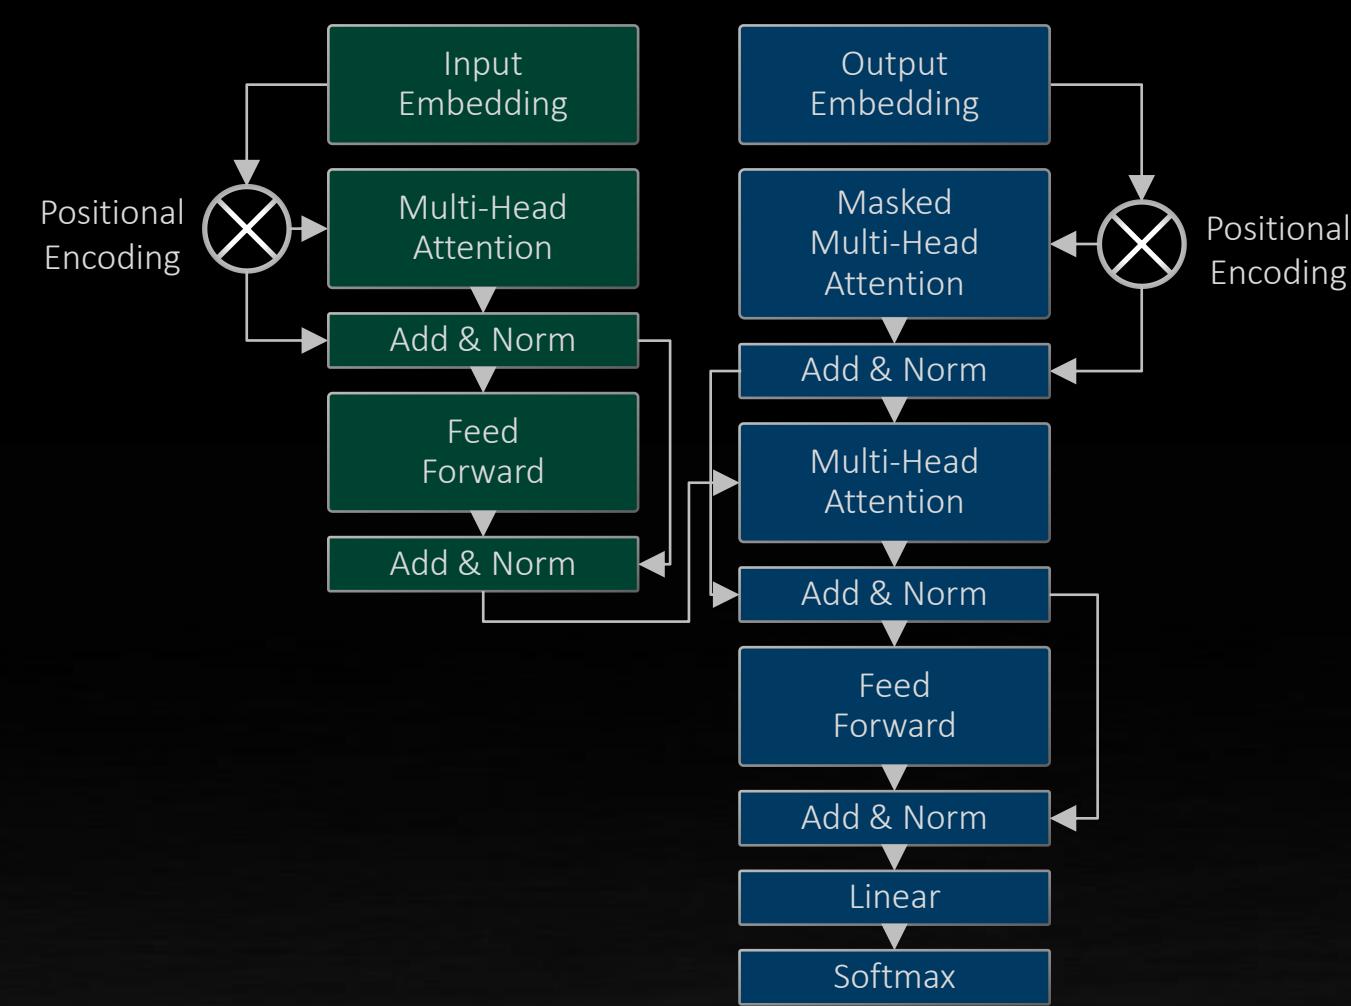


Framework level



Library level

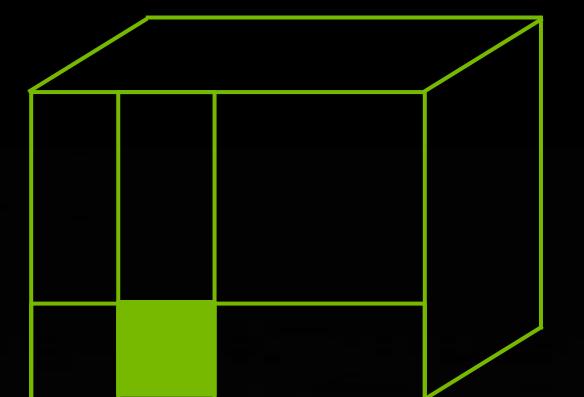
PROGRAMMING TO THE HIERARCHY



Application level

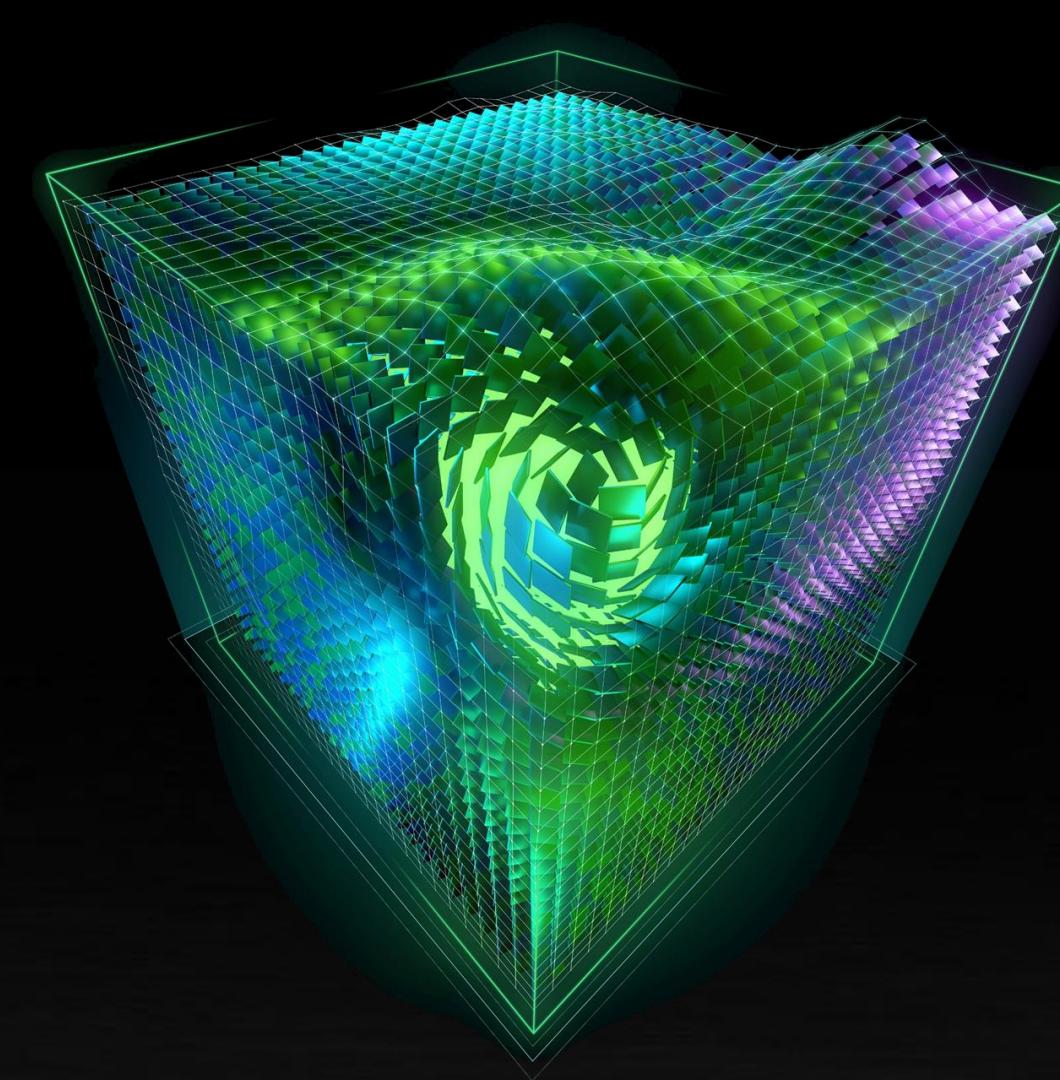


Framework level



cuTENSOR

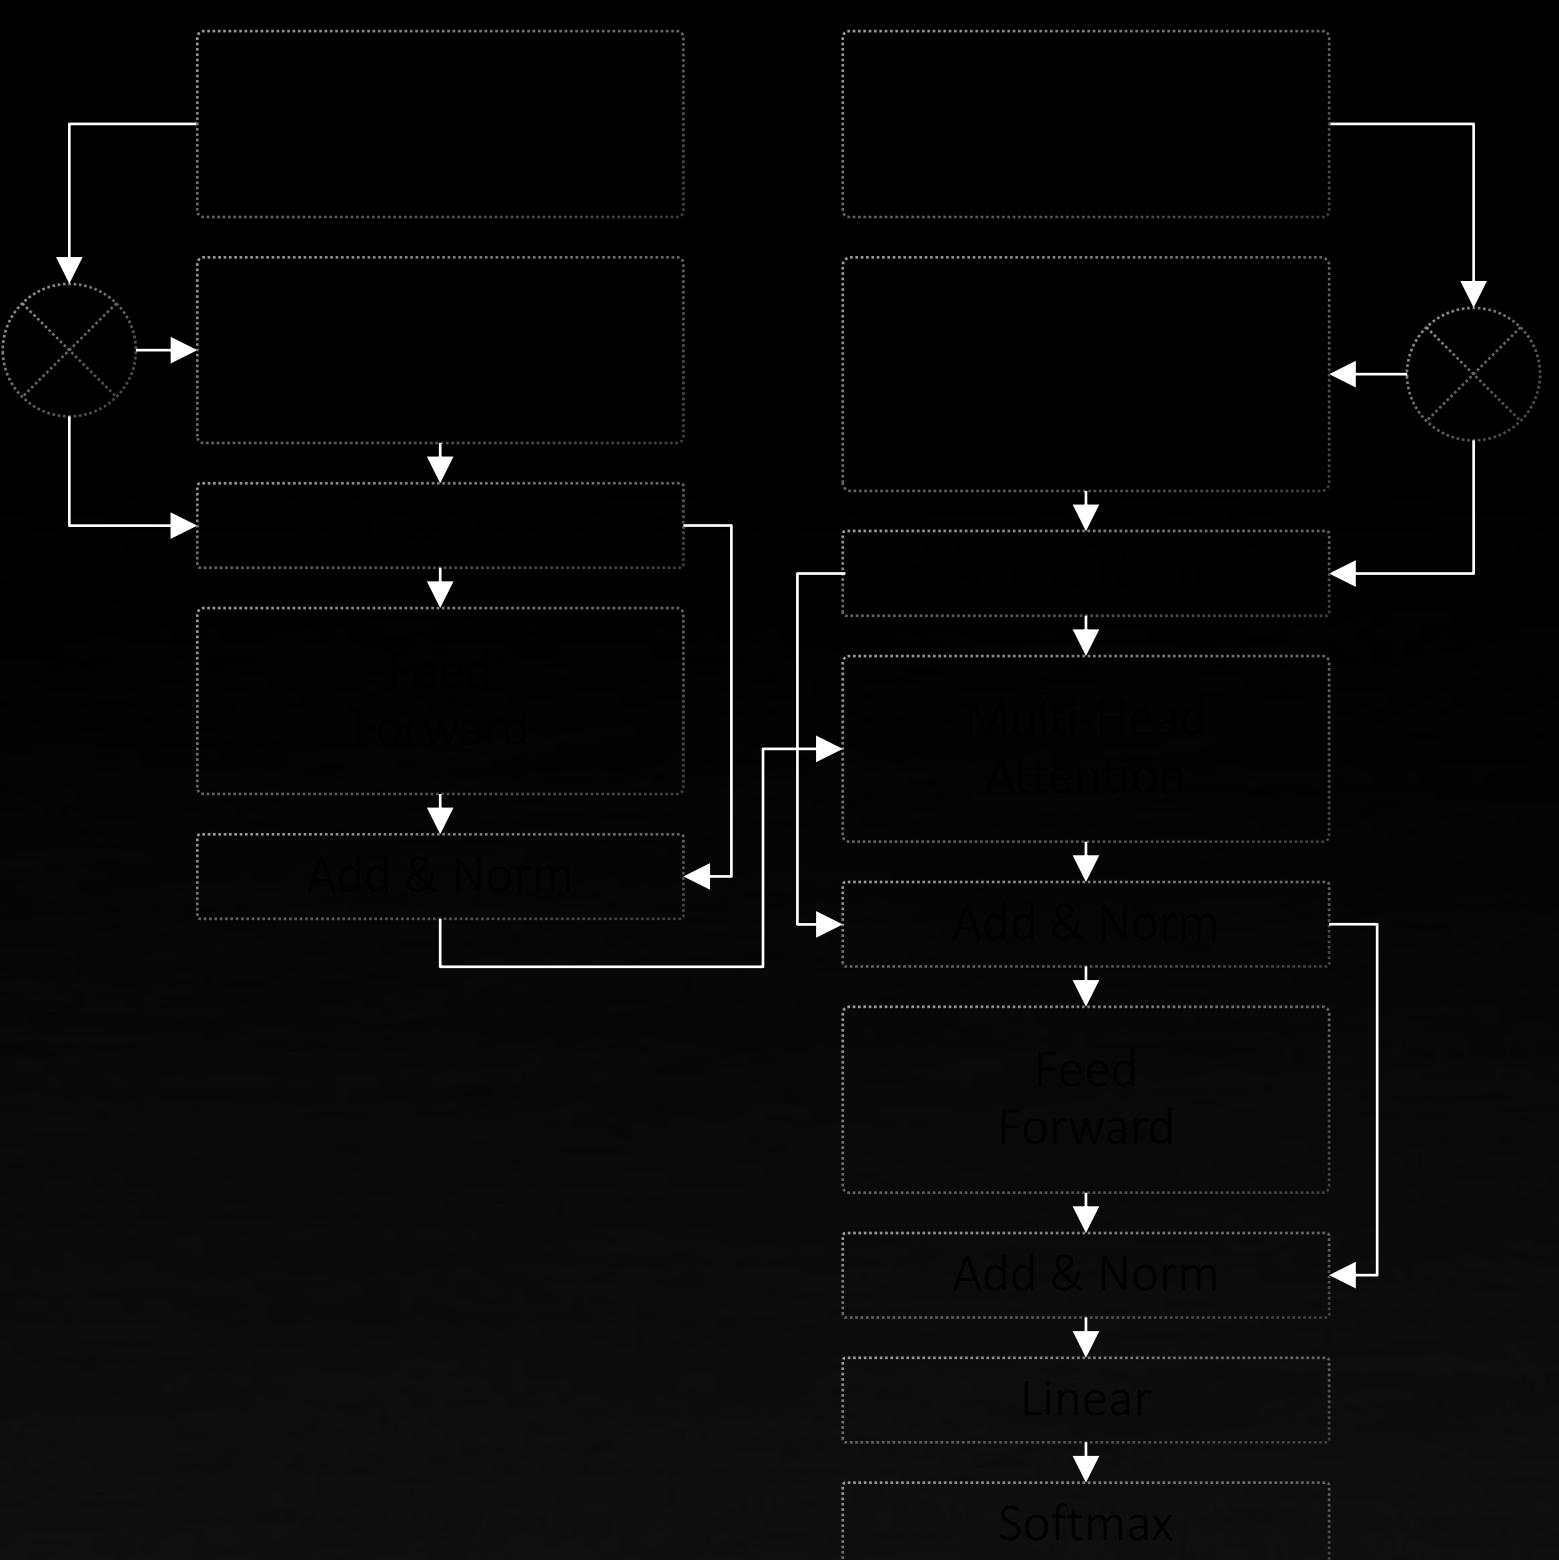
Library level



Runtime level

SCALING: TASK PARALLELISM + LOCALITY OF EXECUTION

Scheduling based on dependencies



Placement based on data exchange

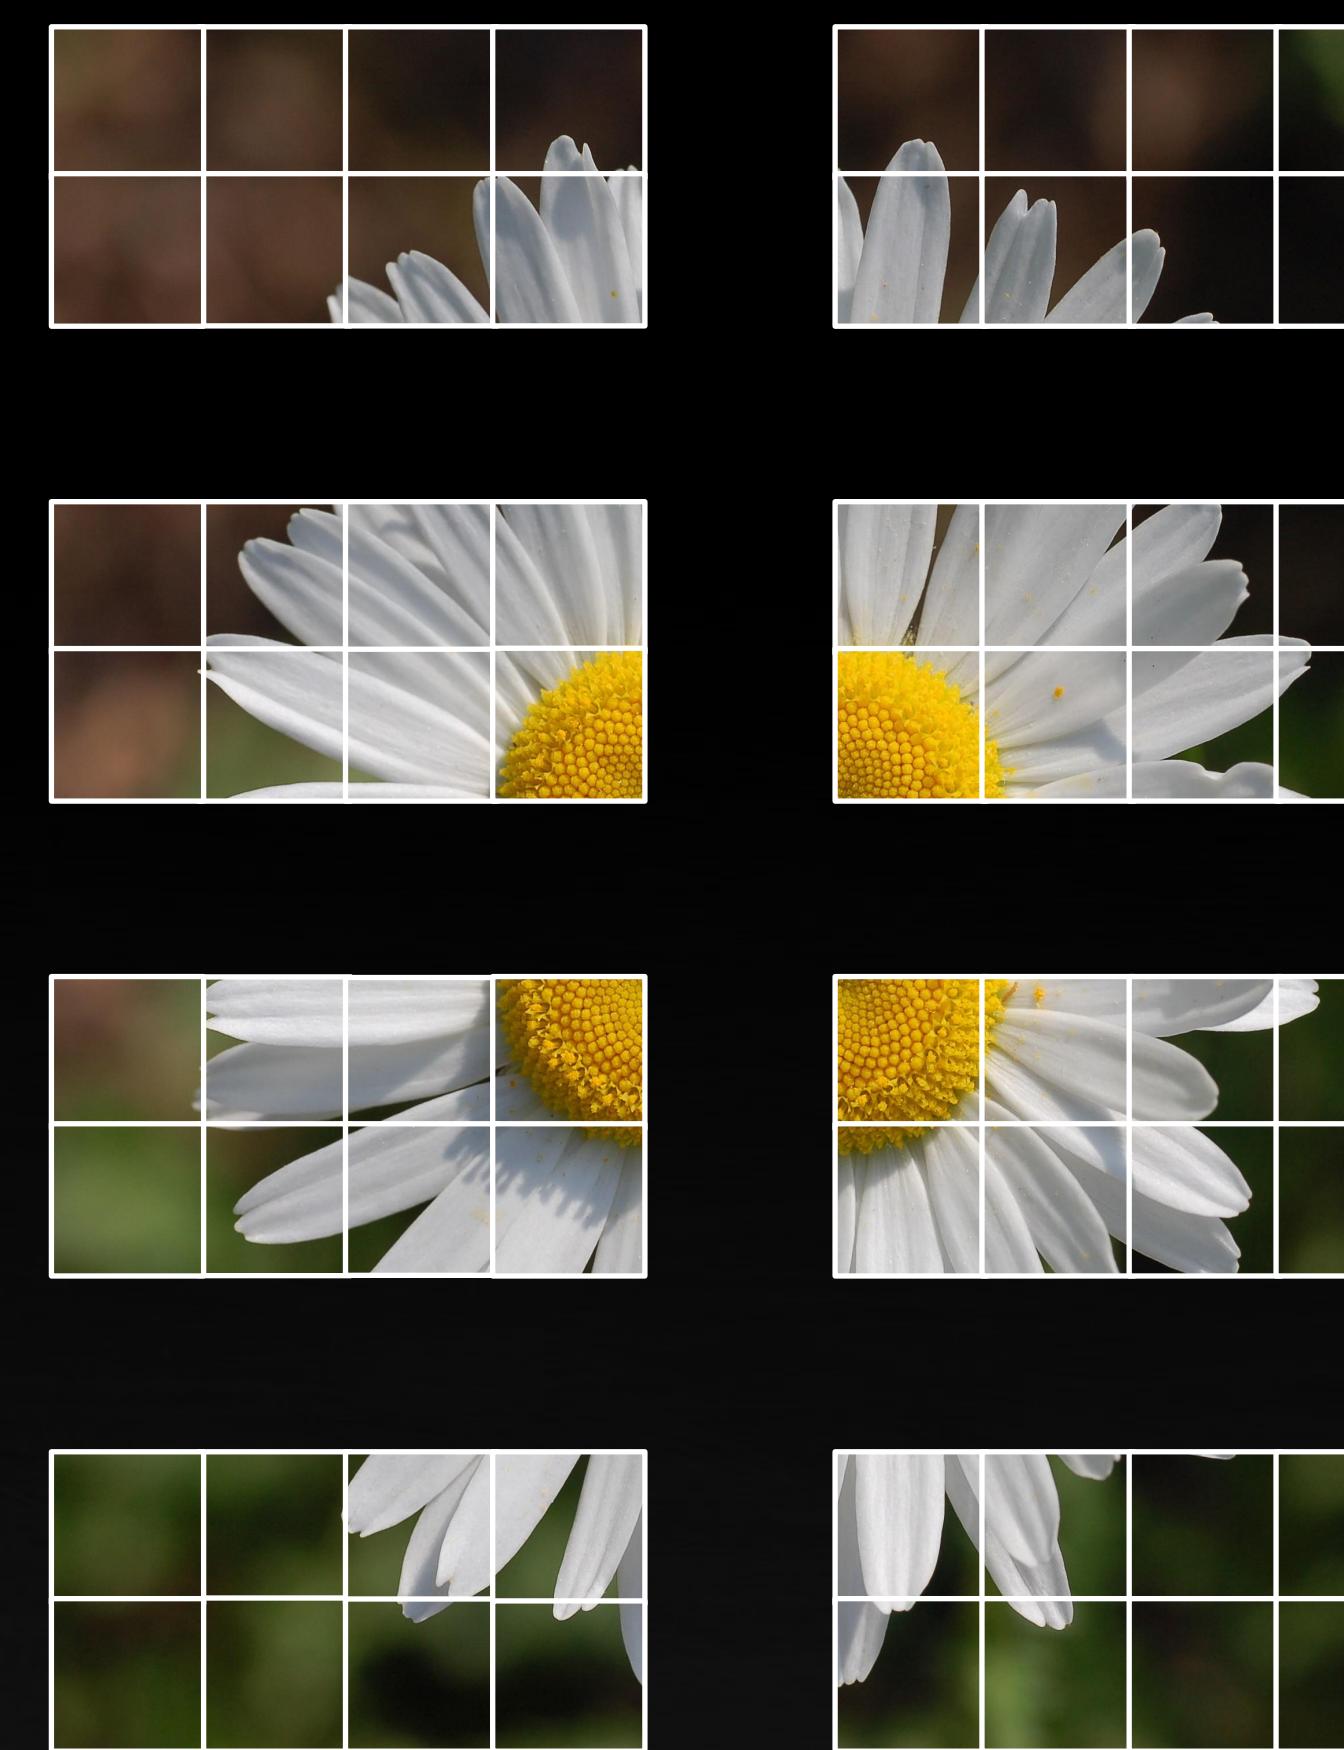


SCALING: DATA PARALLELISM + LOCALITY OF DATA

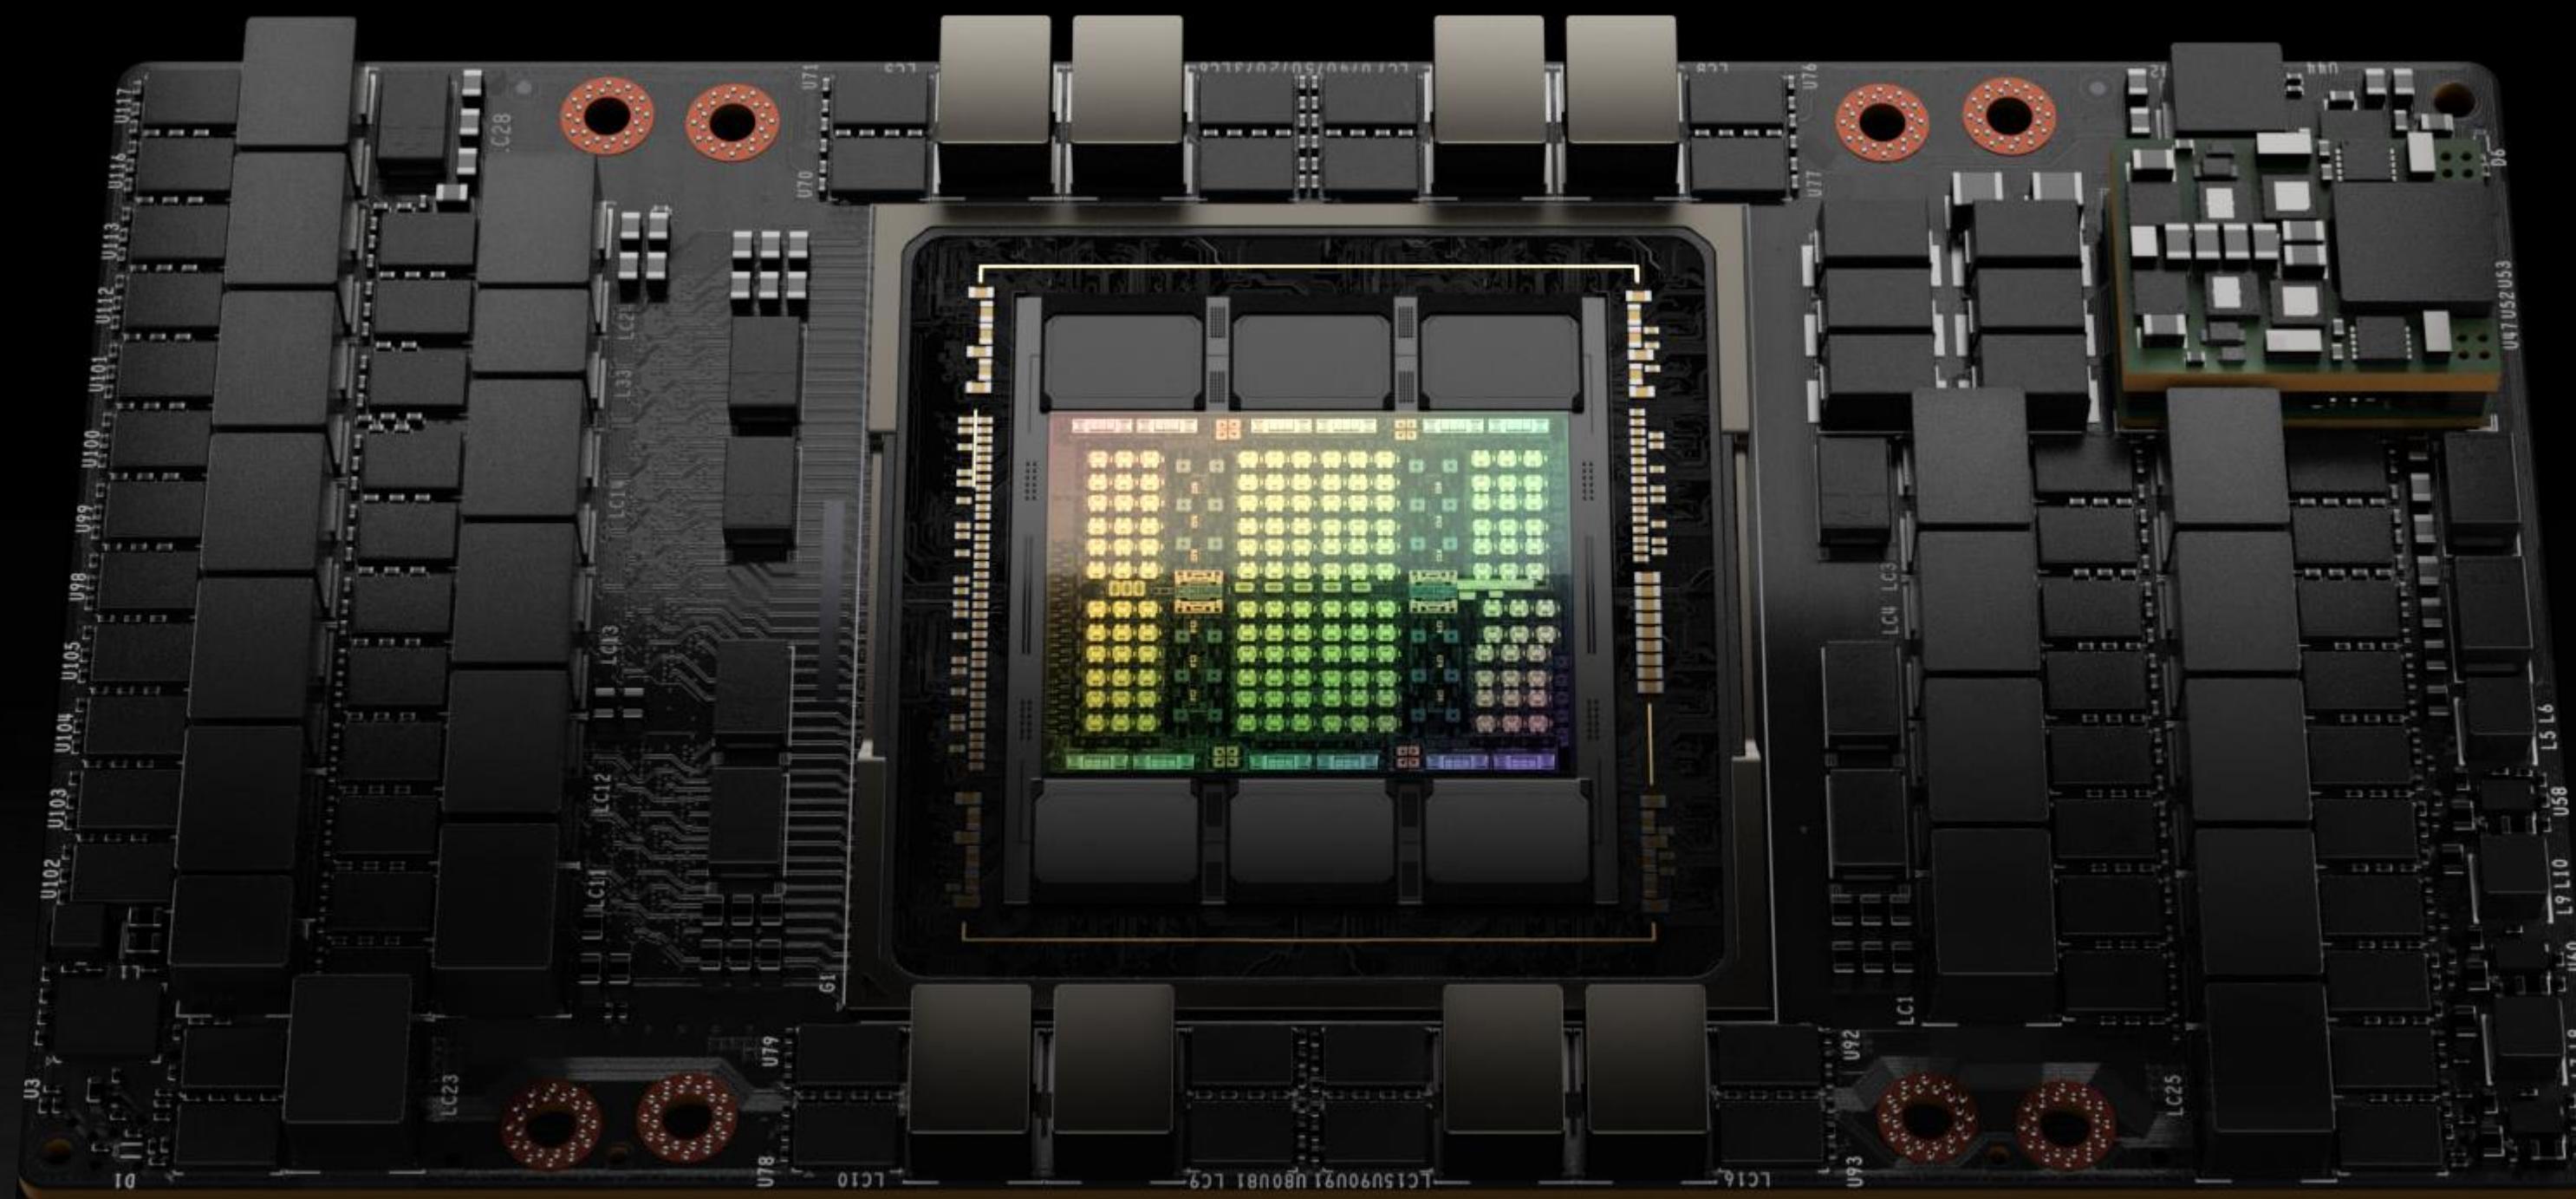
Scheduling based on throughput



Placement based on data reuse



INTRODUCING HOPPER



H100 Physical Architectural Features

132 Streaming Multiprocessors (SMs)

PCIe Gen5 with PCIe Atomics

HBM3 Memory with 3TB/sec Bandwidth

50MB L2 Cache

4th Generation NVLink @ 900GB/sec total bandwidth

H100 Next-Generation Capabilities

Thread Block Clusters

Distributed Shared Memory

Tensor Memory Accelerator (TMA)

Tensor Core Transformer Engine

Confidential Computing Support

Asynchronous Architecture

INTRODUCING HOPPER



H100 SM

256kB Shared Memory / L1-Cache

DPX Instruction Set

4th Generation Tensor Core

Fully-Asynchronous Architecture

Thread Block Clusters

Tensor Memory Accelerator

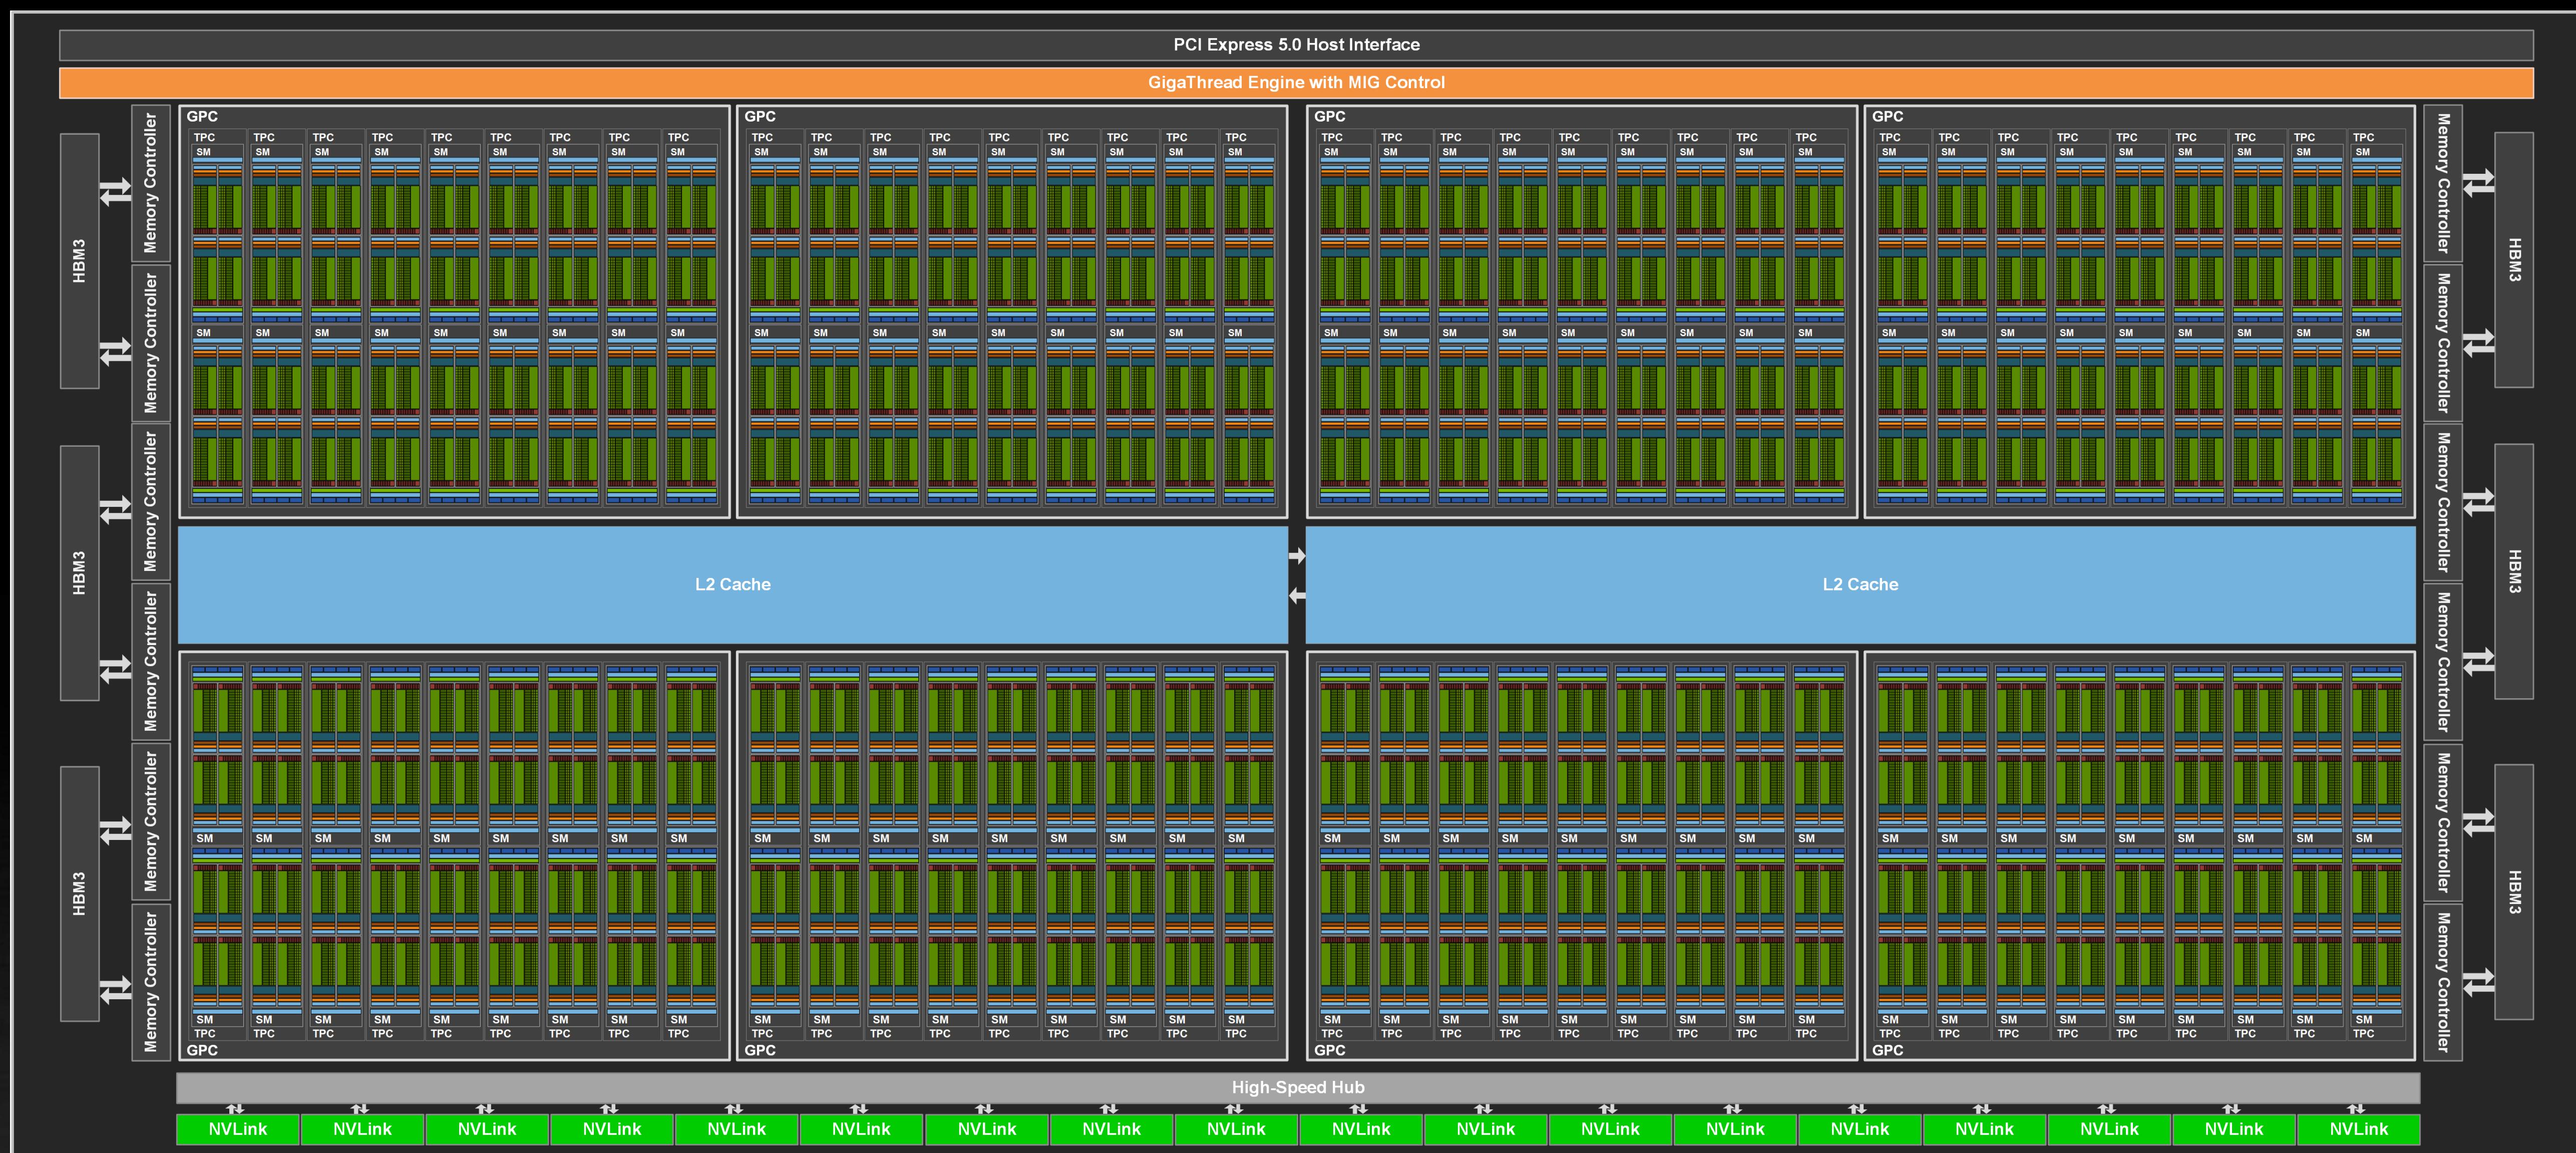
Asynchronous Transaction Barriers

SOME HISTORY: THE KEPLER GK110 GPU, 2012



15 SMs

THE HOPPER H100 GPU, 2022



Hopper H100 Full Chip

132 SMS

9x SMs OVER 10 YEARS



Kepler GK110 Full Chip
15 SMs

Hopper H100 Full Chip

132 SMs

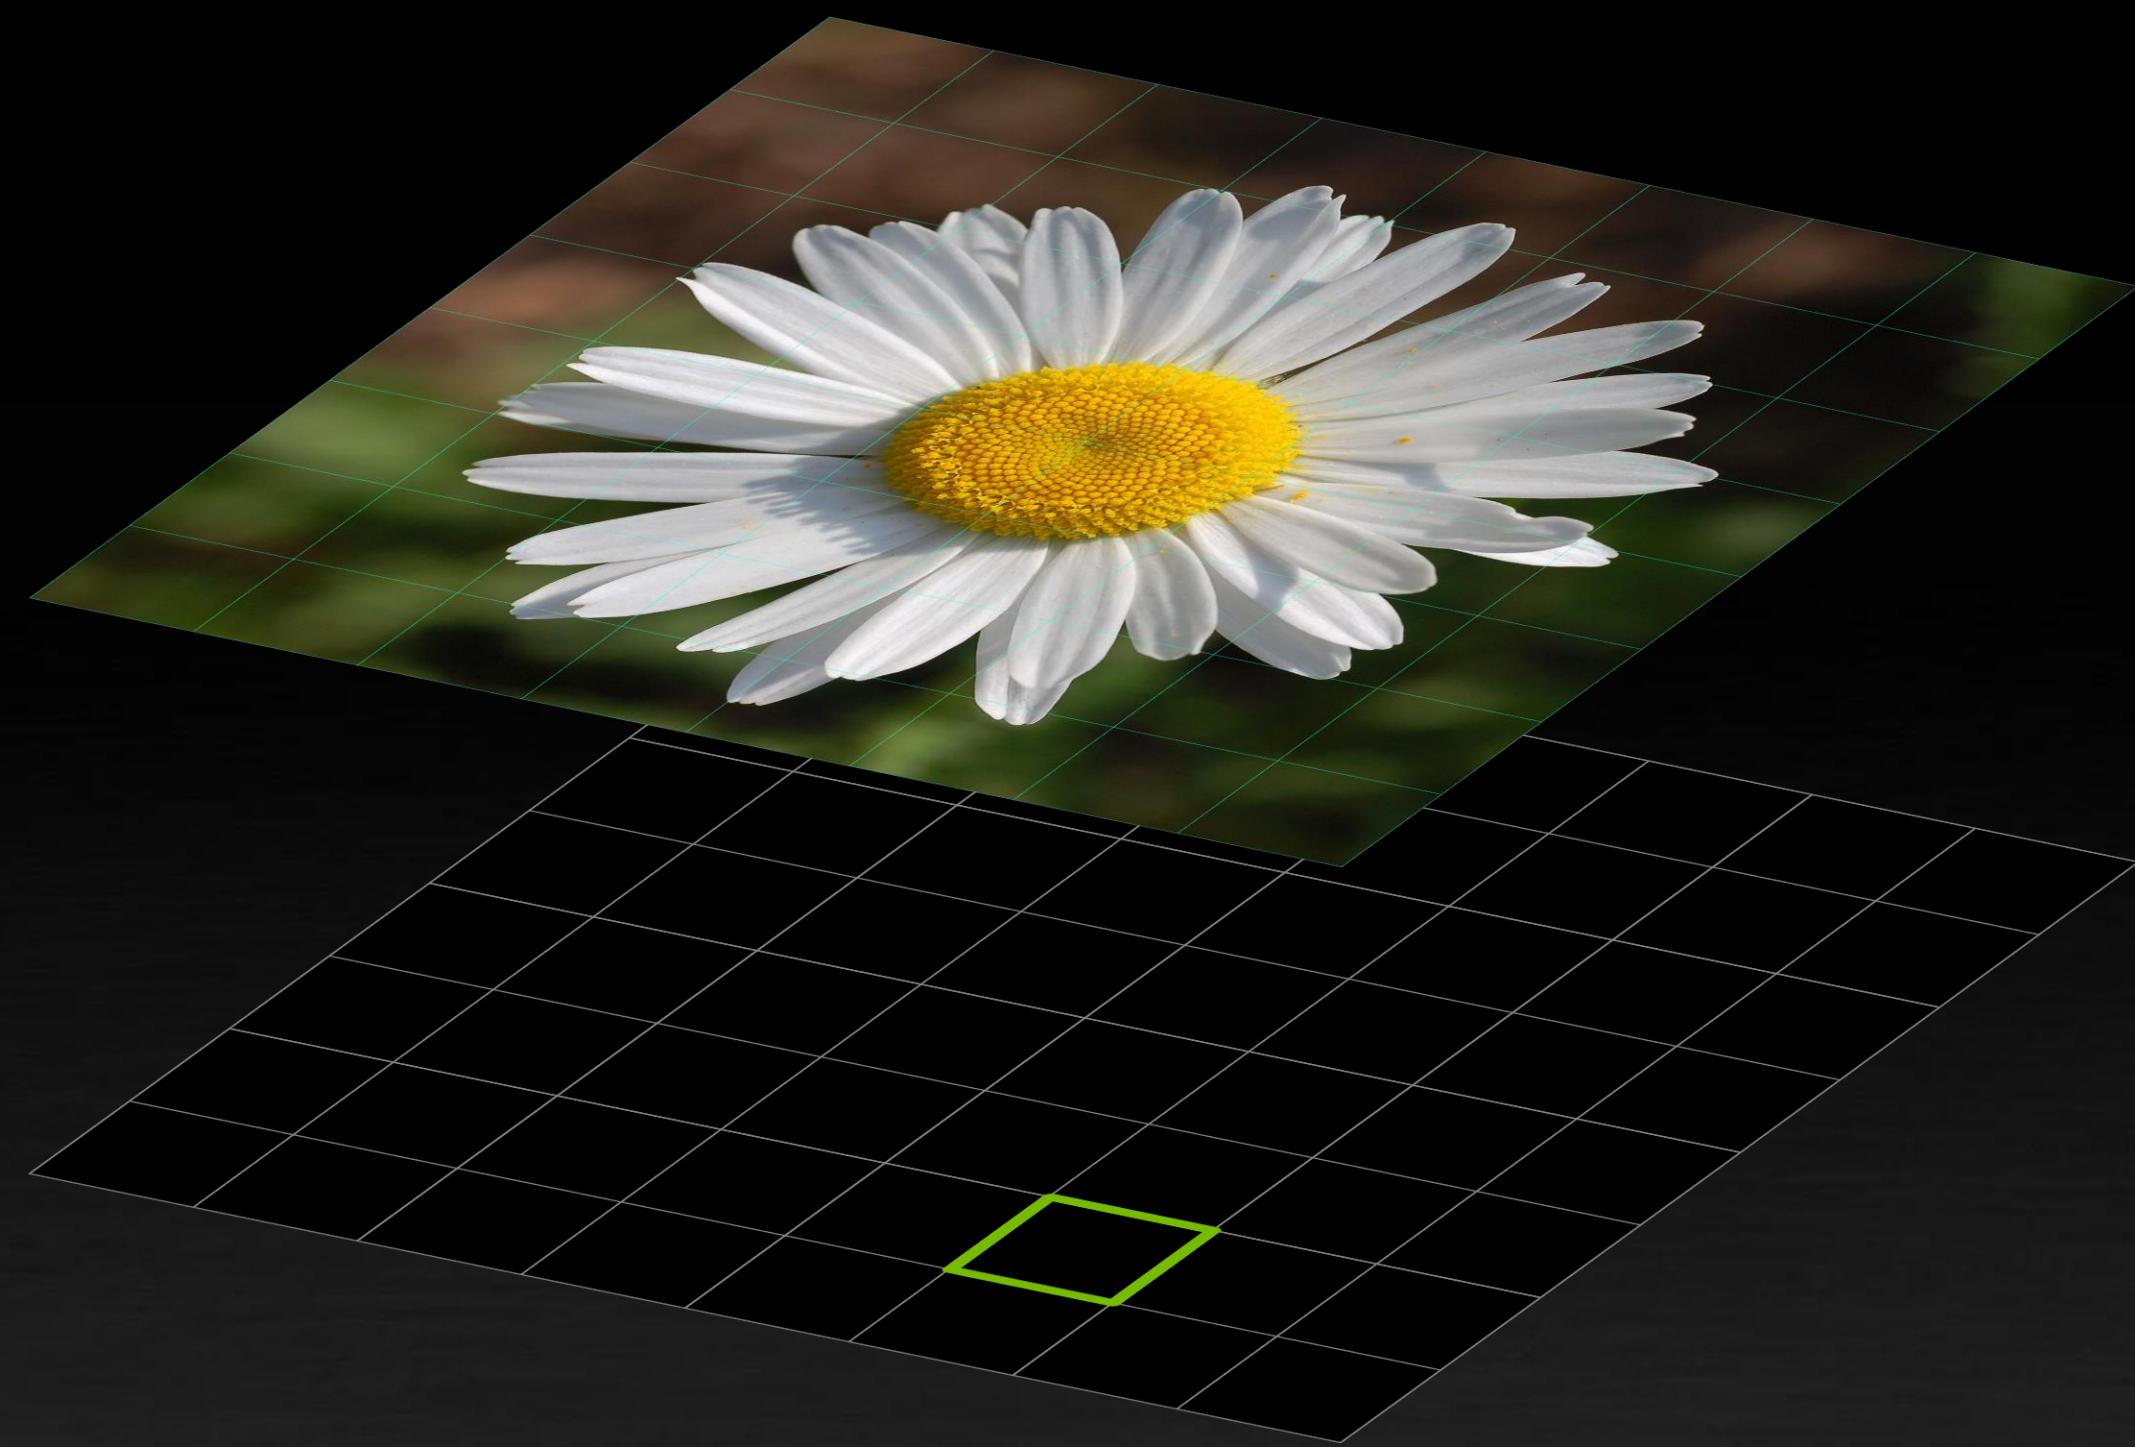
THE CUDA PROGRAMMING MODEL: GRID → BLOCKS → THREADS



Grid
of work

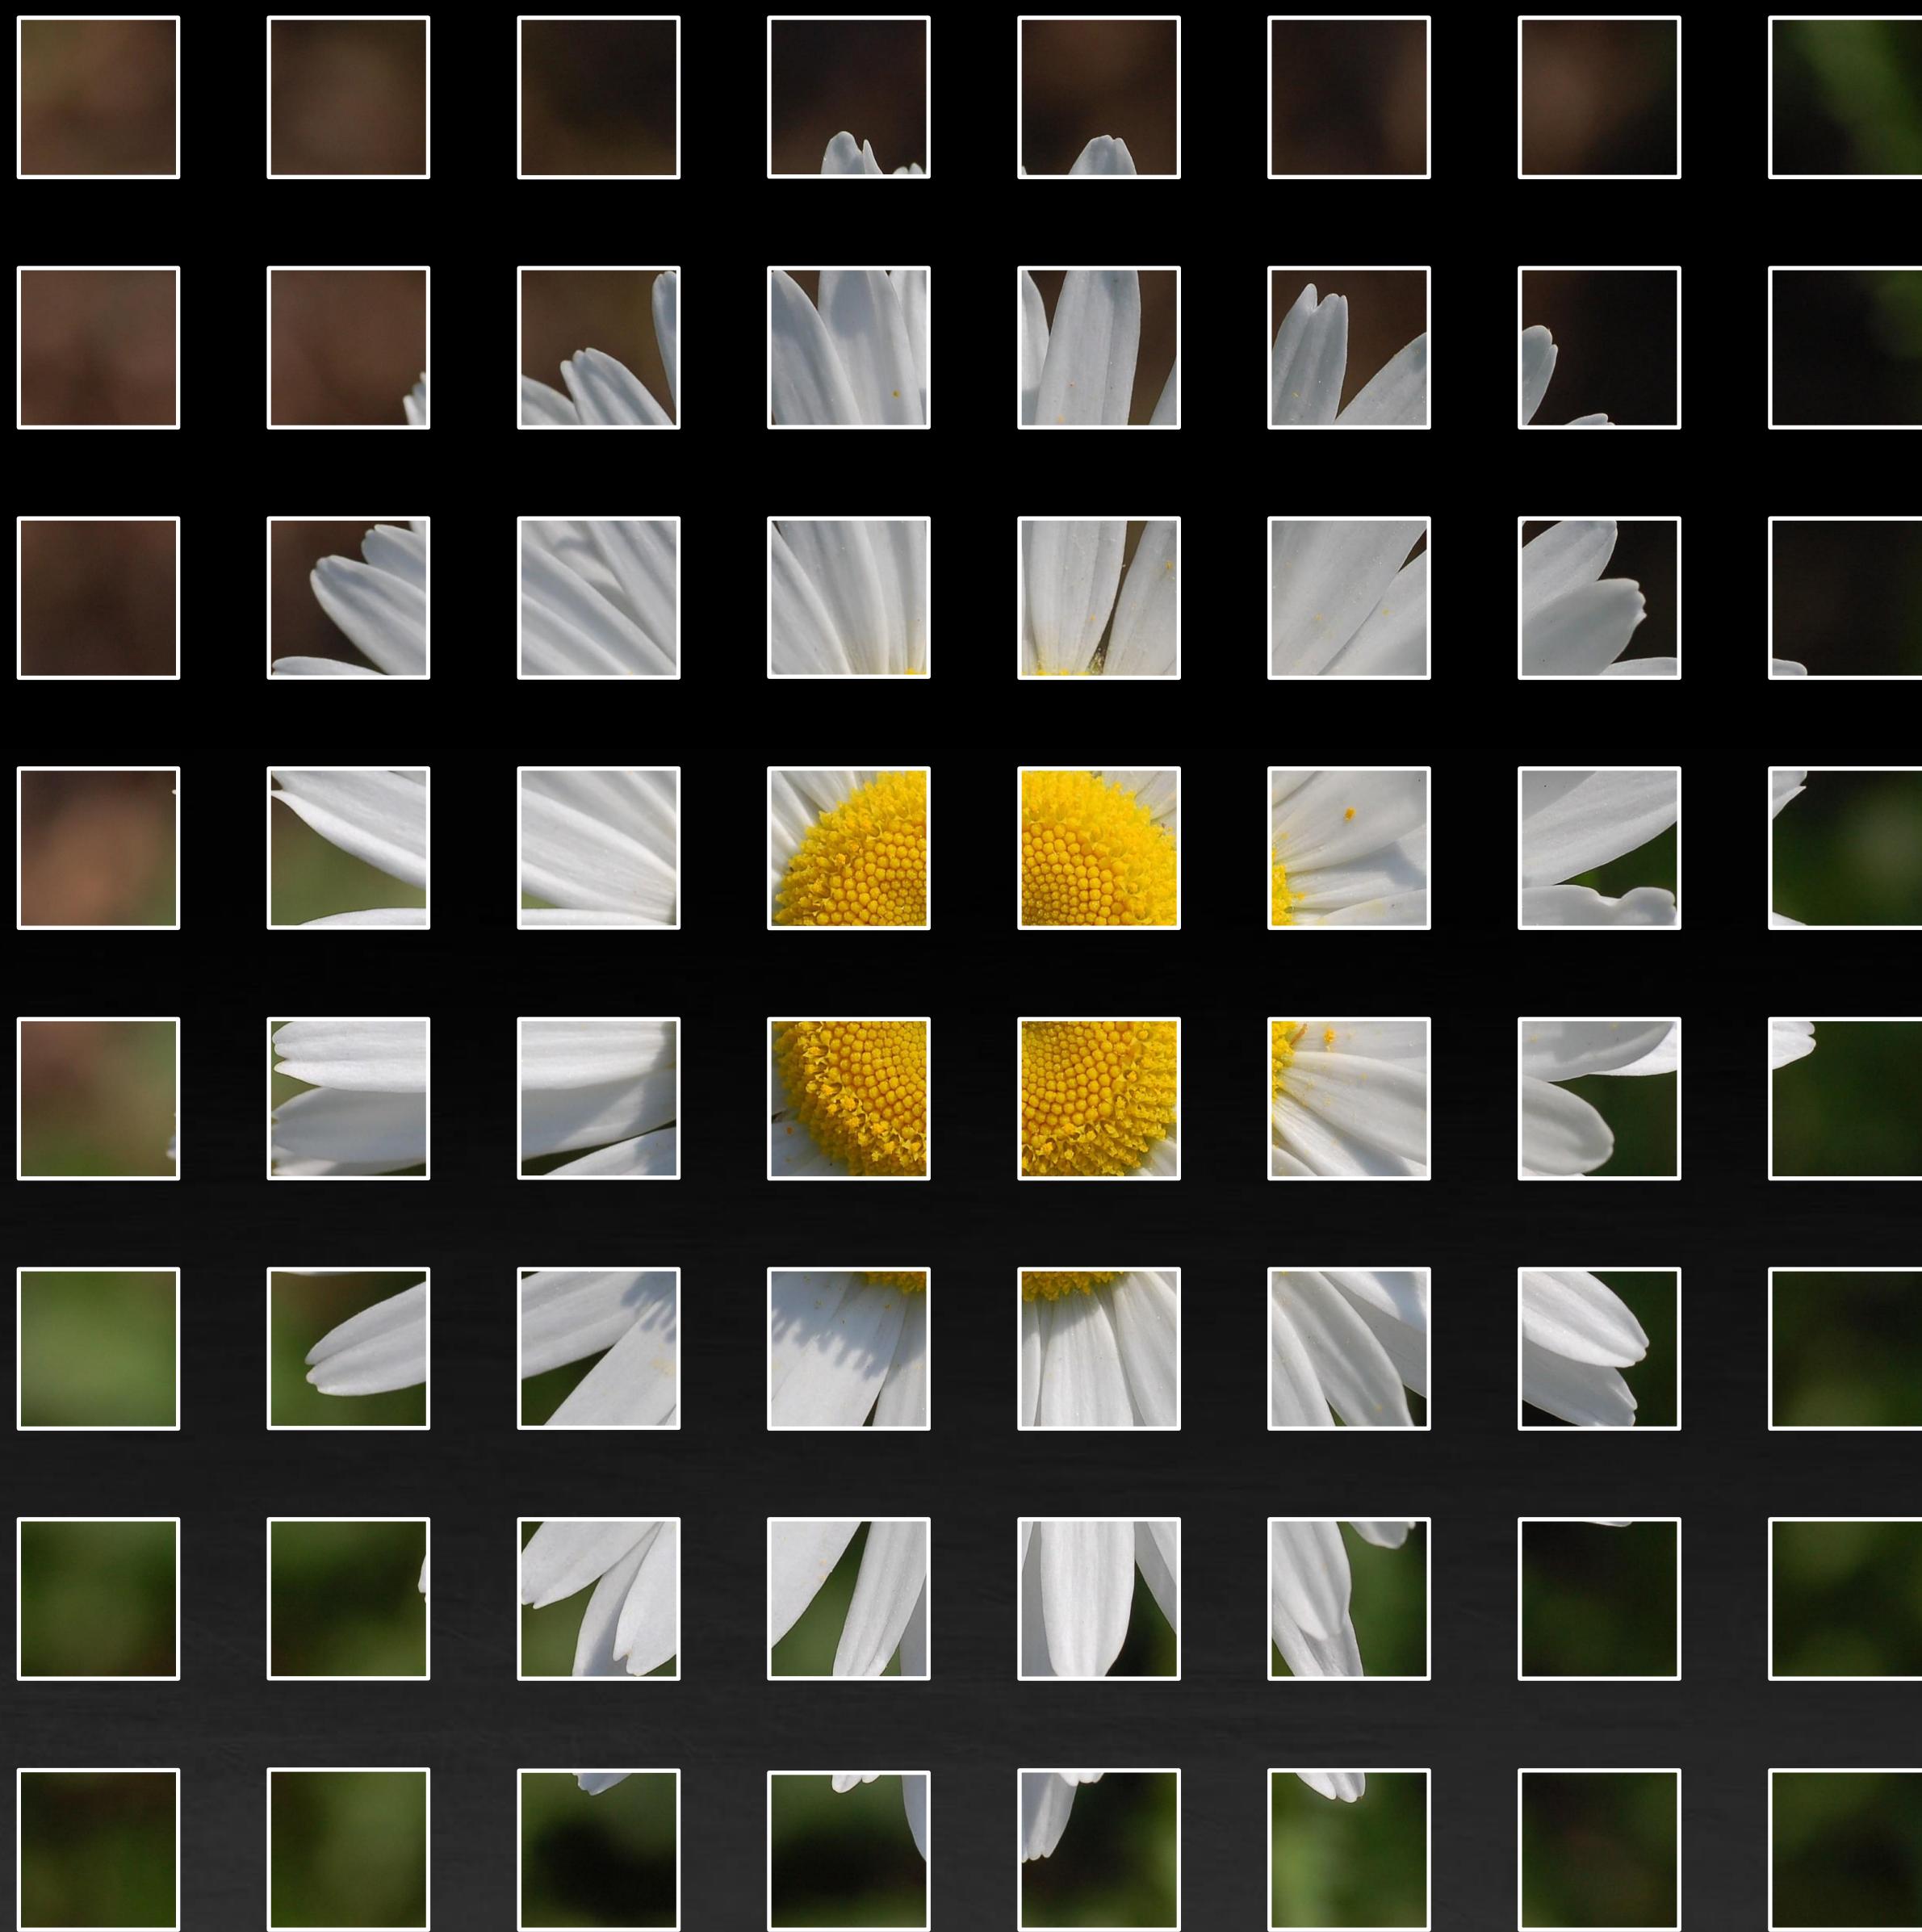


DIVIDE THE WORK INTO A GRID OF EQUAL BLOCKS

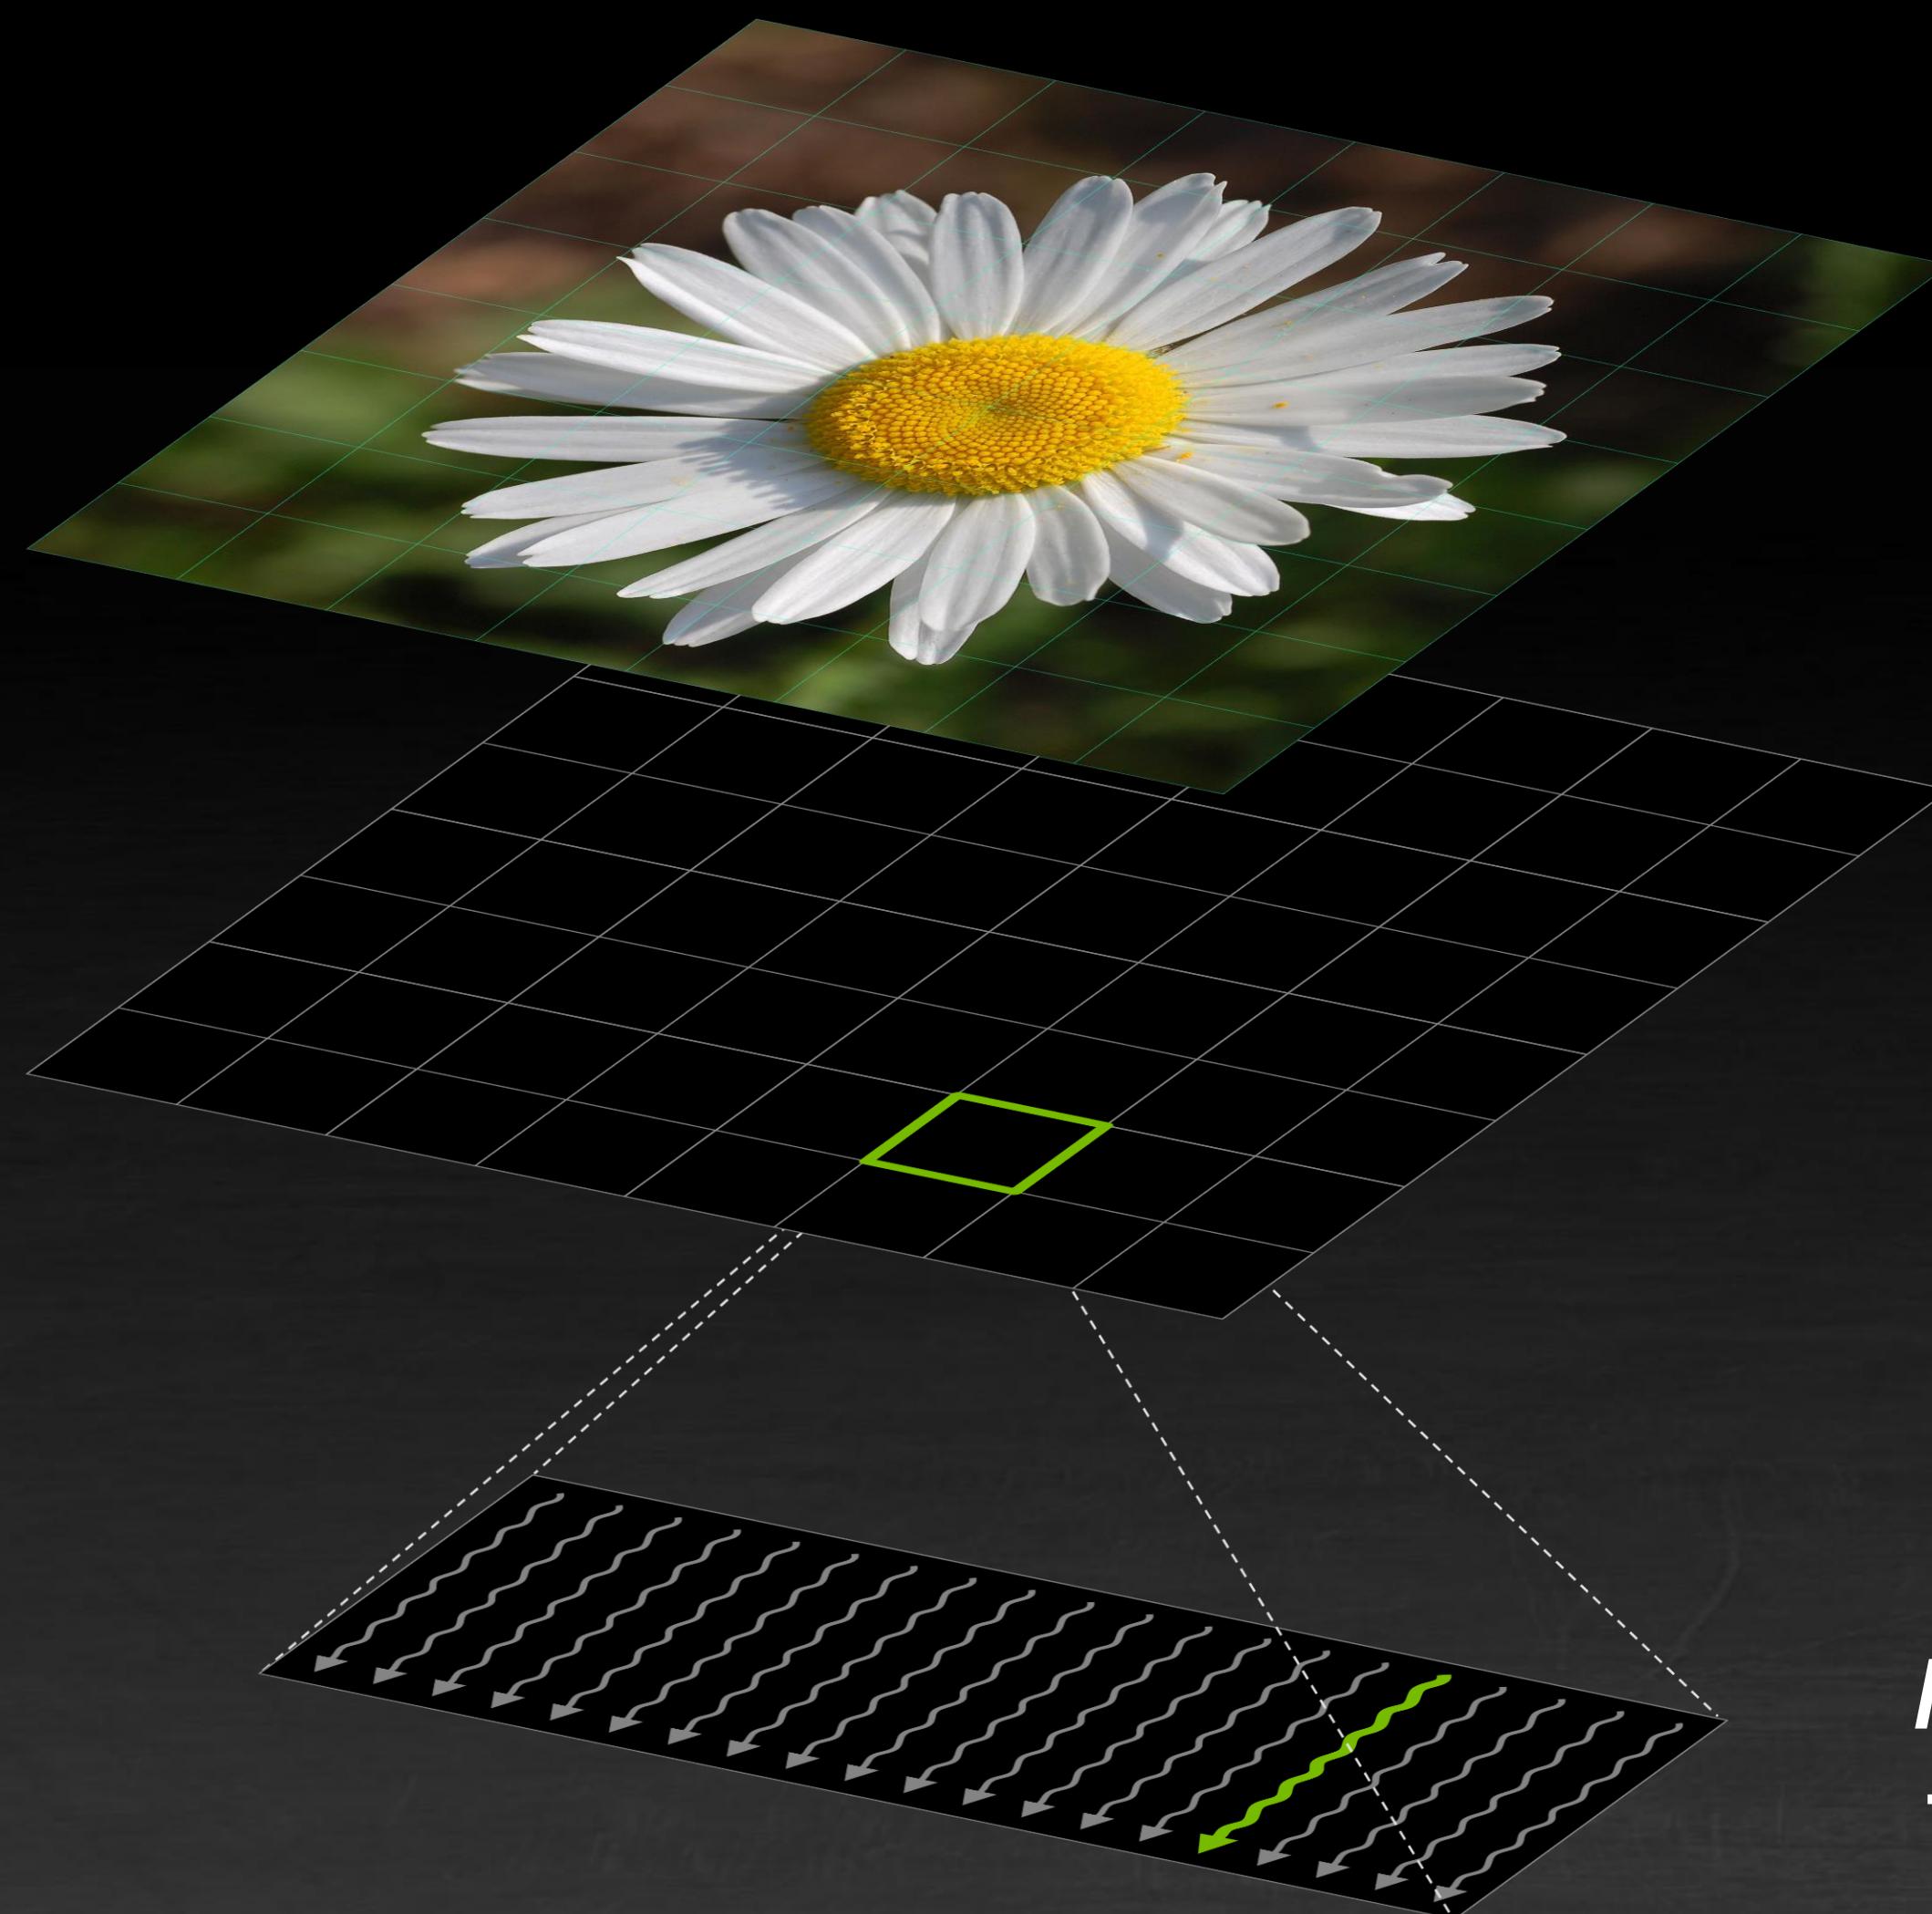


Grid
of work

Divide into
many Blocks



EACH BLOCK RUNS AS IF IT'S AN INDEPENDENT PROGRAM



Grid
of work

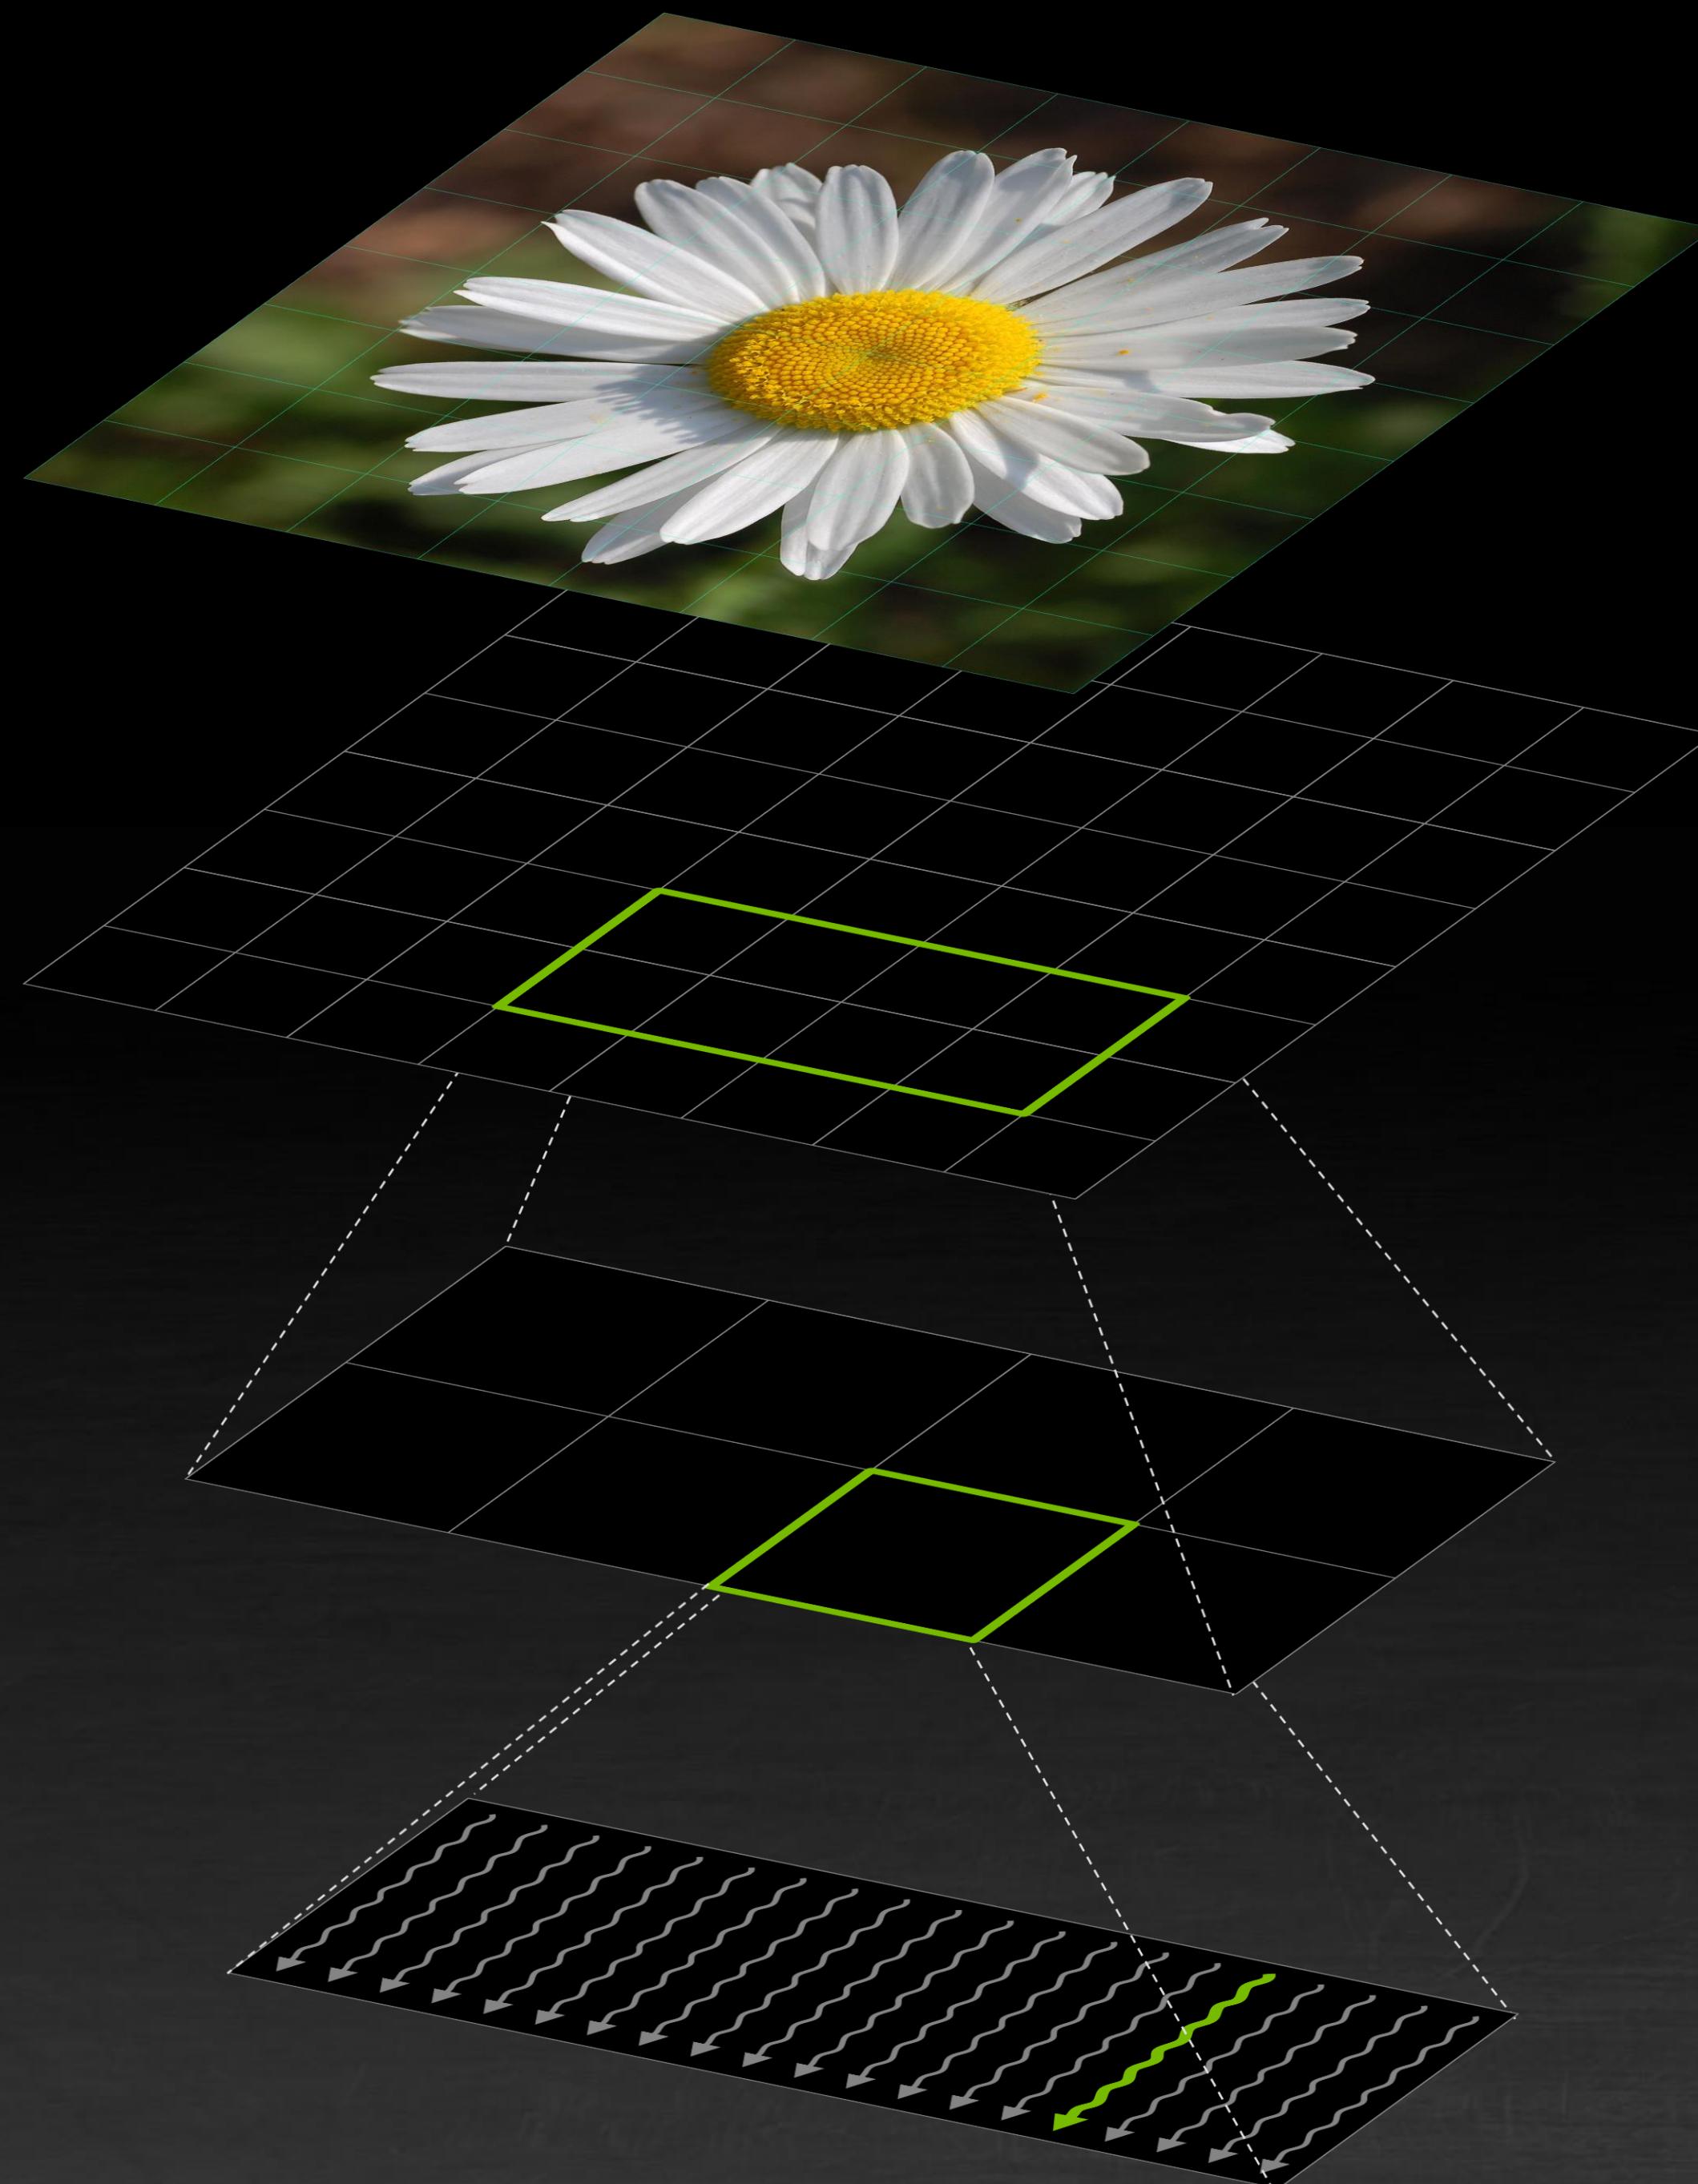
Blocks
of Threads

Many Threads
in each Block



THREAD BLOCK CLUSTER

A collective of blocks, co-scheduled on adjacent multiprocessors

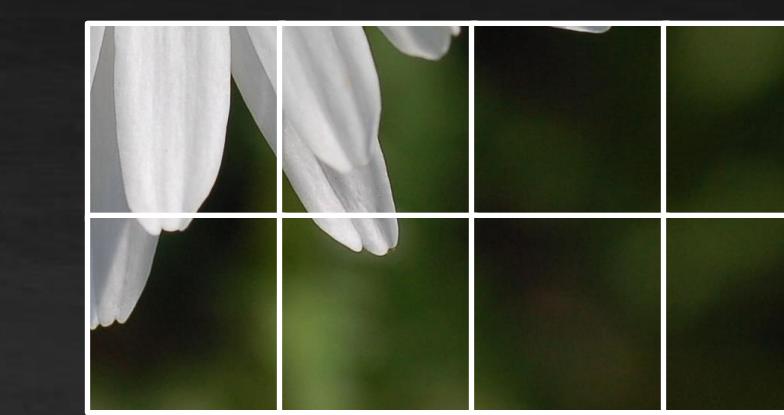
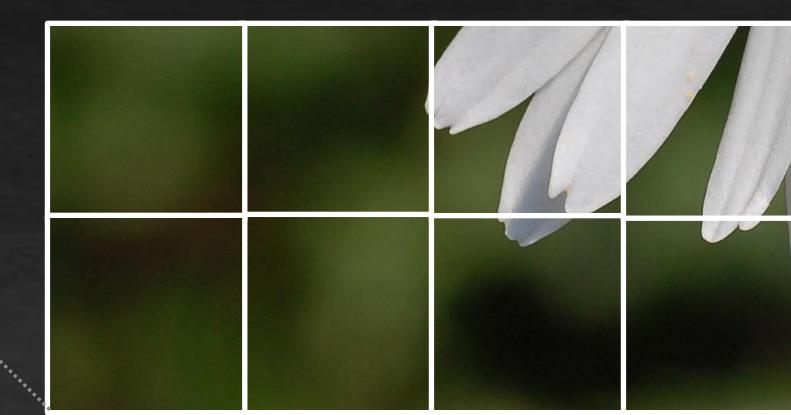
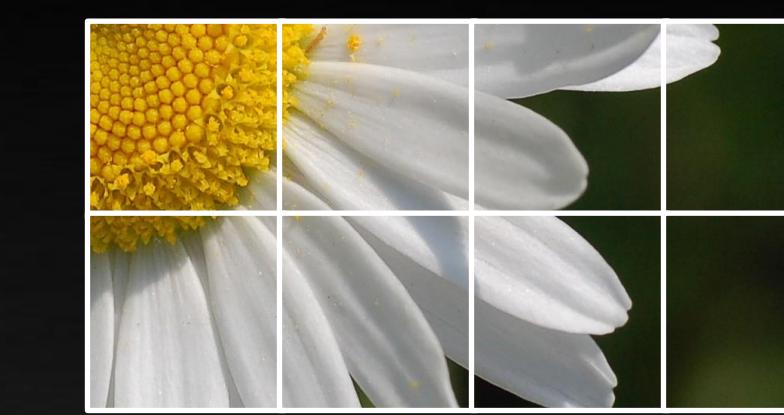
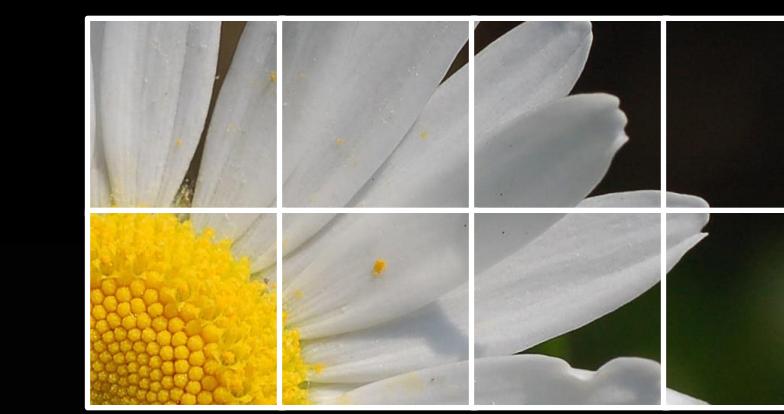
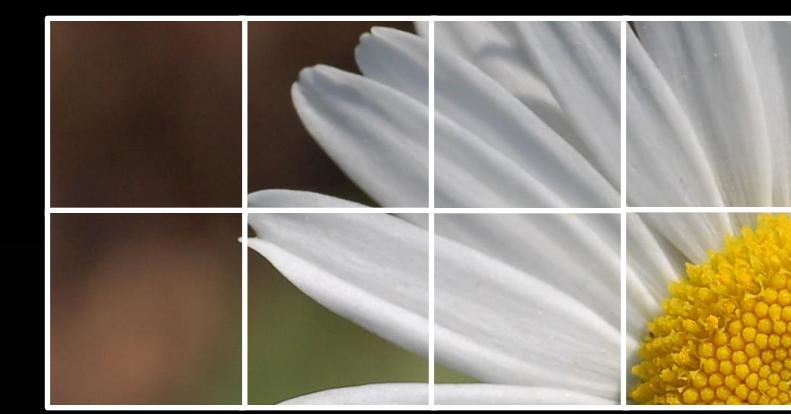
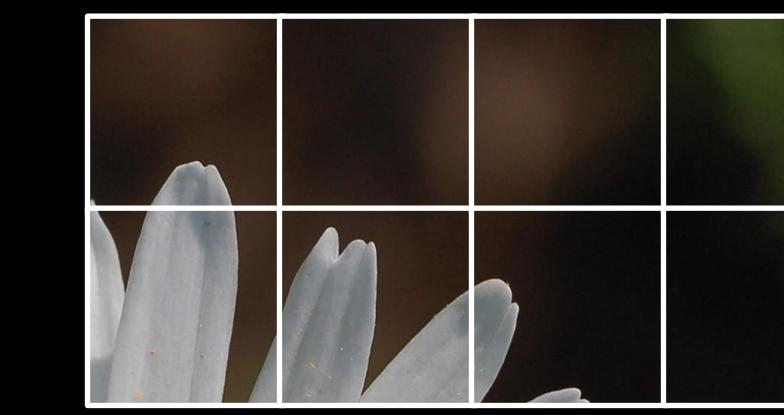
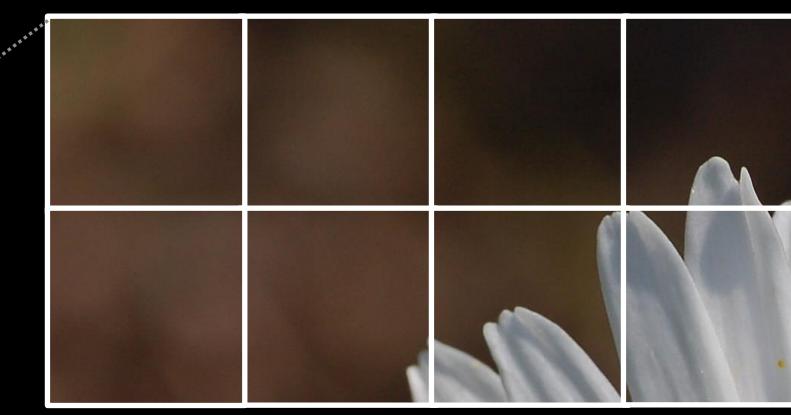


Grid
of work

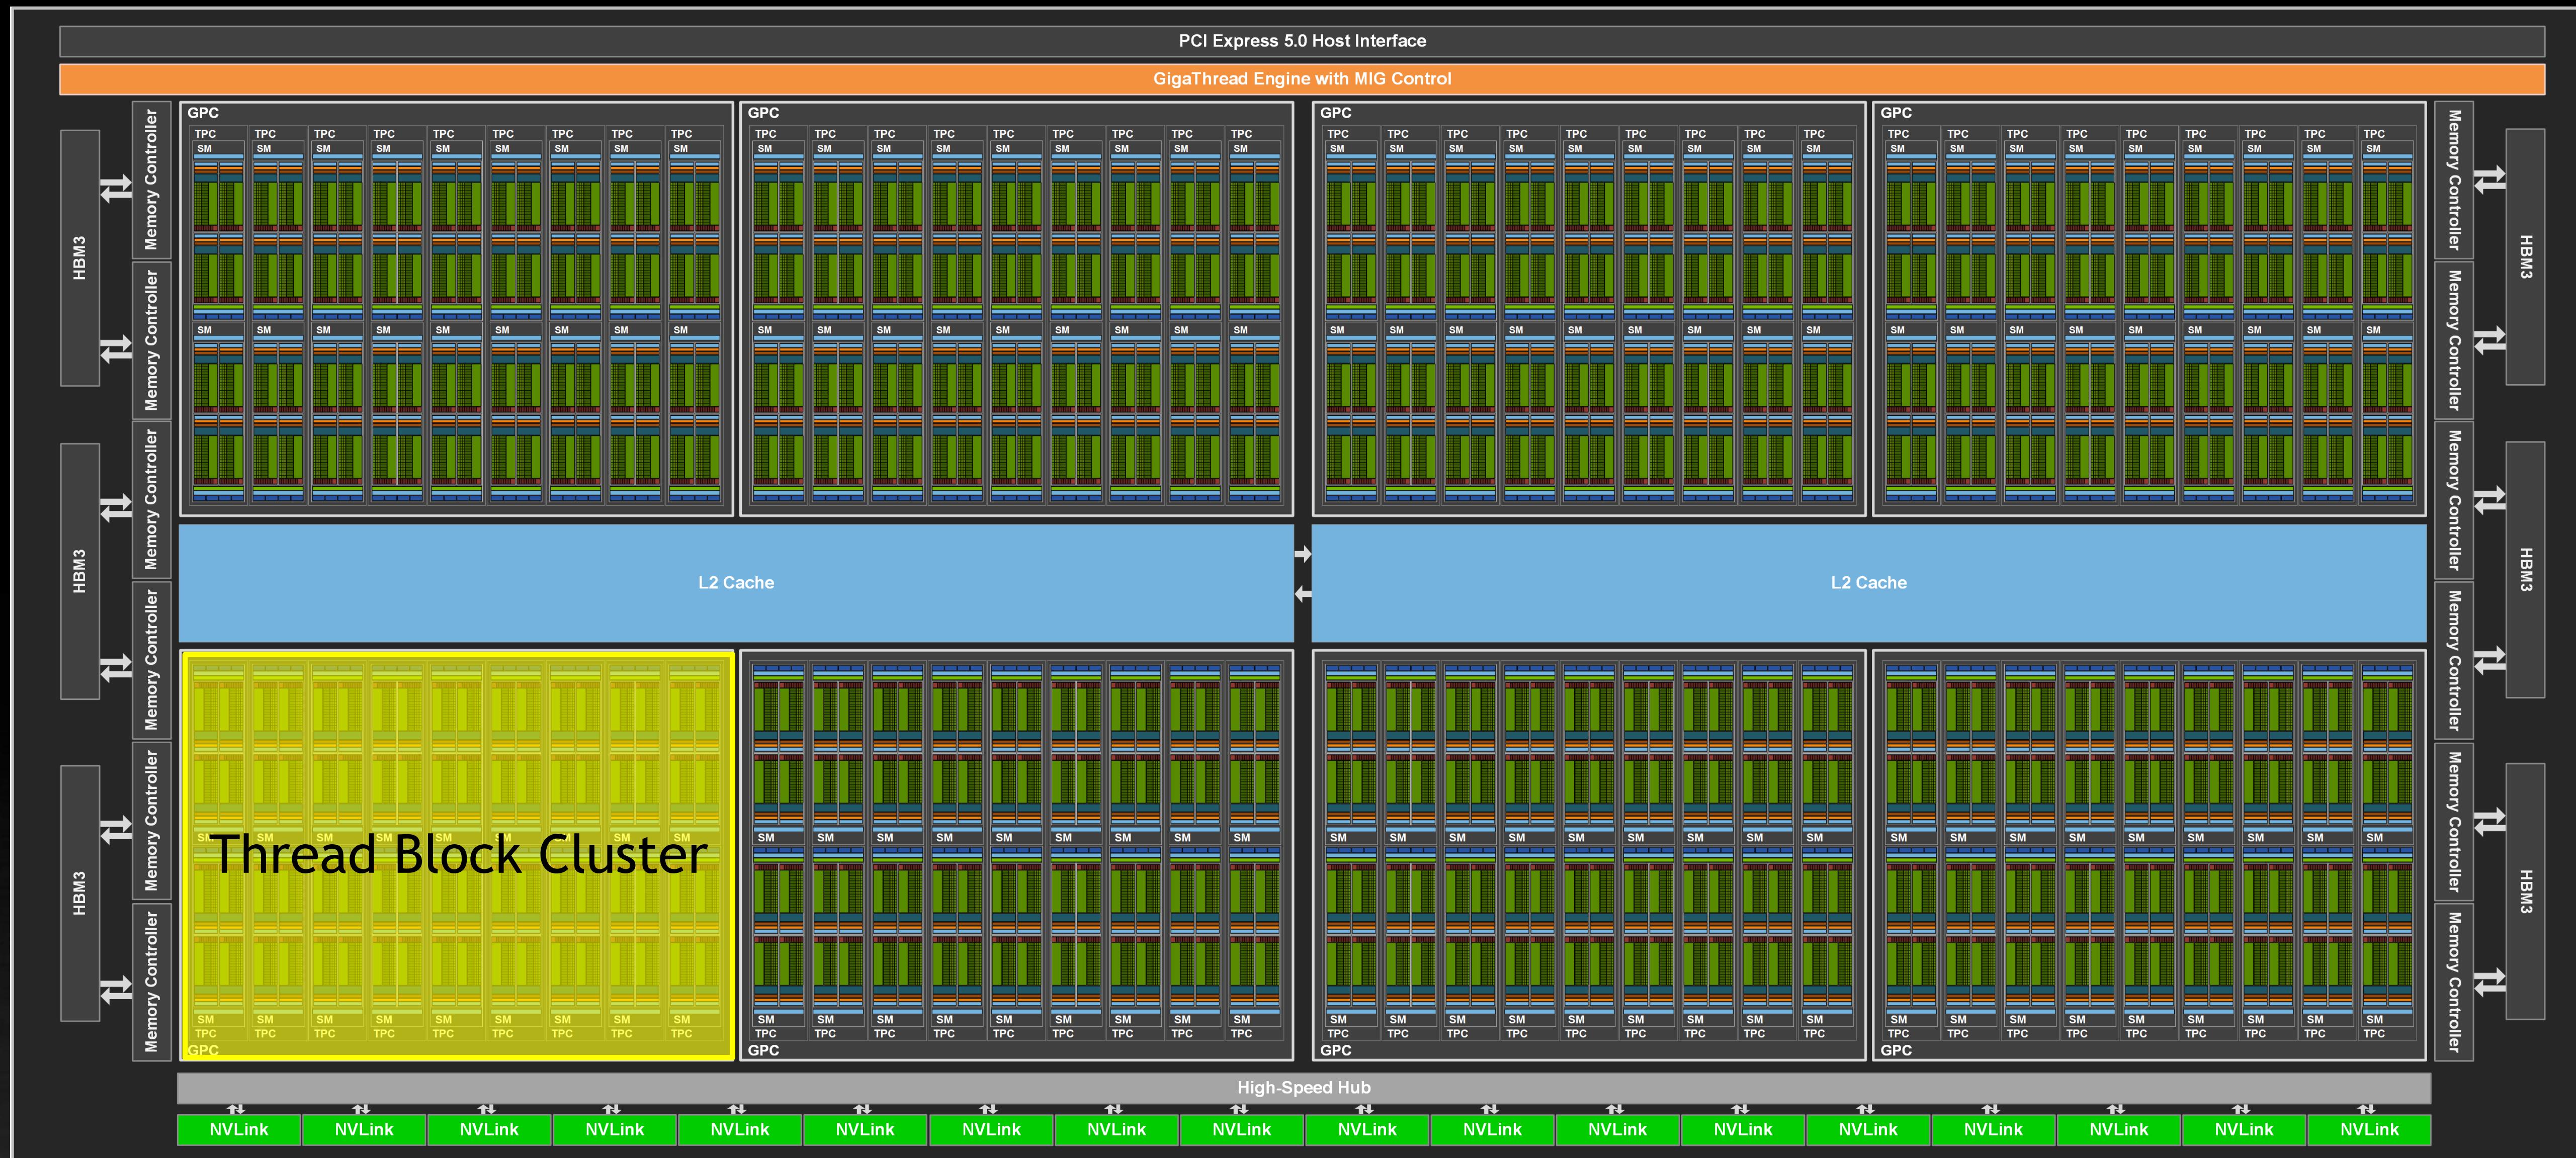
Cluster
of Blocks

Blocks
of Threads

Threads



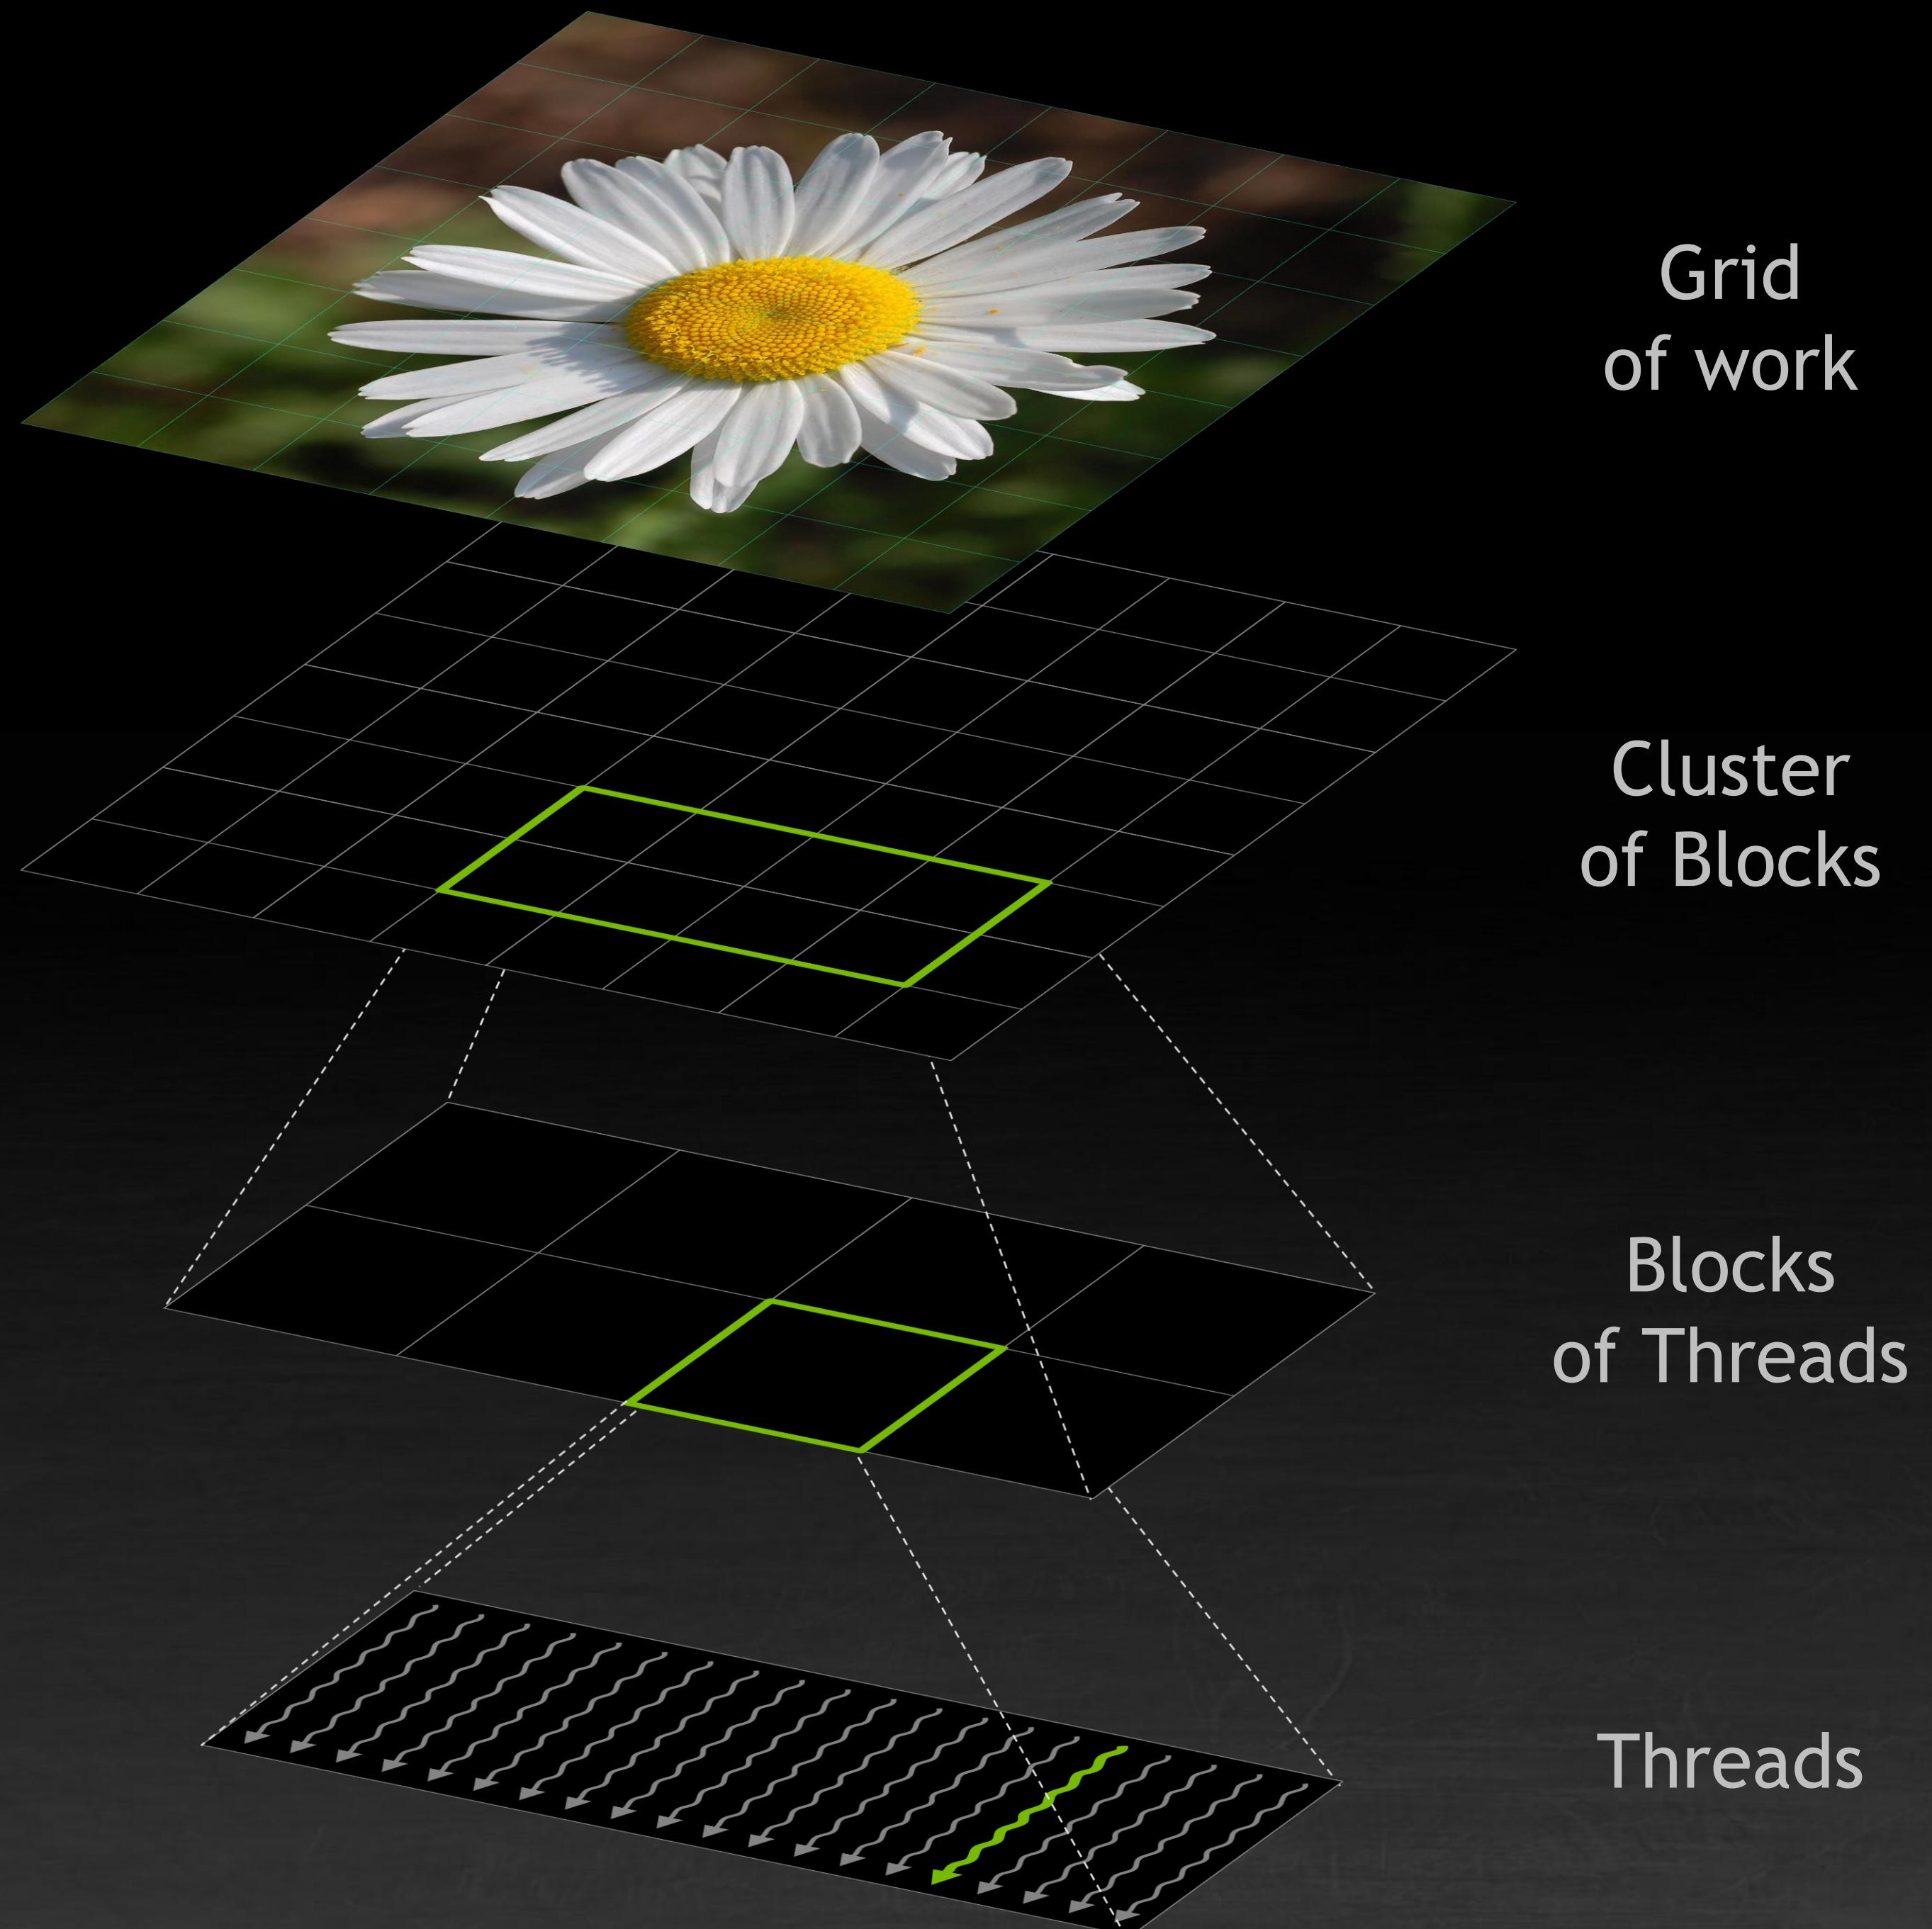
TAKING ADVANTAGE OF LOCALITY AT A GPU SCALE



Guaranteed co-located blocks
New tier of guaranteed concurrency
Fast data exchange & sync

THREAD BLOCK CLUSTER

Building hierarchy into a program



Grid
of work

Cluster
of Blocks

Blocks
of Threads

Threads

A cluster is a collective of up to 16 blocks

Guaranteed to be on different SMs

Guaranteed to be running at the same time

1D, 2D or 3D, just like blocks

Annotate a kernel with its required cluster size

New cluster dimension annotation for `__global__` functions:

`__cluster_dims__(x, [y, [z]])`

```
__cluster_dims__(4, 2, 1)           // 8-block cluster of size 4x2x1
__global__ void helloCluster()
{
    cooperative_groups::cluster_group cluster = this_cluster();
    cluster.sync();

    printf("Hello from cluster elem %d\n", cluster.cluster_rank());
}
```

Plus: New extensible launch API allows configuration at launch time

CLUSTER DISTRIBUTED SHARED MEMORY (DSMEM)

Blocks within a cluster are able to access each others' shared memory directly



Full load/store/atomic access to all shared memory
between blocks within a cluster

CLUSTER DISTRIBUTED SHARED MEMORY (DSMEM)

Blocks within a cluster are able to access each others' shared memory directly

Distributed Shared Memory (DSMEM) Programming Model

All blocks in a cluster can collaborate using DSMEM

Blocks in a cluster can synchronize together via **barriers in DSMEM**

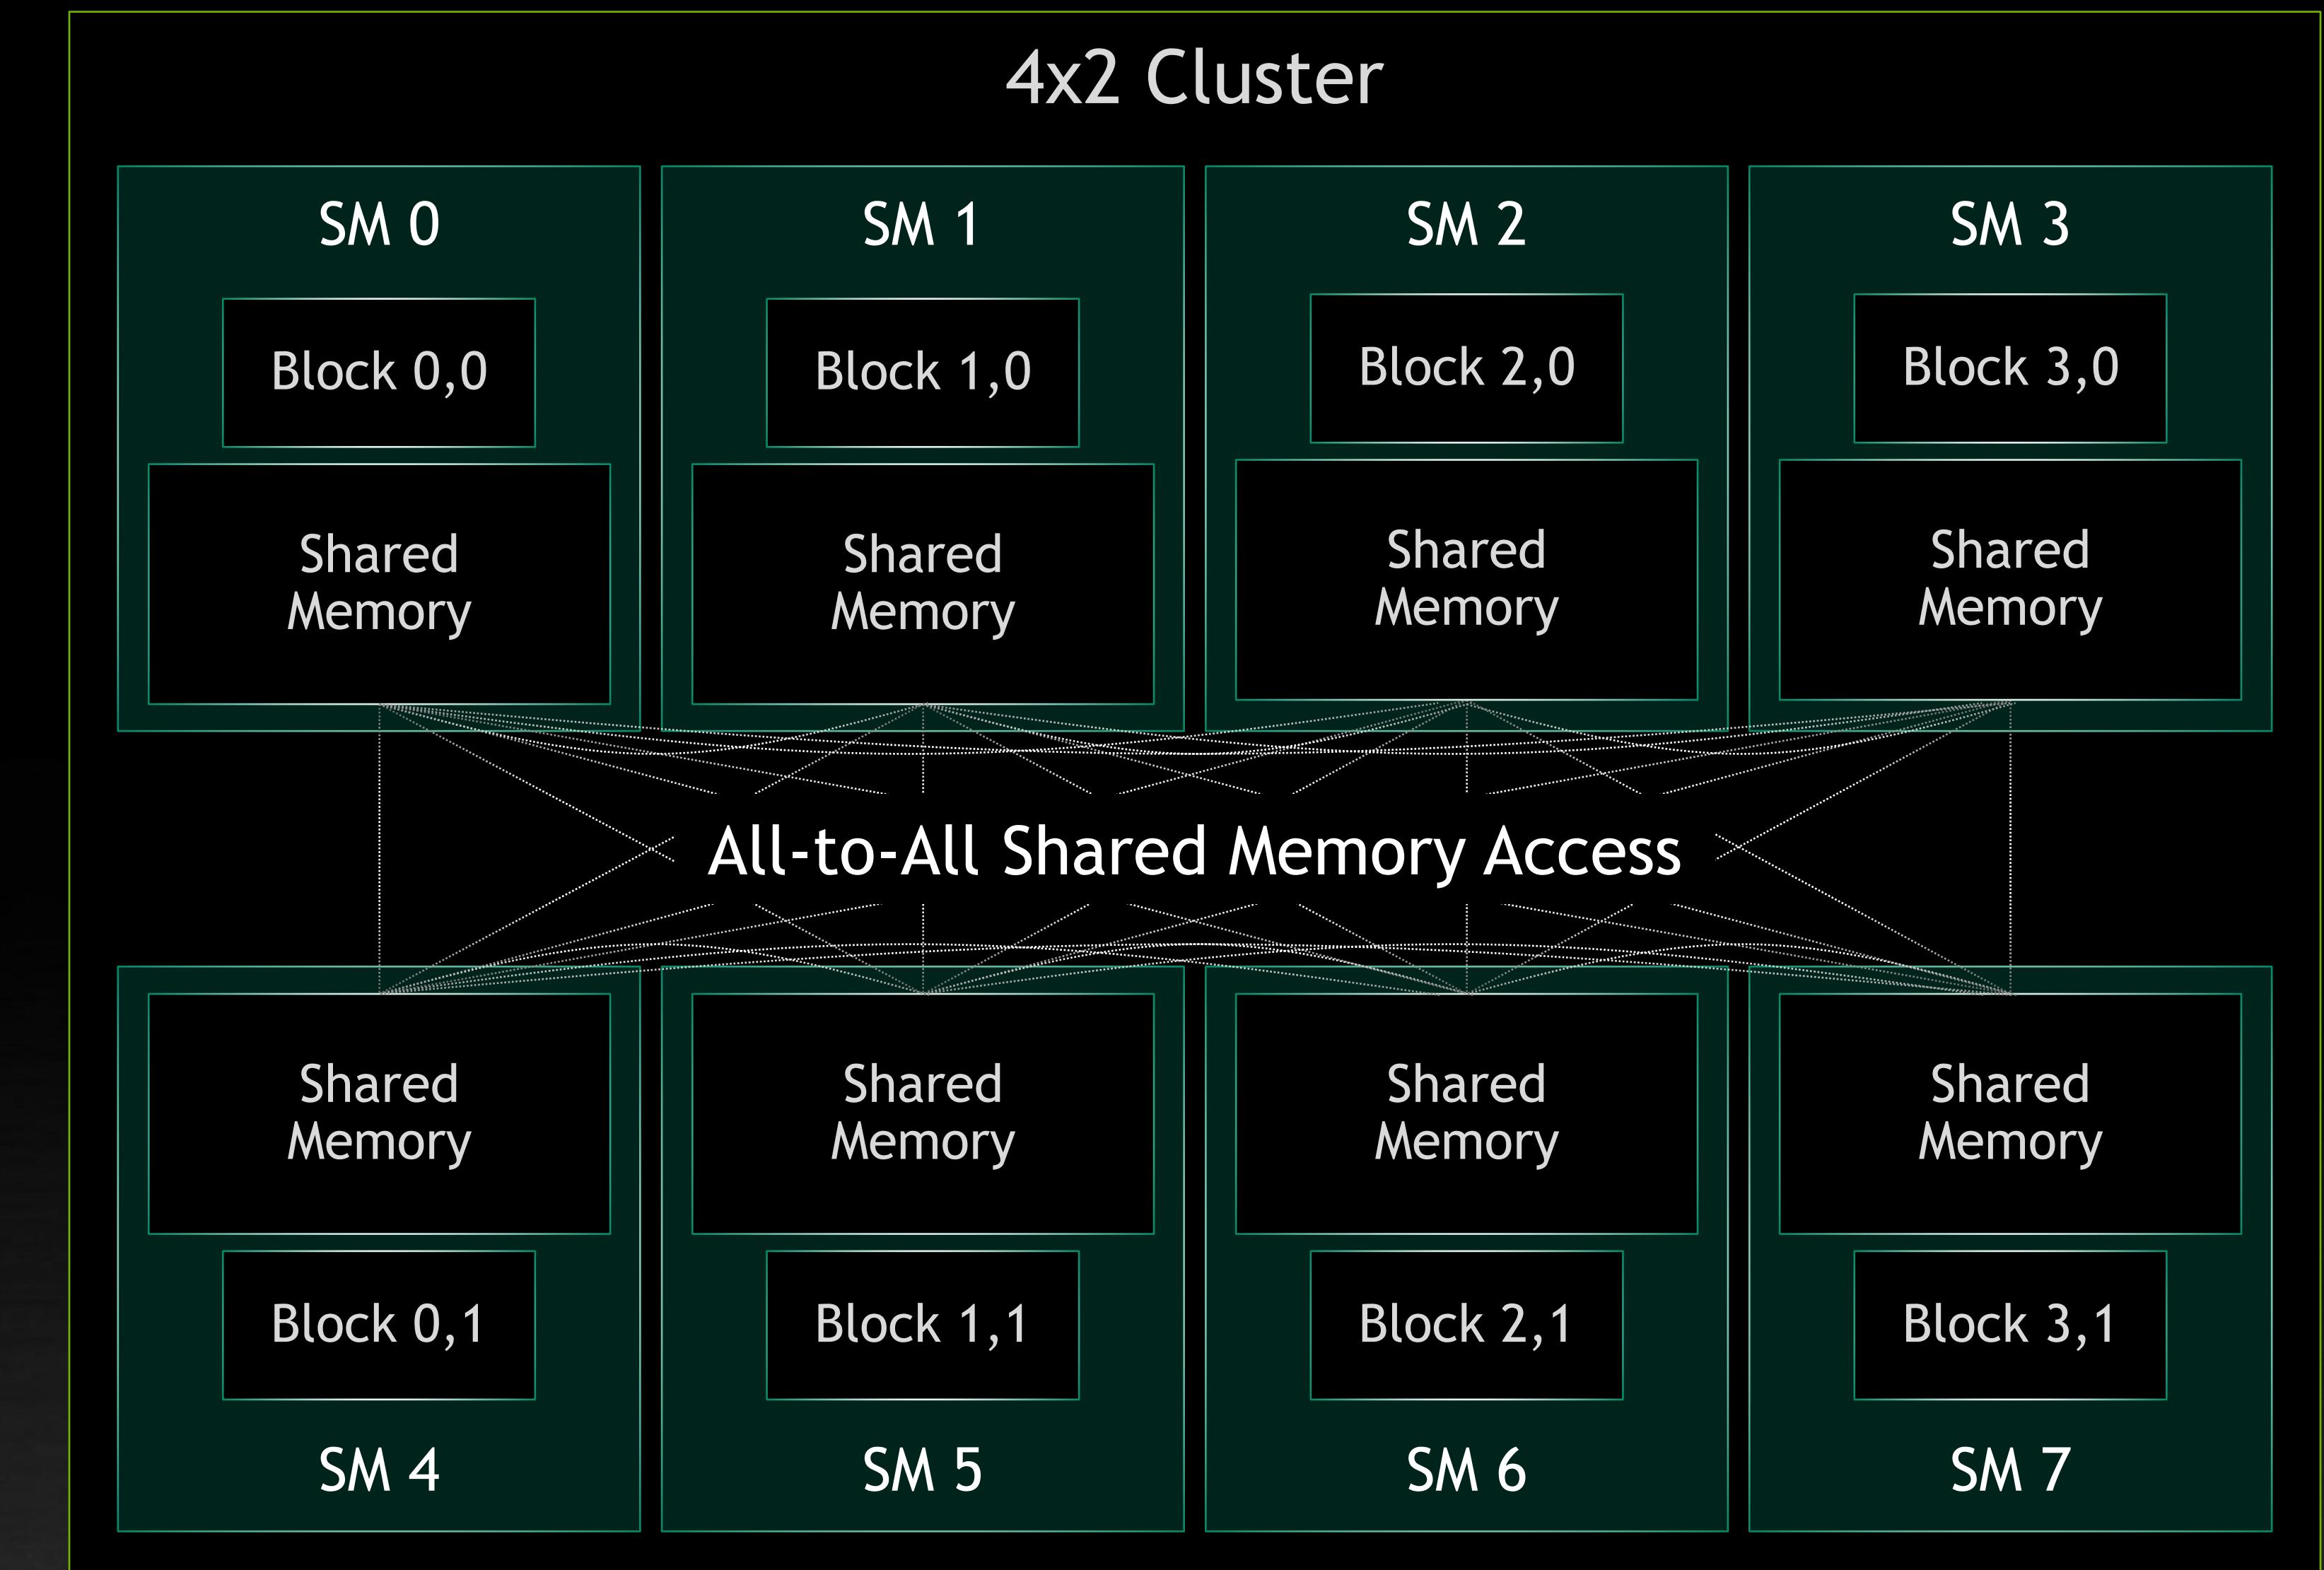
DSMEM is laid out as a **Partitioned Global Address Space**

DSMEM may be indexed by block **rank** - a 1D expansion of cluster index - or by (xyz) index in cluster

```
// Example of getting DSMEM pointer from cluster neighbours
__shared__ int x;

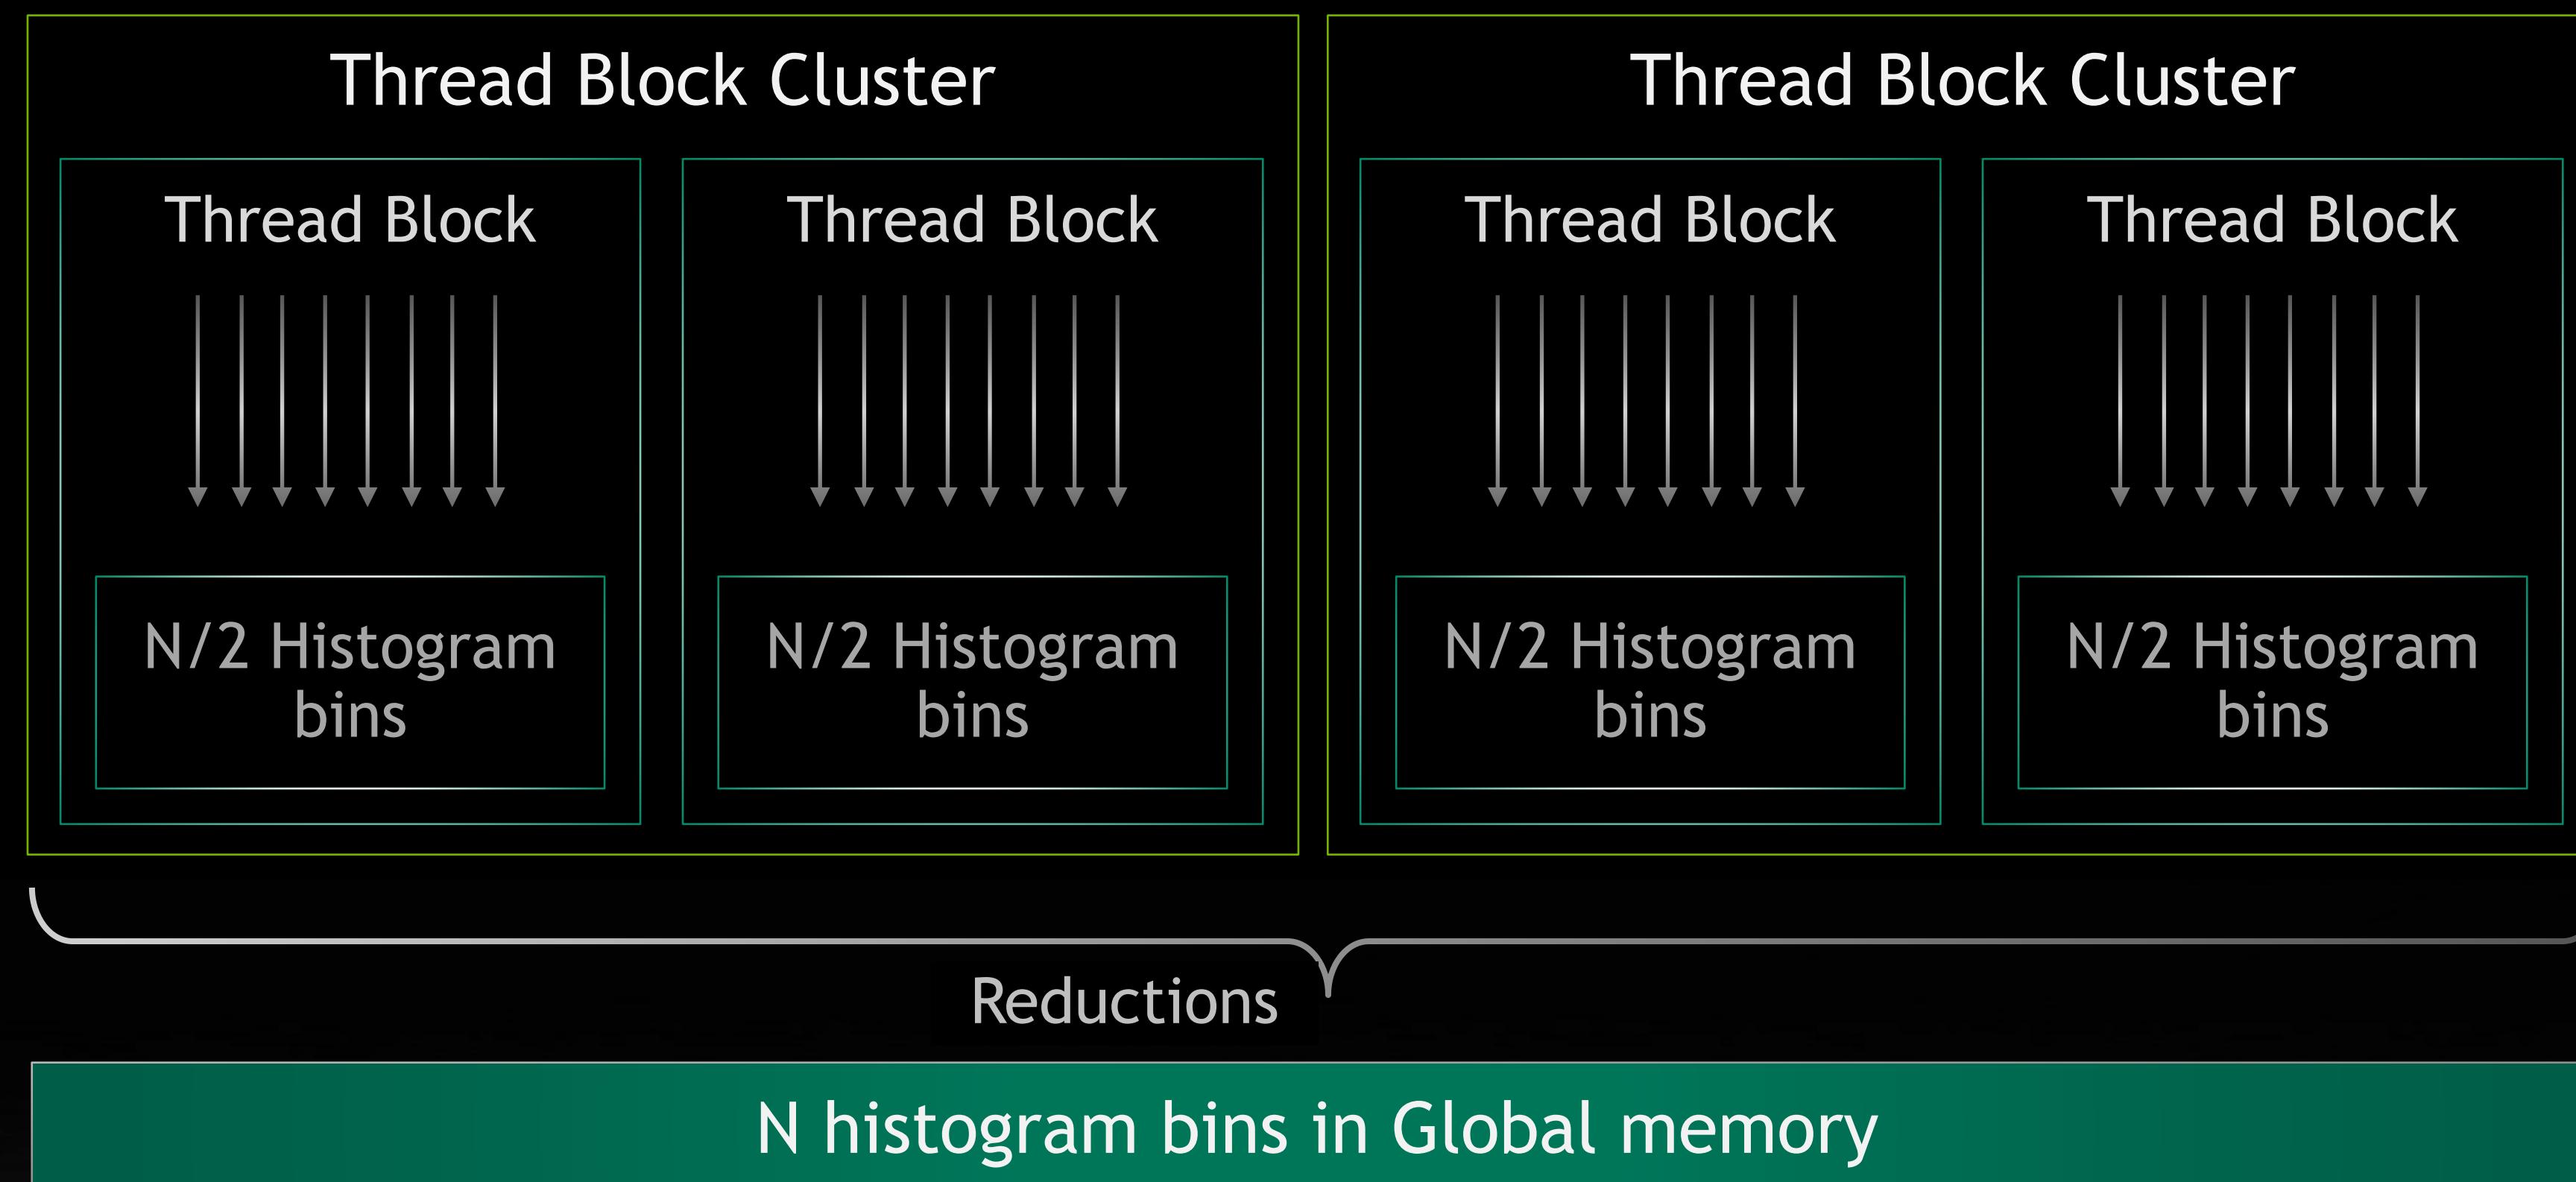
cg::cluster_group = this_cluster();
int *x_ptr = cluster.map_shared_rank(&x, 2);
*x_ptr = 5;

cluster_barrier_arrive();      // Hardware-managed barrier
cluster_barrier_wait();
```



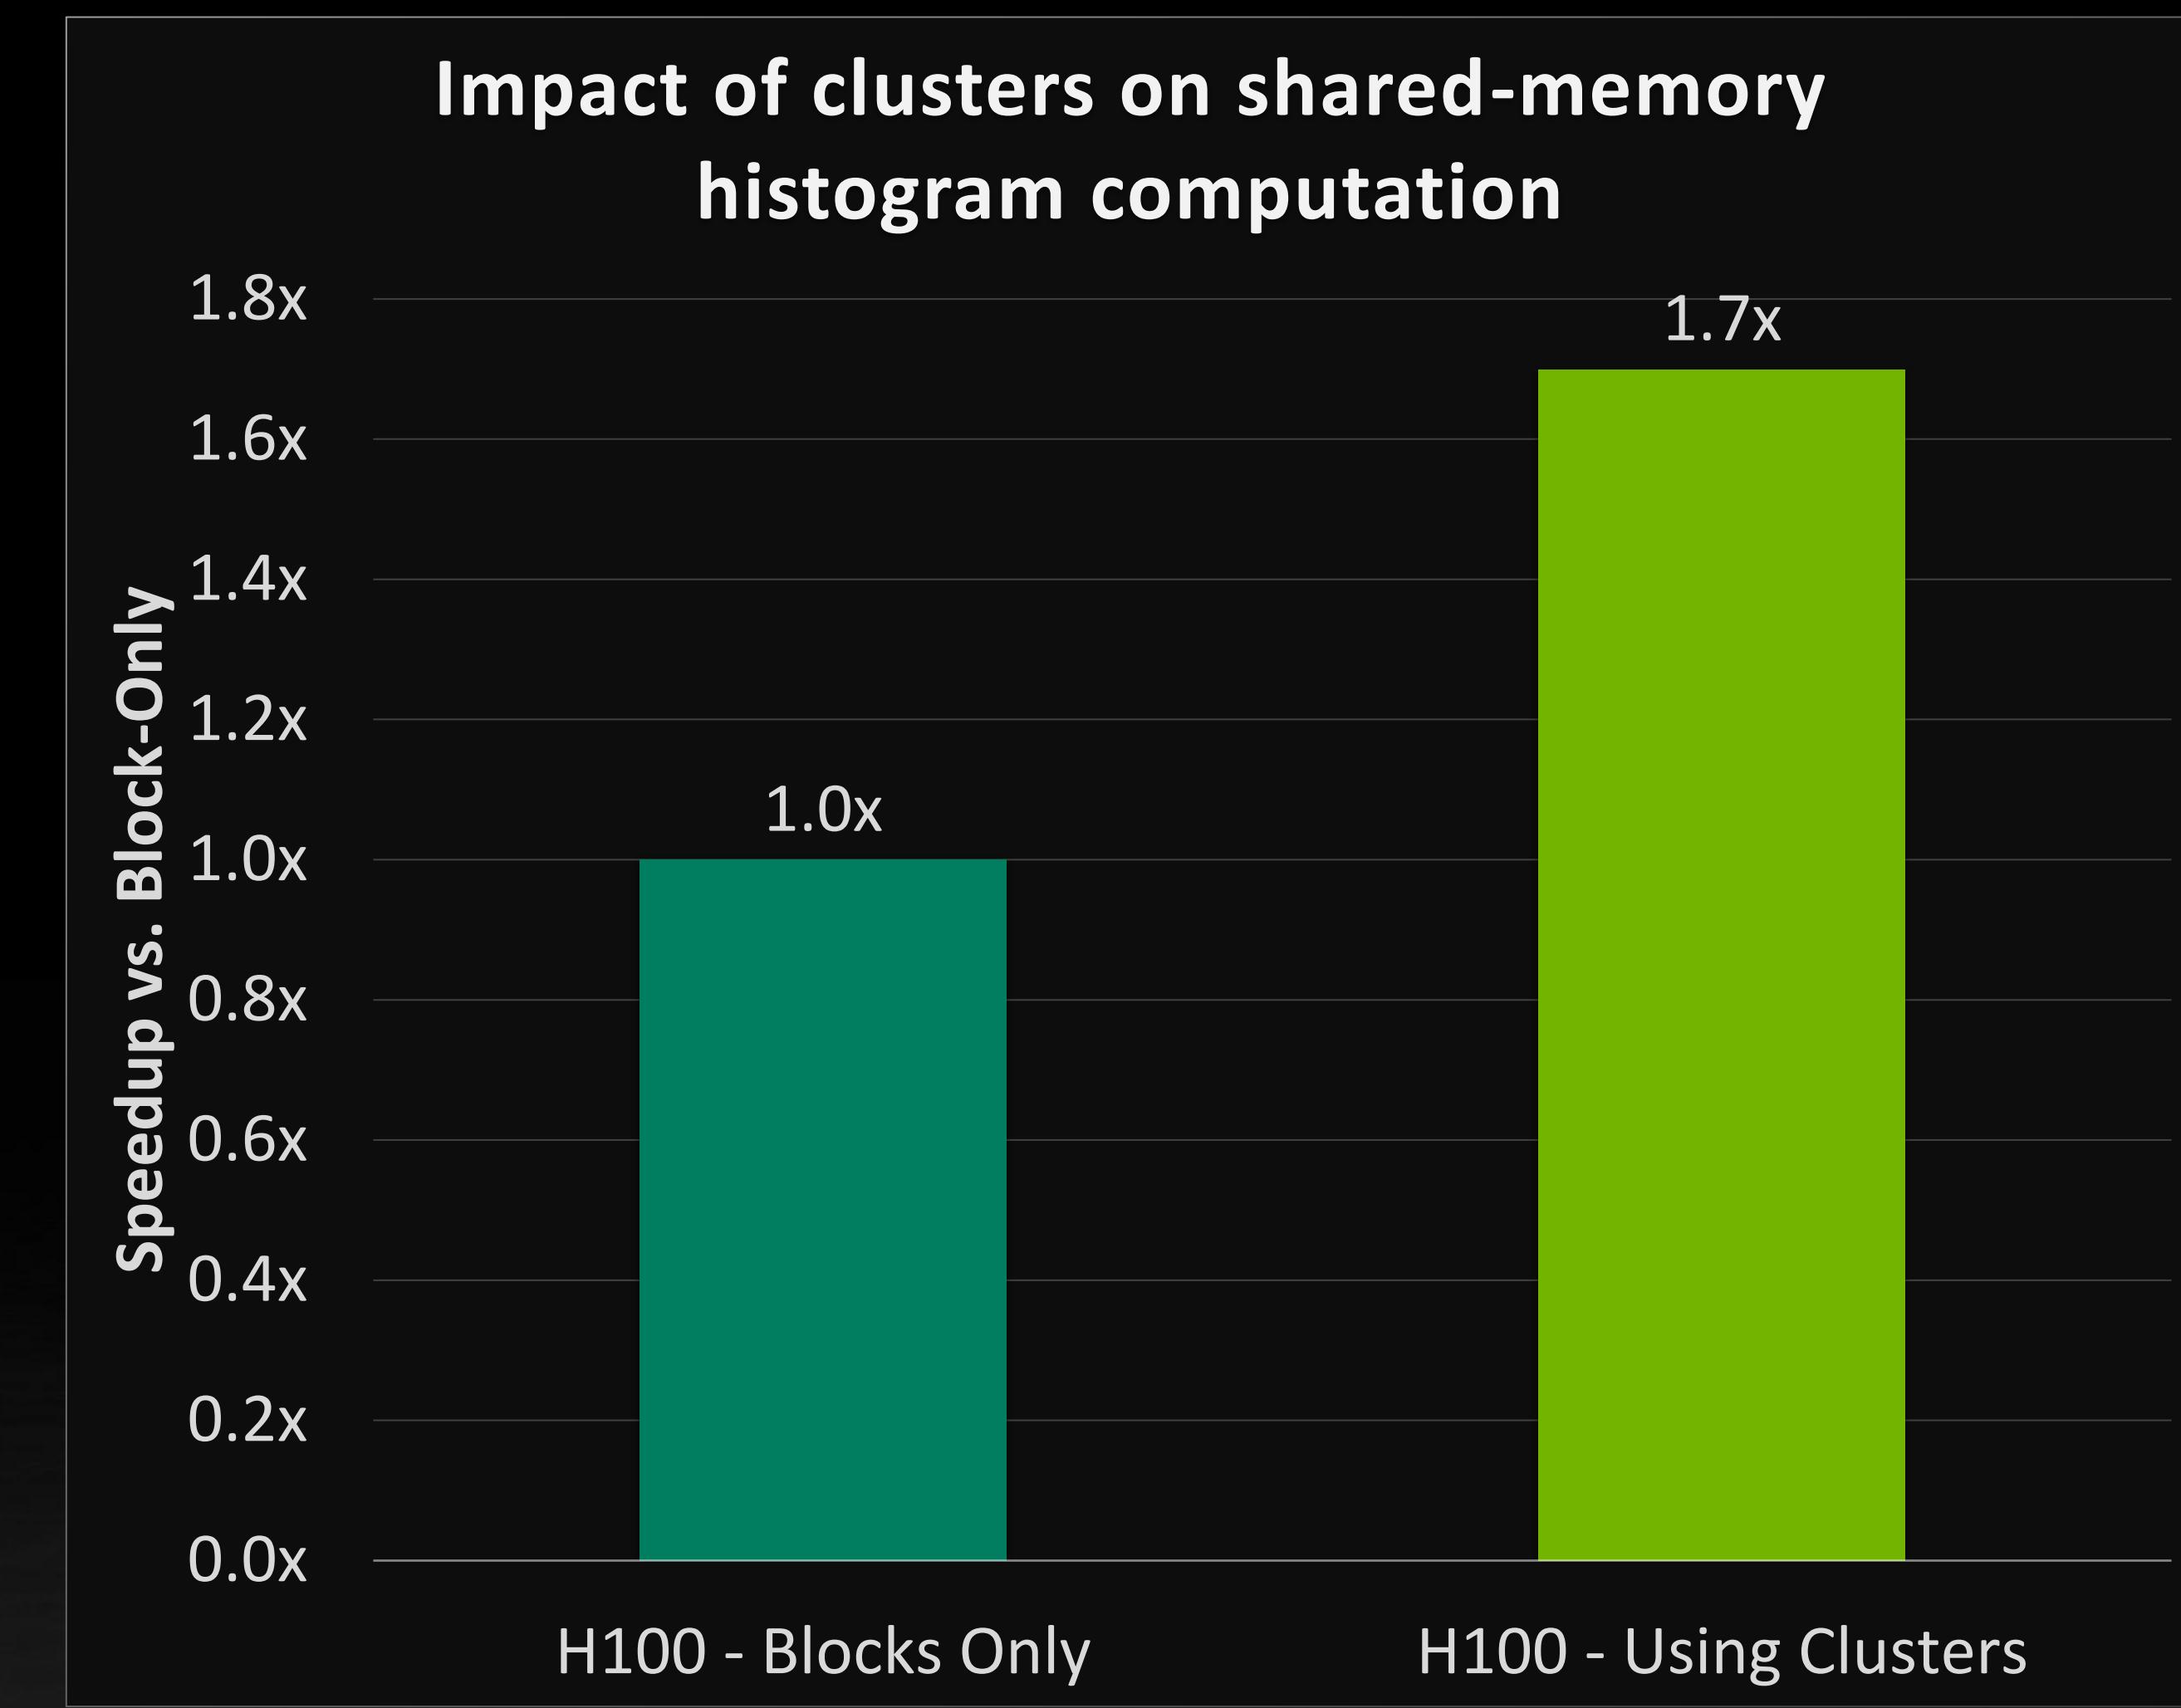
Full load/store/atomic access to all shared memory
between blocks within a cluster

EXAMPLE: HIERARCHICAL HISTOGRAM USING CLUSTER DSMEM



Histograms in CUDA are typically computed in shared memory, followed by reductions in global memory.

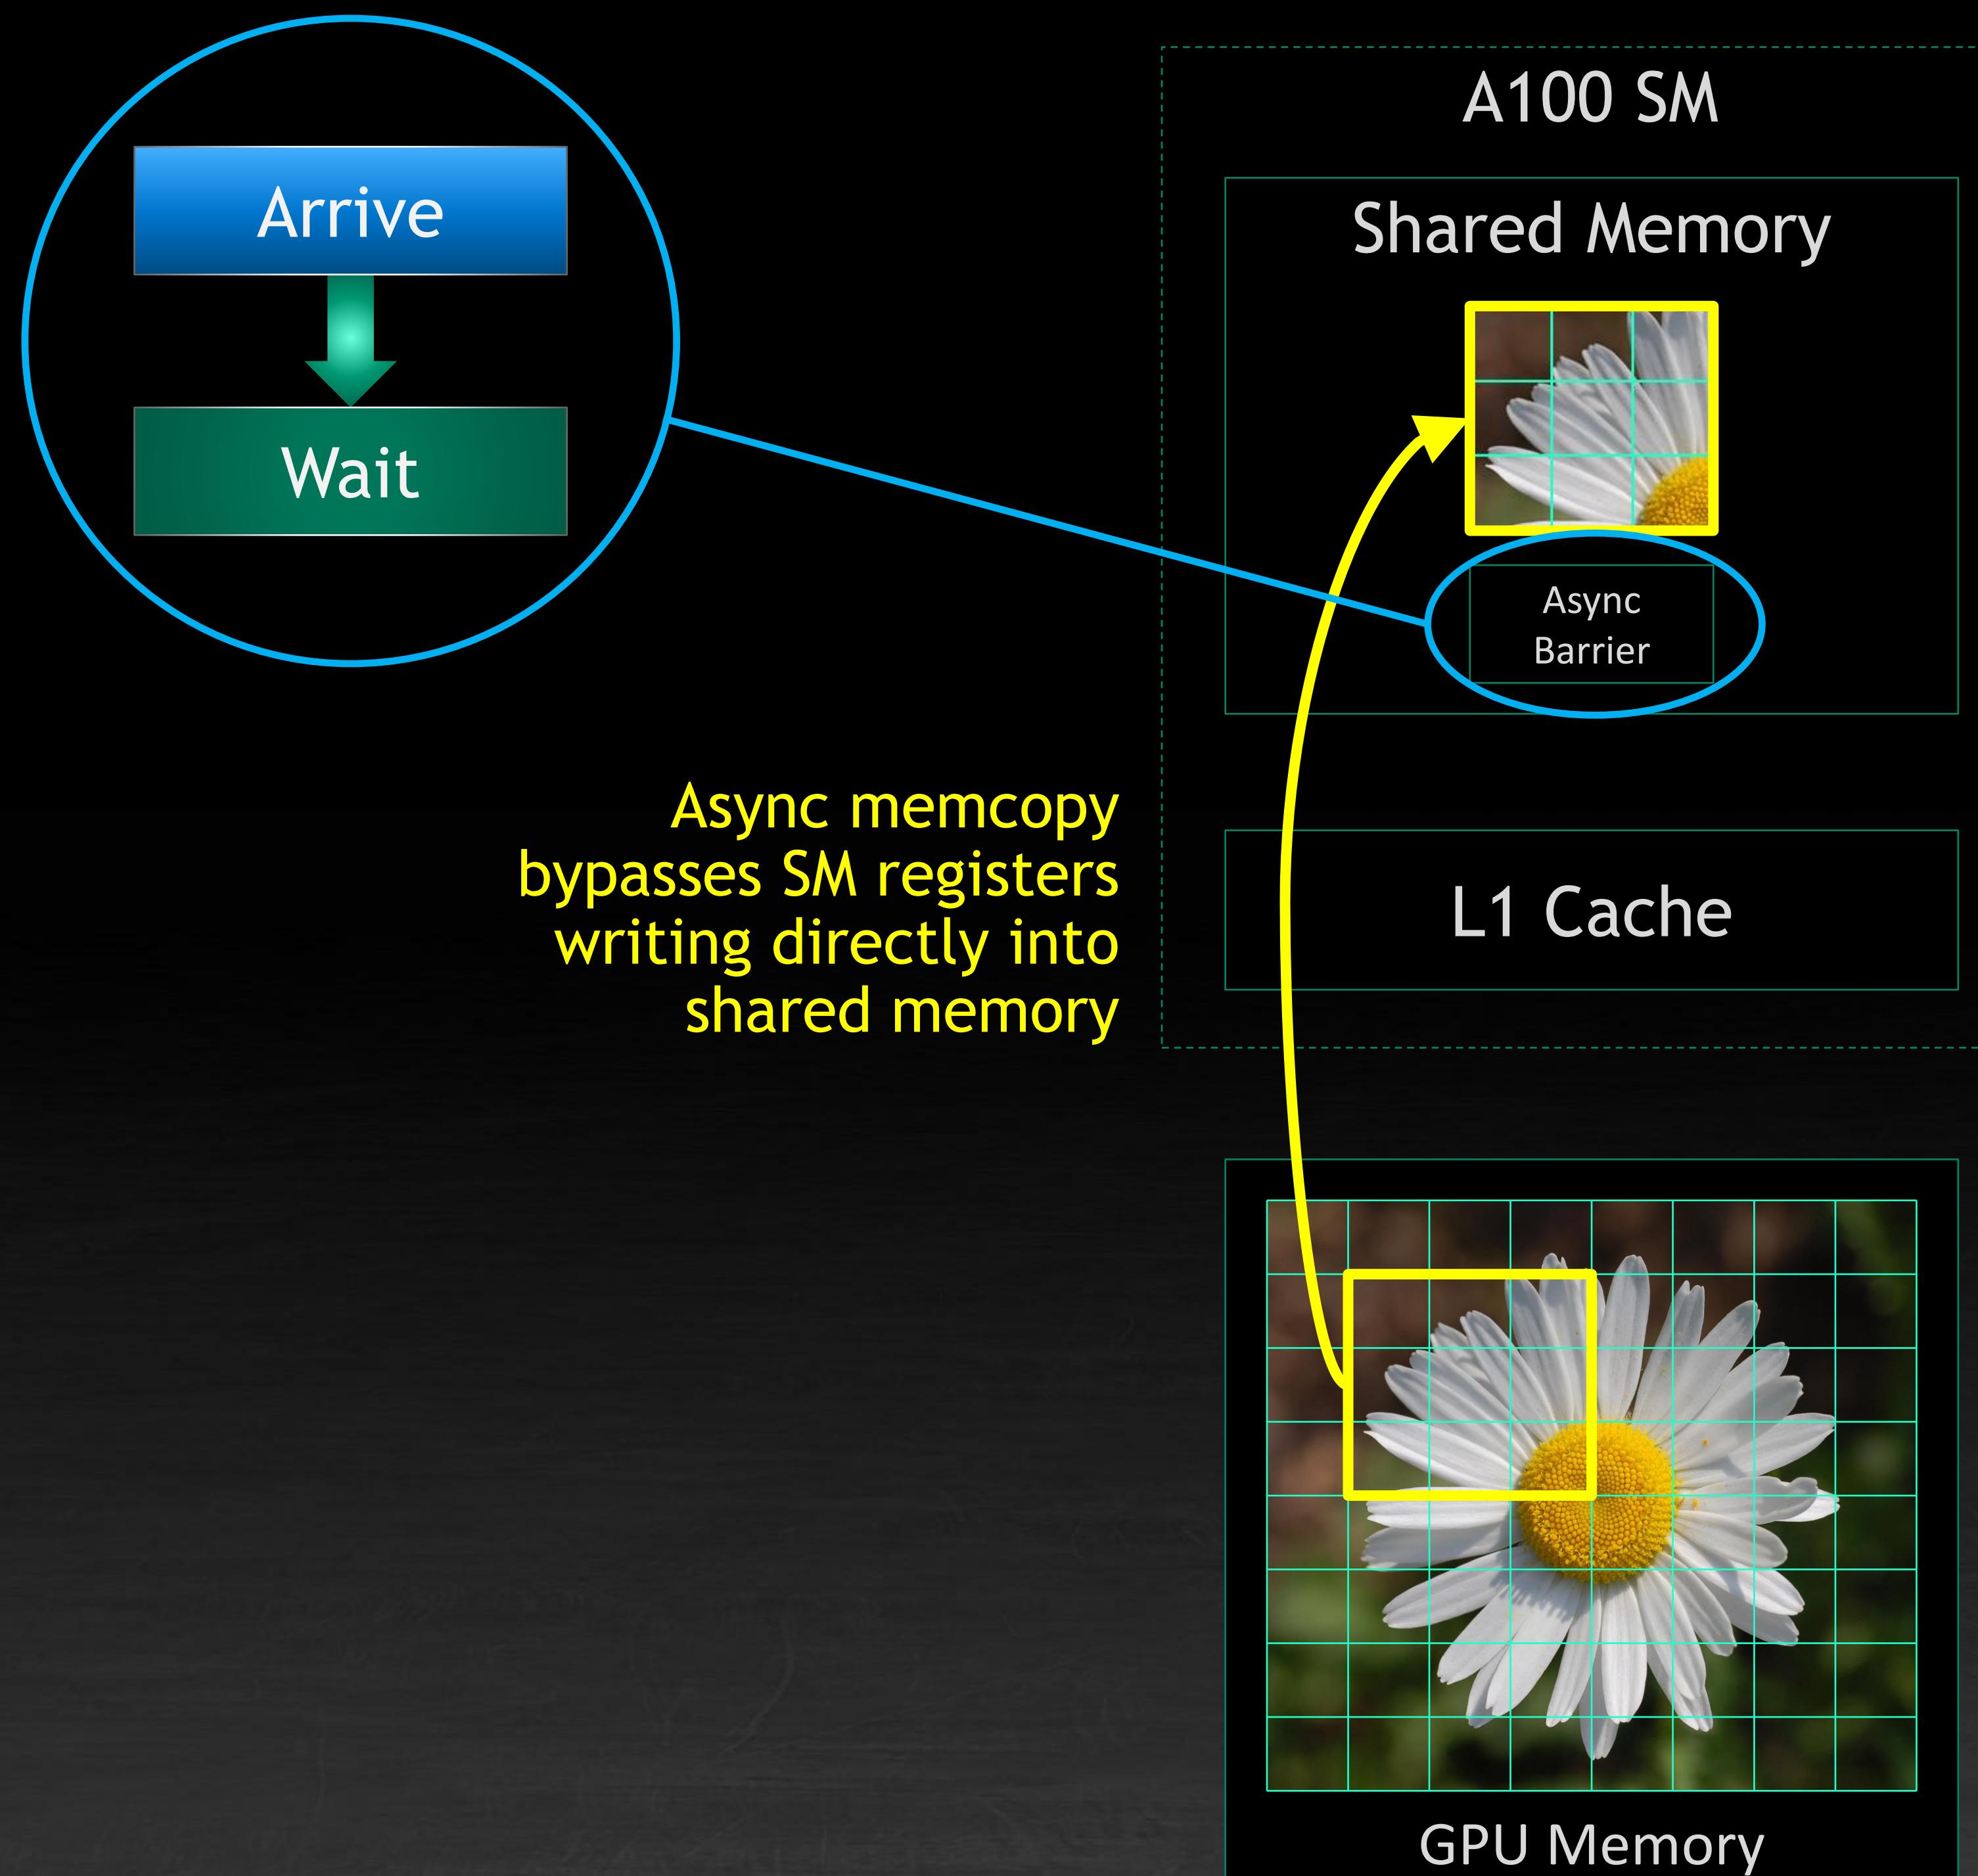
For large histograms, shared memory capacity of a single block is not sufficient.



75K Histogram bins (300KB) fit in distributed shared memory of 2-block clusters → 37.5K (150KB) per thread block

ASYNCHRONOUS COPY TO SHARED MEMORY

cuda::memcpy_async



cuda::memcpy_async()

Bypasses SM registers to write directly to shared
Uses **Asynchronous Barrier** to signal completion

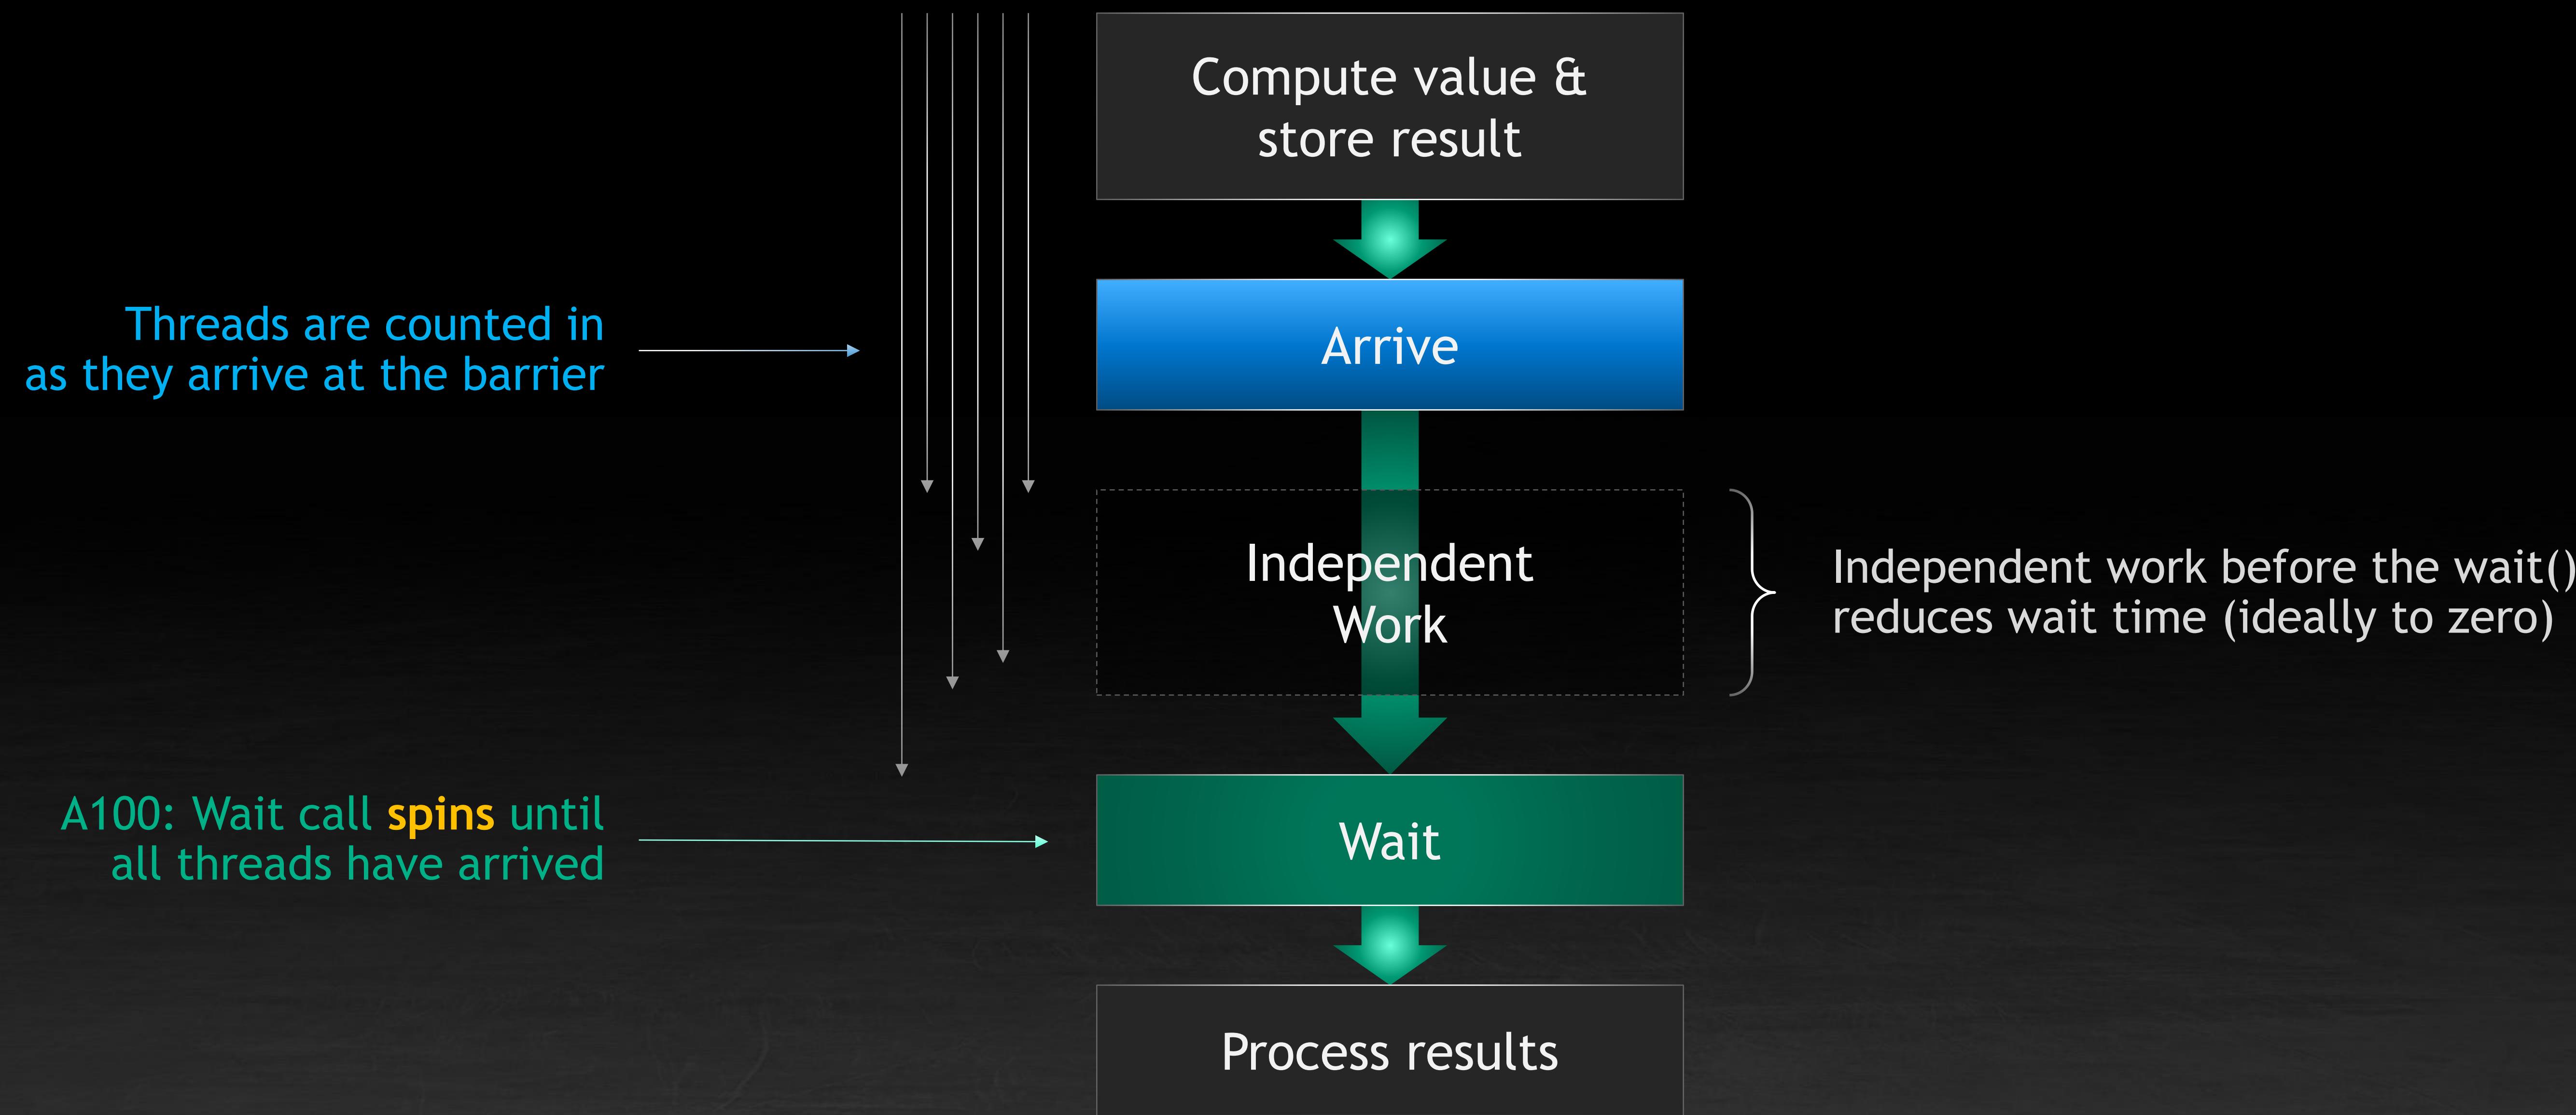
Threads work together to copy data in parallel

```
// Async block-wide memcpy using Cooperative Groups
cg::thread_block tb = cg::this_thread_block();
size_t copy_count, index = 0;
while (index < elementsPerThreadBlock) {
    cg::memcpy_async(
        tb,
        local_smem,
        elementsInShared,
        global_data + index,
        elementsPerThreadBlock - index);
    copy_count = min(elementsInShared,
                    elementsPerThreadBlock - index);
    cg::wait(tb);

    // ... use data copied into shared memory ...
    index += copyCount;
}
```

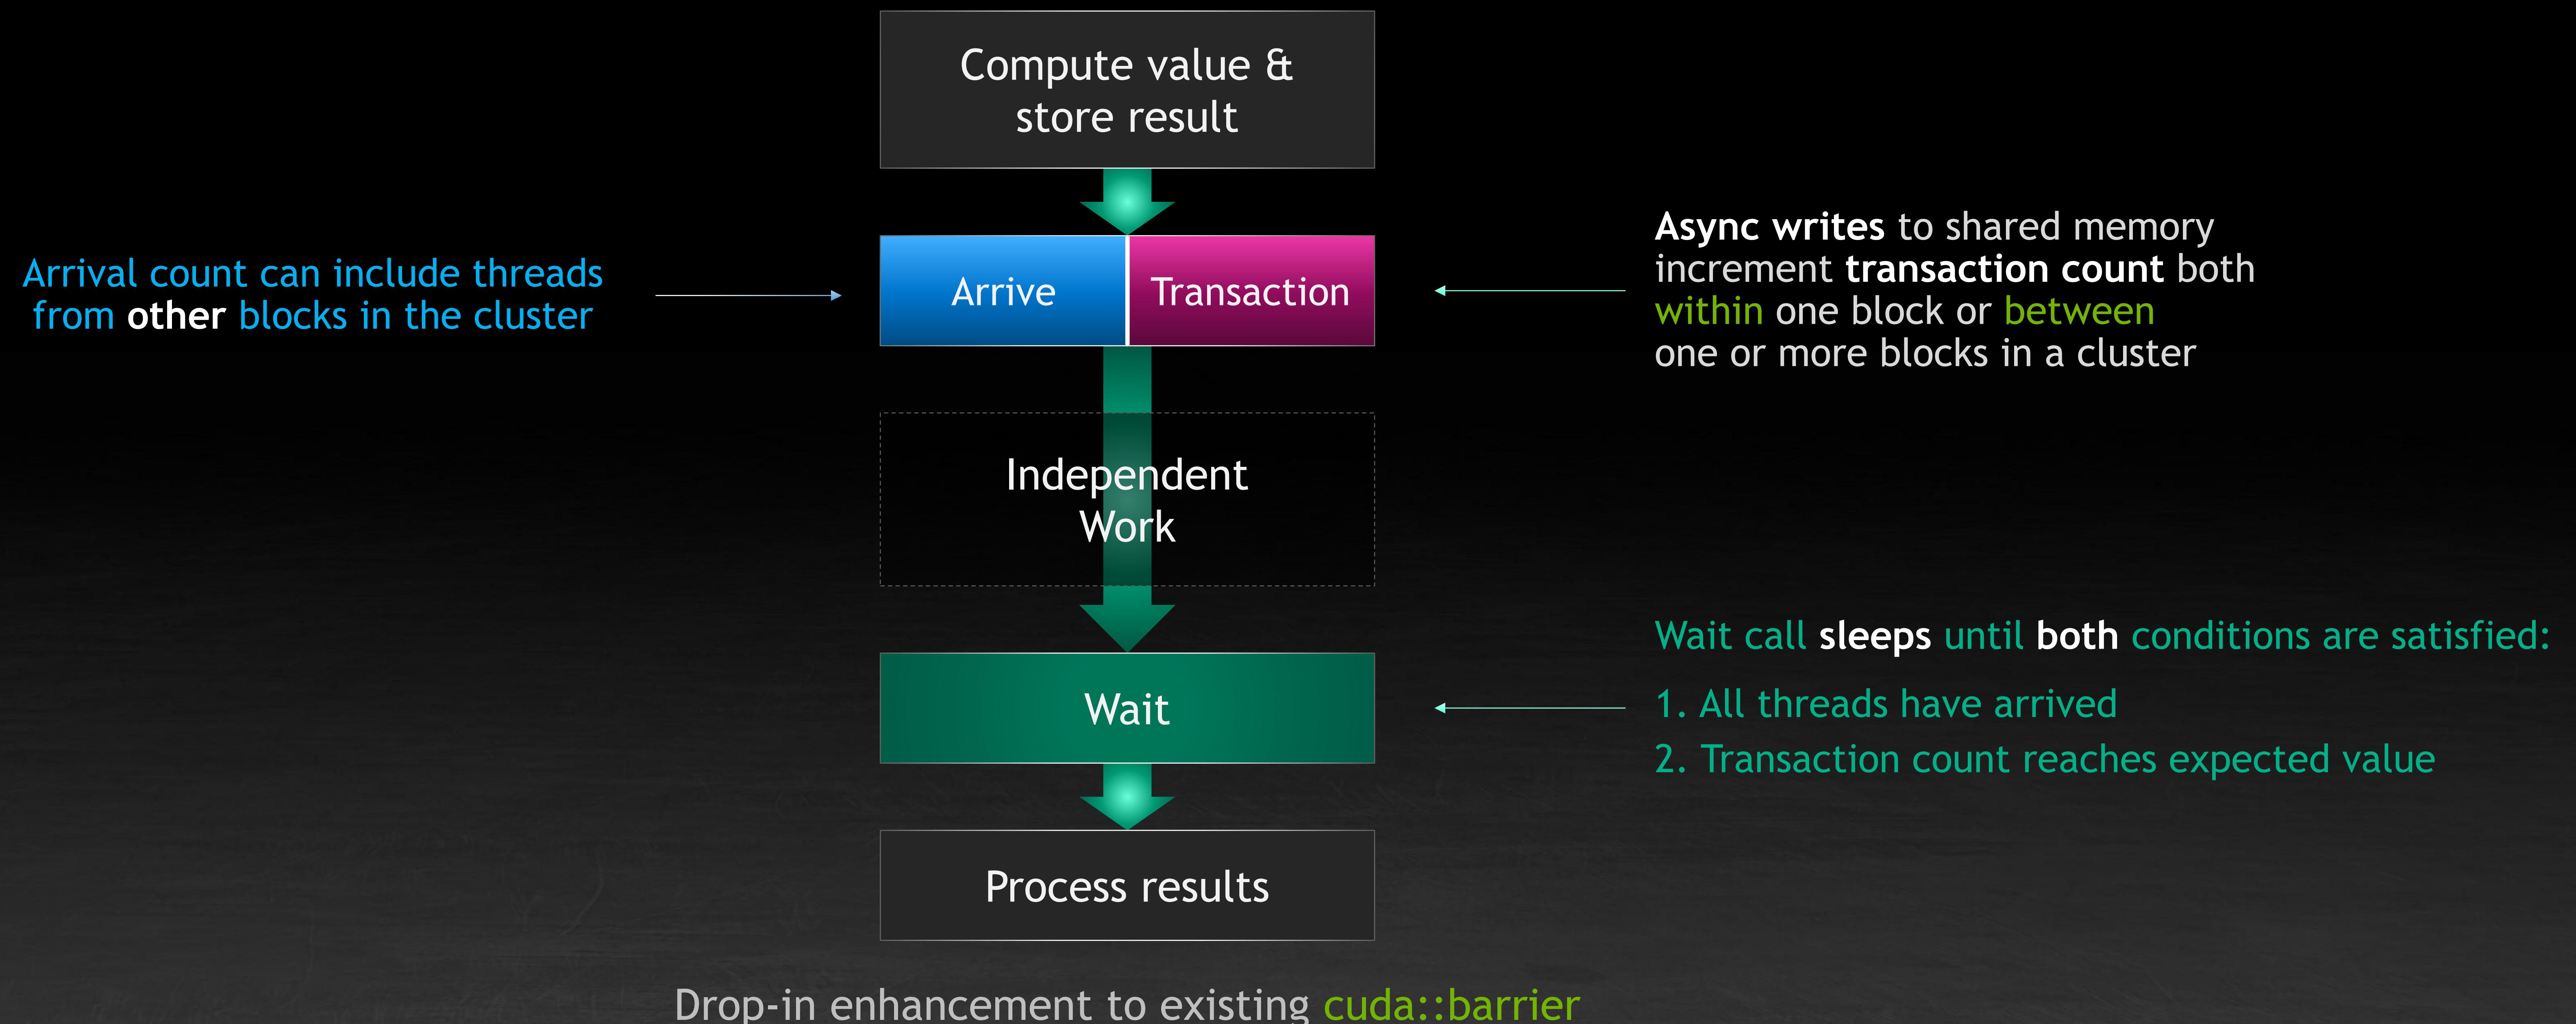
ASYNCHRONOUS BARRIERS

cuda::barrier



ASYNCHRONOUS TRANSACTION BARRIERS

Data can now trigger a barrier as well as just threads



ASYNCHRONOUS ONE-SIDED MEMORY COPIES

`cuda::memcpy_async` within & between blocks in a cluster automatically updates transaction count without handshake

Data arrival increments transaction count

One-sided communication within cluster

Accepts async store **and** atomic operations

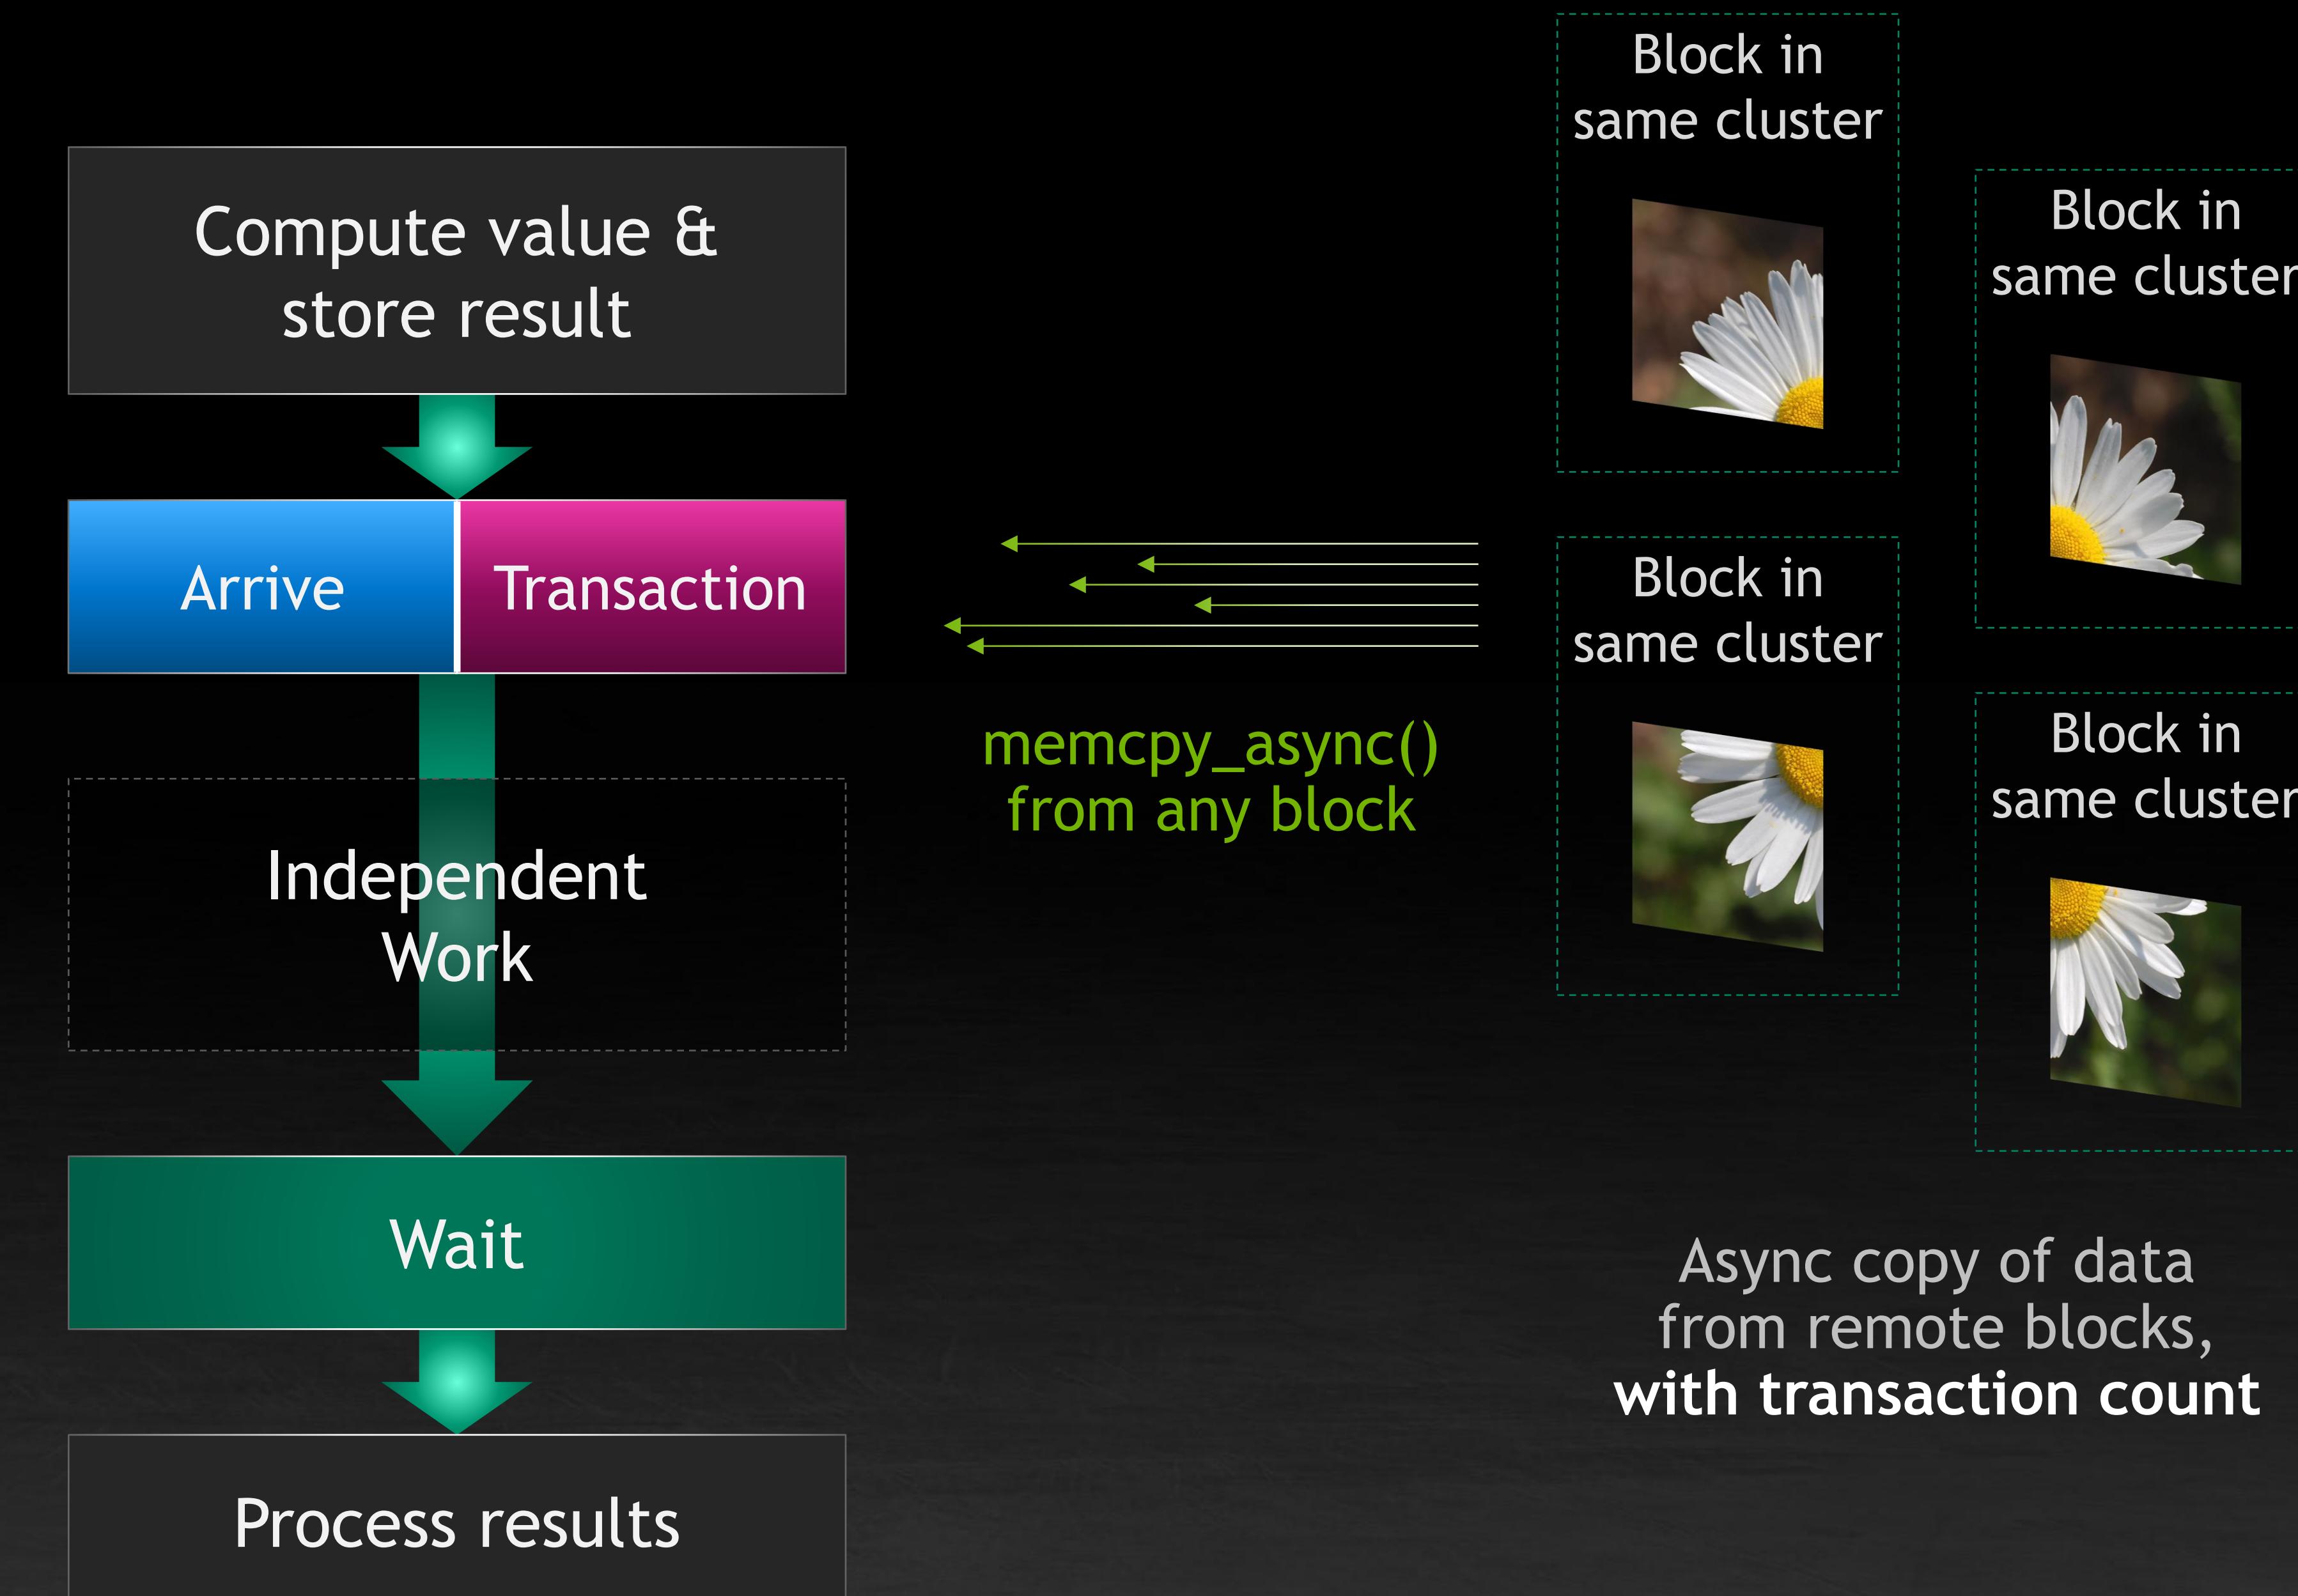
Highly efficient control flow

Self-synchronizing: no memory fence required to update transaction count

7x faster than { data + fence + flag } handshake

Receiver threads **sleep** on wait (instead of busy-waiting) until data arrives

Fast wakeup signal directly from barrier



ASYNCHRONOUS ONE-SIDED COMMUNICATION

No-handshake messaging & data transfer within a block or a cluster

```
cuda::atomic_ref ref;           // Initialization not shown
cuda::barrier admirer;          // Initialization not shown

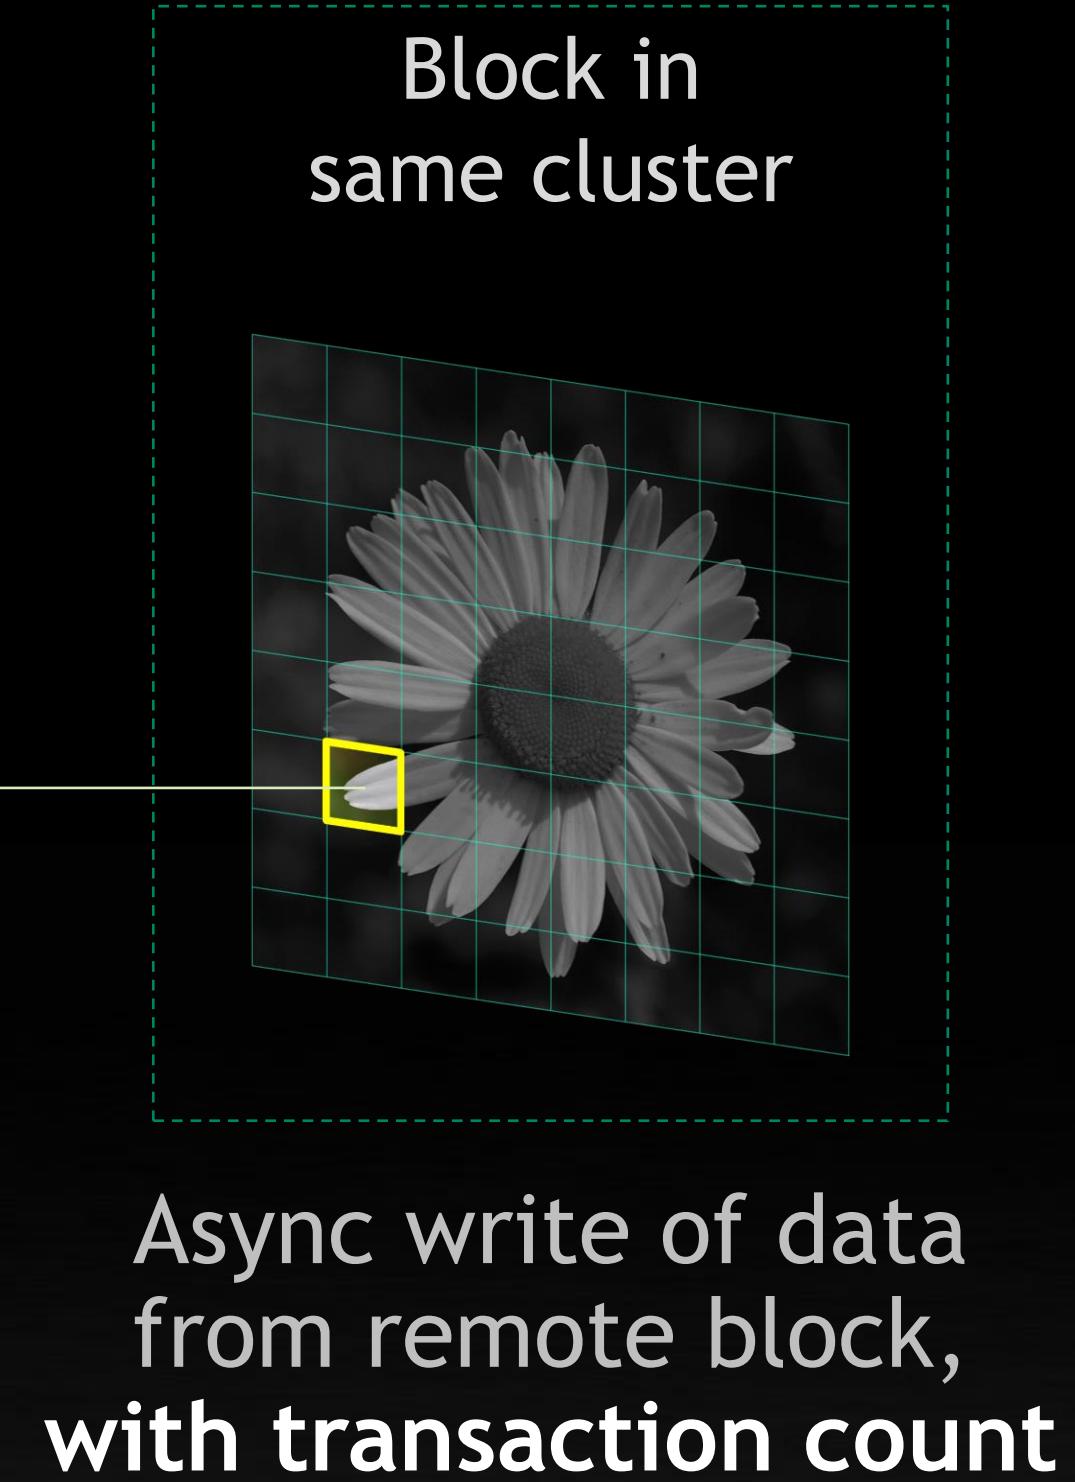
// Remote store with transaction count
void atomic_store_async(ref, value, barrier);

// Remote atomic reduction operations
void atomic_add_async(ref, value, barrier);
void atomic_min_async(ref, value, barrier);
void atomic_max_async(ref, value, barrier);
void atomic_inc_async(ref, value, barrier);
void atomic_dec_async(ref, value, barrier);
void atomic_and_async(ref, value, barrier);
void atomic_or_async(ref, value, barrier);
void atomic_xor_async(ref, value, barrier);

// Note: CAS is not supported as it is not a reduction operation
```

For these individual async store/atomic operations,
initialise barrier to expected message size before sending

Transactions
can also be
individual
async stores
and atomics



TENSOR MEMORY ACCELERATOR UNIT (TMA) FOR ASYNC DATA MOVEMENT

Example: One-sided halo exchange within a cluster

Tensor Memory Accelerator Unit

Hardware-accelerated **bi-directional** bulk copy

Global \leftrightarrow Shared Memory

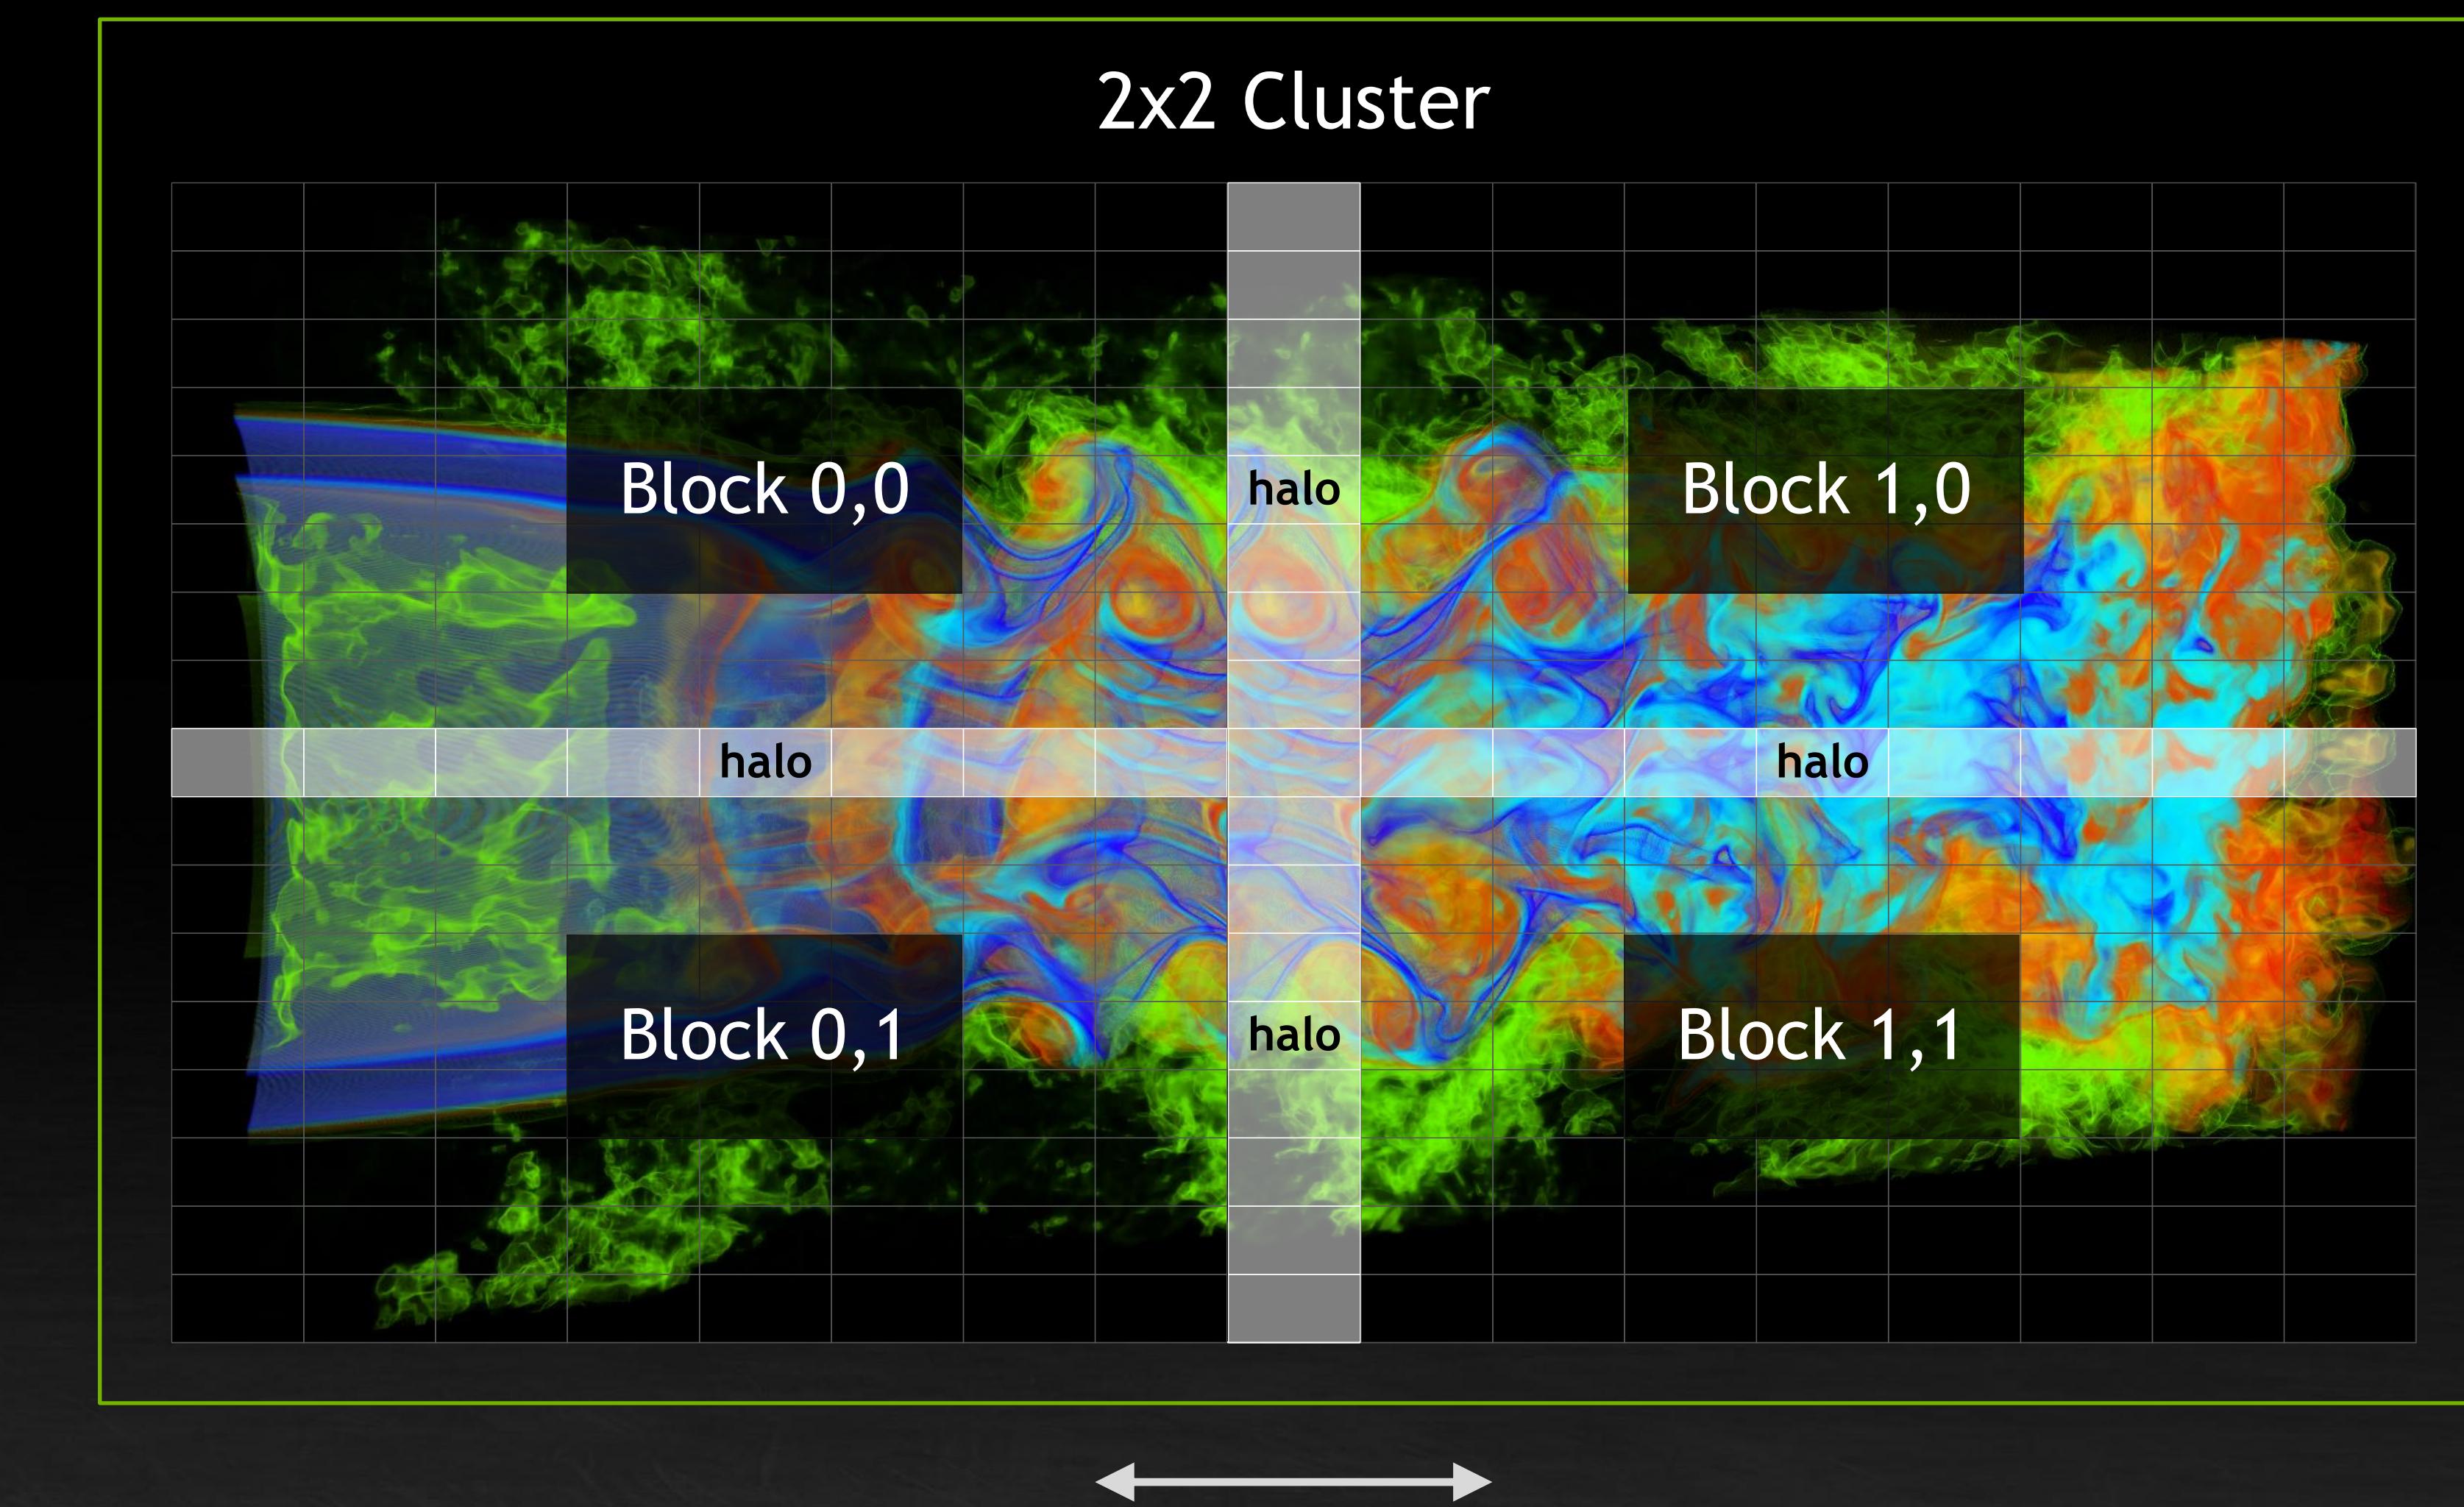
DSMEM \leftrightarrow DSMEM within cluster

Uses **Asynchronous Transaction Barrier** to track completion

Low-Latency One-Sided Data Transfer

Supports one-sided bulk copying within cluster

Supports element-wise reductions when copying to global memory



Self-synchronizing transactions enable 7x faster halo exchange between blocks in the same cluster

TENSOR MEMORY ACCELERATOR UNIT (TMA) FOR ASYNC DATA MOVEMENT

Tensor Memory Accelerator Unit

Hardware-accelerated **bi-directional** bulk copy

Global \leftrightarrow Shared Memory

DSMEM \leftrightarrow DSMEM within cluster

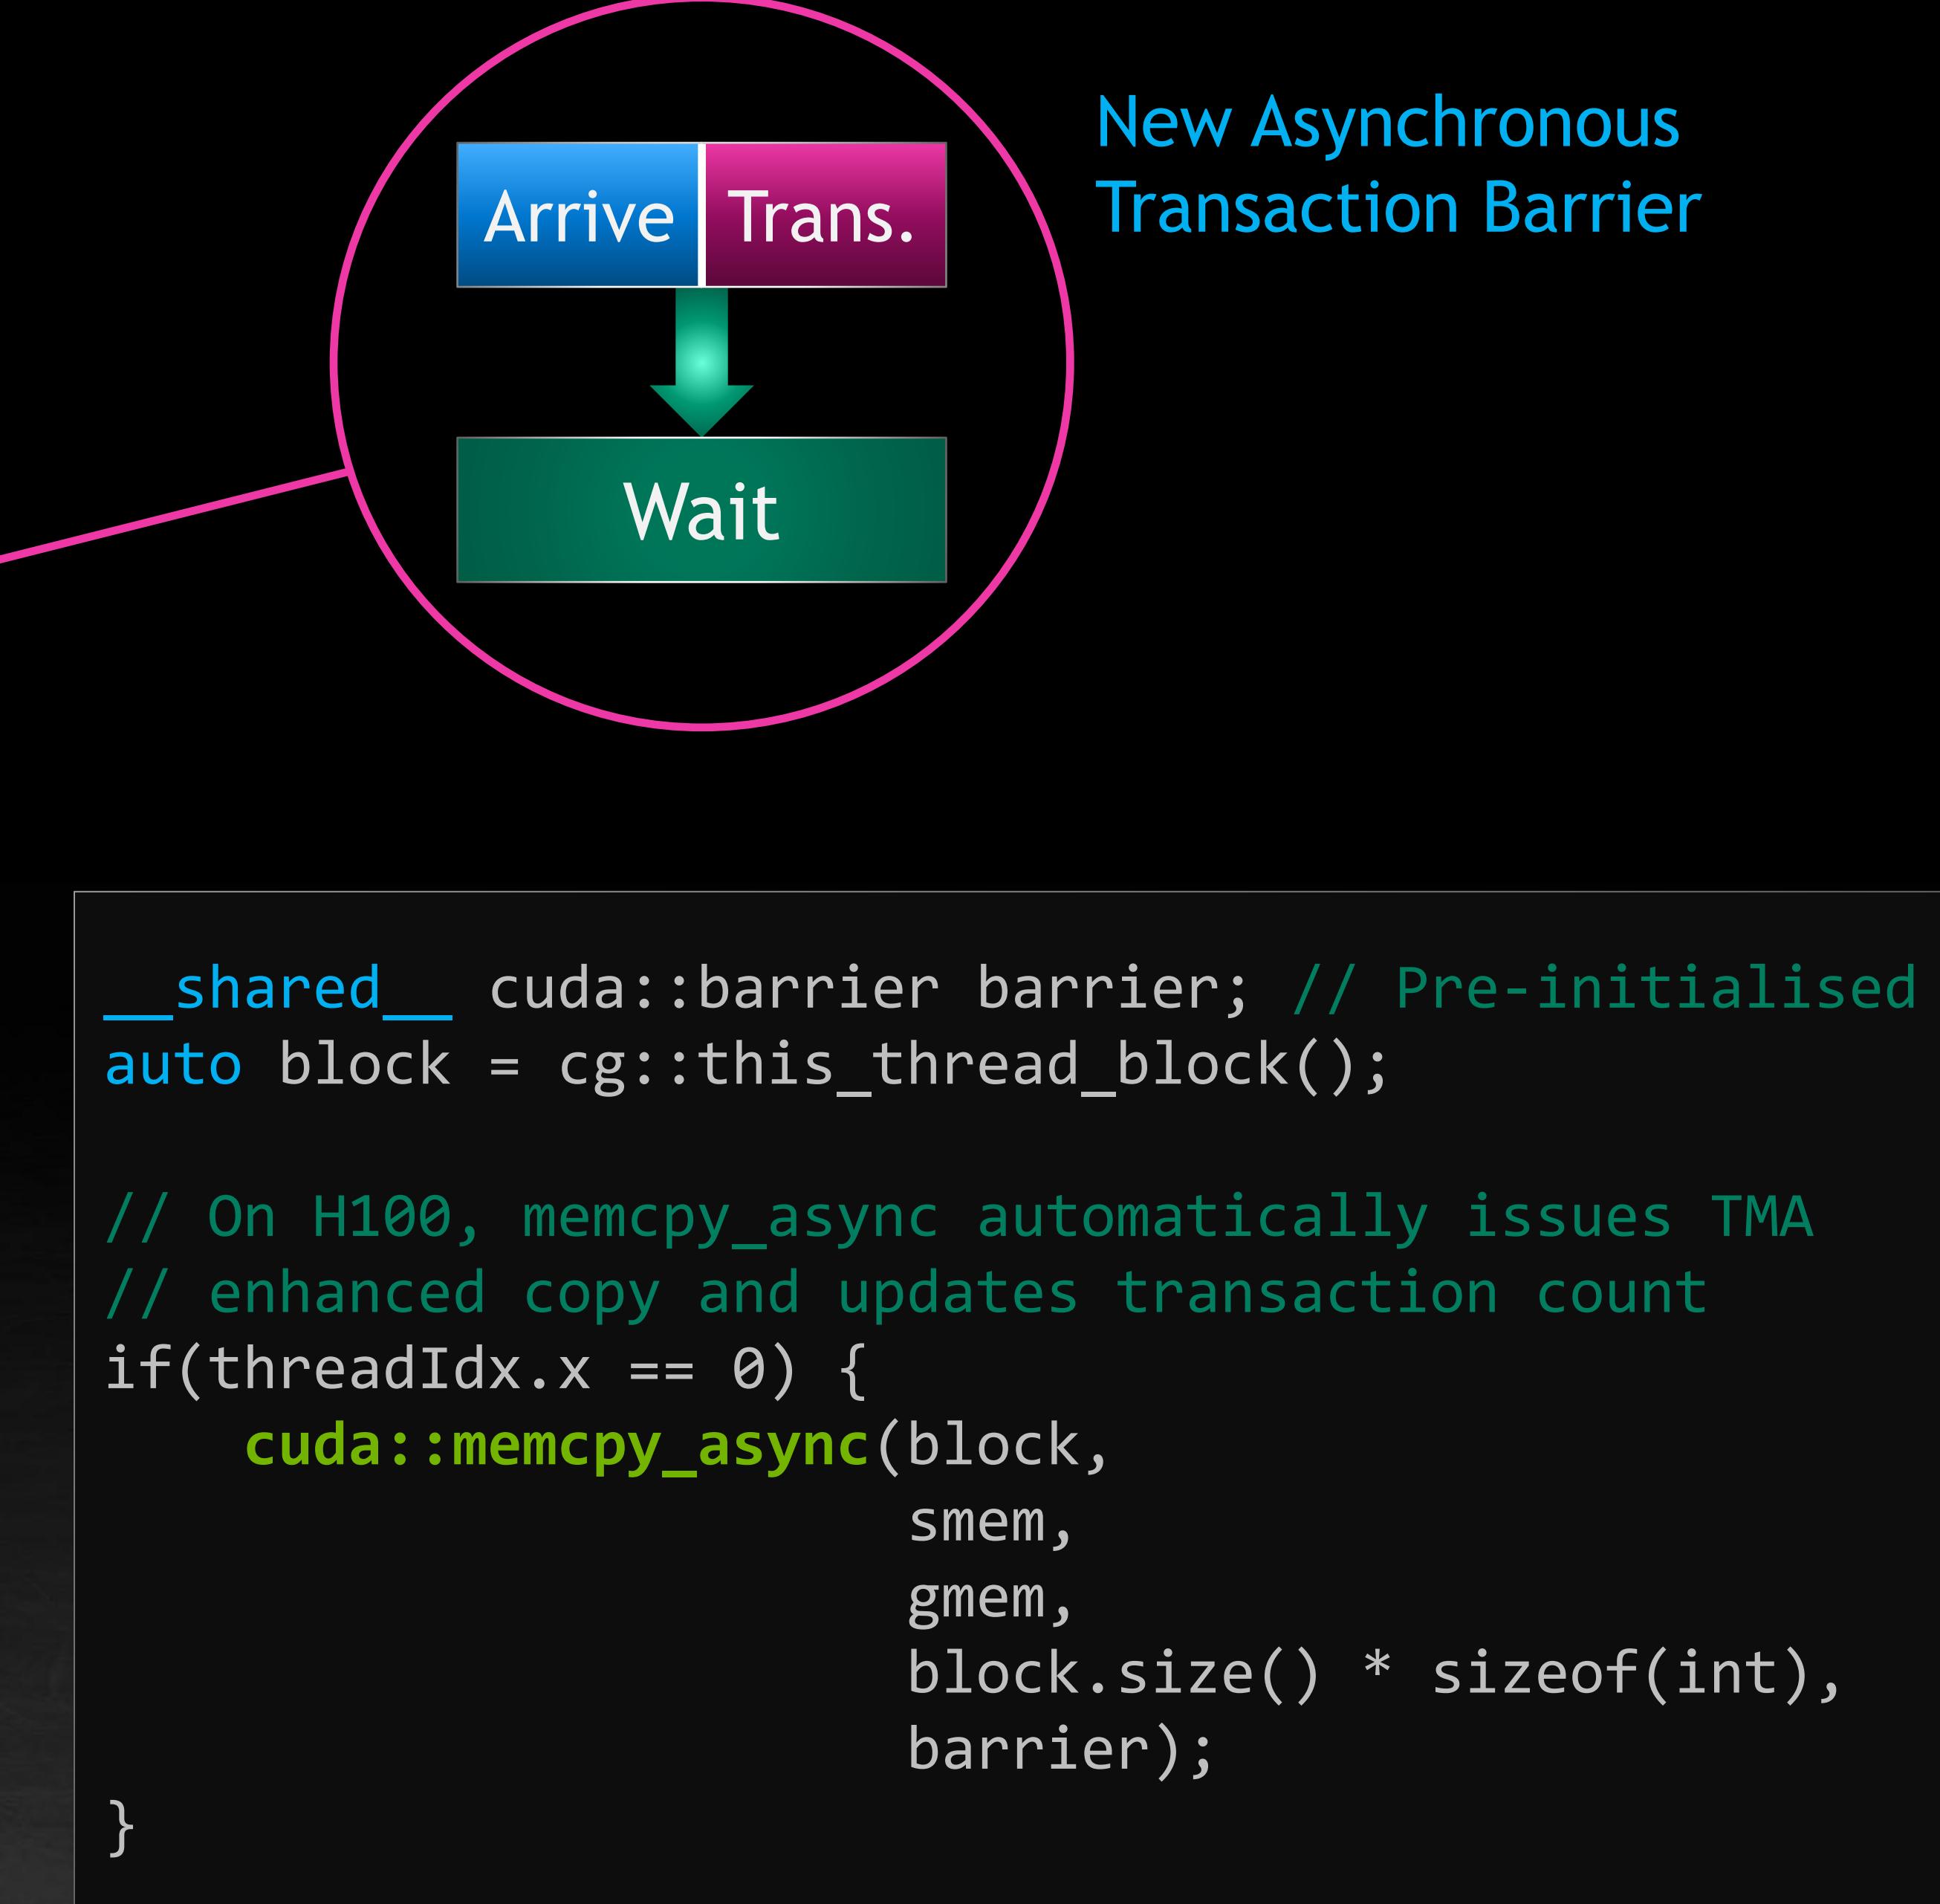
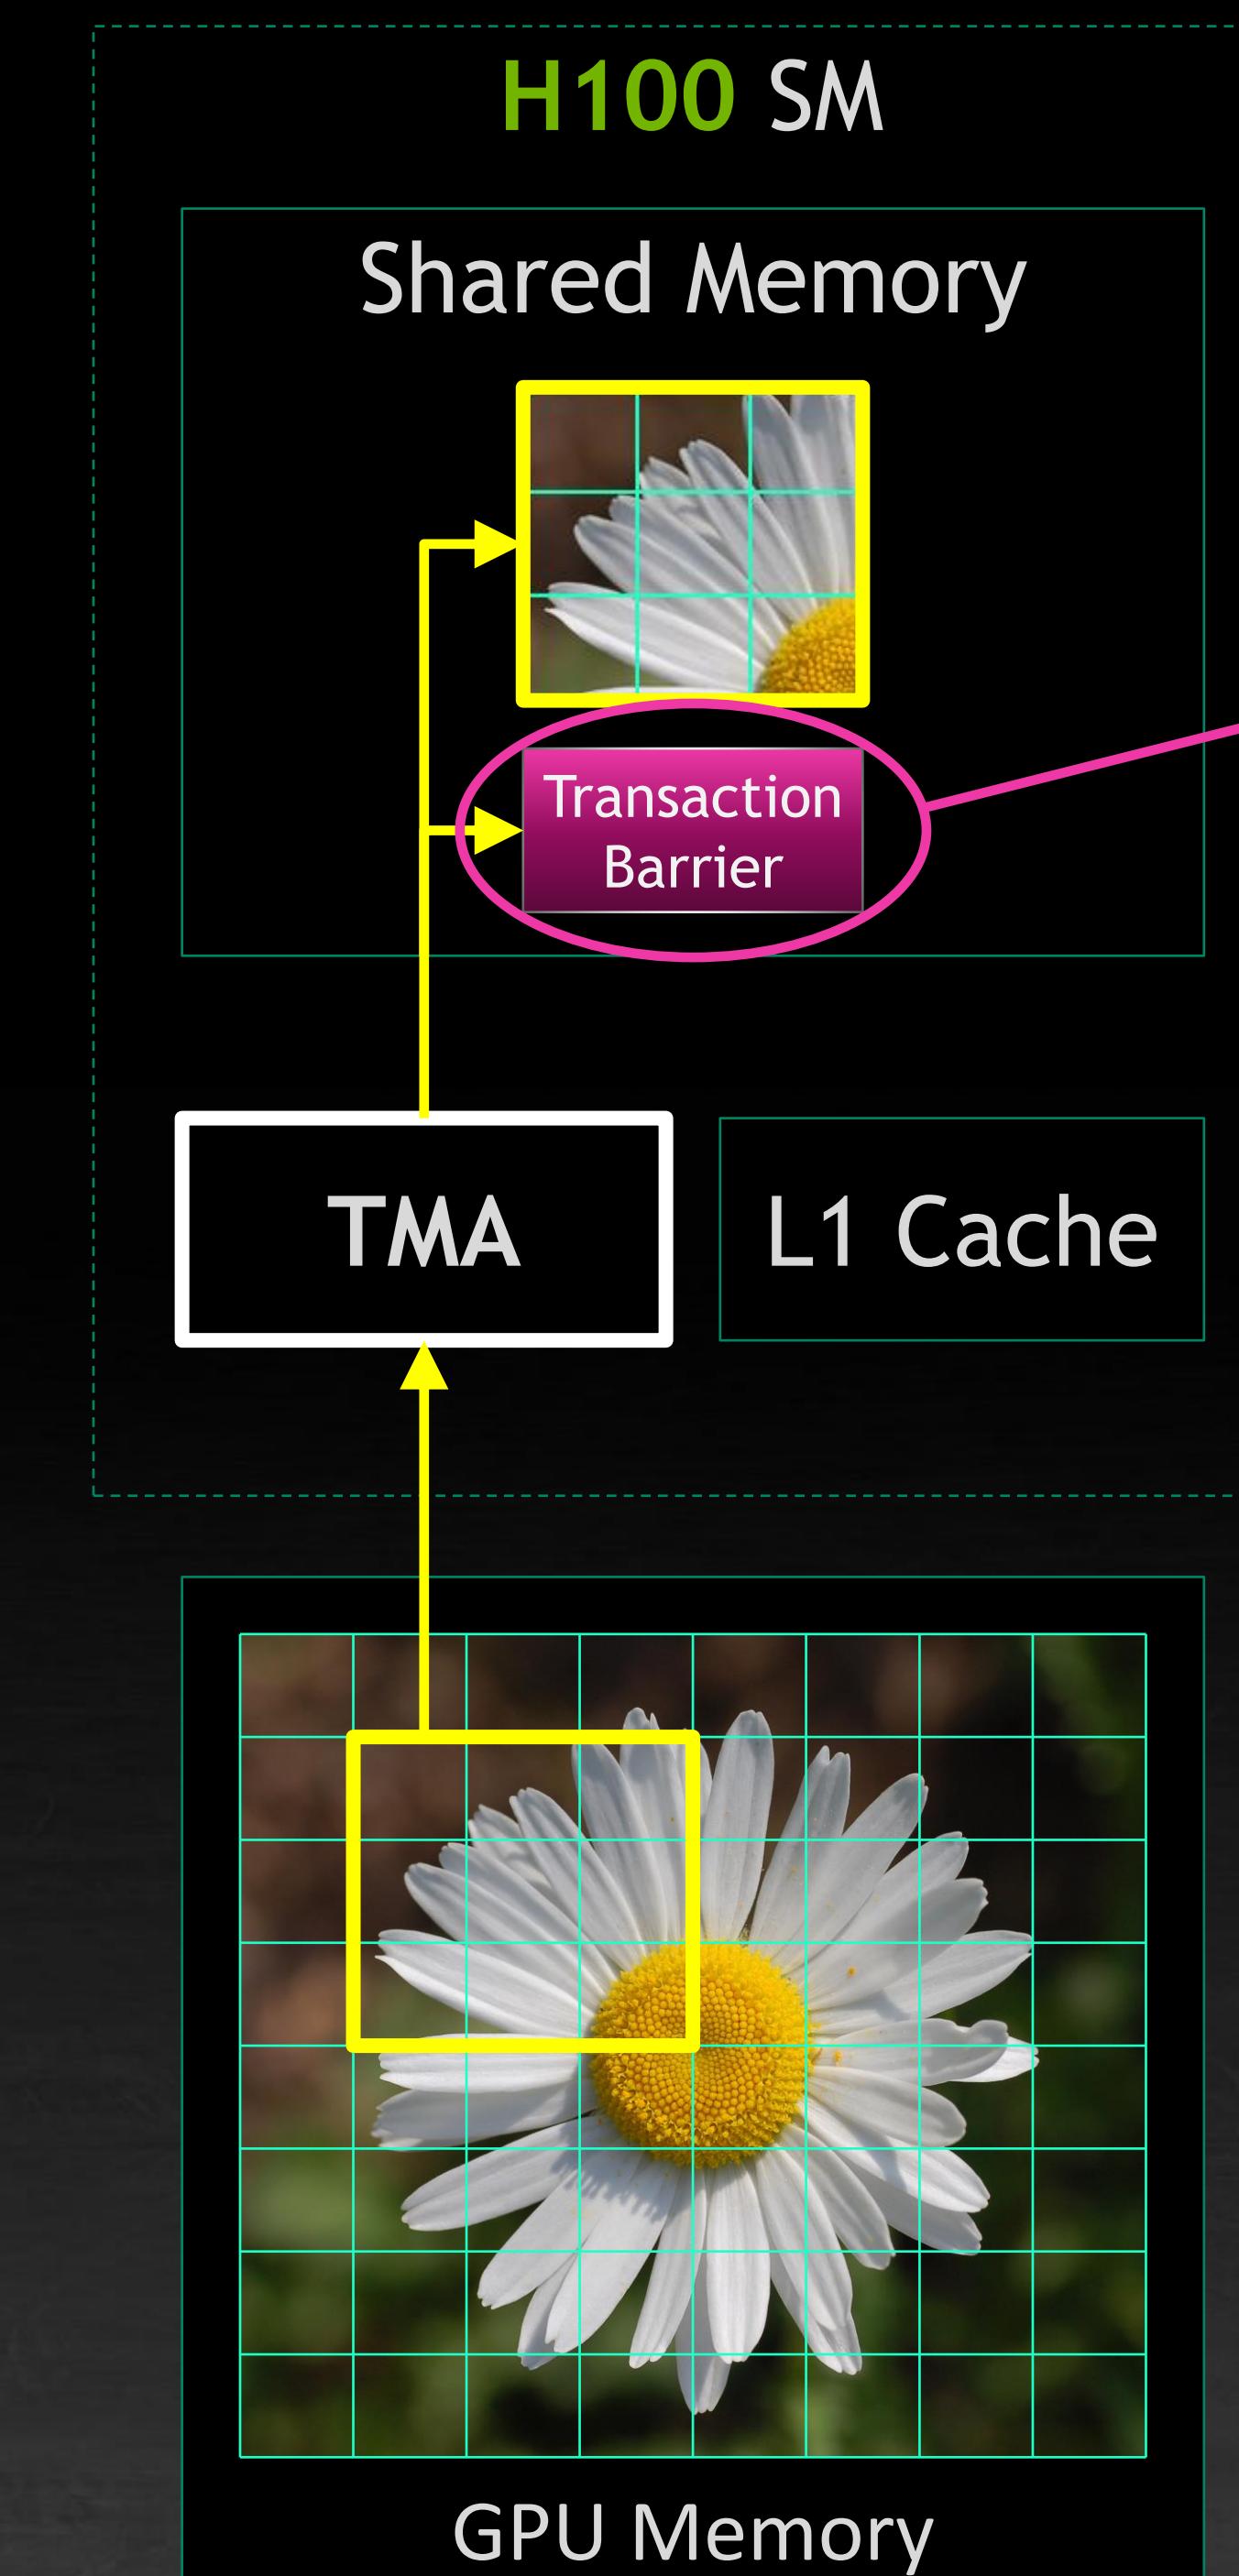
Uses **Asynchronous Transaction Barrier** to track completion

Faster And More Flexible Async Copying

Drop-in enhancement to existing
`cuda::memcpy_async()` function

Single thread triggers any-size copy - no need for
looping or collective copying

Multiple copy operations can contribute to a
single cluster-wide async transaction barrier



Same copy syntax becomes faster and more flexible
New syntax to enable TMA capabilities

HARDWARE ACCELERATED 1D-5D TENSOR MEMORY COPY

Tensor Memory Accelerator Unit

Hardware-accelerated **bi-directional** bulk copy

Global \leftrightarrow Shared Memory

DSMEM \leftrightarrow DSMEM within cluster

Uses **Asynchronous Transaction Barrier** to track completion

Multi-Dimensional Tensor Copying

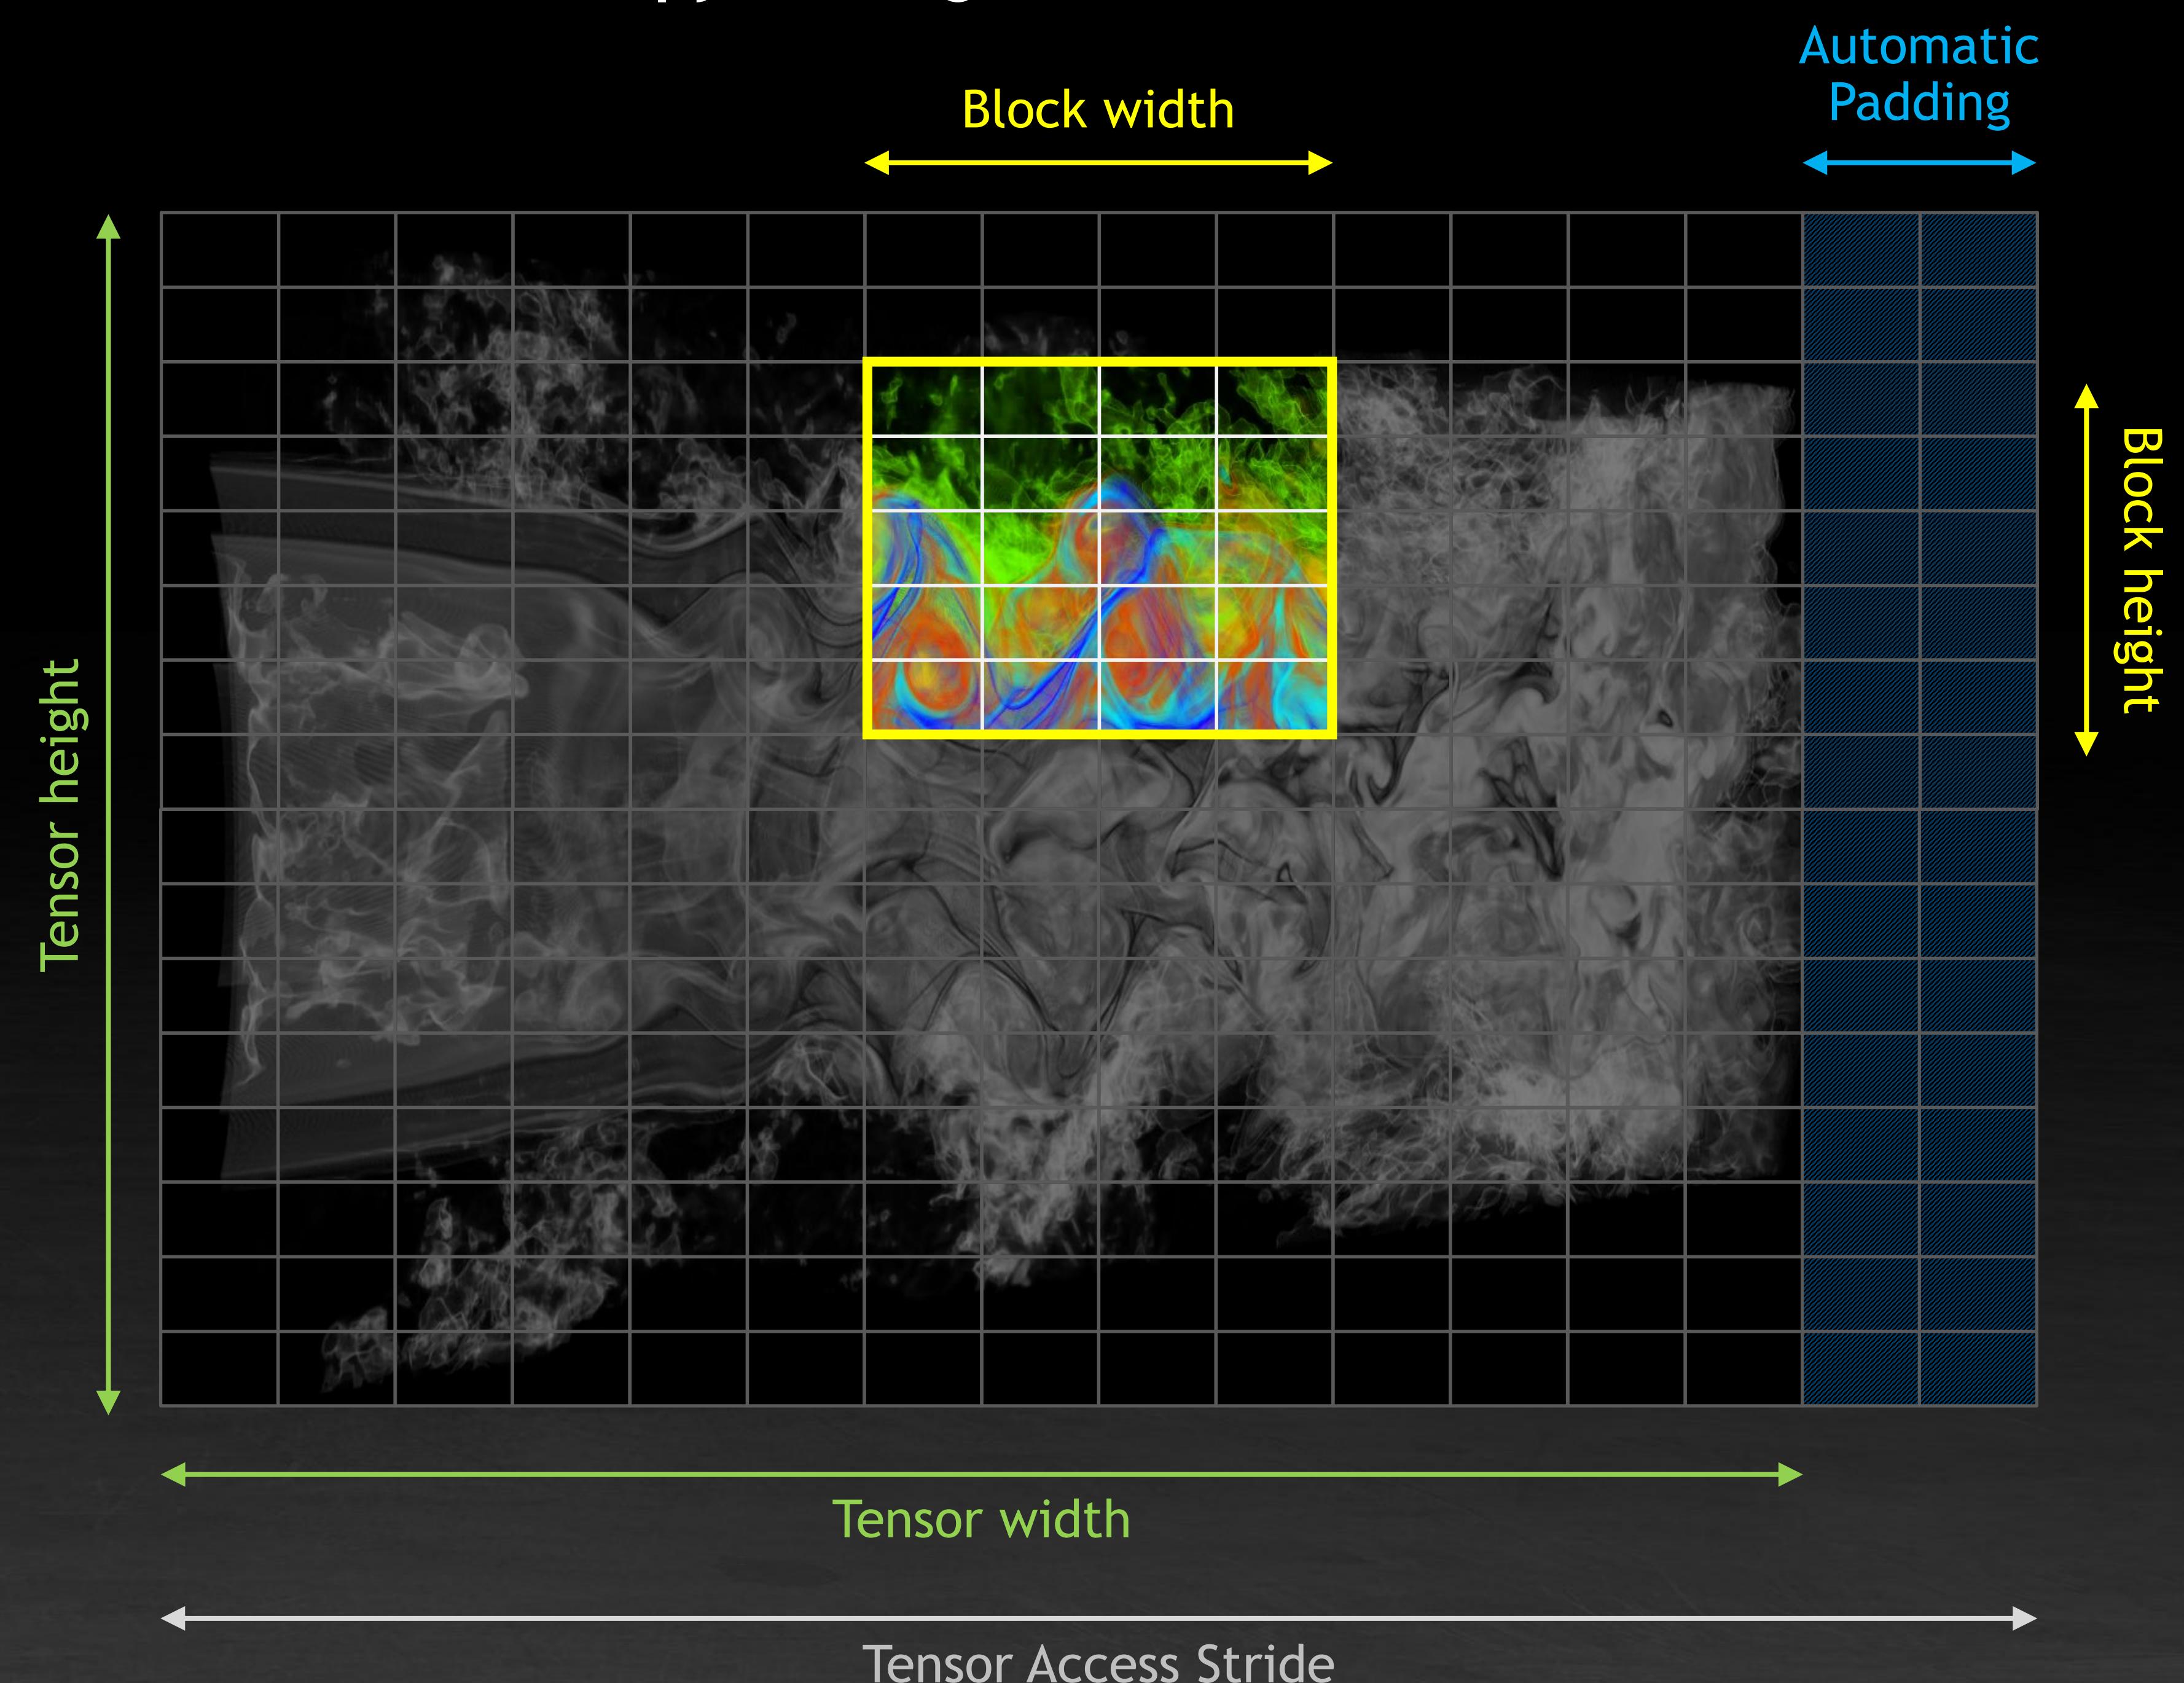
Automatic stride & address generation up to tensors of rank 5

Boundary padding for out-of-bounds accesses

Fire-and-forget from a single thread - everything handled by TMA

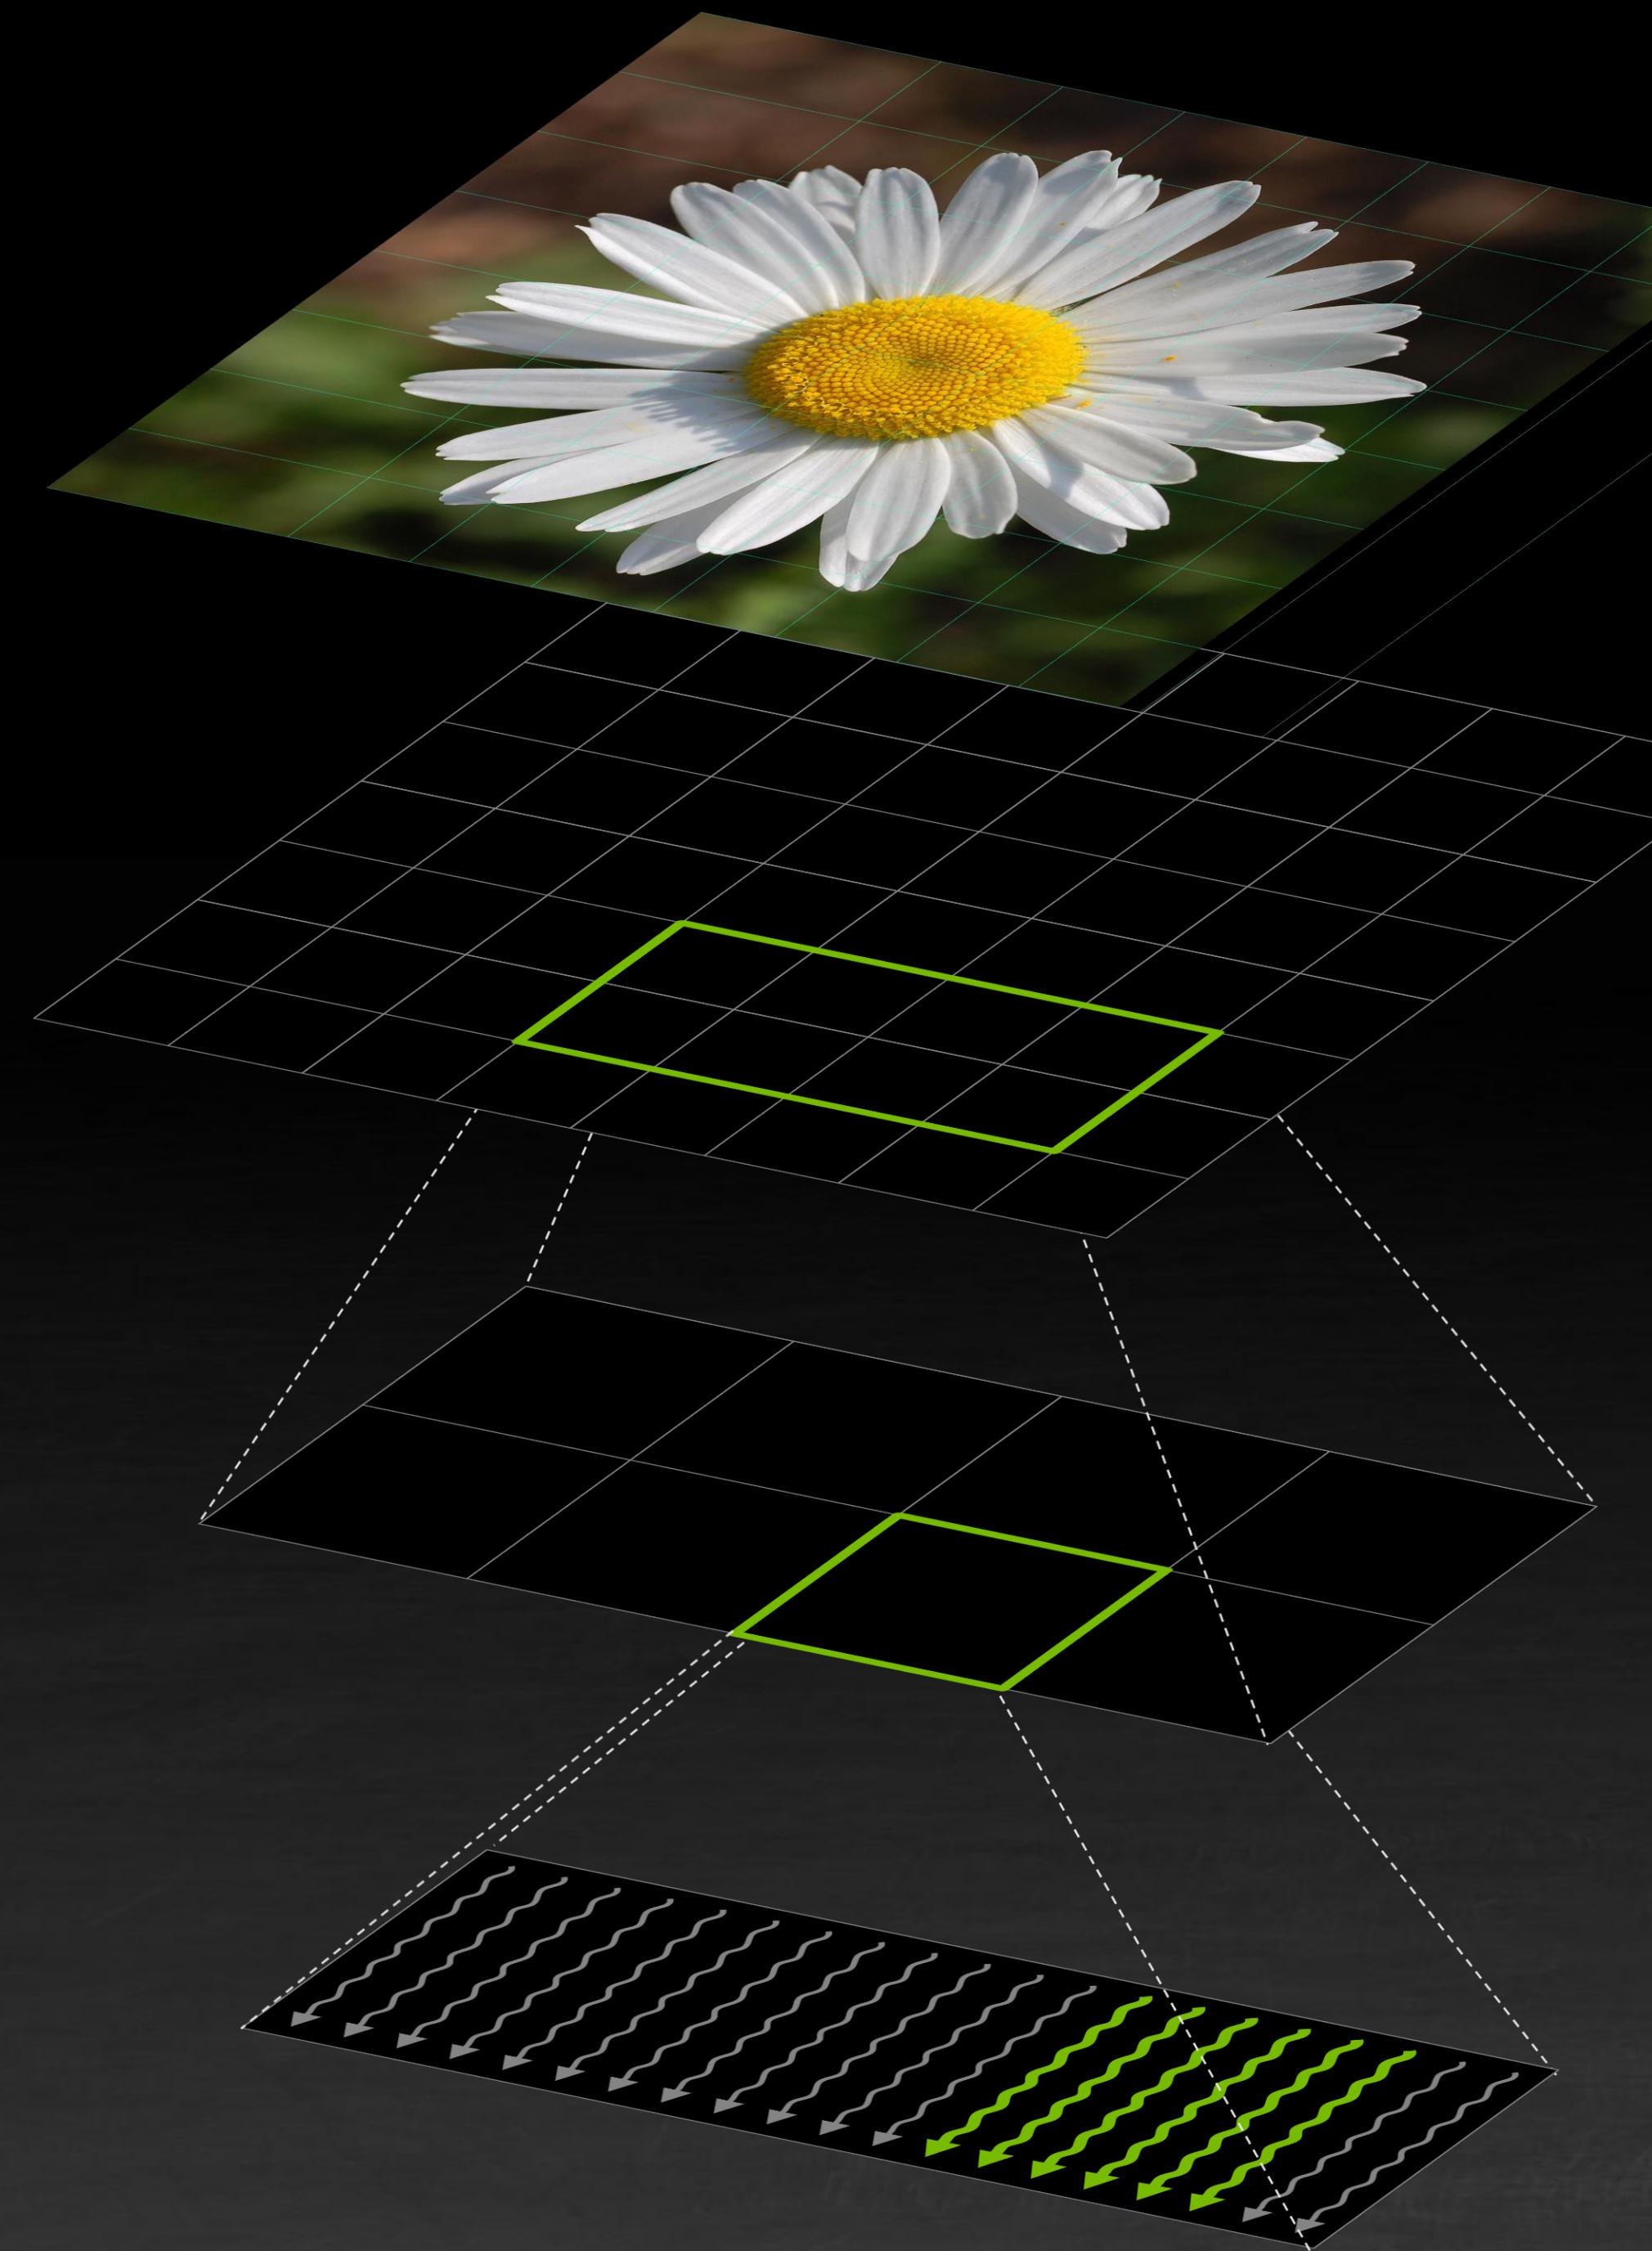
No iteration or bounds-checking code required

The TMA can copy sub-regions of a multi-dimensional tensor



COOPERATIVE GROUPS: PROGRAMMING TO THE NATURAL EXECUTION HIERARCHY

Cooperative Groups is an **explicit** model for expressing cooperation & synchronization among threads



Grid
of work

Cluster
of Blocks

Block
of Threads

Warp
of Threads

```
cooperative_groups::grid_group grid = this_grid();
```

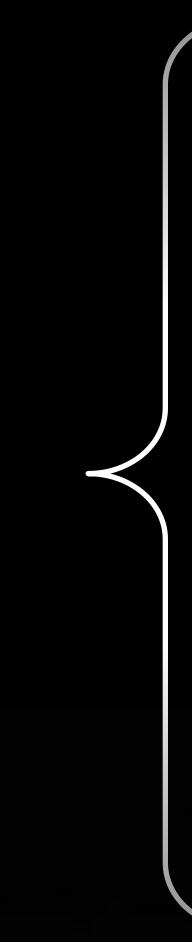
```
cooperative_groups::cluster_group cluster = this_cluster();
```

```
cooperative_groups::thread_block block = this_thread_block();
```

```
cooperative_groups::thread_block_tile<32> warp = tiled_partition<32>(block);
```

COLLECTIVE OPERATIONS AT EVERY LEVEL OF HIERARCHY

Synchronization



Grid must be launched cooperatively in order to use grid-scope collectives

```
cooperative_groups::grid_group grid = this_grid();
```

```
cooperative_groups::cluster_group cluster = this_cluster();
```

Synchronization

Reduction

Prefix-sum



```
cooperative_groups::thread_block block = this_thread_block();
```

Warp & smaller groups can use: shfl, any, all, ballot, match

```
cooperative_groups::thread_block_tile<32> warp = tiled_partition<32>(block);
```

COLLECTIVE OPERATIONS INCLUDING FOR MULTI-WARP GROUPS

Supported `thread_block_tile` sizes of 64, 128, 256, 512

Synchronization

```
cooperative_groups::grid_group grid = this_grid();
```

```
cooperative_groups::cluster_group cluster = this_cluster();
```

Synchronization

Reduction

Prefix-sum

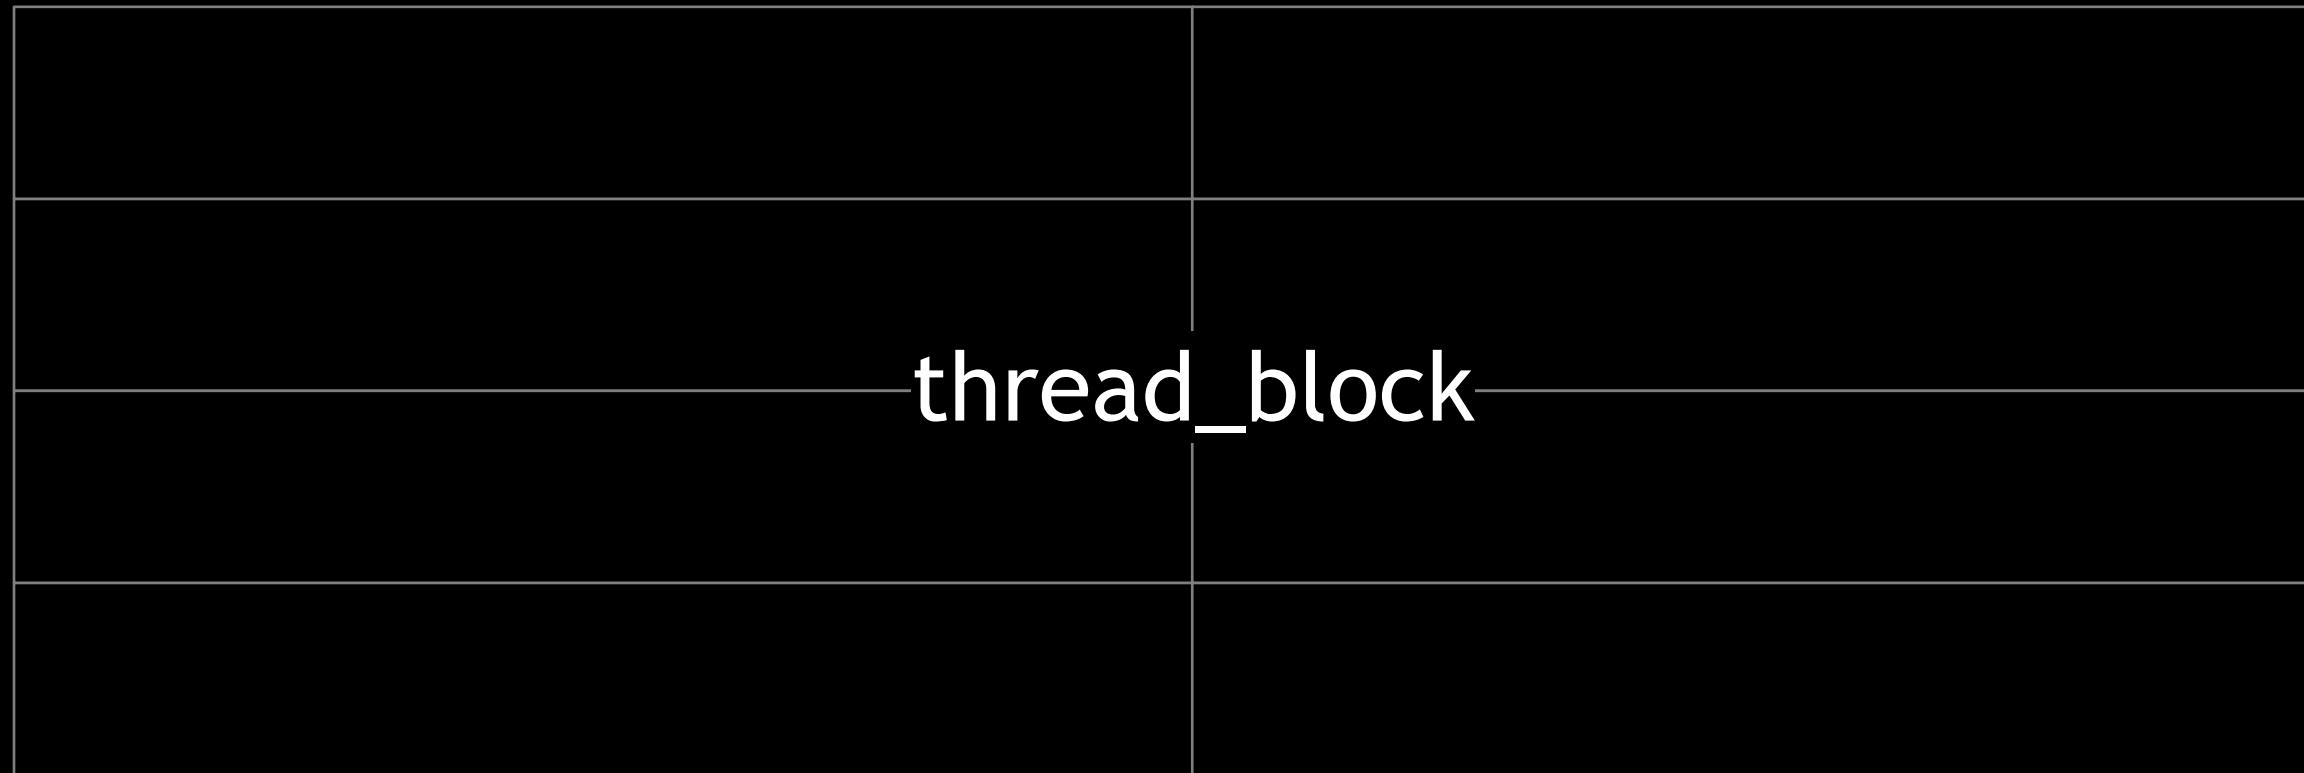
```
cooperative_groups::thread_block block = this_thread_block();
```

```
cooperative_groups::thread_block_tile<128> quadwarp = tiled_partition<128>(block);
```

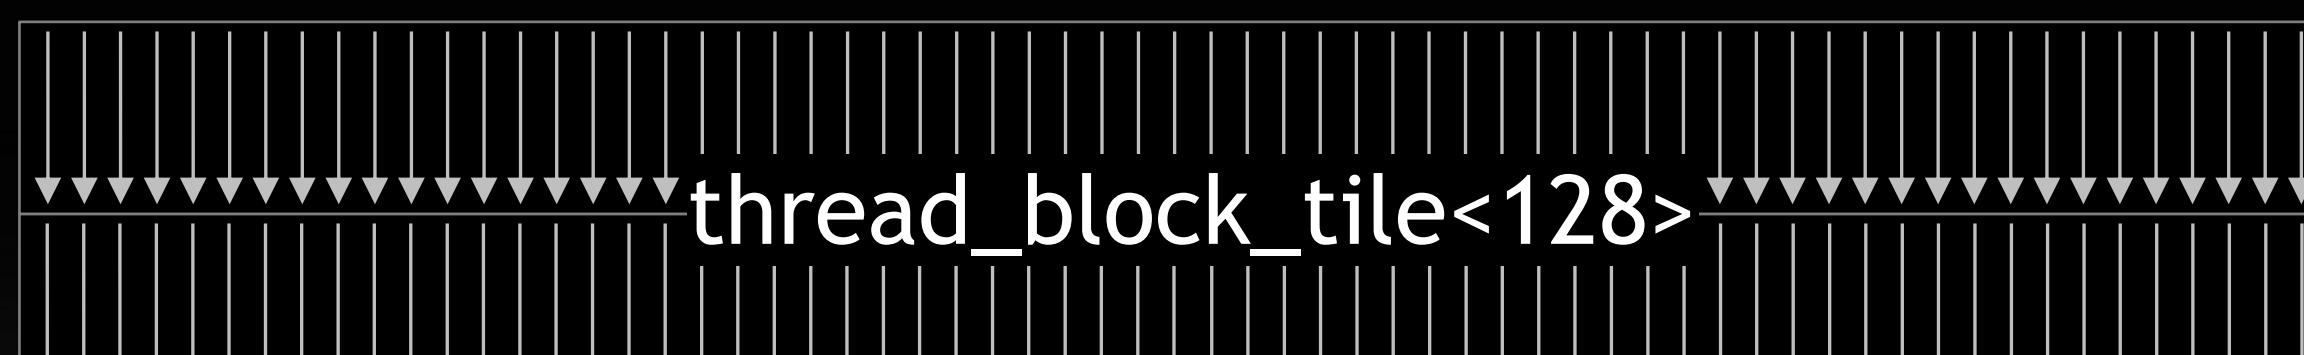
```
cooperative_groups::thread_block_tile<32> warp = tiled_partition<32>(quadwarp);
```

TYPE SAFETY ENABLES COMPOSABLE LIBRARIES OF PARALLEL FUNCTIONS

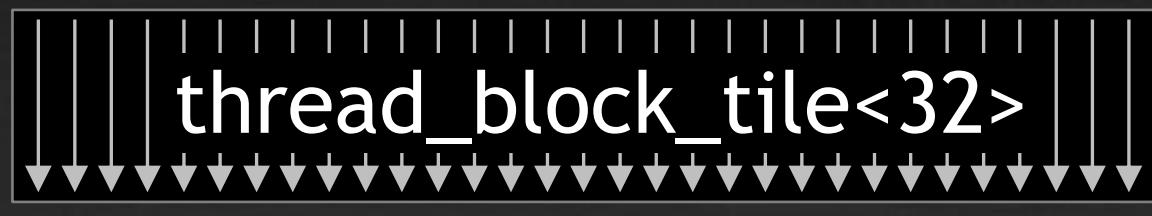
Be explicit about how many threads a function wants - e.g. Radix-128 FFT can enforce 128 threads



Base thread block of 256 threads (8 warps)



Group of 128 threads (4 warps)



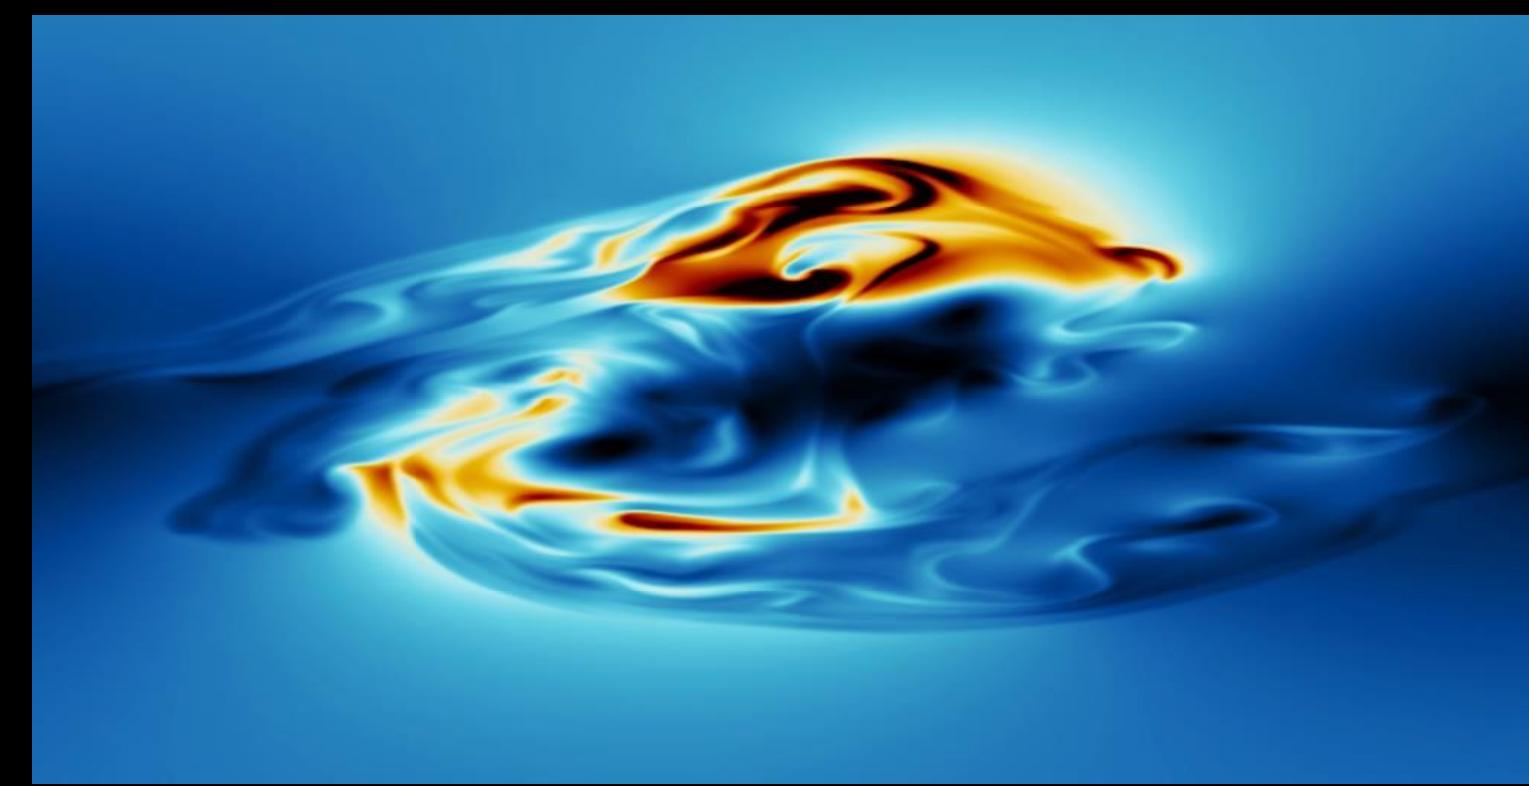
Group of 32 threads (1 warp)

```
using namespace cooperative_groups;
__global__ void kernel(float *in, float *out, ...) {
    // Subdivide my block into 128-thread tiles for the FFT call
    thread_block_tile<128> tile128 = tiled_partition<128>(this_thread_block());
    fft128(tile128, in, out, ...);
}
```

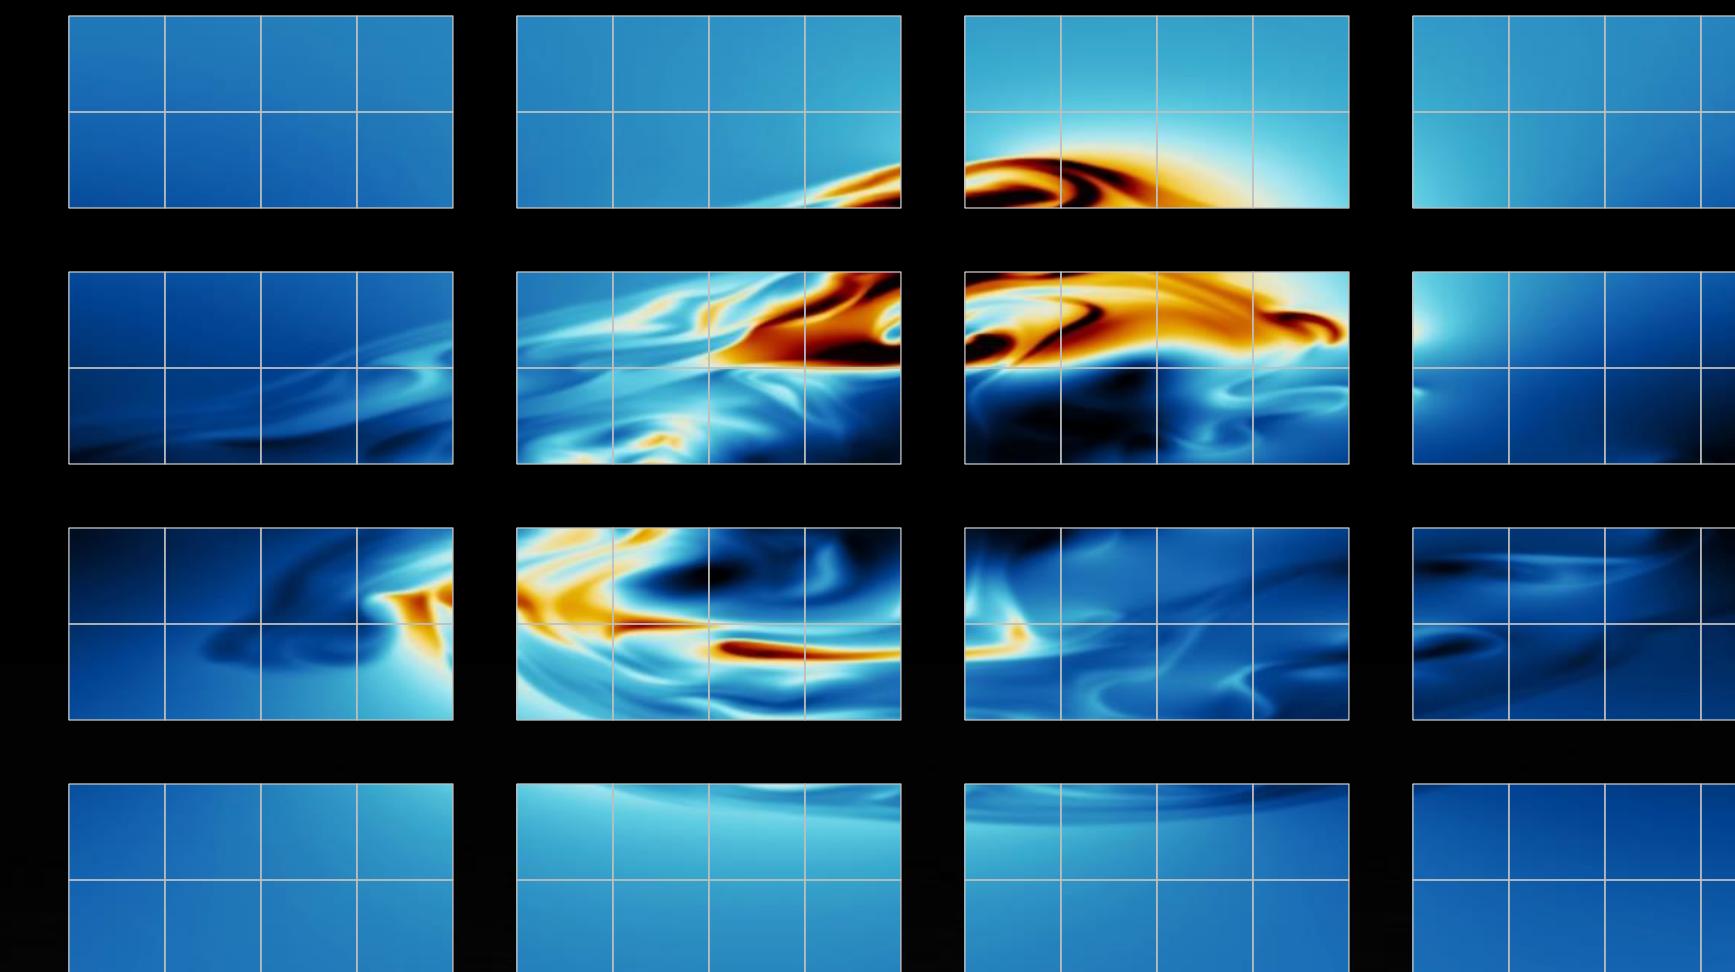
```
// A radix-128 FFT function which *must* be called with 128 threads
__device__ void fft128(const thread_block_tile<128> &group,
                      float *in, float *out, ...) {
    do_fft(in, out, ...);
    group.sync();           // Sync all 128 threads in the group
}
```

```
// Functions can require an explicit thread count via the type system
__device__ int sum(const thread_block_tile<32> &group,
                  float *in, float *out, ...) {
    // Collective operations are common to all block and sub-block groups
    return group.reduce(in[group.thread_rank()]);
}
```

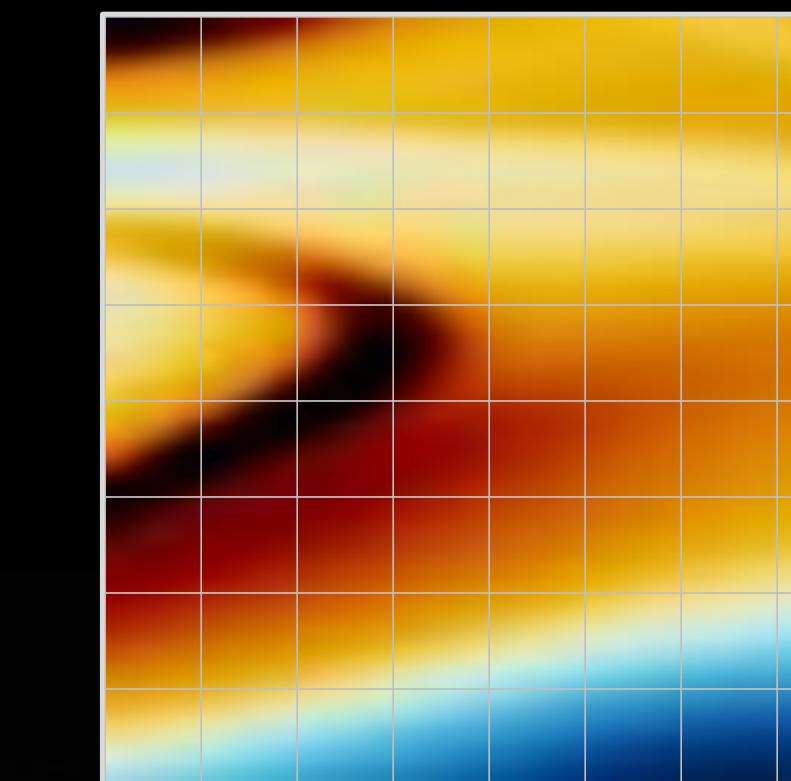
EXAMPLE: HIERARCHICAL EXECUTION, DATA EXCHANGE AND SYNCHRONIZATION



1
Launch cooperative grid
of 4x4 clusters, each
with 4x2 blocks

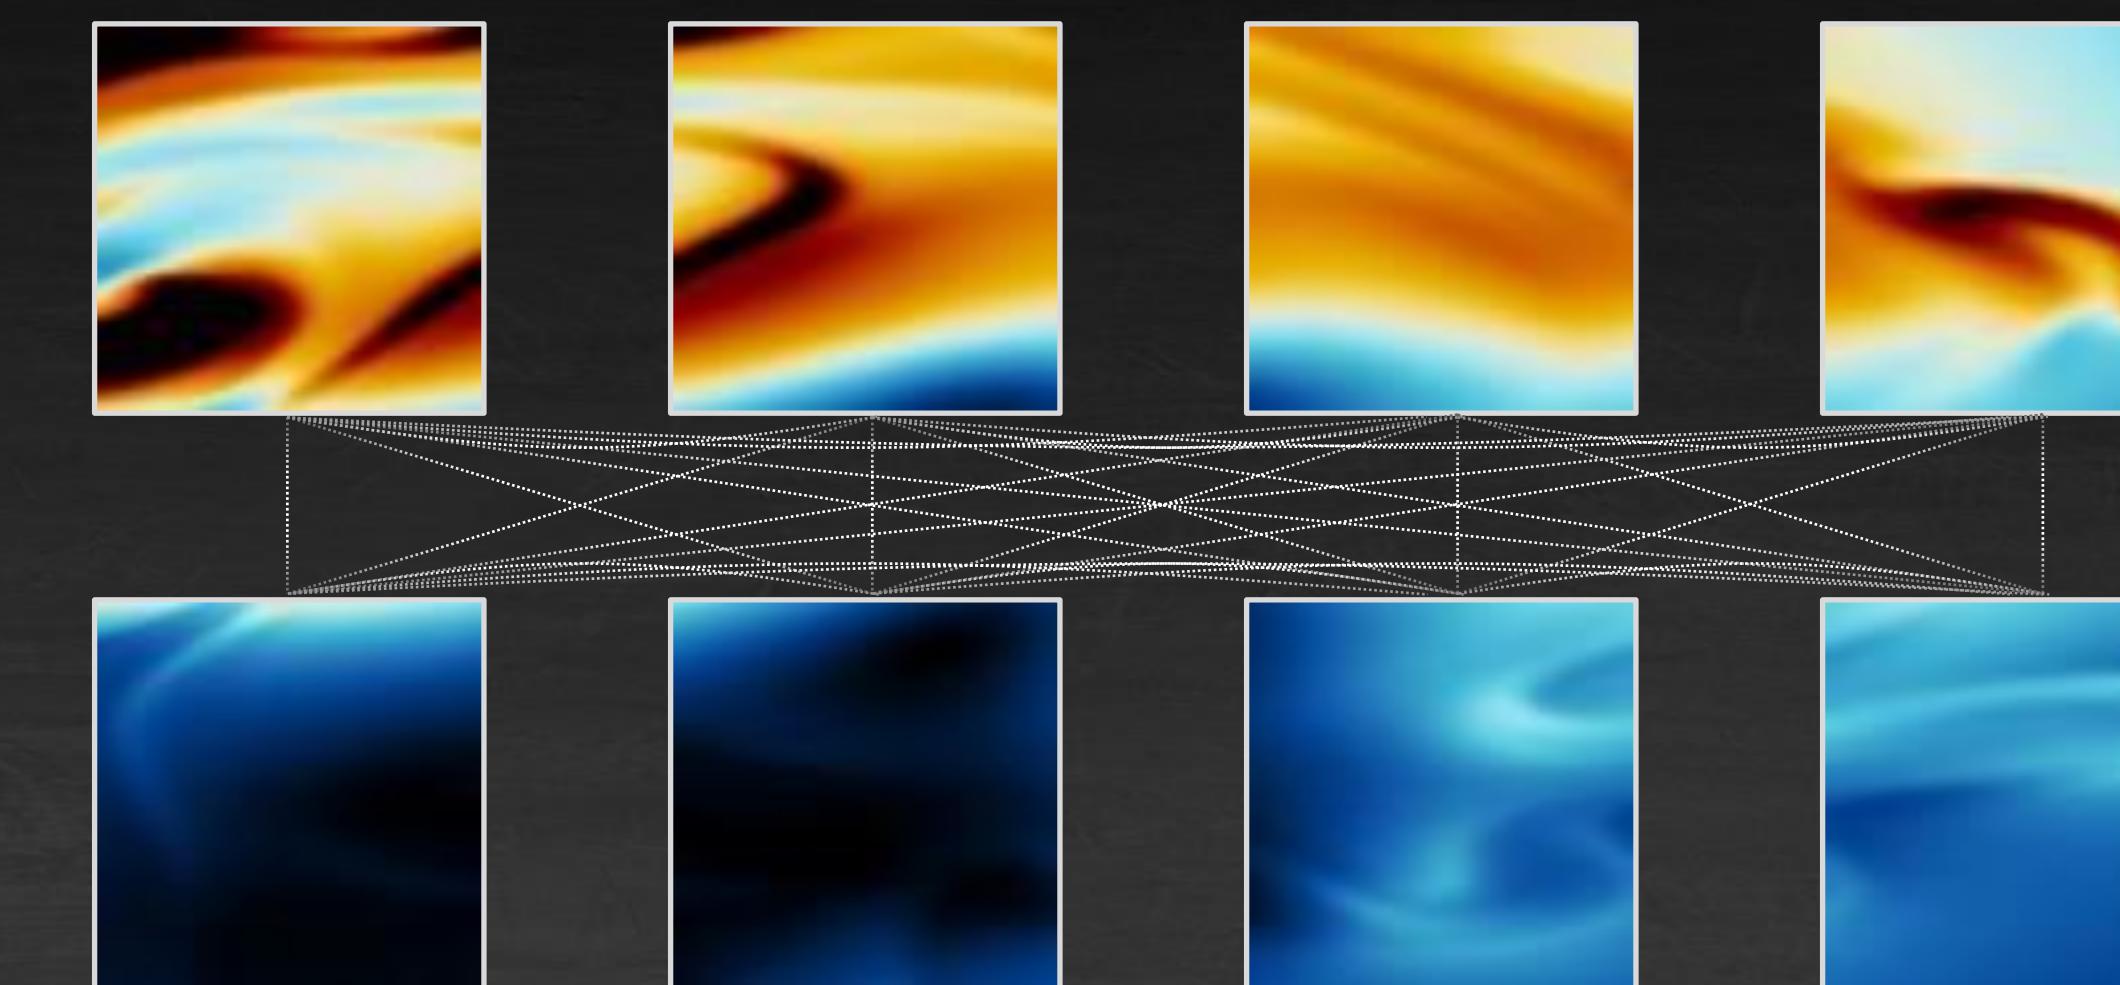


2
Use cooperative cluster indexing
to pull data from neighbours



3
Solve locally
within each
block

5
Synchronize with
grid and solve at
grid level; go to 2

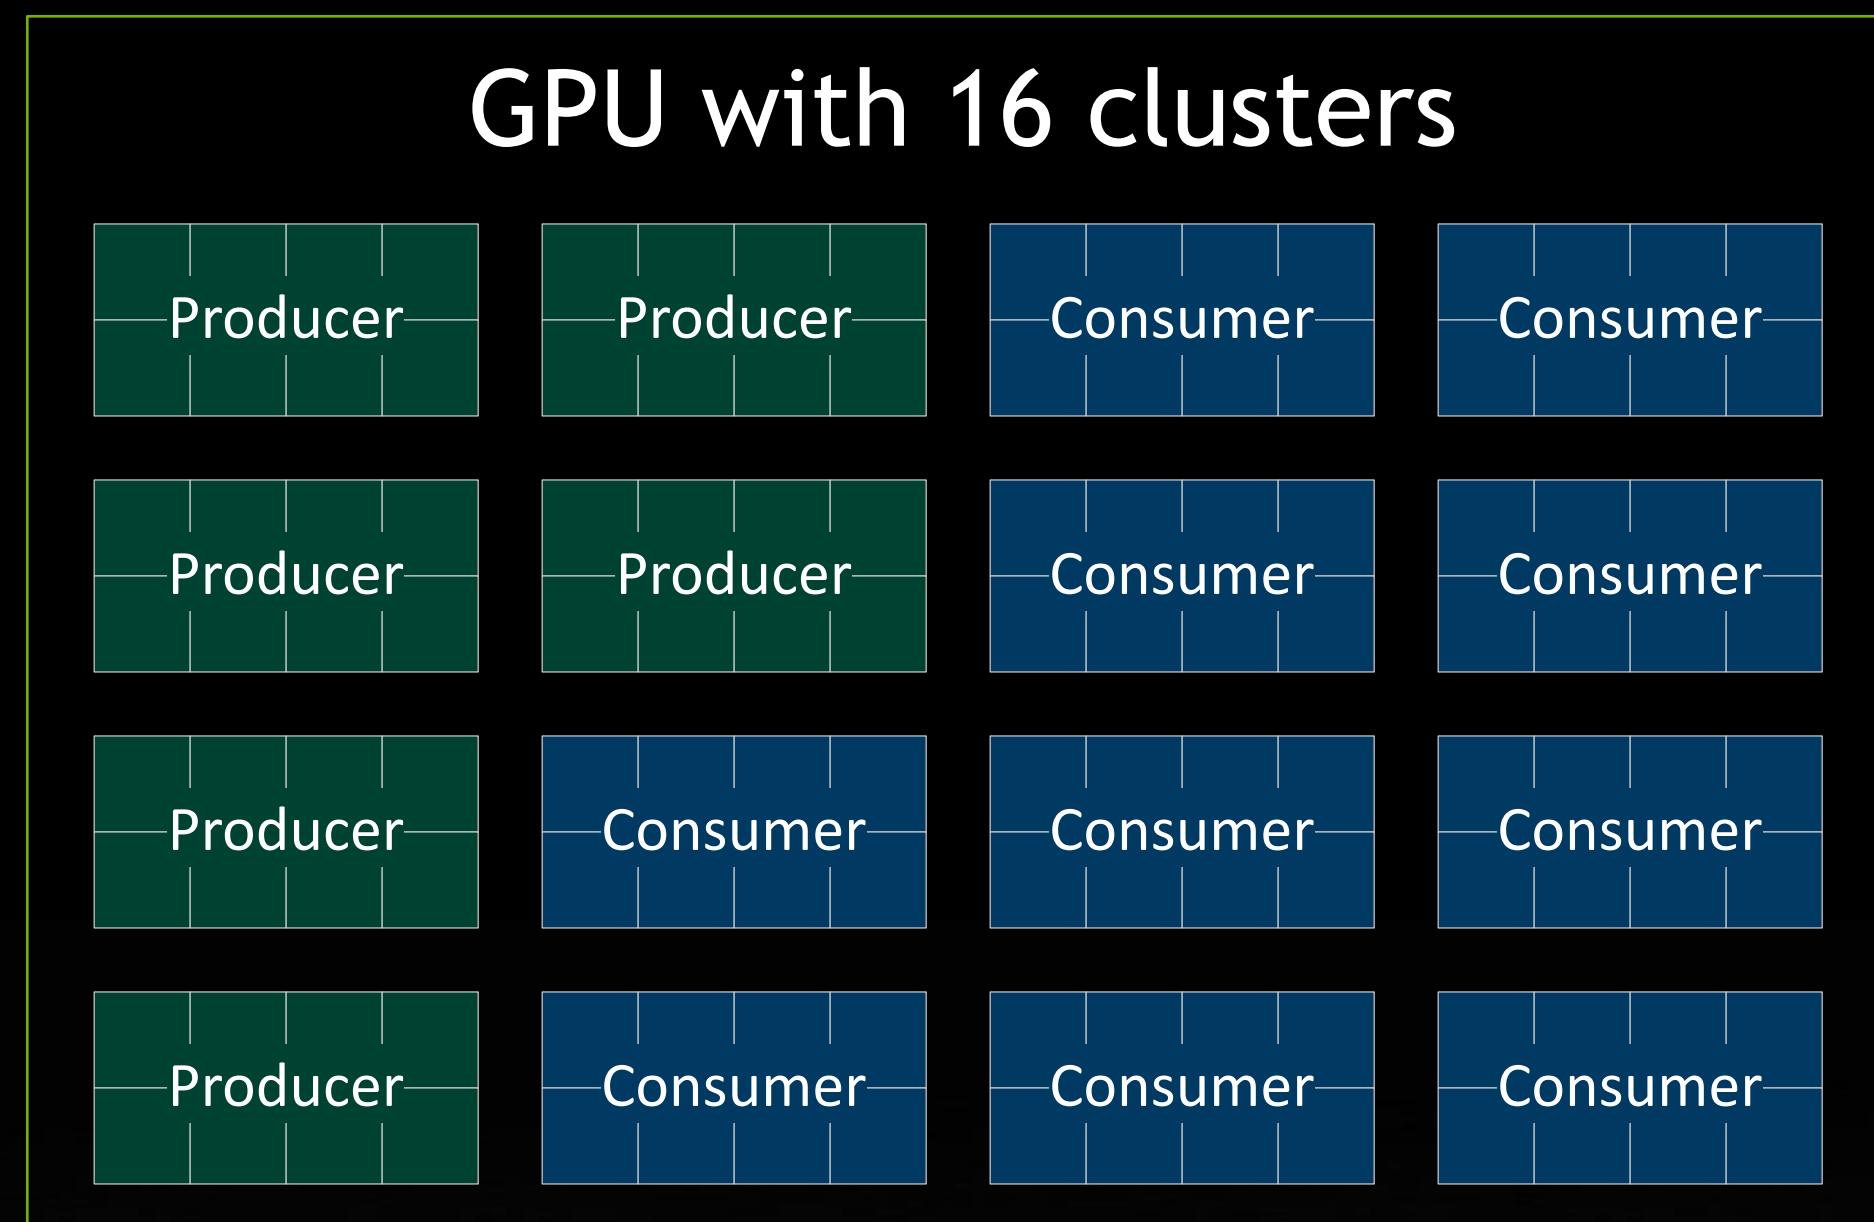


4
Synchronize with cluster and
solve at cluster level

Use distributed shared memory for fast inter-block communication.

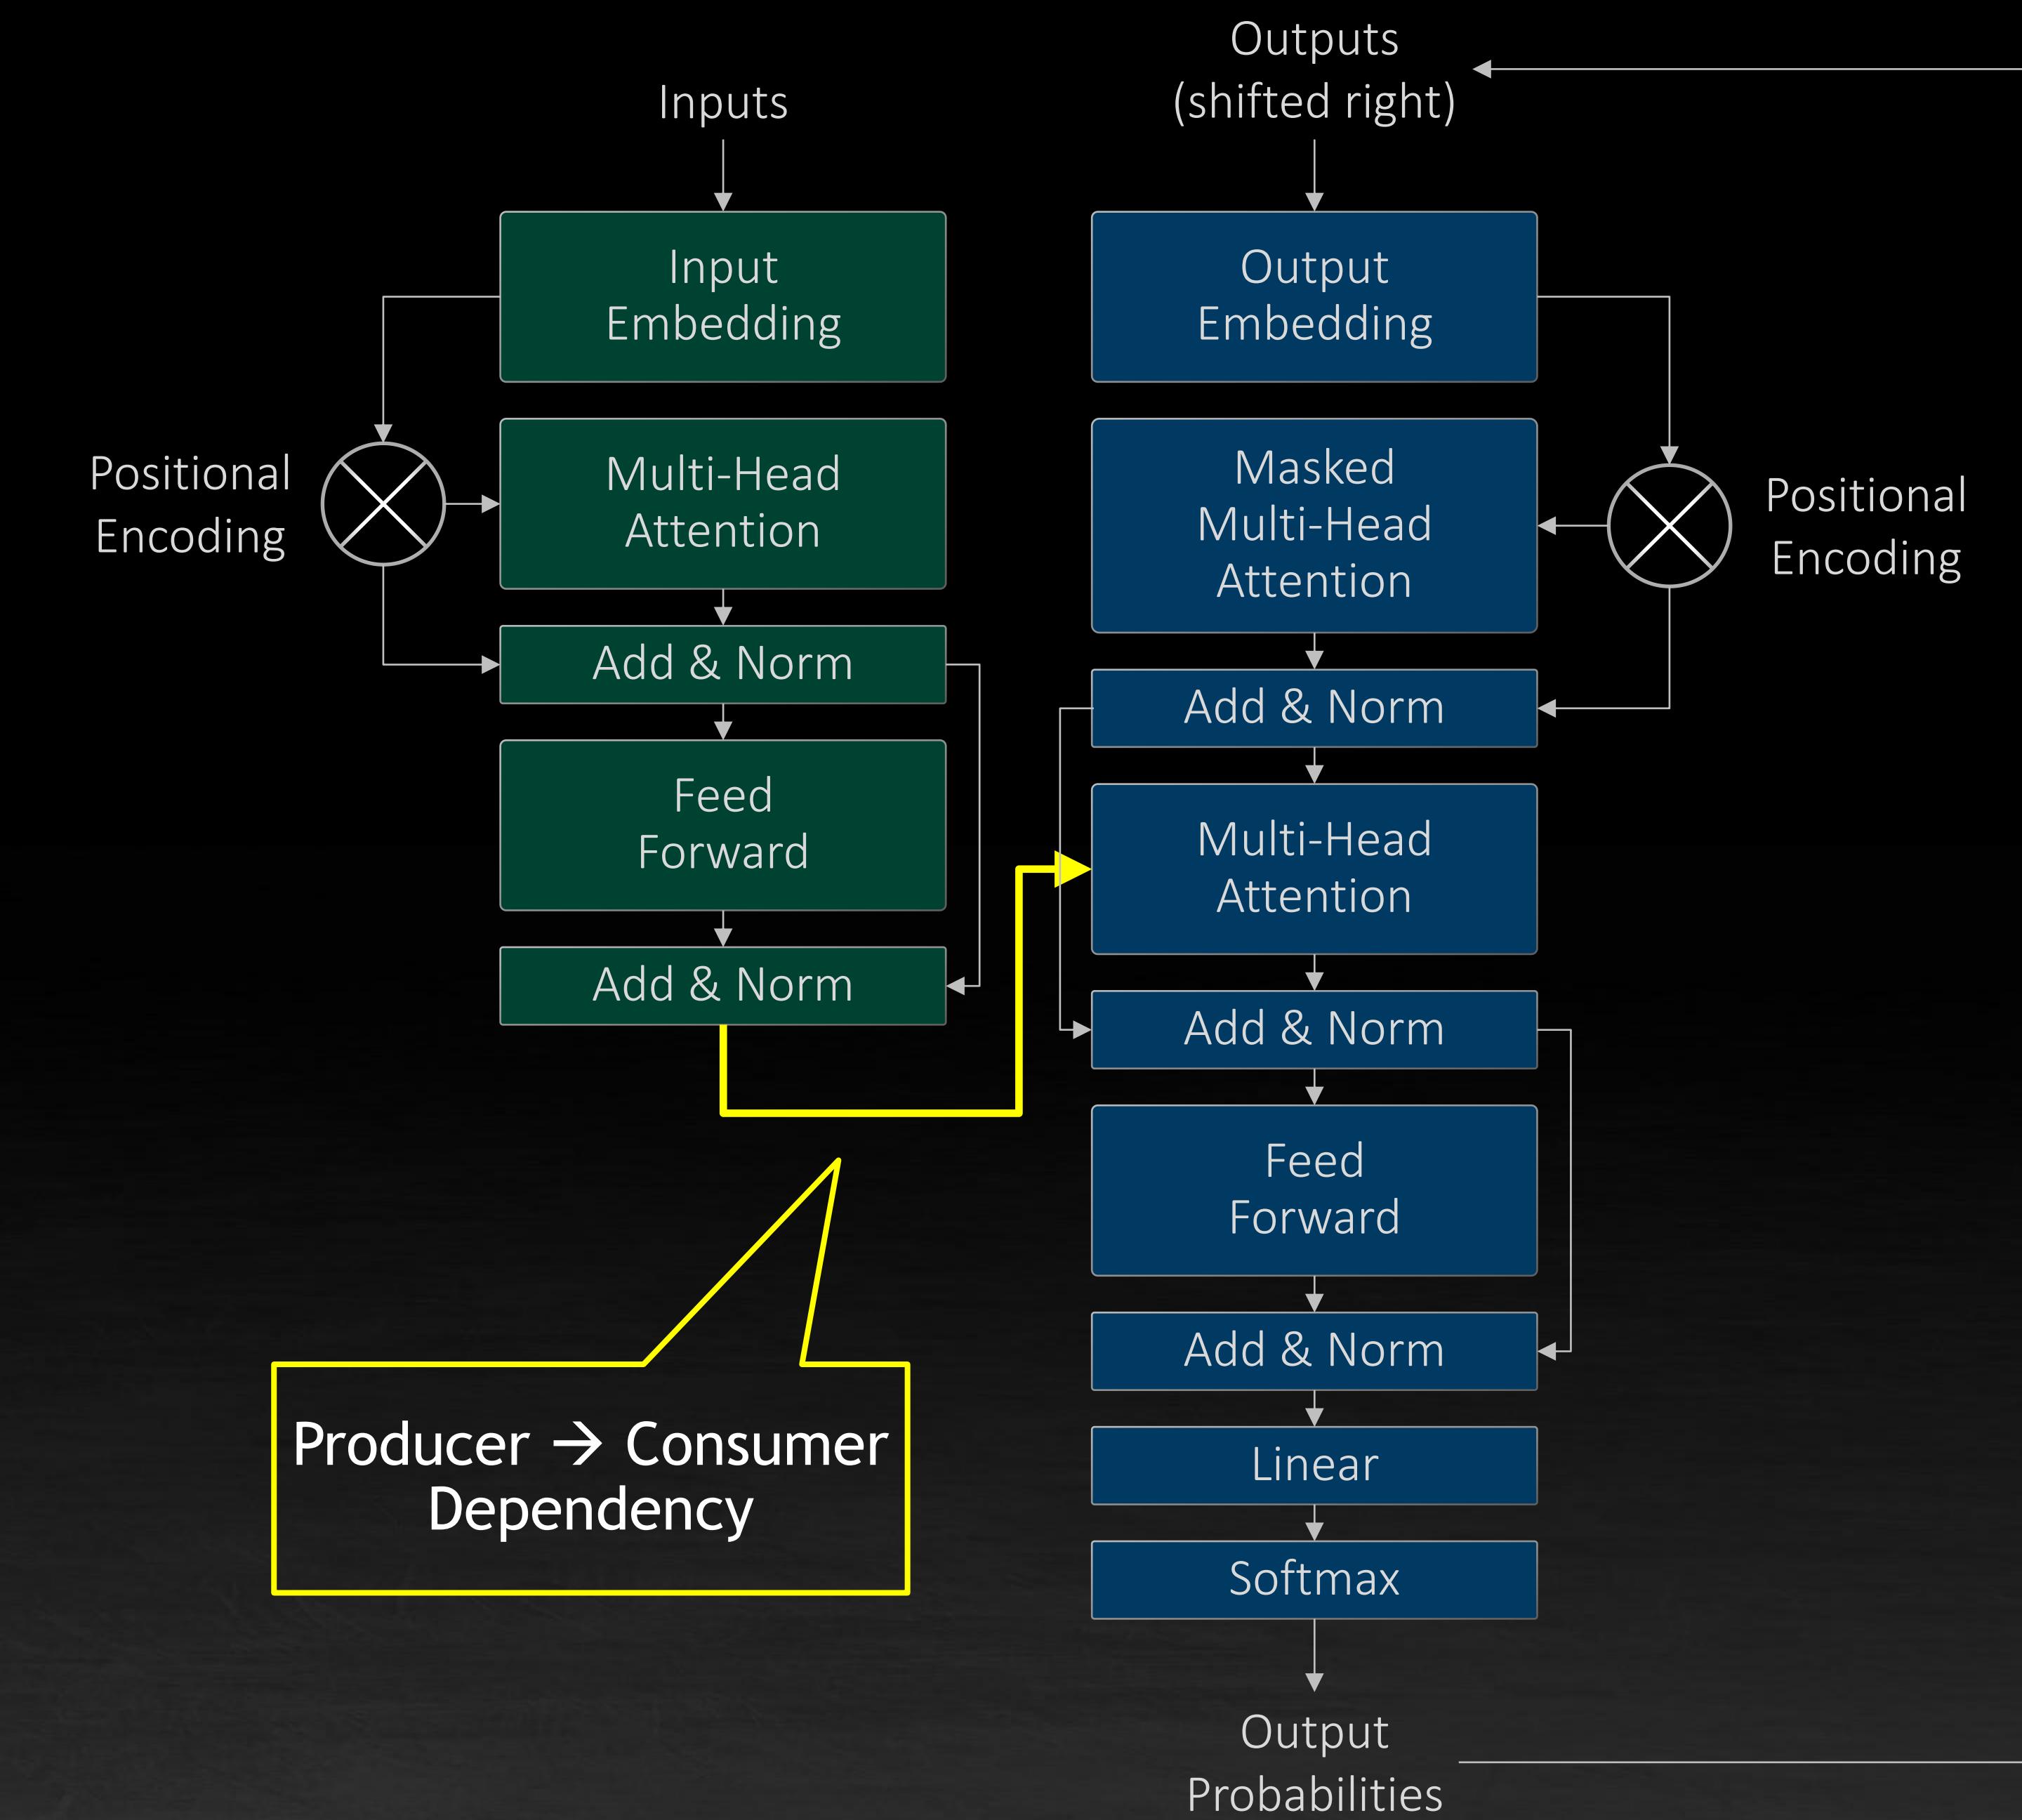
EXAMPLE: PRODUCER/CONSUMER TASK PARALLELISM AT ANY SCALE

Producer/Consumer can operate at every scale - grid, cluster or block



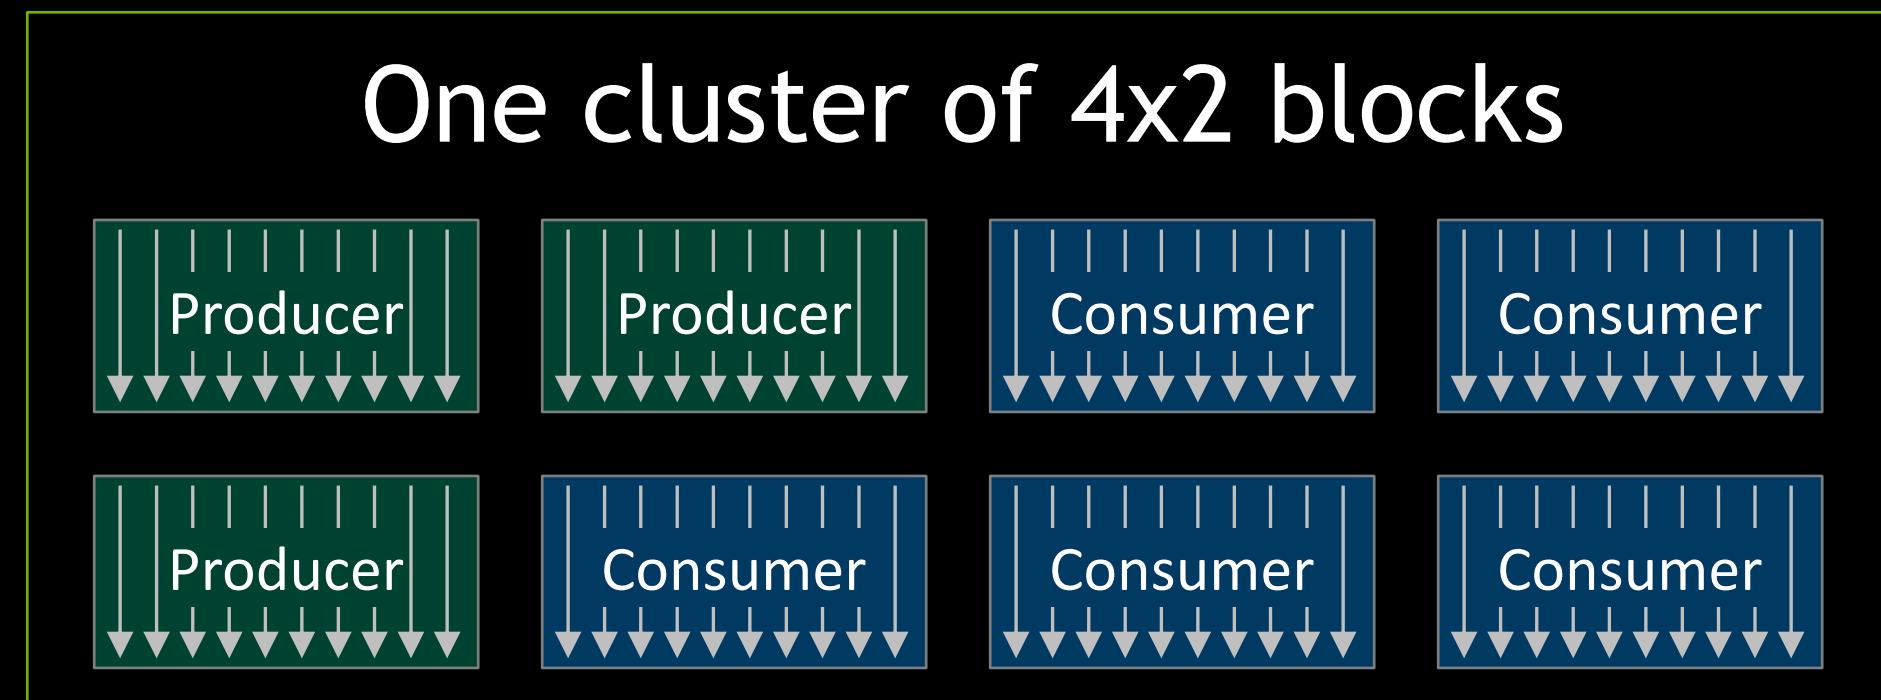
Specialise **clusters** across whole GPU
into “producer” and “consumer” roles

Specialise	Into	Exchange data via	Synchronize using
Grid	Clusters	Global memory	this_grid().sync()



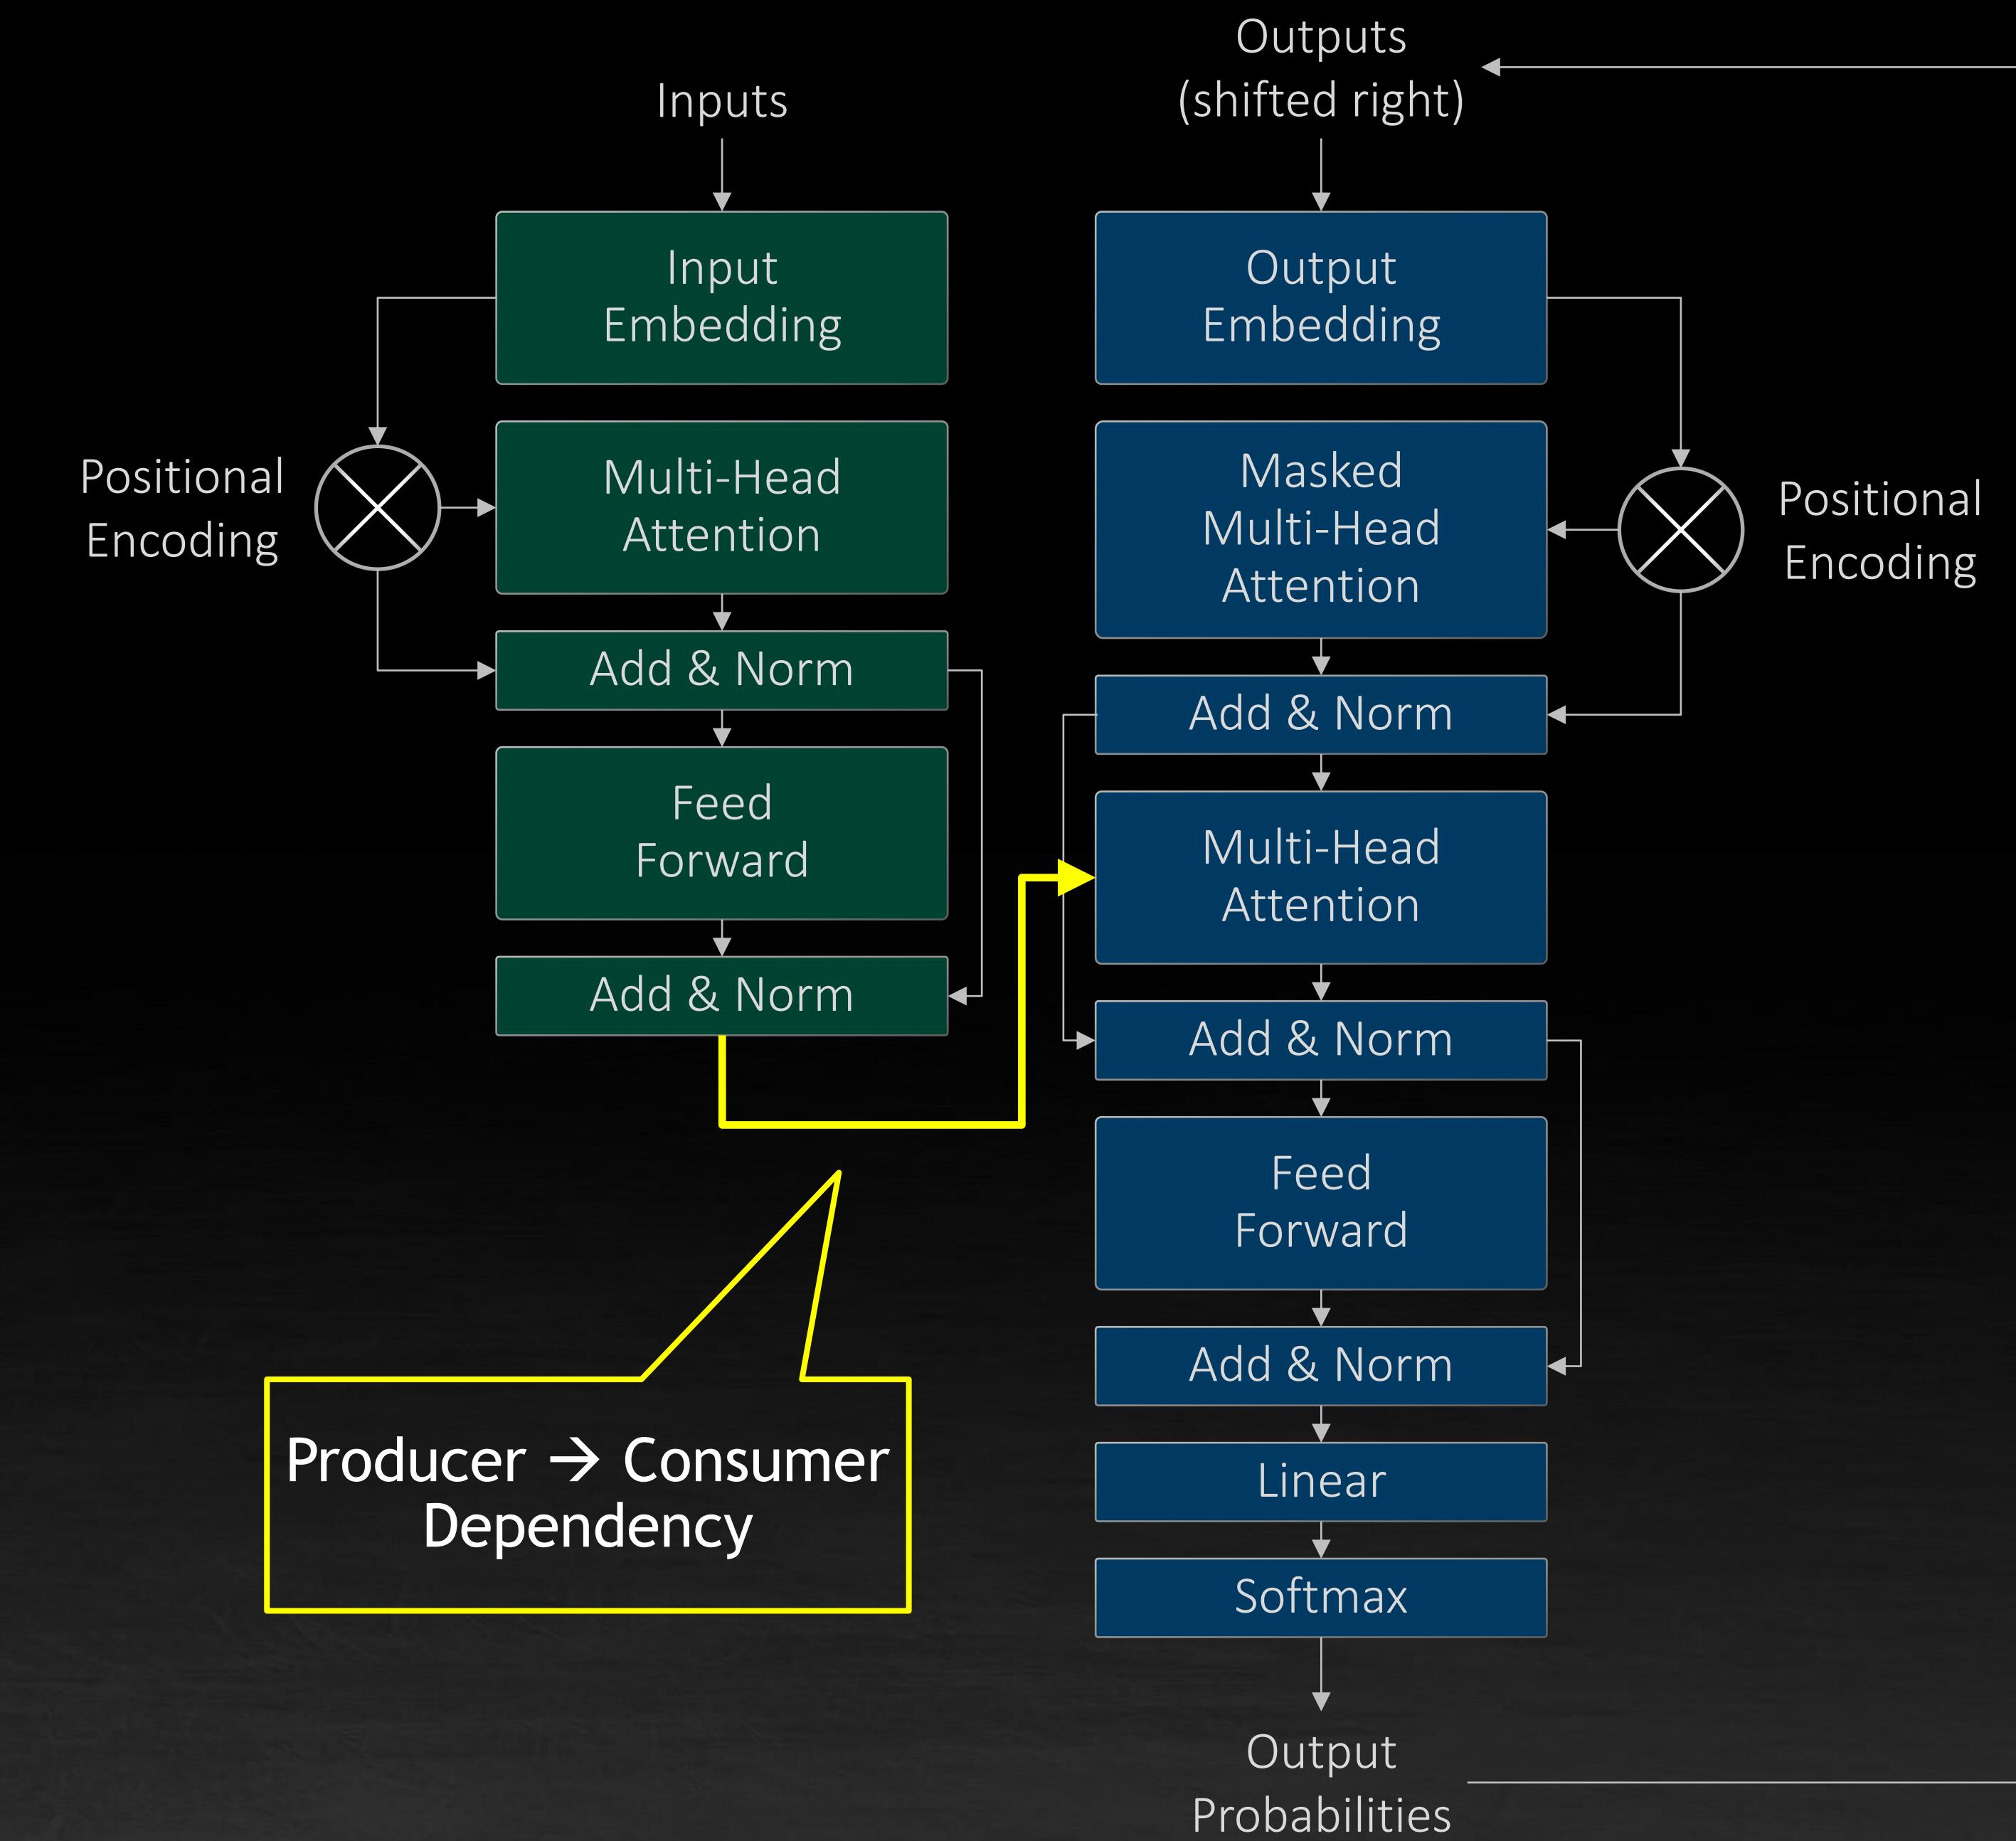
EXAMPLE: PRODUCER/CONSUMER TASK PARALLELISM AT ANY SCALE

Producer/Consumer can operate at every scale - grid, cluster or block



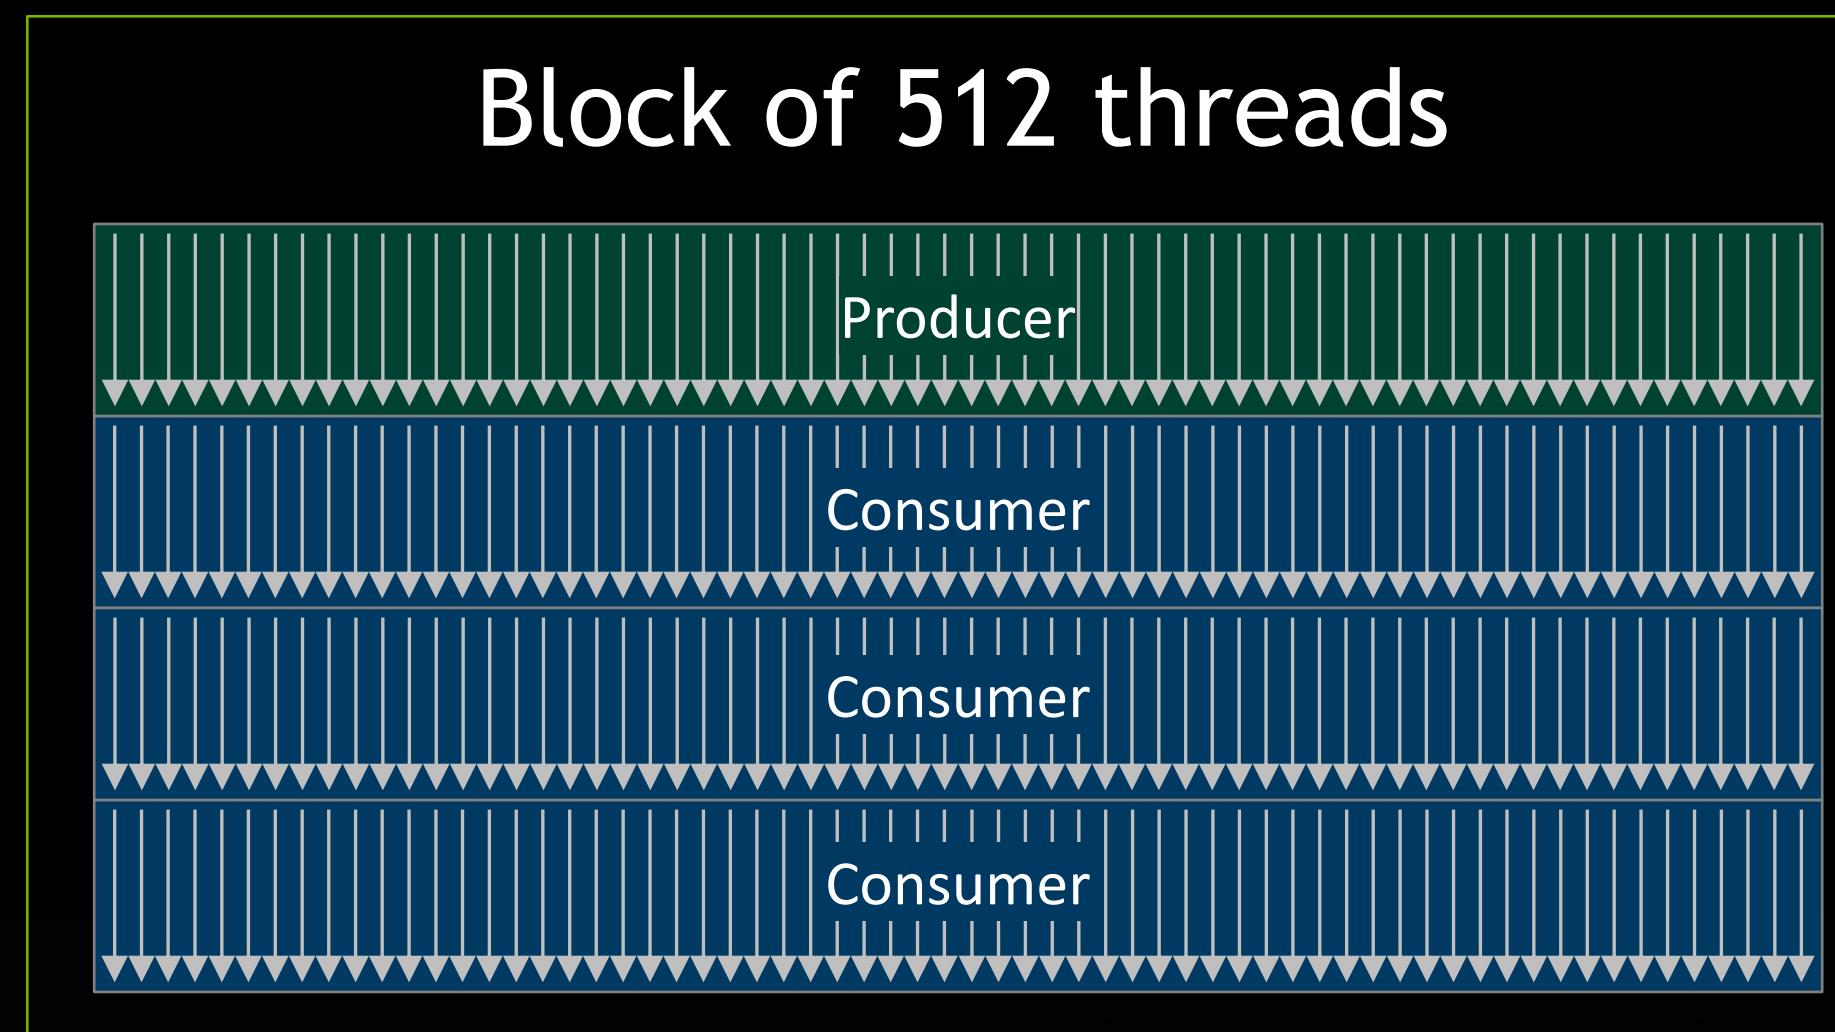
Specialise blocks within each cluster into “producer” and “consumer” roles

Specialise	Into	Exchange data via	Synchronize using
Grid	Clusters	Global memory	<code>this_grid().sync()</code>
Cluster	Blocks	Distributed-shared or global	<code>this_cluster().sync()</code>



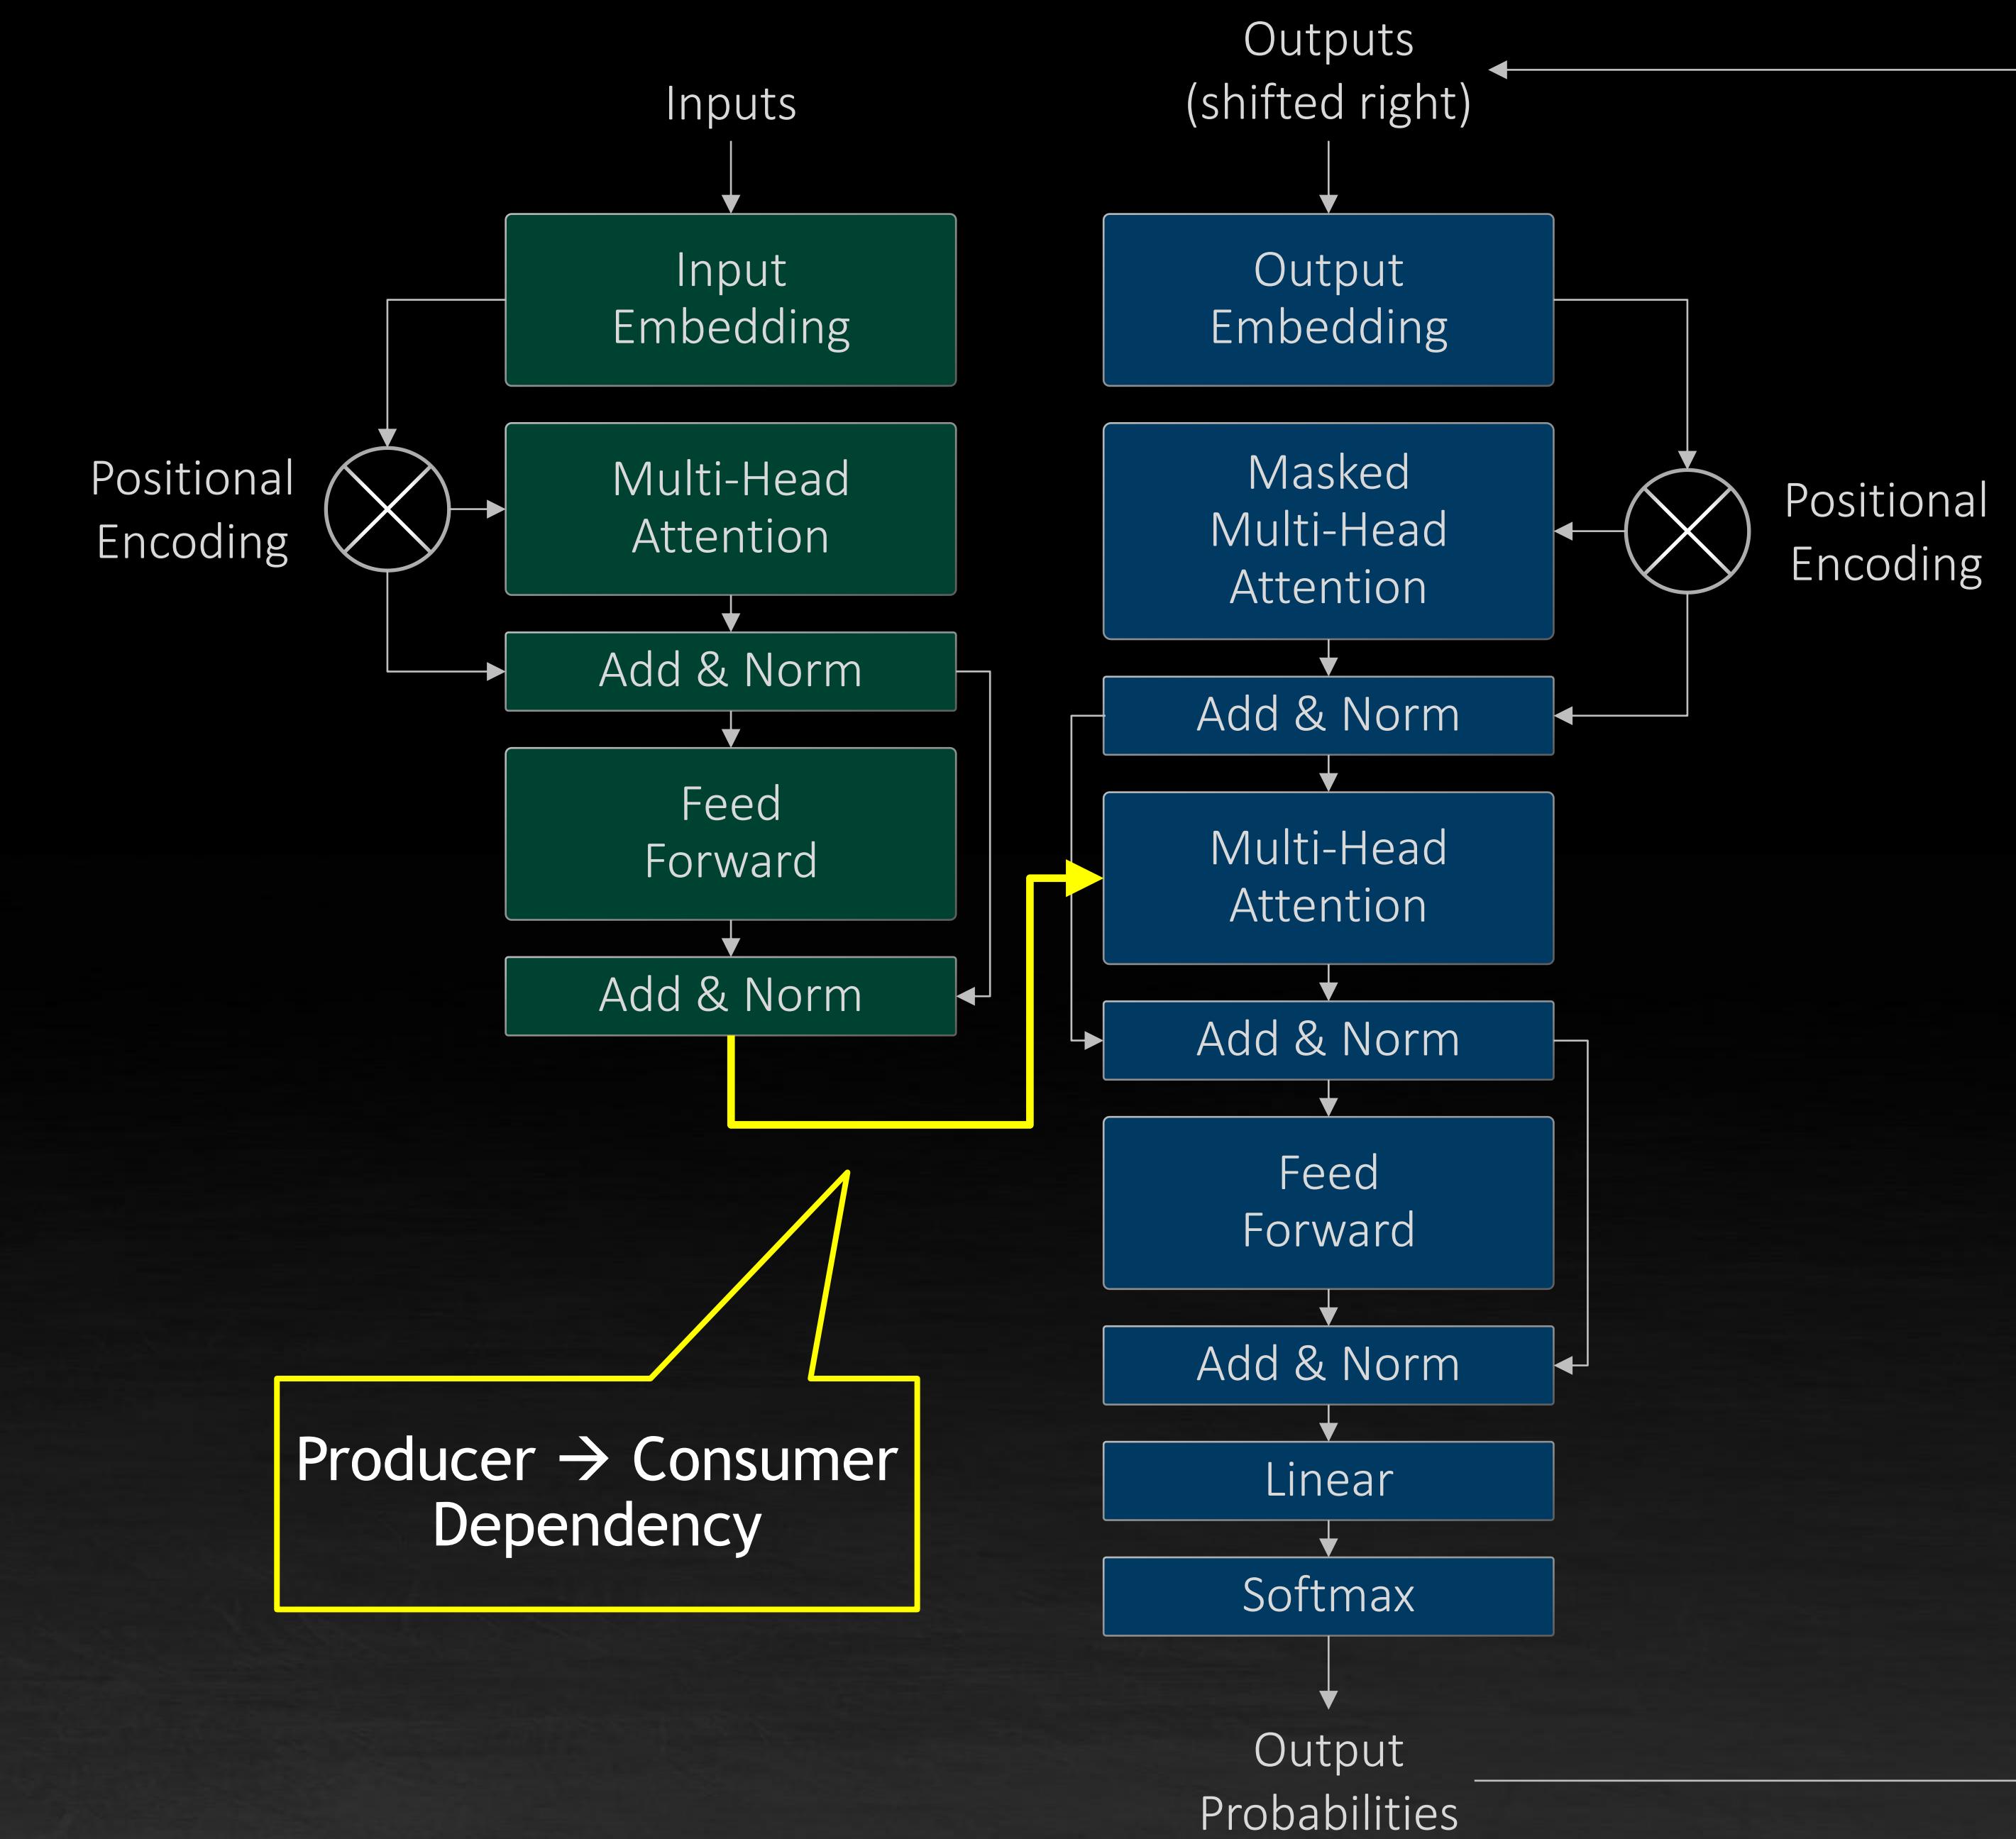
EXAMPLE: PRODUCER/CONSUMER TASK PARALLELISM AT ANY SCALE

Producer/Consumer can operate at every scale - grid, cluster or block



Specialise **groups of warps** within each **block** into “producer” and “consumer” roles

Specialise	Into	Exchange data via	Synchronize using
Grid	Clusters	Global memory	this_grid().sync()
Cluster	Blocks	Distributed-shared or global	this_cluster().sync()
Block	Warps	Shared or global memory	this_thread_block().sync()



EXAMPLE: LONGSTAFF SCHWARTZ PRICING MODEL

Producer/Consumer task parallelism

What is Longstaff Schwartz?

Pricing technique for Quantitative Finance

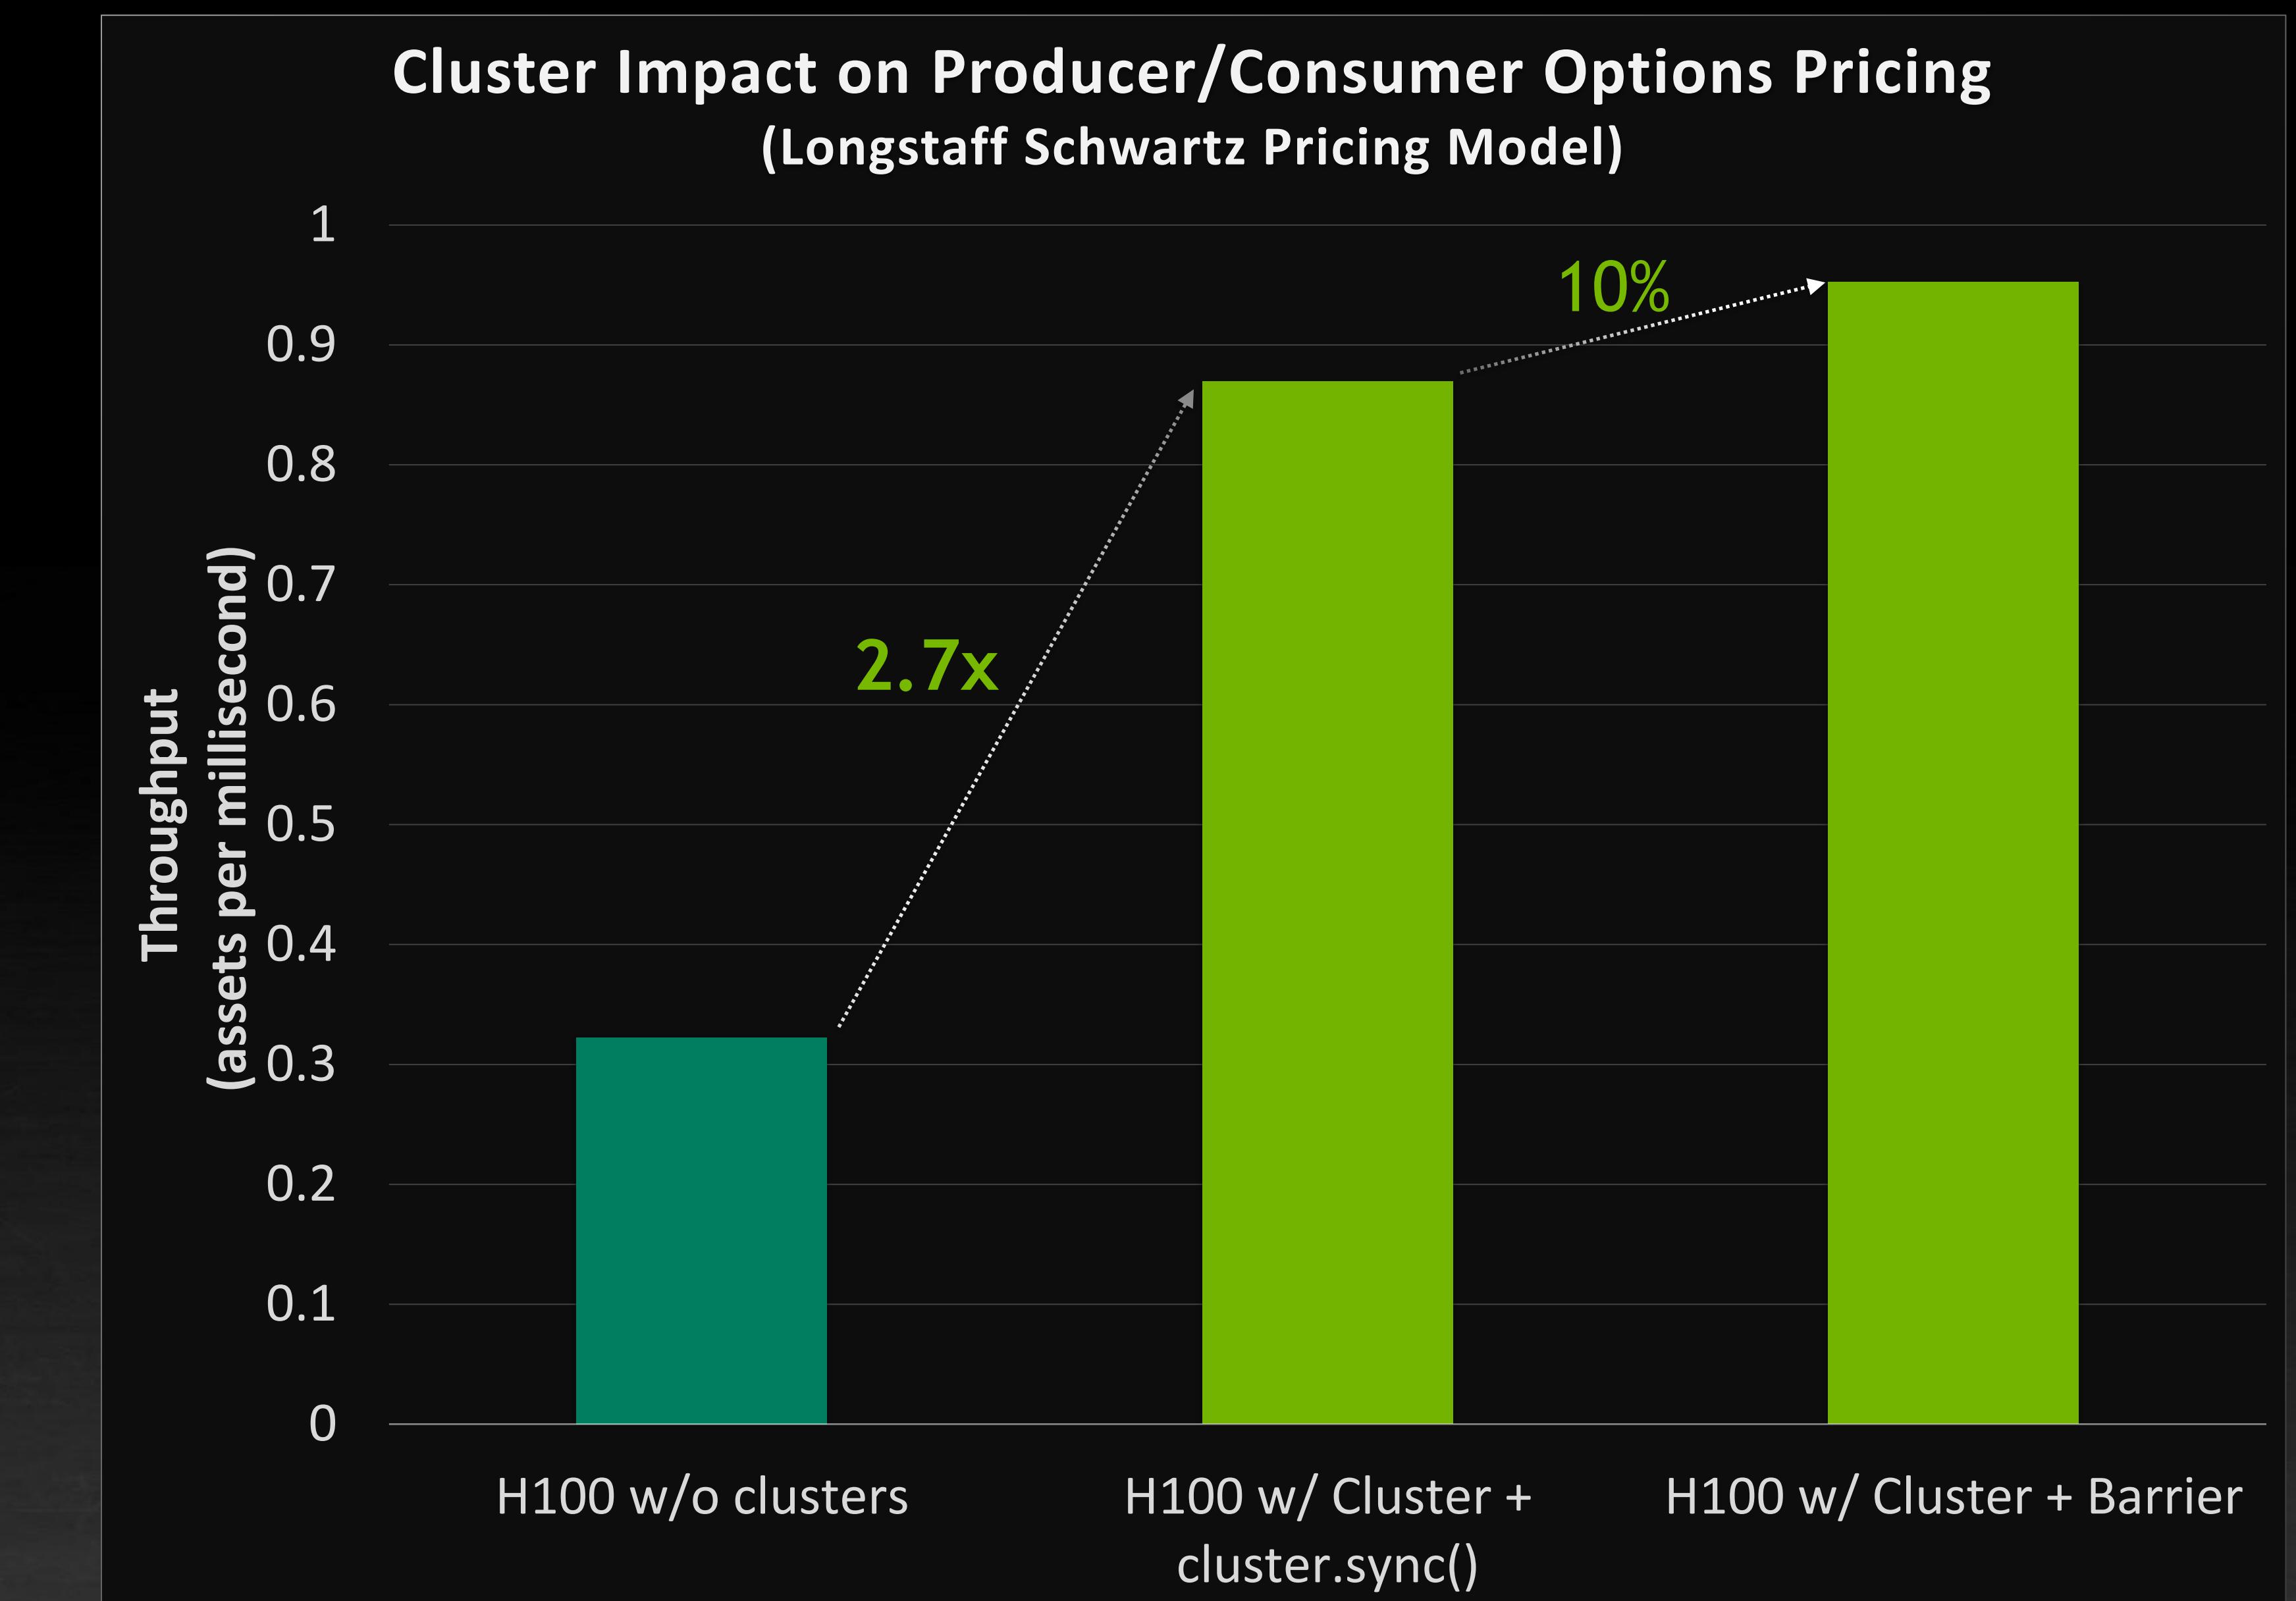
Monte-Carlo simulation with backward iteration to determine financial option values

Analysis of Example

Larger capacity of **Distributed shared memory** enables new approach to programming efficient kernels

Thread Block level **Producer Consumer models can be now implemented in shared memory** with fast synchronization

Hardware accelerated cluster synchronization is easy to use and efficient



CUDA C++ SUPPORT FOR 128-BIT INTEGERS

```
int main() {
    __int128 a, b, c;

    // Init "a" as product of two 64-bit numbers
    a = 197746253418ULL;
    a *= 4060883912384ULL;

    // Perform an arithmetic operation with "a"
    b = a + 10;

    // Launch a kernel, passing in 128-bit int*'s
    cudaMemcpy(device_a, &a, sizeof(a));
    add<<< ... >>>(device_a, device_c);
    cudaMemcpy(&c, device_c, sizeof(c));

    // Verify same result on GPU and CPU
    assert(b == c);
}
```

CUDA 11.7 supports `__int128` across toolchain & libraries
Usable with compatible host compilers (gcc, Clang, ICC)
or with run-time compilation (NVRTC)

Full release

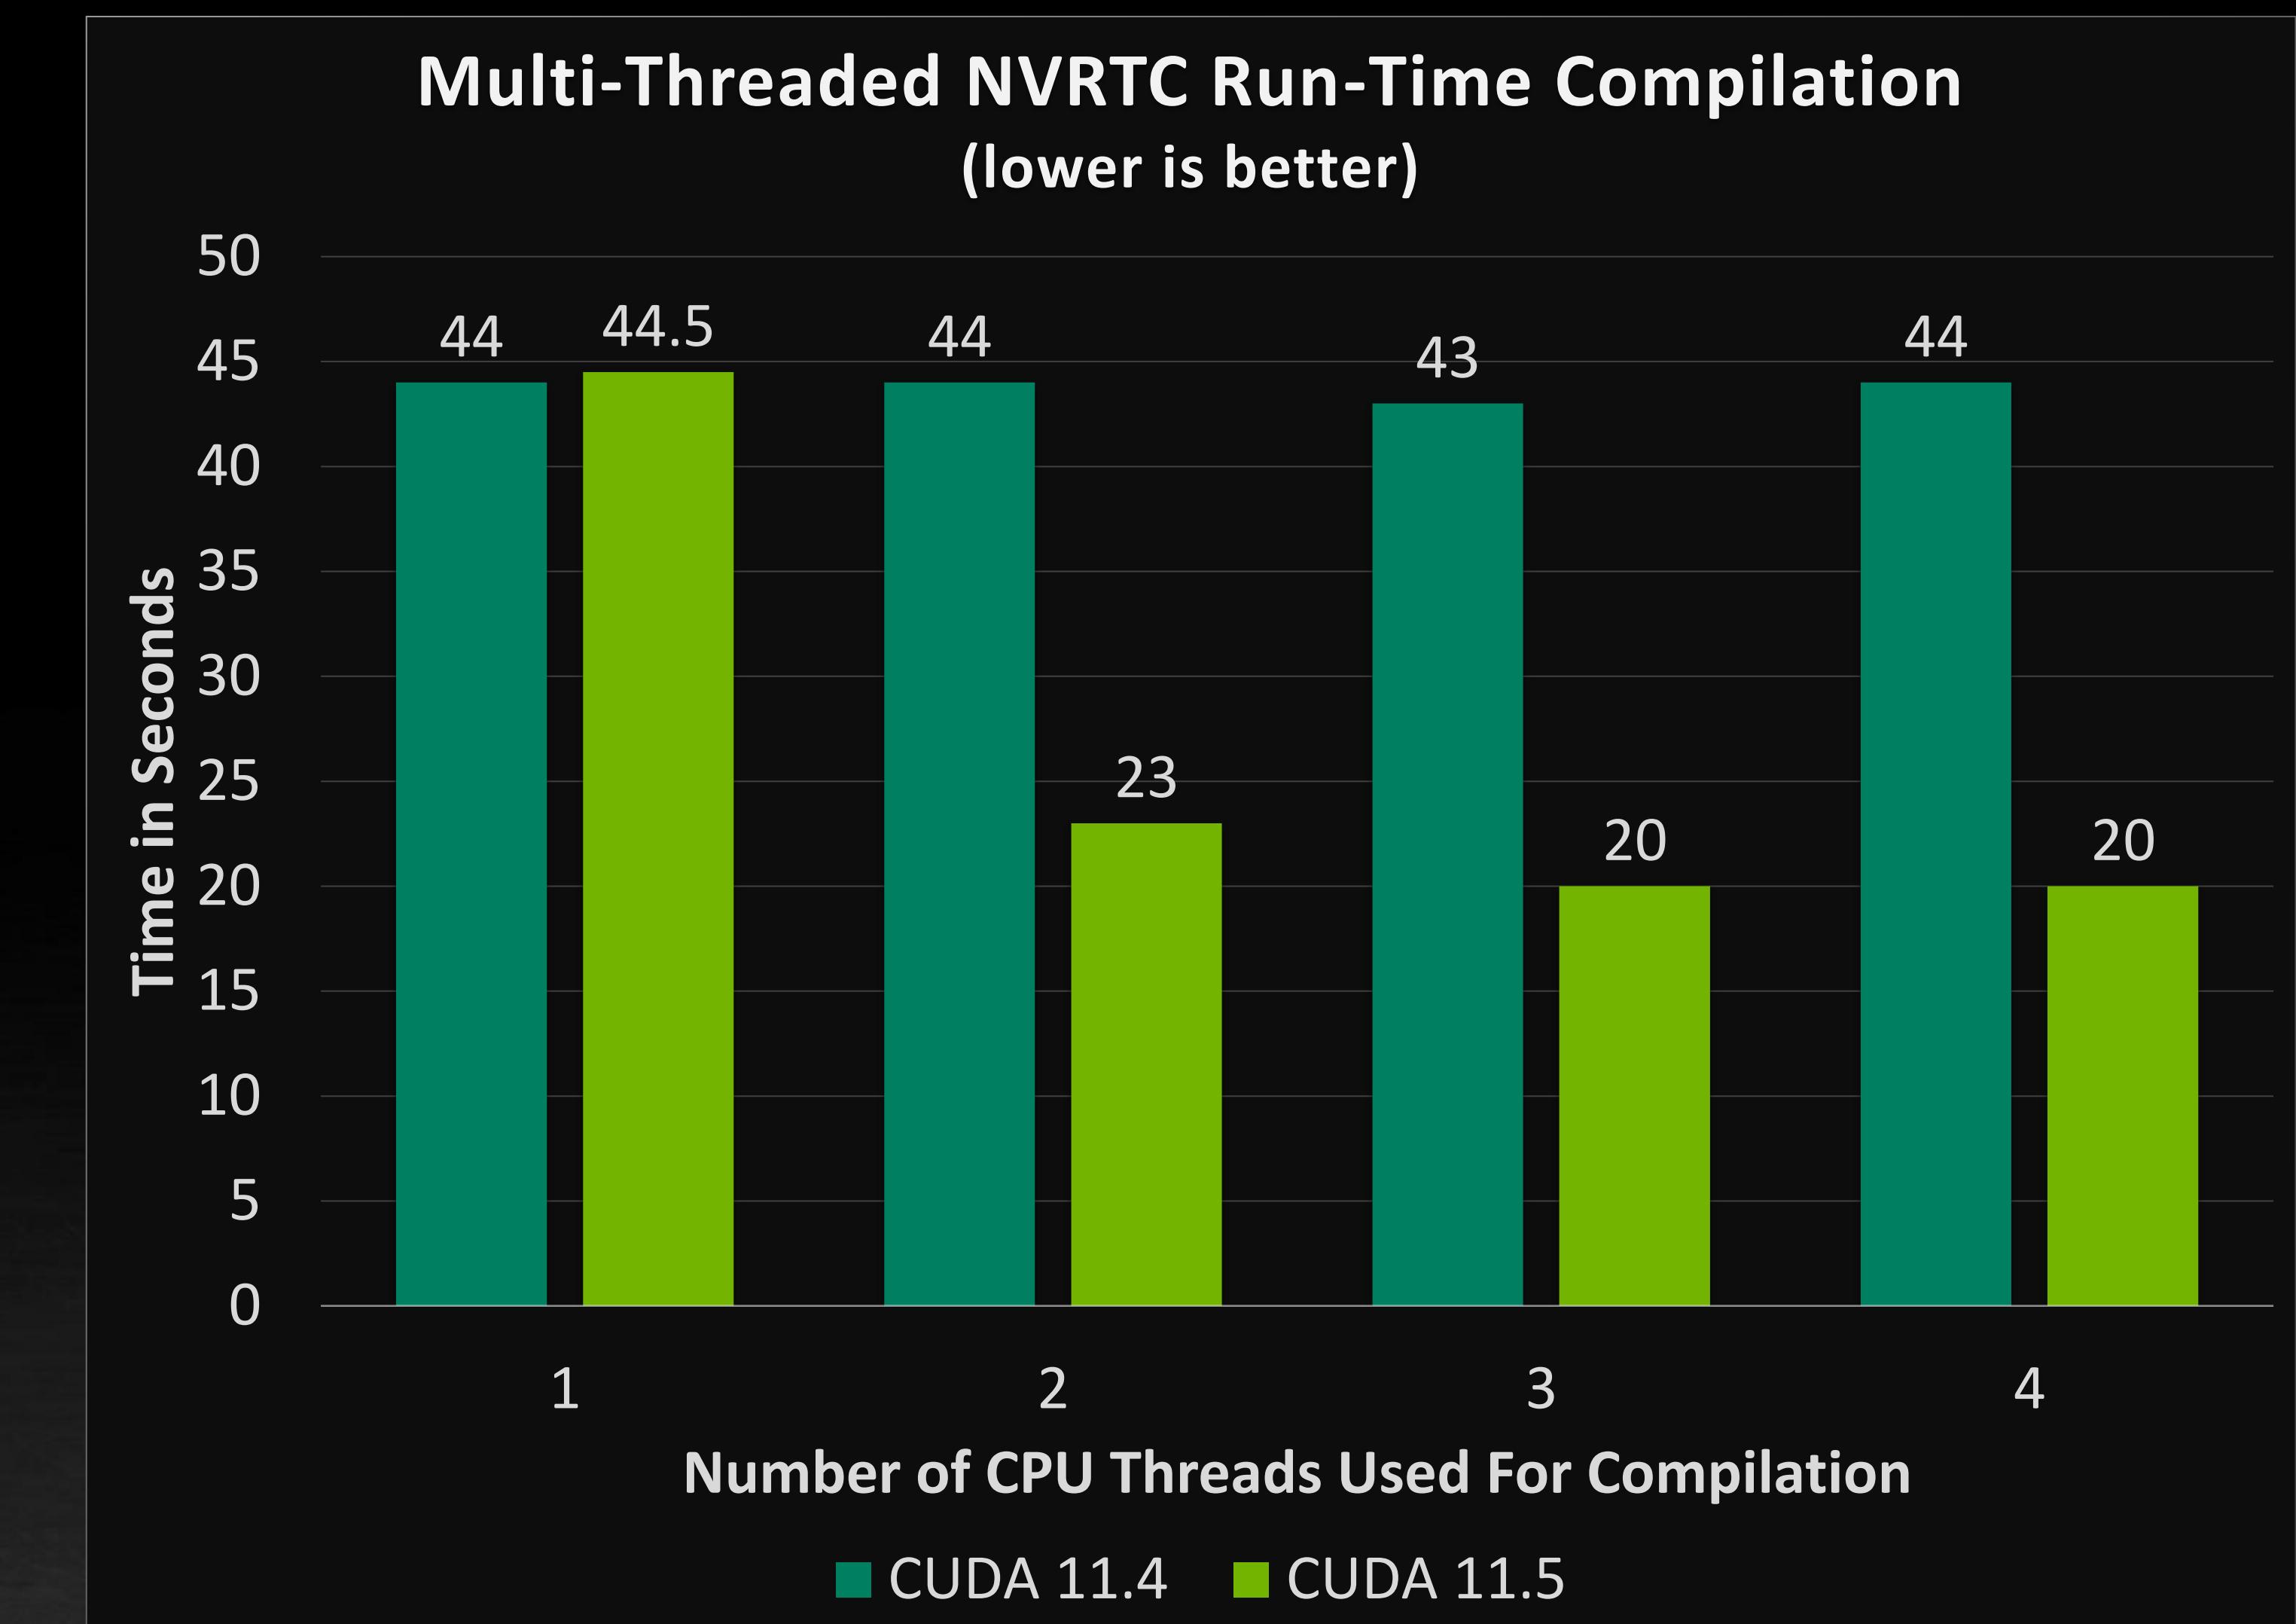
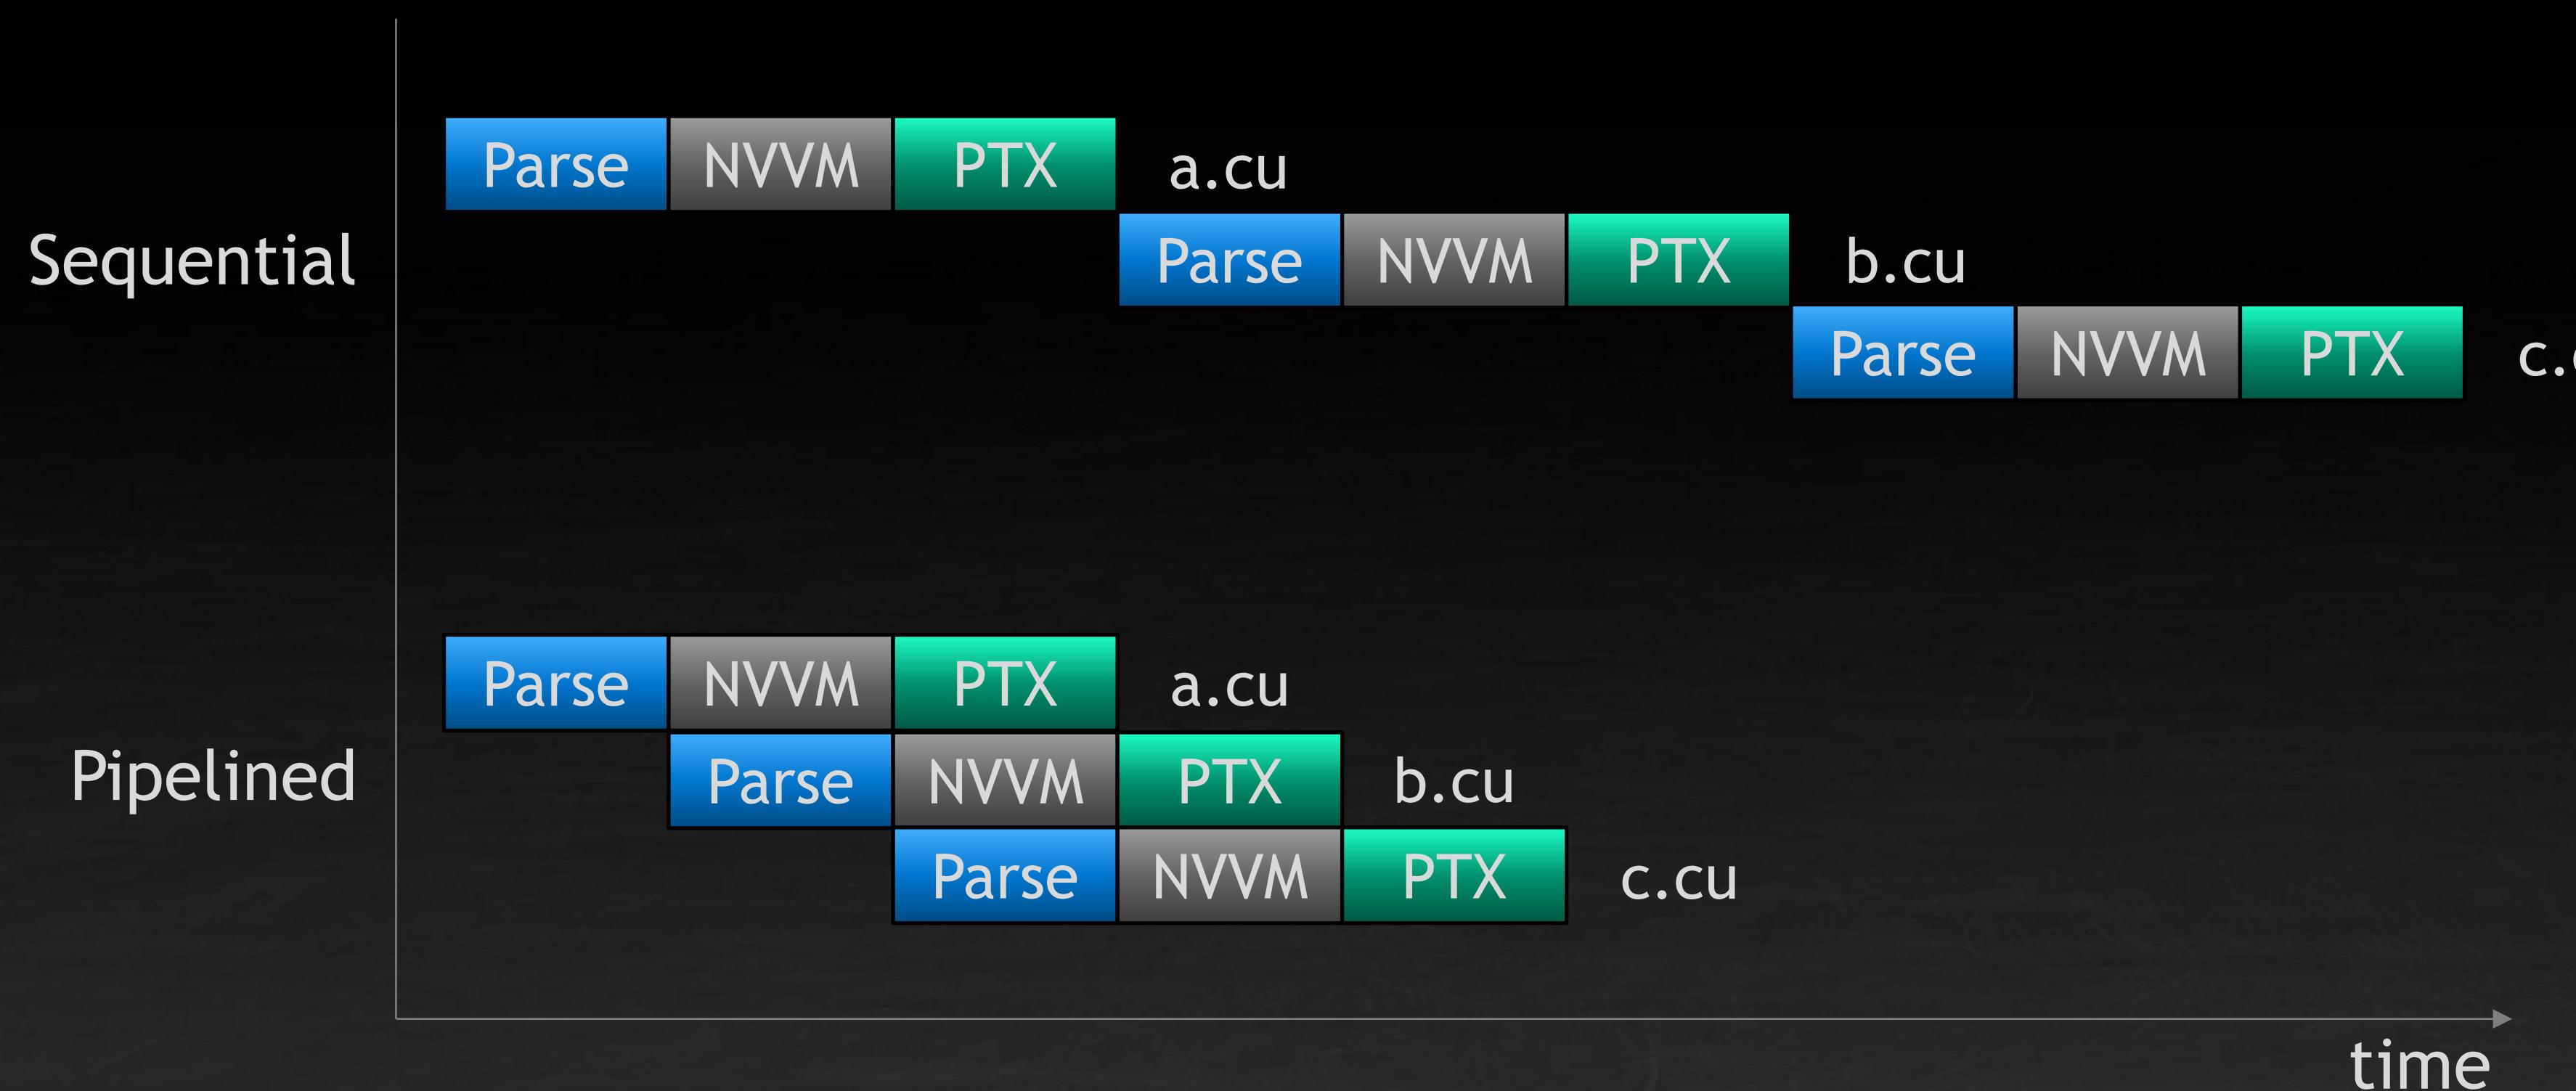
Supports arithmetic, logical & bitwise operations, math operations, library support, dev tools, etc.

```
__global__ void add (__int128 *in1, __int128 *out) {
    // Perform the same arithmetic operation
    // as we did in host code
    *out = *in1 + 10;
}
```

NVRTC MULTI-THREADED COMPIRATION

Moving from **global** to **per-stage** locking in NVRTC & PTX JIT

Enables concurrent compilation when compiling with multiple CPU threads



RECENT COMPILER UPDATES

C++20 Support in CUDA 11.7

C++20 supported as preview with these host compilers:

GCC 10+, Clang 10+, nvc++ 20.7+

MSVC support coming in next release

Programmatic management of diagnostics

```
#pragma nv_diag_suppress error_number[,error_number...]
#pragma nv_diag_warning error_number[,error_number...]
#pragma nv_diag_error error_number[,error_number...]
#pragma nv_diag_default error_number[,error_number...]
#pragma nv_diag_once error_number[,error_number...]
#pragma nv_diag_push
#pragma nv_diag_pop
```

Grid-private “Hard” Constants

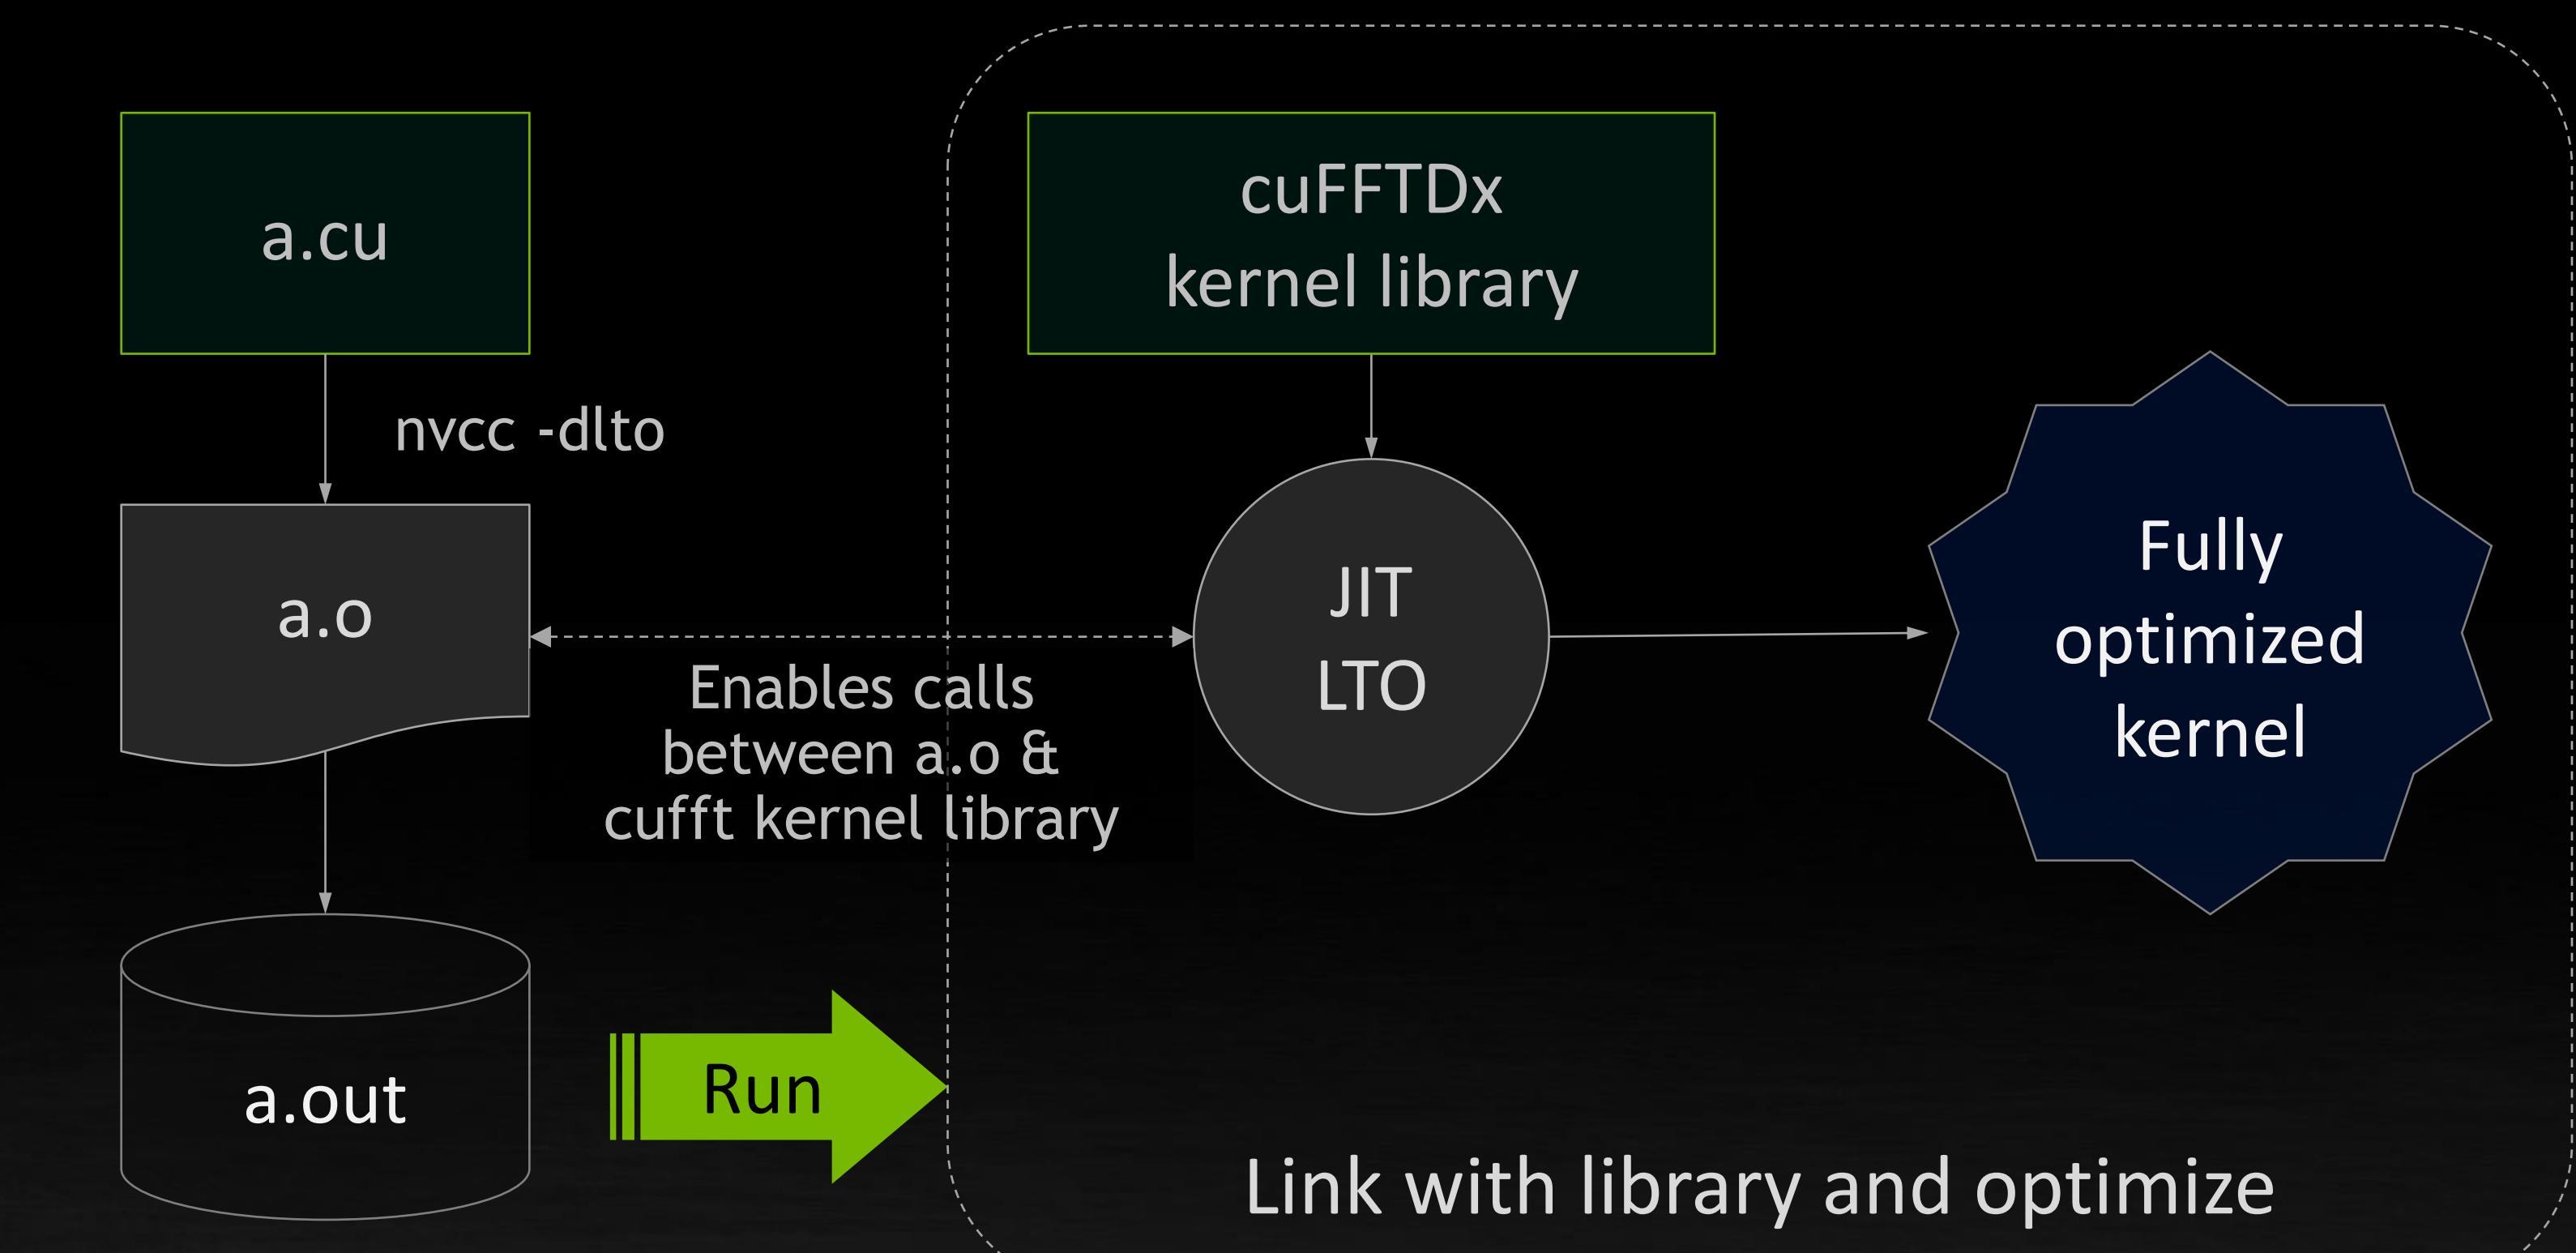
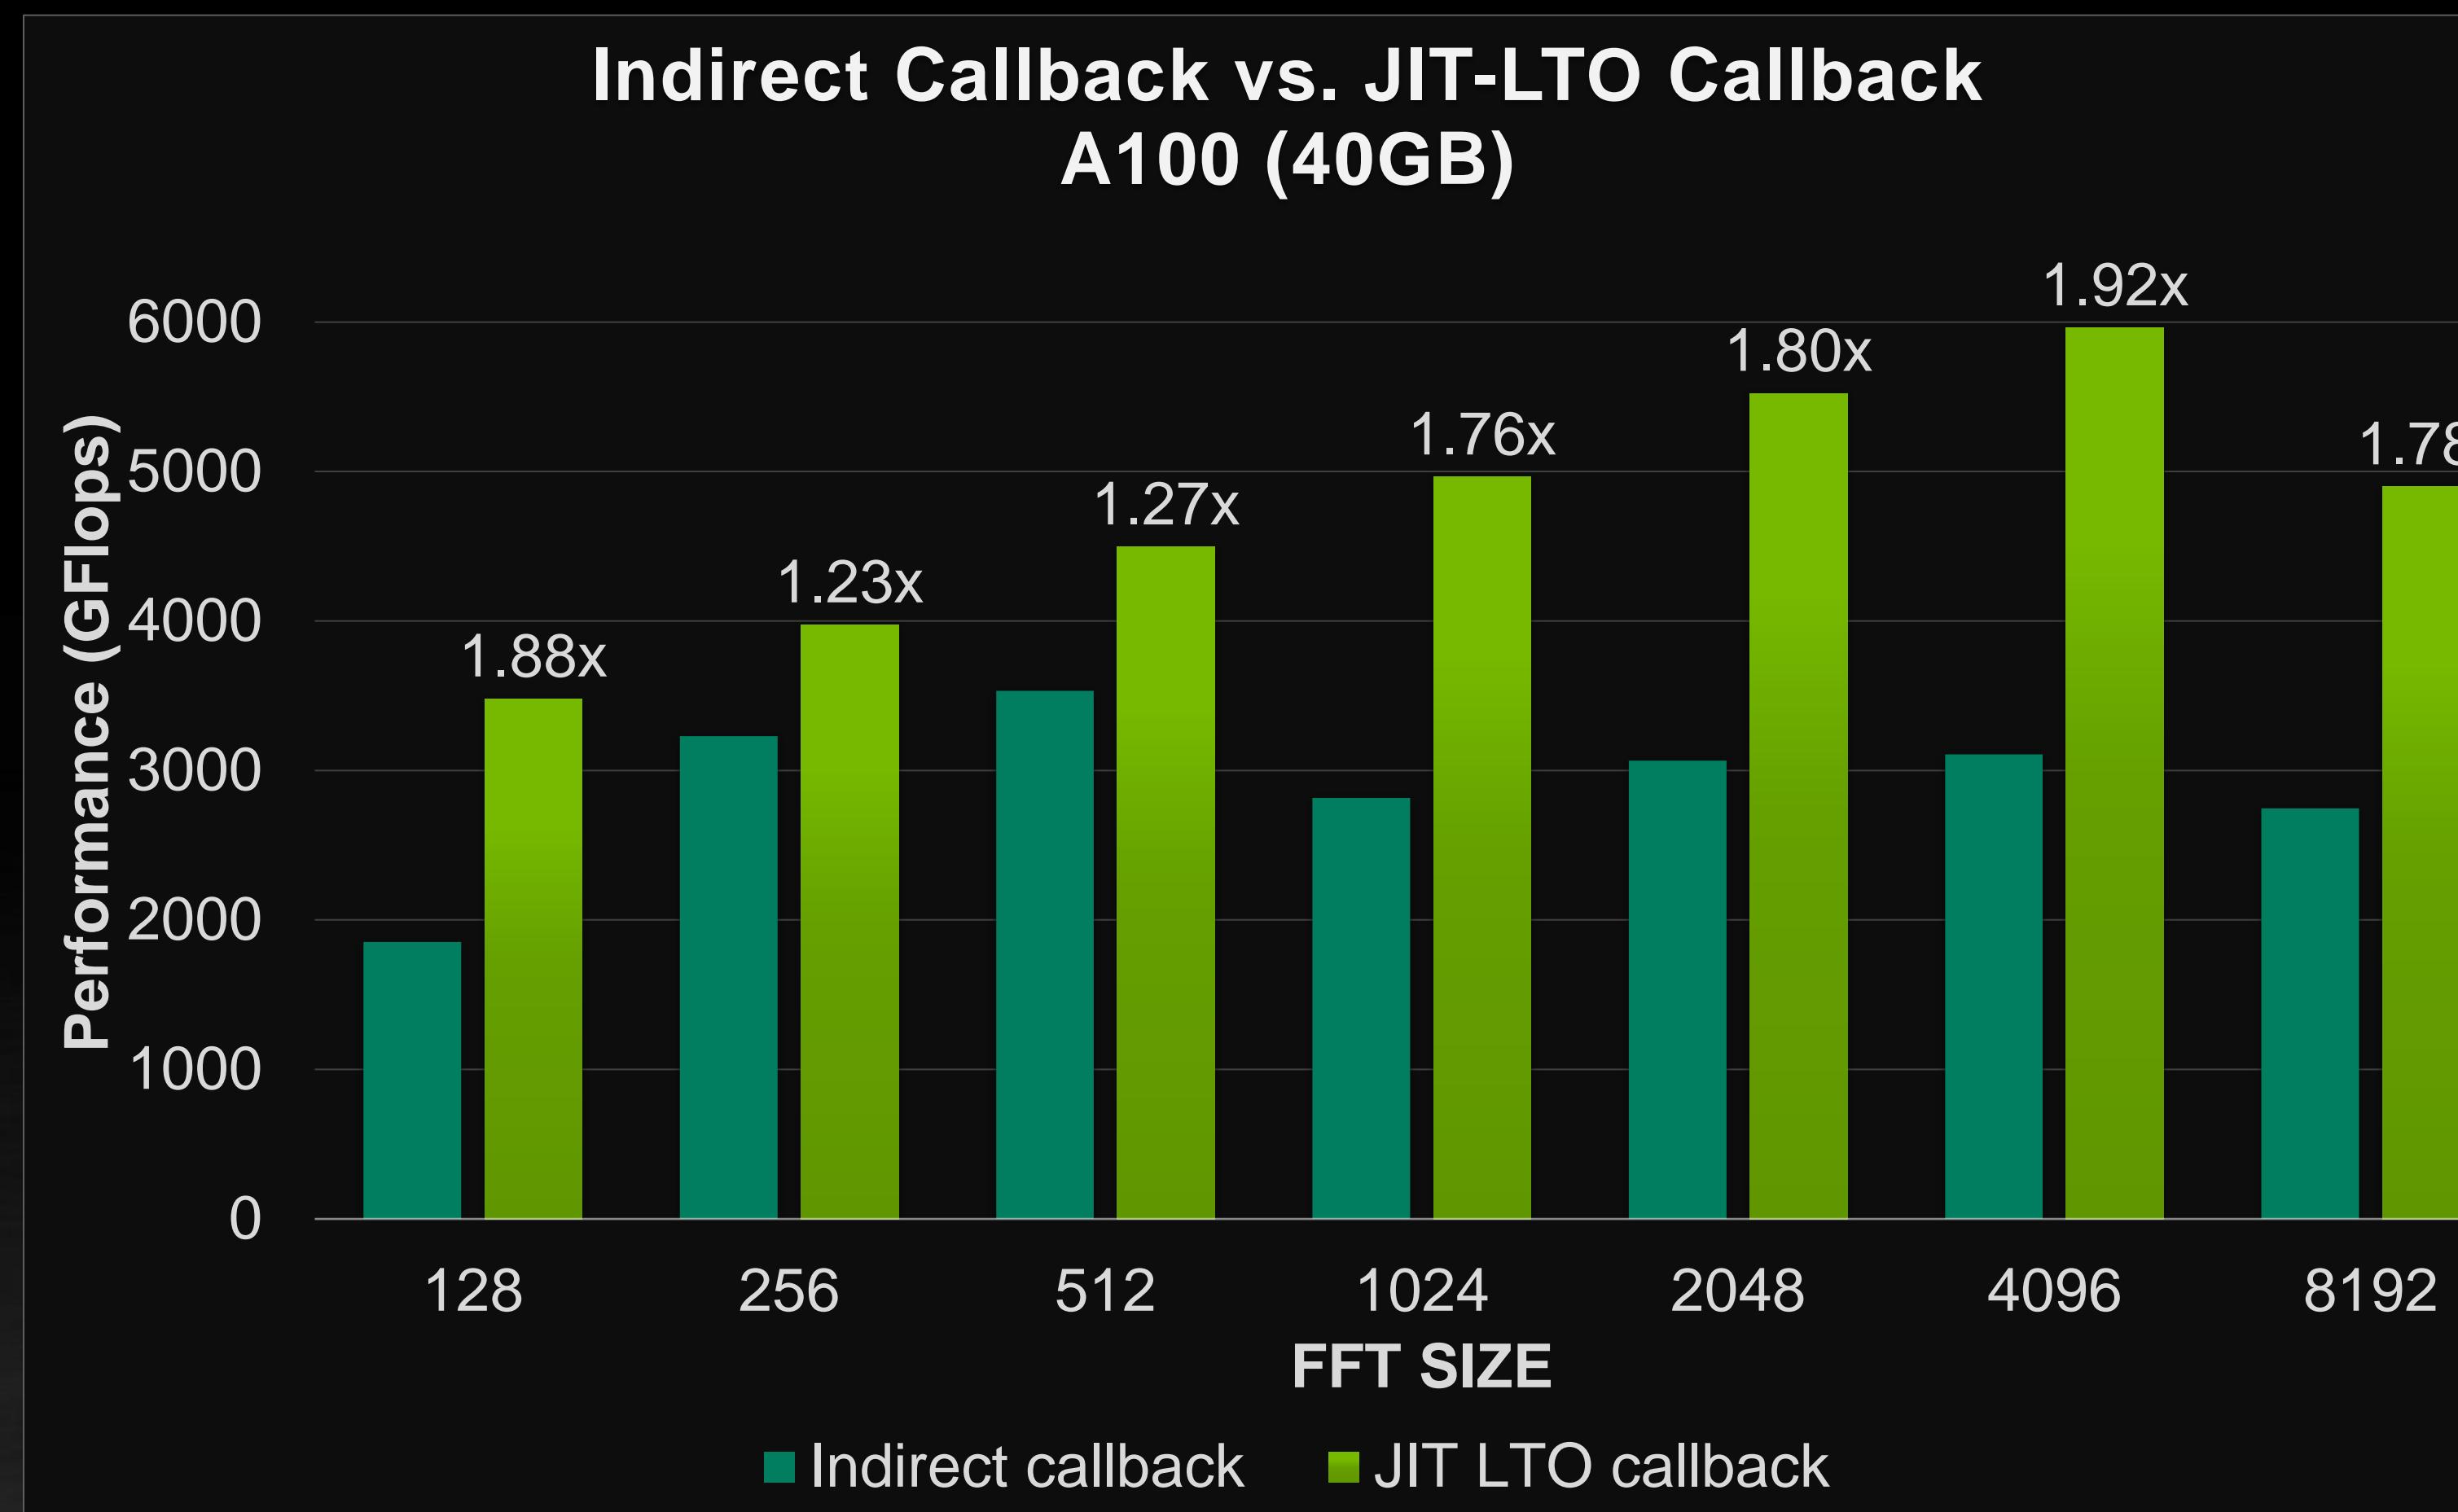
New **`__grid_constant__`** annotation for kernel parameters allows compiler to avoid creating per-thread copies of read-only parameters

```
__global__ void kernel( __grid_constant__ const struct param_t p ) {
    foo( &p );
}
```

New nvcc target -arch options

-
- arch=all generates codes for all supported architectures
 - arch=all-major generates codes for all supported major architectures
 - arch=native generates codes for all visible GPUs on the system
-

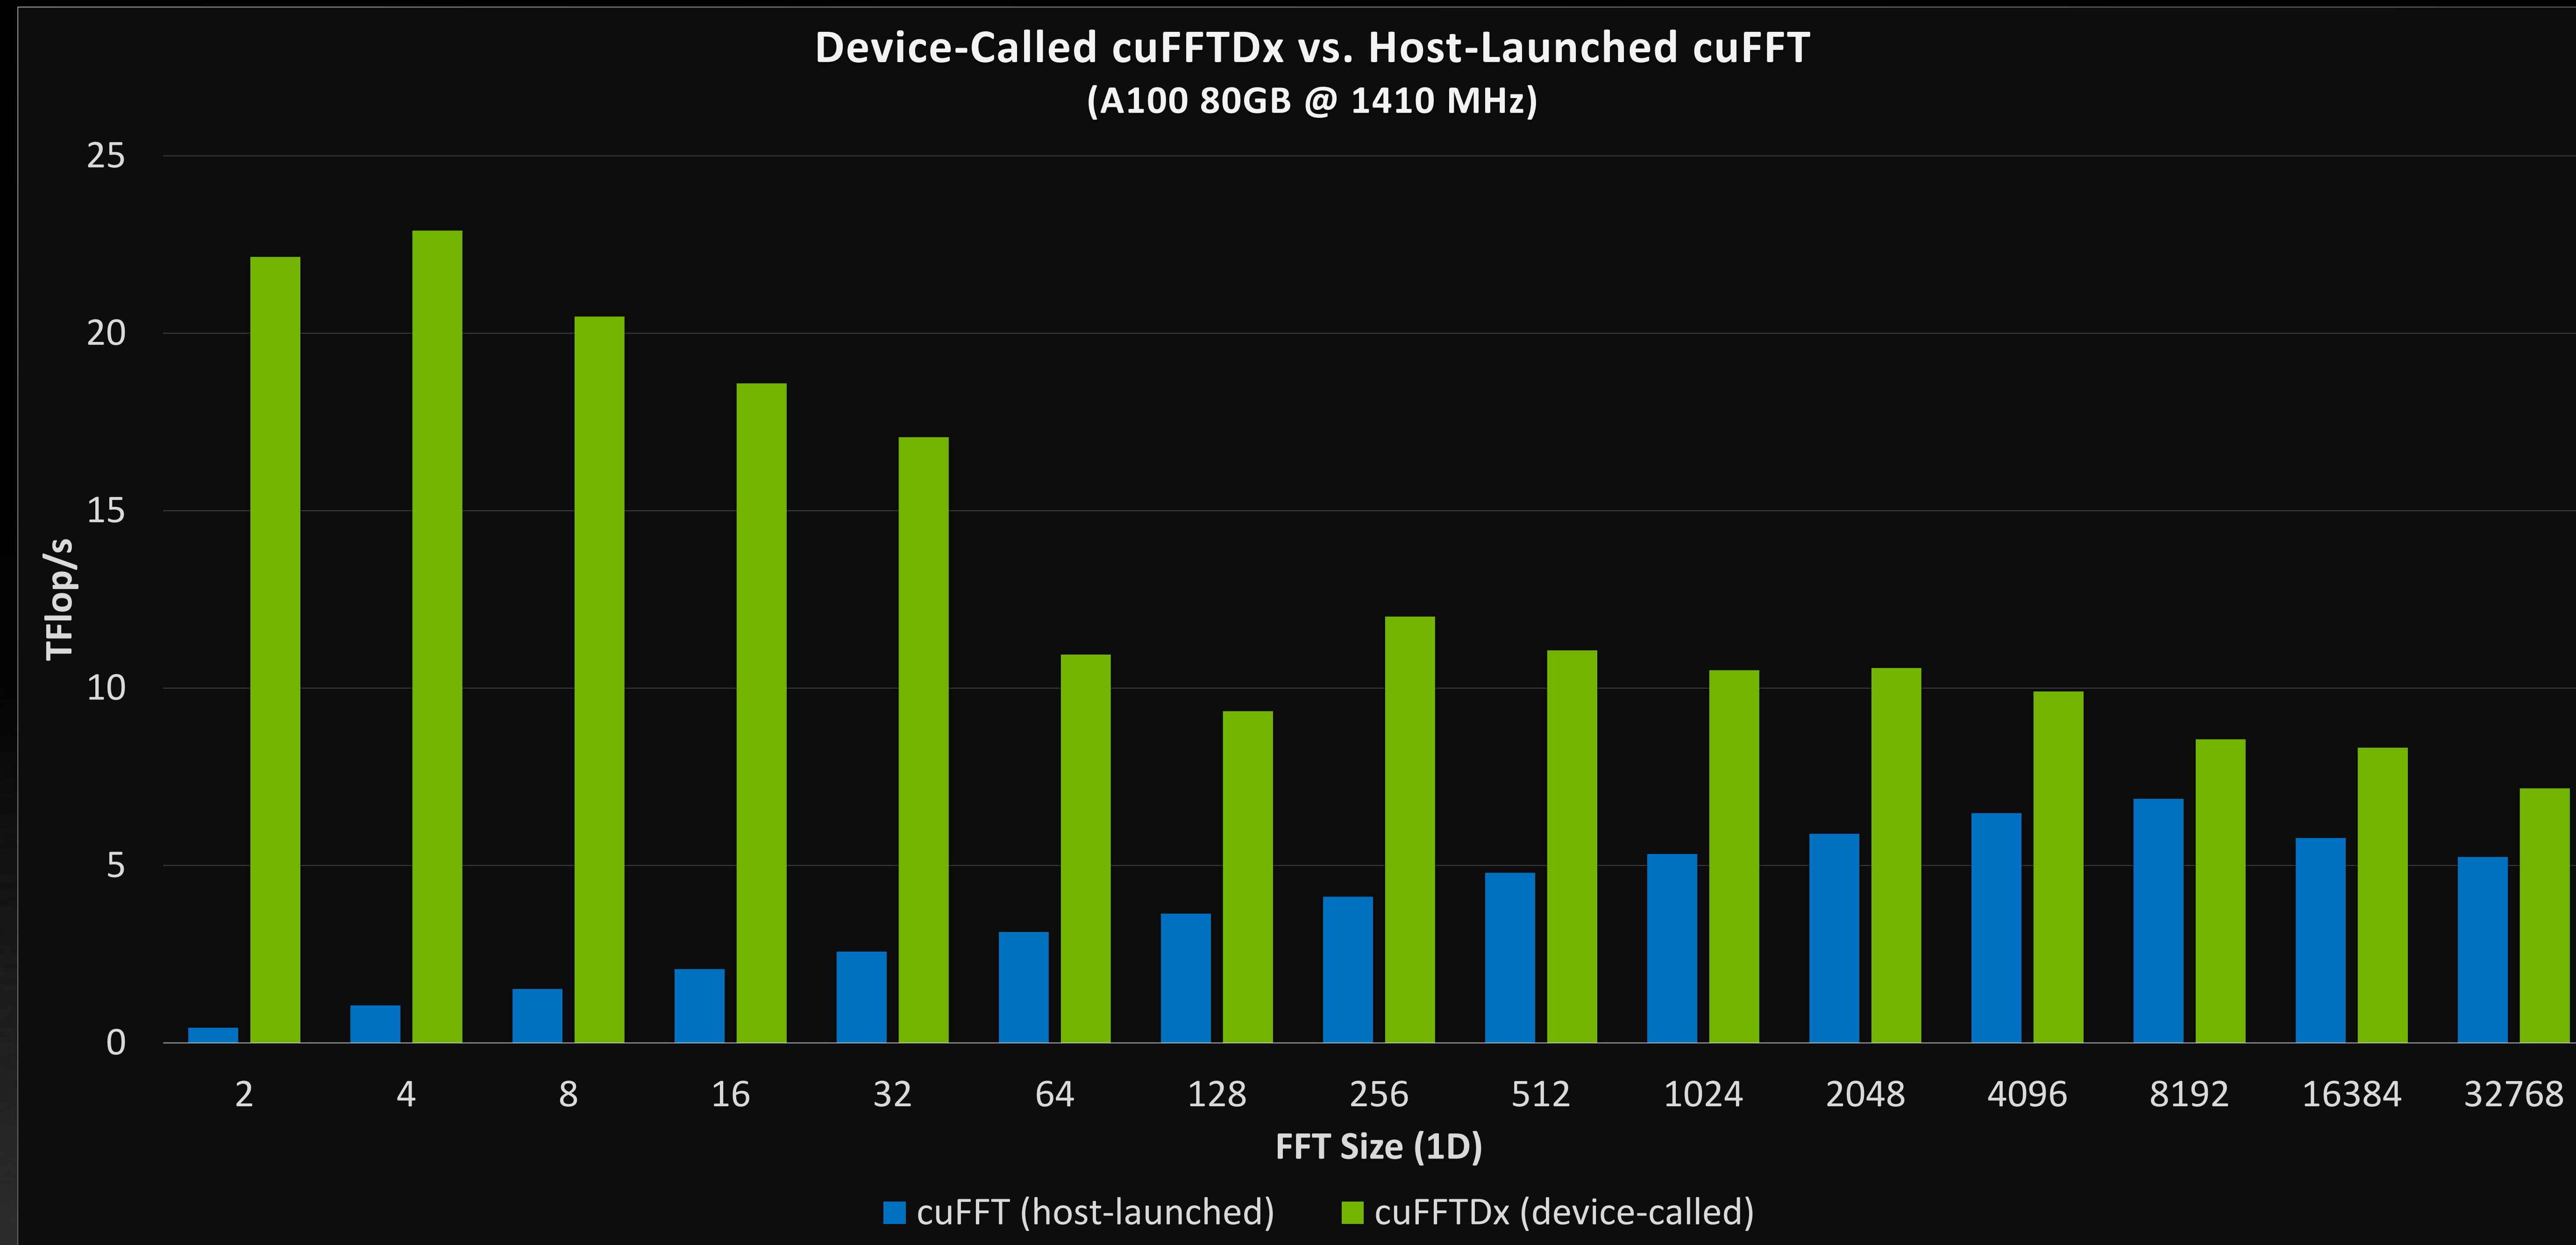
JIT LINKING WITH LINK-TIME OPTIMIZATIONS



JIT-LTO
Object files are combined
and optimized at runtime

MATHS LIBRARIES DEVICE EXTENSIONS

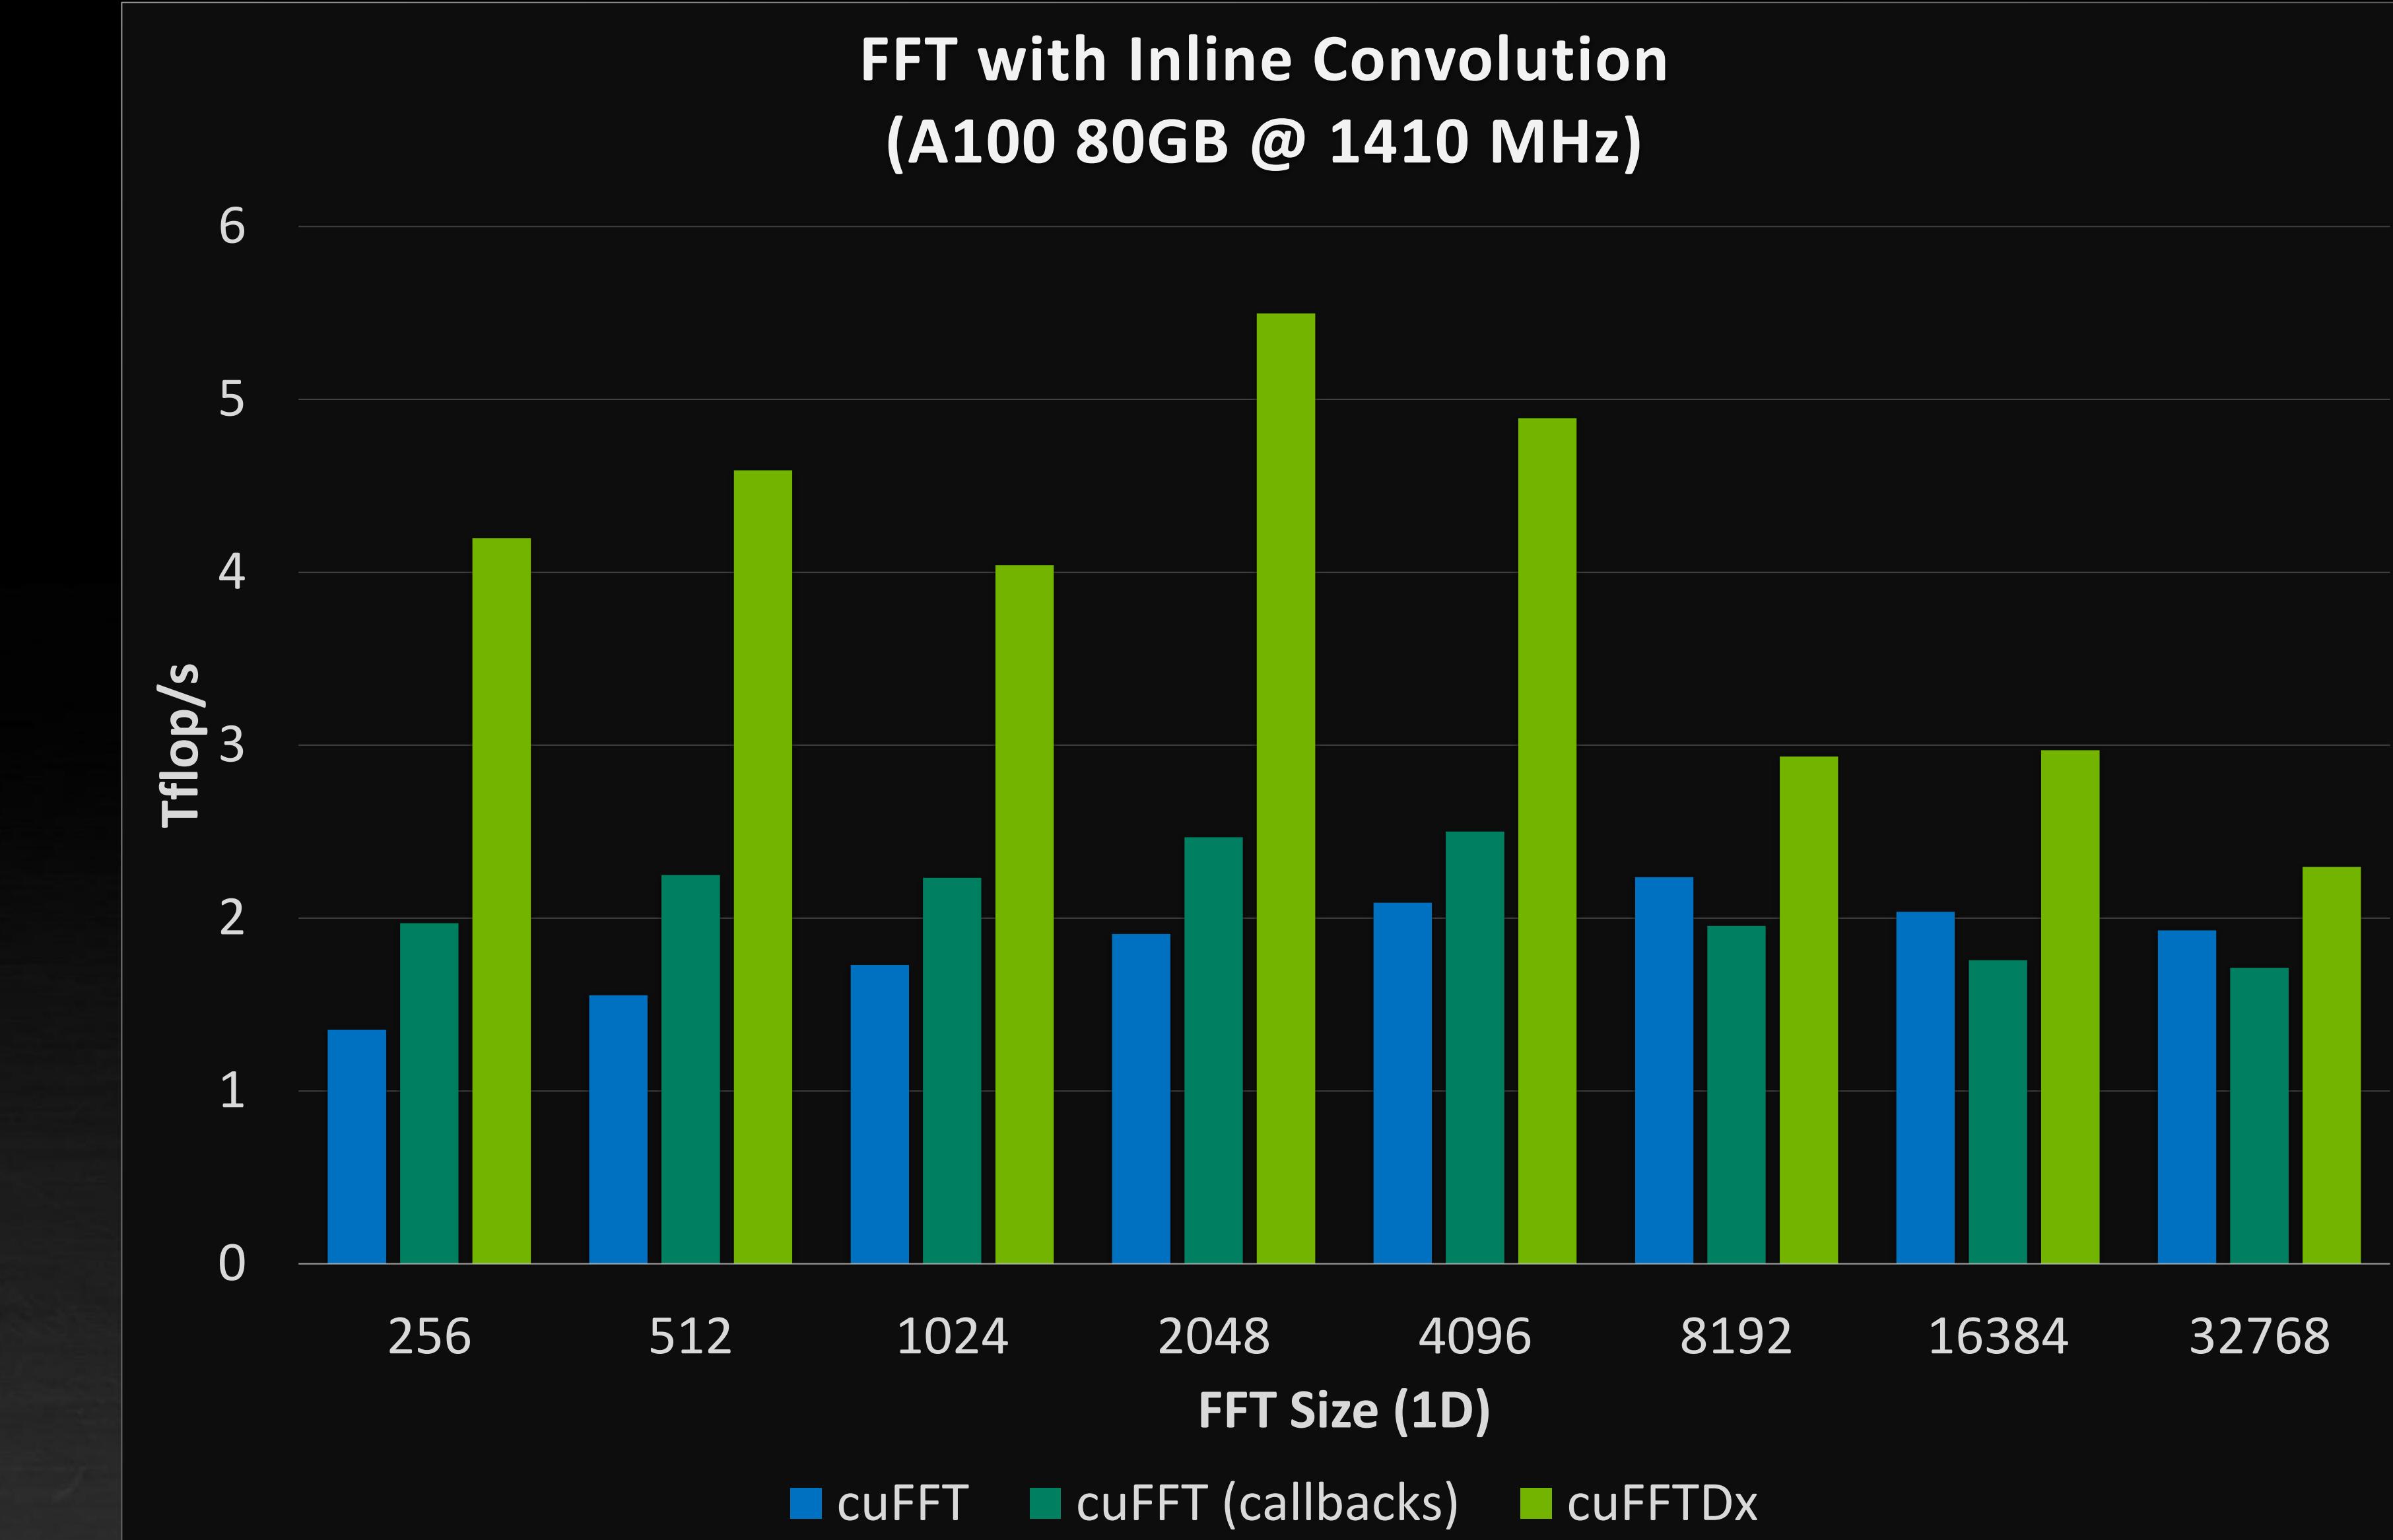
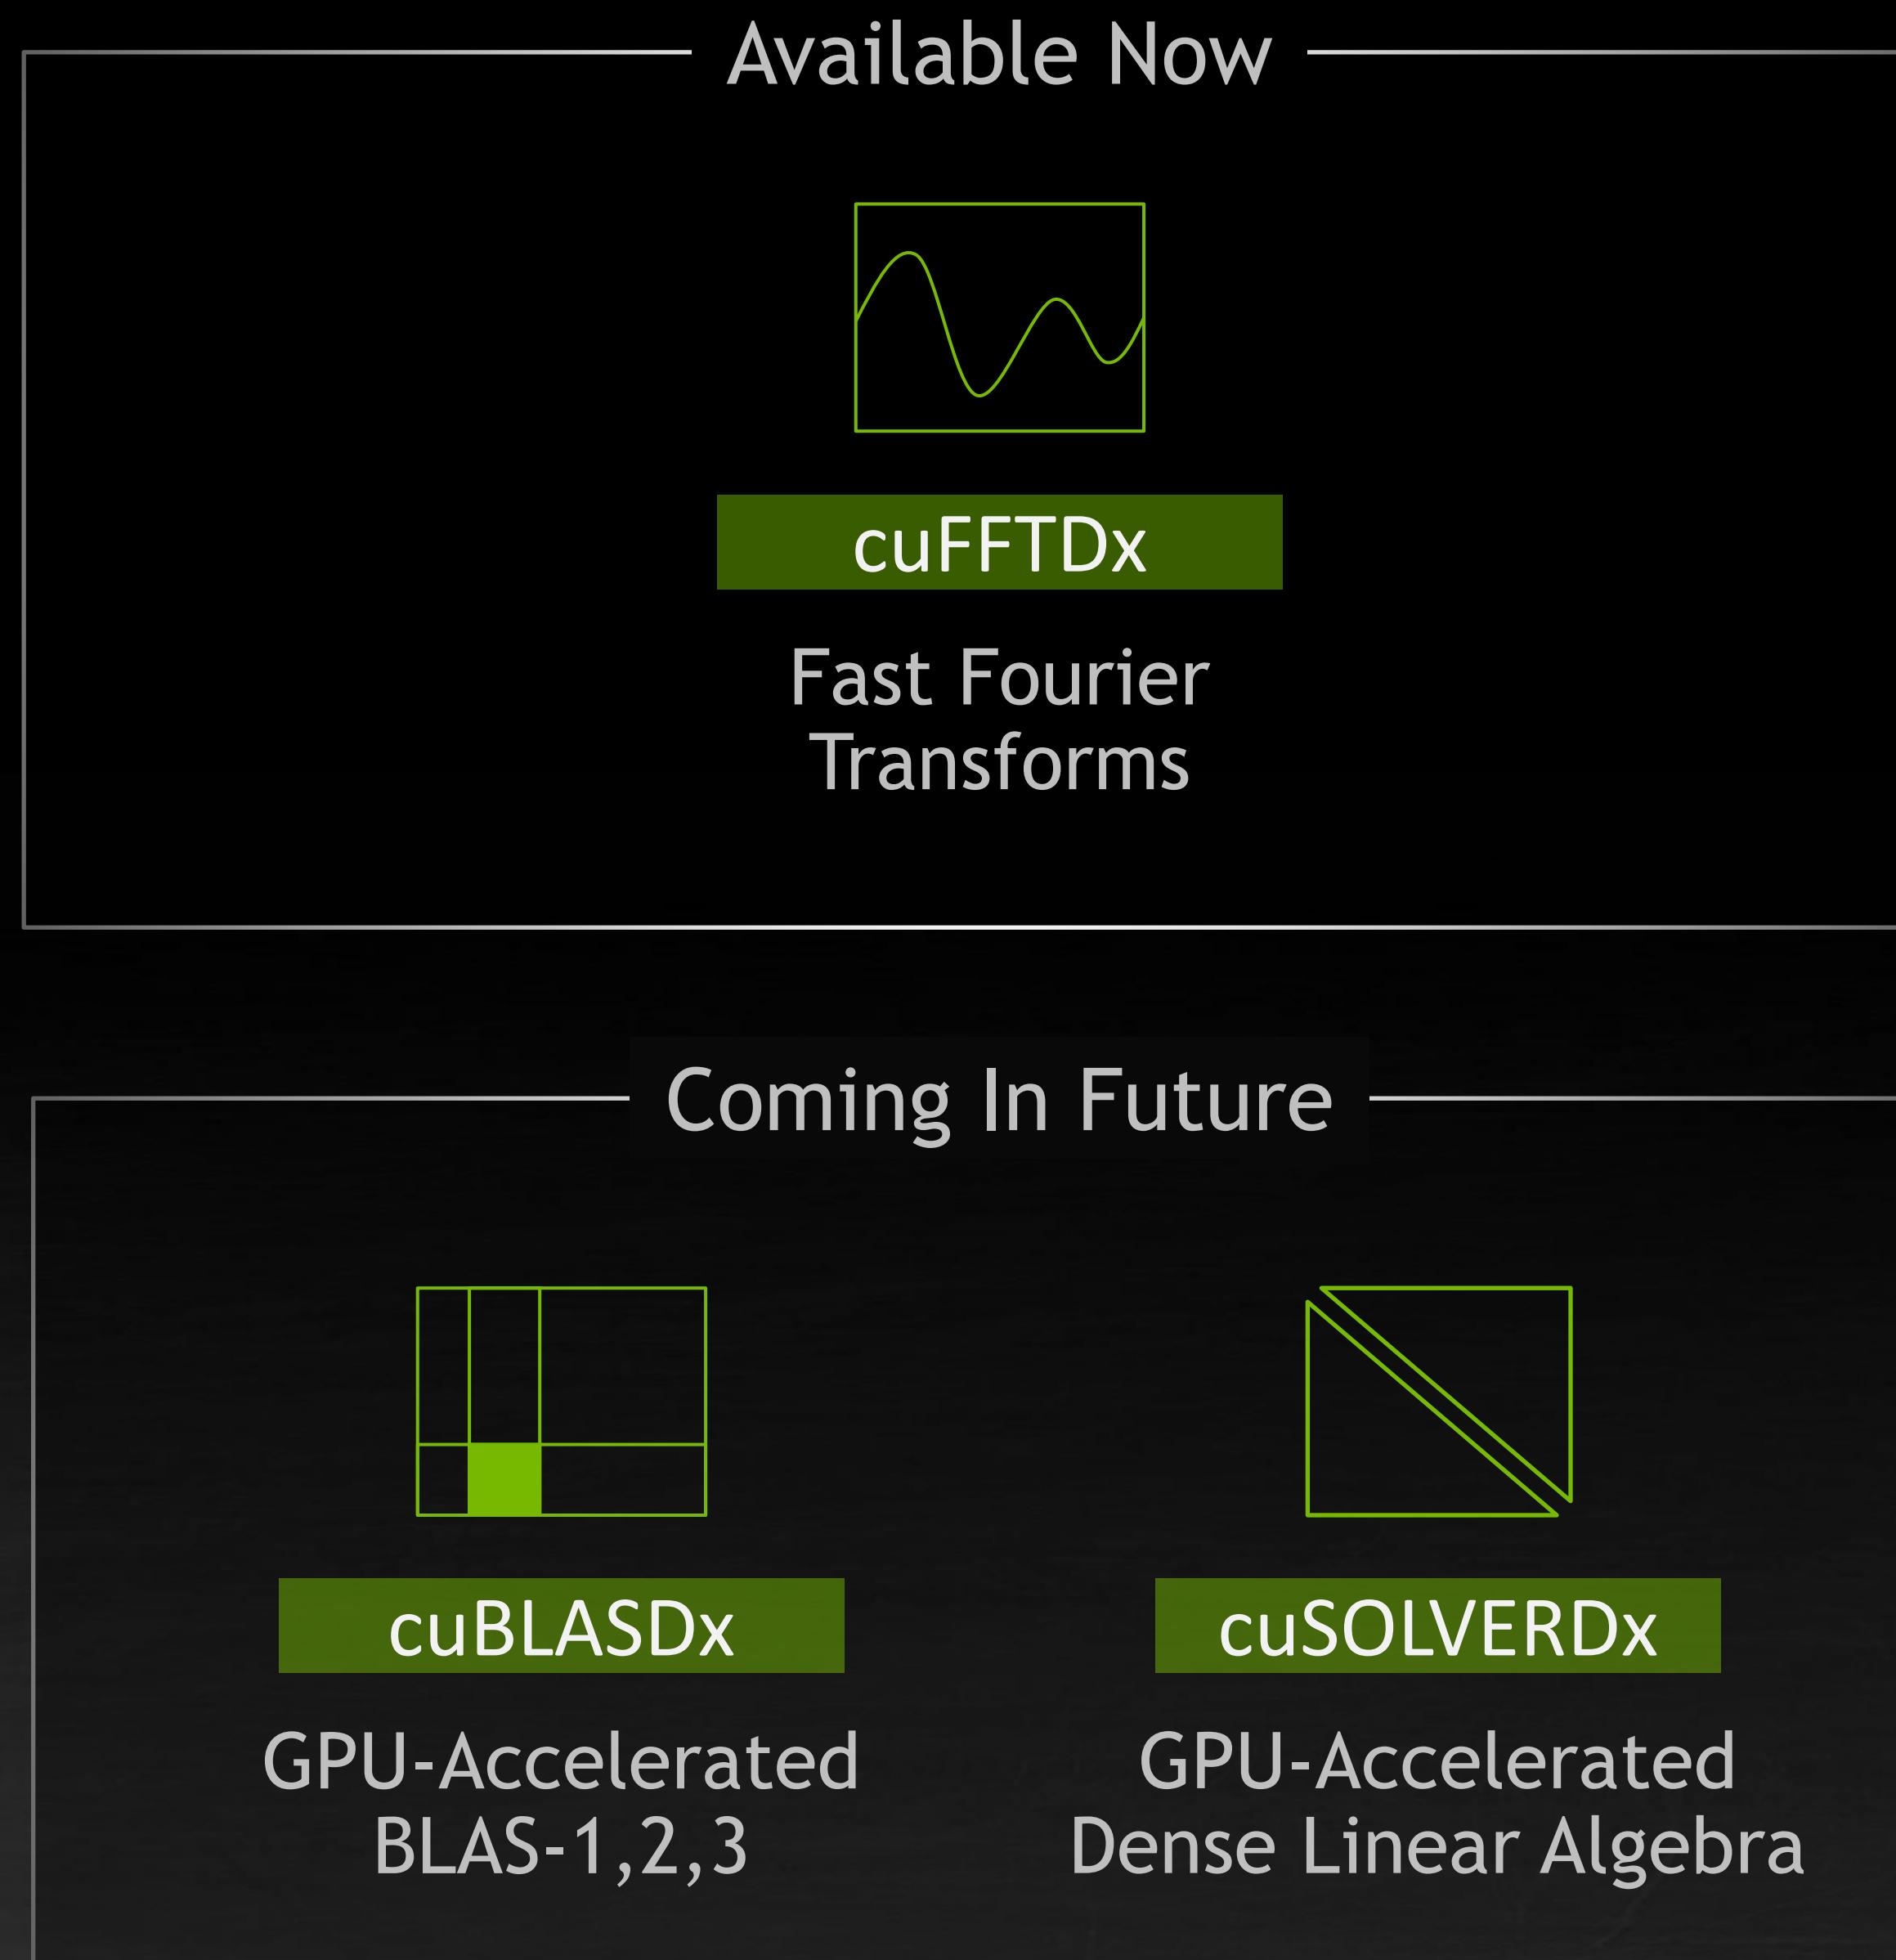
Calling cuFFTDx from device code significantly speeds up small problem sizes (blue bars)



MATHS LIBRARIES DEVICE EXTENSIONS

Invoke custom kernel code from within high performance Math Libraries (**green** bars)

Libraries With Device-Callable Interfaces



Connect with the Experts: NVIDIA Math Libraries [CWE41721]



CUTLASS

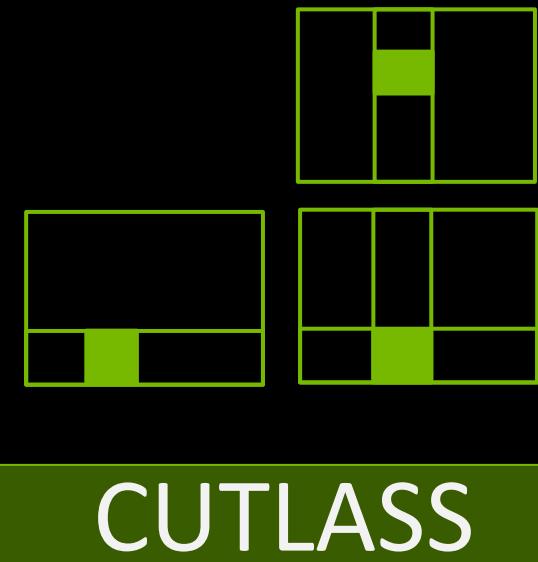
CUDA C++ Template Library for Deep Learning and High Performance Computing

**Optimal CUDA C++ matrix operation templates
at all scopes and scales**

Tensor-core accelerated matrix & tensor operations & more
GEMM, convolution, reduction, fused input & output operations
Open-source header file C++ template library

CUDA Device Code

GPU Hardware



Cutlass is compiled in to CUDA kernels, interfacing directly with GPU hardware

Tensor-core accelerated operations at all granularities

Device	{ GEMM, Convolution, Reductions } x { all data types } x { SIMT, Tensor Cores } x { all architectures }
Kernel	GEMM, Batched GEMM, Convolution, Reduction, Fused output operations, Fused input operations
Thread Block	Pipelined Matrix Multiply, Epilogue, Collective access to tensors, Convolution matrix access
Warp	Tensor Core Multiply-Add operations, Efficient access to permuted tensor layouts
Thread	Numeric conversion, <functional> operators on arrays, complex<T>, fast math algorithms
Architecture Intrinsic	Templates wrapping architecture-specific PTX instructions (e.g. mma, cp.async, ldmatrix, cvt)

CUTLASS: ACCELERATED SINGLE PRECISION USING TENSOR CORES

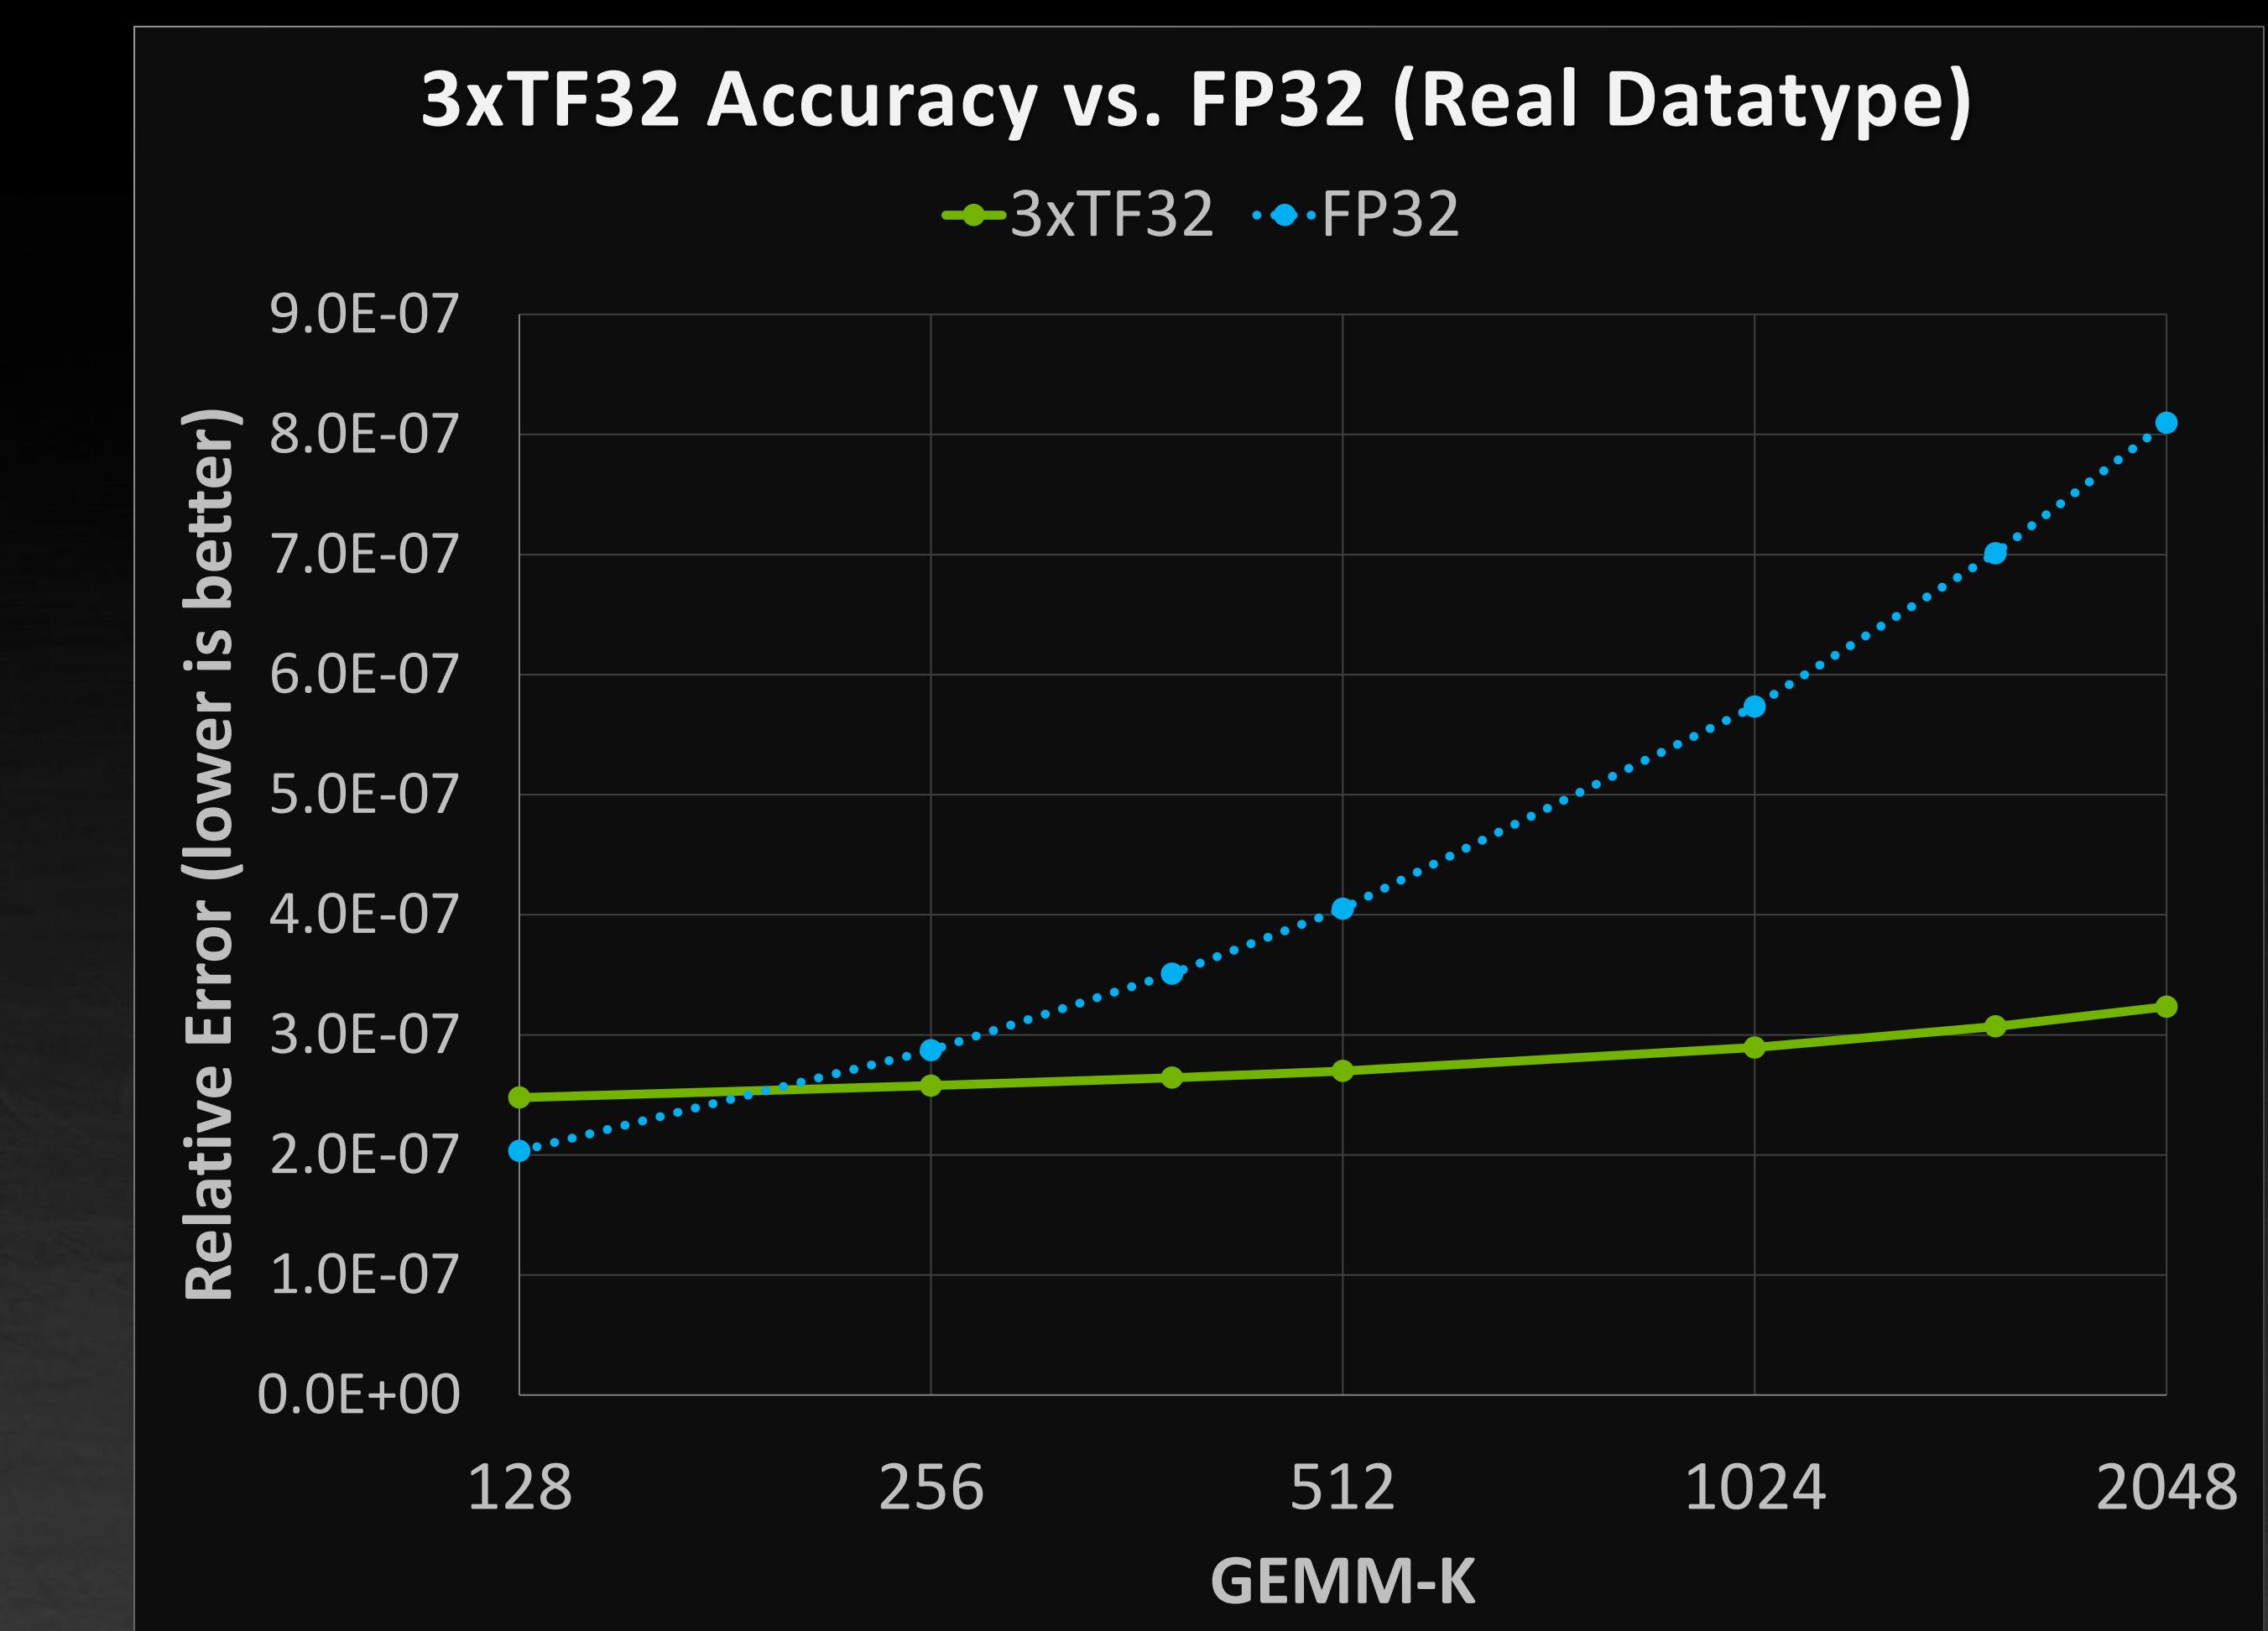
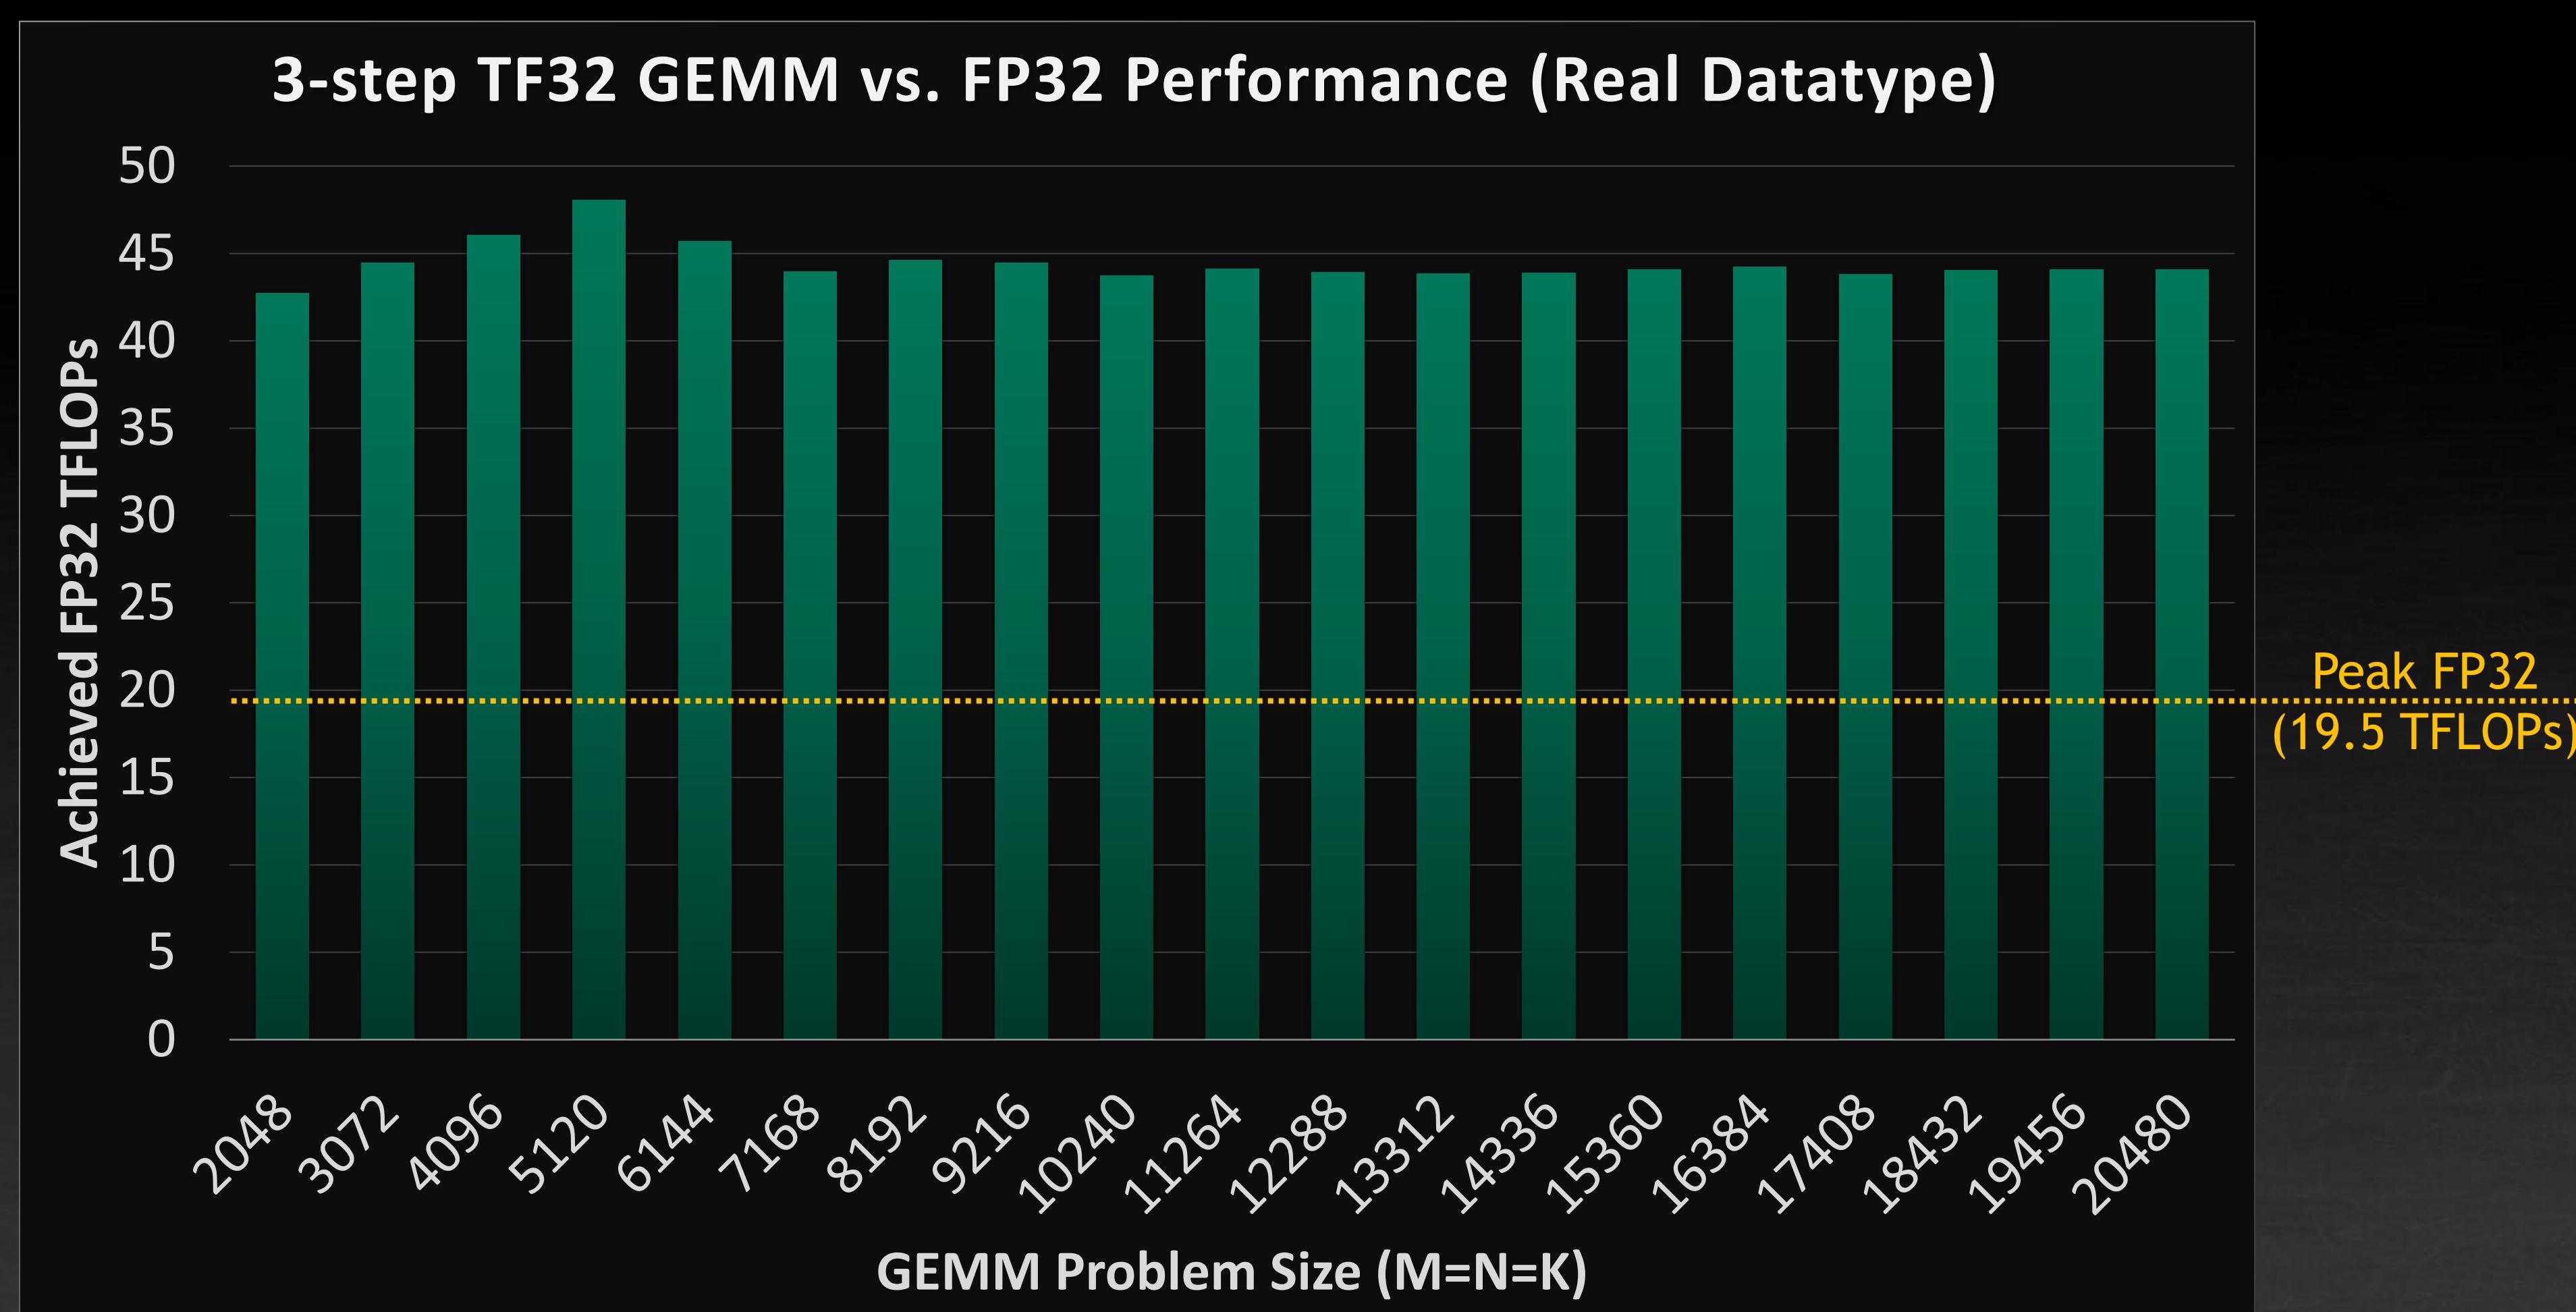
FP32 precision at 48 TFLOPs on A100, using 3-step TF32 sequence

3-step sequence compensates for round-off error, enabling use of TF32 tensor cores

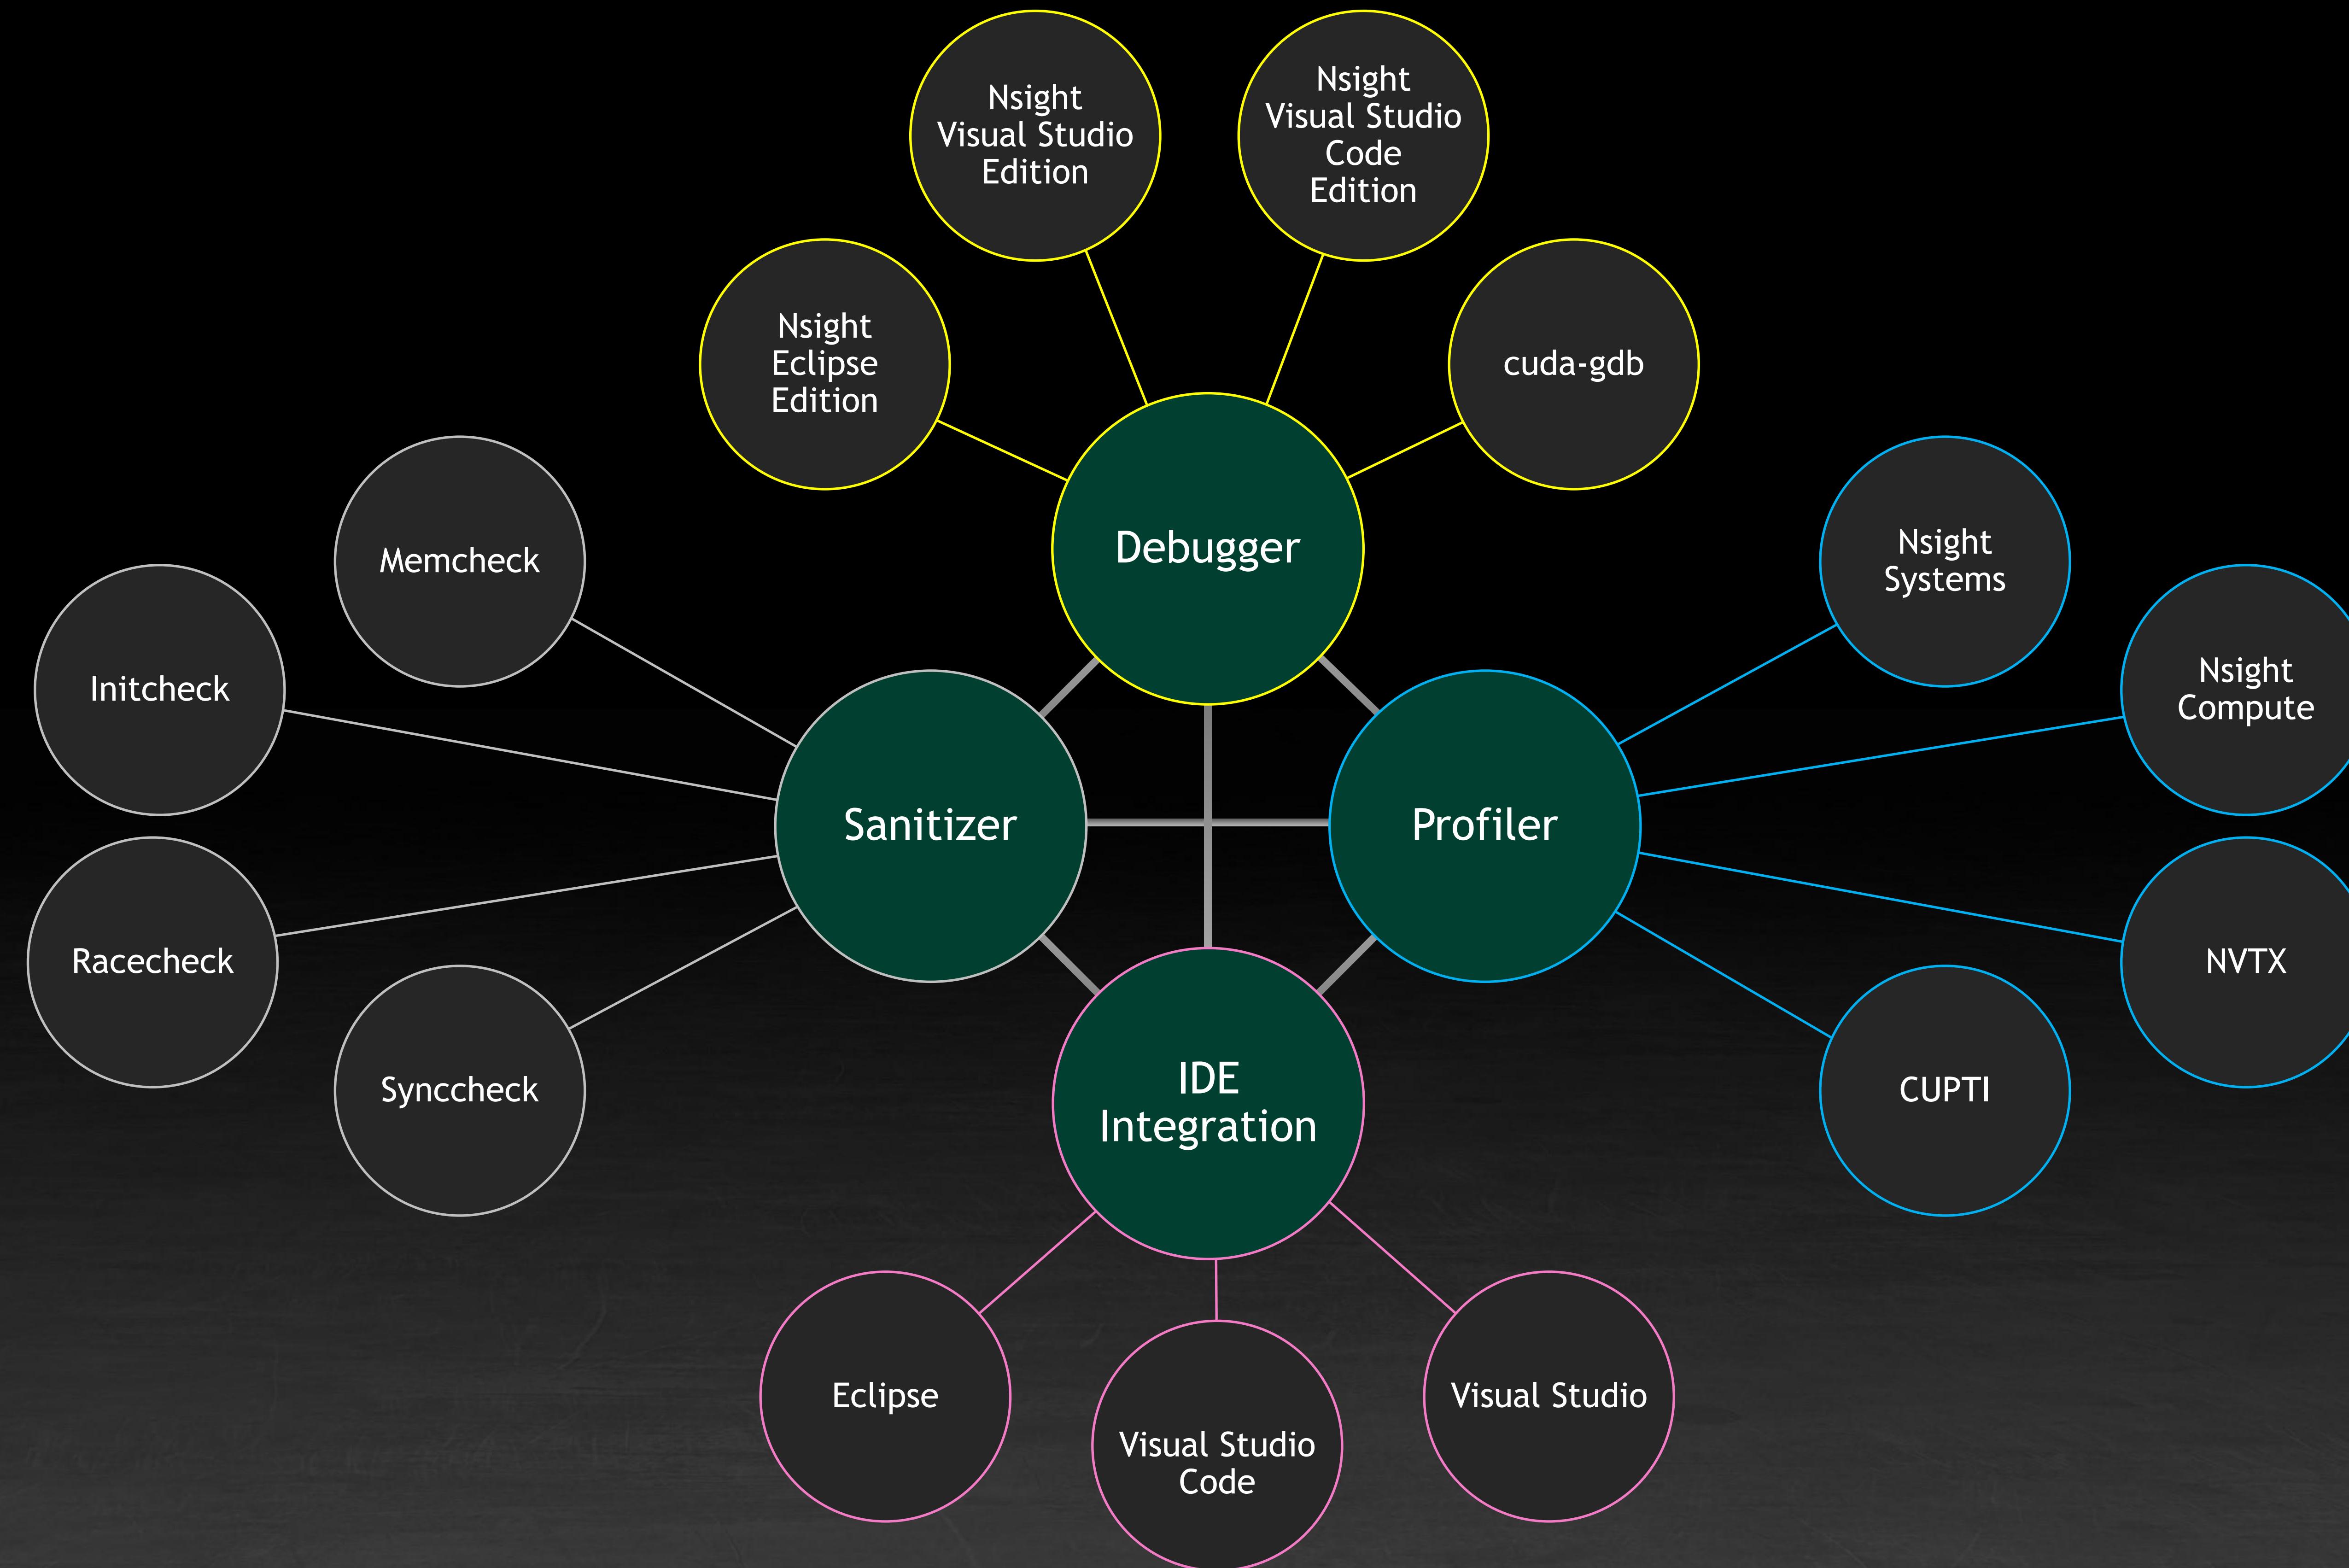
More than 2x performance vs. peak single-precision FFMA

Better accuracy vs. single-precision (not IEEE compliant)

Example implementation for GEMMs and Convolutions

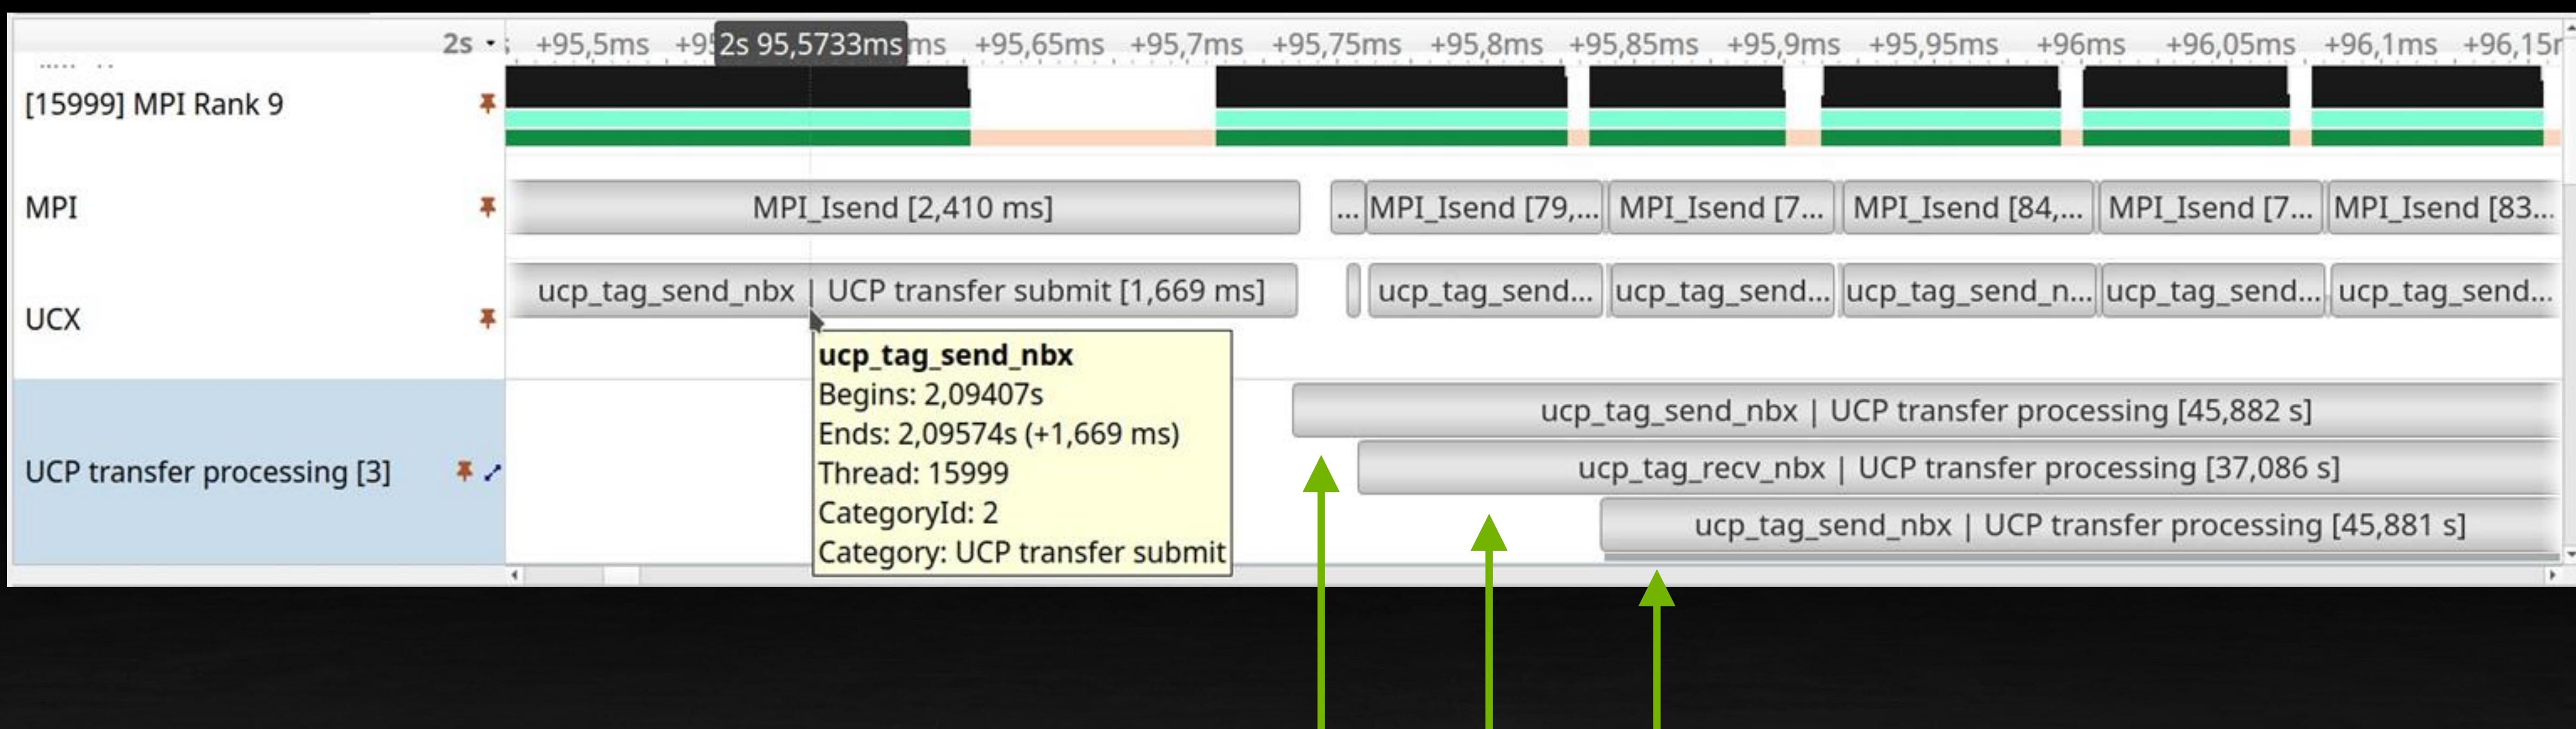


FAMILY OF CUDA DEVELOPER TOOLS



NEW NETWORK PROFILING

Intercept and trace calls into the UCX protocol layer



NIC PERFORMANCE METRICS IN NSIGHT SYSTEMS



```
$ nsys profile -nic-metrics=[true|false] ...
```

NSIGHT COMPUTE: NEW CUDA PLATFORM DEVELOPER TOOLS

NVIDIA Nsight Compute

File Connection Debug Profile Tools Window Help

Connect Disconnect Terminate Profile Kernel Baselines

test_7.ncu-rep

Page: Details Launch Time Cycles Regs GPU SM Frequency CC Process

Current 141 - mergeSortSharedKernel (4096, 1, 1)x(512, 1, 1) 7.01 msecnd 9,283,564 55 0 - NVIDIA GeForce RTX 2080 Ti 1.32 cycle/nsecnd 7.5 [5750] mergeSort

Copy as Image

Occupancy

The occupancy is the ratio of the number of active warps per multiprocessor to the maximum number of possible active warps. Another way to view occupancy is the percentage of the hardware's ability to process warps that is actively in use. Higher occupancy does not always result in higher performance, however low occupancy always reduces the ability to hide latencies, resulting in overall performance degradation. Large discrepancies between the theoretical and the achieved occupancy during execution typically indicates highly imbalanced workloads.

Theoretical Occupancy [%] 100 Block Limit Registers [block]
Theoretical Active Warps per SM [warp] 32 Block Limit Shared Mem [block]
Achieved Occupancy [%] 99.40 Block Limit Warps [block]
Achieved Active Warps Per SM [warp] 31.81 Block Limit SM [block]

Occupancy Limiters This kernel's theoretical occupancy is not impacted by any block limit.

Impact of Varying Register Count Per Thread

Impact of Varying Block Size

Occupancy Calculator

Source: tea_leaf_ppcg.cuhnl

Launch: 0 - 32655 - device_tea_leaf_ppcg_solve_update_r

Time: 1.07 msecnd

Cycles: 1,458,003

Regs: 32

GPU: NVIDIA GeForce RTX 2080 Ti

SM Frequency: 1.36 cycle/nsecnd

CC: 7.5 [10906] tea_leaf

View: Source and SASS

Navigation: Instructions Executed

Source: device_tea_leaf_ppcg_solve_update_r

Navigation: Instructions Executed

Inst E

Address Source

```

145 sd[THARR2D(0, 0, 0)] = alpha[step]*sd[THARR2D(0, 0, 0)]
146           + beta[step]*r[THARR2D(0, 0, 0)];
147     }
148   }
149   */
150   /* New update to rtemp for use in calc_sd */
151   __global__ void device_tea_leaf_ppcg_solve_update_r(
152     kernel_info_t kernel_info,
153     double * __restrict const rtemp,
154     const double * __restrict const Kx,
155     const double * __restrict const Ky,
156     const double * __restrict const sd)
157   {
158     __kernel_indexes;
159
160     if (WITHIN_BOUNDS)
161     {
162       const double result = (1.0
163         + (Ky[THARR2D(0, 1, 0)] + Ky[THARR2D(0, 0, 1)])
164         + (Kx[THARR2D(1, 0, 0)] + Kx[THARR2D(0, 0, 1)])*sd[THARR2D(0, 0, 0)]
165         - (Ky[THARR2D(0, 1, 0)]*sd[THARR2D(0, 1, 0)] + Ky[THARR2D(0, 0, 1)]*sd[THARR2D(0, -1, 0)])
166         - (Kx[THARR2D(1, 0, 0)]*sd[THARR2D(1, 0, 0)] + Kx[THARR2D(0, 0, 1)]*sd[THARR2D(-1, 0, 0)]);
167
168     }
169   }
170 
```

Standalone Source Viewer

Register Dependency Visualization

How To Understand and Optimize Shared Memory Accesses using Nsight Compute [S41723]



NVTX v3 OVERVIEW

```
#include <nvtx3/nvToolsExt.h>

void LaunchKernel() {
    nvtxRangePush(__FUNCTION__);
    Kernel<<<2000, 256>>>();
    nvtxRangePop();
}

void Example() {
    nvtxRangePush("Creating CUDA context");
    cudaFree(0);
    nvtxRangePop();

    nvtxRangePush("Launching & waiting for kernels");
    for (int i = 0; i < 3; ++i) {
        LaunchKernel();
    }
    cudaDeviceSynchronize();
    nvtxRangePop();
}
```

What is it?

Annotation library inserting markers for tools analysis

Negligible overhead

No dependencies - just header files

How does it work?

Markers record a point of time

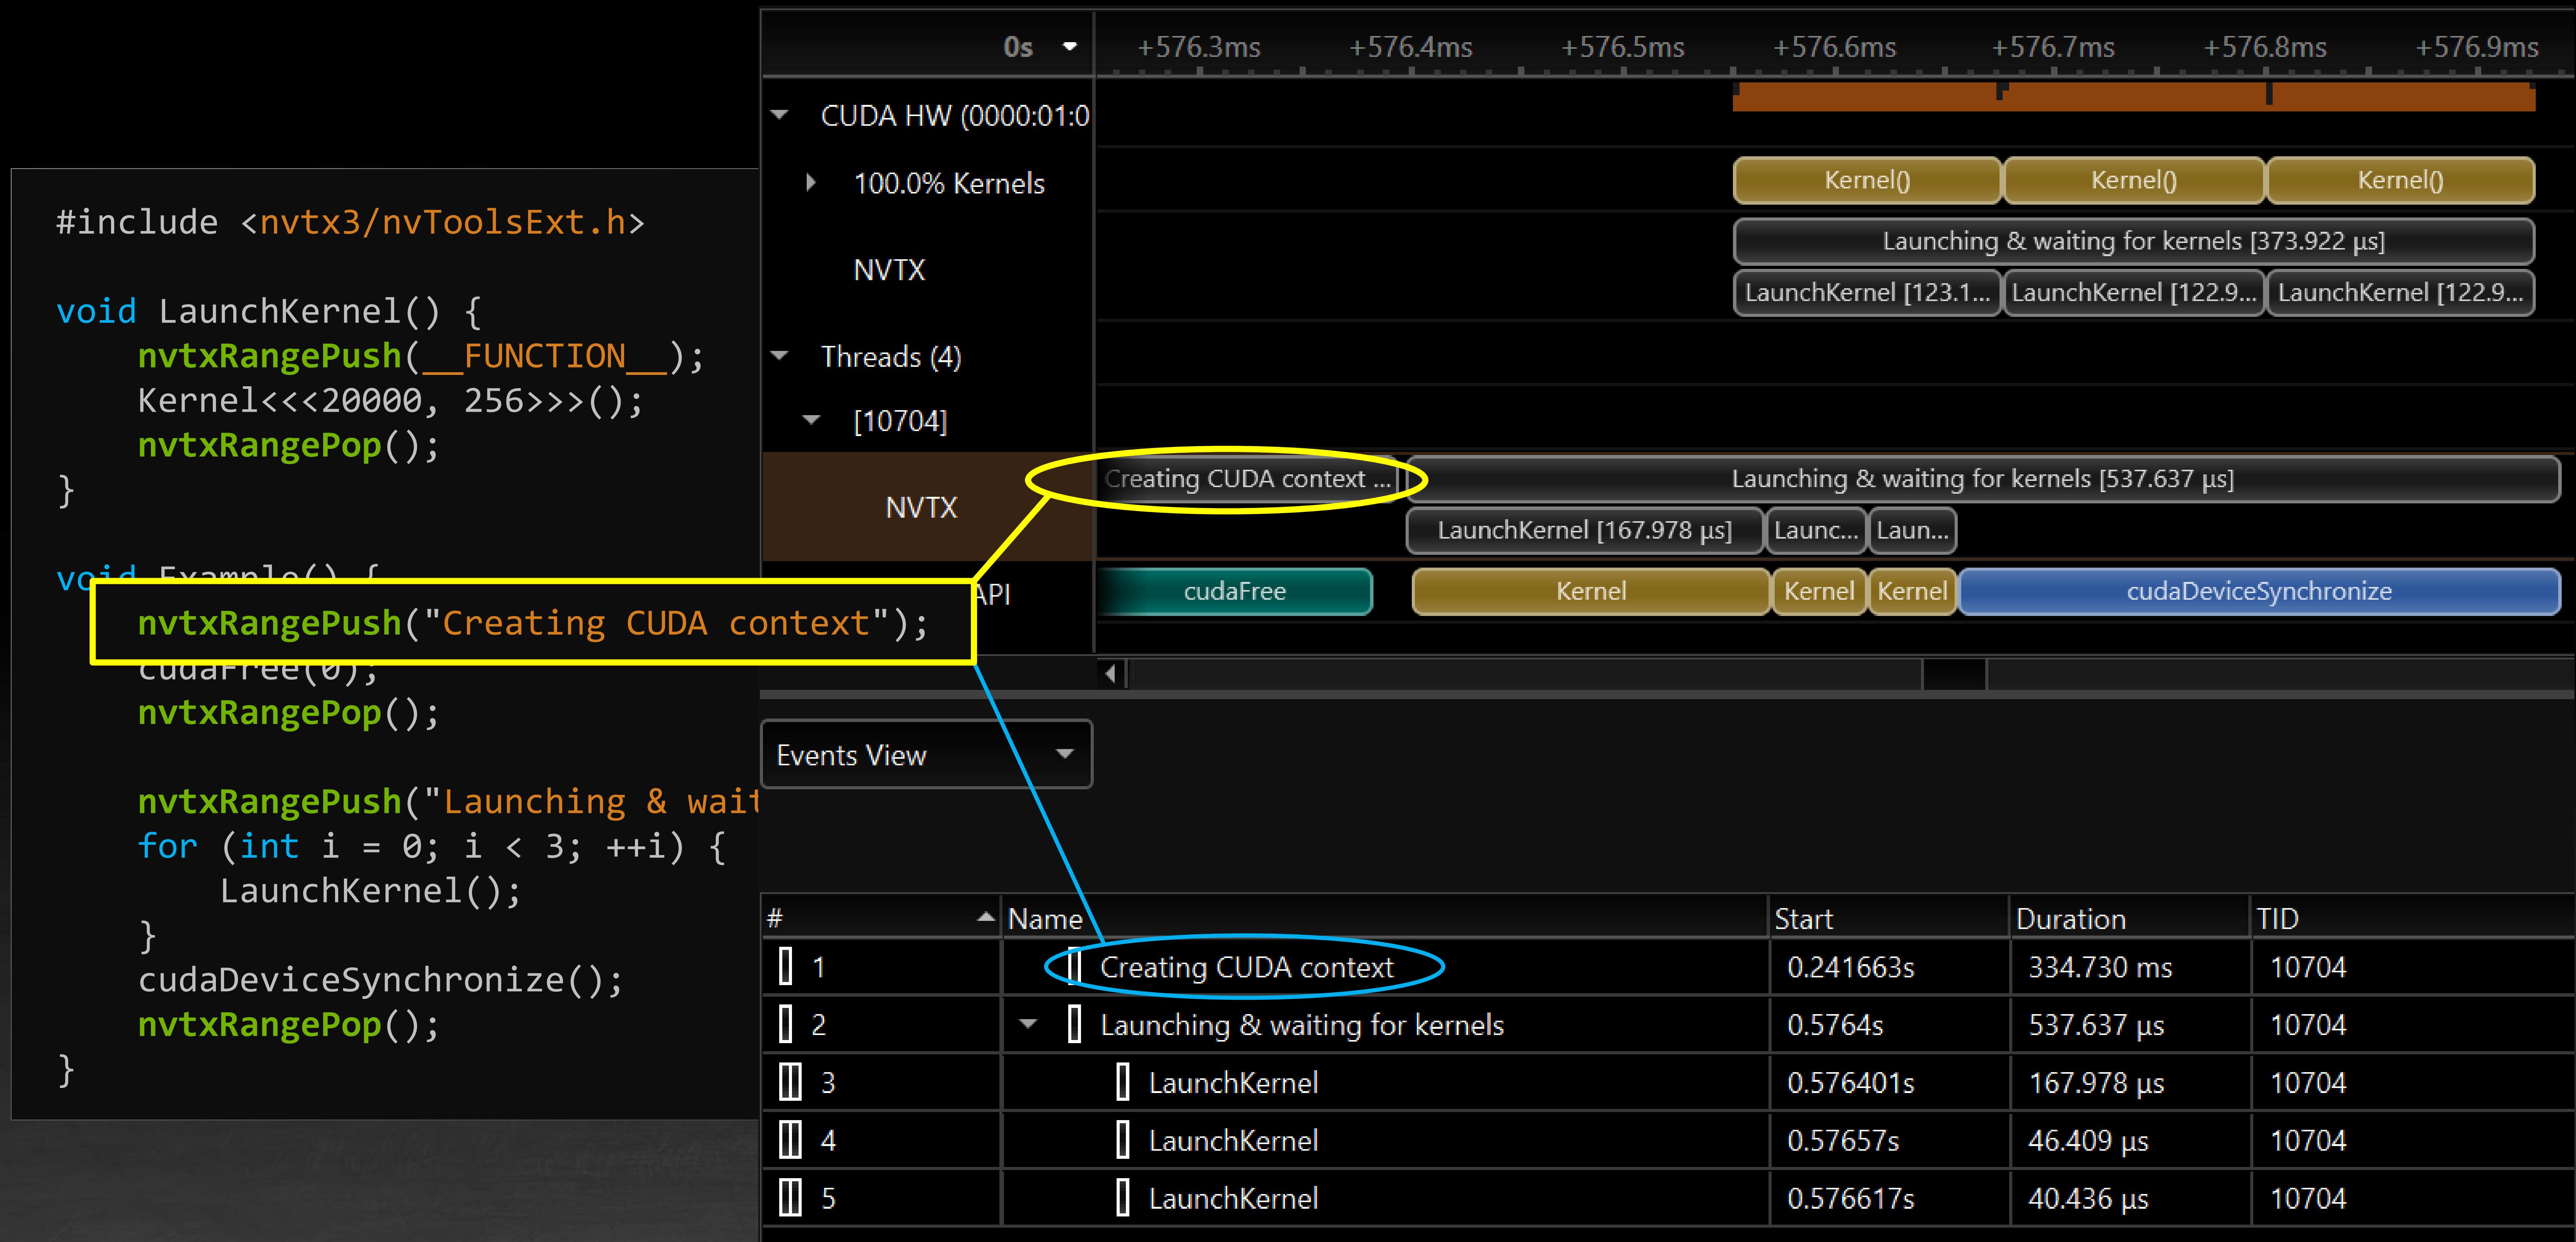
Push/pop ranges to record a span of time

Ranges form a stack to show nested detail

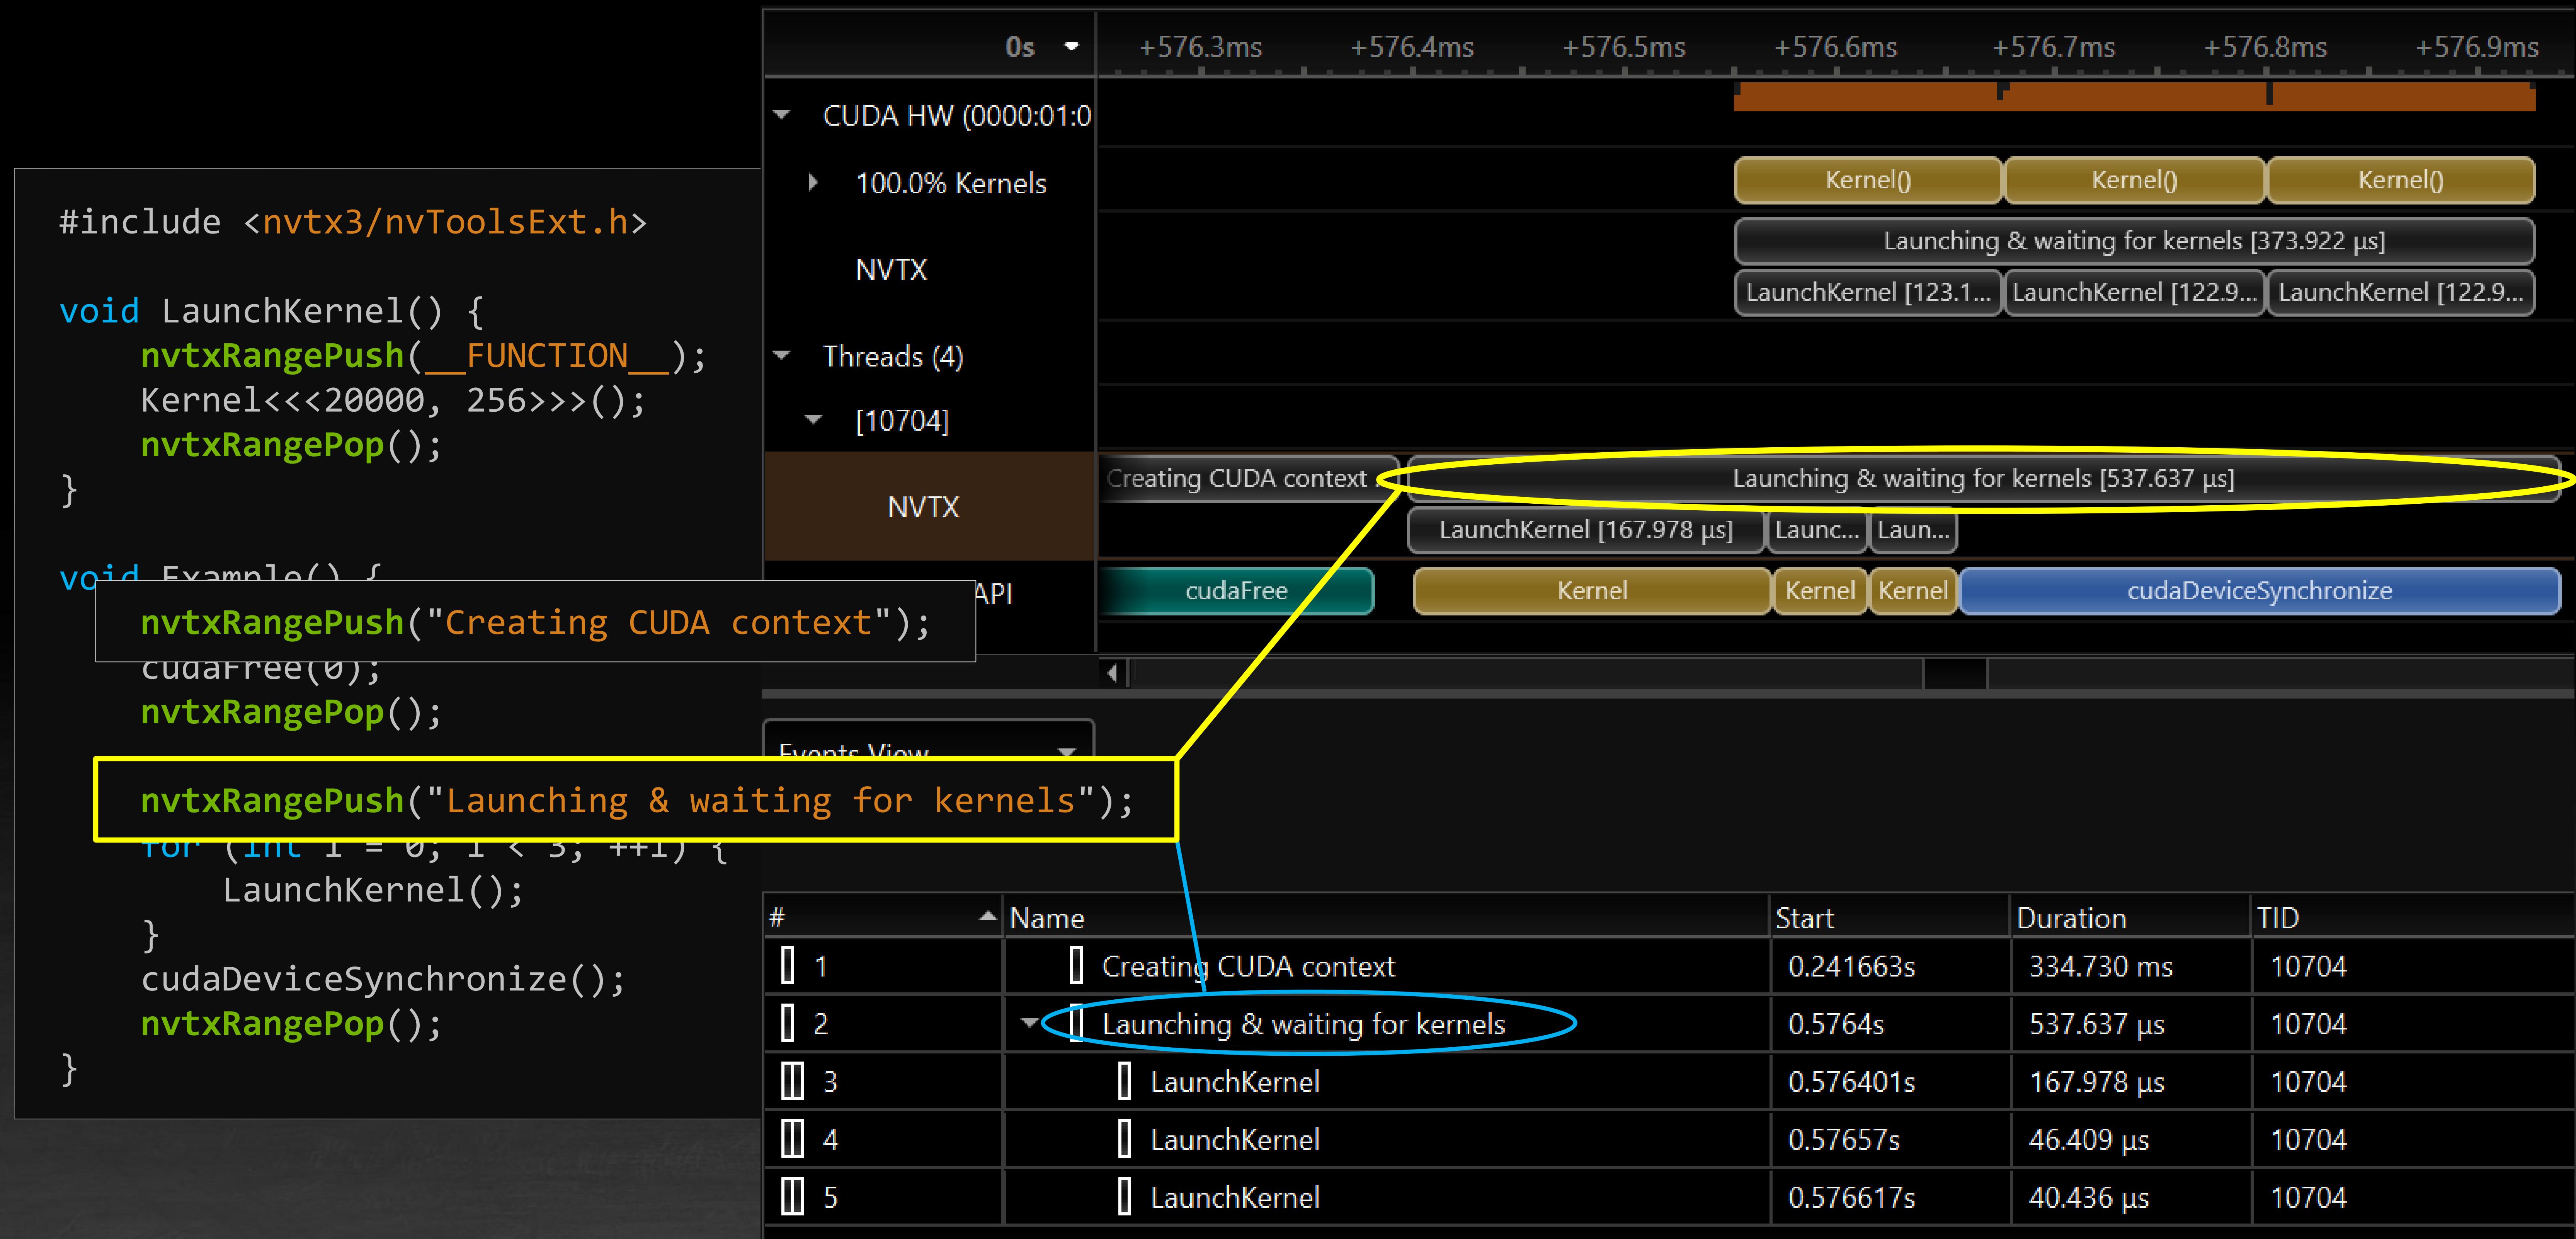
Tools track NVTX calls to display timelines (see next slide)

See link below for details & download

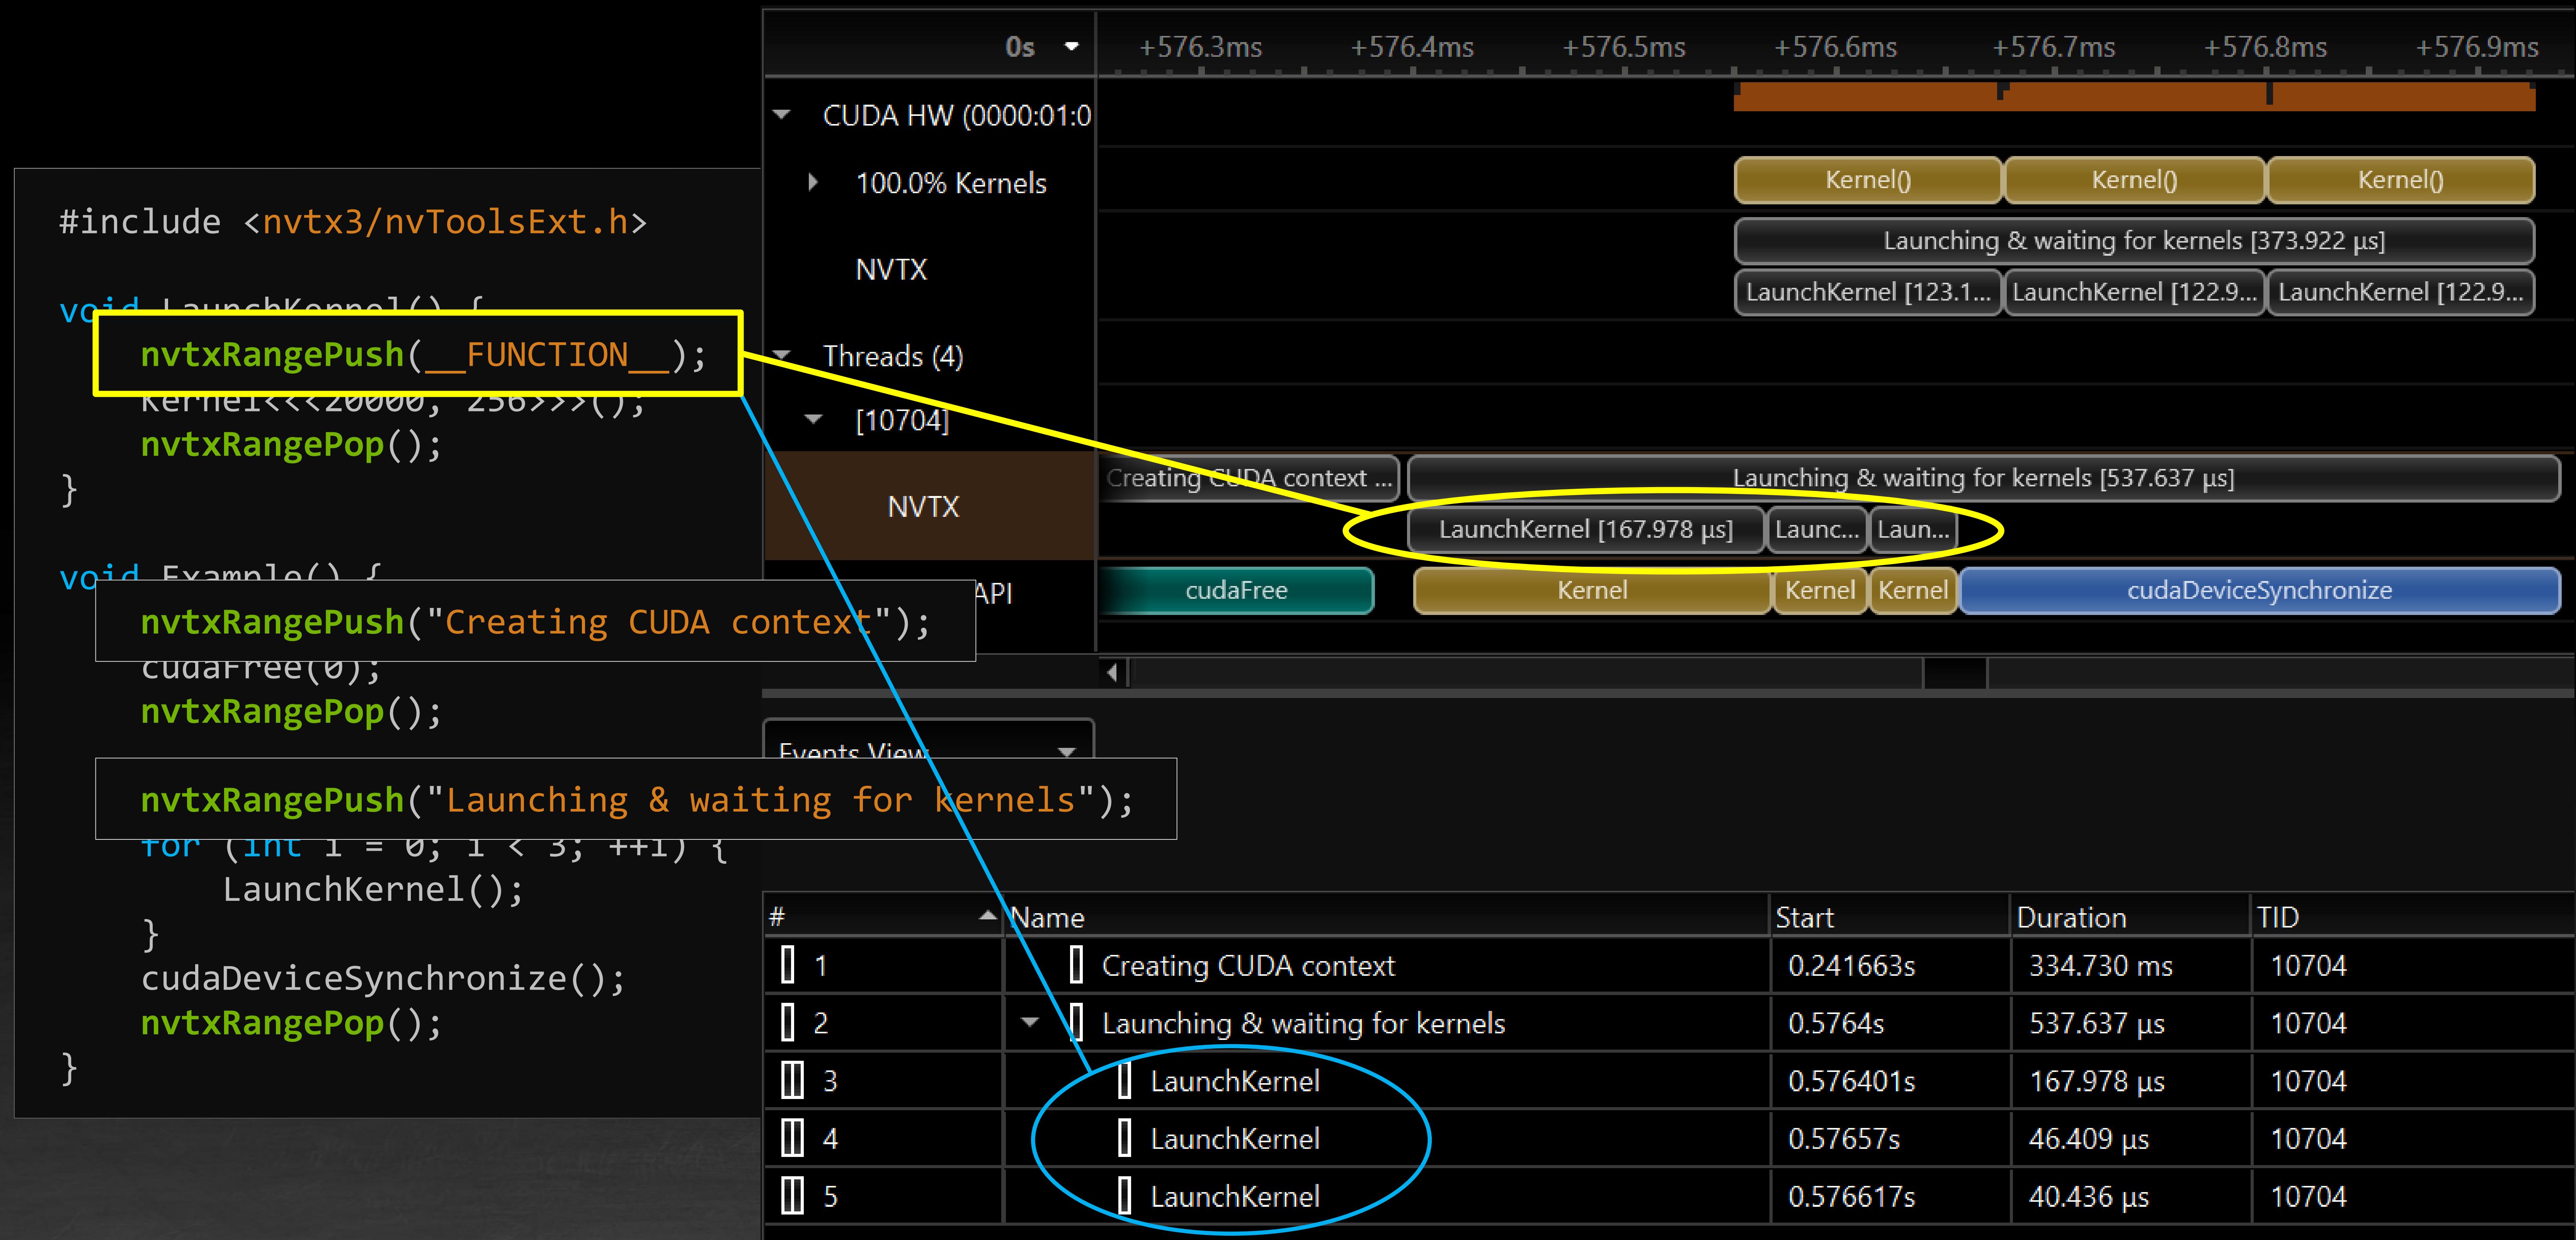
NVTX v3 OVERVIEW



NVTX v3 OVERVIEW



NVTX v3 OVERVIEW



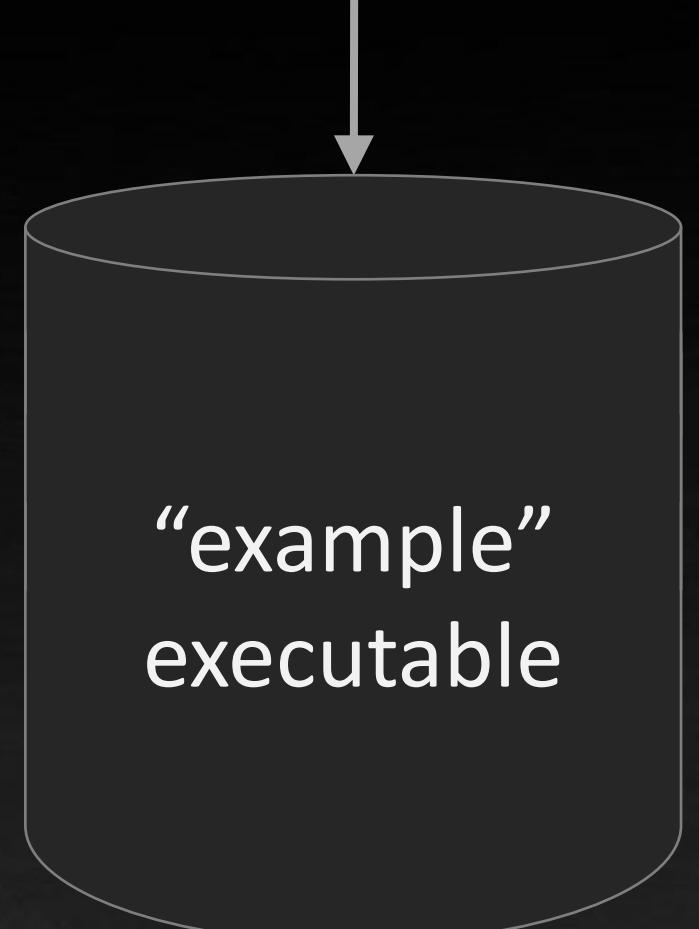
ANATOMY OF A CUDA BINARY

Hello world example

```
__global__ void kernel() {
    printf("Hello, CUDA\n");
}

void main() {
    kernel<<< 1, 1 >>>();
    cudaDeviceSynchronize();
}
```

```
nvcc example.cu -o example
```



“example”
executable

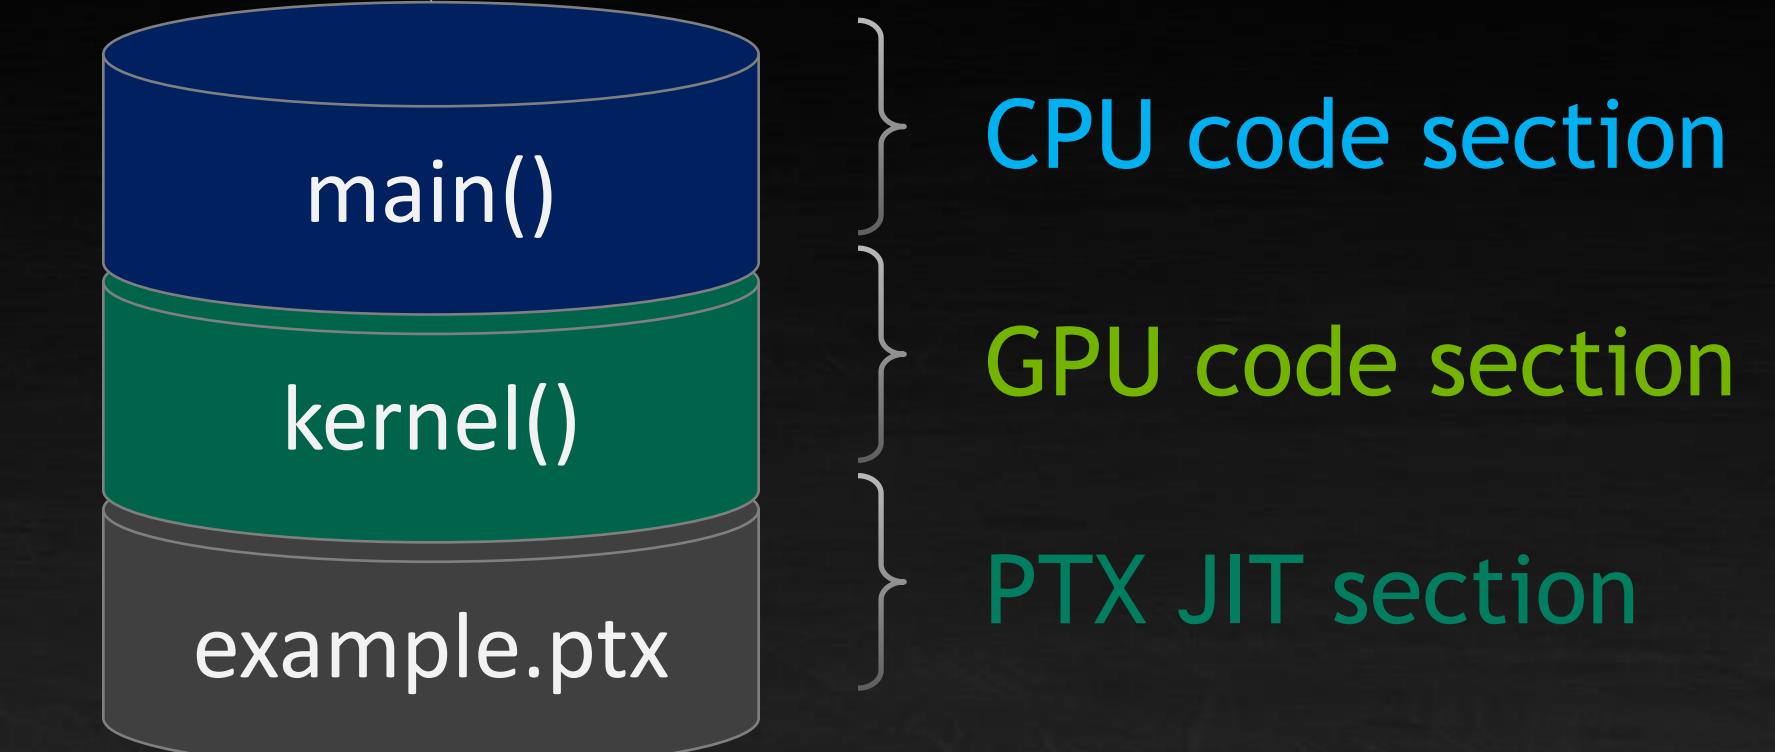
ANATOMY OF A CUDA BINARY

Hello world example

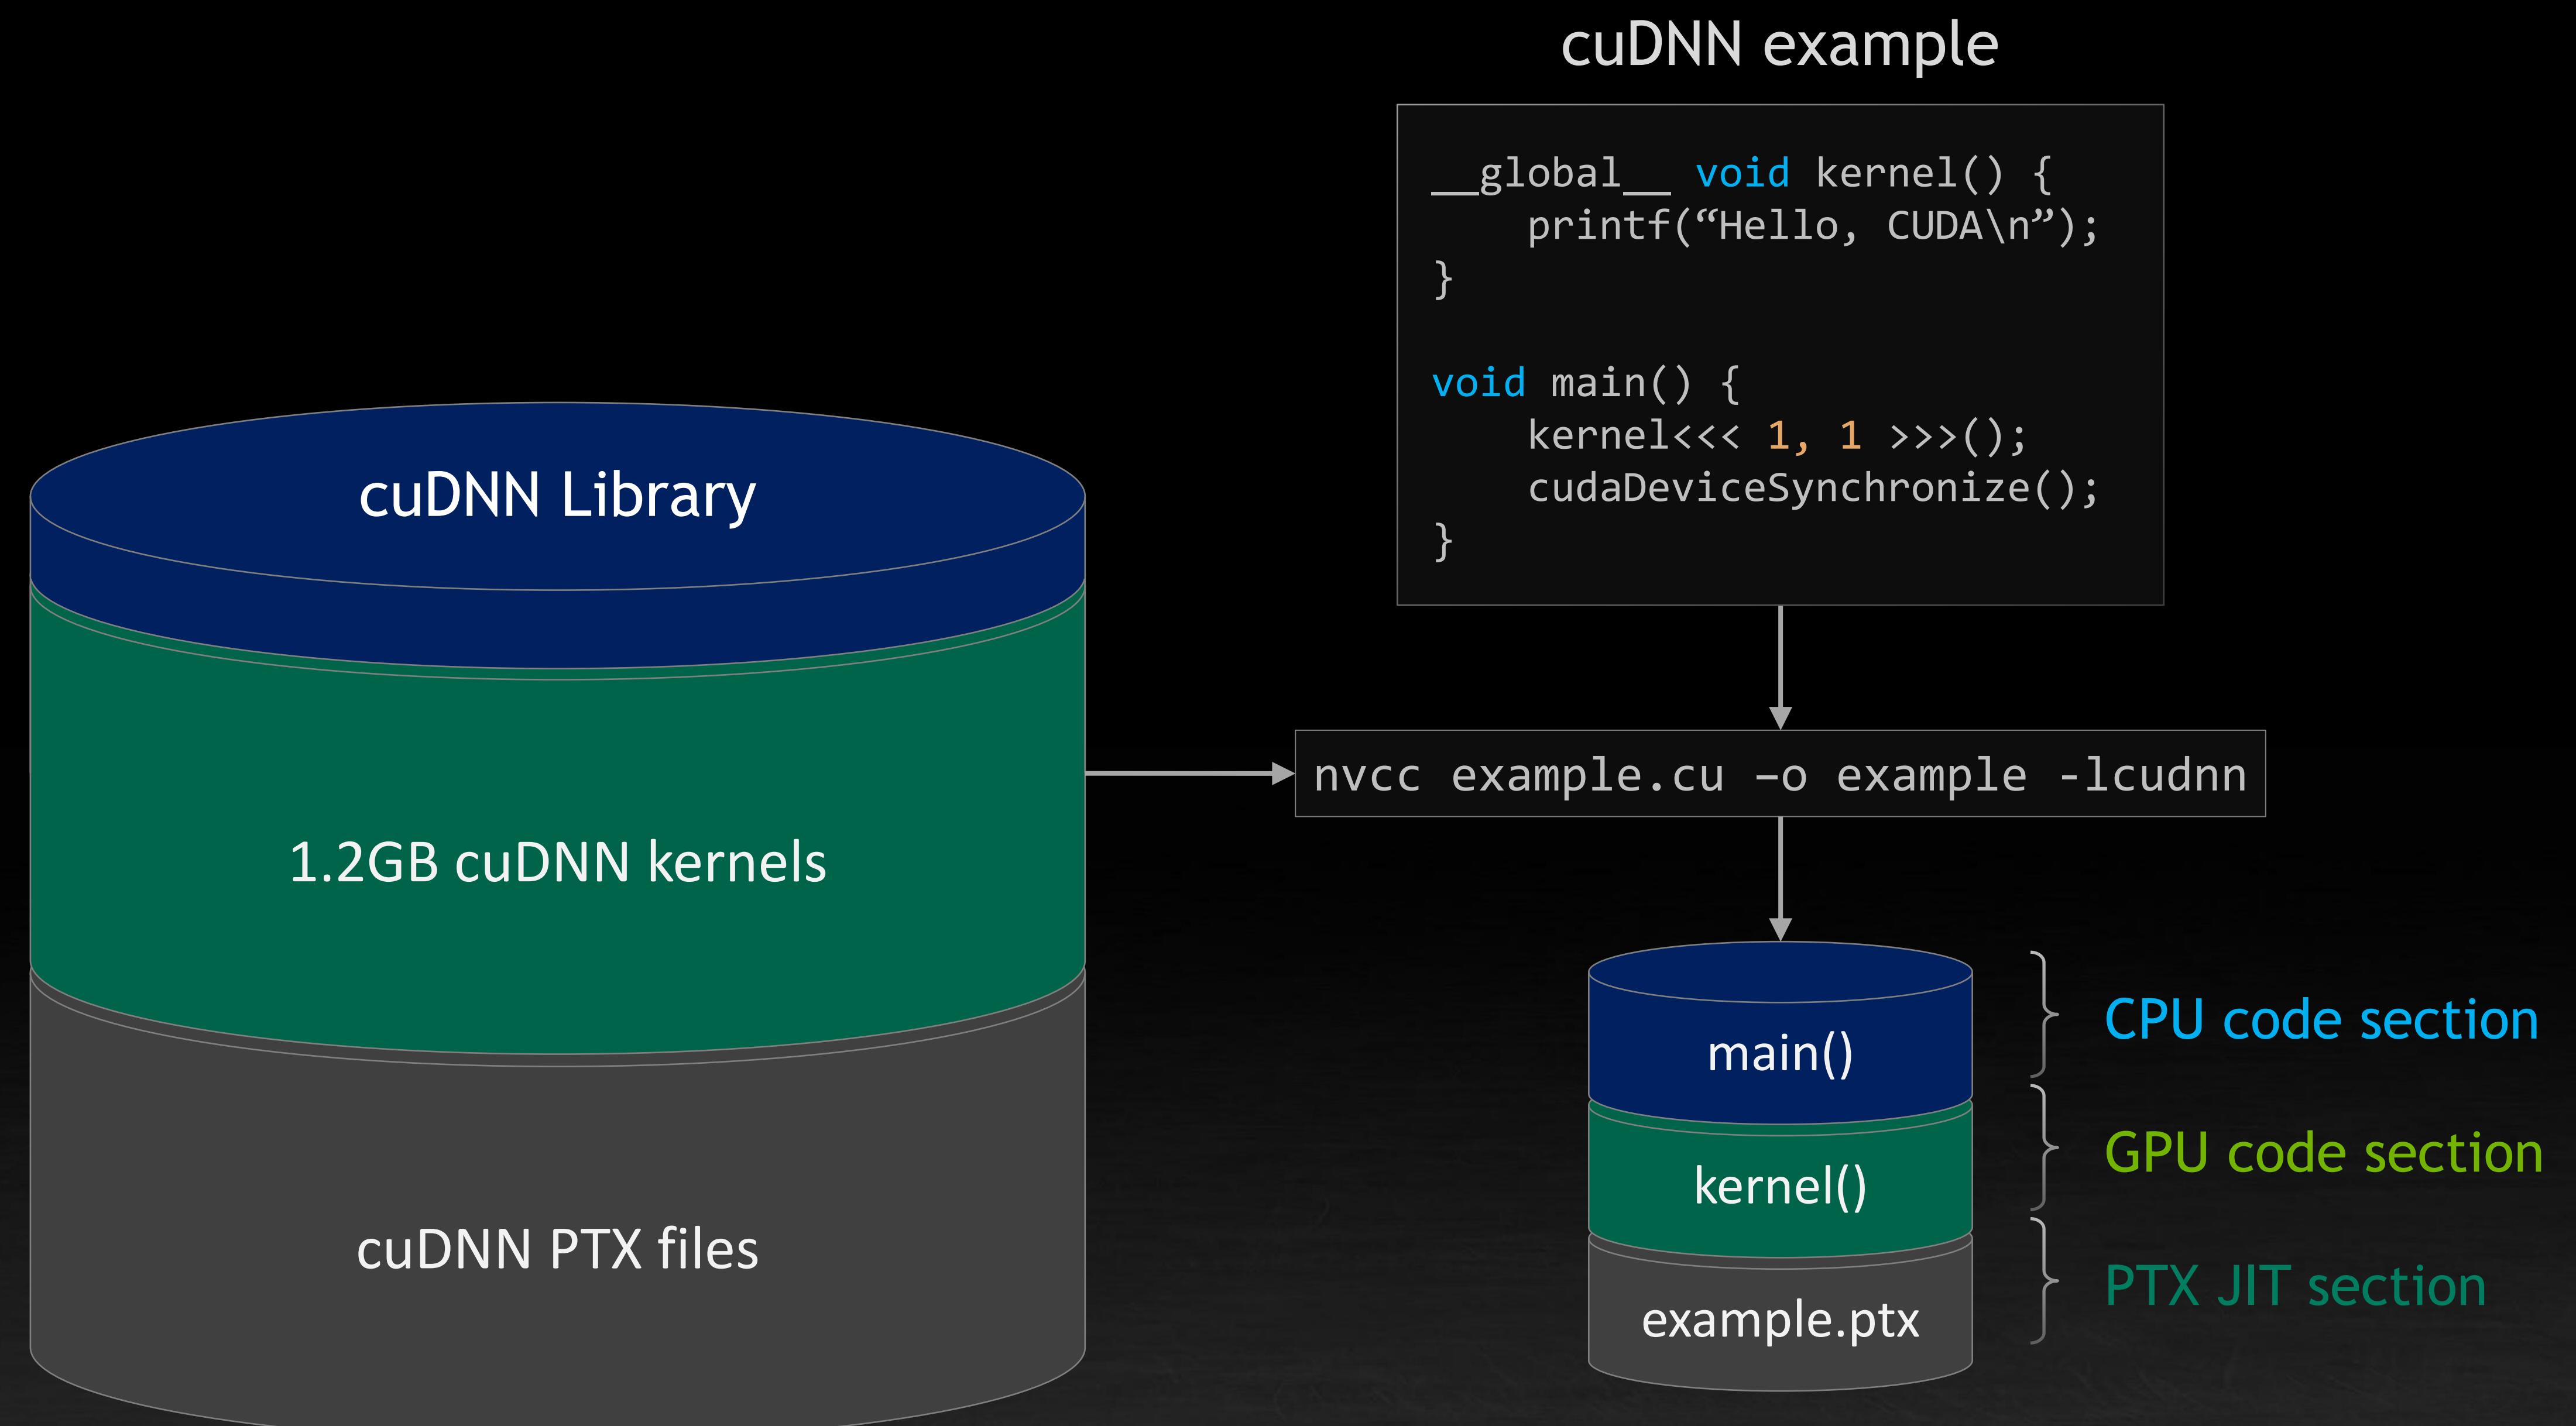
```
__global__ void kernel() {
    printf("Hello, CUDA\n");
}

void main() {
    kernel<<< 1, 1 >>>();
    cudaDeviceSynchronize();
}
```

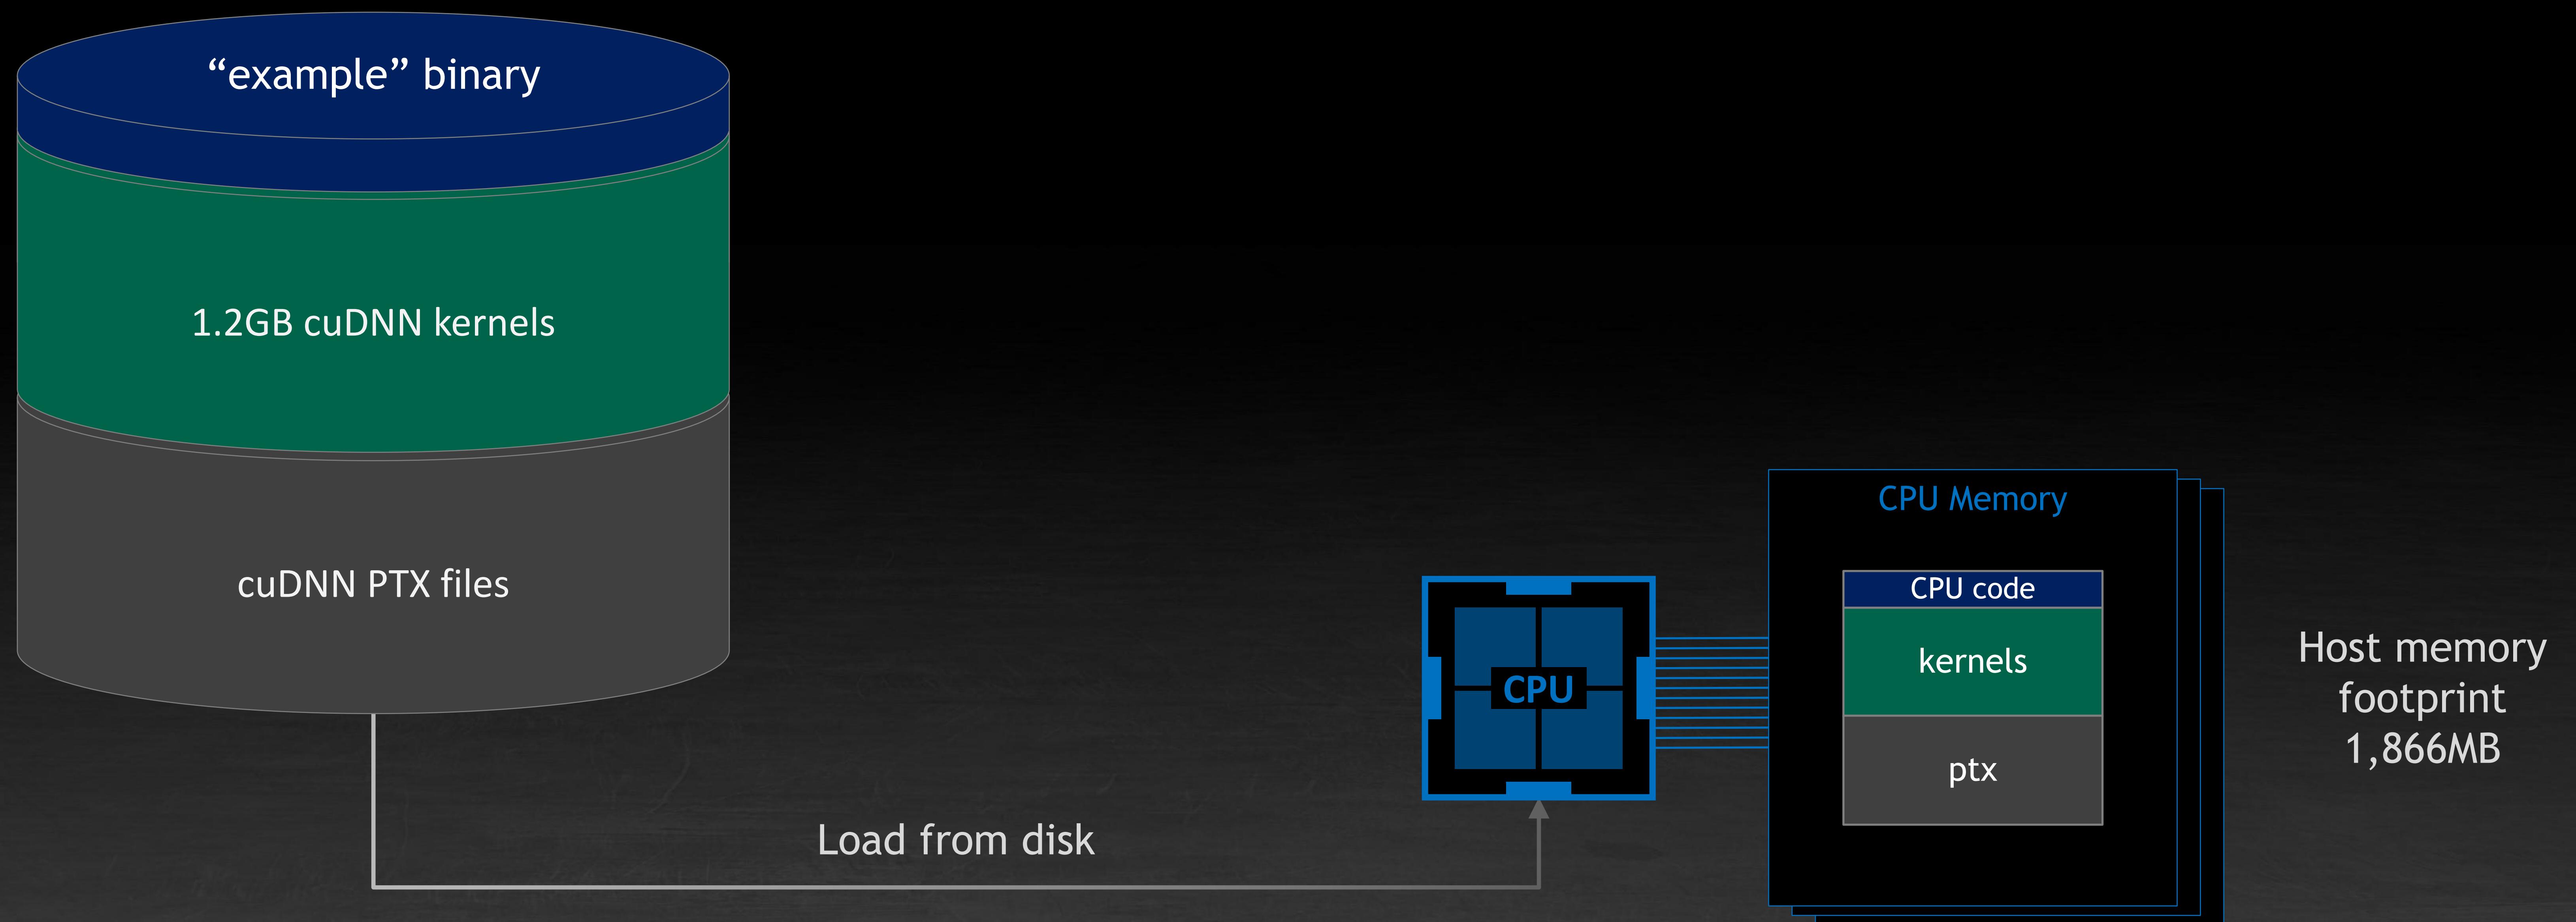
```
nvcc example.cu -o example
```



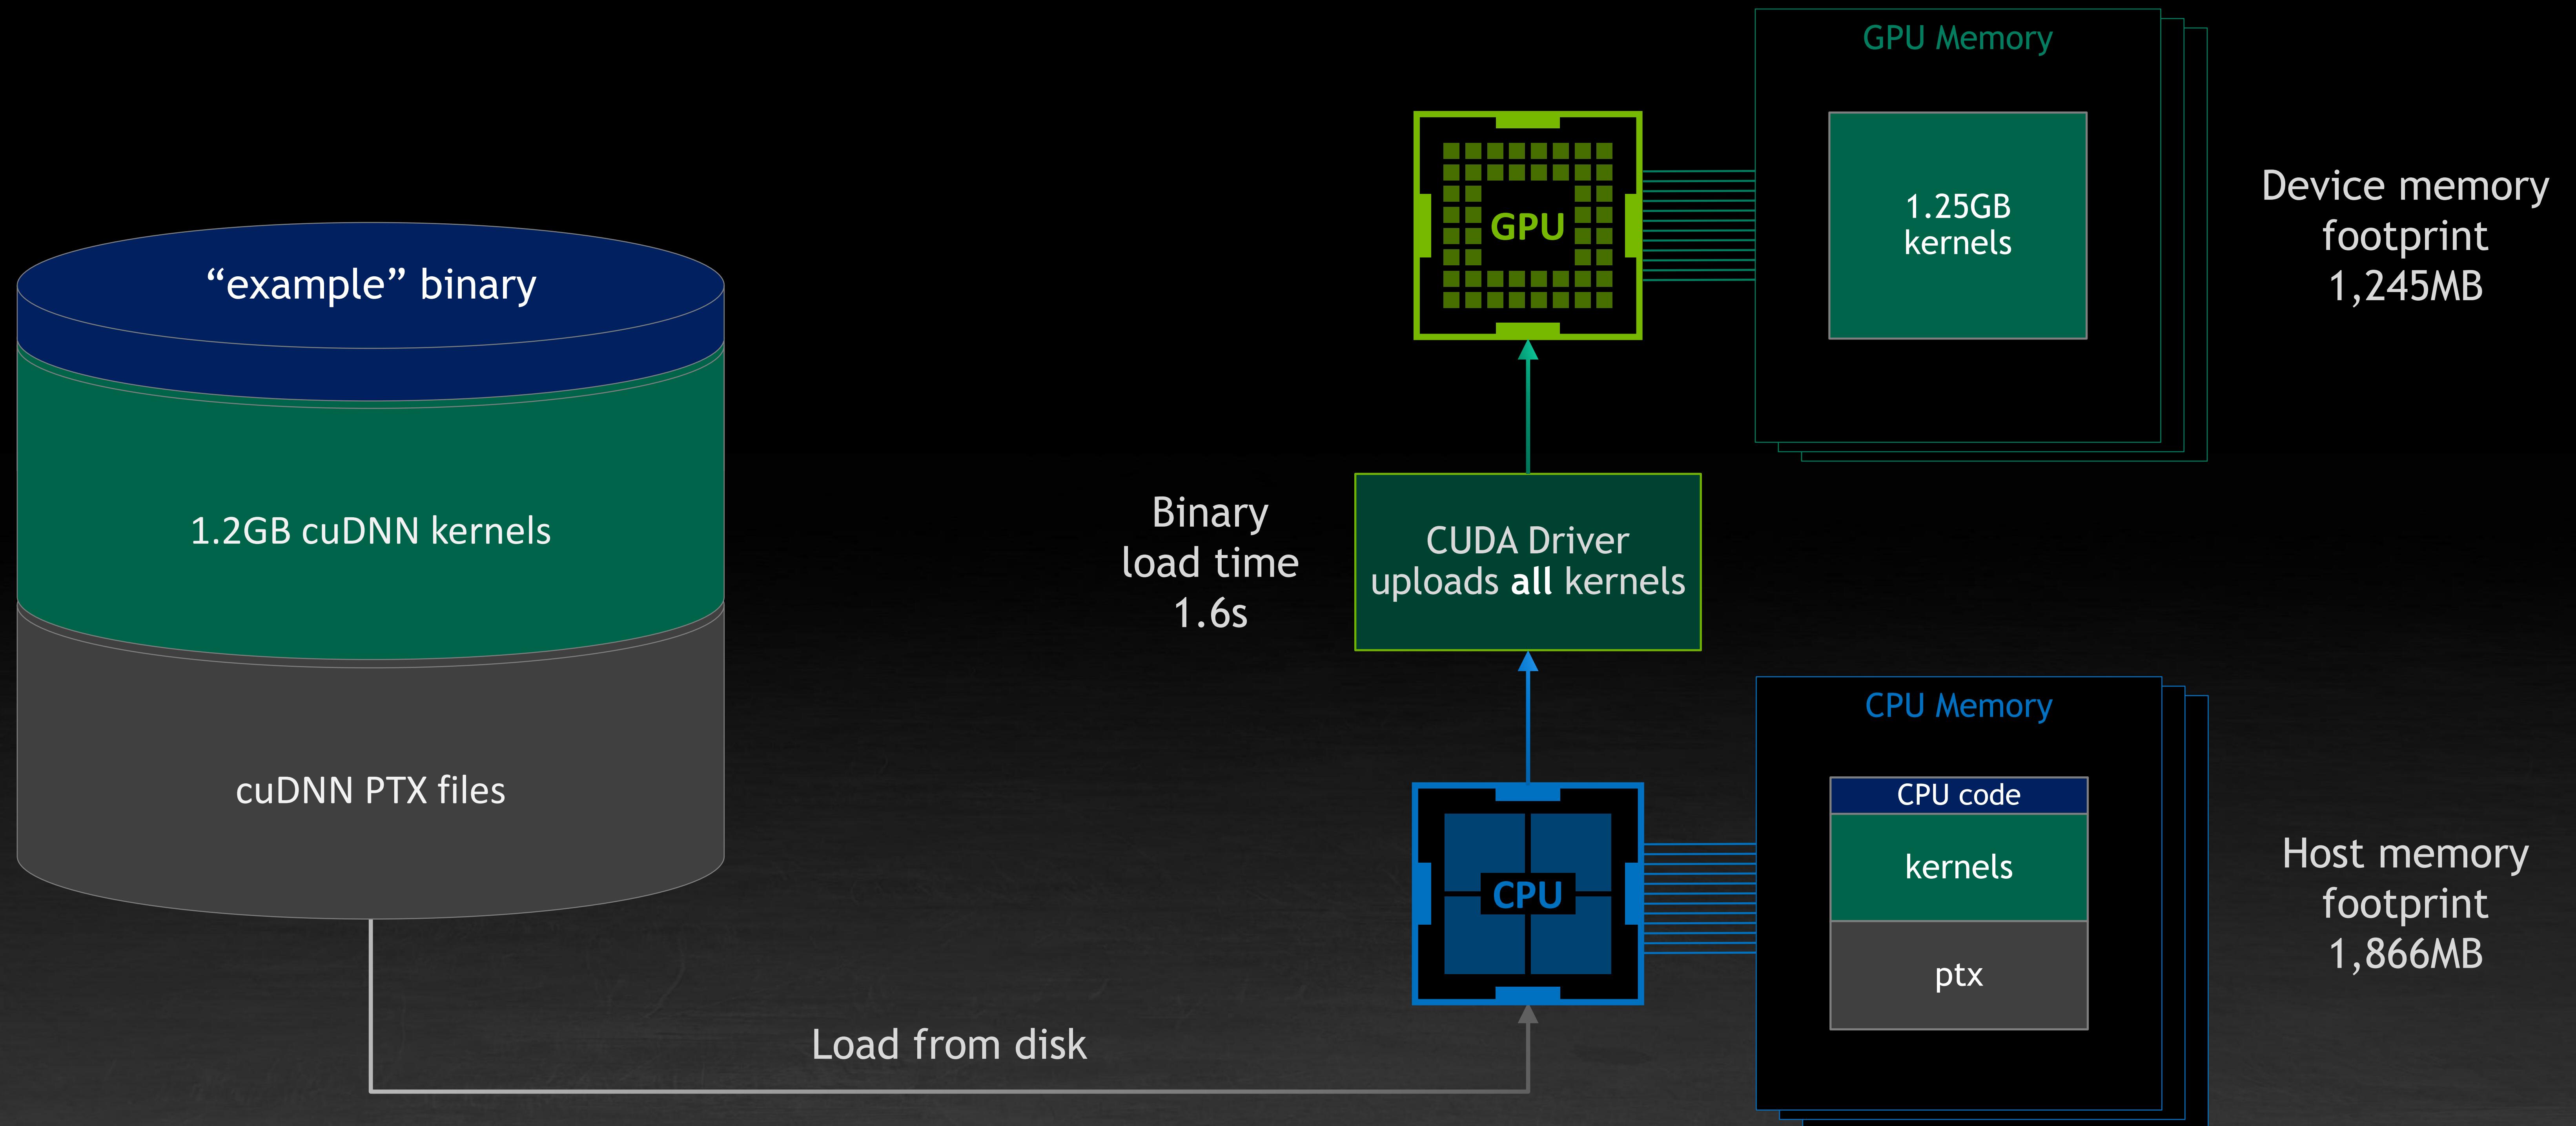
ANATOMY OF A CUDA BINARY



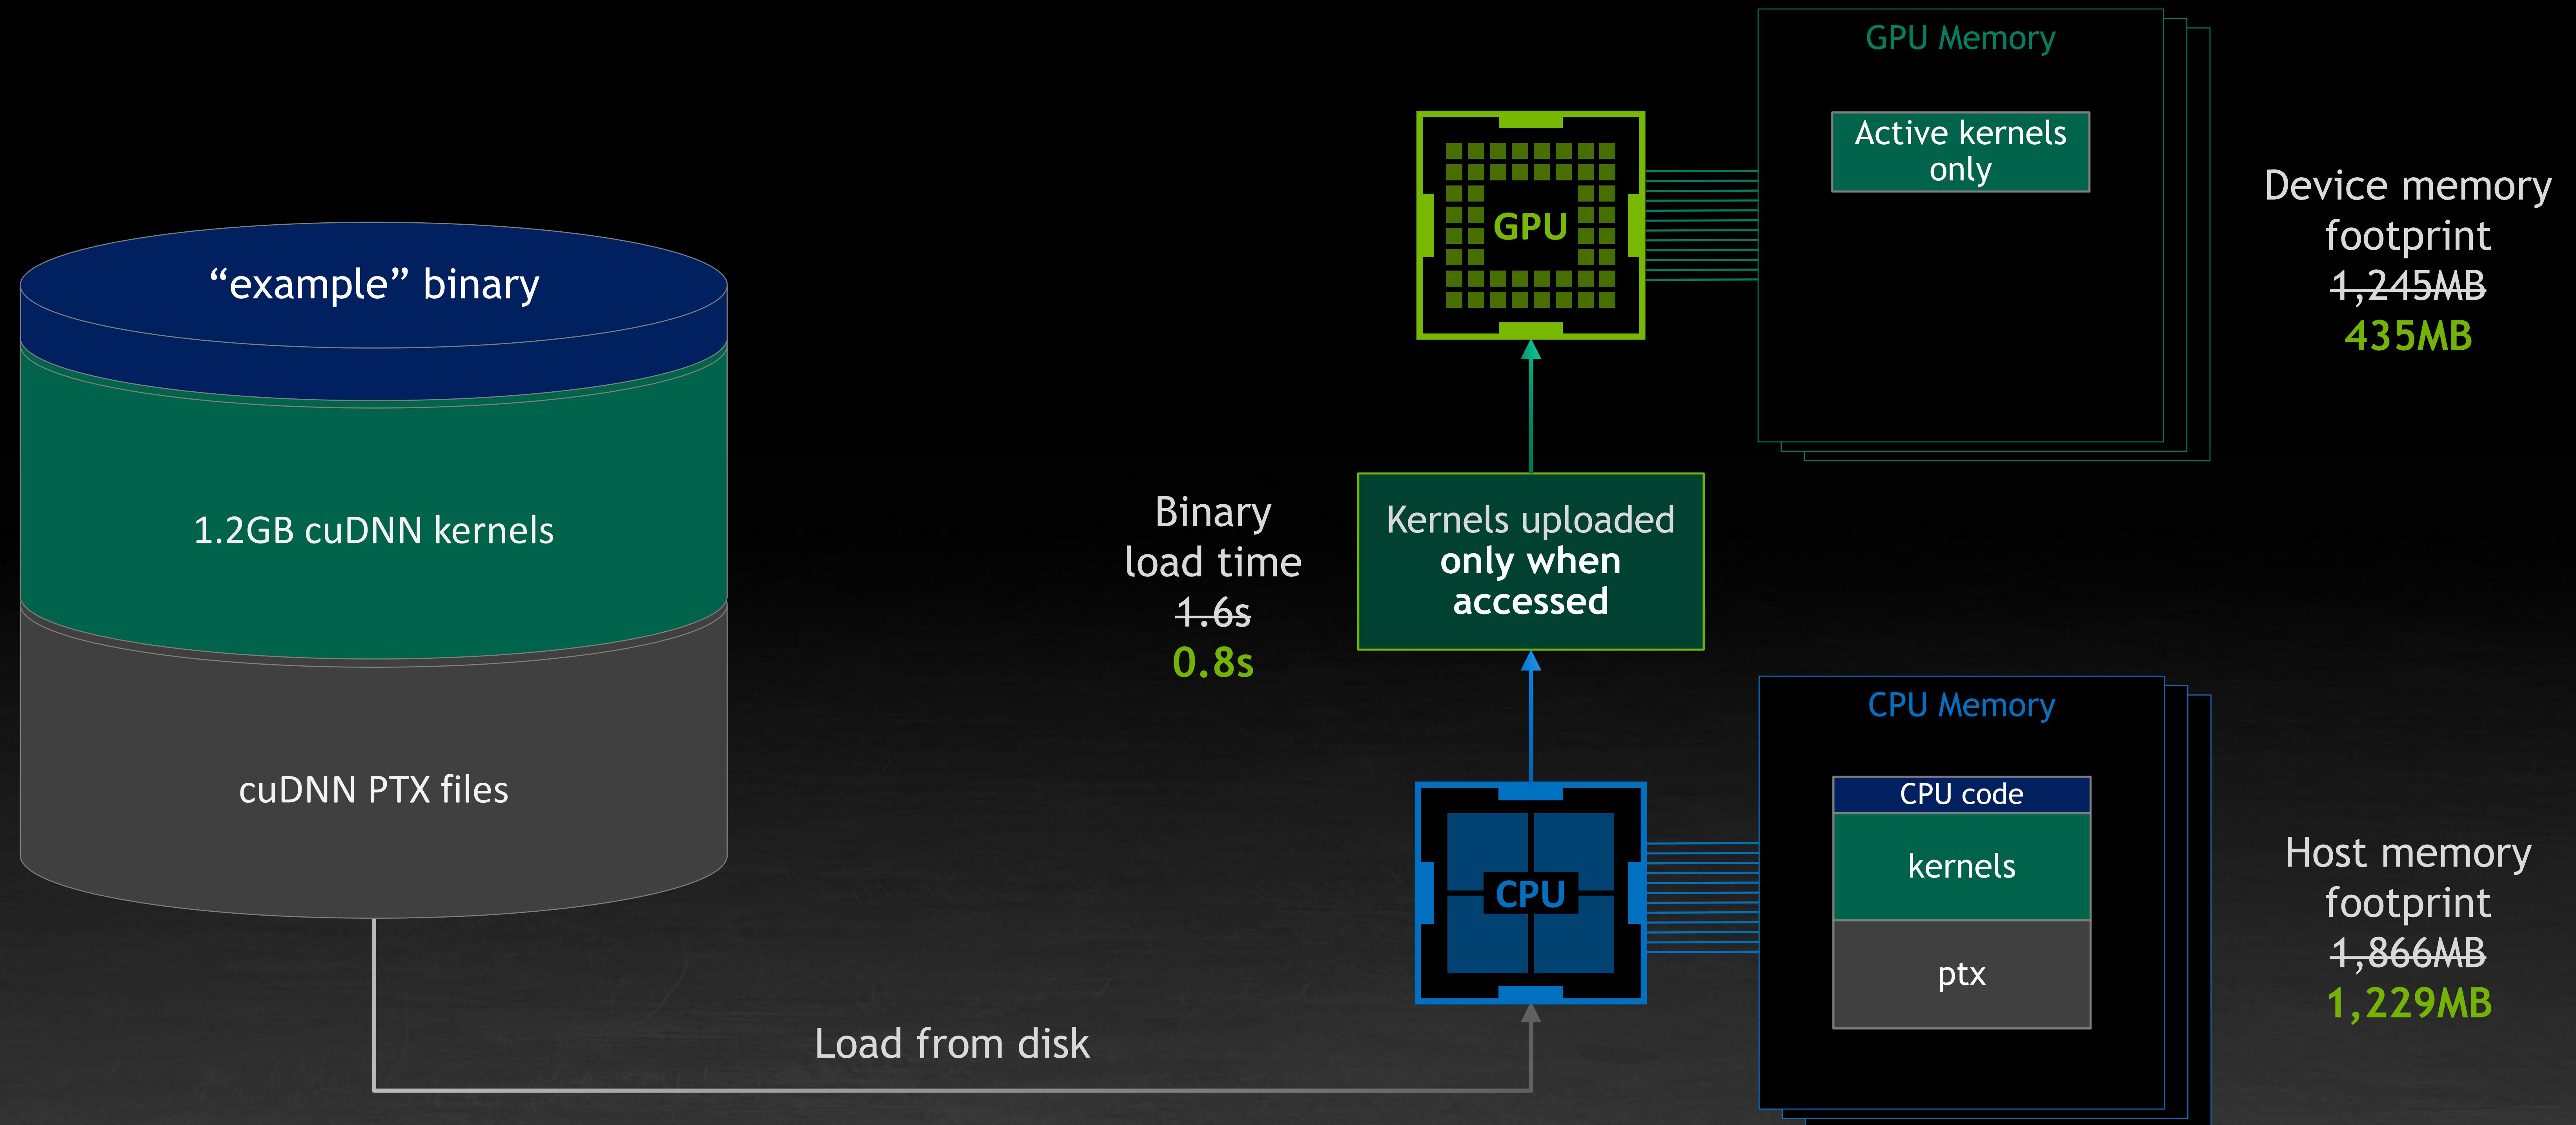
LOADING A CUDA BINARY



LOADING A CUDA BINARY

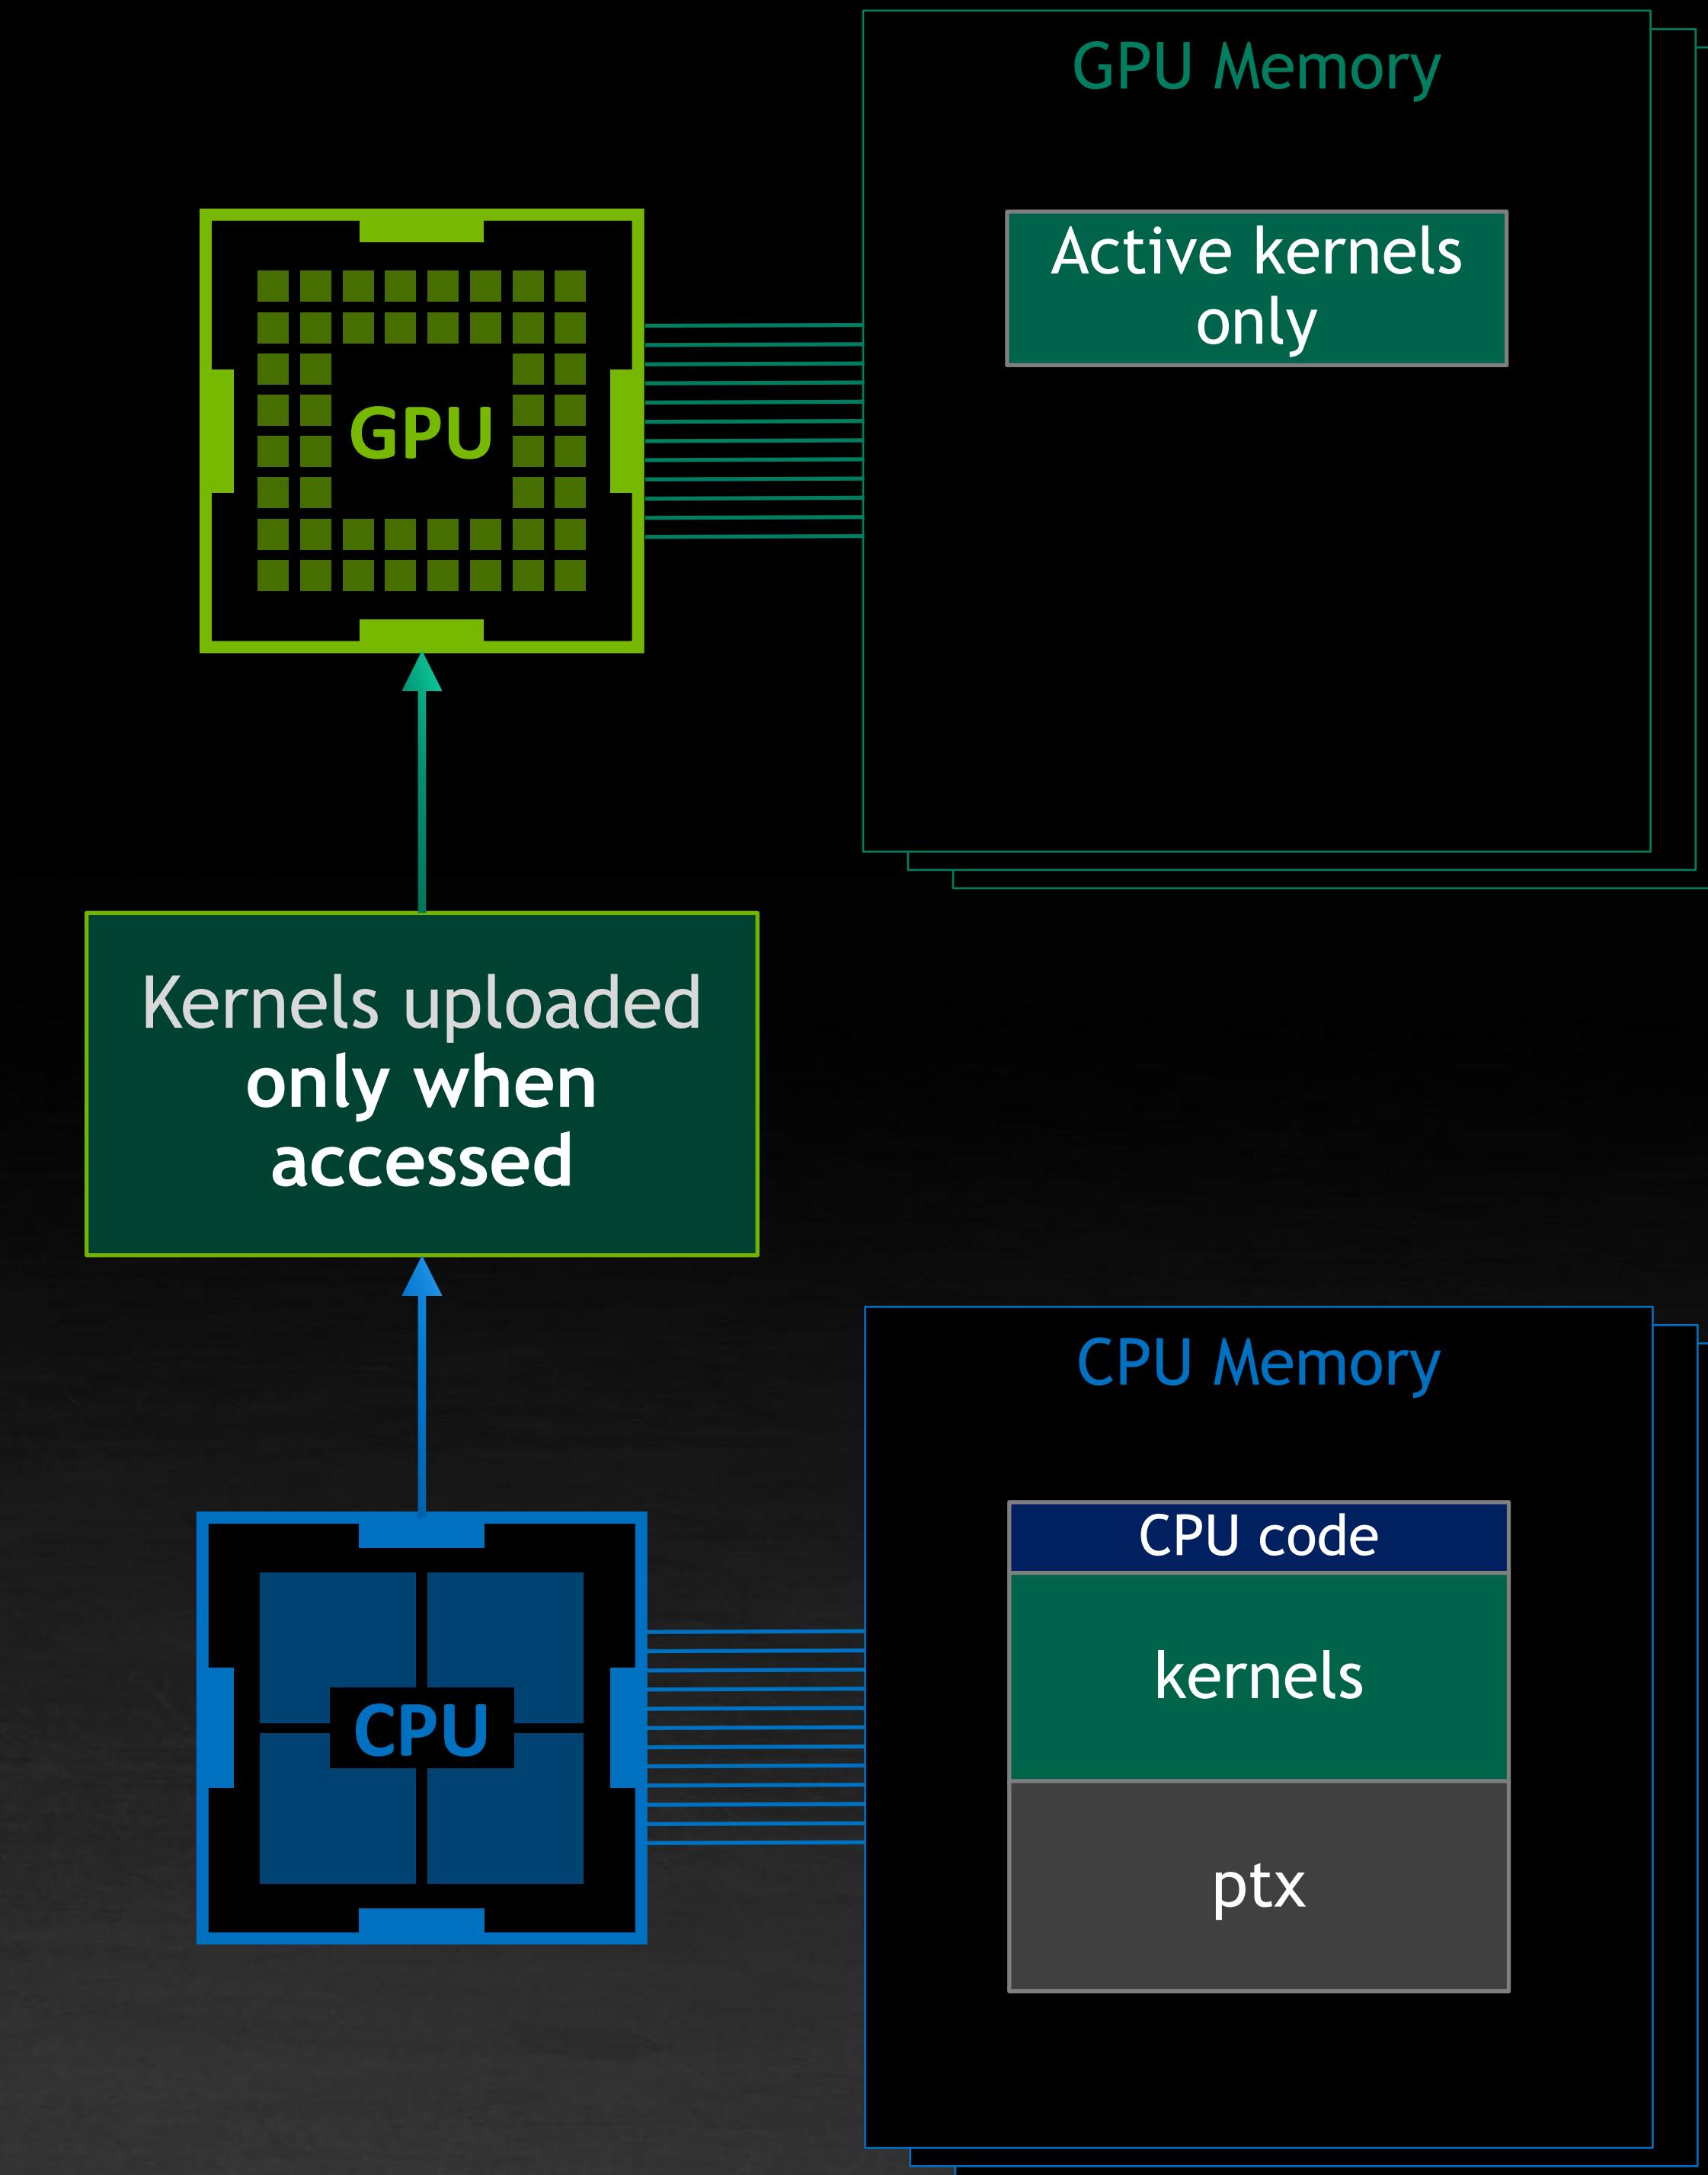
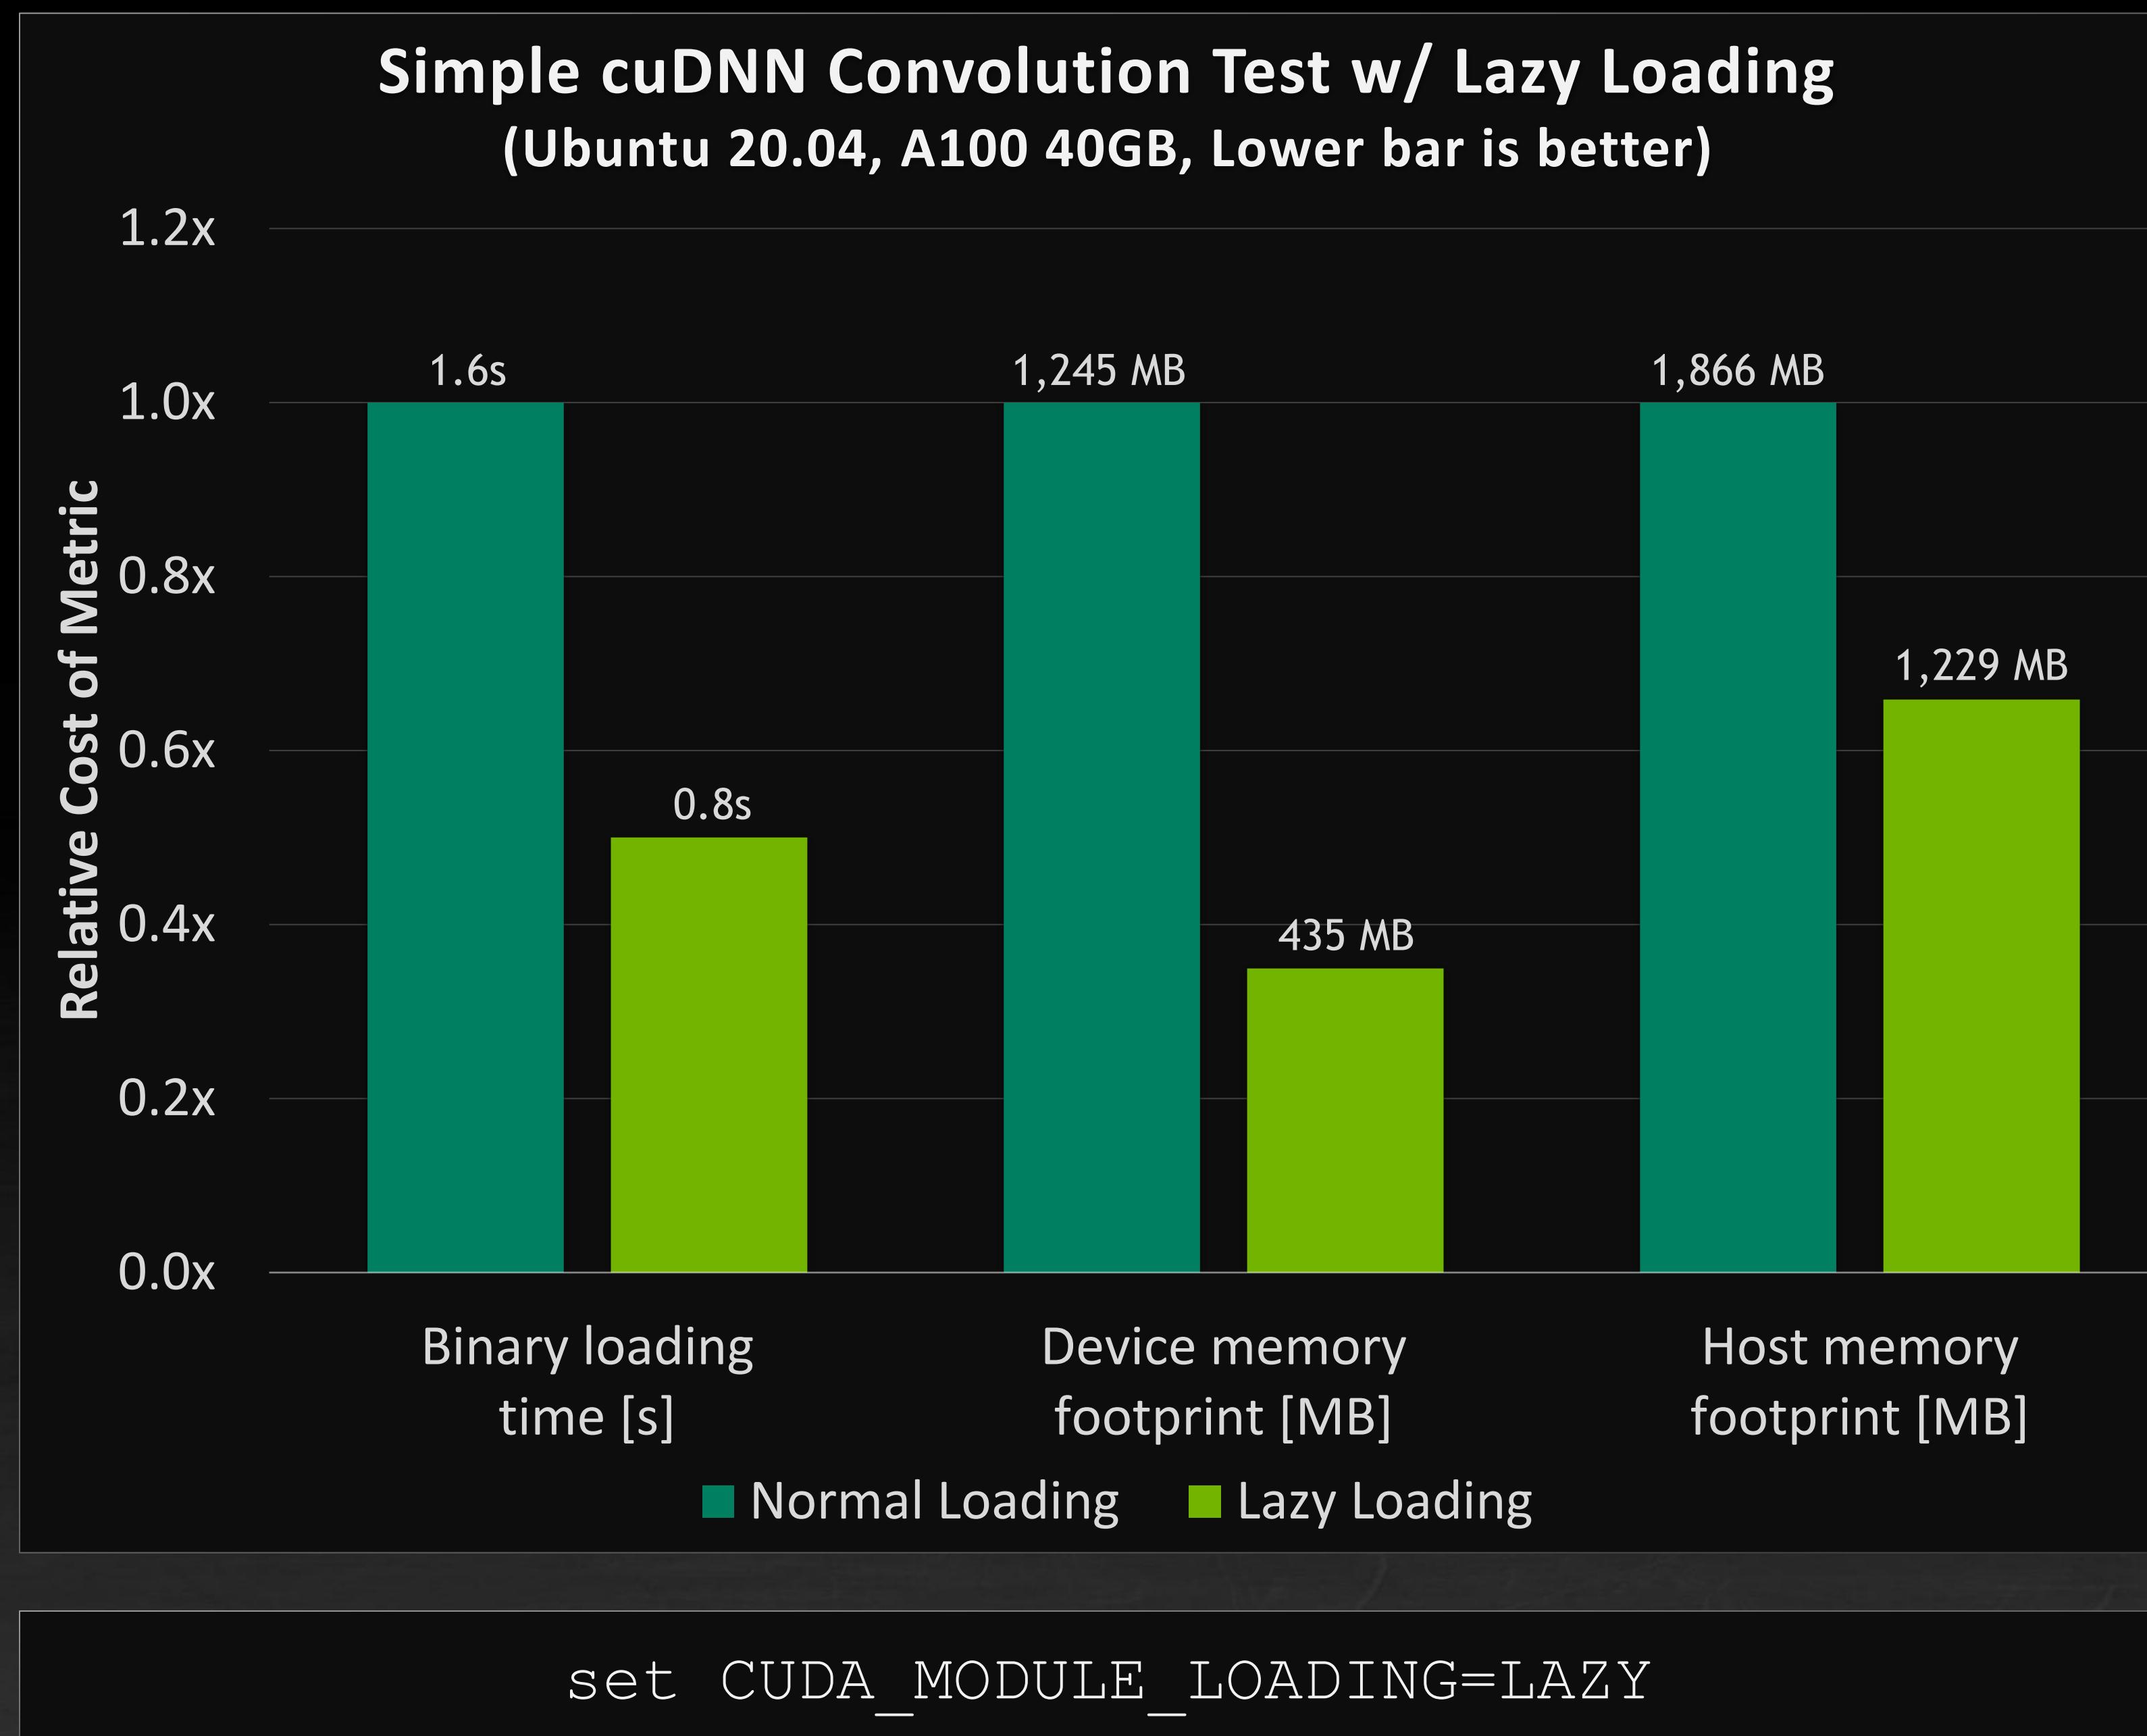


LAZY LOADING: DO NOT UPLOAD A FUNCTION UNTIL REFERENCED



ACTIVATE USING AN ENVIRONMENT VARIABLE

`CUDA_MODULE_LOADING=LAZY`

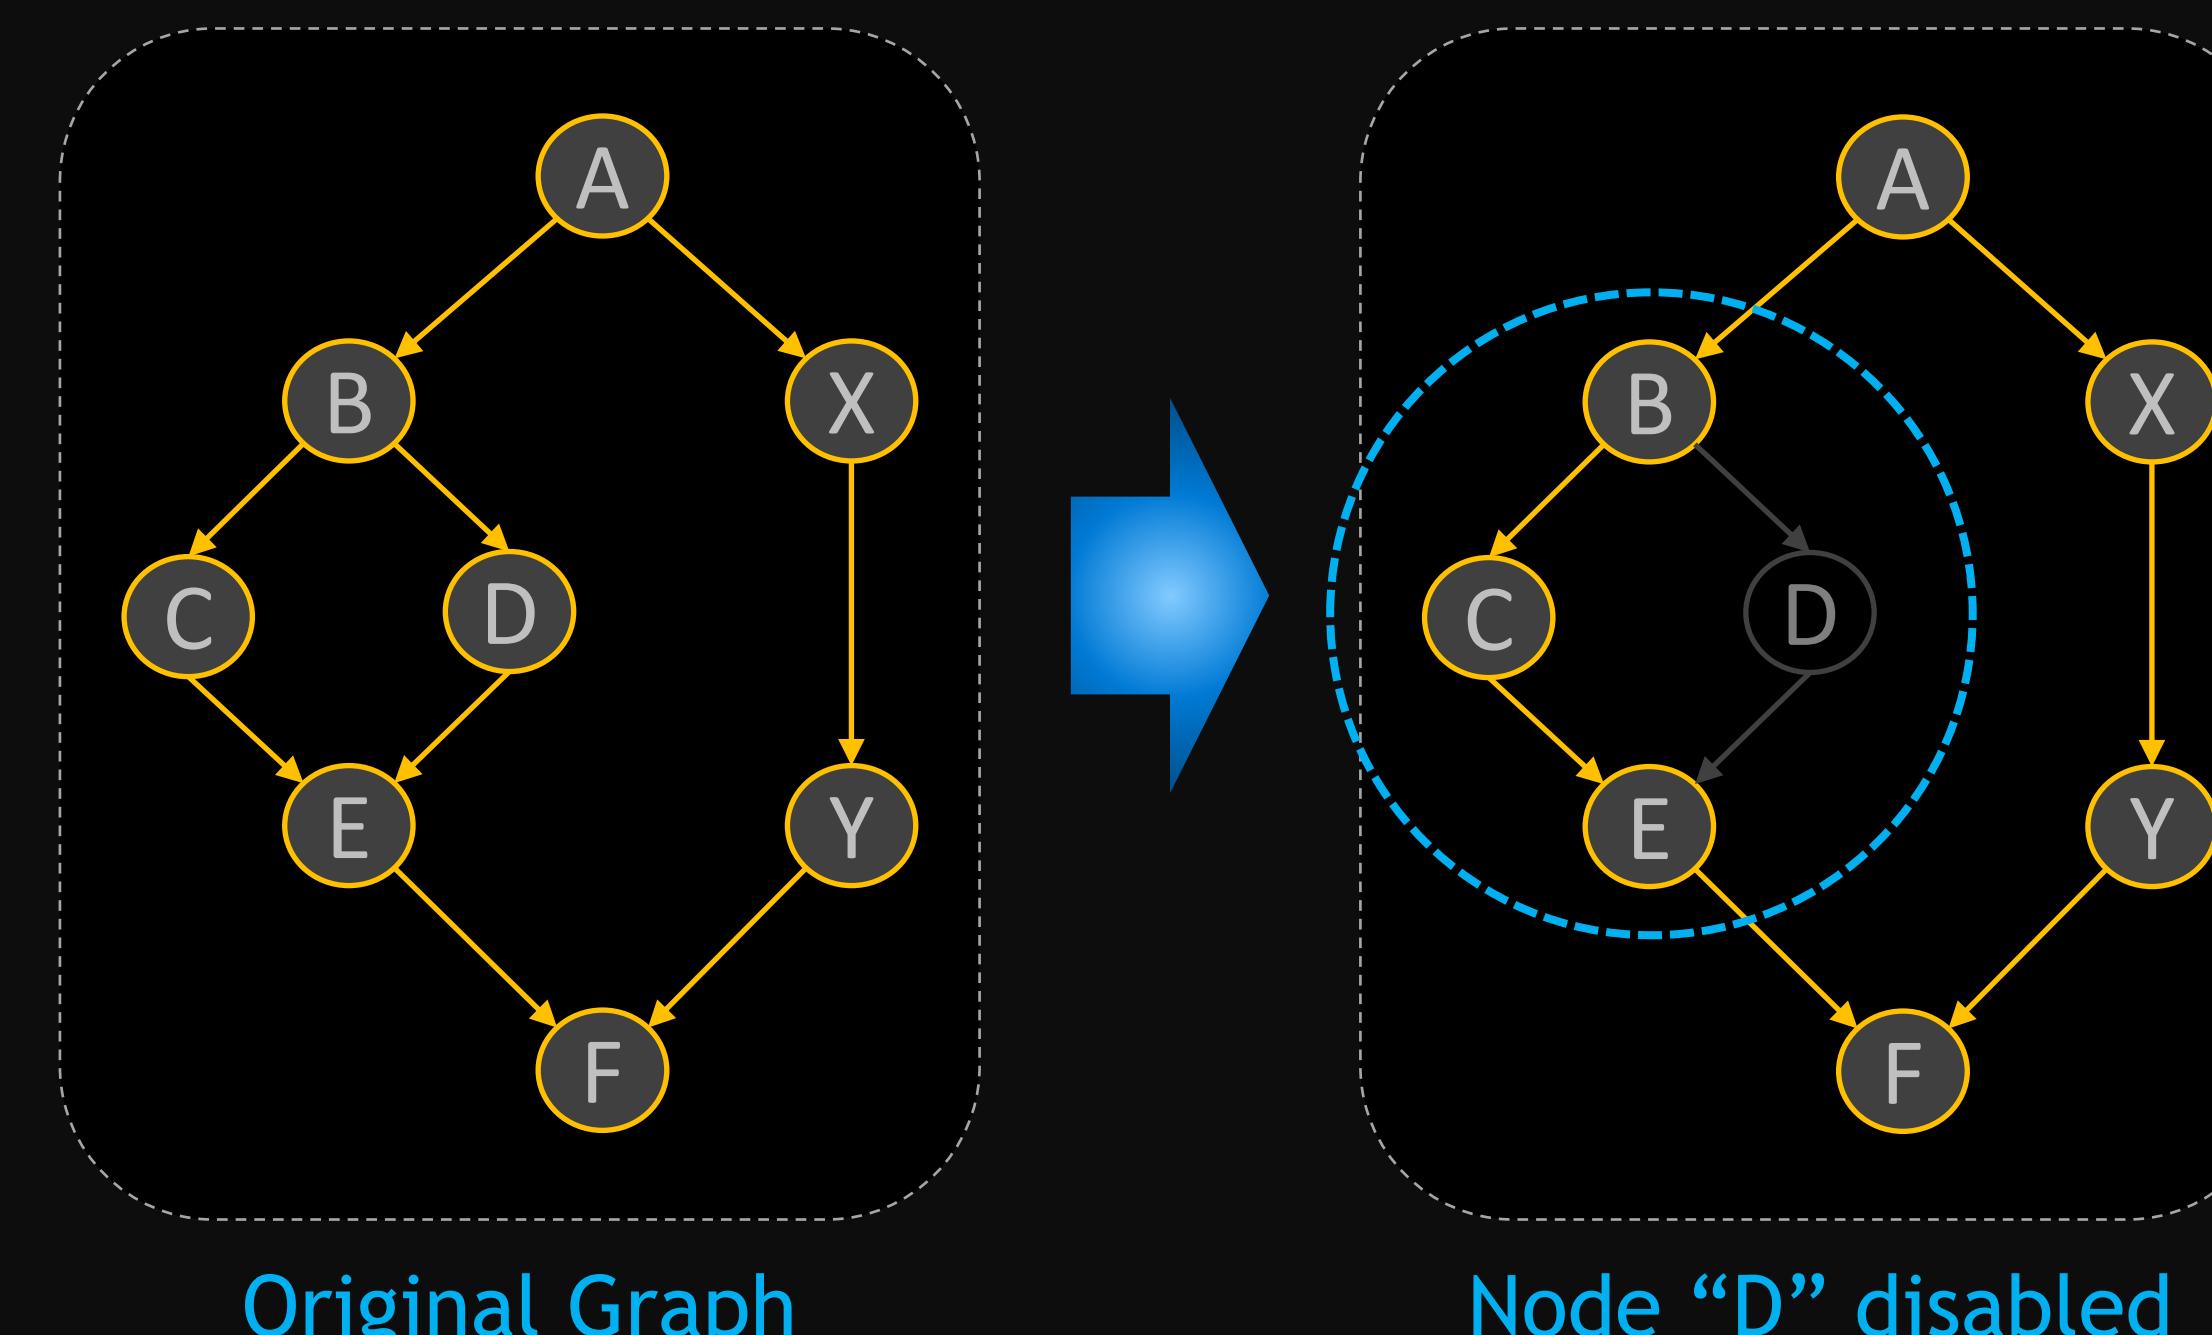


EXECUTION MANAGEMENT IN CUDA GRAPHS

Graph Node Enable/Disable

Modify execution flow prior to launch by disabling nodes in the graph

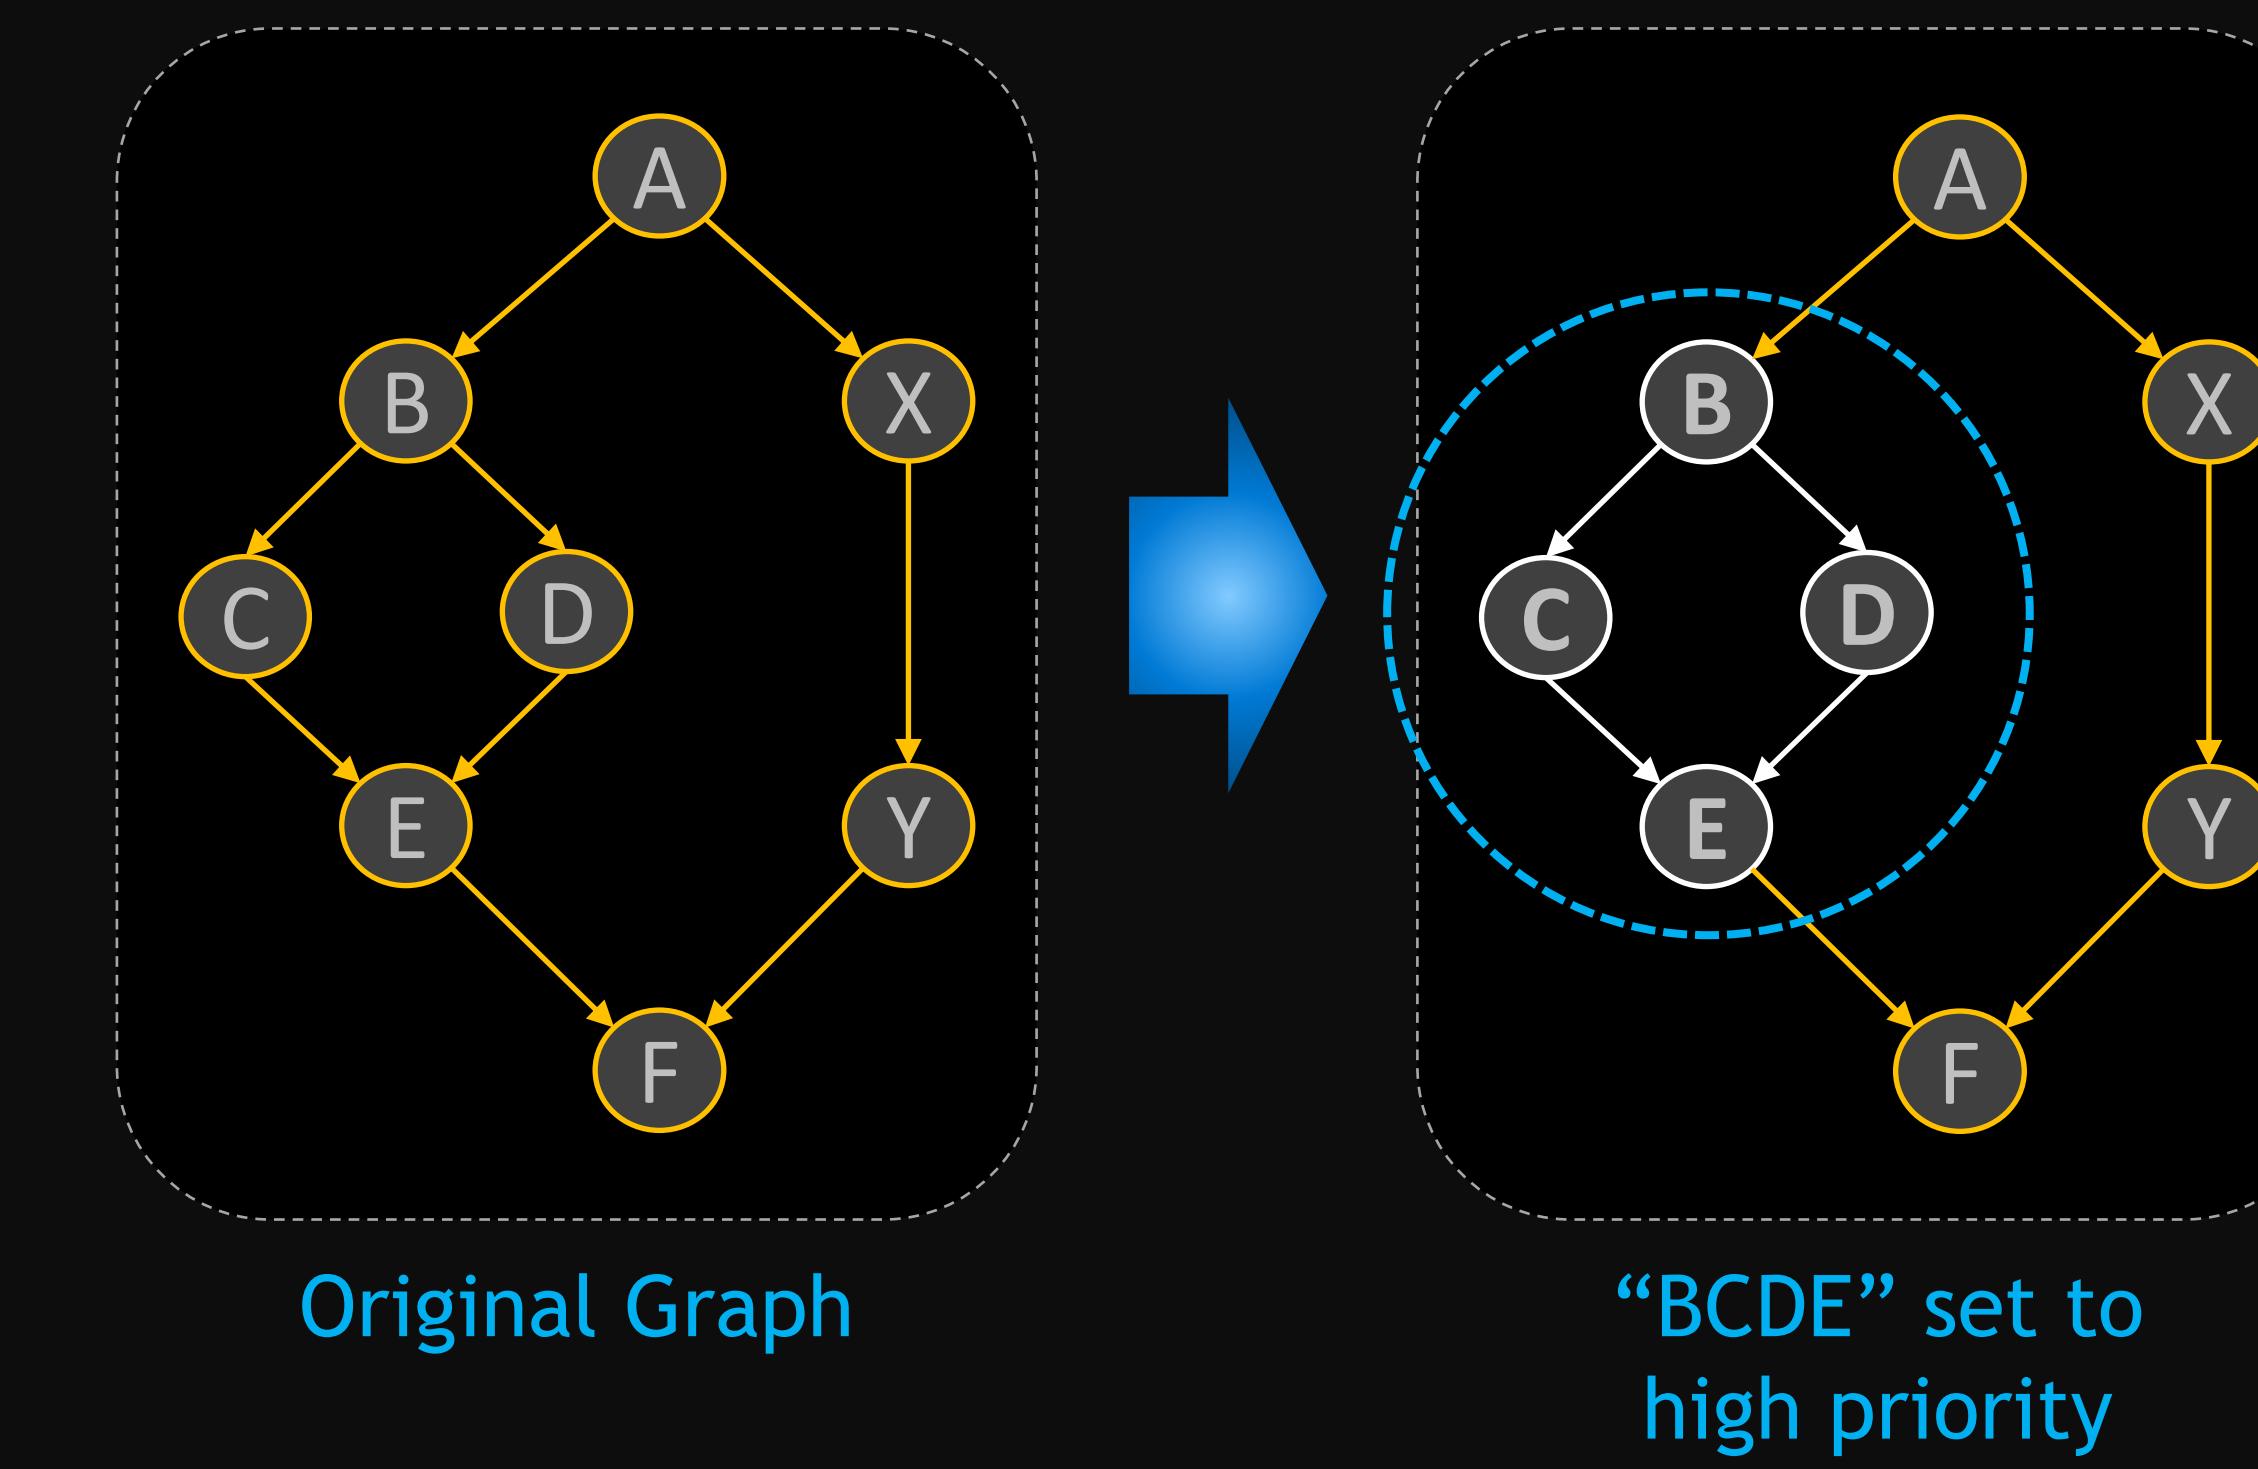
In this example, node “D” will be skipped, so left branch will execute only “BCE”



Node-Specific Priority

Create high priority branches in a task graph to manage scheduling and execution order

In this example, white edges indicate high priority so “BCDE” will run with priority over “XY”



libcu++ : THE CUDA C++ STANDARD LIBRARY

ISO C++ == Language + Standard Library

CUDA C++ == Language + **libcu++**

Strictly conforming to ISO C++, plus conforming extensions

Opt-in, Heterogeneous, Incremental

cuda::std::

Opt-in

Does not interfere with or replace your host standard library

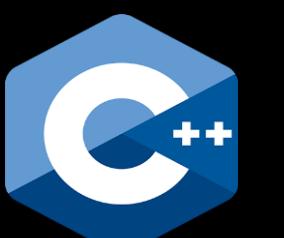
Heterogeneous

Copyable/Movable objects can migrate between host & device
Host & Device can call all member functions
Host & Device can concurrently use synchronization primitives*

Incremental

A subset of the standard library today
Each release adds more functionality

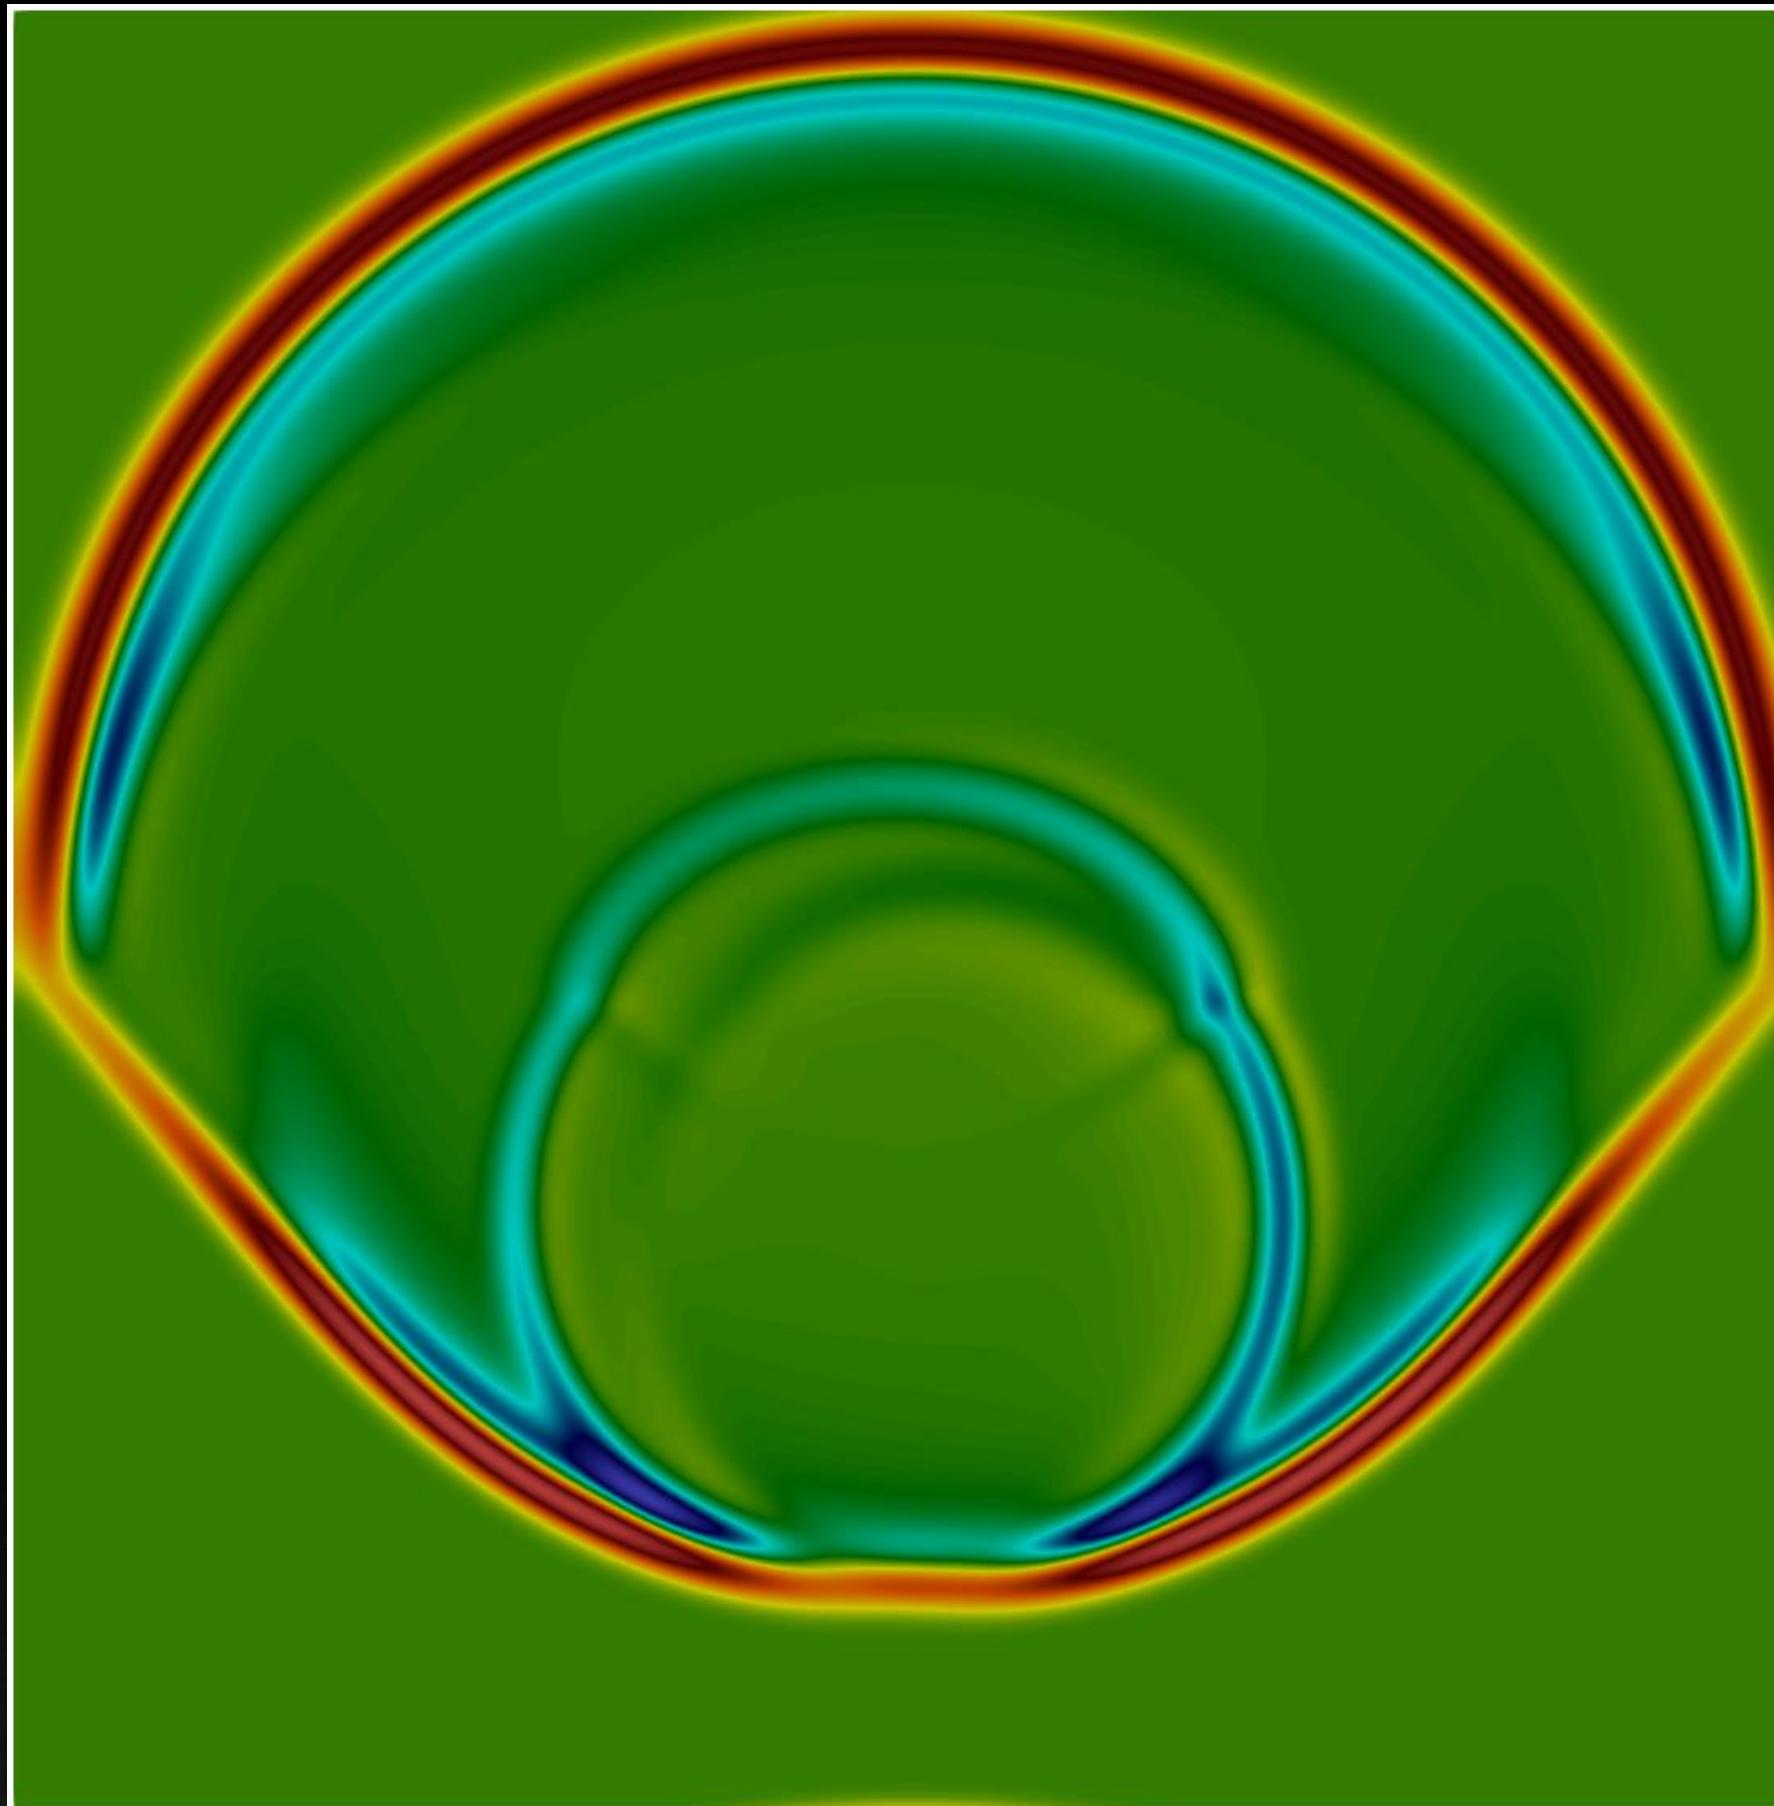
*Synchronization primitives must be in managed memory and be declared with `cuda::std::thread_scope_system`

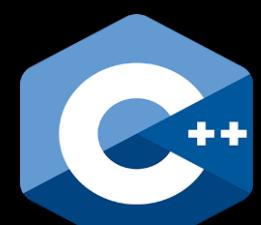


STANDARD SENDERS & RECEIVERS

Looking ahead to asynchronous task graphs in Standard C++

```
sender auto maxwell_eqs(scheduler auto &compute,
                         grid_accessor A, ...) {
    return repeat_n(n_outer_iterations,
                    repeat_n(n_inner_iterations,
                            schedule(compute)
                            | bulk(G.cells, update_h(G))
                            | halo_exchange(G, hx, hy)
                            | bulk(G.cells, update_e(time, dt, G))
                            | halo_exchange(G, hx, hy))
                            | transfer(cpu_serial_scheduler)
                            | then(output_results))
                    );
}
```

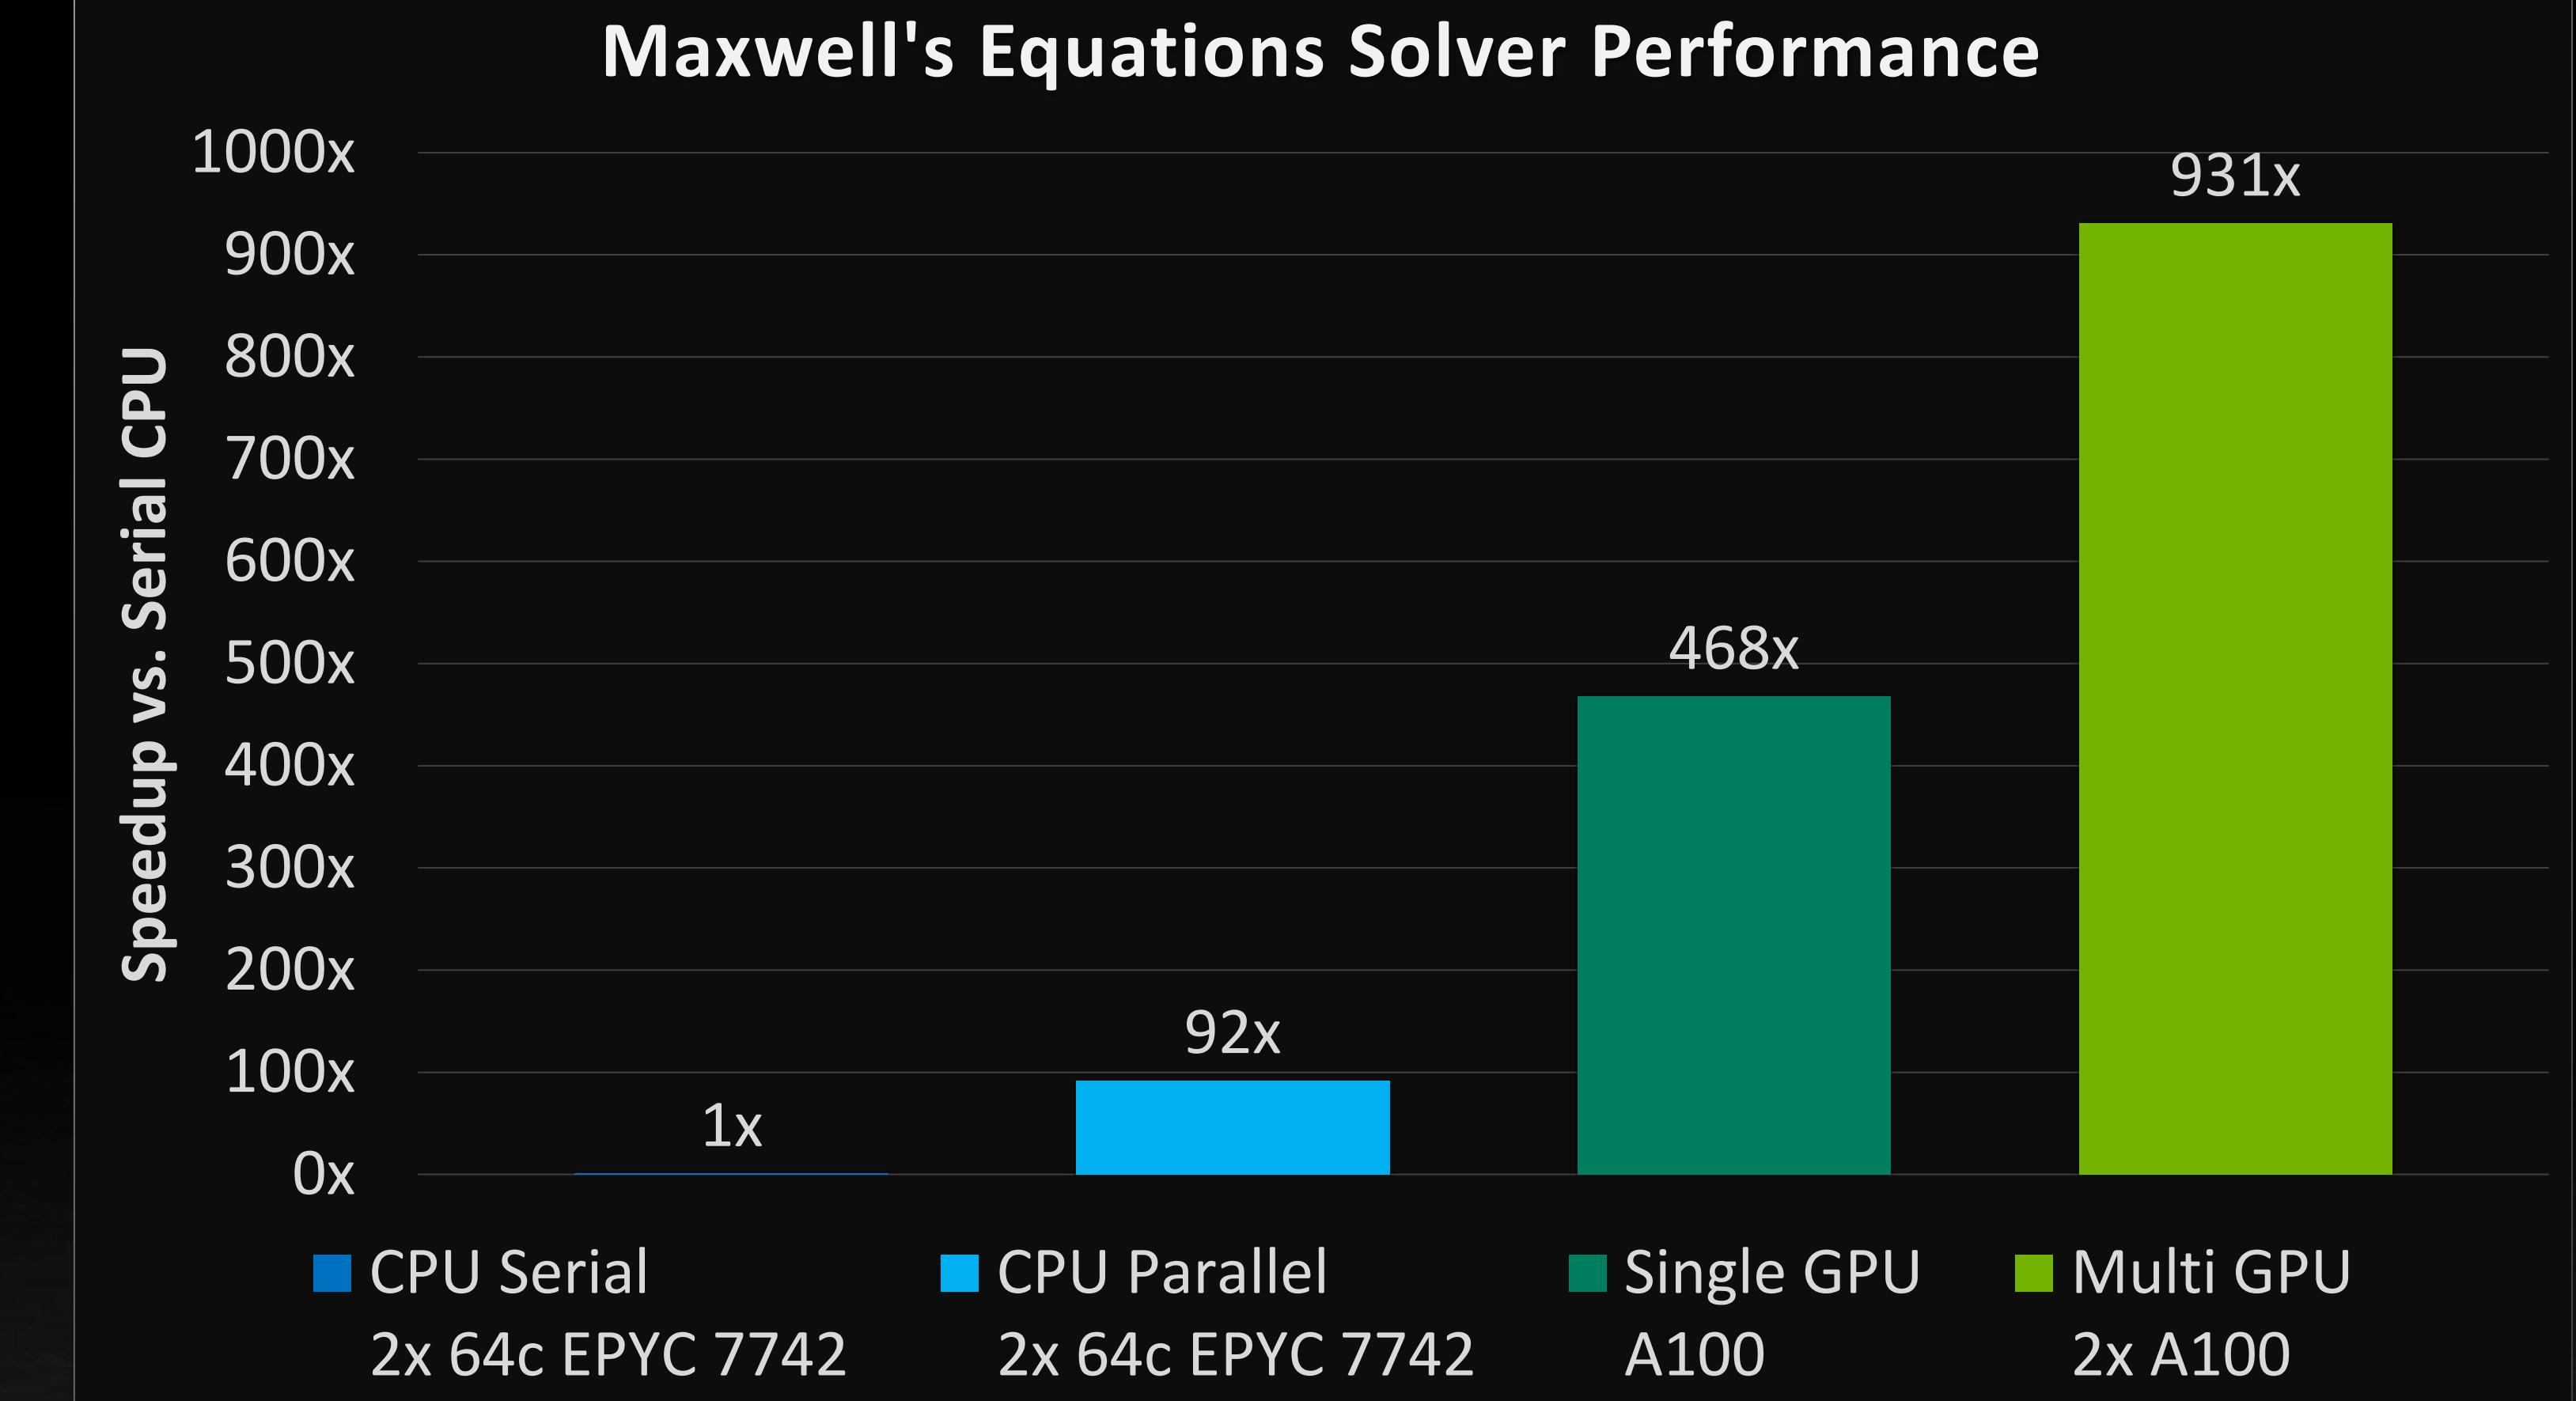




STANDARD SENDERS & RECEIVERS

Looking ahead to asynchronous task graphs in Standard C++

```
sender auto maxwell_eqs(scheduler auto &compute,  
                         grid_accessor A, ...) {  
    return repeat_n(n_outer_iterations,  
                   repeat_n(n_inner_iterations,  
                           schedule(compute)  
                           | bulk(G.cells, update_h(G))  
                           | halo_exchange(G, hx, hy)  
                           | bulk(G.cells, update_e(time, dt, G))  
                           | halo_exchange(G, hx, hy))  
                           | transfer(cpu_serial_scheduler)  
                           | then(output_results))  
};  
}
```



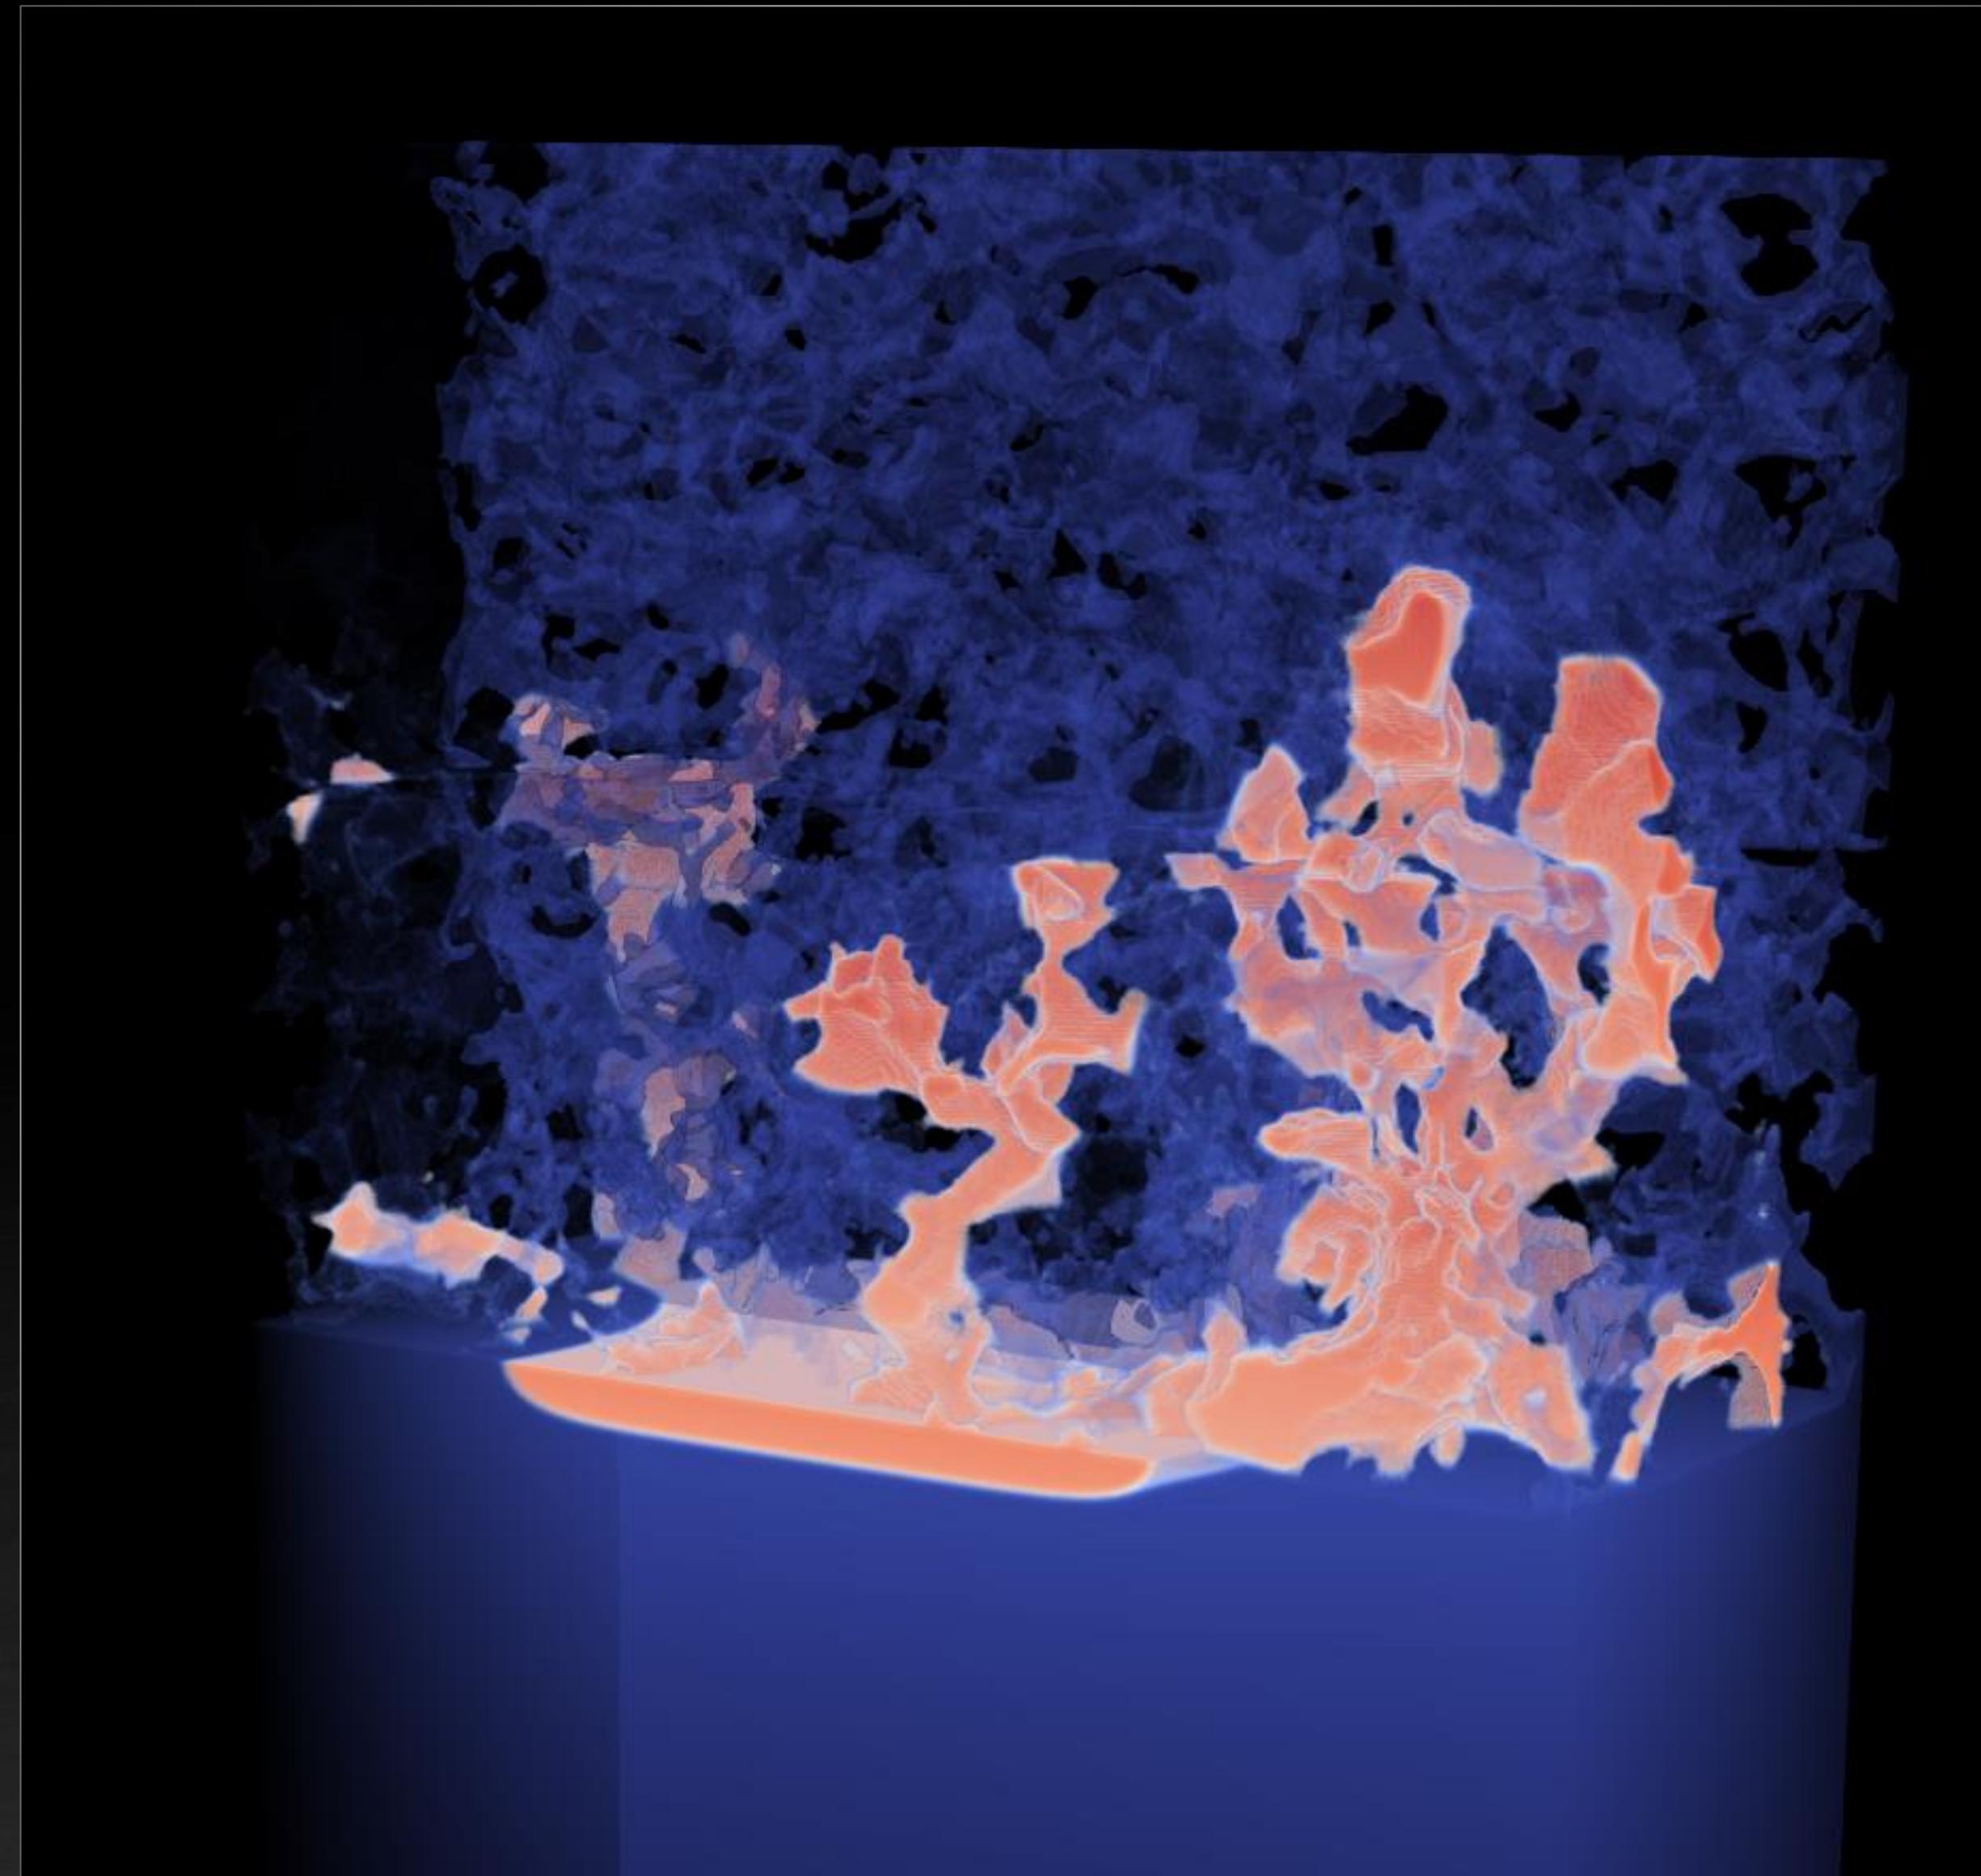
```
sync_wait(maxwell_eqs(cpu_serial_scheduler));  
sync_wait(maxwell_eqs(cpu_parallel_scheduler));  
sync_wait(maxwell_eqs(single_gpu_scheduler));  
sync_wait(maxwell_eqs(multi_gpu_scheduler));
```



Portable standard C++ runs on any target
Change a single line of code to scale
from single-core to multi-GPU

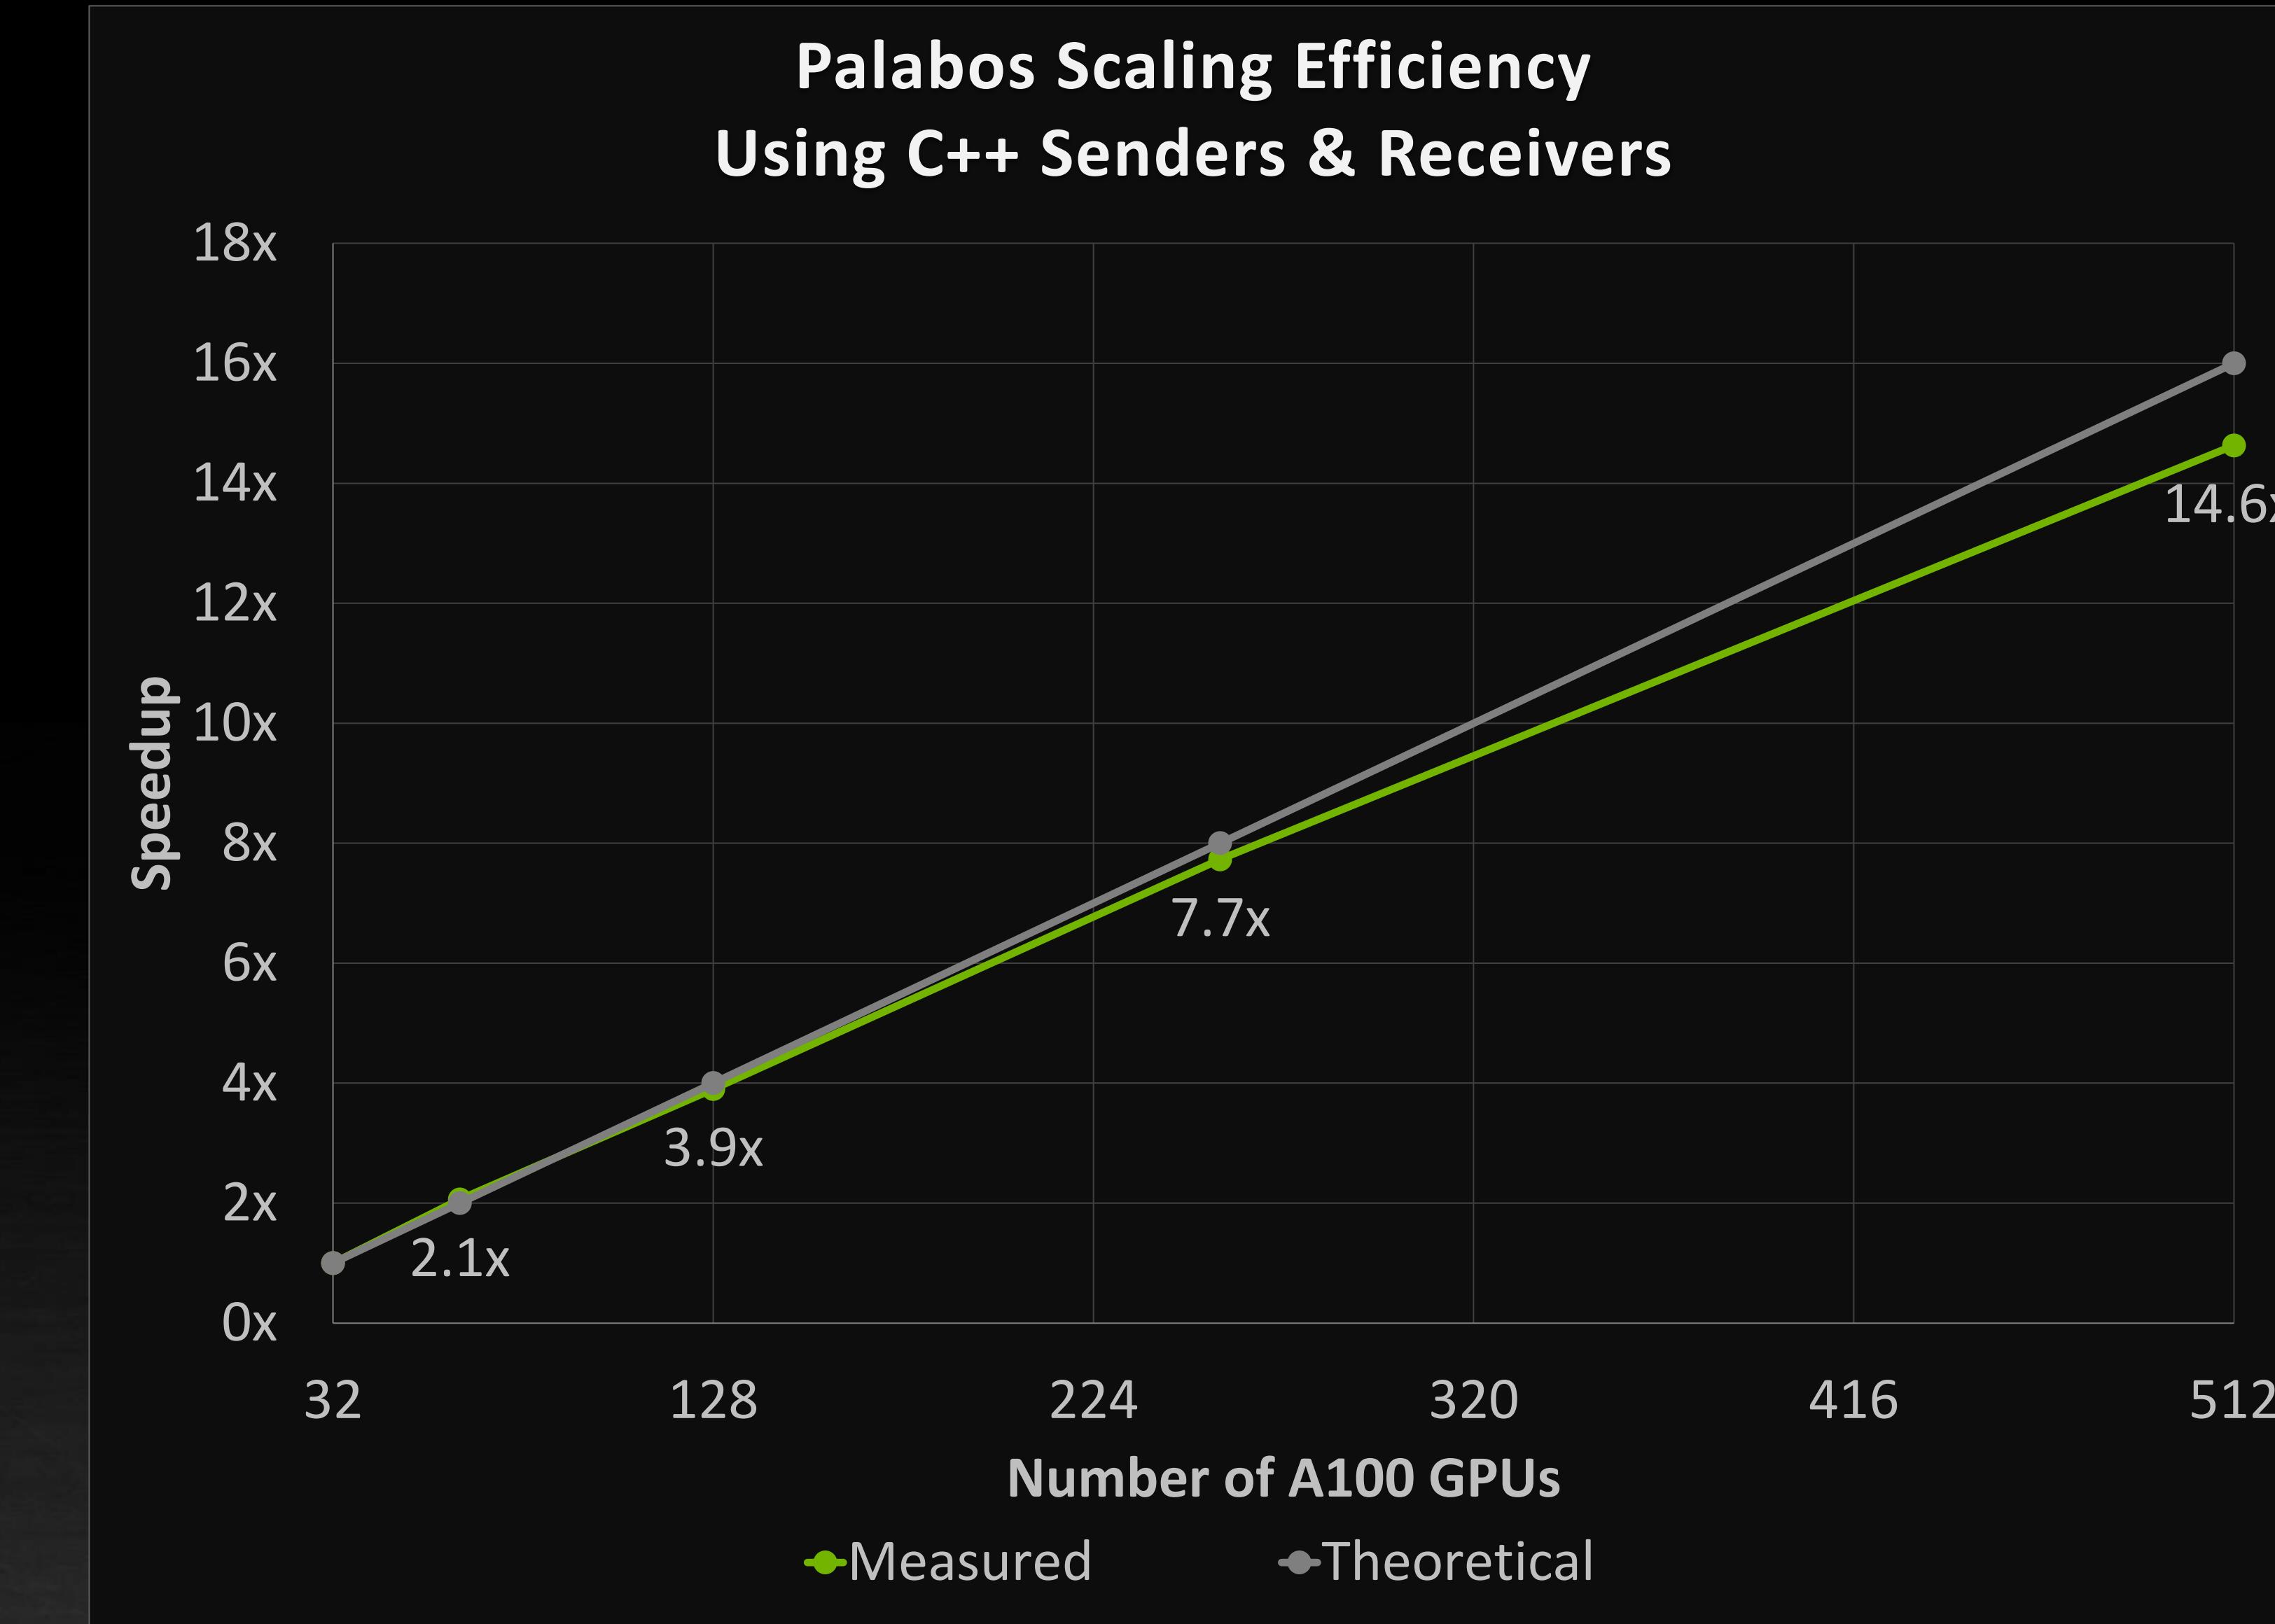
PALABOS CARBON SEQUESTRATION SIMULATION

Scale-Up Example Using Standard C++ Senders & Receivers



Buoyancy-driven diffusion of CO_2 into sandstone

Jonas Latt (University of Geneva), Christian Huber (Brown University)
Georgy Evtushenko (NVIDIA), Gonzalo Brito (NVIDIA)



Palabos is a framework for parallel computational fluid dynamics simulations using Lattice-Boltzmann methods

REFERENCES

S42663 [Inside the NVIDIA Hopper Architecture](#)

S41489 [Optimizing CUDA Applications for NVIDIA Hopper Architecture](#)

(Fall '21) A31155 [JIT LTO Adoption in cuSPARSE/cuFFT: Use Case Overview](#)

S41491 [Recent Developments in NVIDIA Math Libraries](#)

S41493 [What, Where, and Why? Use CUDA Developer Tools to Detect, Locate, and Explain Bugs and Bottlenecks](#)

S41500 [Optimizing Communication with Nsight Systems Network Profiling](#)

S41723 [How To Understand and Optimize Shared Memory Accesses using Nsight Compute](#)

S41690 [C++ Standard Parallelism](#)

(Panel) S41961 [Future of Standard and CUDA C++](#)

CWE41721 [Connect with the Experts: NVIDIA Math Libraries](#)

CWE41541 [What's in your CUDA toolbox: CUDA Profiling, Optimization, and Debugging Tools](#)

Docs [CUDA C Programming Guide: Cooperative Groups](#)

Blog [Implementing High-Precision Decimal Arithmetic with CUDA int128](#)

Blog [Reducing Application Build Times Using CUDA C++ Compilation Aids](#)

Repo <https://github.com/NVIDIA/cutlass>

Repo <https://github.com/NVIDIA/NVTX>